

23

A Long Journey of Changing std::sort Implementation at Scale

DANILA KUTENIN



Cppcon
The C++ Conference

20
23



October 01 - 06

WHO AM I?

- Senior Software Engineer at Google
- DC efficiency

AGENDA FOR TODAY

- History of sorting
- Why have we decided to change anything?
- Bugs, bugs, bugs
- Implementation
- What can you do?

REMINDERS

- Sorting is the ordering of elements
- std::sort, std::stable_sort, ranges::sort, etc

```
std::sort(begin, end);  
std::ranges::sort(cont);
```

REMINDERS

- Sorting is the ordering of elements
- std::sort, std::stable_sort, ranges::sort, etc

```
std::sort(begin, end, comp);  
std::ranges::sort(cont, comp);
```

QUICKSORT

QUICKSORT

- Quick sort

QUICKSORT

- Quick sort
- Take any element

QUICKSORT

- Quick sort
- Take any element
- Less to the left

QUICKSORT

- Quick sort
- Take any element
- Less to the left
- More to the right

QUICKSORT

- Quick sort
- Take any element
- Less to the left
- More to the right
- Recurse on both parts

- Quick sort (STL version)

```
1 template <class RandomAccessIterator>
2 void sort(RandomAccessIterator first, RandomAccessIterator l
3           __quick_sort_loop(first, last);
4           __final_insertion_sort(first, last);
5 }
```

- Quick sort (STL version)

```
1 template <class RandomAccessIterator>
2 void sort(RandomAccessIterator first, RandomAccessIterator last)
3     __quick_sort_loop(first, last);
4     __final_insertion_sort(first, last);
5 }
```

- Quick sort (STL version)

```
1 template <class RandomAccessIterator, class T, class Compare>
2 void __quick_sort_loop(RandomAccessIterator first,
3                         RandomAccessIterator last, Compare comp) {
4     while (last - first > __stl_threshold) {
5         RandomAccessIterator cut = __unguarded_partition(
6             first, last, T(__median(*first,
7                               *(first + (last - first)/2),
8                               *(last - 1),
9                               comp)),
10            comp);
11     if (cut - first >= last - cut) {
12         __quick_sort_loop(cut, last, comp);
13         last = cut;
14     } else {
15         quick_sort_loop(first, cut, comp);
```

- Quick sort (STL version)

```
1 template <class RandomAccessIterator, class T, class Compare>
2 void __quick_sort_loop(RandomAccessIterator first,
3                         RandomAccessIterator last, Compare comp) {
4     while (last - first > __stl_threshold) {
5         RandomAccessIterator cut = __unguarded_partition(
6             first, last, T(__median(*first,
7                         *(first + (last - first)/2),
8                         *(last - 1),
9                         comp)),
10            comp);
11     if (cut - first >= last - cut) {
12         __quick_sort_loop(cut, last, comp);
13         last = cut;
14     } else {
15         quick_sort_loop(first, cut - comp);
```

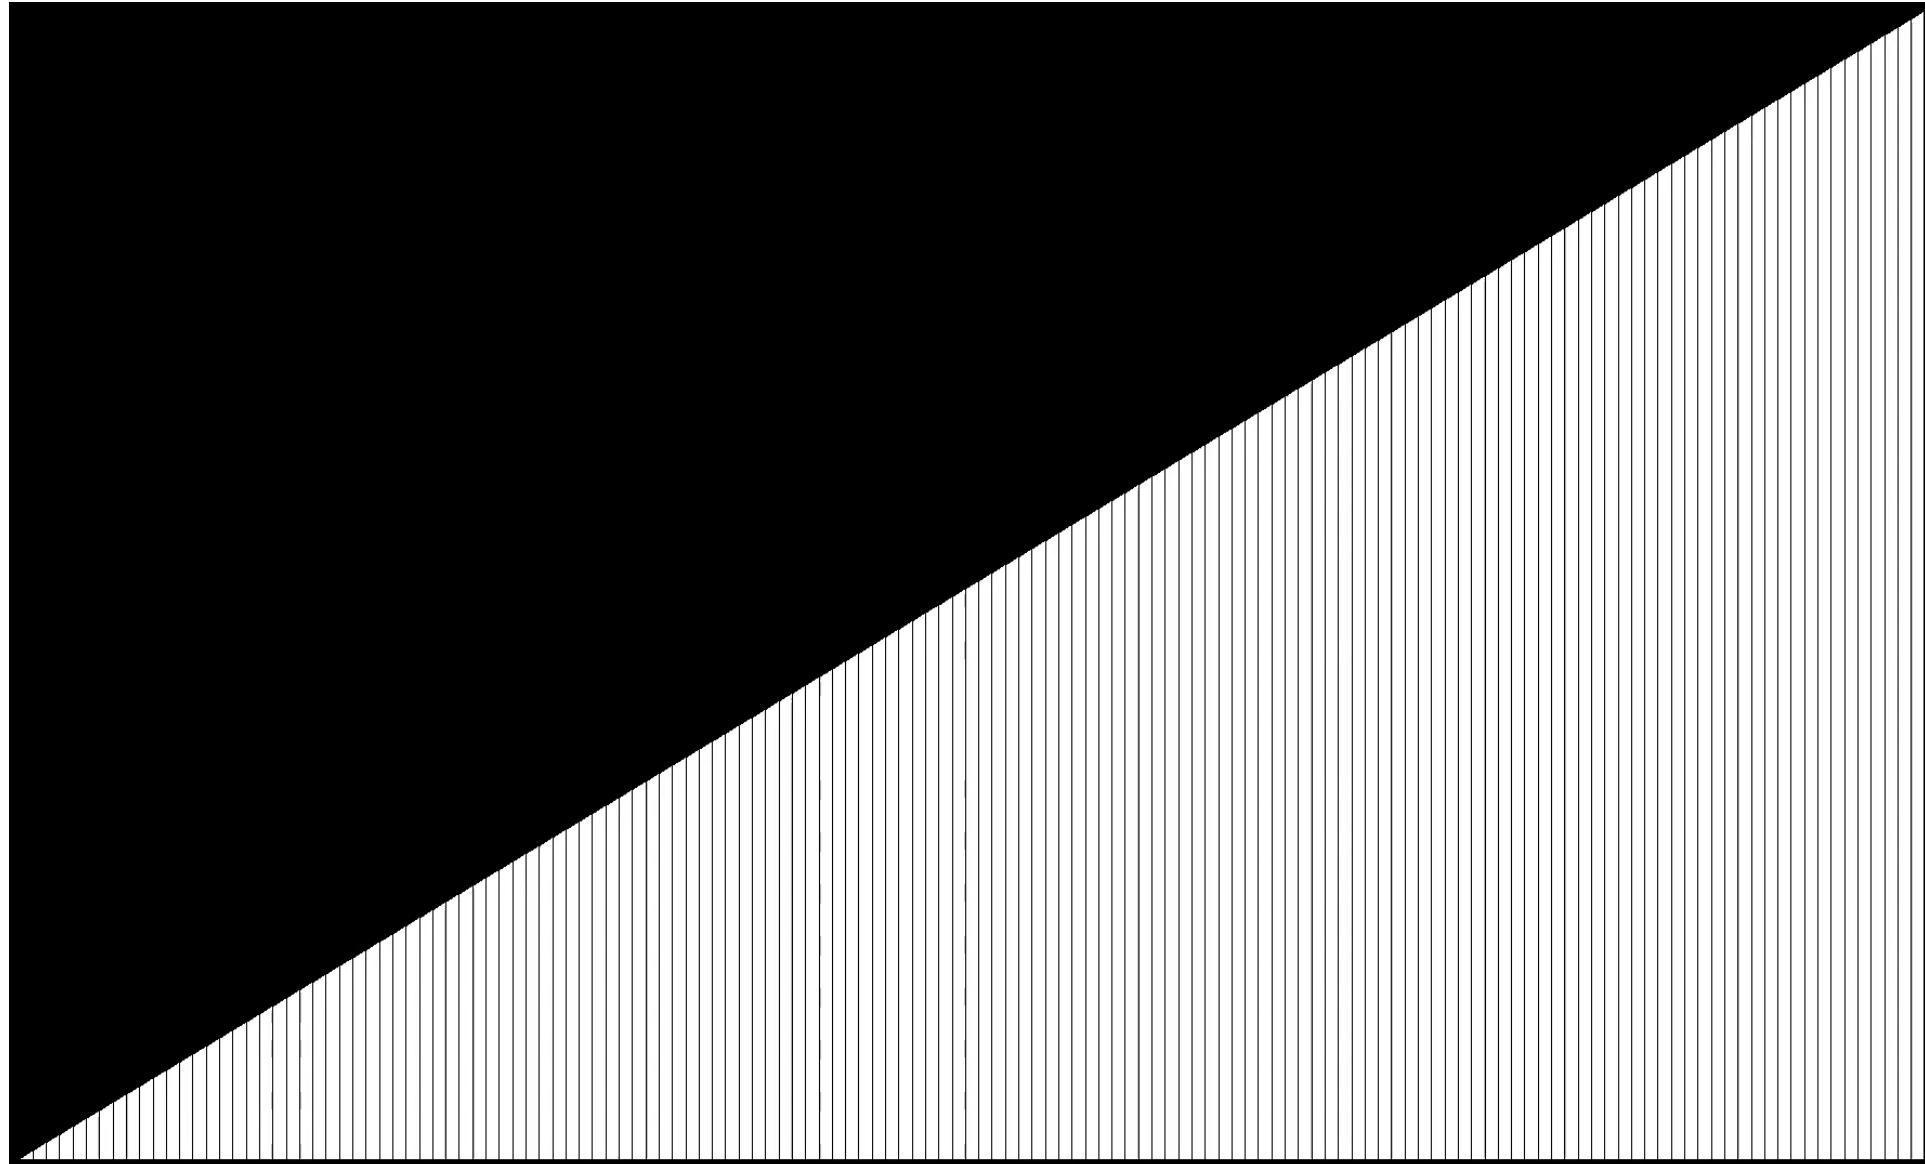
- Quick sort (STL version)

```
1 template <class RandomAccessIterator, class T, class Compare>
2 void __quick_sort_loop(RandomAccessIterator first,
3                         RandomAccessIterator last, Compare comp) {
4     while (last - first > __stl_threshold) {
5         RandomAccessIterator cut = __unguarded_partition(
6             first, last, T(__median(*first,
7                         *(first + (last - first)/2),
8                         *(last - 1),
9                         comp)),
10            comp);
11     if (cut - first >= last - cut) {
12         __quick_sort_loop(cut, last, comp);
13         last = cut;
14     } else {
15         quick_sort_loop(first, cut - comp);
```

- Quick sort (STL version)

```
1 template <class RandomAccessIterator>
2 void sort(RandomAccessIterator first, RandomAccessIterator last)
3     __quick_sort_loop(first, last);
4     __final_insertion_sort(first, last);
5 }
```

- Quick sort (STL version)



STANDARD LIBRARY EVOLVED FROM STL

STD::SORT

std::sort

Defined in header `<algorithm>`

<code>template< class RandomIt ></code>		
<code>void sort(RandomIt first, RandomIt last);</code>	(1)	(until C++20)
<code>template< class RandomIt ></code>		
<code>constexpr void sort(RandomIt first, RandomIt last);</code>		(since C++20)
<code>template< class ExecutionPolicy, class RandomIt ></code>		
<code>void sort(ExecutionPolicy&& policy,</code>	(2)	(since C++17)
<code>RandomIt first, RandomIt last);</code>		
<code>template< class RandomIt, class Compare ></code>		
<code>void sort(RandomIt first, RandomIt last, Compare comp);</code>	(3)	(until C++20)
<code>template< class RandomIt, class Compare ></code>		
<code>constexpr void sort(RandomIt first, RandomIt last, Compare comp);</code>		(since C++20)
<code>template< class ExecutionPolicy, class RandomIt, class Compare ></code>		
<code>void sort(ExecutionPolicy&& policy,</code>	(4)	(since C++17)
<code>RandomIt first, RandomIt last, Compare comp);</code>		

Complexity

$O(N \cdot \log(N))$ comparisons, where N is `std::distance(first, last)`.

STD::SORT

std::sort

Defined in header `<algorithm>`

<code>template< class RandomIt ></code>		
<code>void sort(RandomIt first, RandomIt last);</code>	(1)	(until C++20)
<code>template< class RandomIt ></code>		
<code>constexpr void sort(RandomIt first, RandomIt last);</code>		(since C++20)
<code>template< class ExecutionPolicy, class RandomIt ></code>		
<code>void sort(ExecutionPolicy&& policy,</code>	(2)	(since C++17)
<code>RandomIt first, RandomIt last);</code>		
<code>template< class RandomIt, class Compare ></code>		
<code>void sort(RandomIt first, RandomIt last, Compare comp);</code>	(3)	(until C++20)
<code>template< class RandomIt, class Compare ></code>		
<code>constexpr void sort(RandomIt first, RandomIt last, Compare comp);</code>		(since C++20)
<code>template< class ExecutionPolicy, class RandomIt, class Compare ></code>		
<code>void sort(ExecutionPolicy&& policy,</code>	(4)	(since C++17)
<code>RandomIt first, RandomIt last, Compare comp);</code>		

Proposed resolution:

In 27.8.2.1 [sort], change the complexity to "O(N log N)", and remove footnote 266:

*Complexity: Approximately $\Omega(N \log(N))$ (where $N = last - first$) comparisons on the average.*²⁶⁶⁾

²⁶⁶⁾ If the worst-case behavior is important `stable_sort()` (25.3.1.2) or `partial_sort()` (25.3.1.3) should be used.

Quicksort's worst case is
Just pick the worst element

GCC libstdc++

```
1 template<typename _RandomAccessIterator, typename _Compare>
2 inline void
3     __sort(_RandomAccessIterator __first, _RandomAccessIterator
4     {
5         if (__first != __last)
6         {
7             std::__introsort_loop(__first, __last,
8                 std::__lg(__last - __first) * 2,
9                 __comp);
10            std::__final_insertion_sort(__first, __last, __comp);
11        }
12    }
```

GCC libstdc++

```
1 template<typename _RandomAccessIterator, typename _Compare>
2 inline void
3     __sort(_RandomAccessIterator __first, _RandomAccessIterator
4     {
5     if (__first != __last)
6     {
7         std::__introsort_loop(__first, __last,
8             std::__lg(__last - __first) * 2,
9             __comp);
10        std::__final_insertion_sort(__first, __last, __comp);
11    }
12 }
```

GCC

```
1 template<typename _RandomAccessIterator, typename _Compare>
2     inline void
3         __sort(_RandomAccessIterator __first, _RandomAccessIterat
4     {
5         if (__first != __last)
6     {
7             std::__introsort_loop(__first, __last,
8                 std::__lg(__last - __first) * 2,
9                 __comp);
10            std::__final_insertion_sort(__first, __last, __comp);
11    }
12 }
```

Original

```
1 template <class RandomAccessIterator>
2 void sort(RandomAccessIterator first, RandomAccessIterator l
3     __quick_sort_loop(first, last);
4     __final_insertion_sort(first, last);
5 }
```

GCC

```
1 template<typename _RandomAccessIterator, typename _Compare>
2     inline void
3         __sort(_RandomAccessIterator __first, _RandomAccessIterat
4     {
5         if (__first != __last)
6     {
7             std::__introsort_loop(__first, __last,
8                 std::__lg(__last - __first) * 2,
9                 __comp);
10            std::__final_insertion_sort(__first, __last, __comp);
11    }
12 }
```

Original

```
1 template <class RandomAccessIterator>
2 void sort(RandomAccessIterator first, RandomAccessIterator l
3     __quick_sort_loop(first, last);
4     __final_insertion_sort(first, last);
5 }
```

INTROSORT

INTROSORT

- Almost like a quick sort

INTROSORT

- Almost like a quick sort
- If we suspect a lot of recursion calls, use heap sort

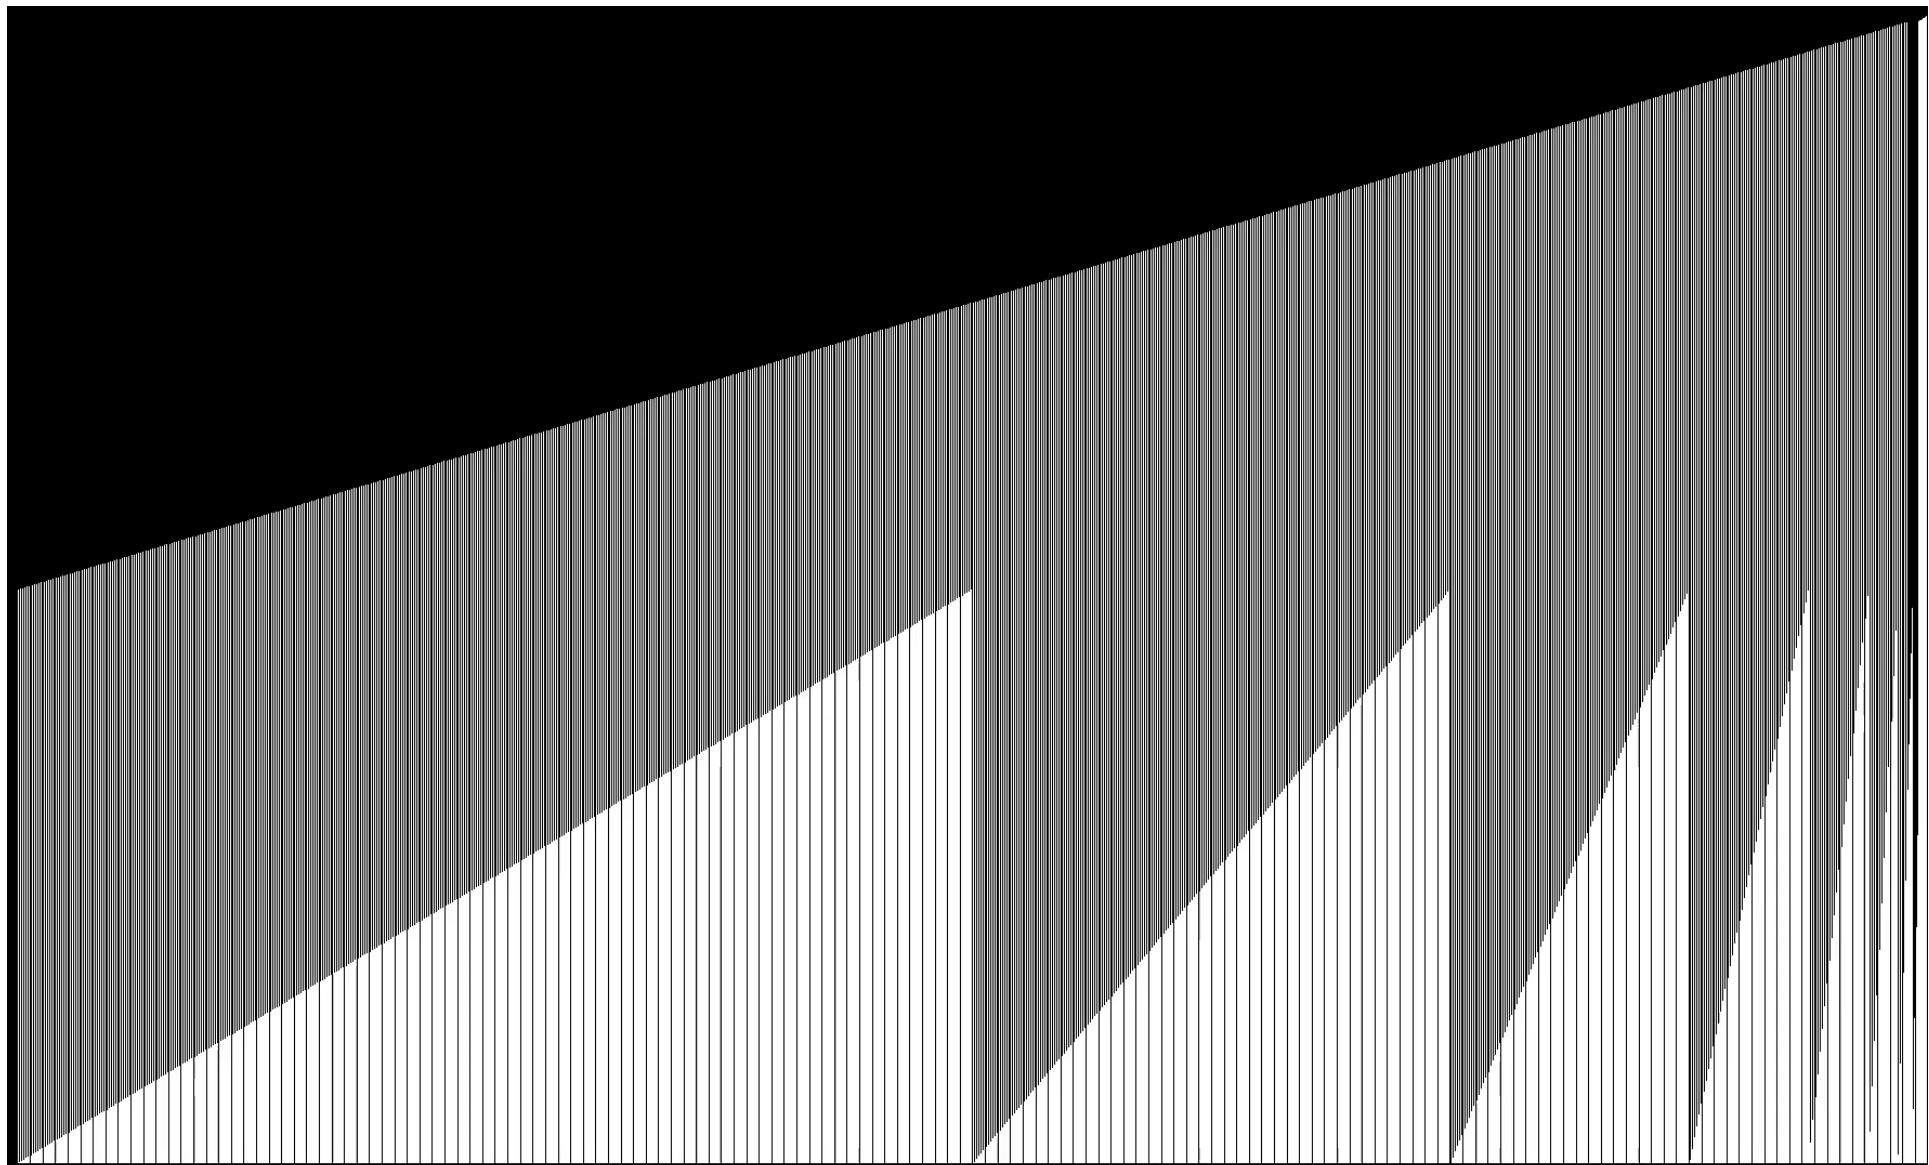
INTROSORT

- Almost like a quick sort
- If we suspect a lot of recursion calls, use heap sort
- GCC libstdc++ uses $2\log_2 n$ depth limit

INTROSORT

- Almost like a quick sort
- If we suspect a lot of recursion calls, use heap sort
- GCC libstdc++ uses $2\log_2 n$ depth limit
- Microsoft standard library does the same

INTROSORT (WORST CASE)



LLVM libc++ did not do anything at all for a long time

LLVM libc++ did not do anything at all for a long time

libc++ std::sort has $O(n^2)$ worst case, standard mandates $O(N \log(N))$
#21211

 Closed Ilvmbot opened this issue on Sep 2, 2014 · 13 comments



Ilvmbot commented on Sep 2, 2014

Member ...

Assignees

mclow

LLVM libc++ did not do anything at all for a long time

libc++ std::sort has $O(n^2)$ worst case, standard mandates $O(N \log(N))$
#21211

 Closed Ilvmbot opened this issue on Sep 2, 2014 · 13 comments



Ilvmbot commented on Sep 2, 2014

Member ...

Assignees
 mclow

 danlark1 commented on Nov 23, 2021



Member ...

Fixed with <https://reviews.llvm.org/D113413>, should land in llvm 14

7 years

Only 0.01% of all calls got into the heap sort fallback in production

LLVM (HISTORY)

What is a good sort?

What is a good sort?

- $O(n \log n)$ comparisons

What is a good sort?

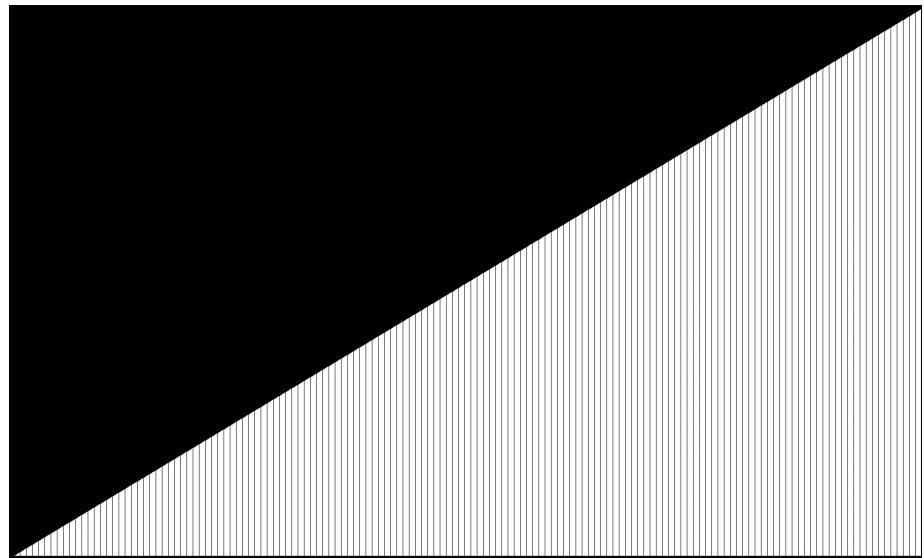
- $O(n \log n)$ comparisons
- Recognizes almost sorted patterns

What is a good sort?

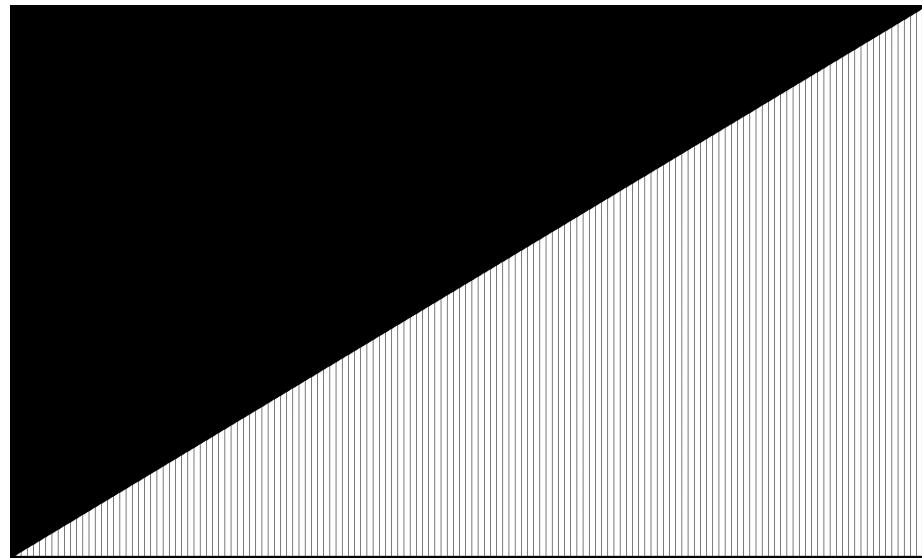
- $O(n \log n)$ comparisons
- Recognizes almost sorted patterns
- Fast for modern hardware

What is a good sort?

- $O(n \log n)$ comparisons
- Recognizes almost sorted patterns
- Fast for modern hardware
- Fewer comparisons for heavy comparison sorting



libcxx on ascending



libstdc++ on ascending

How libcxx has achieved that? Insertion sort on every step up to 8 insertions.

```
1 const unsigned __limit = 8;
2 unsigned __count = 0;
3 for (_RandomAccessIterator __i = __j+1; __i != __last; ++__i)
4     if (__comp(*__i, *__j)) {
5         value_type __t(_VSTD::move(*__i));
6         _RandomAccessIterator __k = __j;
7         __j = __i;
8         // Insert and copy
9         do {
10             *__j = _VSTD::move(*__k);
11             __j = __k;
12         } while (__j != __first && __comp(__t, *--__k));
13         *__j = _VSTD::move(__t);
14         // Exit after 8 trials
15         if (++__count == __limit)
```

How libcxx has achieved that? Insertion sort on every step up to 8 insertions.

```
3  for (_RandomAccessIterator __i = __j+1; __i != __last; ++__i)
4      if (__comp(*__i, *__j)) {
5          value_type __t(_VSTD::move(*__i));
6          _RandomAccessIterator __k = __j;
7          __j = __i;
8          // Insert and copy
9          do {
10              *__j = _VSTD::move(*__k);
11              __j = __k;
12          } while (__j != __first && __comp(__t, *--__k));
13          *__j = _VSTD::move(__t);
14          // Exit after 8 trials
15          if (++__count == __limit)
16              return ++__i == __last;
17      }
```

How libcxx has achieved that? Insertion sort on every step up to 8 insertions.

```
5     value_type __t(_VSTD::move(*__l));
6     _RandomAccessIterator __k = __j;
7     __j = __i;
8     // Insert and copy
9     do {
10         *__j = _VSTD::move(*__k);
11         __j = __k;
12     } while (__j != __first && __comp(__t, *--__k));
13     *__j = _VSTD::move(__t);
14     // Exit after 8 trials
15     if (++__count == __limit)
16         return ++__i == __last;
17     }
18     __j = __i;
19 }
```

Overall every now and then we have a better sorting

Let's change and save some cycles

Let's change and save some cycles

We tried

id_1	id_2
0	3
0	3
0	2
0	2
0	3
1	1
1	3

PostgreSQL output

1	<code>create table example (id_1 integer, id_2 integer);</code>
2	<code>-- Insert lots of equal id_1</code>
3	<code>insert into example (id_1, id_2) values (1, 1);</code>
4	<code>insert into example (id_1, id_2) values (0, 3);</code>
5	<code>insert into example (id_1, id_2) values (0, 2);</code>
6	<code>insert into example (id_1, id_2) values (0, 3);</code>
7	<code>insert into example (id_1, id_2) values (0, 2);</code>
8	<code>insert into example (id_1, id_2) values (0, 3);</code>
9	<code>insert into example (id_1, id_2) values (0, 3);</code>
10	<code>insert into example (id_1, id_2) values (1, 3);</code>
11	
12	<code>-- Order only by the first element, second</code>
13	<code>-- is undefined for equal first elements.</code>
14	<code>select * from example order by id_1;</code>

[diff.sql](#) hosted with ❤ by GitHub

[view raw](#)

id_1	id_2
0	3
0	2
0	3
0	2
0	3
1	1
1	3

MySQL output

PROBLEM: TIES

id_1	id_2
0	3
0	3
0	2
0	2
0	3
1	1
1	3

PostgreSQL output

1	<code>create table example (id_1 integer, id_2 integer);</code>
2	<code>-- Insert lots of equal id_1</code>
3	<code>insert into example (id_1, id_2) values (1, 1);</code>
4	<code>insert into example (id_1, id_2) values (0, 3);</code>
5	<code>insert into example (id_1, id_2) values (0, 2);</code>
6	<code>insert into example (id_1, id_2) values (0, 3);</code>
7	<code>insert into example (id_1, id_2) values (0, 2);</code>
8	<code>insert into example (id_1, id_2) values (0, 3);</code>
9	<code>insert into example (id_1, id_2) values (0, 3);</code>
10	<code>insert into example (id_1, id_2) values (1, 3);</code>
11	
12	<code>-- Order only by the first element, second</code>
13	<code>-- is undefined for equal first elements.</code>
14	<code>select * from example order by id_1;</code>

[diff.sql](#) hosted with ❤ by GitHub

[view raw](#)

id_1	id_2
0	3
0	2
0	3
0	2
0	3
1	1
1	3

MySQL output

PROBLEM: TIES

```
std::vector<std::pair<int, int>> first_elements_equal{
    {1, 1}, {1, 2}
};
std::sort(first_elements_equal.begin(),
          first_elements_equal.end(),
          [](const auto& lhs, const auto& rhs) {
    // Compare only by a part of sorted data.
    return lhs.first < rhs.first;
});
// Serialize or make assumptions about all data.
// Wrong, might be either 1 or 2.
assert(first_elements_equal[0].second == 1);
```

PROBLEM: TIES

It took a **year** to fix all golden tests

PROBLEM: TIES. HOW

```
1 template <class _AlgPolicy, class _RandomAccessIterator, cla
2 inline _LIBCPP_HIDE_FROM_ABI _LIBCPP_CONSTEXPR_SINCE_CXX20
3 void __sort_impl(_RandomAccessIterator __first, _RandomAcces
4     std::__debug_randomize_range<_AlgPolicy>(__first, __last);
5 // ...
```

PROBLEM: TIES. HOW

```
1 template <class _AlgPolicy, class _RandomAccessIterator, cla
2 inline _LIBCPP_HIDE_FROM_ABI _LIBCPP_CONSTEXPR_SINCE_CXX20
3 void __sort_impl(_RandomAccessIterator __first, _RandomAcces
4     std::__debug_randomize_range<_AlgPolicy>(__first, __last);
5 // ...
```

Shuffle the range before

PROBLEM: TIES. HOW

```
_LIBCPP_HIDE_FROM_ABI static uint_fast64_t __seed() {  
    static char __x;  
    return reinterpret_cast<uintptr_t>(&__x);  
}
```

PROBLEM: TIES. HOW

```
_LIBCPP_HIDE_FROM_ABI static uint_fast64_t __seed() {  
    static char __x;  
    return reinterpret_cast<uintptr_t>(&__x);  
}
```

Address space layout randomization

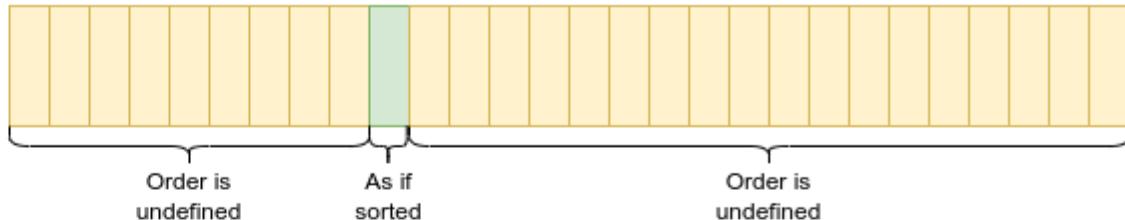
Randomization went beyond std::sort

Randomization went beyond std::sort

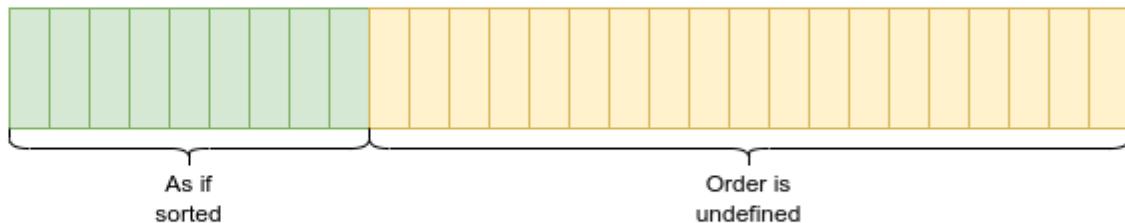
Also to std::nth_element, std::partial_sort

std::nth_element, std::partial_sort

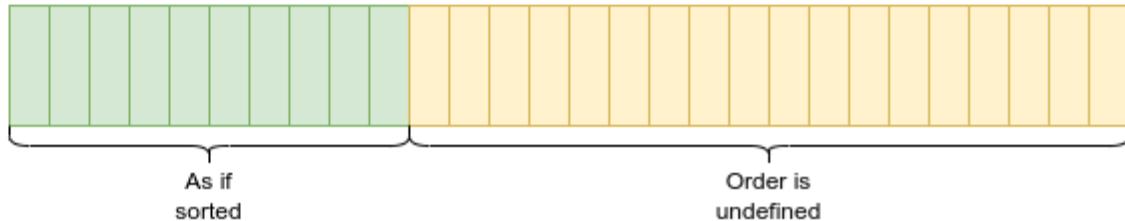
`std::nth_element(begin, begin + 9, end);`



`std::partial_sort(begin, begin + 9, end);`



`std::partial_sort(begin, begin + 10, end);`



Randomize before and randomize after the undefined ranges

std::nth_element

```
1 template <class _AlgPolicy, class _RandomAccessIterator, cl
2 inline _LIBCPP_HIDE_FROM_ABI _LIBCPP_CONSTEXPR_SINCE_CXX20
3 void __nth_elementImpl(_RandomAccessIterator __first, _Ran
4                                     _Compare& __comp) {
5     if (__nth == __last)
6         return;
7     std::__debug_randomize_range<_AlgPolicy>(__first,
8                                                 __last);
9     std::__nth_element<_AlgPolicy, __comp_ref_type<_Compare>
10    std::__debug_randomize_range<_AlgPolicy>(__first,
11                                                 __nth);
12    if (__nth != __last) {
13        std::__debug_randomize_range<_AlgPolicy>(++__nth,
14                                                 __last);
15    }
```

std::partial_sort

```
1 template <class _AlgPolicy, class _Compare, class _RandomAc
2 _LIBCPP_HIDE_FROM_ABI _LIBCPP_CONSTEXPR_SINCE_CXX20
3 _RandomAccessIterator __partial_sort(_RandomAccessIterator
4                                     _Compare& __comp) {
5     if (__first == __middle)
6         return _IterOps<_AlgPolicy>::__next(__middle, __last);
7     std::__debug_randomize_range<_AlgPolicy>(__first, __last)
8     auto __last_iter =
9         std::__partial_sort_impl<_AlgPolicy>(__first, __middle,
10        std::__debug_randomize_range<_AlgPolicy>(__middle, __last
11        return __last_iter;
12 }
```

BUGS

DETERMINISM

DETERMINISM

DETERMINISM

- Assume you sort some data

DETERMINISM

- Assume you sort some data
- Cache it

DETERMINISM

- Assume you sort some data
- Cache it
- Now ties are handled randomly

DETERMINISM

- Assume you sort some data
- Cache it
- Now ties are handled randomly
- Lower hit rate

DETERMINISM

Use std::stable_sort

```
@@ -113,7 +112,8 @@ class ArrayAUCImpl
113         sorted_labels[i].label = label;
114     }
115
116 -     ::sort(sorted_labels.begin(), sorted_labels.end(), [](const
+     auto & lhs, const auto & rhs) { return lhs.score > rhs.score; });
117
112         sorted_labels[i].label = label;
113     }
114
115 +     /// Stable sort is required for labels to apply in same
+     order if score is equal
116 +     std::stable_sort(sorted_labels.begin(), sorted_labels.end(),
+     [](const auto & lhs, const auto & rhs) { return lhs.score >
rhs.score; });
117
```

Let's write a median function

```
1 int64_t median(const std::vector<int64_t>& v) {
2     int64_t med = v.size() / 2;
3     std::nth_element(v.begin(), v.begin() + med, v.end());
4     int64_t result = v[med];
5     if (v.size() % 2 == 0) {
6         std::nth_element(v.begin(), v.begin() + med - 1, v.end());
7         result = (v[med] + v[med - 1]) / 2;
8     }
9     return result;
10 }
```

Let's write a median function

```
1 int64_t median(const std::vector<int64_t>& v) {
2     int64_t med = v.size() / 2;
3     std::nth_element(v.begin(), v.begin() + med, v.end());
4     int64_t result = v[med];
5     if (v.size() % 2 == 0) {
6         std::nth_element(v.begin(), v.begin() + med - 1, v.end());
7         result = (v[med] + v[med - 1]) / 2;
8     }
9     return result;
10 }
```

Let's write a median function

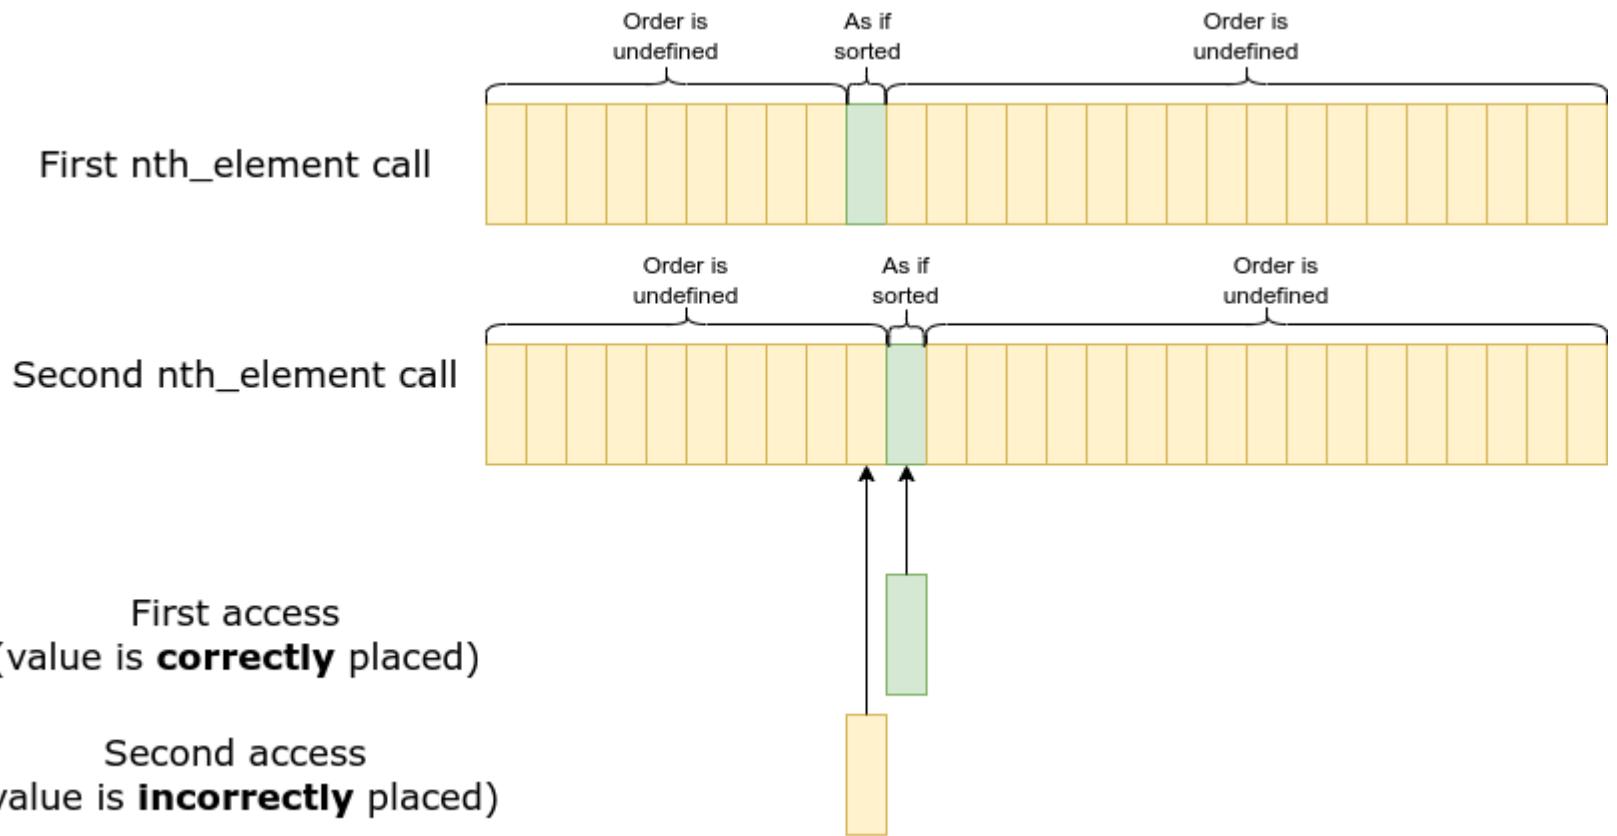
```
1 int64_t median(const std::vector<int64_t>& v) {
2     int64_t med = v.size() / 2;
3     std::nth_element(v.begin(), v.begin() + med, v.end());
4     int64_t result = v[med];
5     if (v.size() % 2 == 0) {
6         std::nth_element(v.begin(), v.begin() + med - 1, v.end());
7         result = (v[med] + v[med - 1]) / 2;
8     }
9     return result;
10 }
```

Let's write a median function

```
1 int64_t median(const std::vector<int64_t>& v) {
2     int64_t med = v.size() / 2;
3     std::nth_element(v.begin(), v.begin() + med, v.end());
4     int64_t result = v[med];
5     if (v.size() % 2 == 0) {
6         std::nth_element(v.begin(), v.begin() + med - 1, v.end());
7         result = (v[med] + v[med - 1]) / 2;
8     }
9     return result;
10 }
```

Let's write a median function

```
1 int64_t median(const std::vector<int64_t>& v) {
2     int64_t med = v.size() / 2;
3     std::nth_element(v.begin(), v.begin() + med, v.end());
4     int64_t result = v[med];
5     if (v.size() % 2 == 0) {
6         std::nth_element(v.begin(), v.begin() + med - 1, v.end());
7         result = (v[med] + v[med - 1]) / 2;
8     }
9     return result;
10 }
```



Before that you might get lucky because `nth_element` puts elements around `nth` close to total order

We found so many of them!

We found so many of them!

Repo: shogun-toolbox/shogun

Stars: 3k

Commit: 9b8d85

File: src/shogun/mathematics/linalg/LinalgNamespace.h

Snippet:

```
1819         int64_t n = a.size() / 2;
1820         std::nth_element(a_copy.begin(), a_copy.begin() + n, a_copy.end());
1821         if (a_copy.size() % 2 == 0)
1822         {
1823             std::nth_element(
1824                 a_copy.begin(), a_copy.begin() + n - 1, a_copy.end());
1825             result = (a_copy[n] + a_copy[n - 1]) / 2;
1826         }
```

We found so many of them!

Repo: [alandefreitas/matplotlibplusplus](#)

Stars: 3.5k

Commit: a40344

File: source/matplot/axes_objects/histogram.cpp

Snippet:

```
384     if (iqr_not_too_small) {  
385         size_t q1_index = static_cast<size_t>(n * 0.25);  
386         size_t q3_index = n - static_cast<size_t>(n * 0.25);  
387         auto x_copy = x;  
388         std::nth_element(x_copy.begin(), x_copy.begin() + q1_index,  
389                         x_copy.end());  
390         std::nth_element(x_copy.begin(), x_copy.begin() + q3_index,  
391                         x_copy.end());  
392         double interquartile_range = x_copy[q3_index] - x_copy[q1_index];  
393         double iq = std::max(interquartile_range, xrange / 10.);  
394         bin_width = 2 * iq * pow(n, -1. / 3.);  
395     }
```

We found so many of them!

Repo: azerothcore/azerothcore-wotlk

Stars: 4.5k

Commit: 1004da

File: src/common/Utilities/MathUtil.h

Snippet:

```
66         // Applying nth_element
67         // on n/2th index
68         std::nth_element(a.begin(),
69                         a.begin() + n / 2,
70                         a.end());
71
72         // Applying nth_element
73         // on (n-1)/2 th index
74         std::nth_element(a.begin(),
75                         a.begin() + (n - 1) / 2,
76                         a.end());
77
78         // Find the average of value at
79         // index N/2 and (N-1)/2
80         return (T)(a[(n - 1) / 2]
81                     + a[n / 2])
82                     / 2.0;
```

We found so many of them!

Repo: [jmzkChain/jmzk](#)

Stars: 1.2k

Commit: 3fbc2d

File: `libraries/chain/include/evt/chain/contracts/evt_contract_utils.hpp`

Snippet:

```
404         // find median
405         if(values.size() % 2 == 0) {
406             auto it1 = values.begin() + values.size() / 2 - 1;
407             auto it2 = values.begin() + values.size() / 2;
408
409             std::nth_element(values.begin(), it1 , values.end());
410             std::nth_element(values.begin(), it2 , values.end());
411
412             nv = ::floor((*it1 + *it2) / 2);
413 }
```

We found so many of them!

Repo: Ableton/link

Stars: 1k

Commit: adbbb5

File: include/ableton/link/Median.hpp

Snippet:

```
36     if (n % 2 == 0)
37     {
38         std::nth_element(begin, begin + n / 2, end);
39         std::nth_element(begin, begin + (n - 1) / 2, end);
40         return (*(begin + (n / 2)) + *(begin + (n - 1) / 2)) / 2.0;
41     }
```

We found so many of them!

Repo: <http://mia.sourceforge.net/>

Stars: N/A. medical image analysis

Commit: a081c3

File: mia/core/fullstats.cc

Snippet:

```
++      m_median = -1,
45 } else {
46     Vector::iterator i1 = tmp.begin() + n / 2 - 1;
47     Vector::iterator i2 = tmp.begin() + n / 2;
48     std::nth_element(tmp.begin(), i1, tmp.end());
49     std::nth_element(tmp.begin(), i2, tmp.end());
50     m_median = (*i1 + *i2) / 2.0;
51 }
```

We found so many of them!

Repo: [zenustech/zeno](#)

Stars: 471

Commit: 4fd158

File: projects/FastFLIP/FLIP_vdb.cpp

Snippet:

```
3644     int nleaf = max_v_per_leaf.size();
3645     int top90 = nleaf * 99 / 100;
3646     std::nth_element(max_v_per_leaf.begin(), max_v_per_leaf.begin() + nleaf - 1,
3647                       max_v_per_leaf.end());
3648     std::nth_element(max_v_per_leaf.begin(), max_v_per_leaf.begin() + top90,
3649                       max_v_per_leaf.end());
3650
3651     /*for (const auto& v : max_v_per_leaf) {
3652         if (v > max_v) {
3653             max_v = v;
3654         }
3655     }*/
3656     max_v = max_v_per_leaf[nleaf - 1];
```

We found so many of them!

Repo: [orfeotoolbox/OTB](#)

Stars: 287

Commit: d422bb

File: Modules/Registration/DisparityMap/include/otbDisparityMapMedianFilter.hxx

Snippet:

```
380         if ((p & 0x1) == 0)
381     {
382         const unsigned int                         medianPosition_low   = p / 2 - 1;
383         const unsigned int                         medianPosition_high = p / 2;
384         const typename std::vector<InputPixelType>::iterator medianIterator_low  = pixels.begin() + medianPosition_low;
385         const typename std::vector<InputPixelType>::iterator medianIterator_high = pixels.begin() + medianPosition_high;
386         std::nth_element(pixels.begin(), medianIterator_low, pixels.end());
387         std::nth_element(pixels.begin(), medianIterator_high, pixels.end());
388         outputIt.Set(static_cast<typename OutputImageType::PixelType>((*medianIterator_low + *medianIterator_high) / 2));
389     }
```

Use

_LIBCPP_DEBUG_RANDOMIZE_UNSPECIFIED_STABILITY

with libcxx

Strict weak ordering

- Irreflexivity: is false
- Asymmetry: $x \prec y$ and $y \prec x$ cannot be both true
- Transitivity: $x \prec y$ and $y \prec z$ imply $x \prec z$
- Transitivity of incomparability: $x \sim y$ and $y \sim z$ imply $x \sim z$, where \sim means neither $x < y$ nor $y < x$.
 $x = y$ $y = x$ $=$ $x = y$
 $x < y < x$

To check this we need to check all triples of elements,
it's
 $O(n^3)$

To check this we need to check all triples of elements,
it's $O(n^3)$

But what are the risks?

On paper, UB. But in practice?

```
new... ▾ Vim CppInsights Quick-bench C++  
  
hs.name < rhs.name;  
  
int main(int argc, char** argv) {  
    vector<PseudoNetwork> example{  
        {prefix_length = 6, name = "a", prefix = "h"},  
        {prefix_length = 1, name = "a", prefix = "a"},  
        {prefix_length = 1, name = "a", prefix = "a"},  
        {prefix_length = 1, name = "a", prefix = "b"},  
        {prefix_length = 1, name = "x", prefix = "a"},  
        {prefix_length = 1, name = "a", prefix = "a"},  
        {prefix_length = 1, name = "a", prefix = "a"},  
    };  
  
    auto t = example.begin(), e = example.end();  
    CompareNetworks());  
    if (sorted  
        ::is_sorted(example.begin(), example.end(),  
                    CompareNetworks)) {  
        cout << "Not sorted :(\n";  
    }  
    cout << "Sorted :)\n";  
}
```

```
x86-64 clang 14.0.0 (C++, Editor #1, Compiler #2) ✓ ×
x86-64 clang 14.0.0 ▾ ✓ -O3 -std=c++11
A ▾ ⚙ Output... ▾ Filter... ▾ Libraries (1) + Add
std::__1::basic_string<char, std::memmove@plt:
operator new(unsigned long)@plt:
operator delete(void*)@plt:
std::__1::ios_base::clear(unsigned long)@plt:
bcmovsd@plt:
std::__1::basic_ostream<char, std::endl@plt:
std::__1::ios_base::getloc() const@plt:
std::__1::basic_ostream<char, std::endl@plt:
std::__1::locale::use_facet(std::locale@plt):
memcmpeq@plt:
std::__1::ios_base::setvbuf@plt:
C ━ Output (0/0) x86-64 clang 14.0.0 i - 3969ms (180560B) ~3
Output of x86-64 clang 14.0.0 (Compiler #2) ✓ ×
A ▾ □ Wrap lines
ASM generation compiler re...
Execution build compiler re...
Program returned: 0
Not sorted :(

```

<https://gcc.godbolt.org/z/c71qzM97f>

Look after yourself, and, if you can, someone else too ✊

x86-64 clang 14.0.0 (Editor #1) □ X

x86-64 clang 14.0.0 □ X

-DNUM_ELEMS=30 -O3 -std=c++17 -stdlib=lib

A Output... Filter... Libraries (1) Overrides + Add new... Add tool...

```
operator new(unsigned long)@plt-0x10:  
ff 35 e2 3f 00 00  
401020 push 0x3fe2(%rip)      # 405008 <_GLOBAL_OFFSET_TABLE_>  
ff 25 e4 3f 00 00  
401026 jmp *0x3fe4(%rip)      # 405010 <_GLOBAL_OFFSET_TABLE_>  
0f 1f 40 00  
40102c nopl 0x0(%rax)  
main:  
53  
401140 push %rbx  
bf 78 00 00 00  
401141 mov $0x78,%edi  
e8 e5 fe ff ff  
401146 call 401030 <operator new(unsigned long)@plt>  
48 89 c3  
40114b mov %rax,%rbx  
66 0f 76 c0  
40114c nopl 0x0(%rax)  
40114d nopl 0x0(%rax)
```

C Output (0/0) x86-64 clang 14.0.0 i -cached (74520B) Compiler License

x86-64 clang 14.0.0 (Editor #1) □ X

x86-64 clang 14.0.0 □ X

A Output... Filter... Libraries (1) Overrides + Add new... Add tool...

```
operator new(unsigned long)@plt-0x10:  
ff 35 e2 3f 00 00  
401020 push 0x3fe2(%rip)      # 405008 <_GLOBAL_OFFSET_TABLE_>  
ff 25 e4 3f 00 00  
401026 jmp *0x3fe4(%rip)      # 405010 <_GLOBAL_OFFSET_TABLE_>  
0f 1f 40 00  
40102c nopl 0x0(%rax)  
main:  
53  
401140 push %rbx  
bf 7c 00 00 00  
401141 mov $0x7c,%edi  
e8 e5 fe ff ff  
401146 call 401030 <operator new(unsigned long)@plt>  
48 89 c3  
40114b mov %rax,%rbx  
66 0f 76 c0  
40114c nopl 0x0(%rax)  
40114d nopl 0x0(%rax)
```

C Output (0/0) x86-64 clang 14.0.0 (Compiler #2) □ X

A Wrap lines Select all

```
ASM generation compiler returned: 0  
Execution build compiler returned: 0  
Program returned: 0
```

Output of x86-64 clang 14.0.0 (Compiler #2) □ X

A Wrap lines Select all

```
ASM generation compiler returned: 0  
Execution build compiler returned: 0  
Program returned: 139  
Program terminated with
```

<https://gcc.godbolt.org/z/qj7EeMohP>

RISKS

- Incorrect comparator for elements, all good
- Incorrect comparator for elements, can fail with
OOB

RISKS

- Incorrect comparator for ≤ 30 elements, all good
- Incorrect comparator for > 30 elements, **can fail with OOB**
- Have you written a well thought test for 30+ elements?

**At Google we consistently (once in months) have
outage somewhere because of that**

OOB read

```
// The search going up is known to be guarded but the search c
// Prime the downward search with a guard.
// __m still guards upward moving __i
while (__comp(*__i, *__m))
    ++__i;
```

To check this we need to check all triples of elements,
it's , sorting is .
 $O(n^3)$ $O(n \log n)$

There is a $O(n^2)$ algorithm to check for strict weak ordering

There is a $O(n^2)$ algorithm to check for strict weak ordering

- Sort the range, if not `std::is_sorted`, return false

There is a $O(n^2)$ algorithm to check for strict weak ordering

- Sort the range, if not `std::is_sorted`, return false
- Find prefix of all equivalent elements

There is a $O(n^2)$ algorithm to check for strict weak ordering

- Sort the range, if not `std::is_sorted`, return false
- Find prefix of all equivalent elements
- Check they are equal

There is a $O(n^2)$ algorithm to check for strict weak ordering

- Sort the range, if not `std::is_sorted`, return false
- Find prefix of all equivalent elements
- Check they are equal
- Check they are less than all next

There is a $O(n^2)$ algorithm to check for strict weak ordering

- Sort the range, if not `std::is_sorted`, return false
- Find prefix of all equivalent elements
- Check they are equal
- Check they are less than all next
- Remove them

There is a $O(n^2)$ algorithm to check for strict weak ordering

- Sort the range, if not `std::is_sorted`, return false
- Find prefix of all equivalent elements
- Check they are equal
- Check they are less than all next
- Remove them
- Proof can be read at

https://github.com/danlark1/quadratic_strict_weak_ordering

is still a lot but we can use it for, say, 100 first elements
 $O(n^2)$

Most sorts are small!

CAVEATS

CAVEATS

- Sort should not fail

CAVEATS

- Sort should not fail
- It only tells if strict weak ordering is met

CAVEATS

- Sort should not fail
- It only tells if strict weak ordering is met
- Arbitrary comparator can do anything

CAVEATS

- Sort should not fail
- It only tells if strict weak ordering is met
- Arbitrary comparator can do anything
- We can check we don't read OOB in std::sort itself

[RFC] Strict weak ordering checks in the debug libc++

■ Runtimes ■ C++



danlark1

1 Apr 26

Summary

The proposal aims to fix strict weak ordering problems in comparison-based algorithms to prevent security issues and uncover bugs in the user code. The proposal includes adding debug checks for strict weak ordering within `_LIBCPP_ENABLE_DEBUG_MODE`.

Background

Not so long ago we [proposed and changed std::sort algorithm](#) 25. However, it was rolled back in 16.0.1 because of failures within the broken comparators. That was even true for the [previous implementation](#) 4, however, the new one exposed the problems more often.

<https://discourse.llvm.org/t/rfc-strict-weak-ordering-checks-in-the-debug-libc/70217>

OOB check in 36d8b4

guarded since we know the last element is greater than the pivot.

```
__comp(__pivot, *++__first)) {  
  
    __rst;  
    P_ASSERT(__first != __end, "Would read out of bounds, is your comparator a va  
    (!__comp(__pivot, *__first));  
  
    + first < last && ! comp( pivot, * first)) {
```

Submitted SWO checks for std::{stable_,partial_}sort and rolled out in LLVM 17 in [7e1ee1](#)

```
25
26     template <class _RandomAccessIterator, class _Comp>
27     _LIBCPP_HIDE_FROM_ABI _LIBCPP_CONSTEXPR_SINCE_CXX14 void
28     __check_strict_weak_ordering_sorted(_RandomAccessIterator __
29     #ifdef _LIBCPP_DEBUG_STRICT_WEAK_ORDERING_CHECK
30         using __diff_t = __iter_diff_t<_RandomAccessIterator>;
31         using _Comp_ref = __comp_ref_type<_Comp>;
32         if (!__libcpp_is_constant_evaluated()) {
33             // Check if the range is actually sorted.
```

-D_LIBCPP_DEBUG_STRICT_WEAK_ORDERING_CHECK
-D_LIBCPP_ENABLE_DEBUG_MODE=1

It took us around **6-7 months** to fix everything

BUGS

We found fewer of them but more dangerous.

```
absl::c_sort(vector, [](const auto& lhs, const auto& rhs) {  
    return lhs >= rhs; // BUG, comp(a, a) is true  
} );
```

We found fewer of them but more dangerous.

```
const bool descending = absl::GetFlag(FLAGS_sort_descending);
absl::c_sort(vector, [descending](const auto& lhs,
                                const auto& rhs) {
    if (lhs < rhs) {
        return descending;
    }
    // BUG, if descending is false, comp(a, a) is true.
    return !descending;
});
```

We found fewer of them but more dangerous.

```
absl::c_sort(vector, [](const auto& lhs, const auto& rhs) {
    return lhs->value() < rhs->value();
} );

// In test

TEST(TestSort, NoValue) {
    std::vector<T> vector = {NoValueElement()};
    // Technically passes but incorrect.
    Sort(vector);
}
```

We found fewer of them but more dangerous.

```
absl::c_sort(vector, [](const auto& lhs, const auto& rhs) {
    if (!lhs.has_value()) {
        return true; // BUG, comp(a, a) can be true.
    }
    if (!rhs.has_value()) {
        return false;
    }
    return lhs.value() < rhs.value();
} );
```

We found fewer of them but more dangerous.

```
absl::c_sort(*response->mutable_elems() ,  
[&] (const Element& a,  
      const Element& b) {  
    int day_diff_a = ...;  
    int day_diff_b = ...;  
    if (day_diff_a != day_diff_b) {  
        // BUG, a = diff, b = -diff, c = diff can be  
        // comp(a, b) is true because of CompareClosely  
        // comp(b, c) is true because of CompareClosely  
        // but comp(a, c) is false because diff is same.  
        return std::abs(day_diff_a) < std::abs(day_diff_b);  
    }  
    return CompareClosely(a, b)  
} );
```

We found fewer of them but more dangerous.

```
absl::c_sort(vector, [](const auto& lhs, const auto& rhs) {  
    // Lots of logic.  
    // ...  
    // ...  
    return true; // BUG, comp(a, a) is true.  
});
```

We found fewer of them but more dangerous.

```
absl::c_sort(vector, [](const auto& lhs, const auto& rhs) {  
    absl::StatusOr<bool> cmp = Compare(lhs, rhs);  
    if (!cmp.ok()) {  
        // BUG, a, b, c can be like comp(a, b) is true,  
        // comp(b, c) is true, comp(a, c) has !cmp.ok().  
        return false;  
    }  
    // ...  
} );
```

We found fewer of them but more dangerous.

```
std::sort(points.begin(), points.end(),
          [](const DataPoint& point_1,
              const DataPoint& point_2) {
    if (MathUtil::AlmostEquals(point_1.coordinates[0],
                               point_2.coordinates[0])) {
        return point_1.id < point_2.id;
    }
    return point_1.coordinates[0] < point_2.coordinates[0];
}) ;
```

We found fewer of them but more dangerous.

(a is equal to $a + \epsilon$), ($a + \epsilon$ equals to $a + 2\epsilon$) but (a is not equal to $a + 2\epsilon$). Transitivity of incomparability is lost

We found fewer of them but more dangerous.

```
std::vector<float> vector_with_possibly_nans = Fill();
absl::c_sort(vector_with_possibly_nans);
// Can be bug:
// As doubles/floats can be NaN which means x < NaN and NaN < x
// are both false and that means x is equivalent to NaN
// thus for every finite x we have x == NaN but clearly
// x == NaN and y == NaN does not imply x == y.
// If you are sure your sort does not have NaNs, no action needed.
```

We found fewer of them but more dangerous.

The screenshot shows a debugger interface with two panes. The left pane is a C++ code editor titled "C++ source #1" containing the following code:

```
1 #include <algorithm>
2 #include <vector>
3
4 void Sort(std::vector<float>& v);
5
6 void Sort(std::vector<float>& v) {
7     std::sort(v.begin(), v.end(), std::greater<int>());
8 }
```

The right pane is a debugger window titled "x86-64 clang (trunk) (C++, Editor #1, Compiler #2)" showing assembly output. The assembly code is:

Line	Assembly
1	Sort(std::__1::vector<float, std::__1::allocator<float> > &v)
2	pushq %rax
3	movq %rdi, %rax
4	movq (%rdi), %rdi
5	movq 8(%rax), %rsi
6	movq %rsi, %rax

Below the assembly window, there is an "Output" tab showing "Compiler returned: 0".

We found fewer of them but more dangerous.

Repo: android/platform/system/unwinding

Stars: N/A

Commit: d96368

File: libunwindstack/tools/unwind_reg_info.cpp

Snippet:

```
85     std::sort(loc_regs.begin(), loc_regs.end(), [](auto a, auto b) {
86         if (a.first == CFA_REG) {
87             return true;
88         } else if (b.first == CFA_REG) {
89             return false;
90         }
91         return a.first < b.first;
92     });

```

We found fewer of them but more dangerous.

Repo: iovisor/bpftrace

Stars: 5.2k

Commit: 384e4d

File: src/utils.cpp

Snippet:

```
465     // Sort paths lexically by name (with the exception of unified, which always
466     // comes first)
467     std::sort(result.begin(), result.end(), [](auto &pair1, auto &pair2) {
468         if (pair1.first == "unified")
469             return true;
470         else if (pair2.first == "unified")
471             return false;
472         else
473             return pair1.first < pair2.first;
474     });

```

We found fewer of them but more dangerous.

Repo: ROCmSoftwarePlatform/rocSPARSE

Stars: 2.4k

Commit: 584393

File: clients/common/rocsparse_host.cpp

Snippet:

```
3535
3536     rocsparse_int* col_entry = &col[row_begin];
3537     T*           val_entry = &val[row_begin];
3538
3539     std::sort(perm.begin(), perm.end(), [&](const rocsparse_int& a, const rocsparse_int& b) {
3540         return col_entry[a] <= col_entry[b];
3541     });
3542
3543     for(rocsparse_int j = 0; j < row_nnz; ++j)
3544     {
3545         csr_col_ind_C[row_begin + j] = col_entry[perm[j]];
3546         csr_val_C[row_begin + j]    = val_entry[perm[j]];
3547     }
3548 }
```

We found fewer of them but more dangerous.

Repo: chromium/webrtc

Stars: N/A

Commit: a7b919

File: rtc_base/network.cc

Snippet:

```
65  bool CompareNetworks(const Network* a, const Network* b) {  
66      if (a->prefix_length() == b->prefix_length()) {  
67          if (a->name() == b->name()) {  
68              return a->prefix() < b->prefix();  
69          }  
70      }  
71      return a->name() < b->name();  
72 }
```

We found fewer of them but more dangerous.

```
// Self modifying comparator can lead to O(n^2) worst case
std::nth_element(tmp.begin(), tmp.end() - 2, tmp.end(), [&](int x,
    if (v[x] == gas && v[y] == gas) {
        if (x == candidate) v[x] = num_solid++;
        else v[y] = num_solid++;
    }
    if (v[x] == gas) candidate = x;
    else if (v[y] == gas) candidate = y;
    return v[x] < v[y];
} );
```

We found fewer of them but more dangerous.

libc++ std::nth_element is quadratic, should be linear
average #52747

SORTING TRIVIA

Can we sort {1.02, 1.01, 1.0}?

Can we sort {1.02, 1.01, 1.0}?

Yes

Can we sort {3.0, NaN, 4.0}?

Can we sort {3.0, NaN, 4.0}?

No

Can we sort {3.0, NaN, 4.0}?

No

NaN thinks it is equal to 3.0 and 4.0

Can we sort {NaN, 3.0, NaN}?

Can we sort {NaN, 3.0, NaN}?

Yes

Can we sort {NaN, 3.0, NaN}?

Yes

All elements are equal

WHAT ARE WE CHANGING?

WHAT ARE WE CHANGING?

- After these cleanups we can change to any sort

WHAT ARE WE CHANGING?

- After these cleanups we can change to any sort
- We chose a mix of BlockQuickSort and pdqsort

```
1 // *m is median. partition [first, m) < *m and
2 // *m <= [m, last)
3 // Special handling for almost sorted targets
4 while (true) {
5     while (comp(*++i, *m));
6     while (!comp(*--j, *m));
7     if (i > j) break;
8     swap(*i, *j);
9 }
10 swap(*i, *m);
```

Data dependent branches are hard to predict

```
1 // *m is median. partition [first, m) < *m and
2 // *m <= [m, last)
3 // Special handling for almost sorted targets
4 while (true) {
5     while (comp(*++i, *m));
6     while (!comp(*--j, *m));
7     if (i > j) break;
8     swap(*i, *j);
9 }
10 swap(*i, *m);
```

Data dependent branches are hard to predict

```
1 // *m is median. partition [first, m) < *m and
2 // *m <= [m, last)
3 // Special handling for almost sorted targets
4 while (true) {
5     while (comp(*++i, *m));
6     while (!comp(*--j, *m));
7     if (i > j) break;
8     swap(*i, *j);
9 }
10 swap(*i, *m);
```

Pack results into blocks

```
1 uint64_t __left_bitset = 0;
2 _RandomAccessIterator __iter = __first;
3 for (int __j = 0; __j < __block_size;) {
4     bool __comp_result = !__comp(*__iter, __pivot);
5     __left_bitset |= (__comp_result << __j);
6     __j++;
7     ++__iter;
8 }
```

Pack results into blocks

```
1 uint64_t __left_bitset = 0;
2 _RandomAccessIterator __iter = __first;
3 for (int __j = 0; __j < __block_size;) {
4     bool __comp_result = !__comp(*__iter, __pivot);
5     __left_bitset |= (__comp_result << __j);
6     __j++;
7     ++__iter;
8 }
```

CTZ - Count Trailing Zeros, BLSR - Reset Lowest Set Bit

```
1 void __swap_bitset_pos(__first, __last, __left_bitset, __rig
2   while (__left_bitset != 0 && __right_bitset != 0) {
3     difference_type tz_left = __ctz(__left_bitset);
4     __left_bitset = __blsr(__left_bitset);
5     difference_type tz_right = __ctz(__right_bitset);
6     __right_bitset = __blsr(__right_bitset);
7     __VSTD::iter_swap(__first + tz_left, __last - tz_right);
8   }
9 }
```

CTZ - Count Trailing Zeros, BLSR - Reset Lowest Set Bit

```
1 void __swap_bitset_pos(__first, __last, __left_bitset, __rig
2   while (__left_bitset != 0 && __right_bitset != 0) {
3     difference_type tz_left = __ctz(__left_bitset);
4     __left_bitset = __blsr(__left_bitset);
5     difference_type tz_right = __ctz(__right_bitset);
6     __right_bitset = __blsr(__right_bitset);
7     __VSTD::iter_swap(__first + tz_left, __last - tz_right);
8   }
9 }
```

CTZ - Count Trailing Zeros, BLSR - Reset Lowest Set Bit

```
1 void __swap_bitset_pos(__first, __last, __left_bitset, __rig
2   while (__left_bitset != 0 && __right_bitset != 0) {
3     difference_type tz_left = __ctz(__left_bitset);
4     __left_bitset = __blsr(__left_bitset);
5     difference_type tz_right = __ctz(__right_bitset);
6     __right_bitset = __blsr(__right_bitset);
7     __VSTD::iter_swap(__first + tz_left, __last - tz_right);
8   }
9 }
```

CTZ - Count Trailing Zeros, BLSR - Reset Lowest Set Bit

```
1 void __swap_bitset_pos(__first, __last, __left_bitset, __rig
2   while (__left_bitset != 0 && __right_bitset != 0) {
3     difference_type tz_left = __ctz(__left_bitset);
4     __left_bitset = __blsr(__left_bitset);
5     difference_type tz_right = __ctz(__right_bitset);
6     __right_bitset = __blsr(__right_bitset);
7     __VSTD::iter_swap(__first + tz_left, __last - tz_right);
8   }
9 }
```

We also improved picking a pivot out of 9 elements. So called "Tukey's ninther"

```
1 sort3(first,
2        first + half_len,
3        last - 1, comp);
4 sort3(first + 1,
5        first + (half_len - 1),
6        last - 2, comp);
7 sort3(first + 2,
8        first + (half_len + 1),
9        last - 3, comp);
10
11 sort3(first + (half_len - 1),
12        first + half_len,
13        first + (half_len + 1), comp);
```

We also improved picking a pivot out of 9 elements. So called "Tukey's ninther"

```
1 sort3(first,
2        first + half_len,
3        last - 1, comp);
4 sort3(first + 1,
5        first + (half_len - 1),
6        last - 2, comp);
7 sort3(first + 2,
8        first + (half_len + 1),
9        last - 3, comp);
10
11 sort3(first + (half_len - 1),
12        first + half_len,
13        first + (half_len + 1), comp);
```



Add... More

Sponsors

C++ source #1 x

A - B + v ↻

C++

```
1 #include <cstdint>
2
3 int64_t* __bitset_partition(int64_t* __first, int64_t* __last,
4                             int size) {
5     int64_t __left_bitset = 0;
6     int64_t __right_bitset = 0;
7
8     // Reminder: length = __last - __first + 1.
9     while (size > 2 * 64 - 1) {
10         if (__left_bitset == 0) {
11             int64_t* __iter = __first;
12             for (int __j = 0; __j < 64;) {
13                 bool __comp_result = *__iter >= __pivot;
14                 __left_bitset |= (static_cast<int64_t>(__comp_result))
15                 __j++;
16                 ++__iter;
17             }
18             size -= 64;
19         }
20         if (__right_bitset == 0) {
21             // Possible vectorization. With a proper "-march" flag,
22             // will be compiled into a set of SIMD instructions.
23             int64_t* __iter = __last - 1;
24             for (int __j = 0; __j < 64;) {
25                 bool __comp_result = *__iter < __pivot;
26                 __right_bitset |= (static_cast<int64_t>(__comp_result))
27                 __j++;
28                 --__iter;
29             }
30             size -= 64;
31         }
32         __first += (__left_bitset == 0) ? 64 : 0;
33         __last -= (__right_bitset == 0) ? 64 : 0;
34     }
35     return __first;
36 }
37
```

x86-64 clang 14.0.0 (C++, Editor #1, Compiler #1) x

x86-64 clang 14.0.0 ✓ -msse4.2 -O3

A - B Output... Filter... Libraries + Add new... Add tool...

138	movdqu	xmm0, ymmword ptr [rax]
139	movdqu	xmm2, ymmword ptr [rax + 16]
140	movdqu	xmm3, ymmword ptr [rax + 32]
141	movdqqa	xmm4, ymm11
142	pcmpgtq	xmm4, ymm0
143	movdqqa	xmm0, ymm11
144	pcmpgtq	xmm0, ymm2
145	movdqu	xmm2, ymmword ptr [rax + 48]
146	pandn	xmm4, ymm5
147	pandn	xmm0, ymm6
148	por	xmm0, ymm4
149	movdqqa	xmm4, ymm11

C Output (0/0) x86-64 clang 14.0.0 - 869ms (861678) ~1334 lines filtered

x86-64 clang 14.0.0 ✓ -mavx2 -O3

A - B Output... Filter... Libraries + Add new... Add tool...

121	.LBB0_3.	# IN LOOP. header =
122	vpcmpgtq	ymm2, ymm0, ymmword ptr [rax]
123	vpandn	ymm2, ymm2, ymm5
124	vpcmpgtq	ymm3, ymm0, ymmword ptr [rax + 32]
125	vpandn	ymm3, ymm3, ymm6
126	vpor	ymm2, ymm3, ymm2
127	vpcmpgtq	ymm3, ymm0, ymmword ptr [rax + 64]
128	vpandn	ymm3, ymm3, ymm7
129	vpcmpgtq	ymm4, ymm0, ymmword ptr [rax + 96]
130	vpandn	ymm4, ymm4, ymm8
131	vpor	ymm3, ymm3, ymm4
132	vpor	ymm2, ymm2, ymm3
133	vpcmpgtq	ymm3, ymm0, ymmword ptr [rax + 128]

<https://godbolt.org/z/nrhT88MsM>

	old	new	
BM_Sort_uint64_Rand_1	4.26ns ± 0%	4.41ns ± 1%	3.44%
BM_Sort_uint64_Rand_4	1.96ns ± 0%	1.95ns ± 1%	-0.34%
BM_Sort_uint64_Rand_16	10.5ns ± 0%	11.6ns ± 1%	10.77%
BM_Sort_uint64_Rand_64	18.8ns ± 0%	17.5ns ± 1%	-6.60%
BM_Sort_uint64_Rand_256	26.4ns ± 0%	21.7ns ± 1%	-17.88%
BM_Sort_uint64_Rand_1024	33.8ns ± 1%	24.6ns ± 1%	-27.15%
BM_Sort_uint64_Rand_16384	48.7ns ± 0%	30.1ns ± 1%	-38.10%
BM_Sort_uint64_Rand_262144	63.4ns ± 1%	35.7ns ± 1%	-43.71%

On other patterns we sometimes regressed but by a little, hardest cases are random.

BM_Sort_string_Random_1	4.66ns ± 6%	4.40ns ± 4%	-5.55%
BM_Sort_string_Random_4	14.9ns ± 3%	15.0ns ± 6%	~
BM_Sort_string_Random_16	45.5ns ± 6%	35.8ns ± 8%	-21.37%
BM_Sort_string_Random_64	66.6ns ± 4%	58.2ns ± 3%	-12.69%
BM_Sort_string_Random_256	86.0ns ± 5%	77.4ns ± 3%	-10.01%
BM_Sort_string_Random_1024	106ns ± 3%	96ns ± 6%	-9.39%
BM_Sort_string_Random_16384	154ns ± 3%	141ns ± 5%	-8.03%
BM_Sort_string_Random_262144	213ns ± 4%	197ns ± 4%	-7.59%

SOME THOUGHTS

SOME THOUGHTS

- Default for standard sorting is unstable

SOME THOUGHTS

- Default for standard sorting is unstable
- Python, Rust, Java have stable default sorts

SOME THOUGHTS

- Default for standard sorting is unstable
- Python, Rust, Java have stable default sorts
- Why do we have `std::sort` that can read OOB even if comparator is broken?

SOME THOUGHTS

- Default for standard sorting is unstable
- Python, Rust, Java have stable default sorts
- Why do we have `std::sort` that can read OOB even if comparator is broken?
- `std::ranges::sort` wanted to have 1 argument, found 0 bugs with different iterators in `std::sort`

SOME MORE THOUGHTS

- `std::weak_ordering` in C++20 allows to make fewer bugs

```
struct Weak {
    bool operator==(Weak const&) const;
    std::weak_ordering operator<=(Weak const&) const;
};

struct Foo {
    Weak w;
    int i;
    auto operator<=(Foo const&) const = default;
};
```

RESULTS

RESULTS

- We called it a **comparator sanitizer**

RESULTS

- We called it a **comparator sanitizer**
- We saw 7-8% perf improvement in production

RESULTS

- We called it a **comparator sanitizer**
- We saw 7-8% perf improvement in production
- Reduced branch mispredictions a lot

RESULTS

- We called it a **comparator sanitizer**
- We saw 7-8% perf improvement in production
- Reduced branch mispredictions a lot
- Fixed thousands of bugs

RESULTS

- We called it a **comparator sanitizer**
- We saw 7-8% perf improvement in production
- Reduced branch mispredictions a lot
- Fixed thousands of bugs
- Fought the Hyrum's Law

RESULTS

- We called it a **comparator sanitizer**
- We saw 7-8% perf improvement in production
- Reduced branch mispredictions a lot
- Fixed thousands of bugs
- Fought the Hyrum's Law
- Even simplest things are hard to do right

RESULTS

- We called it a **comparator sanitizer**
- We saw 7-8% perf improvement in production
- Reduced branch mispredictions a lot
- Fixed thousands of bugs
- Fought the Hyrum's Law
- Even simplest things are hard to do right
- Things take time (2.5 years)

WHAT CAN YOU DO?

WHAT CAN YOU DO?

- If you use libcxx, enable -

D_LIBCPP_DEBUG_RANDOMIZE_UNSPECIFIED_STARTUP
-D_LIBCPP_DEBUG_STRICT_WEAK_ORDERING_CHECKS
D_LIBCPP_ENABLE_DEBUG_MODE=1 in your **debug** builds

WHAT CAN YOU DO?

- If you use libcxx, enable -
`D_LIBCPP_DEBUG_RANDOMIZE_UNSPECIFIED_STARTUP`
`-D_LIBCPP_DEBUG_STRICT_WEAK_ORDERING_CHECKS`
`D_LIBCPP_ENABLE_DEBUG_MODE=1` in your **debug builds**
- Probably randomization can be applied to std::partition_point
std::make_heap

WHAT CAN YOU DO?

- If you use libcxx, enable -
`D_LIBCPP_DEBUG_RANDOMIZE_UNSPECIFIED_STARTUP`
`-D_LIBCPP_DEBUG_STRICT_WEAK_ORDERING_CHECKS`
`D_LIBCPP_ENABLE_DEBUG_MODE=1` in your **debug builds**
- Probably randomization can be applied to std::partition_point
std::make_heap
- SWO checks for std::nth_element, etc

WHAT CAN YOU DO?

WHAT CAN YOU DO?

- Remember Bjarne talked about invalidation of vector iterators?

WHAT CAN YOU DO?

- Remember Bjarne talked about invalidation of vector iterators?
- In debug mode reallocate on **first 100 push_backs**

LINKS AND FURTHER

- [danlark.org blog about sorting](#)
- [LLVM RFC for strict weak ordering checks](#)
- [Sorting change itself with benchmarks, etc](#)

- Personal Email: kutdanila@gmail.com
- Work Email: danielak@google.com
- Github: [@Danlark1](https://github.com/Danlark1)
- Telegram: [@Danlark](https://t.me/Danlark)

