

+ 23

# *Back To Basics* Concurrency

DAVID OLSEN



20  
23













Jetstar Airways (jetstar.com), CC BY-SA 2.0, via Wikimedia Commons






Kentaro Iemoto from Tokyo, Japan, CC BY-SA 2.0, via Wikimedia Commons



# BACK TO BASICS: CONCURRENCY

David Olsen, Software Engineer, NVIDIA  
CppCon, October 5, 2023





# STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

# STD::COUNTING\_SEMAPHORE

Lightweight synchronization primitive that can control access to a shared resource

Maintains an internal counter

Calls to `release()` increment the counter

Calls to `acquire()` decrement the counter, or block if counter == 0

Calls to `acquire()` and `release()` can happen on different threads

[https://en.cppreference.com/w/cpp/thread/counting\\_semaphore](https://en.cppreference.com/w/cpp/thread/counting_semaphore)



# CONCURRENCY

## Definition

“Multiple logical threads of execution with unknown inter-task dependencies”

- Daisy Hollman, “[A Unifying Abstraction for Async](#),” CppCon 2019

# CONCURRENCY

## Definition

“Multiple logical threads of execution with **unknown** inter-task dependencies”

- Daisy Hollman, “[A Unifying Abstraction for Async](#),” CppCon 2019



# CONCURRENCY

## Definition

“Multiple logical threads of execution with **[some]** inter-task dependencies”

# CONCURRENCY

## Definition

“Multiple logical threads of execution with [some] inter-task dependencies”

Doing things at the same time



# CONCURRENCY

## Definition

“Multiple logical threads of execution with [some] inter-task dependencies”

Doing things at the same time

Some things need to happen before other things

# CONCURRENCY

## Definition

“Multiple logical threads of execution with [some] inter-task dependencies”

Doing things at the same time

Some things need to happen before other things

Some things can't happen at the same time



# CONCURRENCY

## Definition

“Multiple logical threads of execution with [some] inter-task dependencies”

Doing things at the same time

Some things need to happen before other things

Some things can't happen at the same time

threads



# CONCURRENCY

## Definition

“Multiple logical threads of execution with [some] inter-task dependencies”

Doing things at the same time

Some things need to happen before other things

Some things can't happen at the same time

synchronization



# PARALLELISM

## Definition

“Multiple logical threads of execution with no inter-task dependencies”

- Daisy Hollman, “[A Unifying Abstraction for Async](#),” CppCon 2019





# CONCURRENCY MADE EASY

# CONCURRENCY MADE EASY

Not easy, but a little less hard

# CONCURRENCY MADE EASY

Not easy, but a little less hard

And it's really parallelism, not concurrency



# CONCURRENCY MADE EASY

Well, less hard. And it's parallelism, not concurrency

## C++ parallel algorithms

Let the compiler do all the hard work for you

Bryce Adelstein Lelbach, “[The C++17 Parallel Algorithms Library and Beyond](#),” CppCon 2016

David Olsen, “[Faster Code Through Parallelism on CPUs and GPUs](#),” CppCon 2019

# C++ PARALLEL ALGORITHMS

## Execution policies

Add execution policy as first argument to algorithm call

`std::execution::seq`

Run sequentially, no parallelism

`std::execution::par`, `std::execution::par_unseq`

Request to compiler to run in parallel

Promise by user that code is safe to run in parallel; no data races

`std::execution::unseq`, `std::execution::par_unseq`

Request to compiler to vectorize

Promise by user that code is safe to vectorize; no data races or locks

# C++ PARALLEL ALGORITHMS

## Existing algorithms

If using standard algorithms already, add execution policy argument

```
std::sort(items.begin(), items.end(), compare_by_price{});
```

```
std::fill(v.begin(), v.end(), -1);
```

# C++ PARALLEL ALGORITHMS

## Existing algorithms

If using standard algorithms already, add execution policy argument

```
std::sort(std::execution::par,  
          items.begin(), items.end(), compare_by_price{});
```

```
std::fill(std::execution::par_unseq,  
          v.begin(), v.end(), -1);
```



# C++ PARALLEL ALGORITHMS

## Loops to algorithms

Some raw loops can be converted into parallel algorithms

```
std::int64_t sum = 0;
for (std::int64_t x : v) {
    sum += x;
}
```

# C++ PARALLEL ALGORITHMS

## Loops to algorithms

Some raw loops can be converted into parallel algorithms

```
std::int64_t sum = 0;
for (std::int64_t x : v) {
    sum += x;
}
```

```
std::int64_t sum = 0;
std::for_each(std::execution::par_unseq, v.begin(), v.end(),
    [&sum](std::int64_t x){ sum += x; });
```

<https://godbolt.org/z/3PWxd79rT>

# C++ PARALLEL ALGORITHMS

## Loops to algorithms

Some raw loops can be converted into parallel algorithms

```
std::int64_t sum = 0;
for (std::int64_t x : v) {
    sum += x;
}
```

```
std::int64_t sum = 0;
std::for_each(std::execution::par_unseq, v.begin(), v.end(),
    [&sum](std::int64_t x){ sum += x; });
```

**Data race!**  
**Don't do this!**



# C++ PARALLEL ALGORITHMS

## Loops to algorithms

Some raw loops can be converted into parallel algorithms

```
std::int64_t sum = 0;
for (std::int64_t x : v) {
    sum += x;
}
```

```
std::int64_t sum = 0;
std::mutex mtx;
std::for_each(std::execution::par, v.begin(), v.end(),
    [&sum, &mtx](std::int64_t x){
        std::scoped_lock lock(mtx);
        sum += x;
    });
```

<https://godbolt.org/z/on4Koejqx>



# C++ PARALLEL ALGORITHMS

## Loops to algorithms

Some raw loops can be converted into parallel algorithms

```
std::int64_t sum = 0;
for (std::int64_t x : v) {
    sum += x;
}
```

Abysmal performance!  
Don't do this!



```
std::int64_t sum = 0;
std::mutex mtx;
std::for_each(std::execution::par, v.begin(), v.end(),
    [&sum, &mtx](std::int64_t x){
        std::scoped_lock lock(mtx);
        sum += x;
    });
```

<https://godbolt.org/z/on4Koejqx>

# C++ PARALLEL ALGORITHMS

## Loops to algorithms

Some raw loops can be converted into parallel algorithms

```
std::int64_t sum = 0;
for (std::int64_t x : v) {
    sum += x;
}
```

```
std::int64_t sum = std::reduce(std::execution::par_unseq,
    v.begin(), v.end(), std::int64_t(0), std::plus<>{});
```

# C++ PARALLEL ALGORITHMS

Some of the algorithms

`for_each` `for_each_n` `transform`

`find` `find_if` `find_end` `search` `count` `any_of` `adjacent_find`

`copy` `copy_if` `move` `fill` `replace` `generate` `rotate`

`sort` `stable_sort` `partial_sort` `nth_element` `is_sorted`

`reduce` `transform_reduce` `inclusive_scan` `exclusive_scan`

# C++ PARALLEL ALGORITHMS

Some of the algorithms

**for\_each for\_each\_n transform**

find find\_if find\_end search count any\_of adjacent\_find

copy copy\_if move fill replace generate rotate

sort stable\_sort partial\_sort nth\_element is\_sorted

reduce transform\_reduce inclusive\_scan exclusive\_scan



# C++ PARALLEL ALGORITHMS

Some of the algorithms

for\_each for\_each\_n transform

find find\_if find\_end search count any\_of adjacent\_find

copy copy\_if move fill replace generate rotate

sort stable\_sort partial\_sort nth\_element is\_sorted

reduce transform\_reduce inclusive\_scan exclusive\_scan

# C++ PARALLEL ALGORITHMS

Some of the algorithms

for\_each for\_each\_n transform

find find\_if find\_end search count any\_of adjacent\_find

copy copy\_if move fill replace generate rotate

sort stable\_sort partial\_sort nth\_element is\_sorted

reduce transform\_reduce inclusive\_scan exclusive\_scan

# C++ PARALLEL ALGORITHMS

Some of the algorithms

for\_each for\_each\_n transform

find find\_if find\_end search count any\_of adjacent\_find

copy copy\_if move fill replace generate rotate

sort stable\_sort partial\_sort nth\_element is\_sorted

reduce transform\_reduce inclusive\_scan exclusive\_scan

# C++ PARALLEL ALGORITHMS

Some of the algorithms

for\_each for\_each\_n transform

find find\_if find\_end search count any\_of adjacent\_find

copy copy\_if move fill replace generate rotate

sort stable\_sort partial\_sort nth\_element is\_sorted

reduce transform\_reduce inclusive\_scan exclusive\_scan

# C++ PARALLEL ALGORITHMS

Some of the algorithms

for\_each for\_each\_n transform

find find\_if find\_end search count any\_of adjacent\_find

copy copy\_if move fill replace generate rotate

sort stable\_sort partial\_sort nth\_element is\_sorted

reduce **transform\_reduce** inclusive\_scan exclusive\_scan

# STD::TRANSFORM\_REDUCE

## Word count

```
std::size_t word_count(std::string_view s) {  
    if (s.empty()) return 0;  
    return std::transform_reduce(std::execution::par_unseq,  
        begin(s), end(s) - 1, begin(s) + 1,  
        std::size_t(!std::isspace(s.front()) ? 1 : 0),  
        std::plus<>{},  
        [] (char l, char r) {  
            return std::isspace(l) && !std::isspace(r);  
        });  
}
```

Bryce Adelstein LeBlach, “[C++17 Parallel Algorithms and Beyond](#),” CppCon 2016



# STD::TRANSFORM\_REDUCE

## Traveling salesman

```
route_cost find_best_route(int const* distances, int N) {  
    return std::transform_reduce(std::execution::par,  
        counting_iterator<long>(0L), counting_iterator<long>(factorial(N)),  
        route_cost(),  
        route_cost::min,  
        [=](long i) {  
            int cost = 0;  
            route_iterator it(i, N);  
            int from = it.first();  
            while (!it.done()) {  
                int to = it.next();  
                cost += distances[from*N + to];  
                from = to;  
            }  
            return route_cost(i, cost);  
        });  
}
```

David Olsen, "[Faster Code Through Parallelism on CPUs and GPUs](#)," CppCon 2019

# C++ PARALLEL ALGORITHMS

## Summary

Use C++ standard parallel algorithms when appropriate

Use the right algorithm for the task

Make sure your code is thread safe with no data races

# STD::THREAD

# STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

## STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

std::thread variable is a handle, not the thread

Exists in the main thread, not in the created thread

## STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

Pass a callable to `std::thread` constructor

Can be lambda, pointer to function, `std::function`, object with call operator



## STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

Arguments for the callable are also passed to the `std::thread` constructor

## STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

`std::thread` constructor creates a new thread and runs the callable on that thread

## STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

`join()` waits for the thread to complete, blocking if necessary

## STD::THREAD

Hello, world!

```
int main() {  
    std::thread my_thread{[](int z){  
        std::cout << "Hello from thread: " << z << "\n";  
    }, 42};  
    my_thread.join();  
}
```

`join()` waits for the thread to complete, blocking if necessary

Must call `join()` or `detach()`, or program will `std::terminate()`

# STD::THREAD

Low-level tool

Don't use `std::thread` directly

Use a threading library or other task management system

Similar characteristics to `std::thread`:

- A way to start a thread or start work concurrently

- A way to wait for the task to complete, or at least check for completion

Might have a way to return a value

# DATA RACES

## Race Conditions



# DATA RACES

## Race Conditions

Conflict over a resource without coordination

Bad things happen as a result

# DATA RACES

## Examples

Counting people in a room while people are going in and out

Two people trying to walk through a door at the same time

# DATA RACES

## Examples

Counting people in a room while people are going in and out

Two people trying to walk through a door at the same time

Outfielders both trying to catch a fly ball



# DATA RACES

## Definition

Data race:

1. Two or more threads access the same memory
2. At least one access is a write
3. The threads do not synchronize with each other

A data race is undefined behavior

# DATA RACES

## Race Conditions

It's all about visibility of memory changes

When threads synchronize, changes made by one thread are guaranteed to be visible in the other thread

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

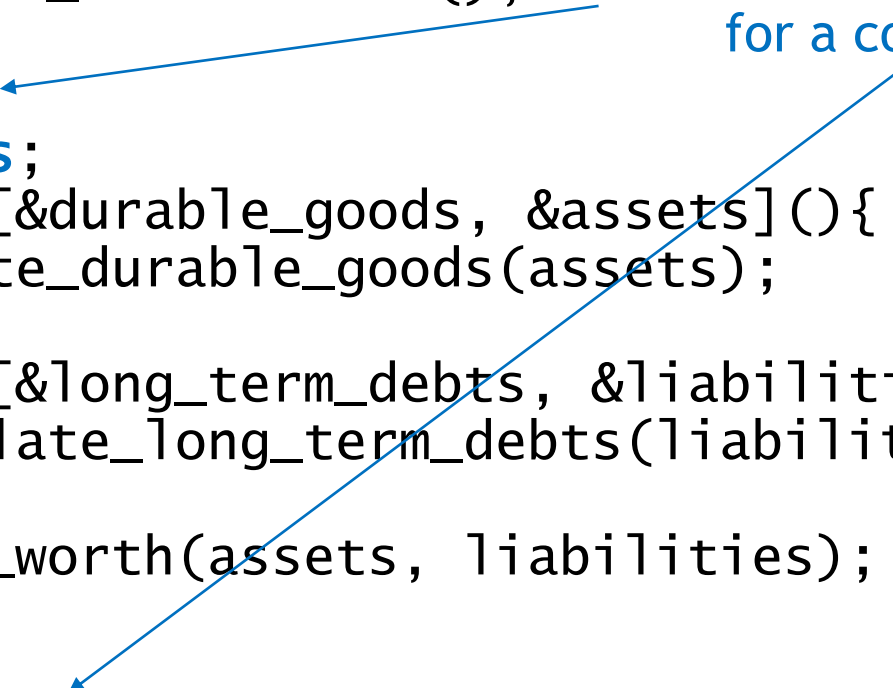


# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Goal: financial report for a company



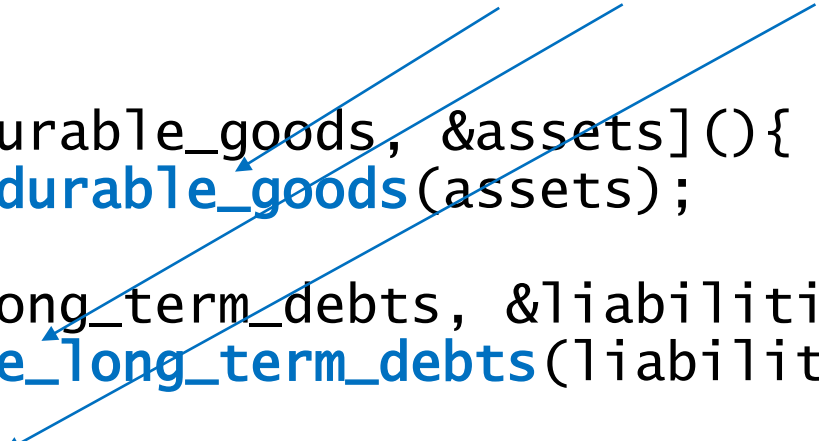
The diagram consists of two blue arrows. One arrow originates from the text 'Goal: financial report for a company' and points to the variable 'durable\_goods' in the code. The other arrow originates from the same text and points to the variable 'long\_term\_debts' in the code. This visualizes that both variables are used in the final 'create\_report' function call, despite being updated by separate threads, which creates a data race.

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Calculations can be slow



# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

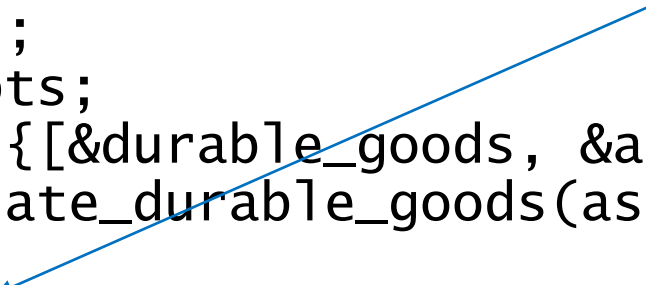
Background thread  
for durable goods

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

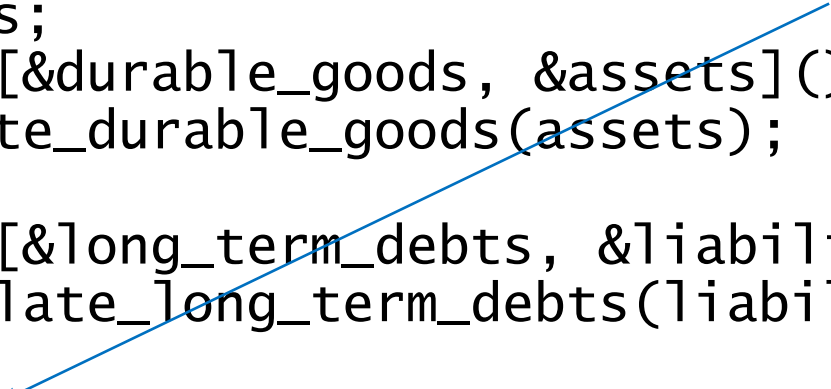
Background thread  
for liabilities



# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```



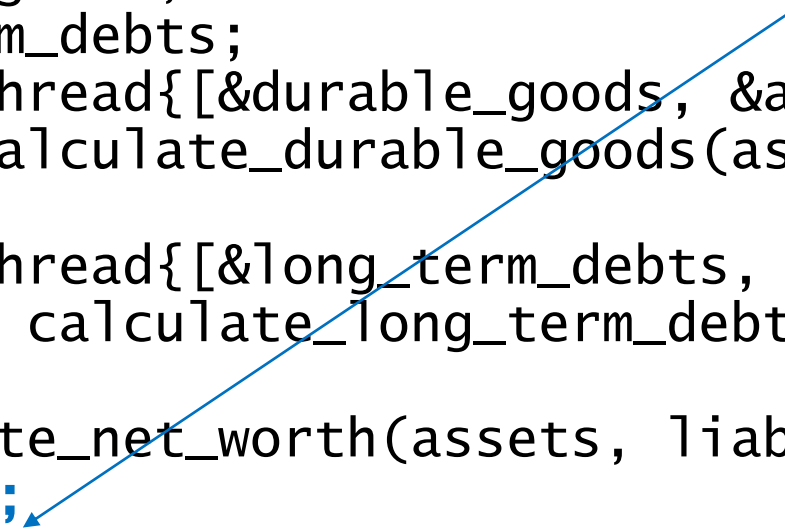
Net worth on  
main thread

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Wait for threads to complete



# DATA RACE ?

Is this a data race?

```
auto assets ← list_all_assets();           Parent thread write
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets](){
    durable_goods = calculate_durable_goods(assets);
}};                                       Child thread read
std::thread debts_thread{[&long_term_debts, &liabilities](){
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

# SYNCHRONIZATION: THREAD CREATION

Creating a thread synchronizes the parent and child threads

All memory changes in parent thread before thread creation are visible in the child thread



# DATA RACE ?

Is this a data race?

```
auto assets ← list_all_assets();           Parent thread write
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets](){      Child thread read
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities](){
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

**Synchronization**

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();  
auto liabilities = list_all_liabilities();  
currency_t net_worth;  
currency_t durable_goods;  
currency_t long_term_debts;  
std::thread goods_thread{[&durable_goods, &assets]() {  
    durable_goods = calculate_durable_goods(assets);  
}};  
std::thread debts_thread{[&long_term_debts, &liabilities]() {  
    long_term_debts = calculate_long_term_debts(liabilities);  
}};  
net_worth = calculate_net_worth(assets, liabilities);  
goods_thread.join();  
debts_thread.join();  
create_report(net_worth, durable_goods, long_term_debts);
```

Parent thread write

Synchronization

Child thread read

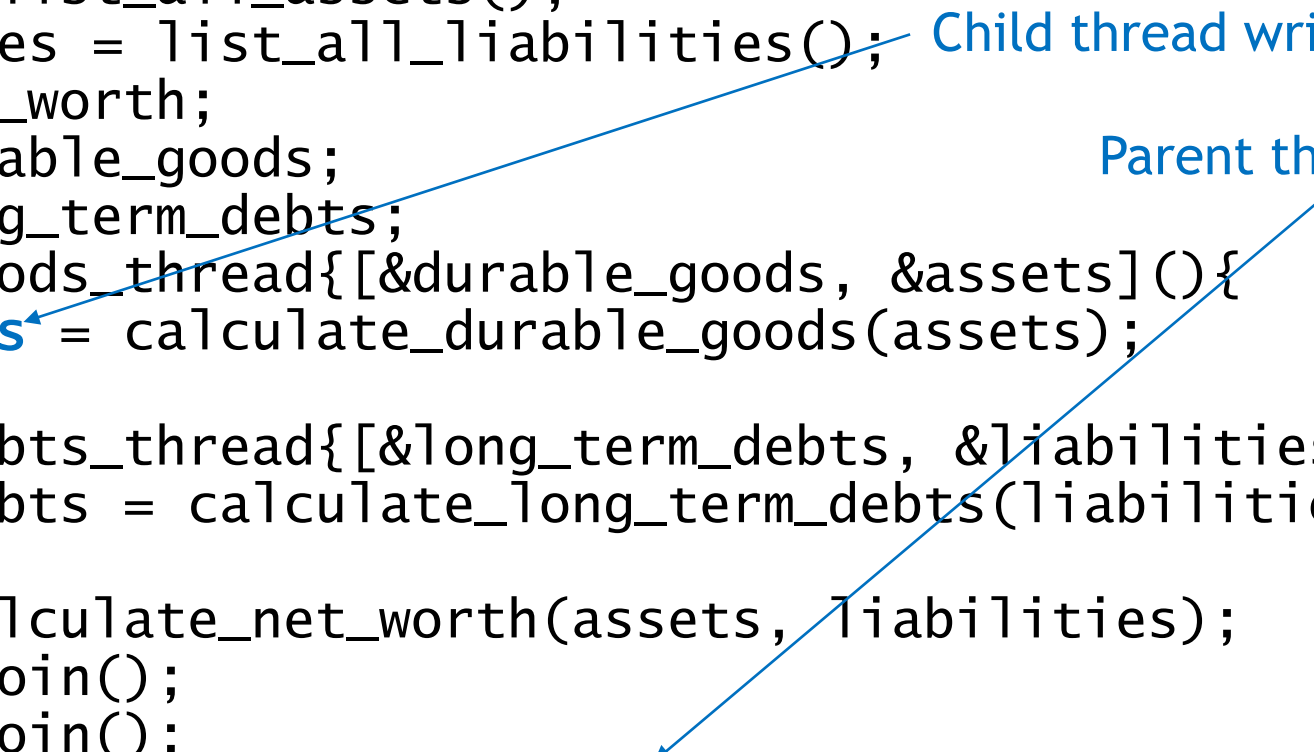
# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Child thread write

Parent thread read



The diagram illustrates a data race between two threads. A blue arrow labeled 'Child thread write' points from the text 'Child thread write' to the assignment of 'durable\_goods' in the 'goods\_thread' lambda function. Another blue arrow labeled 'Parent thread read' points from the text 'Parent thread read' to the use of 'durable\_goods' in the 'create\_report' function call. Both threads are active simultaneously, accessing the same memory location without synchronization, which constitutes a data race.

# SYNCHRONIZATION: THREAD JOIN

Joining a thread synchronizes the parent and child threads

All memory changes in child thread are visible in parent thread after call to `join()`

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Child thread write

Parent thread read

Synchronization

# DATA RACE ?

Is this a data race?

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Child thread

Parent thread

# DATA RACE ?

Is this a data race?

All reads,  
no synchronization needed

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```

Child thread

Parent thread

# NOT A DATA RACE

```
auto assets = list_all_assets();
auto liabilities = list_all_liabilities();
currency_t net_worth;
currency_t durable_goods;
currency_t long_term_debts;
std::thread goods_thread{[&durable_goods, &assets]() {
    durable_goods = calculate_durable_goods(assets);
}};
std::thread debts_thread{[&long_term_debts, &liabilities]() {
    long_term_debts = calculate_long_term_debts(liabilities);
}};
net_worth = calculate_net_worth(assets, liabilities);
goods_thread.join();
debts_thread.join();
create_report(net_worth, durable_goods, long_term_debts);
```



# MUTEXES

MUTual EXclusion

Most basic form of synchronization

Only one thread at a time can lock or acquire the mutex

# MUTEXES

## Example

### An airplane lavatory

Anyone can go in and lock the door

Has exclusive use of lavatory while the door is locked

Unlocking the door and leaving makes it available for the next person



Kristoferb, CC BY-SA 3.0, via Wikimedia Commons

# MUTEXES

## Example

### An airplane lavatory

Anyone can go in and lock the door

Has exclusive use of lavatory while the door is locked

Unlocking the door and leaving makes it available for the next person



Kristoferb, CC BY-SA 3.0, via Wikimedia Commons

# MUTEXES

## Example

A talking stick in a group discussion

Only the person holding the stick is supposed to talk



# MUTEXES

## Example

Air traffic control is the mutex

Pilot gets permission to use runway



# STD::MUTEX

## Member functions

# STD::MUTEX

## Member functions

```
void lock();
```

Blocks until the lock is acquired

## STD::MUTEX

### Member functions

```
void lock();
```

Blocks until the lock is acquired

```
bool try_lock();
```

Returns immediately

true if lock acquired; false if not



## STD::MUTEX

### Member functions

```
void lock();
```

Blocks until the lock is acquired

```
bool try_lock();
```

Returns immediately

true if lock acquired; false if not

```
void unlock();
```

Releases the lock

Undefined behavior if current thread doesn't own the lock

<https://en.cppreference.com/w/cpp/thread/mutex>

# STD::MUTEX

## Synchronization

A call to `lock()` synchronizes with any previous call to `unlock()` on the same mutex

Anything that happened on thread 1 before it called `unlock()` is visible to thread 2 after its call to `lock()` returns

# STD::MUTEX

## Synchronization

### Thread 1

```
// Section 1A  
m.lock();  
  
// Section 1B  
m.unlock();  
  
// Section 1C
```

### Thread 2

```
// Section 2A  
m.lock();  
  
// Section 2B  
m.unlock();  
  
// Section 2C
```

# STD::MUTEX

## Synchronization

### Thread 1

// Section 1A

m.lock();

// Section 1B

m.unlock();

// Section 1C

### Thread 2

// Section 2A

m.lock();

// Section 2B

m.unlock();

// Section 2C

Synchronization



# STD::MUTEX

## Synchronization

### Thread 1

// Section 1A

m.lock();

// Section 1B

m.unlock();

// Section 1C

### Thread 2

// Section 2A

m.lock();

// Section 2B

m.unlock();

// Section 2C

Visible

Synchronization

# STD::MUTEX

## Member functions

Don't call `lock()`, `try_lock()`, or `unlock()` on a mutex

# STD::MUTEX

## Member functions

Don't call `lock()`, `try_lock()`, or `unlock()` on a mutex

Always use a lock guard instead

# LOCK GUARDS

RAII wrapper around mutexes

Constructor calls lock()

Destructor calls unlock()

Guarantee that the mutex is always released



# STD::SCOPED\_LOCK

Constructor takes one or more mutexes

Calls lock() on each of the mutexes

Destructor calls unlock() on each of the mutexes

Not copyable or movable

No member functions or other operations

C++17

Uses CTAD to deduce class template arguments

```
std::scoped_lock lock(mutex_a, mutex_b);
```

[https://en.cppreference.com/w/cpp/thread/scoped\\_lock](https://en.cppreference.com/w/cpp/thread/scoped_lock)

# STD::LOCK\_GUARD

Constructor takes one mutex

Calls lock() on the mutex

Destructor calls unlock() on the mutex

Not copyable or movable

No member functions or other operations

Can't assume CTAD is available

```
std::lock_guard<std::mutex> lock(mutex_a);
```

[https://en.cppreference.com/w/cpp/thread/lock\\_guard](https://en.cppreference.com/w/cpp/thread/lock_guard)

# STD::UNIQUE\_LOCK

Owns a mutex

Destructor calls unlock() on the mutex if the mutex is locked

Has lock(), unlock(), and several other member functions

Movable, but not copyable

Useful when you need more control over when the mutex is locked

[https://en.cppreference.com/w/cpp/thread/unique\\_lock](https://en.cppreference.com/w/cpp/thread/unique_lock)

# MUTEX AND LOCK GUARD EXAMPLE

Needs synchronization

```
if (transaction.amount < account.balance) {  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

# MUTEX AND LOCK GUARD EXAMPLE

Needs synchronization

```
if (transaction.amount < account.balance) {  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

Data race!



# MUTEX AND LOCK GUARD EXAMPLE

## Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

<https://godbolt.org/z/fnovzn8a>

# MUTEX AND LOCK GUARD EXAMPLE

## Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

<https://godbolt.org/z/fnovzn8a>

# MUTEX AND LOCK GUARD EXAMPLE

## Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
}  
else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

Lock guards must be  
named variables!





# MUTEX AND LOCK GUARD EXAMPLE

## Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

} Mutex locked

} Mutex locked

# MUTEX AND LOCK GUARD EXAMPLE

No data race

Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
}  
else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

} Mutex locked

} Mutex locked

# MUTEX AND LOCK GUARD EXAMPLE

## Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

Data race!



# MUTEX AND LOCK GUARD EXAMPLE

## Incomplete synchronization

```
if (transaction.amount < account.balance) {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= transaction.amount;  
    record(account, transaction);  
    accept_transaction(transaction);  
} else {  
    std::scoped_lock lock(account.mutex);  
    account.balance -= overdraft_charge;  
    record(account, overdraft_of(transaction));  
    notify_user(account, overdraft_message);  
    reject_transaction(transaction);  
}
```

Don't need locked mutex



<https://godbolt.org/z/fnovzn8a>

# MUTEX AND LOCK GUARD EXAMPLE

```
bool transaction_ok;
{ std::scoped_lock lock(account.mutex);
  transaction_ok = transaction.amount < account.balance;
  if (transaction_ok) {
    account.balance -= transaction.amount;
    record(account, transaction);
  } else {
    account.balance -= overdraft_charge;
    record(account, overdraft_of(transaction));
    notify_user(account, overdraft_message);
  }
}
if (transaction_ok) {
  accept_transaction(transaction);
} else {
  reject_transaction(transaction);
}
```

<https://godbolt.org/z/aPjejjWxP>

# MUTEX AND LOCK GUARD EXAMPLE

```
bool transaction_ok;
{ std::scoped_lock lock(account.mutex);
  transaction_ok = transaction.amount < account.balance;
  if (transaction_ok) {
    account.balance -= transaction.amount;
    record(account, transaction);
  } else {
    account.balance -= overdraft_charge;
    record(account, overdraft_of(transaction));
    notify_user(account, overdraft_message);
  }
}
if (transaction_ok) {
  accept_transaction(transaction);
} else {
  reject_transaction(transaction);
}
```

<https://godbolt.org/z/aPjejjWxP>

# MUTEX AND LOCK GUARD EXAMPLE

```
bool transaction_ok;
{ std::scoped_lock lock(account.mutex);
  transaction_ok = transaction.amount < account.balance;
  if (transaction_ok) {
    account.balance -= transaction.amount;
    record(account, transaction);
  } else {
    account.balance -= overdraft_charge;
    record(account, overdraft_of(transaction));
    notify_user(account, overdraft_message);
  }
}
if (transaction_ok) {
  accept_transaction(transaction);
} else {
  reject_transaction(transaction);
}
```

<https://godbolt.org/z/aPjejjWxP>

# MUTEX AND LOCK GUARD EXAMPLE

```
bool transaction_ok;
{ std::scoped_lock lock(account.mutex);
  transaction_ok = transaction.amount < account.balance;
  if (transaction_ok) {
    account.balance -= transaction.amount;
    record(account, transaction);
  } else {
    account.balance -= overdraft_charge;
    record(account, overdraft_of(transaction));
    notify_user(account, overdraft_message);
  }
}
if (transaction_ok) {
  accept_transaction(transaction);
} else {
  reject_transaction(transaction);
}
```

<https://godbolt.org/z/aPjejjWxP>



# MUTEX GOTCHAS

## Deadlock

When multiple mutexes are locked at the same time, they must always be locked in the same order

If not, deadlock will happen

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l2(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l1(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l2(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l1(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1a(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l1b(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l2b(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2a(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l2(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l1(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l2(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l1(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l2(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l1(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l2(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l1(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```



# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1a(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l1b(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l2b(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2a(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l_a1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l_b1(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l_b2(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l_a2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1a1(mutex_a);  
  change_data(data_a);  
  { std::scoped_lock l1b1(mutex_b);  
    change_data(data_a, data_b);  
  }  
}
```

### Thread 2

```
{ std::scoped_lock l2b2(mutex_b);  
  change_data(data_b);  
  { std::scoped_lock l2a2(mutex_a);  
    change_data(data_a, data_b);  
  }  
}
```

# MUTEX GOTCHAS

## Deadlock

`std::scoped_lock` is useful for avoiding deadlock

If given multiple mutexes, always locks them in the same order

```
std::scoped_lock lockA{mutexA, mutexB, mutexC};
```

```
std::scoped_lock lockB(mutexC, mutexA, mutexB);
```

No deadlock

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(  
    mutex_a, mutex_b);  
  change_data(data_a);  
  change_data(data_a, data_b);  
}
```

### Thread 2

```
{ std::scoped_lock l2(  
    mutex_b, mutex_a);  
  change_data(data_b);  
  change_data(data_a, data_b);  
}
```

# MUTEX GOTCHAS

## Deadlock

### Thread 1

```
{ std::scoped_lock l1(  
    mutex_a, mutex_b);  
  change_data(data_a);  
  change_data(data_a, data_b);  
}
```

### Thread 2

```
{ std::scoped_lock l2(  
    mutex_b, mutex_a);  
  change_data(data_b);  
  change_data(data_a, data_b);  
}
```

# STD::ATOMIC

# STD::ATOMIC

Atomic operations happen as a unit

Things can't change in the middle

```
std::atomic<int> x;
```

Thread 1  
x += 5;

Thread 2  
x += 7;



# STD::ATOMIC

Atomic operations happen as a unit

Things can't change in the middle

```
std::atomic<int> x;
```

Thread 1  
x += 5;

Thread 2  
x += 7;

No interference.  
No data race.

# STD::ATOMIC

## Synchronization

Synchronizes between threads that access the same atomic object

```
std::atomic<int> x;
```

### Thread 1

```
// Section 1A  
x.store(5);  
  
// Section 1B  
  
// Section 1C
```

### Thread 2

```
// Section 2A  
  
// Section 2B  
y = x.load();  
  
// Section 2C
```

# STD::ATOMIC

## Synchronization

Synchronizes between threads that access the same atomic object

```
std::atomic<int> x;
```

### Thread 1

```
// Section 1A
```

```
x.store(5);
```

```
// Section 1B
```

```
// Section 1C
```

### Thread 2

```
// Section 2A
```

```
// Section 2B
```

```
y = x.load();
```

```
// Section 2C
```

Synchronization



# STD::ATOMIC

## Synchronization

Synchronizes between threads that access the same atomic object


```
std::atomic<int> x;
```

### Thread 1

```
// Section 1A  
x.store(5);  
// Section 1B  
// Section 1C
```

### Thread 2

```
// Section 2A  
// Section 2B  
y = x.load();  
// Section 2C
```



# DATA RACES

## Examples

# DATA RACES

Loop example - has data race

```
bool flag = false;
int main() {
    std::thread t([](){
        std::printf("waiting...\n");
        while (not flag) { }
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    flag = true;
    t.join();
    std::printf("Done\n");
}
```

<https://godbolt.org/z/Gj8dsq9n6>

# DATA RACES

Loop example - has data race

```
bool flag = false;
int main() {
    std::thread t([]() {
        std::printf("waiting...\n");
        while (not flag) { }
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    flag = true;
    t.join();
    std::printf("Done\n");
}
```

Read

No synchronization

Write

<https://godbolt.org/z/Gj8dsq9n6>

# DATA RACES

Loop example - has data race

```
bool flag = false;
int main() {
    std::thread t([]() {
        std::printf("waiting...\n");
        while (not flag) { }
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    flag = true;
    t.join();
    std::printf("Done\n");
}
```

Compiler assumes  
infinite loop

<https://godbolt.org/z/Gj8dsq9n6>



# DATA RACES

## Loop example - deadlock

```
bool flag = false;
std::mutex flag_mutex;
int main() {
    std::thread t([](){
        std::printf("Waiting...\n");
        { std::scoped_lock lock(flag_mutex);
            while (not flag) { }
        }
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    { std::scoped_lock lock(flag_mutex);
        flag = true;
    }
    t.join();
    std::printf("Done\n");
}
```

# DATA RACES

## Loop example - deadlock

```
bool flag = false;
std::mutex flag_mutex;
int main() {
    std::thread t([](){
        std::printf("Waiting...\n");
        { std::scoped_lock lock(flag_mutex);
            while (not flag) { }
        }
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    { std::scoped_lock lock(flag_mutex);
        flag = true;
    }
    t.join();
    std::printf("Done\n");
}
```

Holds lock indefinitely

Can't acquire lock

# DATA RACES

## Loop example - mutex

```
bool flag = false;
std::mutex flag_mutex;
int main() {
    std::thread t([](){
        std::printf("Waiting...\n");
        bool local_flag;
        do {
            std::scoped_lock lock{flag_mutex};
            local_flag = flag;
        } while (not local_flag);
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    { std::scoped_lock lock(flag_mutex);
      flag = true;
    }
    t.join();    std::printf("Done\n");
}
```

<https://godbolt.org/z/b5ProbeMa>

# DATA RACES

## Loop example - mutex

```
bool flag = false;
std::mutex flag_mutex;
int main() {
    std::thread t([](){
        std::printf("waiting...\n");
        bool local_flag;
        do {
            std::scoped_lock lock{flag_mutex};
            local_flag = flag;
        } while (not local_flag);
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    { std::scoped_lock lock(flag_mutex);
      flag = true;
    }
    t.join();    std::printf("Done\n");
}
```

<https://godbolt.org/z/b5ProbeMa>

# DATA RACES

## Loop example - mutex

```
bool flag = false;
std::mutex flag_mutex;
int main() {
    std::thread t([](){
        std::printf("waiting...\n");
        bool local_flag;
        do {
            std::scoped_lock lock{flag_mutex};
            local_flag = flag;
        } while (not local_flag);
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    { std::scoped_lock lock(flag_mutex);
      flag = true;
    }
    t.join();    std::printf("Done\n");
}
```

<https://godbolt.org/z/b5ProbeMa>

# DATA RACES

## Loop example - atomic

```
std::atomic<bool> flag{false};
int main() {
    std::thread t([](){
        std::printf("waiting...\n");
        while (not flag) { }
        std::printf("Flag changed\n");
    });
    std::this_thread::sleep_for(500ms);
    std::printf("Setting flag\n");
    flag = true;
    t.join();
    std::printf("Done\n");
}
```

<https://godbolt.org/z/f4Tv967KP>

# DATA RACES

Increment example - has data race

```
int counter = 0;
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = []() {
        for (int i = 0; i < 50; ++i) {
            ++counter;
        }
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter);
}
```

# DATA RACES

Increment example - has data race

```
int counter = 0;
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = []() {
        for (int i = 0; i < 50; ++i) {
            ++counter;
        }
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter);
}
```



# DATA RACES

Increment example - has data race

```
int counter = 0;
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = []() {
        for (int i = 0; i < 50; ++i) {
            ++counter;
        }
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter);
}
```

# DATA RACES

Increment example - has data race

```
int counter = 0;
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = []() {
        for (int i = 0; i < 50; ++i) {
            ++counter;
        }
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter);
}
```

# DATA RACES

## Increment example - mutex

```
int counter = 0;
std::mutex counter_mutex;
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = [](){
        for (int i = 0; i < 50; ++i) {
            std::scoped_lock lock{counter_mutex};
            ++counter;
        }
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter);
}
```

# DATA RACES

## Increment example - atomic

```
std::atomic<int> counter{0};
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = [](){
        for (int i = 0; i < 50; ++i) {
            ++counter;
        }
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter.load());
}
```

# DATA RACES

Increment example - fewer updates

```
std::atomic<int> counter{0};
int main() {
    std::vector<std::thread> threads;
    threads.reserve(100);
    auto increment = [](){
        int local_counter = 0;
        for (int i = 0; i < 50; ++i) {
            ++local_counter;
        }
        counter += local_counter;
    };
    for (int i = 0; i < 100; ++i) {
        threads.push_back(std::thread(increment));
    }
    for (std::thread& t : threads) { t.join(); }
    std::printf("%d\n", counter.load());
}
```



# STD::RECURSIVE\_MUTEX

Same thread can lock the mutex multiple times

Number of calls to unlock() must match number of calls to lock()

Useful when some thread-safe APIs call other thread-safe APIs

[https://en.cppreference.com/w/cpp/thread/recursive\\_mutex](https://en.cppreference.com/w/cpp/thread/recursive_mutex)

## STD::TIMED\_MUTEX

Additional locking functions:

`try_lock_for(duration)`

`try_lock_until(time)`

[https://en.cppreference.com/w/cpp/thread/timed\\_mutex](https://en.cppreference.com/w/cpp/thread/timed_mutex)

# STD::RECURSIVE\_TIMED\_MUTEX

`recursive_mutex + timed_mutex`

[https://en.cppreference.com/w/cpp/thread/recursive\\_timed\\_mutex](https://en.cppreference.com/w/cpp/thread/recursive_timed_mutex)



# STD::SHARED\_MUTEX

a.k.a. read-write mutex

Many threads can get shared ownership, or read access

```
sm.lock_shared();
```

Calls to `lock_shared()` from other threads will succeed; calls to `lock()` will block

Only one thread can have exclusive ownership, or write access

```
sm.lock();
```

Calls to `lock_shared()` or `lock()` from other threads will block

Useful when there are lots of readers, but few writers

[https://en.cppreference.com/w/cpp/thread/shared\\_mutex](https://en.cppreference.com/w/cpp/thread/shared_mutex)

## STD::SHARED\_TIMED\_MUTEX

shared\_mutex with additional operations:

`try_lock_for(duration)`

`try_lock_until(time)`

`try_lock_shared_for(duration)`

`try_lock_shared_until(time)`

[https://en.cppreference.com/w/cpp/thread/shared\\_timed\\_mutex](https://en.cppreference.com/w/cpp/thread/shared_timed_mutex)

# STD::SHARED\_LOCK

Constructor takes one shared\_mutex

Calls lock\_shared() on the shared\_mutex

Destructor calls unlock\_shared() on the shared\_mutex

API is similar to unique\_lock

For shared ownership (read access) of a shared mutex, use shared\_lock

For exclusive ownership (write access) of a shared mutex, use scoped\_lock

# STD::CONDITION\_VARIABLE

Complicated to use correctly

Useful when some threads are waiting for a condition and other threads make that condition true

[https://en.cppreference.com/w/cpp/thread/condition\\_variable](https://en.cppreference.com/w/cpp/thread/condition_variable)

# STD::COUNTING\_SEMAPHORE

Lightweight synchronization primitive that can control access to a shared resource

Maintains an internal counter

Calls to `release()` increment the counter

Calls to `acquire()` decrement the counter, or block if counter == 0

Calls to `acquire()` and `release()` can happen on different threads

[https://en.cppreference.com/w/cpp/thread/counting\\_semaphore](https://en.cppreference.com/w/cpp/thread/counting_semaphore)

# STD::LATCH

Synchronize the completion of a shared task

Each thread calls `arrive_and_wait()`

Blocks until all threads have called `arrive_and_wait()`

<https://en.cppreference.com/w/cpp/thread/latch>

# STD::BARRIER

## Reusable latch

Once all threads have called `arrive_and_wait()` and been unblocked, the process starts over

<https://en.cppreference.com/w/cpp/thread/barrier>

# THINGS OTHER THAN THREADS

Processes - same system

Processes on the same system working together

Communicate through file system

How do they synchronize?



# THINGS OTHER THAN THREADS

Processes - different systems

Processes working together across a network

How do they communicate?

Which process has which information?

# THINGS OTHER THAN THREADS

## GPU threads

GPU threads have different synchronization rules than CPU threads

Synchronization rules depend on type of memory and relationship between threads

# CONCLUSION

Concurrency is hard

Use parallel algorithms when appropriate

Avoid data races at all costs

- Share less data

- Mutexes and lock guards

# STD::MUTEX

## Synchronization

### Thread 1

// Section 1A

m.lock();

// Section 1B

m.unlock();

// Section 1C

### Thread 2

// Section 2A

m.lock();

// Section 2B

m.unlock();

// Section 2C

Visible

Synchronization

# CONCLUSION

Concurrency is hard

Use parallel algorithms when appropriate

Avoid data races at all costs

- Share less data

- Mutexes and lock guards







