


Coroutine Patterns and How to Use Them:

Problems and Solutions Using Coroutines
In a Modern Codebase

FRANCESCO ZOFFOLI



About Me

- Software engineer building monitoring systems at  Meta
- Passionate about C++
- Author of the book “C++ Fundamentals” – Packt
- I like writing and talking about C++


makers.f.dev@gmail.com

Outline

- Motivation
- Overview
- Patterns
 - Lifetime
 - Exceptions
 - RAI
 - Synchronization
- Conclusions

Motivation

- 40 years of
 - Experience
 - Accumulated knowledge



Motivation

Coroutines introduce new paradigm



Motivation

- Build shared knowledge

Overview

- Folly - github.com/facebook/folly
- Most libraries are similar

Overview

Key concepts

- Task
- Executor

Overview – Task

- Owns the coroutine state
- Lazy
 - Convertible to eager
- Propagates executor
- Sticky to executor

Overview – Task - Lazy

```
Task<> foo() {  
    println("Hello");  
    co_await sleep(1);  
}
```

```
auto t = foo();  
println("world");  
co_await move(t);
```

Overview – Executor

- Executes synchronous functions
- Could be multi-threaded
- Coroutines are split in synchronous functions

Overview – Executor

```
Task<> foo() {  
    int a = 42;  
    a -= 10;  
    int bytes = co_await send(a);  
    if (bytes == -1 ) {  
        handle_error();  
    }  
}
```

Overview – Executor

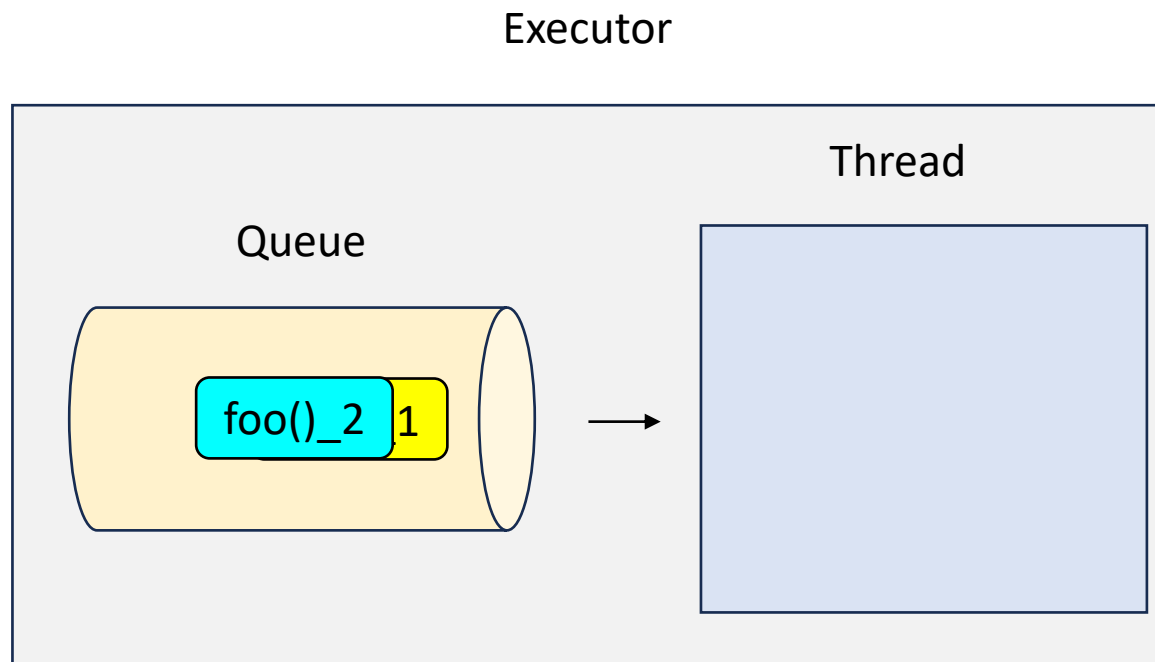
```
Task<> foo() {  
    int a = 42;  
    a -= 10;  
    int bytes = co_await send(a);  
    if (bytes == -1) {  
        handle_error();  
    }  
}
```

Executes until continuation

co_await

Executes after continuation

Overview - Executor



Patterns

- Lifetime
- Exceptions
- RAII
- Synchronization

Lifetime

Lifetime

- Lots of worries
- Use Structured Concurrency^[1]
- ASAN works great

[1] <https://ericniebler.com/2020/11/08/structured-concurrency/>

Lifetime

Member coroutines implicitly capture this

Lifetime

Let's play a game

Lifetime

```
struct Bar {  
    int data = 0;  
  
    Task<int> foo() {  
        co_await sleep(1);  
        co_return data;  
    }  
};
```

```
Task<int> mul_2(Task<int> t){  
    int val = co_await move(t);  
    co_return 2 * val;  
}
```

Lifetime

Let's play a game



Lifetime

```
std::vector<int> vec;  
vec.front() += 1;
```

Lifetime

```
std::vector<int> vec;  
vec.clear();
```

Lifetime

```
Bar a;  
co_await a.foo();
```

```
Bar a;  
Task<int> t = a.foo();  
co_await move(t);
```



Lifetime

```
Task<int> t = Bar{}.foo();  
co_await move(t);
```

```
co_await Bar{}.f();
```



Lifetime

```
Bar a;
```

```
Task<int> t = mul_2(a.foo());
```

```
co_await move(t);
```

```
Task<int> t = mul_2(Bar{}.foo());
```

```
co_await move(t);
```

```
co_await mul_2(A{}.foo());
```



Lifetime – Shared knowledge

- Keep object alive until all methods have completed

Lifetime

When does this concretely happens?

Lifetime

When does this concretely happens?

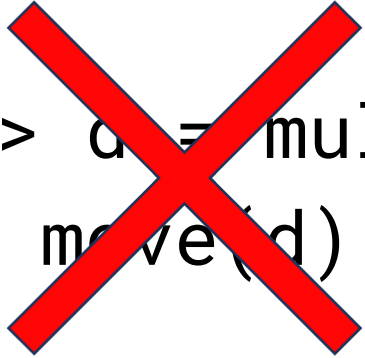
In loops!

```
vector<Task<>> tasks;  
for(int index = ...) {  
    Foo foo{index};  
    tasks.push_back(foo.bar());  
}  
co_await collectAll(tasks);
```

Lifetime

When does this concretely happens?

Lambdas!!



```
Task<int> d = mul_2([=]() { co_return i; }());  
co_await move(d);
```

Lifetime

When does this concretely happens?

Lambdas!!

```
co_await mul_2([=]() { co_return i; }());
```

```
auto lam = [=]() { co_return i; };  
co_await mul_2(lam());
```

Lifetime

Use `co_invoke`!

```
Task<int> d = mul_2(co_invoke([=]() { co_return i; }));  
co_await move(d);
```


Lifetime – Shared knowledge

- Keep object alive until all methods have completed
- Ensure lambda objects are kept alive (using `co_invoke`)

Lifetime

Nothing new!

Think: task reference this pointer

```
vector<tuple<int&>> vec;  
for(int index = ...) {  
    vec.emplace_back(index);  
}
```

Who wouldn't spot this?

Lifetime

Scheduling in the background

Lifetime

```
{  
    Foo a;  
    startInBackground(a.bar());  
    co_await do_other_stuff();  
}
```

Big no no!

Need to join background work

Lifetime

Wait for all work started in the same scope

```
{  
    Foo a;  
    auto bgWork = startInBackground(a.bar());  
    co_await do_other_stuff();  
    co_await move(bgWork);  
}
```

Lifetime

Waiting for many background Tasks – AsyncScope

```
Foo a;  
AsyncScope s;  
for(int index = ...) {  
    co_await s.schedule(a.bar(index));  
}  
co_await do_other_stuff();  
co_await s.join();
```

Lifetime - Shared knowledge

- Keep object alive until all methods have completed
- Ensure lambda objects are kept alive (using `co_invoke`)
- Always join work before leaving the scope

Lifetime

Anyone sees a problem with the code above

What if `do_other_stuff()` throws?

Still need to join!

Exceptions

Exceptions

Await in catch

```
try {  
    co_await do_other_stuff();  
} catch (...) {  
    co_await asyncScope.join();  
}
```

Compiler error!

Not allowed to co_await inside a catch block

Exceptions

Capture the exception

```
exception_ptr eptr;  
try {  
    co_await do_other_stuff();  
} catch (...) {  
    eptr = current_exception();  
}  
if (eptr) {  
    co_await asyncScope.join();  
    rethrow_exception(eptr);  
}
```

Exceptions

Capture the exception (2)

```
auto res = co_await co_awaitTry(do_other_stuff());  
if (res.hasException()) {  
    co_await asyncScope.join();  
}  
// Rethrows if it doesn't have a value  
res.value();
```

Exceptions - Shared knowledge

- To do async work in catch block, capture the exception in a wrapper
 - Use libraries to help

RAII

Different solutions for

- Class hierarchies and members
- Automatic variables

RAII – Classes

Async cleanup pattern

- Define async cleanup method
- Parents/owners call the method, recursively

RAII – Classes

```
class Foo {  
    Task<> cleanup() {  
        co_await collectAll(as.join(), m1.cleanup());  
    }  
  
    Bar m1;  
    AsyncScope as;  
};
```


RAII – Classes

- `cleanup()` shouldn't throw
- Tip: assert in the destructor that `cleanup()` was called

A bit of manual work 😞

RAII – Scopes

Need to ensure `cleanup()` is called

RAII – Scopes

- Avoid exceptions
- Testing

RAII - Shared Knowledge

- Use cleanup pattern to join async work in classes

RAII

“This is bad! I’ll use `blockingWait()` in the destructor”

```
~MyClass() {  
    blockingWait(collectAll(as.join(), m1.cleanup()));  
};
```

RAII


DON'T!!



RAII

Why?

1. `blockingWait()` blocks a thread
2. continuation is scheduled
3. no thread available to run it



RAII

```
struct Foo {  
    Task<> do_work() {  
        co_await  
        as.schedule(sleep(1s));  
    }  
  
    Task<> cleanup() {  
        co_await as.joinAsync();  
    }  
};  
  
~Foo() {  
    blockingWait(cleanup());  
}  
  
AsyncScope as;  
};
```


RAII – Shared Knowledge

- Use cleanup pattern to join async work in classes
- Don't use `blockingWait()` in a destructor

Synchronization

Synchronization

Needed?

Synchronization

No shared access? → No

```
int a = 10;  
co_await reschedule_on_current_executor;  
// Resume in different thread  
do_stuff(a);
```

Synchronization

Shared Access? Depends

- ST executor → No
- MT executor → Yes

```
int a = 10;  
co_await collectAll(  
    increaseBy(a, 1),  
    increaseBy(a, 2));
```

Synchronization

Which tools?

- Regular mutexes
- Coro mutex

Synchronization – Regular

DON'T HOLD ACROSS SUSPENSION POINTS!!!

- UB if held across suspension points
- Only for small sections
- Useful for advanced features/non-coro code

Synchronization – Regular

Use lambda format

```
synchronized.withWLock( [&](auto&) {  
    ...  
} );
```

- `co_await` → compile error
- guard patten is risky

Synchronization – Shared Knowledge

- Do not hold regular lock across suspension points

Synchronization – Coro

Support RAI!

Synchronization – Coro

Access mutex protected data in destructor

- `try_lock()`

Synchronization – Shared Knowledge

- Do not hold regular lock across suspension points
- Use coro mutexes, including guards
- Use `try_lock()` to for coro mutexes in destructor

Conclusions

- Saw patterns to do and avoid
- Build shared knowledge
- Invite people to share experiences
- Identify unsolved needs

Thank you