

23

# Tracy: A Profiler You Don't Want to Miss

MARCOS SLOMP



**Cppcon**  
The C++ Conference

20  
23



October 01 - 06

# Tracy: a profiler you don't want to miss

Marcos Slomp

Senior Research Engineer II



**Cppcon**

The C++ Conference

2023   
October 01-06  
Aurora, Colorado, USA



RESEARCH ENGINEERING & DESIGN

Adobe

# Why Tracy?

**Real-time** workflow  
(analyze application **while it is running**)

**Precise** measurements  
(**nanosecond** resolution)

**Negligible overhead**  
(a few **nanoseconds** per zone)

**Cross-platform**  
(Windows, Linux, macOS, iOS, Android, WASM\*)

**Hybrid profiling** capabilities  
(**sampling** and/or **instrumentation**)  
(CPU and GPU instrumentation)

**Tracing** capabilities  
(values, messages, plots, allocations, ...)

**Hassle-free integration**  
(single source\* file, and header\*)

**Free and Open source**



**PAY WHAT YOU WANT**

**\$0.00**

Exc. Internet fees  
Inc. source code

**1-555-PROFILE**

**FREE Trial! Instant Delivery!**  
**Lifetime Money Back Guarantee (less S&H).**

**<https://github.com/wolfpld/tracy>**

# Why Tracy?

**Real-time** workflow  
(analyze application **while it is running**)

**Precise** measurements  
(**nanosecond** resolution)

**Negligible** overhead  
(a few **nanoseconds** per zone)

**Cross-platform**  
(Windows, Linux, macOS, iOS, Android, WASM\*)

**Hybrid profiling** capabilities  
(**sampling** and/or **instrumentation**)

**Tracing** capabilities  
(values, messages, plots, allocations, ...)

**Free** and **Open source**



**PAY WHAT YOU WANT**

**\$0.00**

Exc. Internet fees  
Inc. source code

**1-555-PROFILE**

**FREE Trial! Instant Delivery!**  
**Lifetime Money Back Guarantee (less S&H).**  
**<https://github.com/wolfpld/tracy>**

# Why Tracy?

**Real-time** workflow  
(analyze application **while it is running**)

**Precise** measurements  
(**nanosecond** resolution)

**Negligible** overhead  
(a few **nanoseconds** per zone)

**Cross-platform**  
(Windows, Linux, macOS, iOS, Android, WASM\*)

**Hybrid profiling** capabilities  
(**sampling** and/or **instrumentation**)

**Tracing** capabilities  
(values, messages, plots, allocations, ...)

**Free** and **Open source**



**PAY WHAT YOU WANT**

**\$0.00**

Exc. Internet fees  
Inc. source code

**1-555-PROFILE**

**FREE Trial! Instant Delivery!**  
**Lifetime Money Back Guarantee (less S&H).**  
**<https://github.com/wolfpld/tracy>**

# DISCLAIMER: this talk is **NOT** about “Tracy vs. other tools”

Telemetry	<a href="http://www.radgametools.com/telemetry.htm">http://www.radgametools.com/telemetry.htm</a>
Microprofile	<a href="https://github.com/jonasmr/microprofile">https://github.com/jonasmr/microprofile</a>
Microprofile 2	<a href="https://github.com/zeux/microprofile">https://github.com/zeux/microprofile</a>
Optick (Brofiler)	<a href="https://github.com/bombomby/optick">https://github.com/bombomby/optick</a>
Palanteer	<a href="https://github.com/dfeneyrou/palanteer">https://github.com/dfeneyrou/palanteer</a>
Orbit Profiler	<a href="https://github.com/google/orbit">https://github.com/google/orbit</a>
SuperLuminal	<a href="https://superluminal.eu">https://superluminal.eu</a>
Perfetto	<a href="https://perfetto.dev">https://perfetto.dev</a>
MicroProfiler	<a href="https://github.com/tyoma/micro-profiler">https://github.com/tyoma/micro-profiler</a>
EasyProfiler	<a href="https://github.com/yse/easy_profiler">https://github.com/yse/easy_profiler</a>
Oprofile	<a href="https://oprofile.sourceforge.io/examples/">https://oprofile.sourceforge.io/examples/</a>
Coz	<a href="https://github.com/plasma-umass/coz">https://github.com/plasma-umass/coz</a>
VerySleepy	<a href="https://github.com/VerySleepy/verysleepy">https://github.com/VerySleepy/verysleepy</a>
LukeStackwalker	<a href="http://lukestackwalker.sourceforge.net">http://lukestackwalker.sourceforge.net</a>
dprofiler	<a href="https://github.com/xwlan/dprofiler">https://github.com/xwlan/dprofiler</a>
hotspot	<a href="https://github.com/KDAB/hotspot">https://github.com/KDAB/hotspot</a>
heaptrack	<a href="https://github.com/KDE/heaptrack">https://github.com/KDE/heaptrack</a>
spall	<a href="https://handmade.network/p/333/spall/">https://handmade.network/p/333/spall/</a>
geiger	<a href="https://github.com/david-grs/geiger">https://github.com/david-grs/geiger</a>
Intel IACA	<a href="https://www.intel.com/content/www/us/en/developer/articles/tool/architecture-code-analyzer.html">https://www.intel.com/content/www/us/en/developer/articles/tool/architecture-code-analyzer.html</a>

# The Obligatory Quote

**premature optimization is the root of all evil.**

*Donald E. Knuth*

# The Obligatory Quote, in proper context

“There is no doubt that the grail of efficiency leads to abuse. Programmers **waste enormous amounts of time** thinking about, or **worrying** about, the **speed of noncritical parts** of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We *should forget about small efficiencies*, say about 97% of the time: premature optimization is the root of all evil. Yet we **should not pass up our opportunities in that critical 3%**.”

*Donald E. Knuth. 1974. Structured Programming with go to Statements. ACM Computing Surveys 6, 4 (Dec. 1974), 261–301.*



Mathieu Ropert

20  
21 |   
October 24-29

<https://www.youtube.com/watch?v=dToaepIXW4s>

*“The real problem is that programmers have spent far too much time worrying about efficiency in the wrong places and at the wrong times”*

 “

“C++ programmers talk all the time about **efficiency**,  
but they generally don’t know how to **measure**.”

- Bjarne Stroustrup  
CppCon 2023 Keynote  
(literally three days ago!)

# Closing the Profiling Zoo

## Tracing

program execution logging

- function calls (+ args & return)
- system events
- state changes
- debug/warning/error messages

**mostly for debugging: not necessarily concerned with performance**

## Sampling

callstack snapshot probing

- system-wide
- automatic (regular intervals)
- OS kernel events
- hardware counters

**issues: information overload; sampling overhead; constrained by sampling frequency**

## Instrumentation

code annotations (markup)

**holistic (automatic)**  
**issues: information overload; profiling overhead (skew)**

**selective (manual)**

## Analysis

pre-recorded? | real-time?  
log files? | command-line tool? | GUI?, ...

An **interactive**, **responsive** and **visual** tool transforms the profiling **experience!**

# Closing the Profiling Zoo



Tracy Profiler GUI



## Frame Info

# GPU Timeline (per “device”)

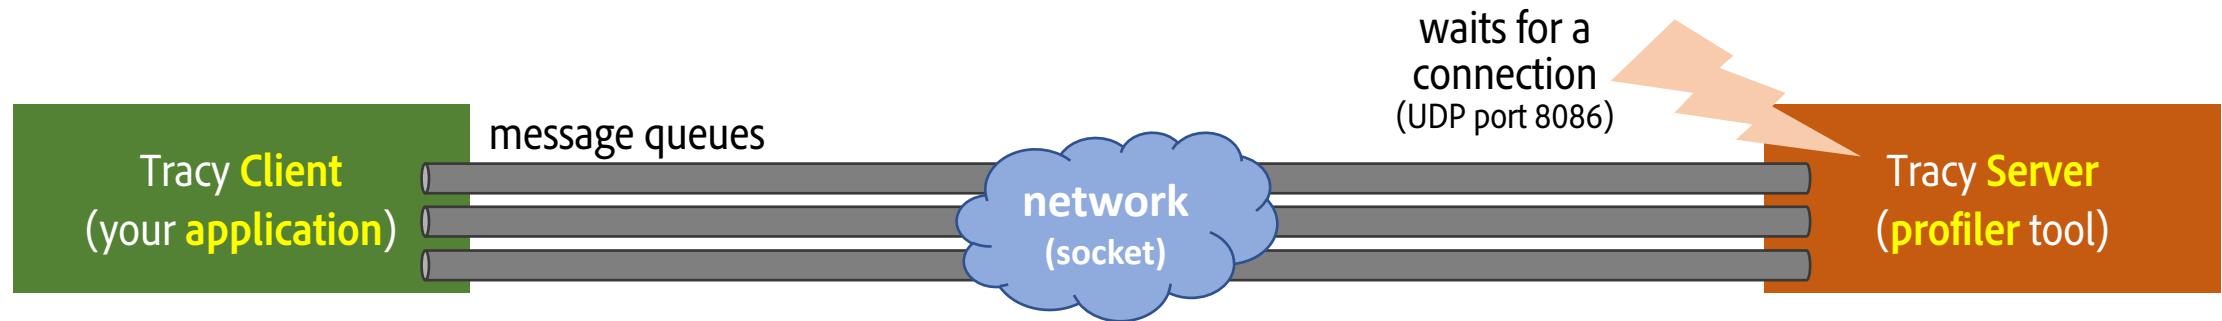


## CPU Timeline (per-thread)

# Allocation Trackers



# Tracy Architecture



Basic integration (the *responsible* way):

1. simply `#include <Tracy.hpp>` whenever needed
2. compile `TracyClient.cpp` alongside your application
3. enable Tracy by defining `TRACY_ENABLE` in your build

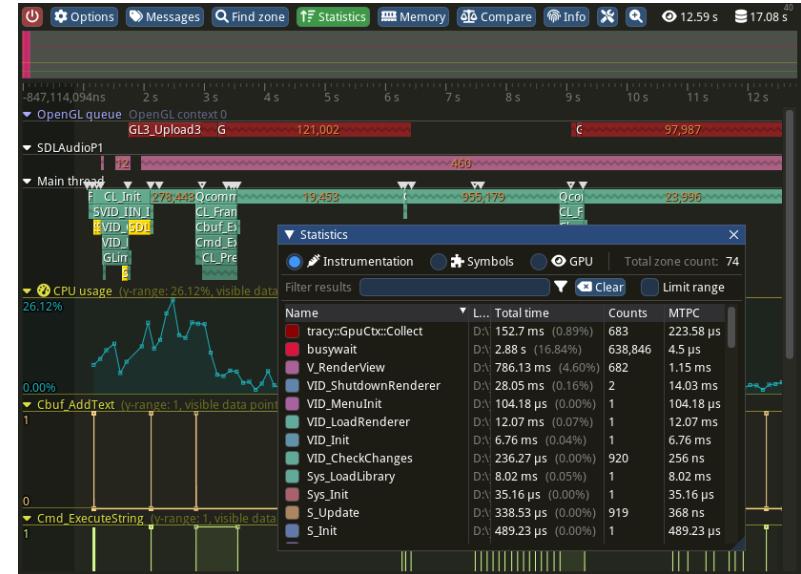
Quick integration (the *YOLO* way):

```
#define TRACY_ENABLE  
#include <Tracy.hpp>
```

anywhere where needed

```
#define TRACY_ENABLE  
#include <TracyClient.cpp>
```

in *one* of the source files of your project

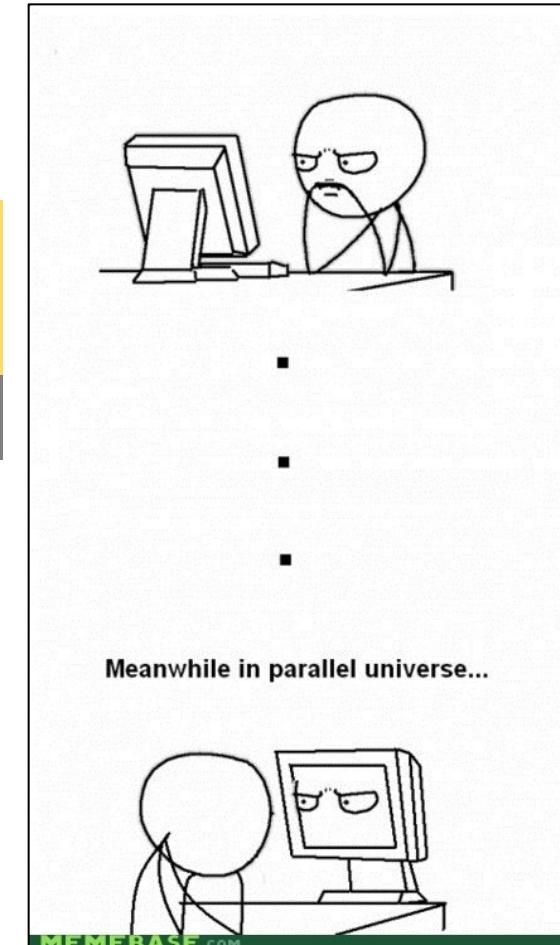
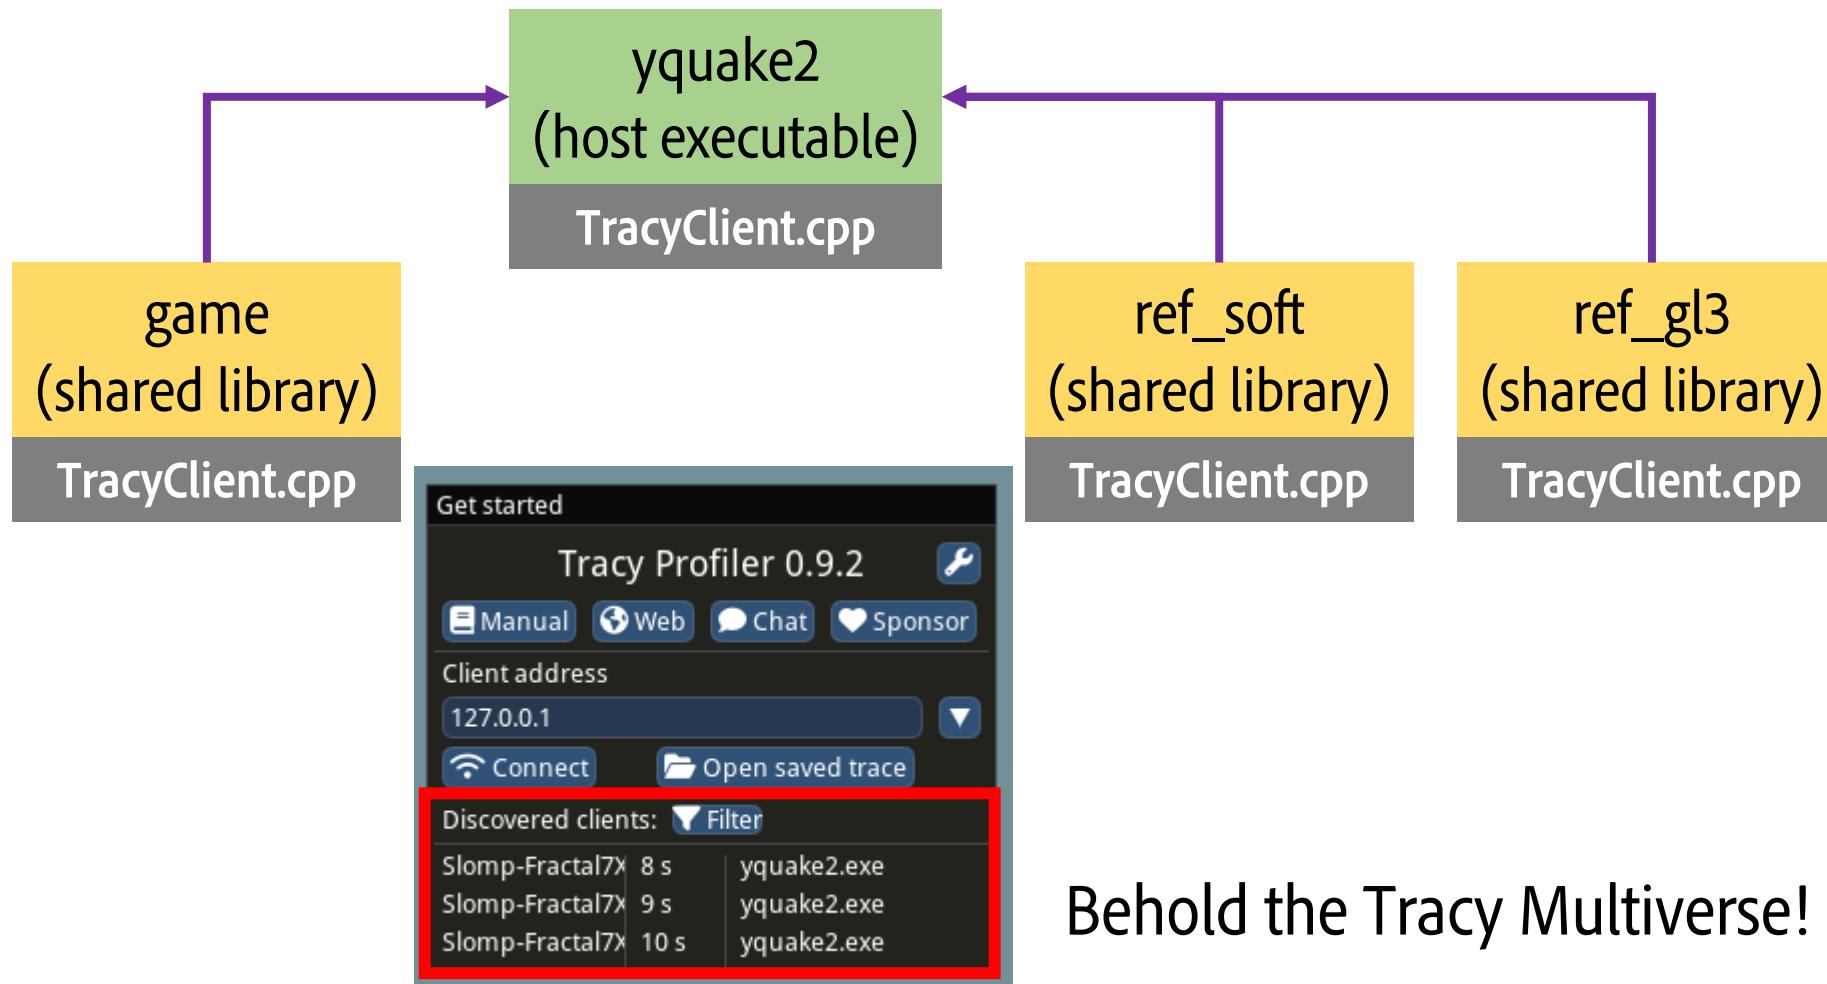


download pre-built binary (**Windows-only**):

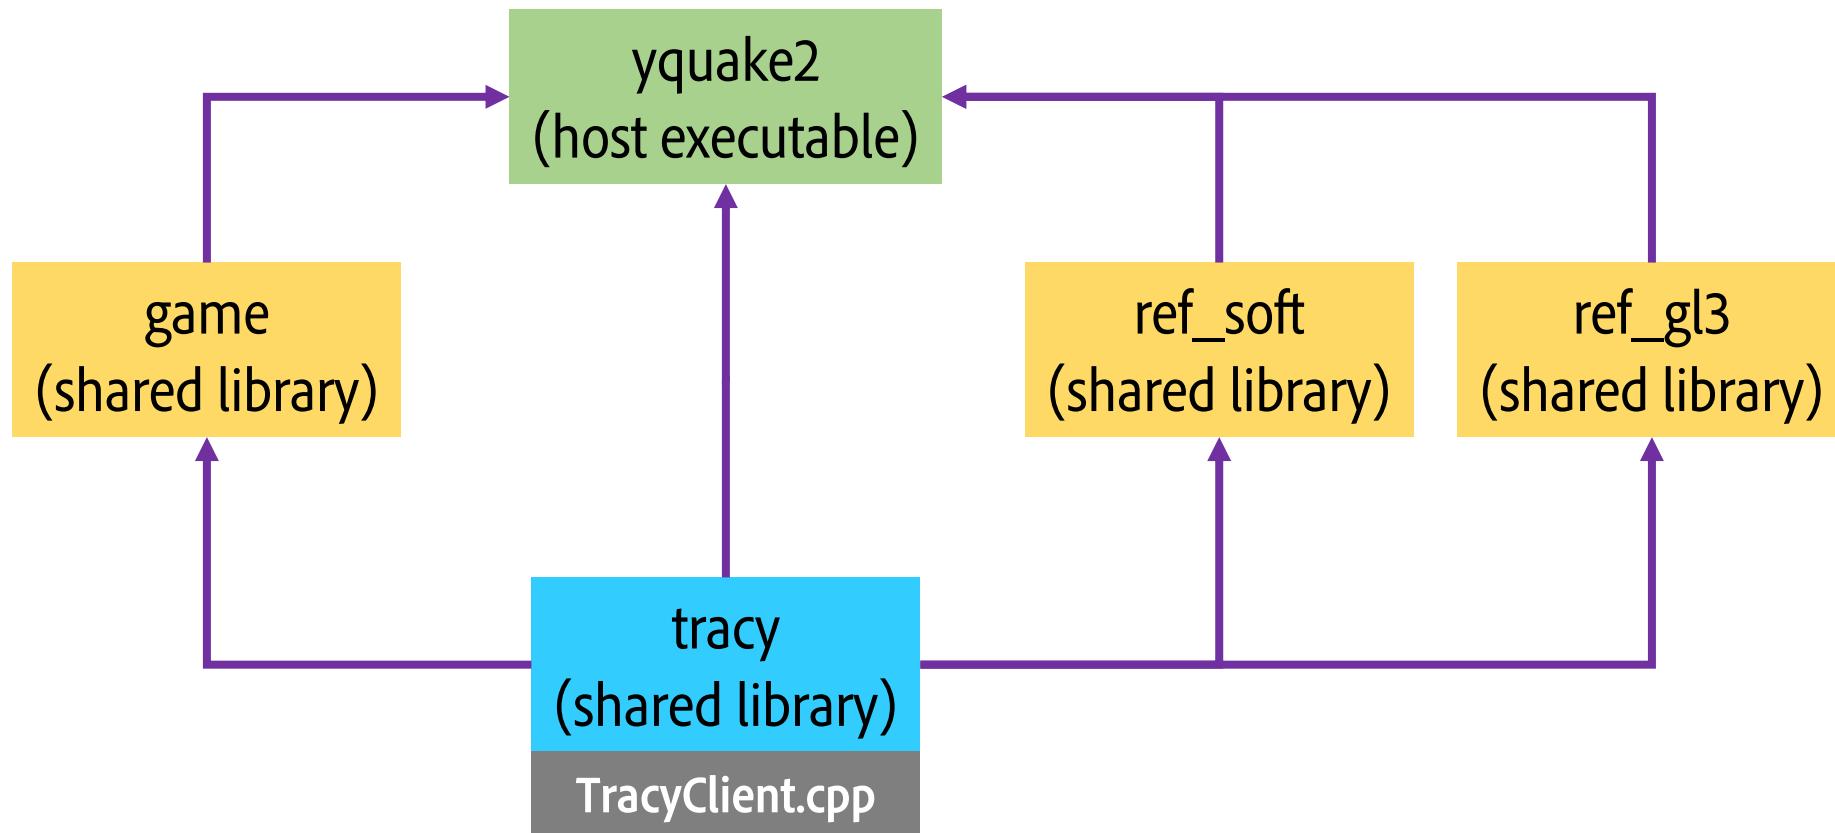
<https://github.com/wolfpld/tracy/releases>

or **build it yourself!**

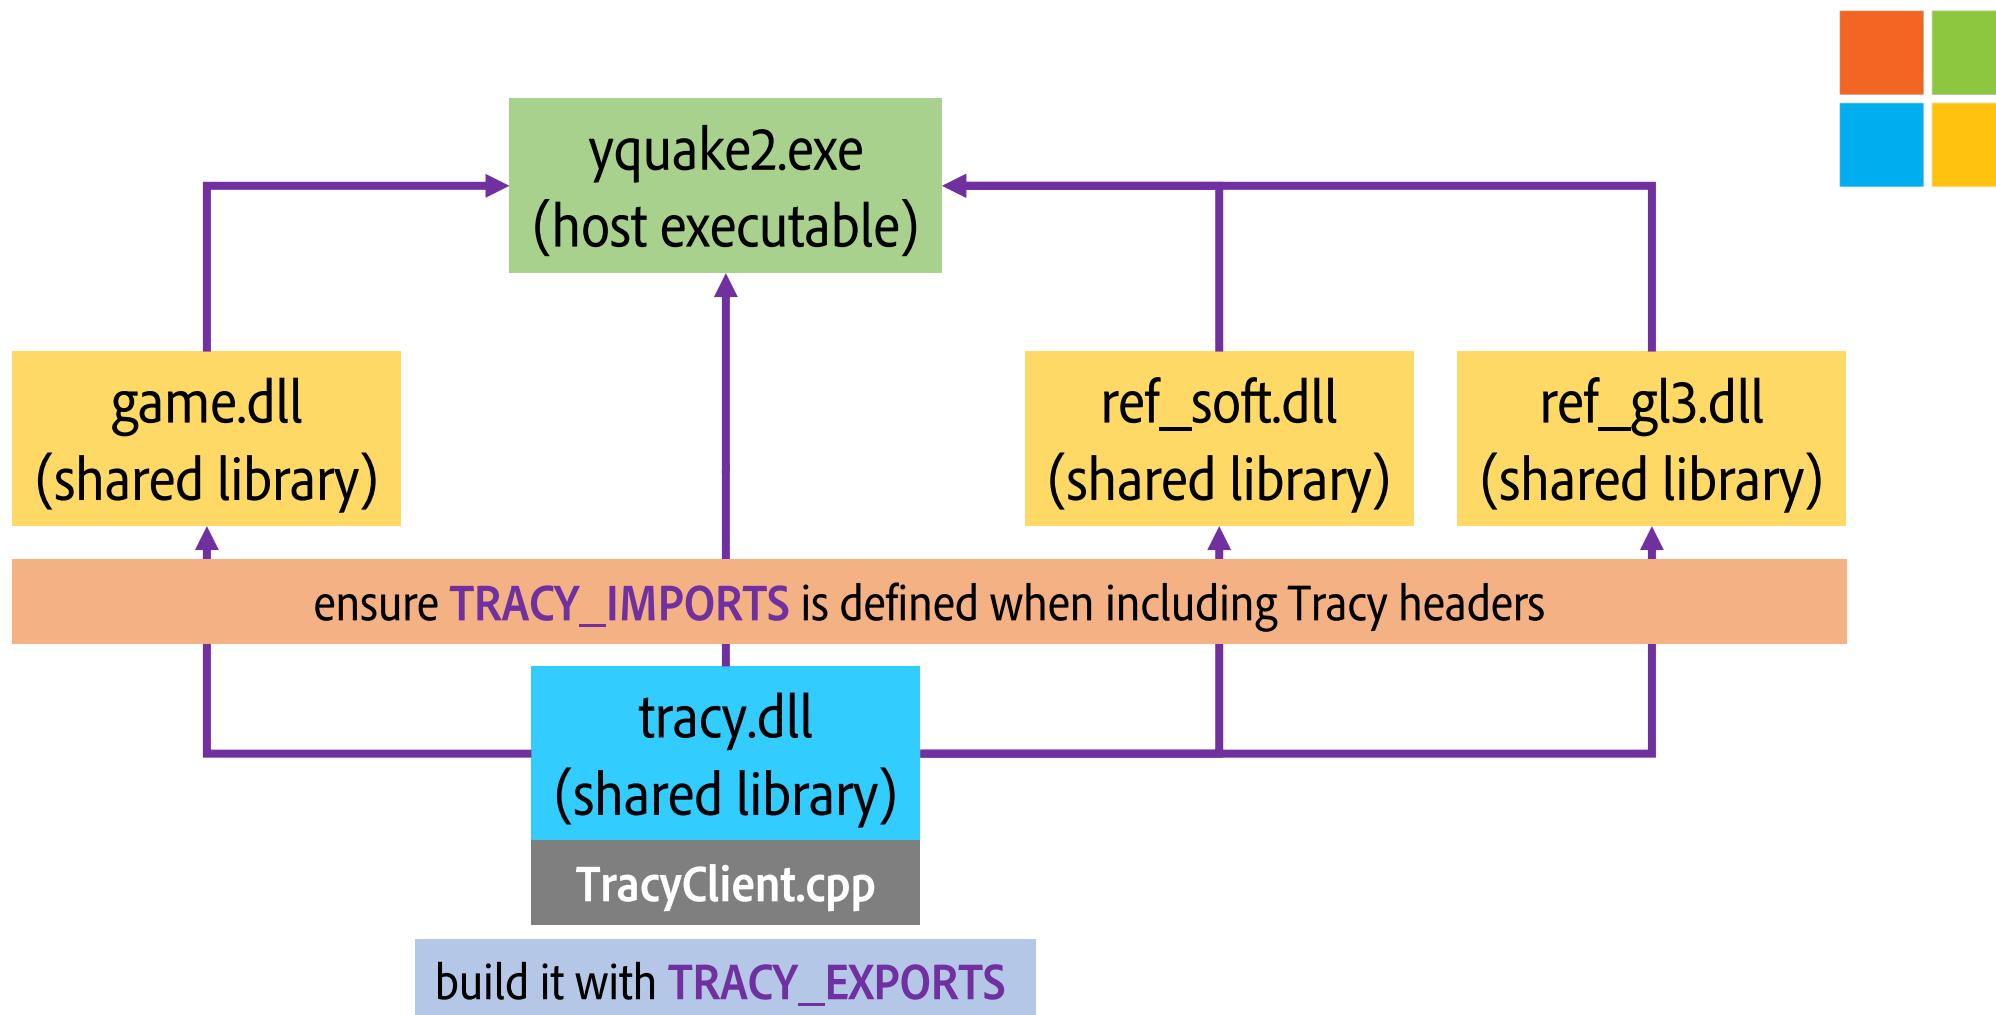
# Tracy integration across shared libraries



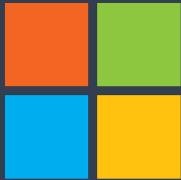
# Tracy integration across shared libraries



# Tracy integration across shared libraries (**Windows**)



# Building the Tracy Profiler from source



1. Obtain the **dependencies** (`capstone`, `glfw`, `freetype`):  
`<tracy-repo>\vcpkg\install_vcpkg_dependencies.bat`

2. Build with the **Visual Studio solution**:  
`<tracy-repo>\profiler\build\win32\Tracy.sln`

**Tracy.exe**



1. Obtain the **dependencies** (`capstone`, `freetype`) (`glfw` or `libxkbcommon`, `wayland`, `libglvnd`, and maybe `libegl` too):  
(use whatever package manager your distro provides)

2. Build with **make**:  
`<tracy-repo>/profiler/build/unix $ make -j8`

**Tracy-release**



1. Obtain the **dependencies** (`capstone`, `glfw`, `freetype`):  
`$ brew install capstone glfw freetype`

2. Build with **make**:  
`<tracy-repo>/profiler/build/unix $ make -j8`

**Tracy-release**

# Zoning...

```
{ this is a zone scope (RAII) }
```

```
int foo()
{
    ...
    ...
    if (!spoon) {
        ...
        ...
    }
    ...
    {
        ...
        ...
    }
}
```



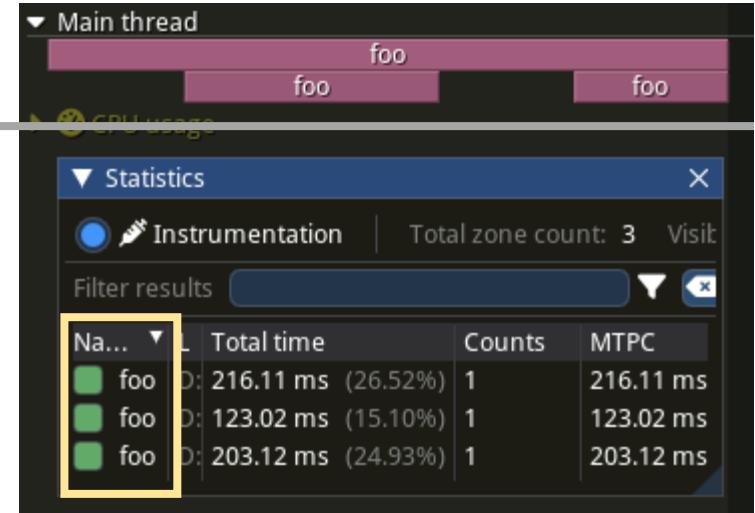
Source: [https://www.mobygames.com/game/657/simcity-2000/promo/group\\_17786/image-185051/](https://www.mobygames.com/game/657/simcity-2000/promo/group_17786/image-185051/)

# Tracy Zoning 101

```
int foo()
{
    ZoneScoped;
    ...
    if (!spoon) {
        ZoneScoped;
        ...
    }
    ...
    {
        ZoneScoped;
        ...
    }
}
```

thread [timeline](#)

zone [statistics](#)

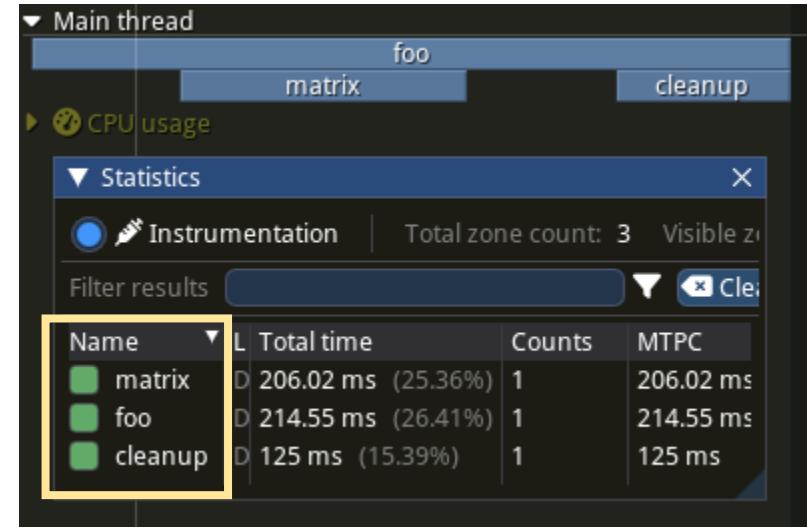


- **ZoneScoped**

- the “[bread-and-butter](#)” of Tracy
- only **1 zone per scope**
- zone [name is inferred](#) (enclosing function name)
- zone [color is automatic](#) (per-thread and depth)

# Tracy Zoning 101

```
int foo()
{
    ZoneScoped;
    ...
    if (!spoon) {
        ZoneScopedN("matrix");
        ...
    }
    ...
    {
        ZoneScopedN("cleanup");
        ...
    }
}
```



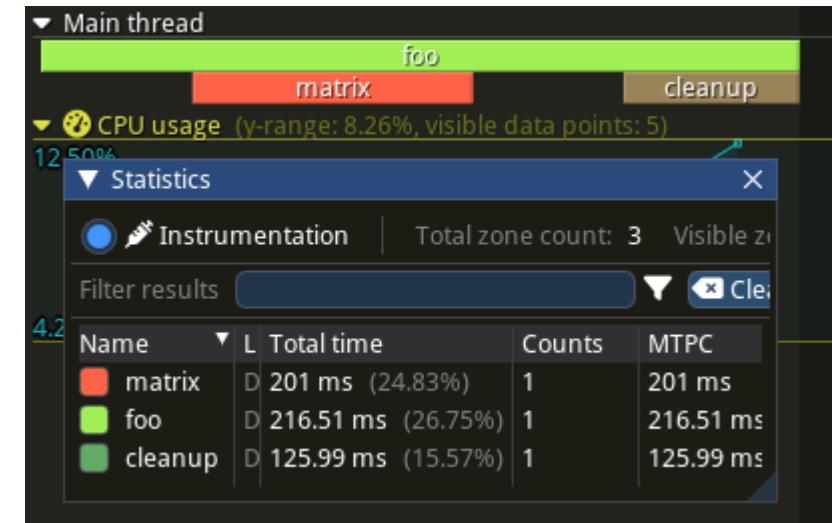
- **ZoneScopedN**
  - only 1 zone per scope
  - zone color is automatic (per-thread and depth)
  - **you name it!** (must be *persistent string literal*)

Tracy manual, Section 3.1

the provided string data must be accessible at any time in program execution (*this also includes the time after exiting the main function*). The string also cannot be changed. This basically means that the only option is to use a string literal

# Tracy Zoning 101

```
int foo()
{
    ZoneScopedC(0xA0F055);  
    ...  
    if (!spoon) {  
        ZoneScopedNC("matrix", Tomato);  
        ...  
    }  
    ...  
    {  
        ZoneScopedN("cleanup");  
        ...  
    }
}
```

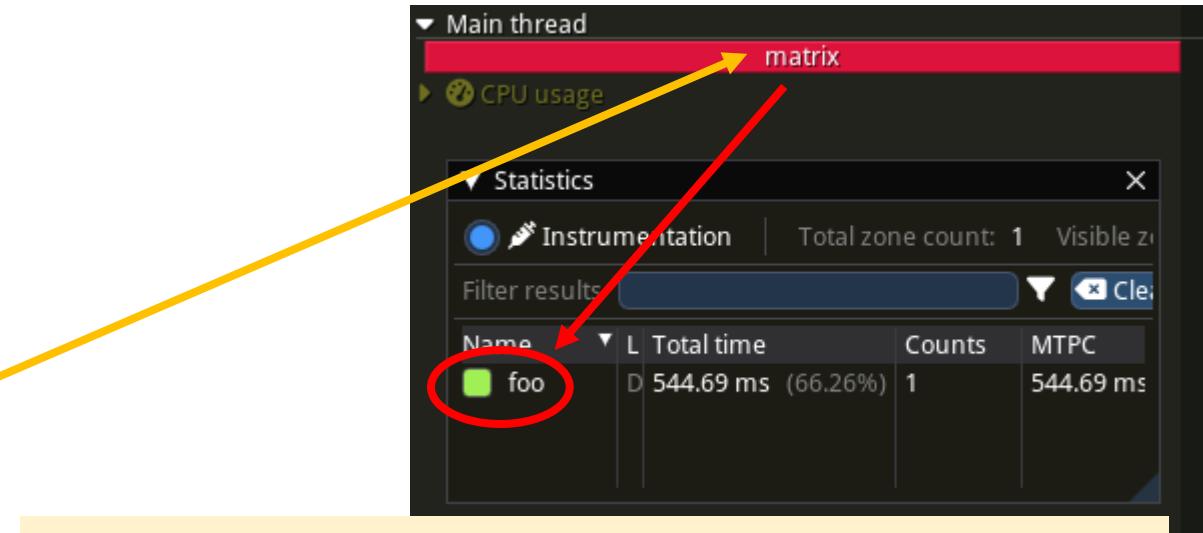


- **ZoneScoped[N]C**
  - only 1 zone per scope
  - you [may] name it
  - **you pick the color**
    - hexadecimal RGB triplet: 0xRRGGBB
    - or [X11 color names](#) (through **tracy::Color**)



# Tracy Zoning 201

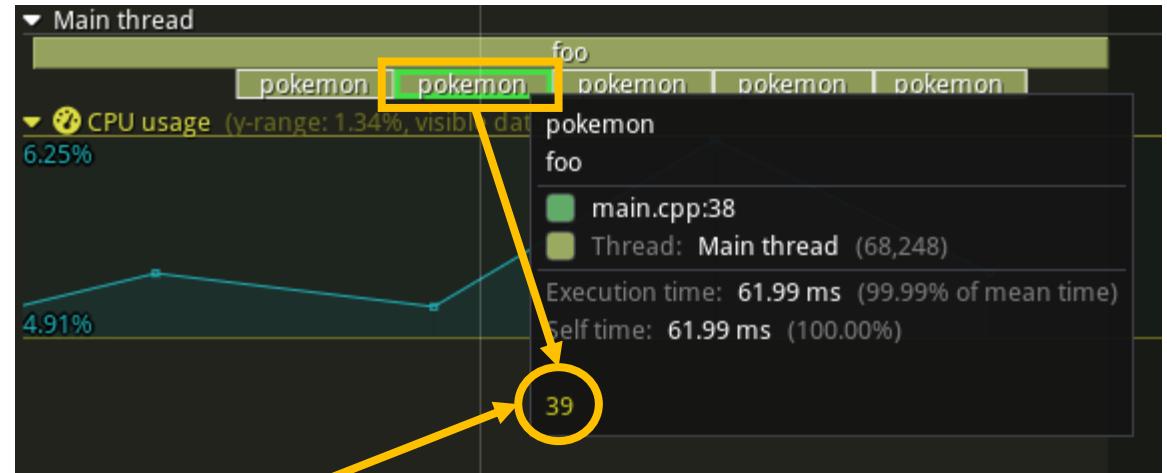
```
int foo()
{
    ZoneScoped;
    ...
    if (!spoon) {
        ZoneName("matrix", 6);
        ...
    }
    ...
    if (error != 0) {
        ZoneColor(Crimson);
        ...
    }
}
```



- **ZoneName**
  - programmatically change the zone name
  - can be **non-persistent string** (e.g., an std::string in the stack)  
**Avoid!** (it will confuse people unfamiliar with your code)
  - new name **WON'T appear** in profiler statistics!
- **ZoneColor**
  - programmatically change the zone color

# Tracy Zoning 201

```
int foo()
{
    ZoneScoped;
    ...
    Pokedex pokedex;
    ...
    ZoneValue(pokedex.size());
    for (auto& pokemon : pokedex)
    {
        ZoneScopedN("pokemon");
        ZoneValue(pokemon.id);
        ...
    }
}
```



- **ZoneValue**
  - logs a value to the current zone scope
  - accepts `int64` or `double` `uint64`

# Tracy Zoning 201

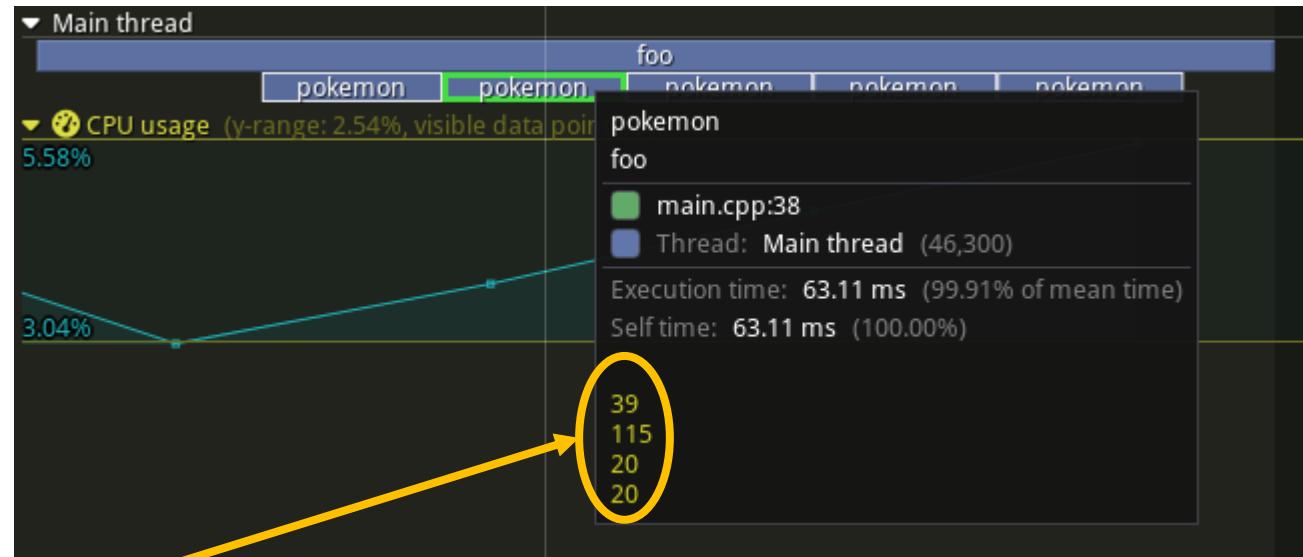
```
int foo()
{
    ZoneScoped;
    ...
    Pokedex pokedex;
    ...
    ZoneValue(pokedex)
    for (auto& pokemo
    {
        ZoneScopedN("Pokedex")
        ZoneValue(pok
        ...
    }
}
```



ne scope

# Tracy Zoning 201

```
int foo()
{
    ZoneScoped;
    Pokedex pokedex;
    for (auto& pokemon : pokedex)
    {
        ZoneScopedN("pokemon");
        ZoneValue(pokemon.id);
        ZoneValue(pokemon.hp);
        ZoneValue(pokemon.def);
        ZoneValue(pokemon.speed);
        ...
    }
}
```

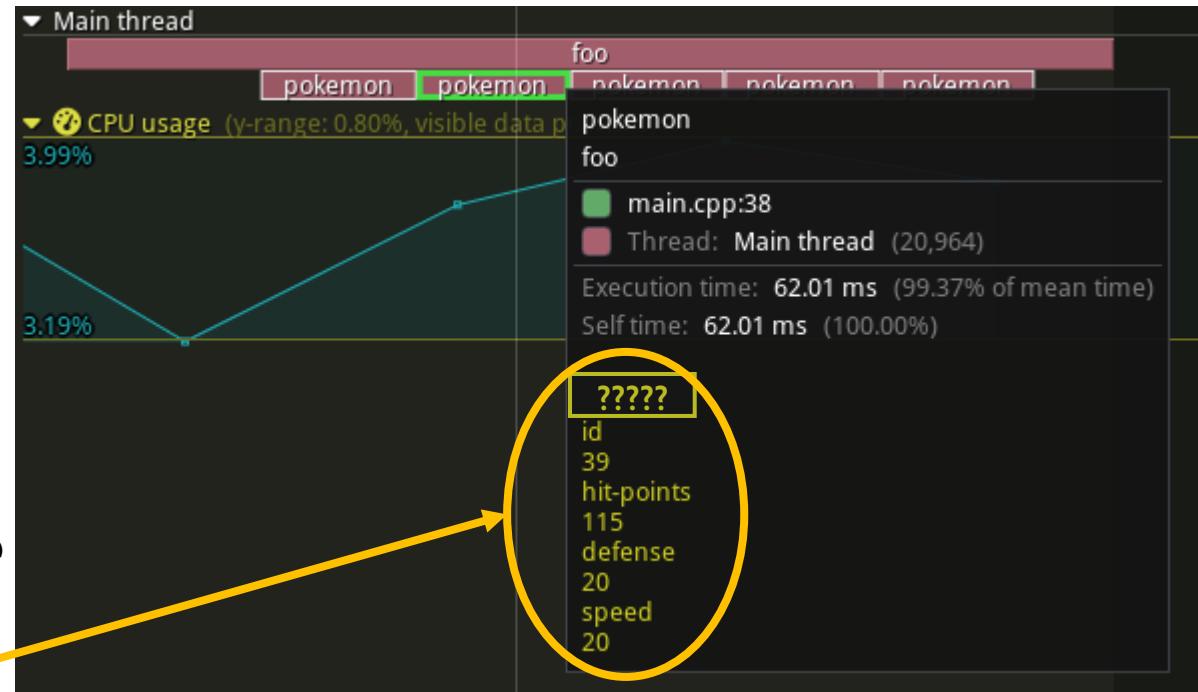


- **ZoneValue**

- logs a value to the current zone scope
- accepts `int64` or `double` `uint64`
- **multiple** values allowed!

# Tracy Zoning 201

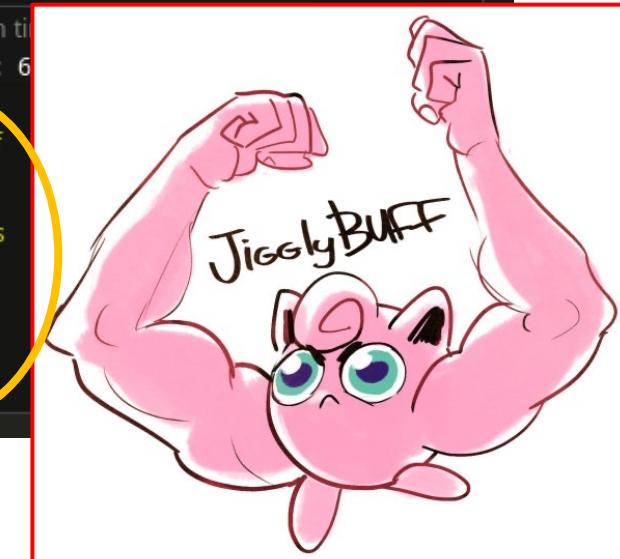
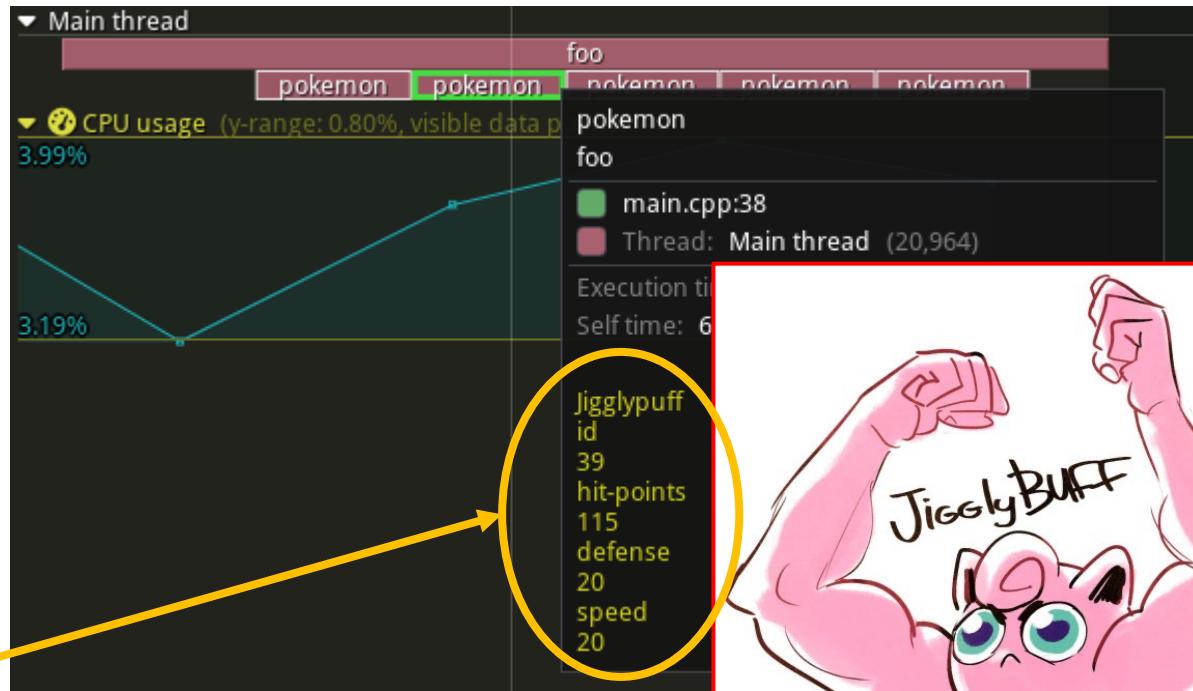
```
int foo()
{
    ZoneScoped;
    Pokedex pokedex;
    for (auto& pokemon : pokedex)
    {
        ZoneScopedN("pokemon");
        ...
        ZoneText("hit-points", 10);
        ZoneValue(pokemon.hp);
        ZoneText("defense", 7);
        ZoneValue(pokemon.def);
        ...
    }
}
```



- **ZoneText**
  - logs a string to the current zone scope
  - non-persistent string allowed
  - multiple texts allowed!

# Tracy Zoning 201

```
int foo()
{
    ZoneScoped;
    Pokedex pokedex;
    for (auto& pokemon : pokedex)
    {
        ZoneScopedN("pokemon");
        ...
        ZoneText("hit-points", 10);
        ZoneValue(pokemon.hp);
        ZoneText("defense", 7);
        ZoneValue(pokemon.def);
        ...
    }
}
```



Source: <https://www.reddit.com/r/pokemon/comments/3o5je7/jigglybuff/>

- **ZoneText**
  - logs a string to the current zone scope
  - non-persistent string allowed
  - multiple texts allowed!

# Tracy Zoning 301

```
int foo()
{
    ZoneScoped;
    ...
    ZoneNamedN(initzone, "init", true);
    ...
    ZoneNamedN(setupzone, "setup", true);
    ...
    ZoneNamedN(simzone, "simulate", true);
    ...
    ZoneNamedN(logzone, "log", false);
    ...
    ZoneNamedN(cleanzone, "cleanup", true);
    ...
}
```



- **ZoneNamed[N|C]**

- have **multiple zones** in the **same scope**
- each "**zone identifier**" must be **unique** in that scope
- creates a zone **stack, not a partition**
- offers "**active**" argument to **toggle** a zone (on/off) **programmatically**

# Tracy Zoning 301

**ZoneScoped[N|C|S]**

**ZoneColor(...)**

**ZoneValue(...)**

**ZoneText(...)**

**ZoneName(...)**

**ZoneIsActive()**

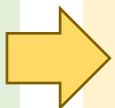
**ZoneNamed[N|C|S]**

**ZoneColorV(varname, ...)**

**ZoneValueV(varname, ...)**

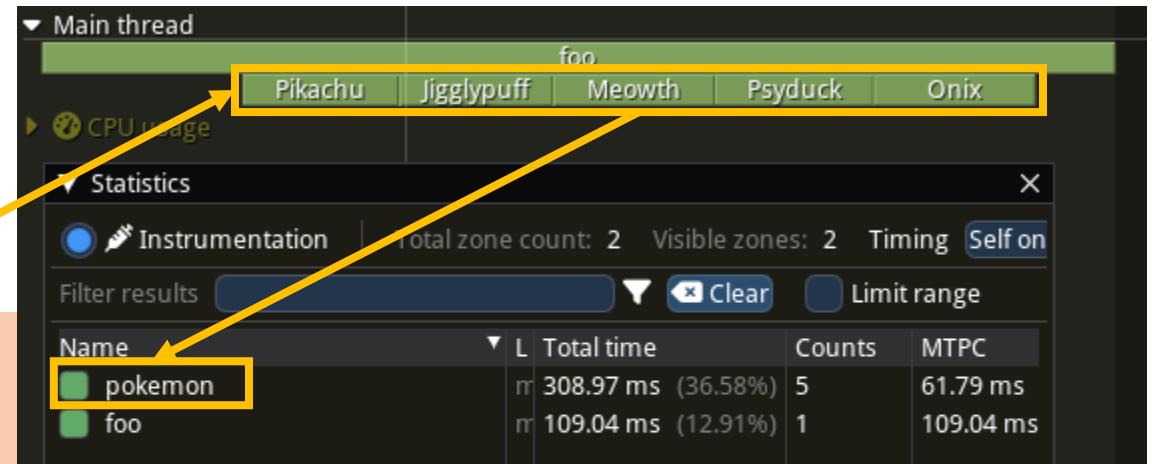
**ZoneTextV(varname, ...)**

**ZoneNameV(varname, ...)**



# Tracy Zoning 401

```
int foo()
{
    ZoneScoped;
    Pokedex pokedex;
    for (auto& pokemon : pokedex)
    {
        ZoneScopedN("pokemon");
        ZoneName(pokemon.name, strlen(pokemon.name));
        ...
    }
}
```



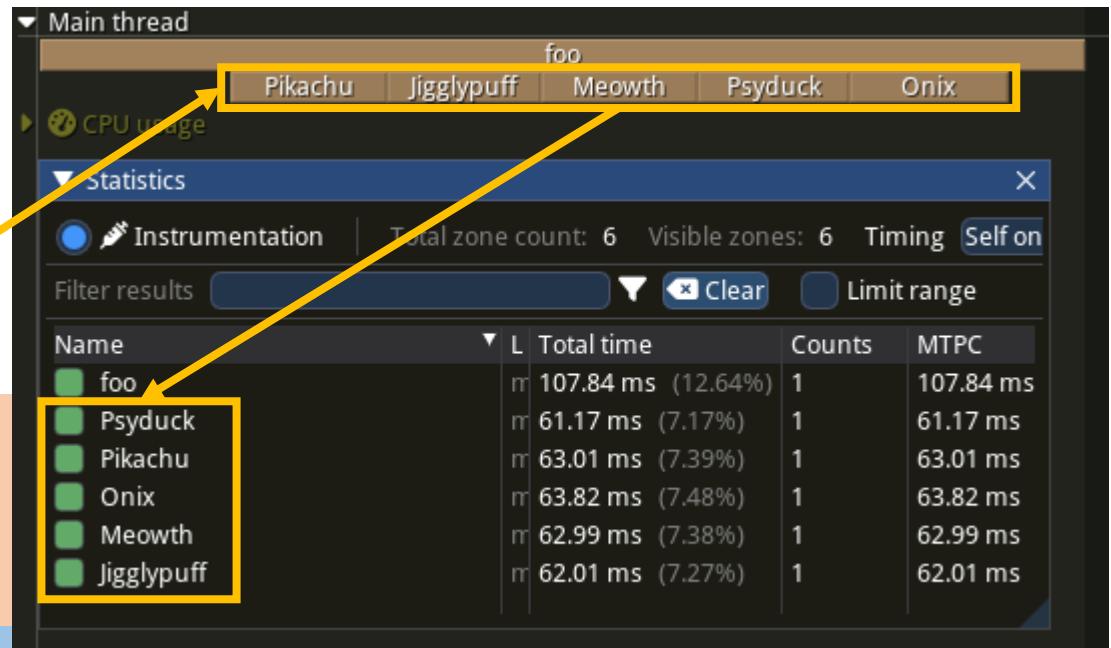
# Tracy Zoning 401

```
int foo()
{
    ZoneScoped;
    Pokedex pokedex;
    for (auto& pokemon : pokedex)
    {
        ZoneScopedN("pokemon");
        ZoneTransientN(pokezone, pokemon.name, true);
        ...
    }
}
```

Transient zones **may be necessary** when instrumenting a module that can be loaded **and unloaded** throughout the lifecycle of the host application (e.g., a plugin)

#### 3.4.4 Transient zones

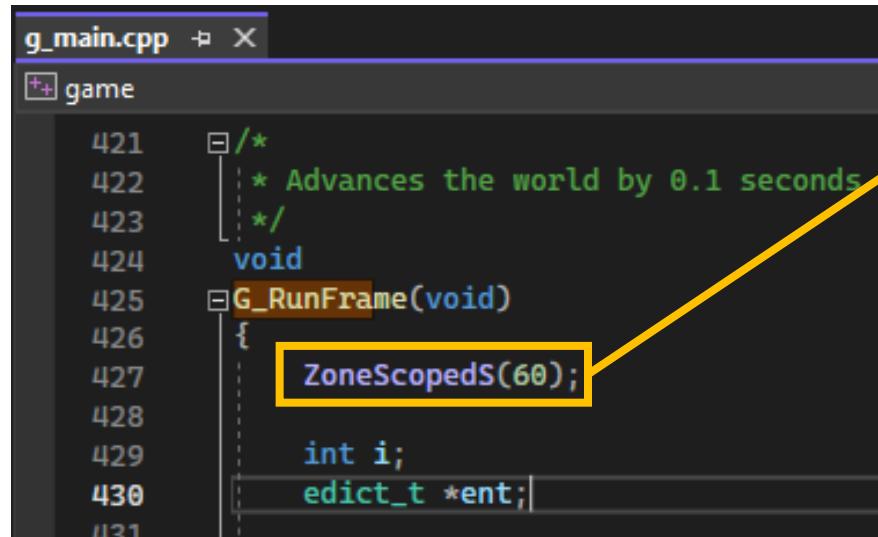
In order to prevent problems caused by unloadable code, described in section 3.1.1, transient zones copy the source location data to an on-heap buffer. This way, the requirement on the string literal data being accessible for the rest of the program lifetime is relaxed, at the cost of increased memory usage.



- **ZoneTransientN**

- name a zone with **non-persistent string**
- the **same zone** can have **multiple names**
- will **show up in statistics!**
- increased **runtime cost** to record zone

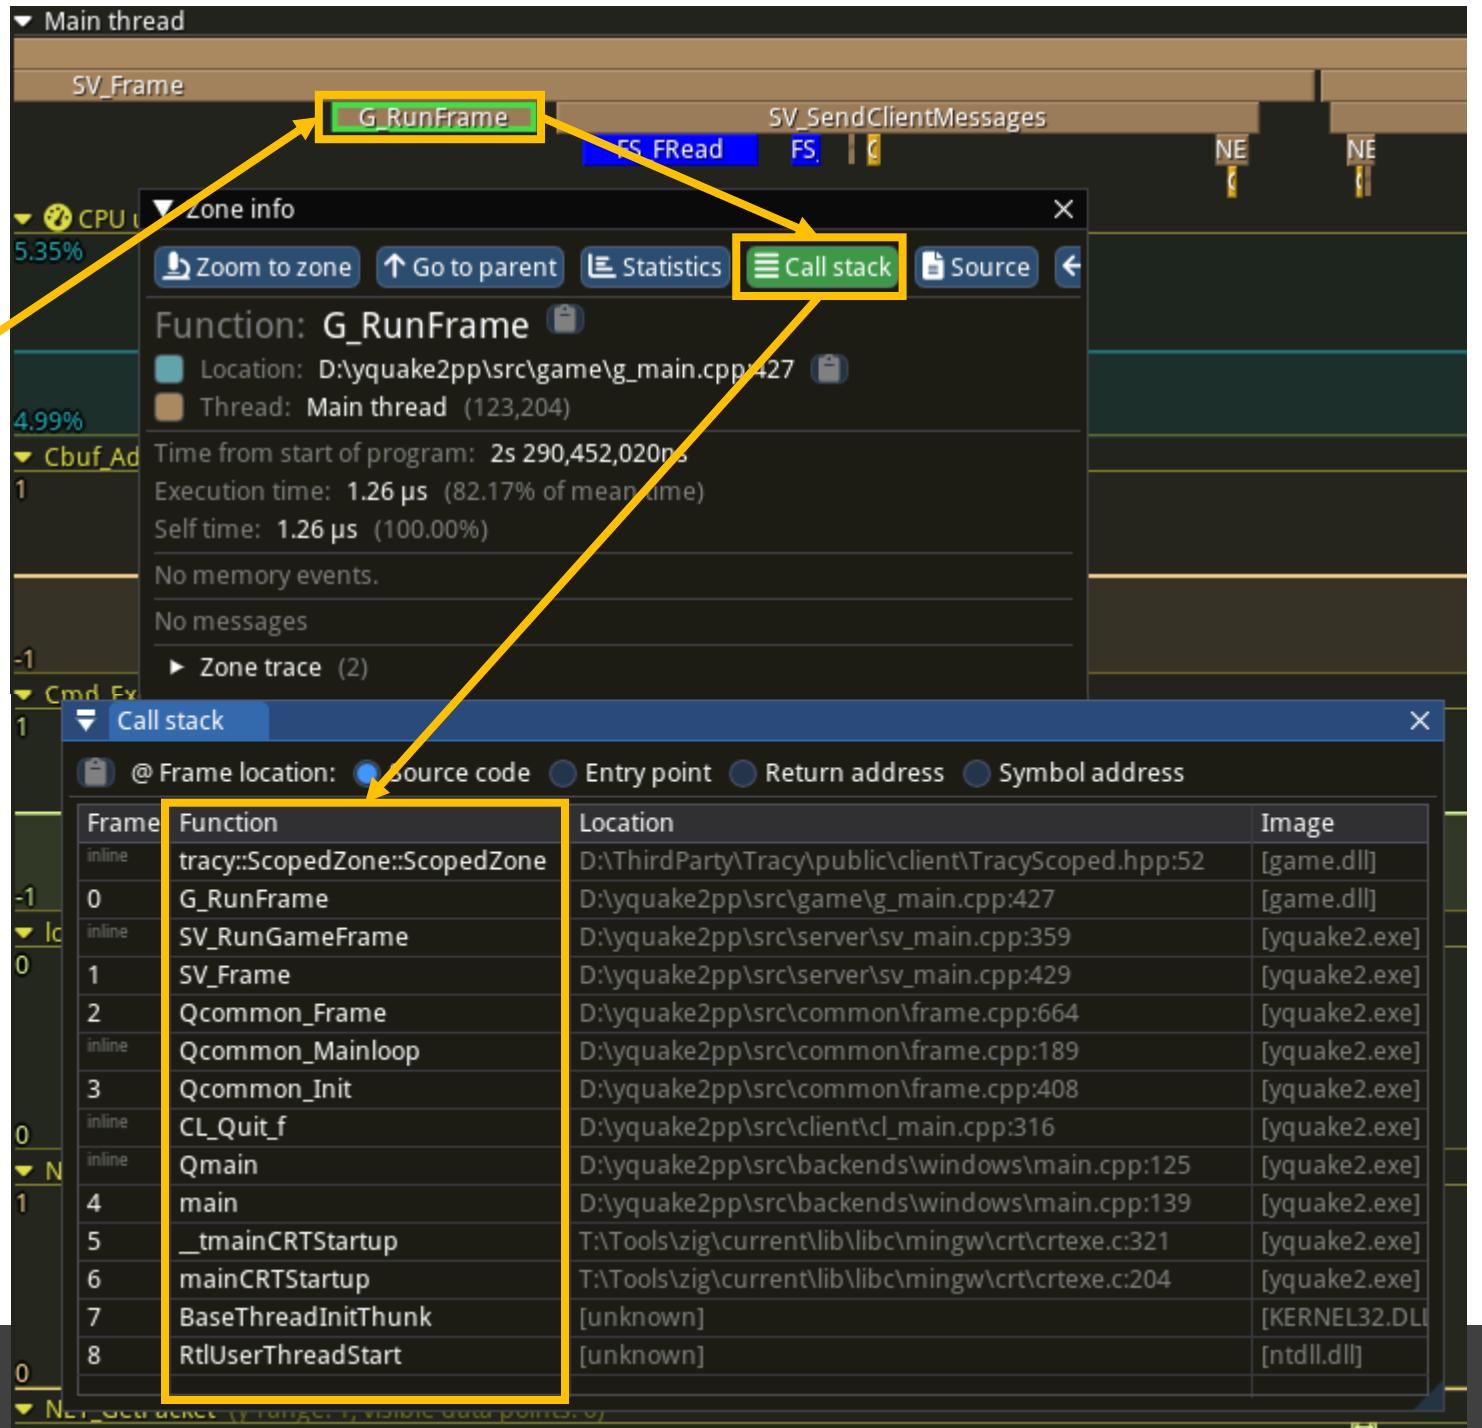
# Tracy Zoning 501



```
g_main.cpp  X
game

421  /*
422   * Advances the world by 0.1 seconds
423   */
424
425 void G_RunFrame(void)
426 {
427     ZoneScopedS(60);
428
429     int i;
430     edict_t *ent;
```

- **ZoneScoped(depth)**
  - captures **callstack**
  - adds some **overhead**
  - useful for **code exploration**
  - **max** callstack depth: **~60**
  - (enables selective sampling profiling)
- **ZoneScoped[N|CS]**

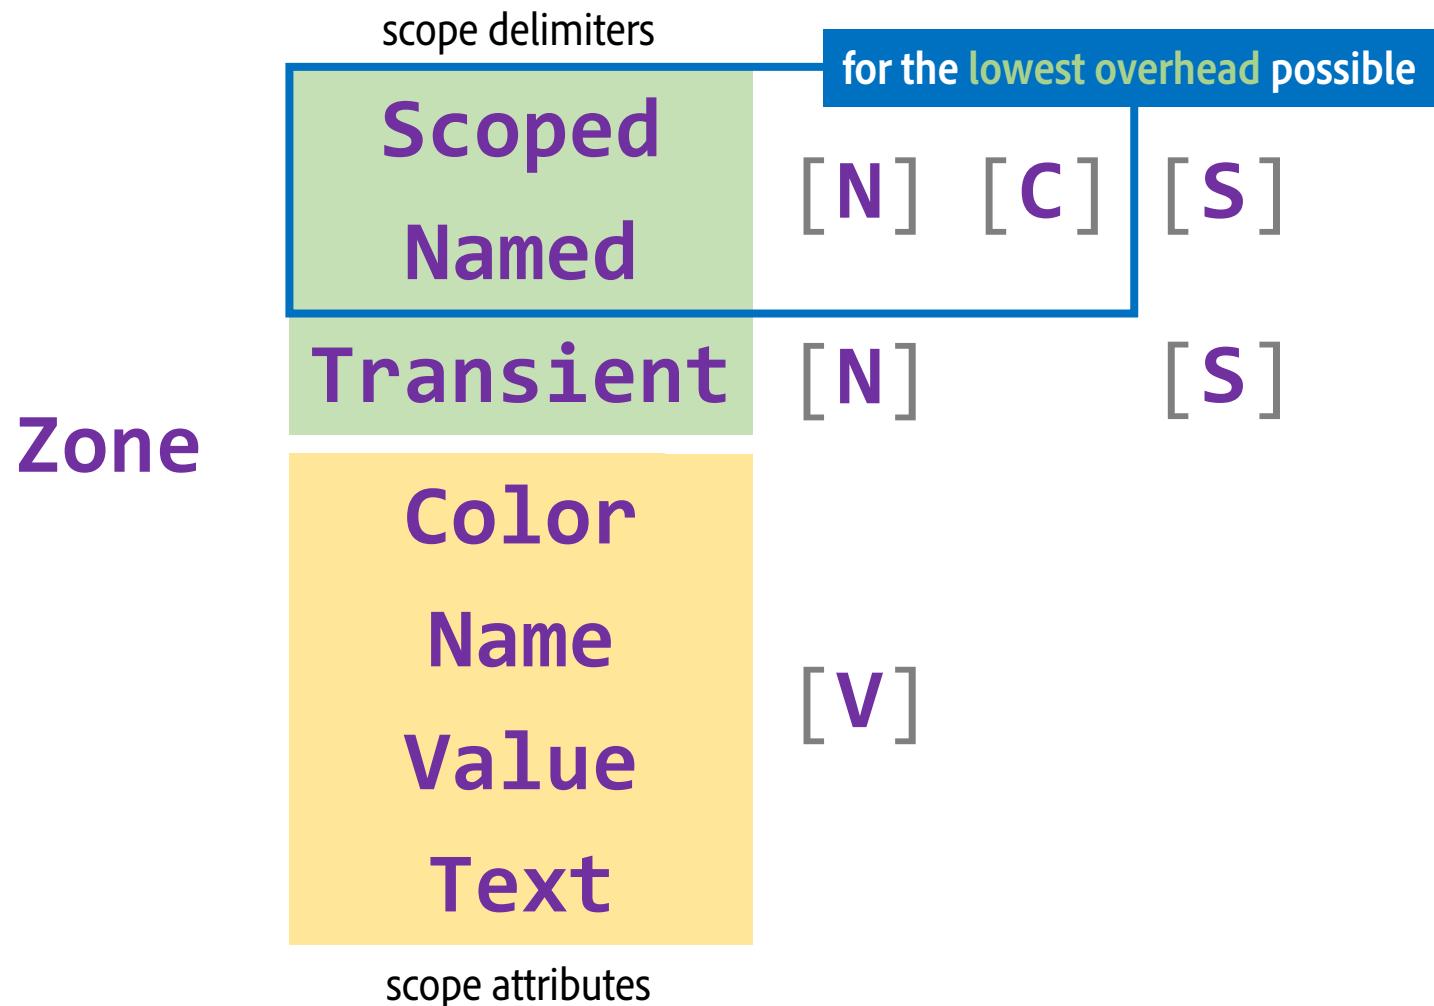


# Callstack capture cost

Depth	x86	x64	ARM	ARM64
1	34 ns	98 ns	6.62 µs	6.63 µs
2	35 ns	150 ns	8.08 µs	8.25 µs
3	36 ns	168 ns	9.75 µs	10 µs
4	39 ns	190 ns	10.92 µs	11.58 µs
5	42 ns	206 ns	12.5 µs	13.33 µs
10	52 ns	306 ns	19.62 µs	21.71 µs
15	63 ns	415 ns	26.83 µs	30.13 µs
20	77 ns	531 ns	34.25 µs	38.71 µs
25	89 ns	630 ns	41.17 µs	47.17 µs
30	109 ns	735 ns	48.33 µs	55.63 µs
35	123 ns	843 ns	55.87 µs	64.09 µs
40	142 ns	950 ns	63.12 µs	72.59 µs
45	154 ns	1.05 µs	70.54 µs	81 µs
50	167 ns	1.16 µs	78 µs	89.5 µs
55	179 ns	1.26 µs	85.04 µs	98 µs
60	193 ns	1.37 µs	92.75 µs	106.59 µs

Tracy Manual: Section 3.11 – Collecting Timestamps

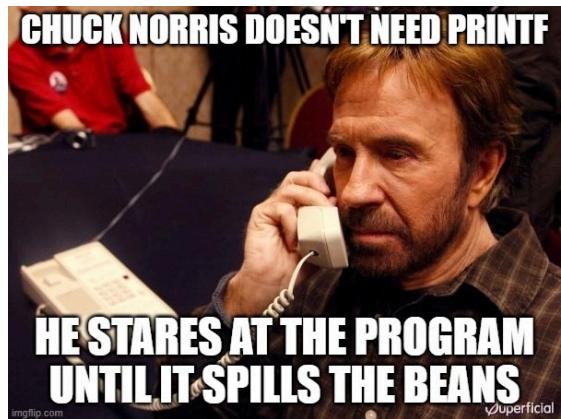
# Recap (cheat sheet)



## Suffixes:

- **N** : explicit name
- **C** : explicit color
- **S** : capture callstack
- **V** : explicit zone scope identifier  
(when using **ZoneNamed[N|C|S]**)  
(or **ZoneTransient[N|S]**)

# TracyMessage() is the new printf()



*persistent string literal*

## TracyMessage[ L | C | S ]

- console replacement!
- time-aware!
- thread-aware!
- message colors!
- browsable log!
- filterable log!
- no thread sync needed!

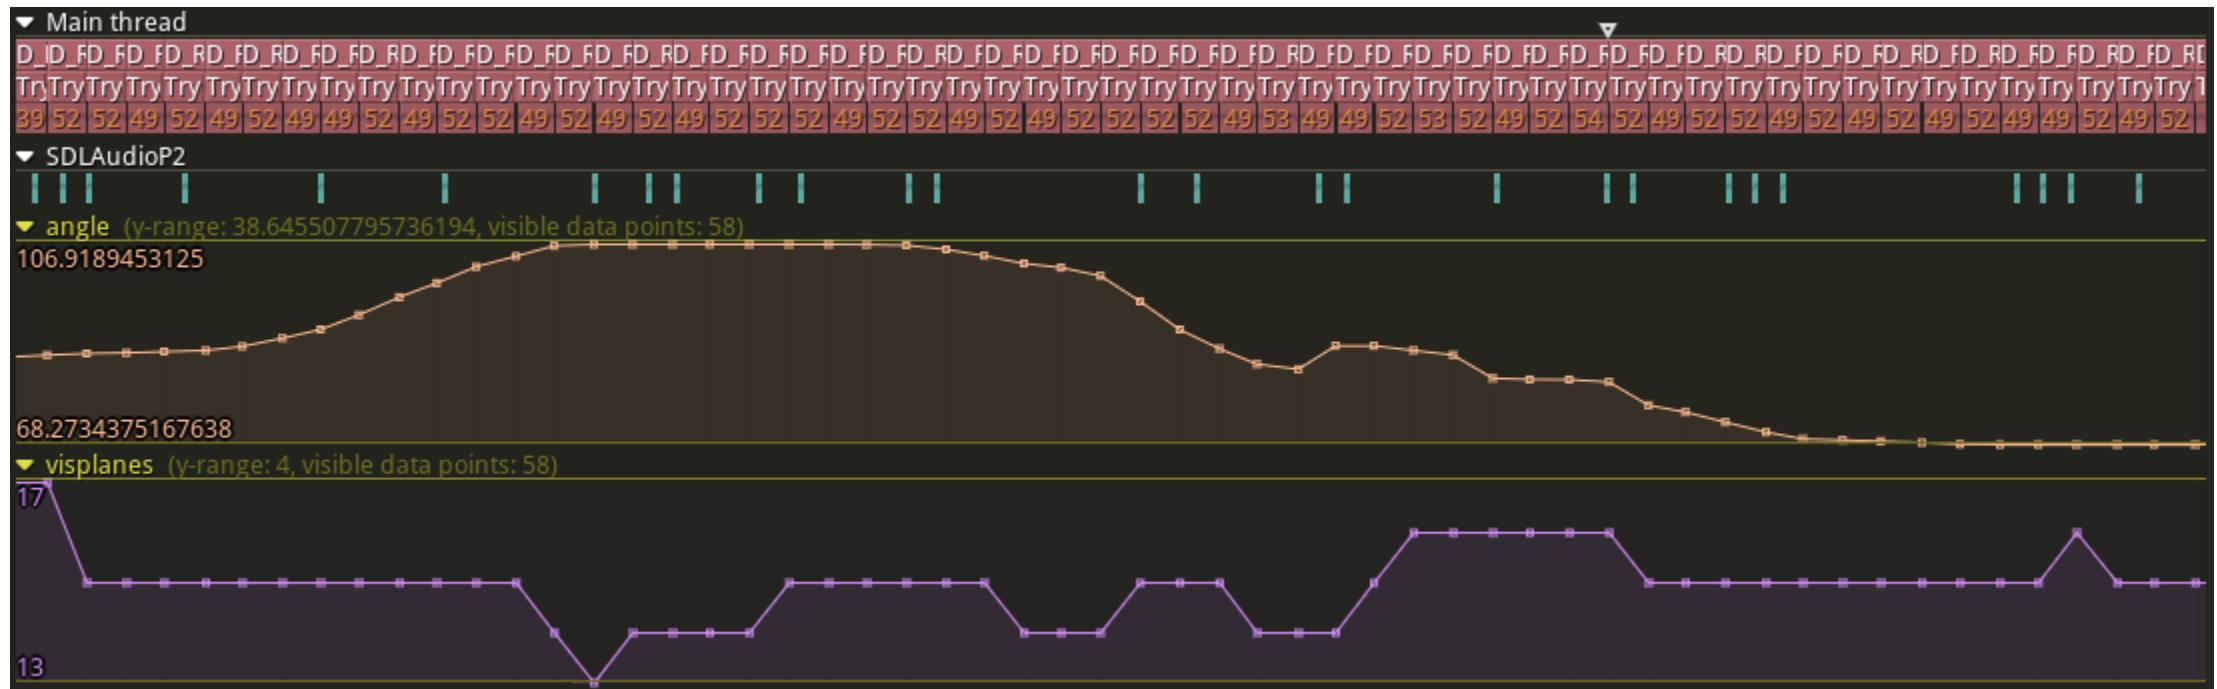
The screenshot shows the Tracy debugger interface. On the left, a timeline displays various events: Swap BuU, Swap Buffers, Stop Lantern, Start Lantern, GUI Update, and Load Scene. Below the timeline, a tree view shows threads: Main thread, Lantern Renderer (GPU), and multiple tbb worker threads (7, 1, 63, 62, 61, 60, 59, 58, 57, 56, 55, 54, 53, 52). A large yellow arrow points from the Chuck Norris meme towards the 'Messages' window on the right. This window is titled 'Messages' and contains a list of log entries. It includes a search bar, a 'Filter messages' button, and buttons for 'Visible threads' (65) and 'Clear'. The log entries show messages from the Lantern Renderer (GPU) and the Main thread, including timing information like '3s 890,905,210ns' and '6s 069,887,760ns', and various system and application logs.

Time	Thread	Message
3s 890,905,210ns	Lantern Renderer (GPU) (58,292)	ecclair (timing): Found 400 light-emitting triangle(s)
3s 909,796,918ns	Lantern Renderer (GPU) (58,292)	ecclair: Creating bottom-level acceleration structure
3s 913,235,212ns	Lantern Renderer (GPU) (58,292)	ecclair (timing): Created 8 bottom-level acceleration
3s 913,677,780ns	Lantern Renderer (GPU) (58,292)	ecclair: Compacting 8 bottom-level accelera
3s 915,435,222ns	Lantern Renderer (GPU) (58,292)	ecclair (timing): Compacted 8 bottom-level accelera
3s 915,834,133ns	Lantern Renderer (GPU) (58,292)	ecclair: Creating top-level acceleration structure...
3s 921,000,822ns	Lantern Renderer (GPU) (58,292)	ecclair: Scene:
3s 921,393,717ns	Lantern Renderer (GPU) (58,292)	ecclair: Bounding box: [-57.46, 6.0180078] x [0, 37
3s 921,999,181ns	Lantern Renderer (GPU) (58,292)	ecclair: Center: (-25.720995, 18.991982, 0)
3s 922,576,876ns	Lantern Renderer (GPU) (58,292)	ecclair: Diameter: 108.95981
3s 925,504,940ns	Lantern Renderer (GPU) (58,292)	ecclair: Creating bottom-level acceleration structure
3s 926,510,549ns	Lantern Renderer (GPU) (58,292)	ecclair (timing): Created 0 bottom-level acceleration
3s 927,105,899ns	Lantern Renderer (GPU) (58,292)	ecclair: Creating top-level acceleration structure...
3s 928,081,567ns	Lantern Renderer (GPU) (58,292)	ecclair (timing): Found 400 light-emitting triangle(s)
3s 931,773,179ns	Lantern Renderer (GPU) (58,292)	ecclair_vk_framebuffer: Creating persistent pixels, k
3s 933,207,281ns	Lantern Renderer (GPU) (58,292)	ecclair_vk_framebuffer: Creating persistent pixels, k
6s 069,887,760ns	Main thread (78,860)	Lantern CPU Check: Warning: No AVX available, Fal
6s 071,344,285ns	Main thread (78,860)	Lantern CPU Check: Warning: No FP16 (HALF-PREC
6s 072,370,240ns	Main thread (78,860)	TBB engage
6s 086,996,305ns	tbb worker 1 (72,712)	TBB engage
6s 087,021,622ns	tbb worker 7 (78,812)	TBB engage
6s 088,598,012ns	Main thread (78,860)	reading runtime_texture_binary_header_v4015
6s 106,293,508ns	Main thread (78,860)	reading runtime_texture_binary_header_v4015
6s 124,248,815ns	Main thread (78,860)	reading runtime_texture_binary_header_v4015
6s 141,332,404ns	tbb worker 29 (84,036)	TBB engage
6s 141,335,770ns	tbb worker 24 (58,764)	TBB engage
6s 141,336,011ns	tbb worker 57 (60,140)	TBB engage

# Plots

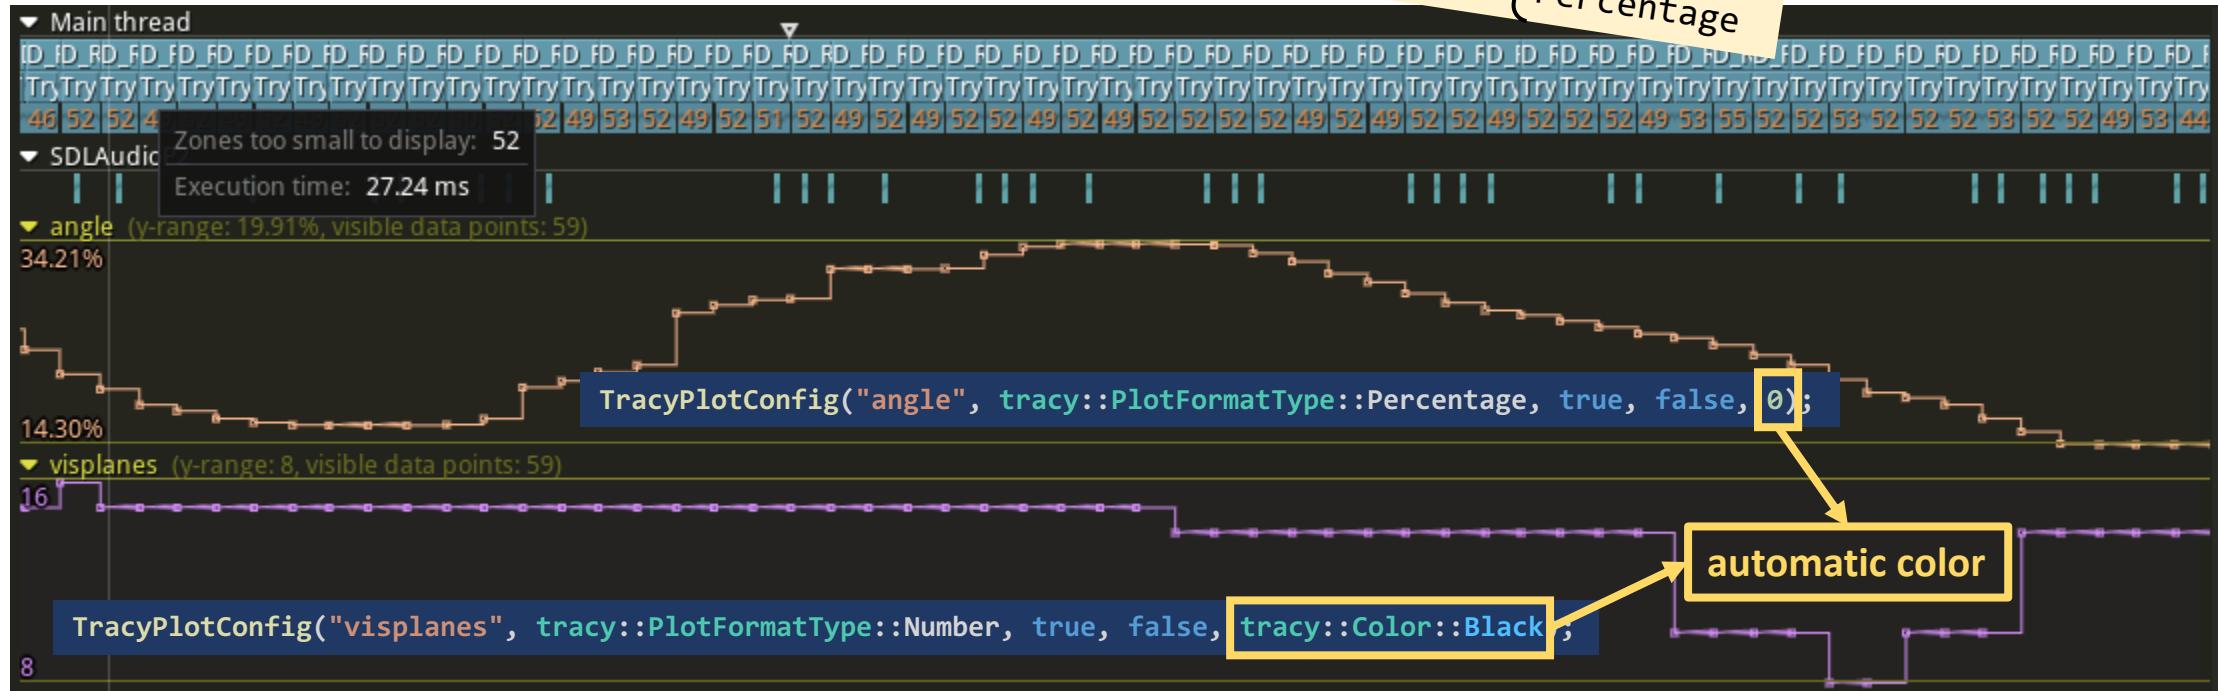
- `TracyPlot(name, value)`  
*persistent* string literal      int64, float or double

plots are **shared** among threads  
**no per-thread plot segregation**  
but profiler **sorts on timestamp**



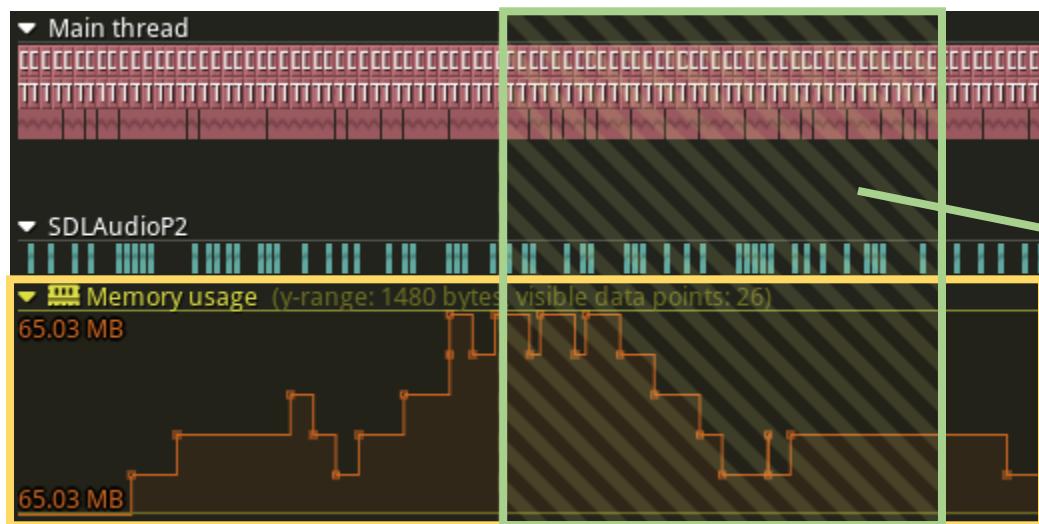
# Plots

- `TracyPlotConfig(name, type, step, fill, color)`



# Tracking Memory Allocations

- `TracyAlloc[N|S](ptr, size, ...)`
- `TracyFree[N|S](ptr, ...)`



TracyAlloc and TracyFree  
won't do memory **allocation**:  
it's merely for **tracking**

The screenshot shows a table of active memory allocations. The header row includes columns for Address, Size, Appearance, Duration, Thread, Zone alloc, and Zone free. The data rows show various allocations made by the "Main thread" using the "doom\_strdup" and "doom\_malloc" functions. All allocations are marked as "active".

Address	Size	Appeare...	Duration	Thread	Zone alloc	Zone free
0x2657fa3d5f	6 bytes	275,887,006ns	22.66 s	Main thread	doom_strdup	active
0x2657fa56af	24 bytes	275,918,470ns	22.66 s	Main thread	doom_malloc	active
0x2650c0f504	32 MB	3s 756,820,177	19.18 s	Main thread	doom_malloc	active
0x2650b9a94a	0 bytes	3s 756,900,165	19.18 s	Main thread	doom_strdup	active
0x2650bf6819	11 bytes	3s 757,173,674	19.18 s	Main thread	doom_strdup	active
0x2650bf682f	15 bytes	3s 757,223,345	19.18 s	Main thread	doom_strdup	active
0x2650b9a957	2 bytes	3s 757,446,604	19.18 s	Main thread	doom_strdup	active
0x2650bf6855	23 bytes	3s 757,454,096	19.18 s	Main thread	doom_strdup	active
0x2650b9a925	7 bytes	3s 757,455,801	19.18 s	Main thread	doom_strdup	active
0x2650be6dea	25 bytes	3s 757,458,045	19.18 s	Main thread	doom_strdup	active
0x2650b9a920	5 bytes	3s 757,459,964	19.18 s	Main thread	doom_strdup	active
0x2650bf6807	9 bytes	3s 757,461,717	19.18 s	Main thread	doom_strdup	active
0x2650bf6831	21 bytes	3s 757,463,483	19.18 s	Main thread	doom_strdup	active
0x2650bf6833	10 bytes	3s 757,465,088	19.18 s	Main thread	doom_strdup	active

# Lock instrumentation

```
std::mutex my_mutex;  
TracyLockable(std::mutex, my_mutex);  
  
and then  
std::lock_guard<LockableBase(std::mutex)> lock (my_mutex);
```

There's also **TracySharedLockable[N]** (and **SharedLockableBase**)  
for reader-writer mutex

# What's Next?

Tips & Tricks  
(& Gotchas)

Study Cases  
from Adobe

Explore  
Together

&

Time for  
Questions

GPU  
instrumentation  
(very briefly)

A Look at Tracy's  
Sampling Mode

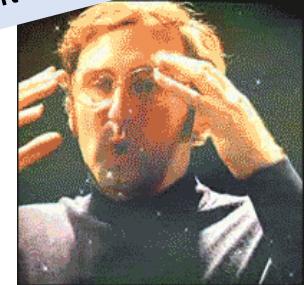
# Tips & Tricks



Apple users...



...get a **real mouse\***!



holding right button  
while moving pointer!

\*or help fix Apple trackpad gestures in ImGui

# Tips & Tricks

## Instrument the “**usual suspects**”

- Idling usually a good thing!
- Waits implicit ones too: e.g. `future.get()`
- Sleeps please, just don’t...
- Lock acquisition (contention) `TracyLockable`
- Memory operations `memcpy()`, `malloc()` et al.
- Explicit system calls (or heavy API calls) shader compilation
- Big-O activity (high-level) sorting, searching, container manipulations, ...
- I/O activity file, network, logging, ...
- Unpredictable events interrupts, user input activity, spurious wake-ups, ...

# Tips & Tricks

## Color zones appropriately

- Idling usually a good thing!
- Waits *implicit* ones too: e.g. `future.get()`
- Sleeps better yet, avoid `sleep()` in production code!
- Lock acquisition (contention) `TracyLockable`
- Memory operations `memcpy()`, `malloc()` et al.
- Explicit system calls (or heavy API calls) shader compilation
- Big-O activity (high-level) sorting, searching, container manipulations, ...
- I/O activity file, network, logging, ...
- Unpredictable events interrupts, user input activity, spurious wake-ups, ...

# Tips & Tricks

## Color zones appropriately

- Idling
- Waits
- Sleep
- Locks
- Memory
- Explicit
- Big-O
- I/O activity
- Unpredicted events

Just to be clear...  
Rely on Tracy's AUTOMATIC coloring EVERYWHERE  
**EXCEPT** for these **KNOWN-TO-BE-PROBLEMATIC** cases  
(they usually appear at the bottom of the zone stack)  
pick vibrant colors: you want them to pop up in the profiler!  
(Tracy's automatic coloring never chooses such saturated colors)

good thing!

future.get()

ion code!

lockable

() et al.

compilation

, searching, container manipulations, ...

file, network, logging, ...

interrupts, user input activity, spurious wake-ups, ...

# Tips & Tricks

**Standardize** the colors (don't think about them!)

Semantic Palette

```
enum Tier
{
    Idle      = tracy::Color::DimGray,
    Wait      = tracy::Color::Crimson,
    Sleep     = tracy::Color::Tomato,
    Lock      = tracy::Color::Coral,
    Memory    = tracy::Color::Goldenrod,
    System    = tracy::Color::Gold,
    BigO      = tracy::Color::LimeGreen,
    IO        = tracy::Color::Blue,
    Fickle    = tracy::Color::Magenta,
    ...
};
```

X11 color names

(consider color blindness!)



# Tips & Tricks

Name your **threads!**

```
tracy::SetThreadName("thread-name")
```

**TBB workers** : use `task_scheduler_observer`

```
class tracy_tbb_observer : public oneapi::tbb::task_scheduler_observer {
public:
    tracy_tbb_observer(tbb::task_arena& arena) : oneapi::tbb::task_scheduler_observer(arena) {
        observe(true); // activate the observer
    }
    void on_scheduler_entry(bool is_worker) override {
        if (is_worker) {
            int tid = tbb::this_task_arena::current_thread_index();
            tracy::SetThreadName(std::string("tbb worker #").append(std::to_string(tid)).c_str());
        }
    }
    void on_scheduler_exit(bool is_worker) override { }
};
```



- ▶ Main thread
- ▶ 38464
- ▶ 24620
- ▶ 46664
- ▶ 20456
- ▶ 36208
- ▶ 34724
- ▶ 27336
- ▶ 27236
- ▶ 42764

- ▶ Main thread
- ▶ tbb worker #7 [CompCore]
- ▶ tbb worker #9 [CompCore]
- ▶ tbb worker #6 [CompCore]
- ▶ tbb worker #5 [CompCore]
- ▶ tbb worker #11 [CompCore]
- ▶ tbb worker #10 [CompCore]
- ▶ tbb worker #8 [CompCore]
- ▶ tbb worker #1 [CompCore]
- ▶ tbb worker #3 [CompCore]

# Tips & Tricks

Consider “siloing” the **usual suspects**

example: don't call **memcpy()** directly... call some **my\_memcpy()** which then calls **memcpy()**

```
void* my_memcpy(void* dst, const void* src, size_t size)
{
    ZoneScopedC(tracy::Color::Blue);
    return ::memcpy(dst, src, size);
}
```

**unified point** for instrumenting the usual suspects!

whole-program **symbol hijacking**

(e.g., to replace malloc without siloing across all process modules)

[LD\\_PRELOAD](#)

[DYLD\\_INSERT\\_LIBRARIES](#)

[Microsoft Detours](#)

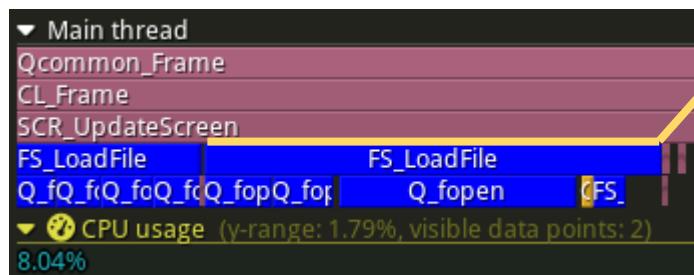
[weak symbol linkage interposition](#)

[MSVC linker /alternatename](#)

# Tips & Tricks

Capture **callstack** on “terminal” zones for “code reconnaissance”

```
048  /*  
649  int  
650  FS_LoadFile(char *path, void **buffer)  
651  {  
652      ZoneScopedCS(tracy::Color::Blue, 64);  
653      ZoneText(path, strlen(path));  
654  
655      byte *buf; /* Buffer. */  
656      int size; /* File size. */  
657      fileHandle_t f; /* File handle. */  
658  }
```



A screenshot of a debugger's call stack window. The stack trace is identical to the one shown in the CPU usage monitor above. A yellow arrow points from the highlighted row in the CPU usage monitor to the corresponding row in the call stack window.

Frame	Function	Location	Image
0	FS_LoadFile	D:\yquake2pp\src\common\filesystem.cpp:652	[yquake2.exe]
1	LoadPCX	D:\yquake2pp\src\client\refresh\files\pcx.cpp:150	[ref_soft.dll]
2	R_LoadImage	D:\yquake2pp\src\client\refresh\soft\sw_image.cpp:593	[ref_soft.dll]
3	R_FindImage	D:\yquake2pp\src\client\refresh\soft\sw_image.cpp:697	[ref_soft.dll]
4	RE_Draw_PicScaled	D:\yquake2pp\src\client\refresh\soft\sw_draw.cpp:392	[ref_soft.dll]
5	M_Main_Draw	D:\yquake2pp\src\client\menu\menu.cpp:621	[yquake2.exe]
6	M_Draw	D:\yquake2pp\src\client\menu\menu.cpp:4973	[yquake2.exe]
7	SCR_UpdateScreen	D:\yquake2pp\src\client\cl_screen.cpp:1656	[yquake2.exe]
8	CL_Frame	D:\yquake2pp\src\client\cl_main.cpp:843	[yquake2.exe]
9	Qcommon_Frame	D:\yquake2pp\src\common\frame.cpp:680	[yquake2.exe]
10	Qcommon_Mainloop	D:\yquake2pp\src\common\frame.cpp:189	[yquake2.exe]
11	Qcommon_Init	D:\yquake2pp\src\common\frame.cpp:408	[yquake2.exe]
12	CL_Quit_f	D:\yquake2pp\src\client\cl_main.cpp:316	[yquake2.exe]
13	Qmain	D:\yquake2pp\src\backends\windows\main.cpp:125	[yquake2.exe]
14	main	D:\yquake2pp\src\backends\windows\main.cpp:139	[yquake2.exe]
15	_tmainCRTStartup	T:\Tools\zig\current\lib\libc\mingw\crt\crtexe.c:321	[yquake2.exe]
16	mainCRTStartup	T:\Tools\zig\current\lib\libc\mingw\crt\crtexe.c:204	[yquake2.exe]
17	BaseThreadInitThunk	[unknown]	[KERNEL32.DLL]
18	RtlUserThreadStart	[unknown]	[ntdll.dll]



- the “**usual suspects**” are usually **leaf/terminal** calls
- they appear at the **bottom** of the **zone stack**
- ideal **candidates** for **code recon** via **callstack** inspection

# Tips & Tricks

Capture **callstack** on **allocation** to hunt down **leaks** : **TracyAlloc[N]S()**

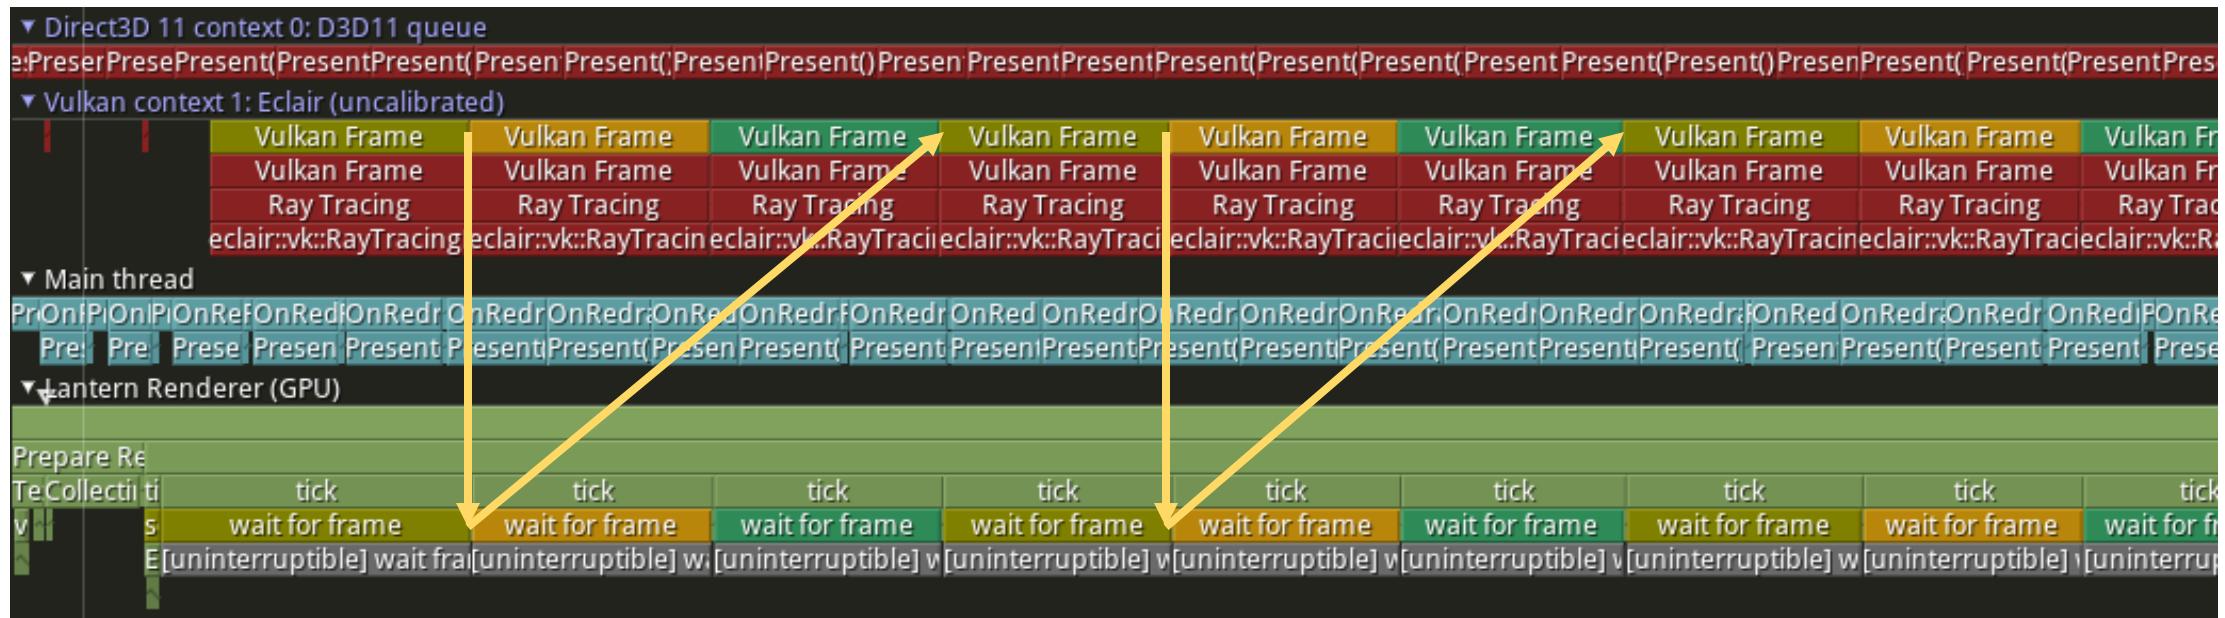
The screenshot shows the TracyAlloc memory profiler interface. The main window displays a table of active allocations, each with a unique address, size (845 bytes), appeared at timestamp (e.g., 4s 200,359,212), duration (1.02 s), thread (Main thread), zone alloc (Qmalloc), zone free (Qfree), and a call stack. The call stack for the first allocation is expanded, showing a chain of function calls from Qmalloc in ref\_soft.dll up to RtlUserThreadStart in ntdll.dll. A yellow box highlights the 'alloc' entry in the call stack column for the first allocation. A yellow callout bubble points to this entry with the text: "For known object leaks, consider instrumenting constructors and destructors with TracyAlloc/Free".

Address	Size	Appeared at	Duration	Thread	Zone alloc	Zone free	Call stack
0x15d1d16ebc	845 bytes	4s 200,359,212	1.02 s	Main thread	Qmalloc	Qfree	alloc [free]
0x15d1d6061f	845 bytes	4s 201,298,125	1.02 s	Main thread	Qmalloc	Qfree	alloc [free]
0x15d1d60655	845 bytes	4s 202,266,676	1.02 s	Main thread	Qmalloc	Qfree	alloc [free]
0x15d1d6068b	845 bytes	4s 203,251,339	1.02 s	Main thread	Qmalloc	Qfree	alloc [free]
0x15d1d57c50	845 bytes	4s 204,230,647	1.02 s	Main thread	Qmalloc	Qfree	alloc [free]
0x15d1d57c86	845 bytes	4s 205,185,405	1.02 s	Main thread	Qmalloc	Qfree	alloc [free]

For known object leaks, consider  
instrumenting constructors and  
destructors with TracyAlloc/Free

# Tips & Tricks

Use **color cycling** to visually **correlate** zones



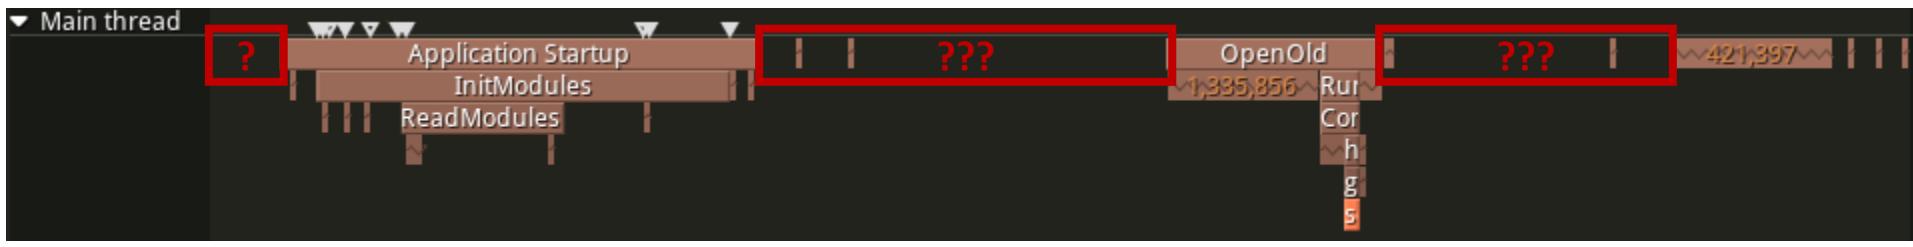
# Tips & Tricks

Use `TracyIsConnected()` to branch on “expensive” tracing stuff

```
if (TracyIsConnected()) {
    // expensive stuff I want to see in the profiler
    // (now that I know the profiler is connected to the host app)
    // ...
    ZoneScopedS(60);      // let's capture the callstack
    // ...
    ZoneText(...);
    // ...
    TracyMessage(...);
    // ...
}
```

# Tips & Tricks

Avoid “**empty gaps**” between zones (at least in the **top-most** level)



# Tips & Tricks

Add API prefix to zone name when no prefix/namespace exists

```
▼ Main thread
Main::MessageLoop::getMessages
TApplication::ProcessEvent
TApplication::OpenOld
CompositeSession::RunTasks
CompositeSession::CompositeTarget

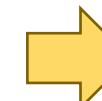
ha h pack anonymous-namespace'::get_kernel
make_kernel_key sc
```



```
▼ Main thread
Main::MessageLoop::getMessages
TApplication::ProcessEvent
TApplication::OpenOld

ha h pack [halide_manta]::get_kernel
make_kernel_key sc
```

Use the :: separator – Tracy Profiler  
will “split” on that for visualization



useful for filtering in statistics window

```
▼ Statistics
Instrumentation Symbols Total zone
Filter results halide_manta

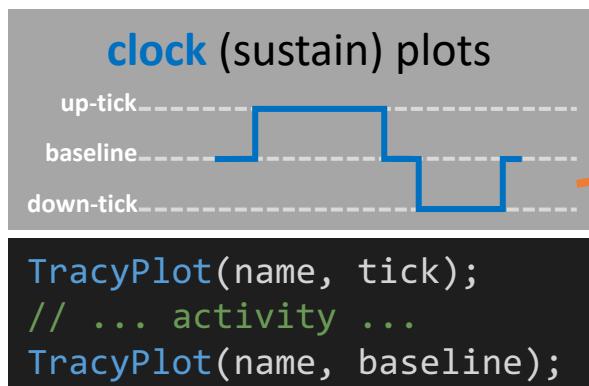
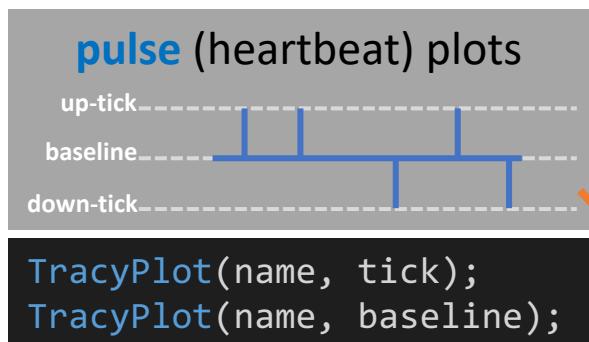
Name
halide_manta_run_internal::commands
halide_manta_run
halide_manta_release_context
halide_manta_initialize_kernels_internal::new
halide_manta_initialize_kernels
halide_manta_acquire_context
`anonymous-namespace'::halide_manta_device_relea
[halide_manta_d3d12]::compile_shader
[halide_manta]::get_kernel
```

# Tips & Tricks

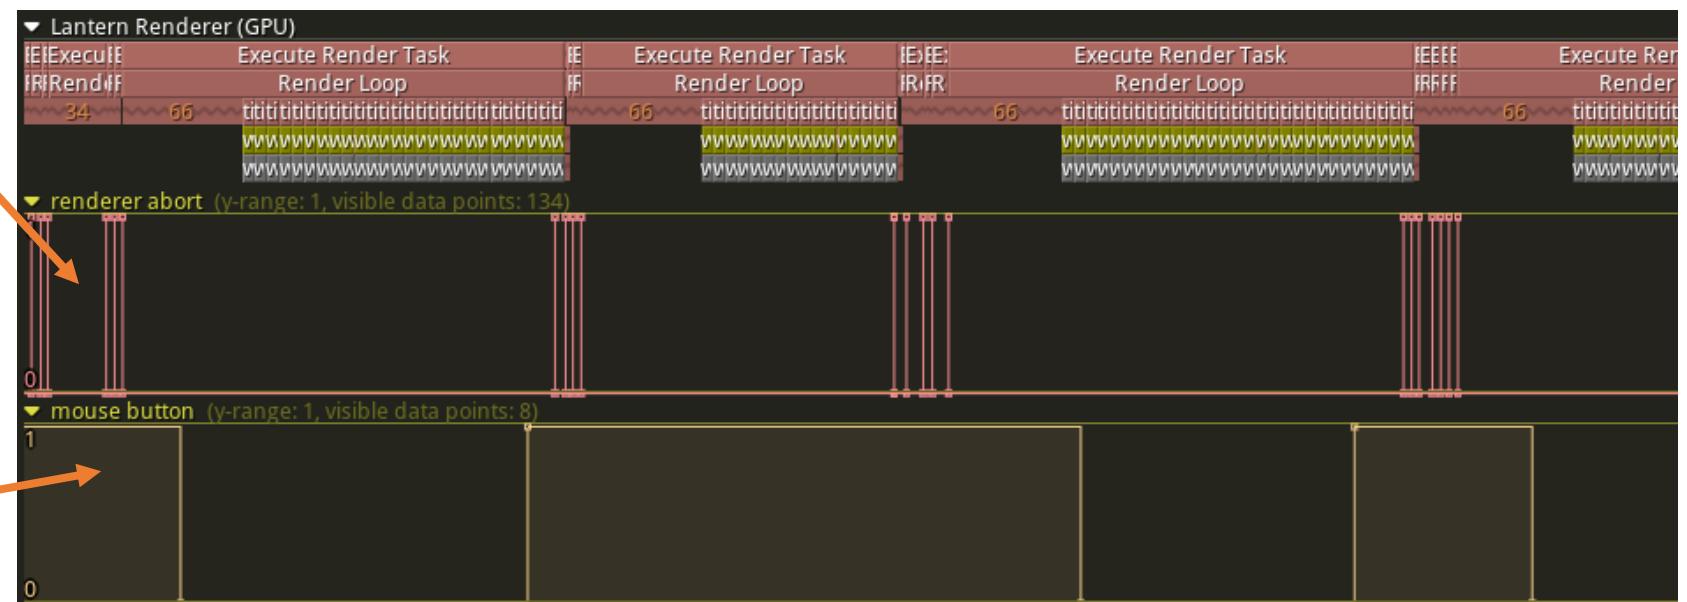
Use **tick signals** to visualize **critical activity**

- user input
- critical callback
- interrupt signals
- IPC signals
- completion signals
- wake-up signals

`TracyPlotConfig(name, type, true, fill, color)`



stairstep



# Tips & Tricks

Use **TracyAlloc()** and **TracyFree()** to track *anything* (objects, resources, handles, ...)

```
bytes...  
TracyAllocN(&frodo, 1, "party members");  
TracyAllocN(&gandalf, 1, "party members");  
TracyAllocN(&aragorn, 1, "party members");  
// ...  
TracyMessageL("[Gandalf] You shall not pass!");  
TracyFreeN(&gandalf, "party members");
```



# Tips & Tricks

## Instrumenting lock contention

too intrusive...

```
std::mutex my_mutex;  
TracyLockable(std::mutex, my_mutex);  
and then  
std::lock_guard<LockableBase(std::mutex)> lock (my_mutex);
```

do this siloing instead!

```
auto scoped_mutex_lock(std::mutex& m) {  
    ZoneScopedC(tracy::Color::Coral);  
    return std::unique_lock(m);  
}
```

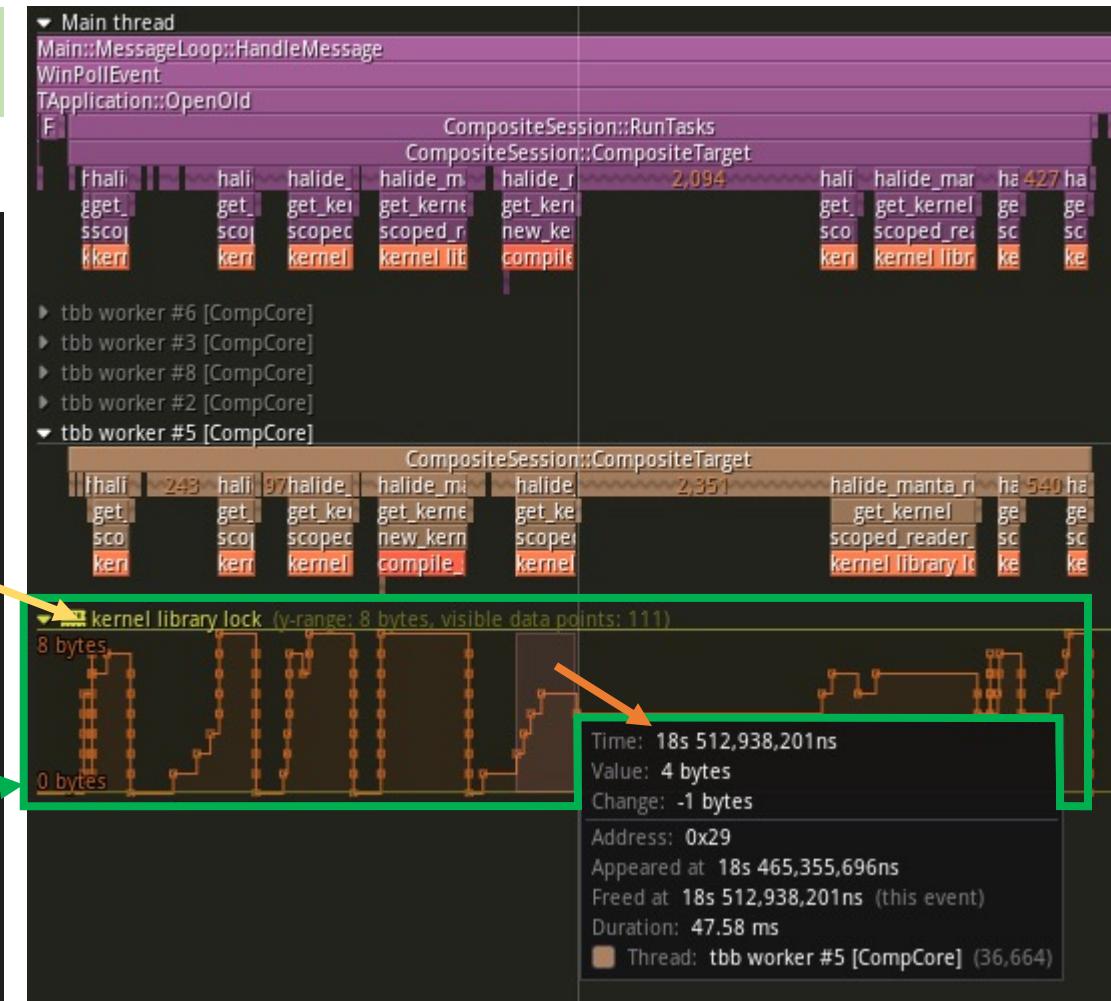
easily adaptable to reader-writer mutex as well  
(std::shared\_mutex and std::shared\_lock)

# Tips & Tricks

## Measuring lock contention

```
static std::atomic<int> lock_id_generator = 0;

auto scoped_mutex_lock(std::mutex& m,
                      const char* name)
{
    if (m.try_lock())
    {
        // lucky! no contention!
        std::unique_lock lock(m, std::adopt_lock);
        return lock;
    }
    // uh-oh... possible contention...
    ZoneTransientN(mutex_lock_scope, name, true);
    ZoneColorV(mutex_lock_scope, tracy::Color::Coral);
    auto token = reinterpret_cast<void*>(++lock_id_generator);
    TracyAllocN(token, 1, name);
    std::unique_lock lock(m);
    TracyFreeN(token, name);
    return lock;
}
```



# Tips & Tricks

Have an in-app **GUI** element to **launch & connect** the Tracy profiler

Team must agree on a **common location** for the Tracy profiler executable

David Farrell's  
pro-tip!  
**Md**

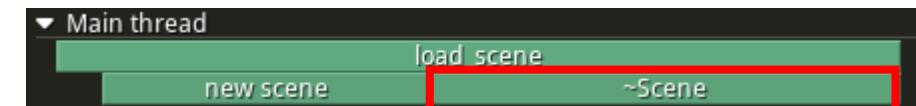
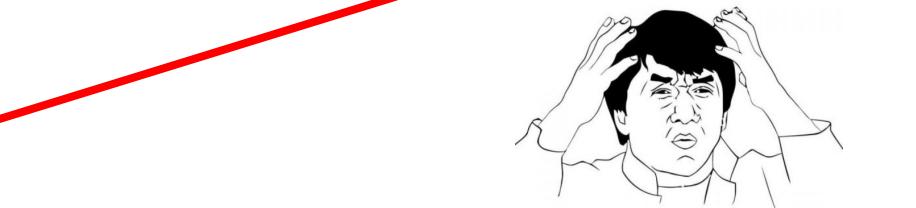
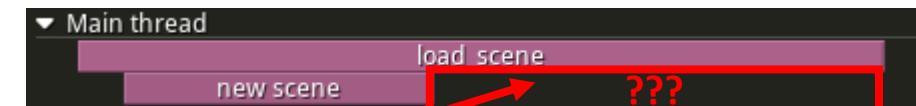


# Gotchas

## Unexplainable “tail” gaps:

- Implicit **deferred execution** (e.g. smart-pointer deleting an object)
- Instrument expensive **destructors**
- or **explicitly release** smart objects at the end of a scope

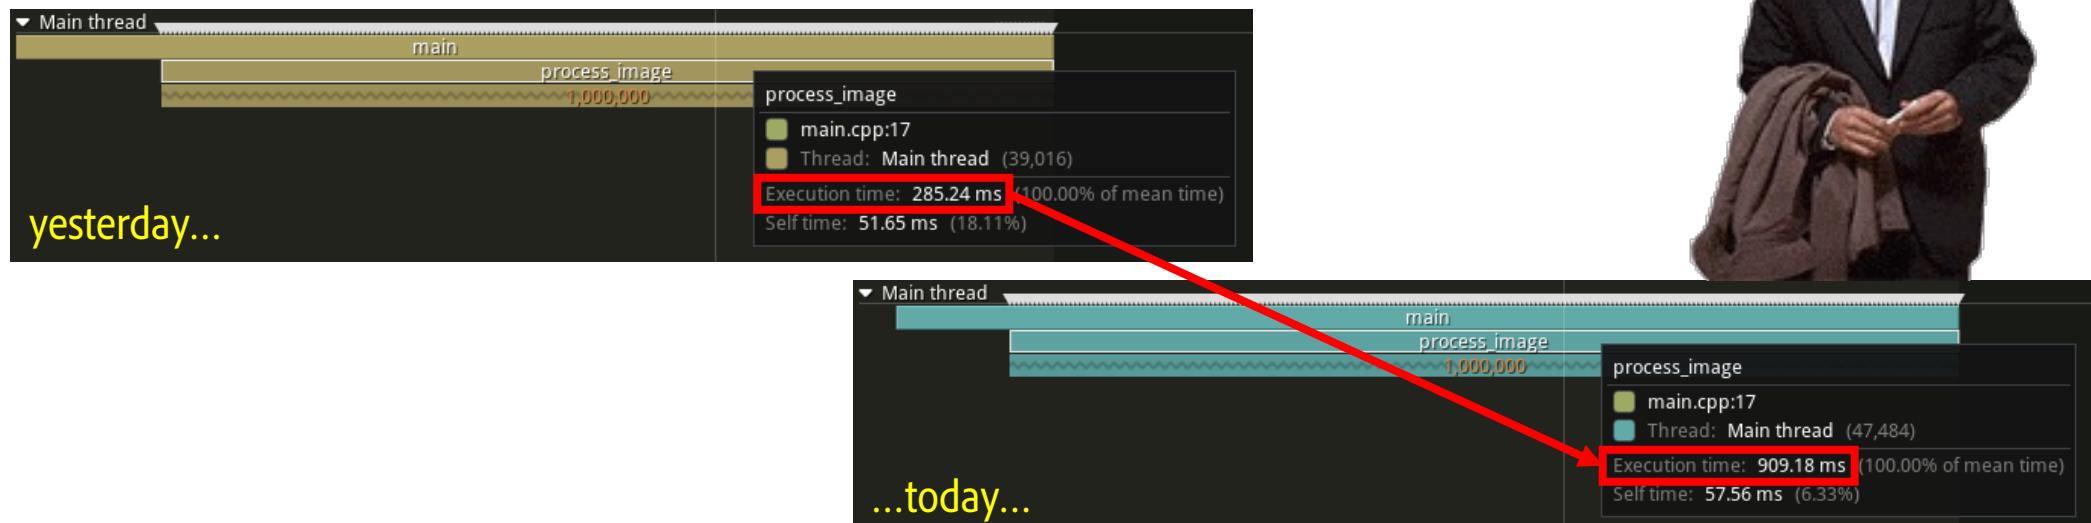
```
bool Viewport::load_scene(const char* filename)
{
    ZoneScoped;
    // ...
    std::shared_ptr<Scene> new_scene;
    {
        ZoneScopedN("new scene");
        new_scene = std::make_shared<Scene>(filename);
        if (!new_scene) return false;
    }
    m_scene = new_scene;
    return true;
}
```



# Gotchas

Zones are **suddenly** taking a lot **more time**

- Have you recently switched to a **Debug build**?
- Only profile **Release builds** (preferably **detached from debugger**)
- Run profiler tool on a **different machine**
- Have you since started **capturing callstacks** (too frequently)?



# Gotchas

Beware **excessive** instrumentation!

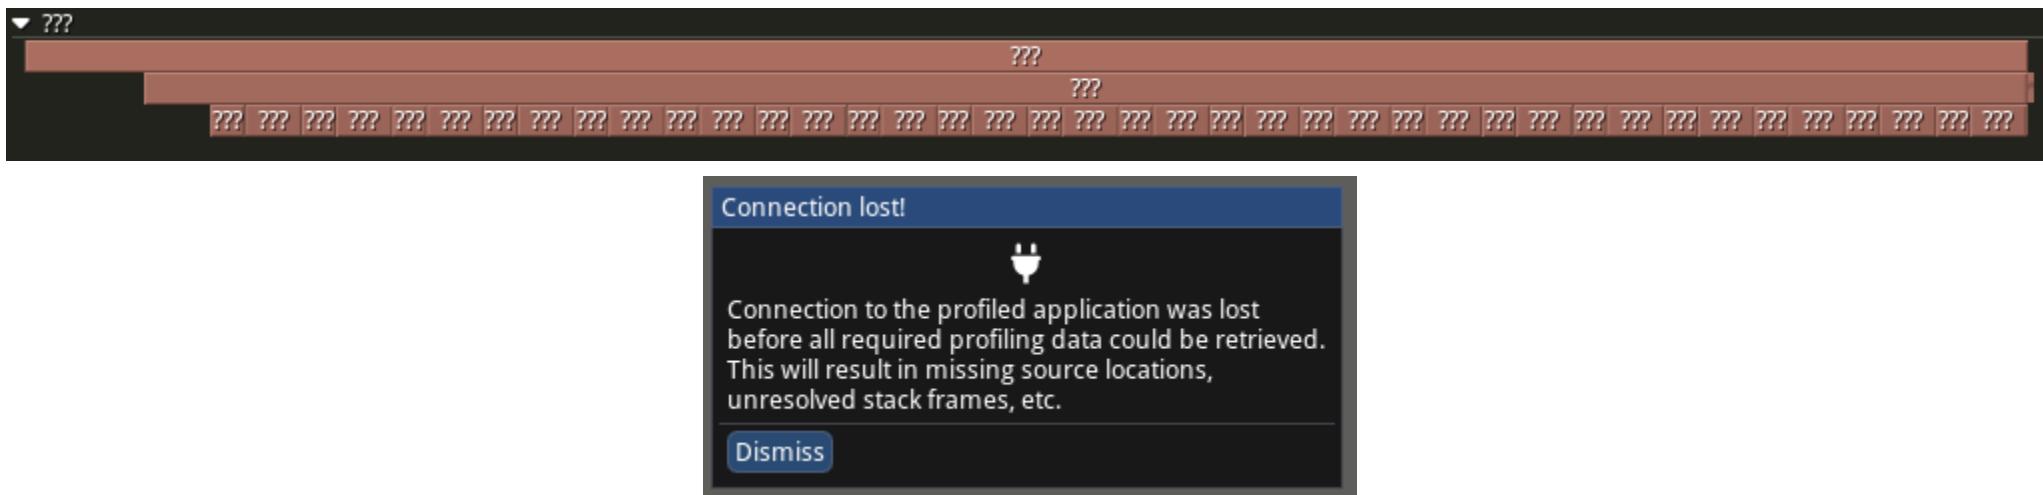
- hot **inner-loops**, small/inlined routines called frequently, ...
- **Profiling skew**: when profiling itself starts affecting performance
- May also affect **compiler optimizations**!

*skewed profiling is still useful,  
but requires a different mindset*

# Gotchas

Profiler **not showing the name** of some zones: getting **???** instead

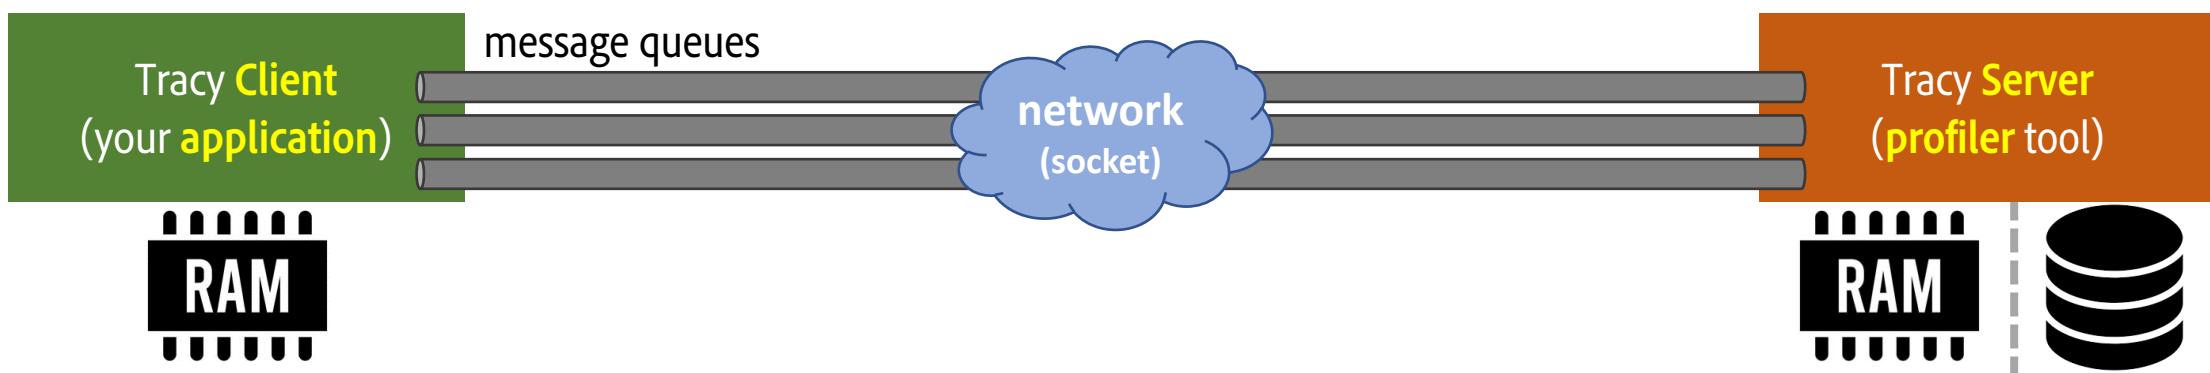
- just **wait** a little bit more (host application may not have sent this data yet)
- the host application may have **crashed** or terminated abruptly
- ensure zone name is a **persistent string**, or use **ZoneTransient()**
- if zone belongs to **plugin code**, make sure the plugin has **not been unloaded**



# Gotchas

## High RAM Usage

- Tracy **never drops** events! (there's no circular/LRU event queue)
- If profiler is **not connected**, application **memory keeps growing!**
- If profiler is running on the **same machine**, it is **competing for RAM**



- `#define TRACY_ON_DEMAND` to delay producer until profiler gets connected
- Use **capture** utility to “**sink**” Tracy events to **disk**
- Consider running the profiler GUI on a **different machine**

# Study Cases

- [Photoshop] Win32 API Detours
- [Photoshop] Multi-threading opportunity
- [Photoshop] Eliminating redundant computations
- [Photoshop] Addressing lock contention
- [Halide] 100x faster HLSL code generation!
- [Halide] Yeah... std::stringstream is kinda slow...
- [Modeler] Improved parallel shader compilation
- [Éclair] Avoiding GPU<->CPU copies with GPU interop

halide / Halide

Type / to search

Code Issues Pull requests Discussions Actions Projects Wiki Security Insights

# Making HLSL code-gen a couple orders of magnitude faster... #7719

Merged steven-johnson merged 1 commit into main from slomp/hlsl-codegen-speedup on Jul 28

Conversation 15

Commits 1

Checks 3

Files changed 1

+0 -1



slomp commented on Jul 27 • edited

Member

...

Title says it all!

The redundant call to print\_expr() ends up triggering an exponential chain of events that ultimately gets discarded.

Reviewers

shoaibkamil ✓

steven-johnson ✓

abadams ●

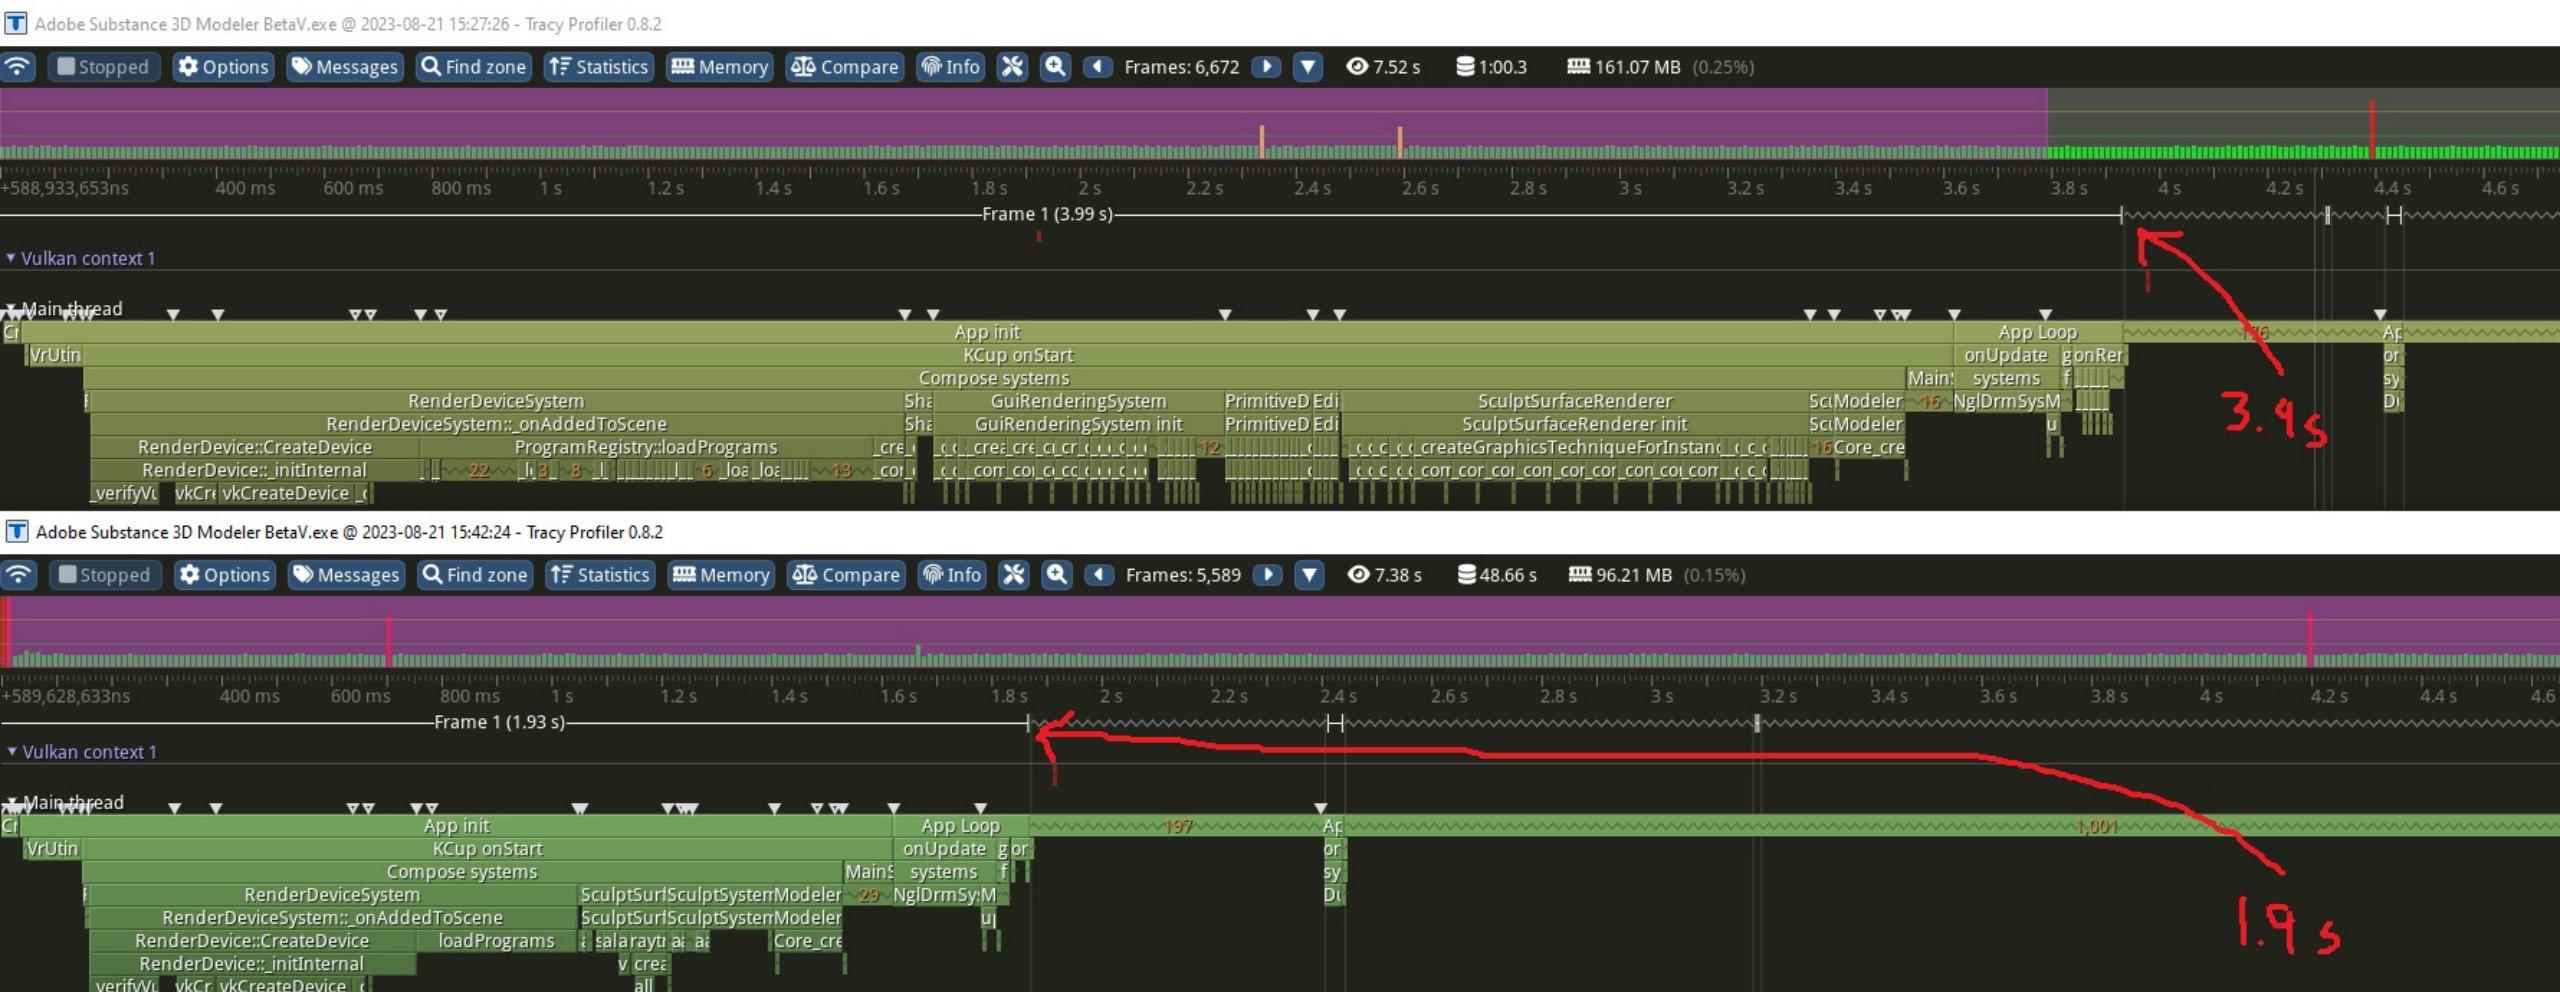
1 src/CodeGen\_D3D12Compute\_Dev.cpp

Viewed

...

```
@@ -947,7 +947,6 @@ string CodeGen_D3D12Compute_Dev::CodeGen_D3D12Compute_C::print_cast(Type target_
947     void CodeGen_D3D12Compute_Dev::CodeGen_D3D12Compute_C::visit(const Cast *op) {
948         Type target_type = op->type;
949         Type source_type = op->value.type();
950 -        string value_expr = print_expr(op->value);
951
952         string cast_expr = print_cast(target_type, source_type, print_expr(op->value));
953     }
```

# Recent story (Adobe Substance 3D Modeler)



# A Taste of Tracy's Sampling Mode

**TracyClient.cpp** is all you need!

then launch your app with **admin privileges**  
and connect it to the Tracy Profiler

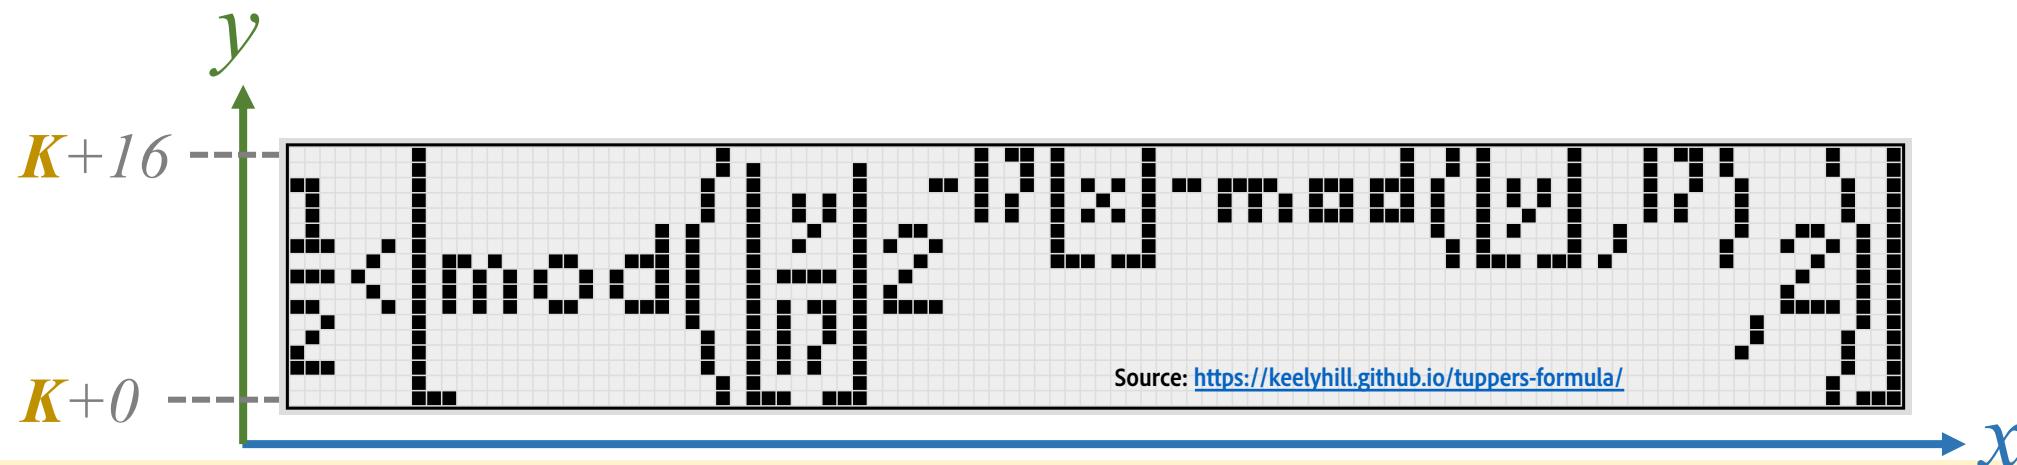
Manual instrumentation is **not necessary!**  
(but you can have **both** coexist **together!**)



# A Taste of Tracy's Sampling Mode

## Tupper's Self-Referential Formula

$$\frac{1}{2} < \left\lfloor \text{mod} \left( \left\lfloor \frac{y}{17} \right\rfloor 2^{-17 \lfloor x \rfloor - \text{mod}(\lfloor y \rfloor, 17)}, 2 \right) \right\rfloor,$$



K =

48584506361897134235820959624942020445814005879832445494830930850619347047088000284506447698655243648499972470249151191104116057391774078569197543265718  
55442057210445735883681829823754139634338225199452191651284348332905131031913302413758765239264874613394906870130562295813219481113685339535565290850  
0238750928568926945559742815463865107300491067230589335860525440966633126534936364395712556569593681518433485760526694016125126695142155053955451915378  
5457525756590740540157929001765967965480064427829131488548259914721248506352686630476300

544 digits

# A New Speedrun Category!

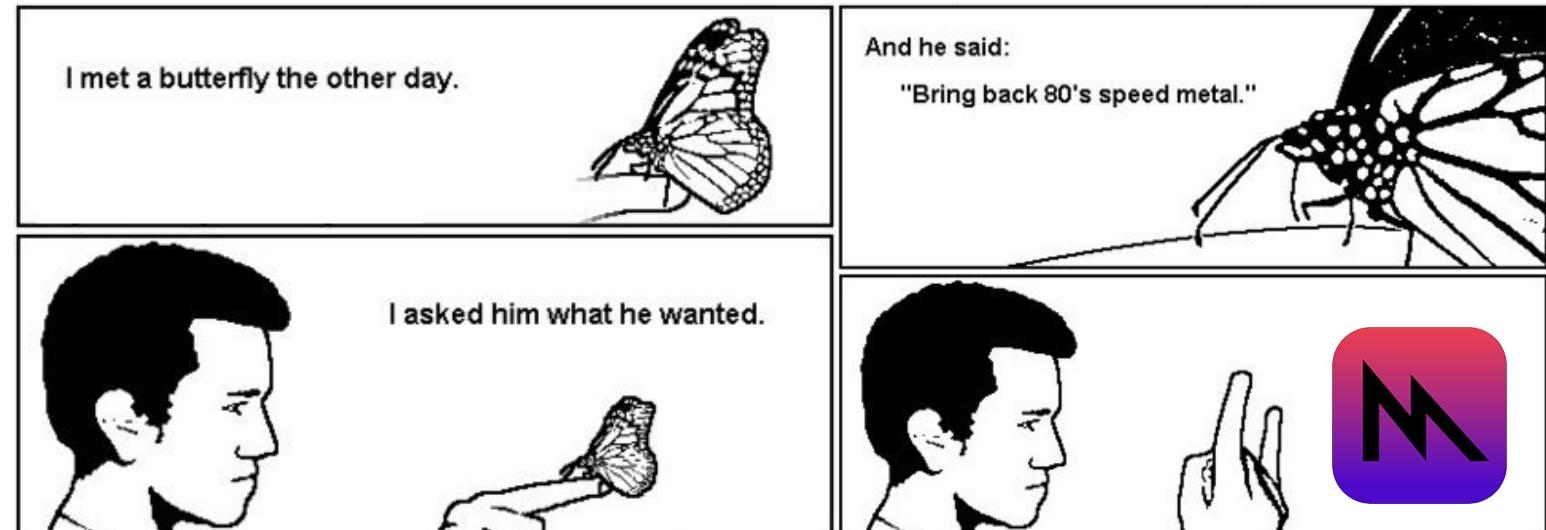
## Quake 2 Demo First Frame Any%

World Record: Marcos Slomp 1.32s

# Known Issues

## macOS/iOS Limitations

- Profiler may get **interrupted** when **application finishes**
  - It's Apple, not Tracy! Ongoing workaround discussion: <https://github.com/wolfpld/tracy/issues/8>
- Vertical scrollbar missing on main profiler window
  - It's ImGUI, not Tracy!
- No GPU **Metal API instrumentation** yet
  - I'm working on it!



# Known Issues

Program **hangs** when calling **exit()**

- Try **TRACY\_MANUAL\_LIFETIME**
- or just **don't call exit()** or similar :-)



imgflip.com

# Known Issues

Tracy **capture files** are **not backwards compatible**

- Try using the **update** utility to convert trace files

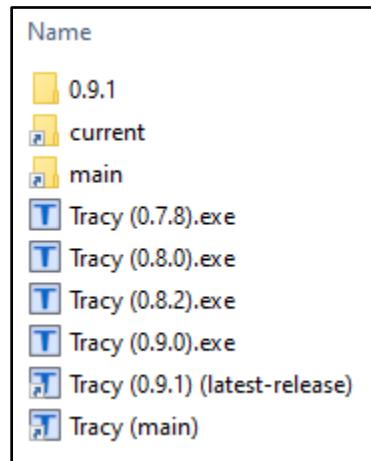
download pre-built binary (Windows-only):

<https://github.com/wolfpld/tracy/releases>

or build it yourself:

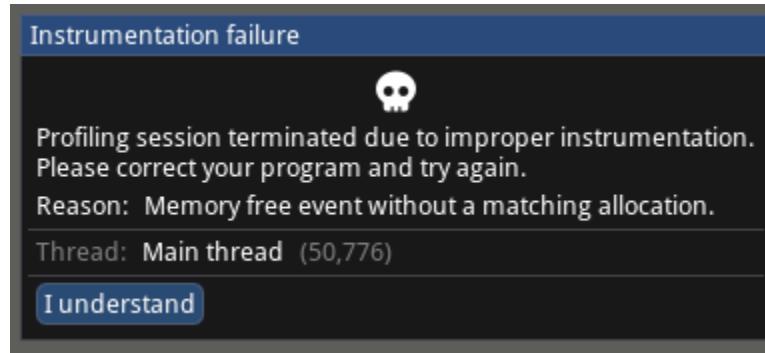
<tracy-repo>/update/build

- or **keep older versions** of Tracy around



# Known Issues

TracyFree[N|S]() may **disconnect** profiler if **no matching allocation** exists

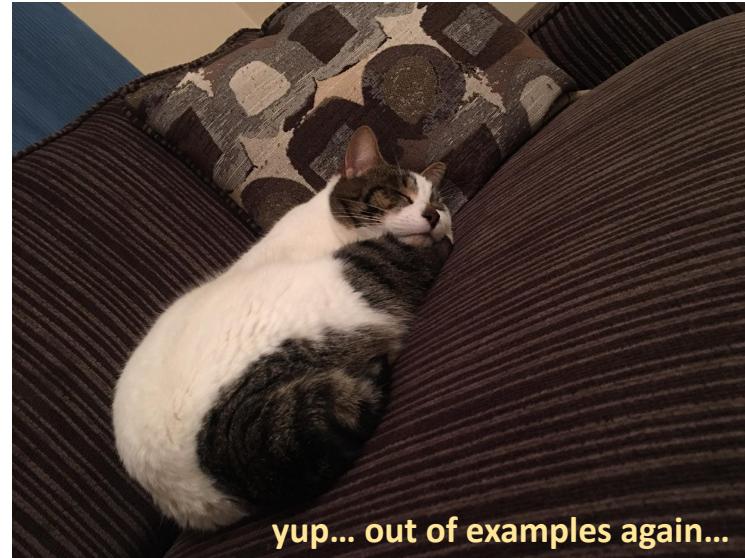


- Workaround:
  - force `m_ignoreMemFreeFaults` to `true` in `TracyWorker.cpp`
  - then make sure to `rebuild` the profiler tool

# Known Issues

My program **won't connect** to the profiler

- may be the case for **short-lived** programs
- ... that said, you may `#define TRACY_NO_EXIT` in the host program (client) to **prevent** it from exiting until a **connection** to the profiler is established



yup... out of examples again...

# More Limitations

## 2.1.8 Limitations

When using Tracy Profiler, keep in mind the following requirements:

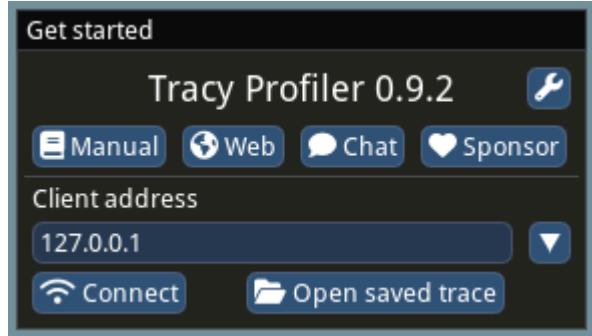
- The application may use each lock in no more than 64 unique threads.
- There can be no more than 65534 unique source locations<sup>16</sup>. This number is further split in half between native code source locations and dynamic source locations (for example, when Lua instrumentation is used).
- If there are recursive zones at any point in a zone stack, each unique zone source location should not appear more than 255 times.
- Profiling session cannot be longer than 1.6 days ( $2^{47}$  ns). This also includes on-demand sessions.
- No more than 4 billion ( $2^{32}$ ) memory free events may be recorded.
- No more than 16 million ( $2^{24}$ ) unique call stacks can be captured.

also: **callstack depth** limited to **about 60** stack frames

The following conditions also need to apply but don't trouble yourself with them too much. You would probably already know if you'd be breaking any.

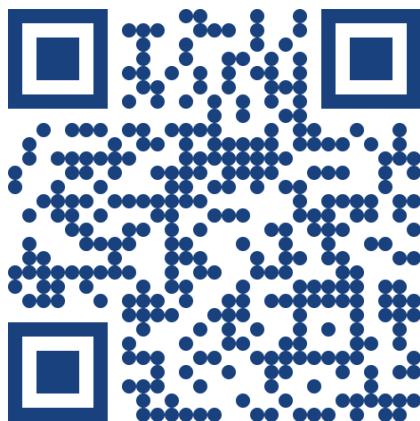
- Only little-endian CPUs are supported.
- Virtual address space must be limited to 48 bits.
- Tracy server requires CPU which can handle misaligned memory accesses.

# Help



RTTM: Read the Tracy Manual (it's *really good!*)  
tracy.pdf (grab it from <https://github.com/wolfpld/tracy/releases>)

Discord server



<https://discord.gg/pk78auc>

# C++ at Adobe!

[developer.adobe.com/cpp](https://developer.adobe.com/cpp)



- Careers
- Events
- Training
- Blog Posts



[research.adobe.com](https://research.adobe.com)

- Careers
- Internships
- Events
- Blog Posts



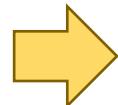
# Tips & Tricks

## Instrumenting **dynamic dispatch** (e.g., virtual calls)

Need to **instrument all overrides...**

Consider using **static delegate** interface to call the **virtual** method

```
class Shape
{
public:
    virtual float area();
};
```



```
class Shape
{
public:
    float area() {
        ZoneScoped;
        return calculate_area();
    }
protected:
    virtual float calculate_area();
};
```

# Tips & Tricks

Begin/End (and track) **zone** on **different scopes** or **across threads**

e.g.: task starts in main thread, moves between worker threads, and gets retired on a reclaimer thread

- A few **palliative** workarounds:
  - Use `FrameMarkStart` and `FrameMarkEnd`
  - Use “matching” `TracyMessage`’s
  - Use `clock/sustain` `TracyPlot`
  - Assign unique `ZoneColor` to each task
- My recommendation: use `TracyAlloc` and `TracyFree`  
(it’s basically a **stackable clock/sustain plot!**)

# Tracy C API

Closing an outer zone inside an inner zone is bad!

- Tracy C-style API:
  - **explicit** scoping: must **store/carry** scope context object around
  - `#include <TracyC.h>`
  - `struct TracyCZoneCtx`
  - `TracyCZone[N|C|S](ctx, ..., active)` : start/begin a zone
  - `TracyCZoneEnd(ctx)` : zone **MUST end** on the **SAME thread** that started it!
  - `TracyCZoneColor(ctx, ...)`
  - `TracyCZoneValue(ctx, ...)`
  - `TracyCZoneText(ctx, ...)`
  - `TracyCZoneName(ctx, ...)`



scoped zones are almost always enough  
(and intuitive, and have less overhead)