

+ 23

Designing Fast and Efficient List-like Data Structures

YANNIC BONENBERGER



20
23



List-like data structures

- `std::vector`
- `std::list`
- `std::deque`

std::vector

- C++ version of the array-list data structure
- Backed by a C-style array
- Automatically allocates a new backing array when inserting into a "full" std::vector

std::vector

```
1 template <class T>
2 class vector {
3     public:
4         // constructor, accessors, ...
5
6     private:
7         size_t size_;
8         size_t capacity_;
9         T* data_;
10 };
```

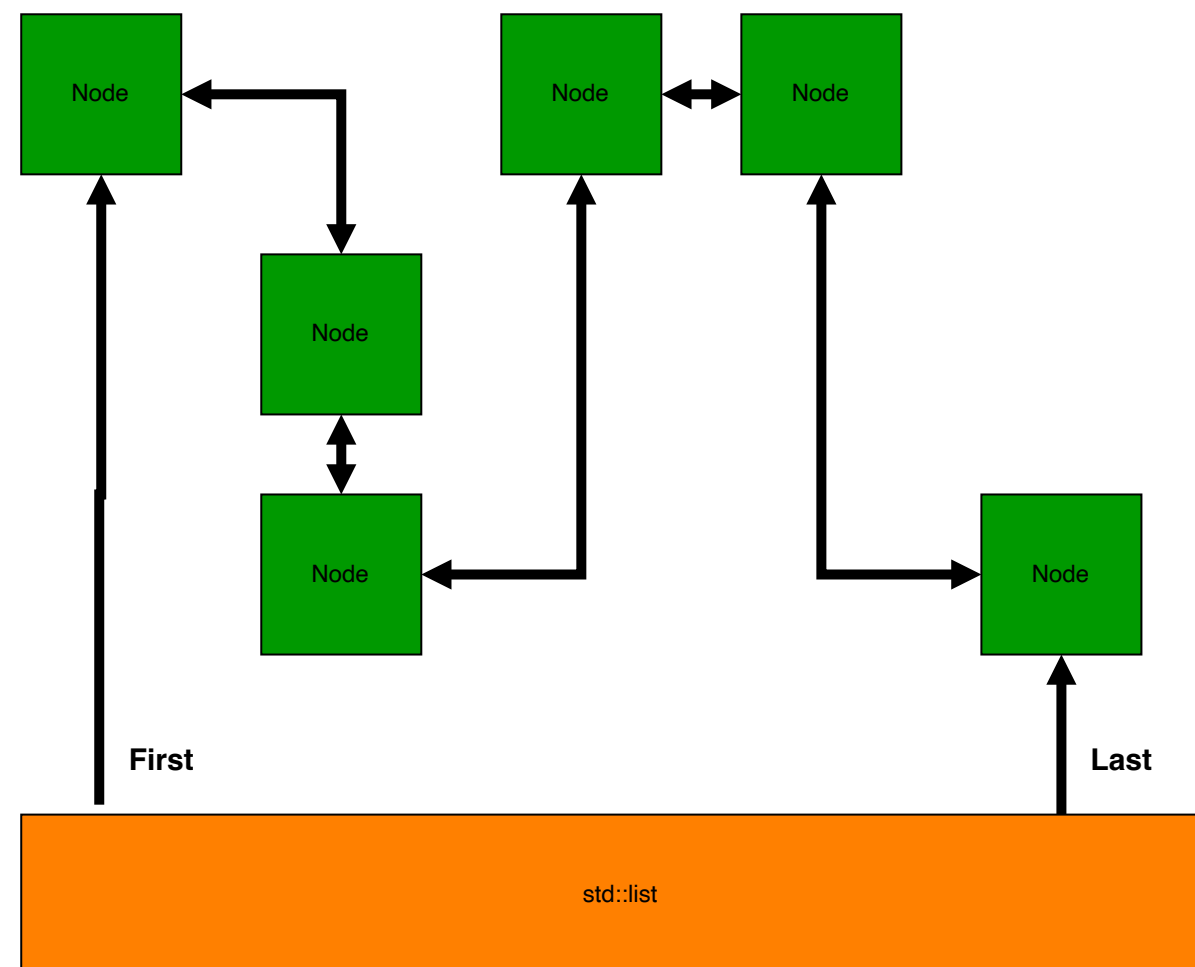
`std::vector`

- Compact layout – elements stored in contiguous memory
- Inserting or removing elements at the end is very cheap ($O(1)$)
- Accessing elements in random order is very cheap ($O(1)$)
- Inserting or removing elements at positions other than the end is expensive ($O(n)$)

`std::list`

- C++ version of the linked-list data structure
- Elements stored in nodes referencing the next node

std::list



`std::list`

- Inserting or removing elements in the front or at the end is very cheap ($O(1)$)
- Accessing elements in random order is expensive ($O(n)$)


`std::deque`

- Conceptually a table of vectors

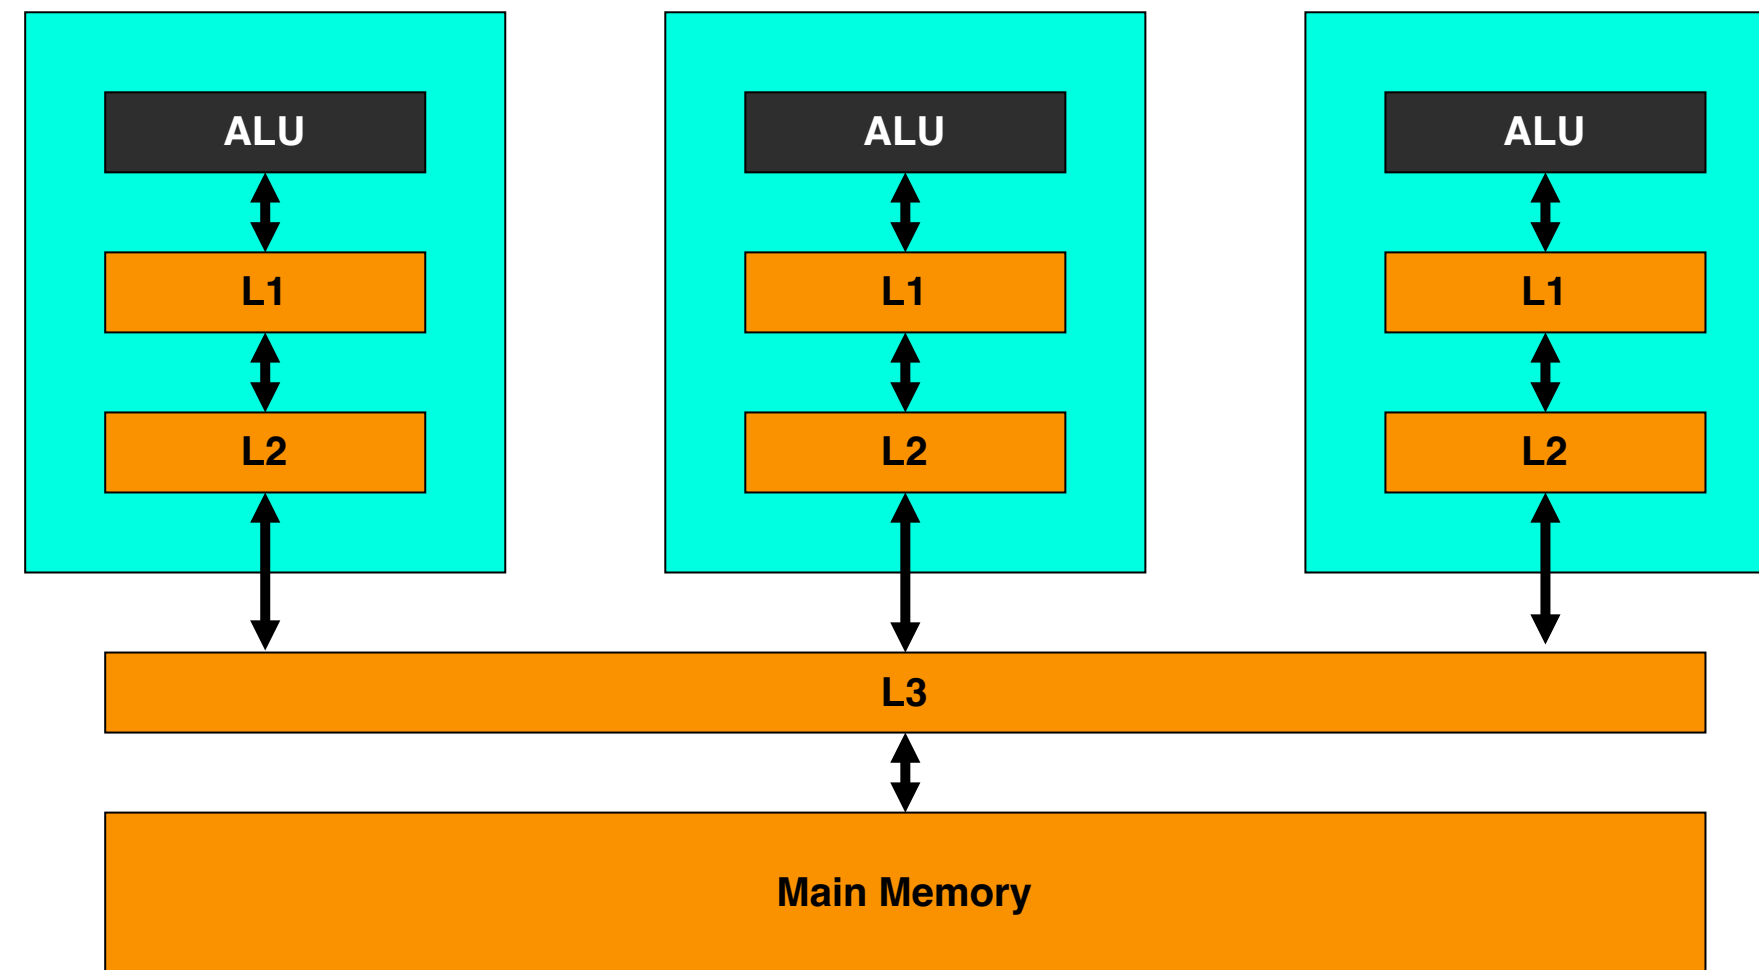
`std::deque`

- Stores elements in chunks of contiguous memory
- Inserting or removing elements in the front or at the end very cheap ($O(1)$)
- Accessing elements in random order is cheap ($O(1)$)
- Inserting or removing elements at positions other than the front or end is expensive ($O(n)$)
- Requires additional indirection for every access

Modern CPU architectures



Modern CPU architectures



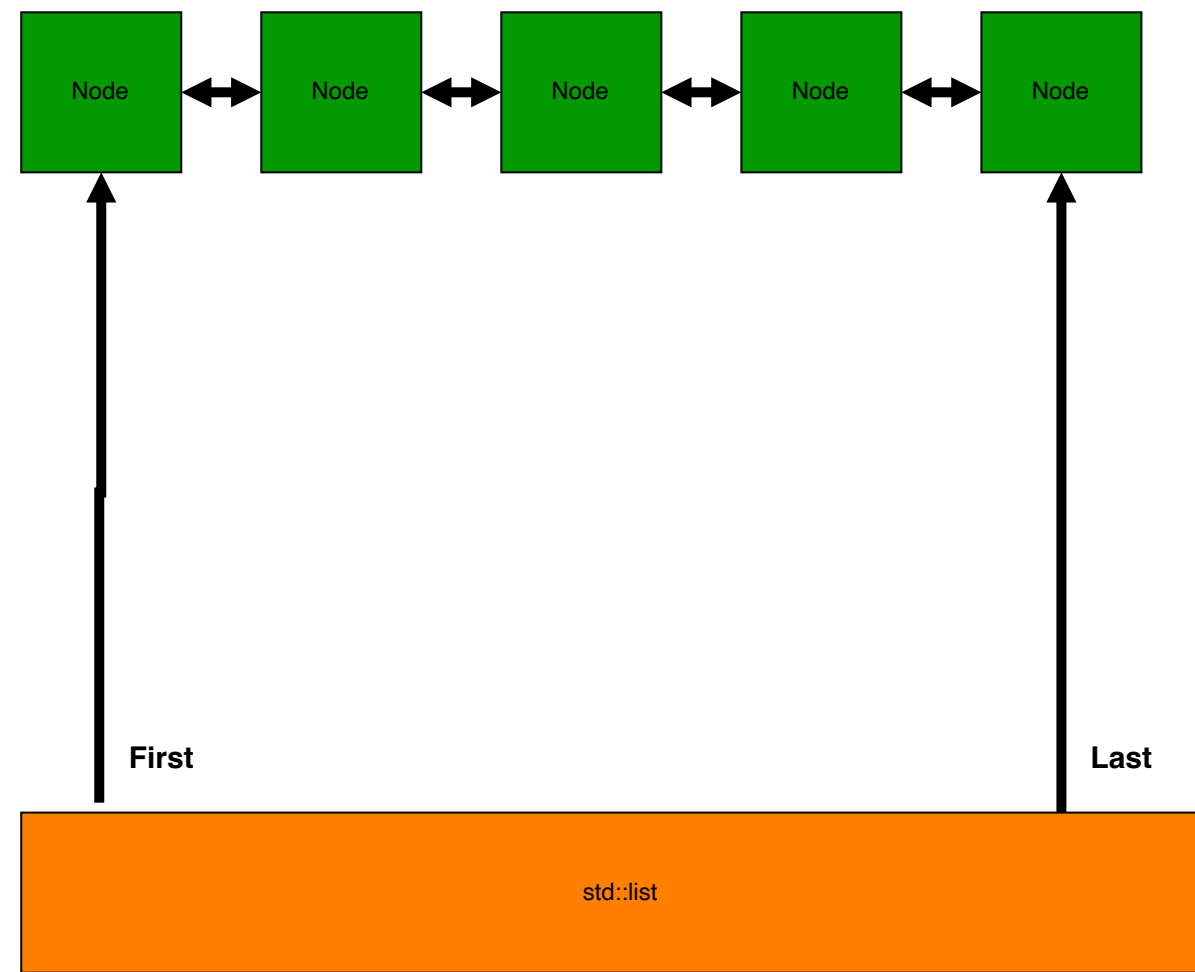
Modern CPU architectures



std::list with custom allocator

- Idea: minimize cache misses by grouping consecutive nodes into contiguous memory

std::list with custom allocator



std::list with custom allocator

- Packing sequences of nodes into a contiguous chunk of memory provides better cache locality
- However: having a wrapper for each element containing a pointer to the previous and next node is high overhead

FixedStack

- Stack implementation with a (compile-time) fixed upper bound for the number of elements it can hold
- Completely avoids heap allocations for small amount of small objects

FixedStack

```
1 template <class T, size_t capacity>
2 class FixedStack {
3     public:
4         FixedStack& Push(const T& element);
5         ElementContainer Pop();
6
7     private:
8         Page<T, capacity> data_;
9 };
```

Memory Pages

```
1 template <class T, size_t page_size>
2 class Page {
3     public:
4         T& Front();
5         T& Back();
6         T& At(size_t position);
7
8         Slot<T>::ElementContainer RemoveFront();
9         Slot<T>::ElementContainer RemoveBack();
10
11     private:
12         size_t front_slot_;
13         size_t back_slot_;
14         Slot<T> slots_[page_size];
15 };
```

Memory Slots

```
1 template <class T>
2 using Slot = SlotImpl<
3     T,
4     (sizeof(T) <= (ABSL_CACHELINE_SIZE / 8)) ||
5     ((sizeof(T) <= (ABSL_CACHELINE_SIZE / 2)) &&
6     ((ABSL_CACHELINE_SIZE % sizeof(T)) == 0))>;
```

Memory Slots

```
1 template <class T>
2 class alignas(T) SlotImpl<T, true> {
3     public:
4         // constructor, accessors, ...
5
6     private:
7         alignas(T) char storage_[sizeof(T)];
8 };
```

Memory Slots

```
1 template <class T>
2 class alignas(T) SlotImpl<T, false> {
3     public:
4         // constructor, accessors, ...
5
6     private:
7         T* storage_;
8 };
```

Queue

```
1 template <class T, size_t page_size = 128>
2 class Queue {
3     public:
4         Queue& Push(const T& element);
5         ElementContainer Pop();
6
7     private:
8         PageList<T, page_size> pages_;
9         PageList* first_;
10        PageList* last_;
11 };
```

Real-World use-case of Queue

```
1 class TaskQueue {
2     public:
3         Add(OnceCallback<void()> task) {
4             task_queue_.Push(std::move(task));
5         }
6
7     private:
8         void Run() {
9             while (true) {
10                 std::move(task_queue_.Pop()).Run();
11             }
12         }
13
14         ThreadSafeQueue<OnceCallback<void()>> task_queue_;
15 };
```


Outlook and Future Work

Early results show that TaskQueue can pass tasks between threads with 50% * less overhead than `std::deque` and avoids most of the synchronization needed for `std::deque`.

* Measured on a MacBook Pro, Apple M1 Max with 64GiB of memory

Outlook and Future Work

However: Especially the dynamic-sized and thread-safe containers including TaskQueue are still experimental.

Outlook and Future Work

Publishing on GitHub still in progress. Planned for later this year.

Questions?

Thank You!