

+ 23

Iteration Revisited

TRISTAN BRINDLE



Cppcon
The C++ Conference

20
23



October 01 - 06

ITERATION BASICS

```
int arr[] = {1, 2, 3, 4, 5};  
size_t sz = 5;  
  
for (size_t i = 0; i < sz; i++) {  
    do_something(arr[i]);  
}
```

ITERATION BASICS

```
for (size_t i = 0; i < sz; i++) {  
    do_something(arr[i]);  
}
```

ITERATION BASICS

```
for (int* p = arr; p != arr + sz; ++p) {  
    do_something(*p);  
}
```

ITERATION BASICS

```
for (int* p = std::begin(arr);  
     p != std::end(arr);  
     p++) {  
    do_something(*p);  
}
```

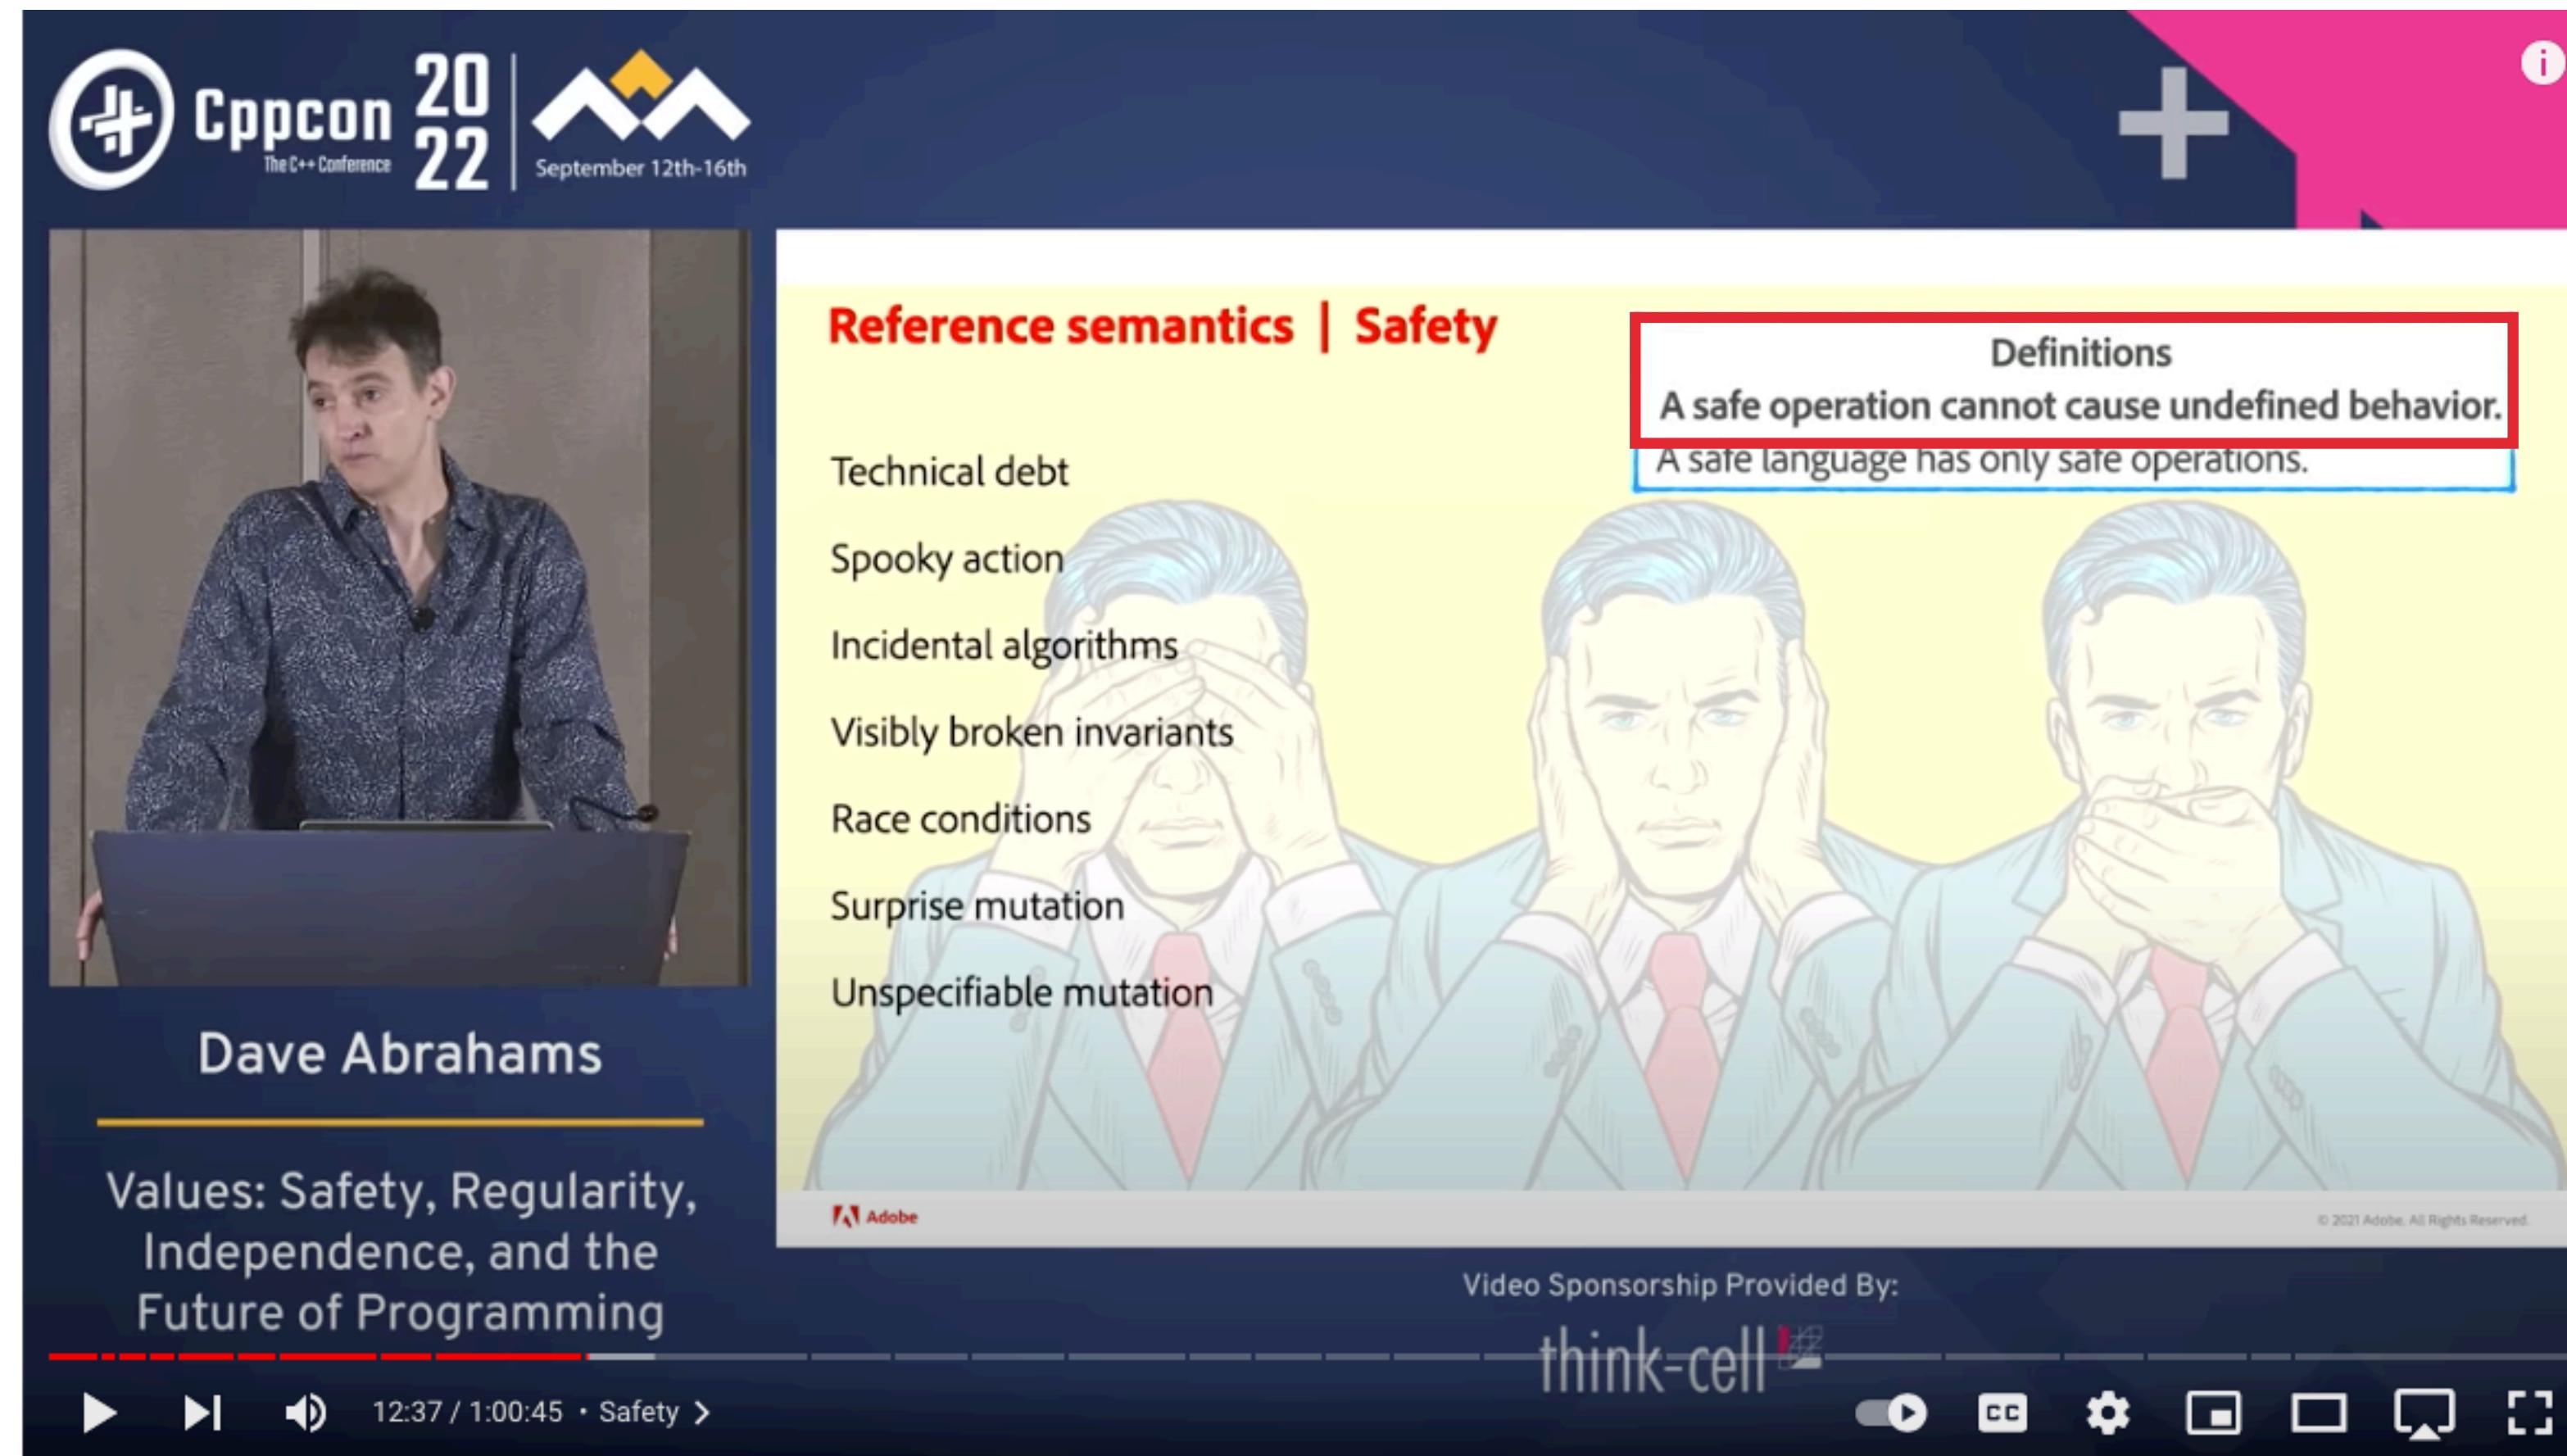
ITERATION BASICS

```
for (auto iter = std::begin(arr);  
     iter != std::end(arr);  
     ++iter) {  
    do_something(*iter);  
}
```

STL ITERATORS

- ▶ STL iterators are a *generalisation* of array pointers
- ▶ Advantages:
 - ▶ Powerful
 - ▶ Low overhead
 - ▶ “Natural” syntax
- ▶ Disadvantages:
 - ▶ Iterators are **just as unsafe** as raw pointers

WHAT DO I MEAN BY SAFETY?



Value Semantics: Safety, Independence, Projection, & Future of Programming - Dave Abrahams CppCon 22

A SAFE OPERATION CANNOT CAUSE
UNDEFINED BEHAVIOUR.

Dave Abrahams

ITERATORS AND UB

```
auto iter = std::end(rng);  
do_something(*iter);
```

- ▶ *Dereferencing a past-the-end iterator is UB*

ITERATORS AND UB

```
auto iter = std::end(rng);  
++iter;
```

- ▶ *Incrementing a past-the-end iterator is UB*
 - ▶ So is decrementing a begin() iterator
 - ▶ So are out-of-bounds random-access jumps

ITERATORS AND UB

```
auto min_iter(const std::vector<int>& vec)
{
    return std::min_element(vec.begin(), vec.end());
}

auto iter = min_iter(get_vector());
std::cout << "Minimum is " << *iter << std::endl;
```

- ▶ **Dangling iterators** can occur easily
- ▶ Any operation on a dangling iterator is UB

ITERATORS AND UB

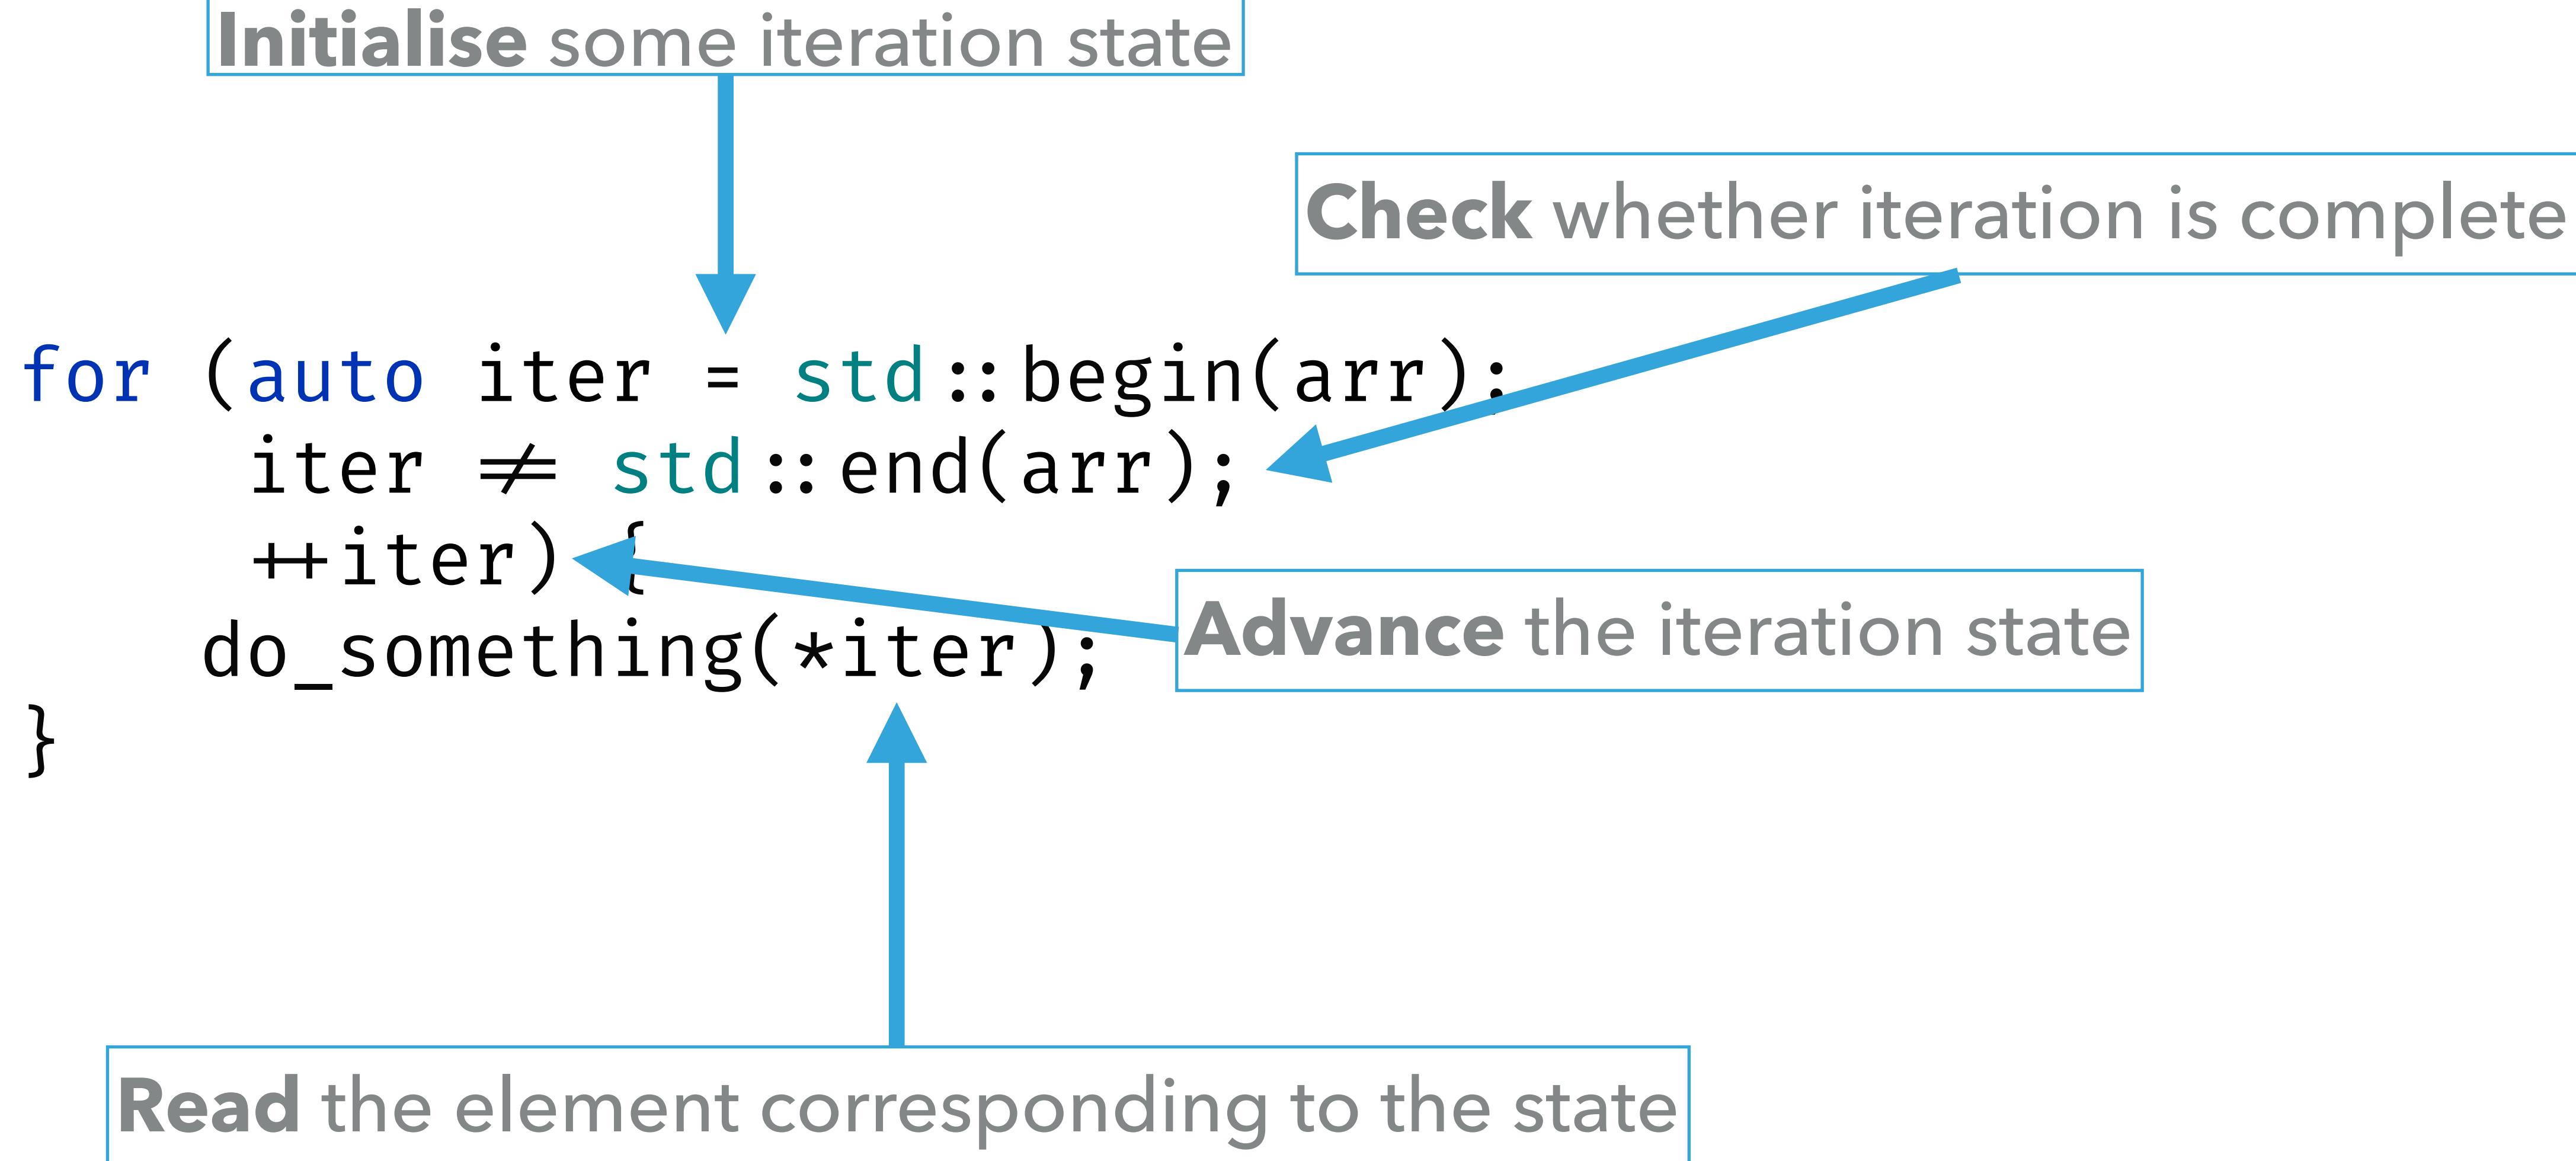
```
auto iter = vec.begin();  
  
vec.push_back(99);  
  
while (iter != vec.end()) {  
    ...  
}
```

- ▶ Iterators can become dangling in non-obvious ways
 - ▶ All part of the C++ experience!

ITERATORS AND UB

- ▶ “A safe operation cannot cause undefined behaviour”
- ▶ Most iterator operations can cause UB
- ▶ => STL iterators are a *fundamentally unsafe* abstraction

ITERATION FUNDAMENTALS



ITERATION FUNDAMENTALS

- ▶ **Initialise** some iteration state
- ▶ **Check** whether iteration is complete
- ▶ **Read** the element corresponding to the state
- ▶ **Advance** the iteration state



Rust iterators:
`next() -> optional<T>`

ITERATION FUNDAMENTALS

Iterators and Ranges: Comparing C++ to D, Rust, and Others

CPPP 2021 PARIS

C++ Barry Revzin MUREX CODE RECKONS

Implementing map in C++ (transform)

```
template <input_range V, copy_constructible F>
    requires view<V> &&
        regular_invocable<F&, range_reference_t<V>>
class map_view {
    struct Iterator;
    struct Sentinel;

    V base_;
    F fun_;

public:
    map_view(V, F);
    auto begin() -> Iterator;
    auto end() -> Sentinel;
};
```

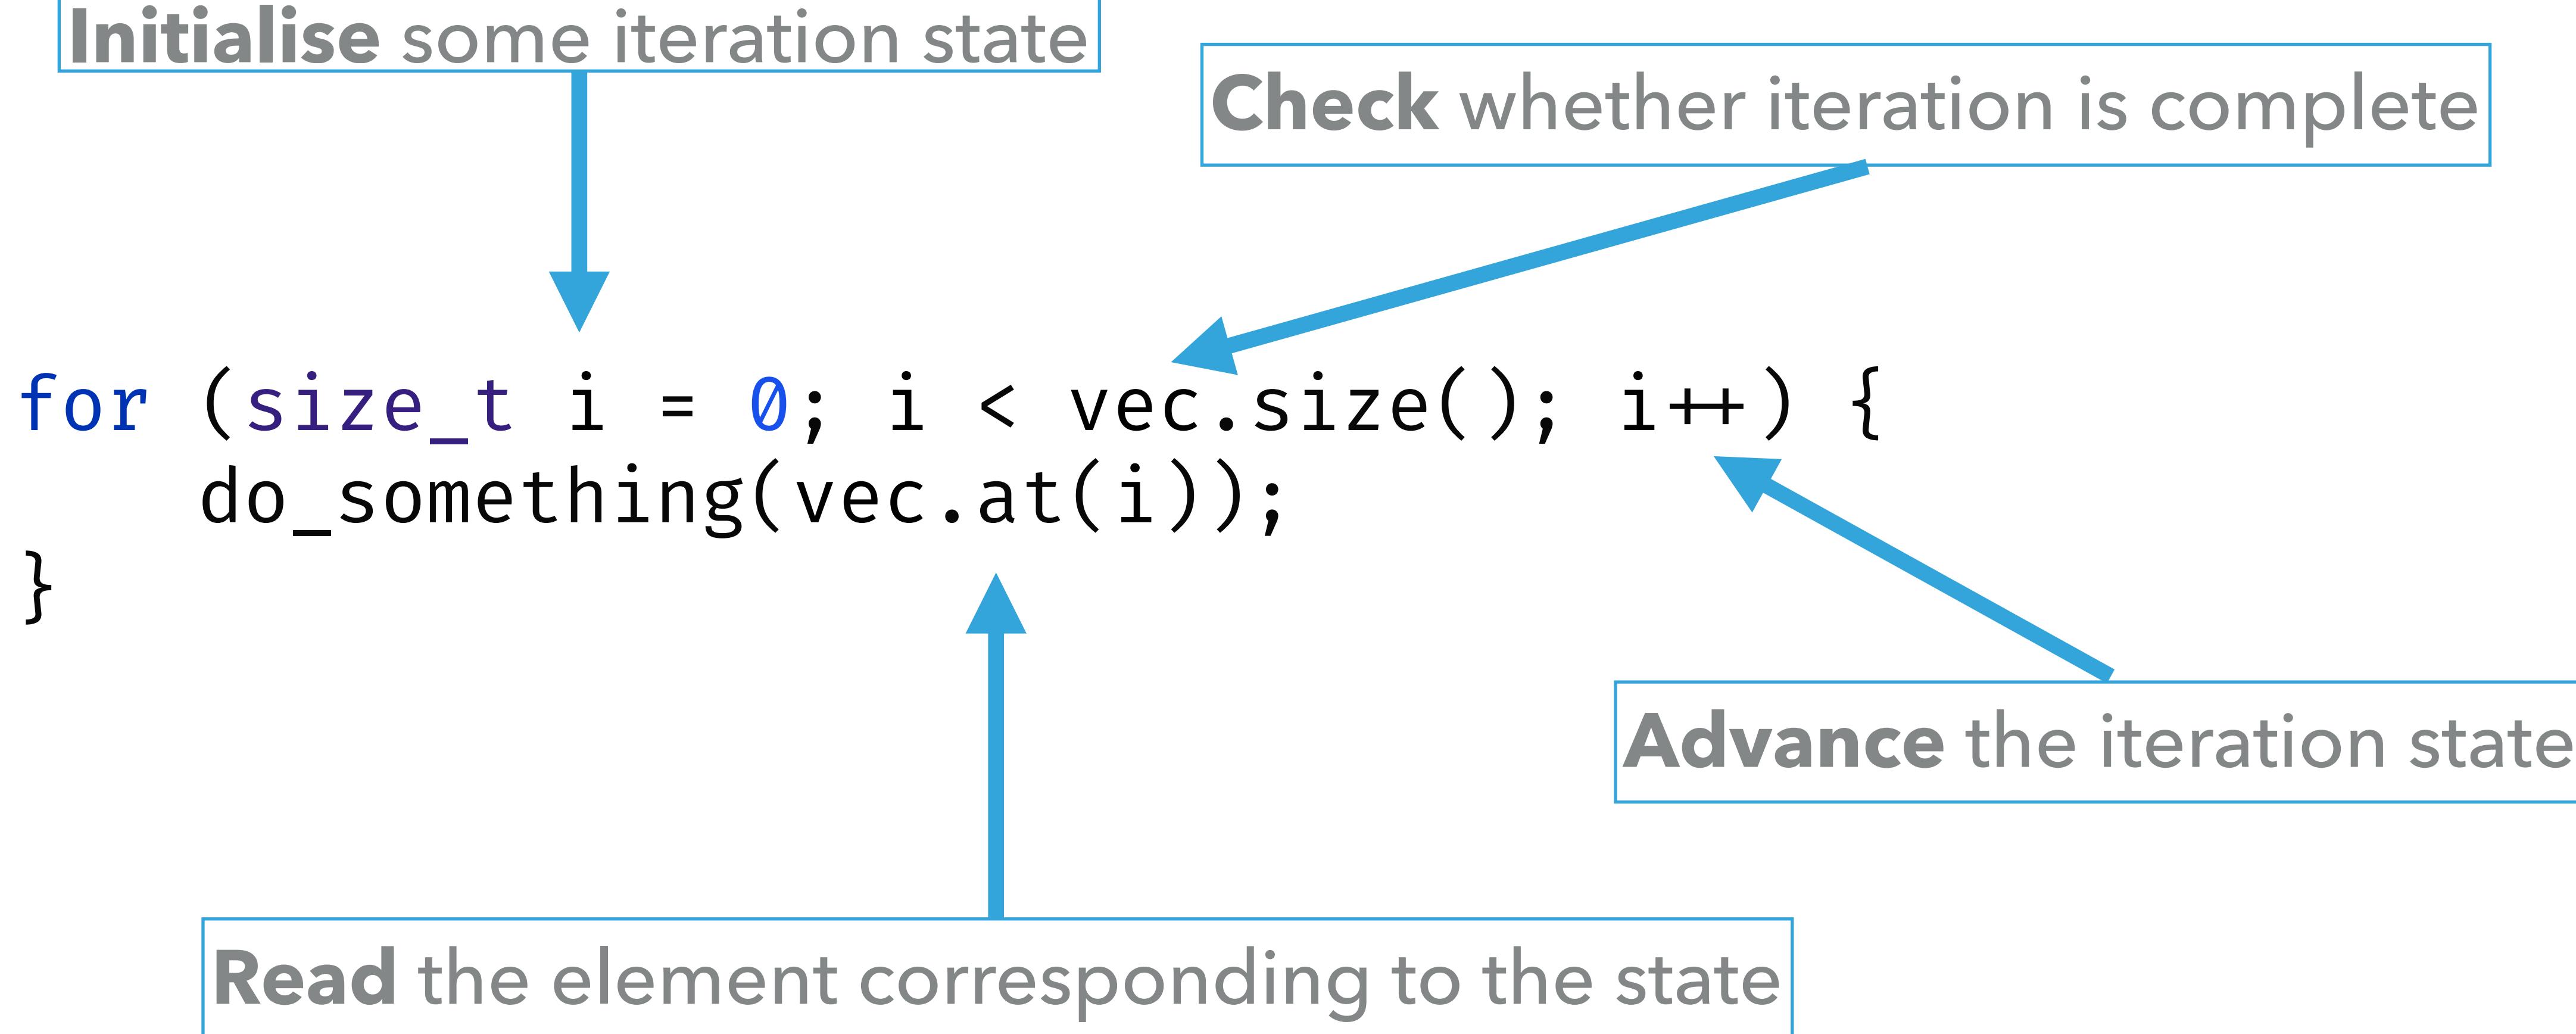
8:24 / 1:31:27

Keynote: Iterators and Ranges: Comparing C++ to D, Rust, and Others - Barry Revzin - CPPP 2021

ITERATION WITH INDICES

```
std::vector<int> vec{1, 2, 3, 4, 5};  
  
for (size_t i = 0; i < vec.size(); i++) {  
    do_something(vec.at(i));  
}
```

ITERATION WITH INDICES



ITERATION WITH INDICES

```
auto idx = std::size(vec);  
do_something(vec.at(idx));
```

- ▶ Reading via a past-the-end index is bounds checked
- ▶ No UB!

ITERATION WITH INDICES

```
auto idx = std::size(vec);  
++idx;
```

- ▶ *Incrementing a past-the-end index is fine*
- ▶ *(Assuming unsigned indices or checked operations on signed indices)*

ITERATION WITH INDICES

```
size_t min_idx(const std::vector<int>& vec) {  
    ...  
}  
  
auto idx = min_idx(get_vector());  
// std::cout << "Minimum is " << ??? << std::endl;
```

- ▶ “Dangling indices” are safe **by design**

ITERATION WITH INDICES

```
size_t idx = 0;  
  
vec.push_back(99);  
  
while (idx < vec.size()) {  
    ...  
}
```

- ▶ “Index invalidation” problems are no longer UB
- ▶ Caught by bounds check at next attempted read

ITERATION WITH INDICES

- ▶ Index-based iteration **dramatically reduces** the potential for UB
- ▶ ...but is still **just as powerful**
- ▶ The Big Idea: generalise iteration by *index*, rather than by *pointer*

FLUX

- ▶ Flux (github.com/tcbrindle/flux) is a C++20 library exploring these ideas
- ▶ Aims:
 - ▶ Much improved safety by default vs STL iterators
 - ▶ Improved ease-of-use and fewer “gotchas”
 - ▶ Equal or better runtime performance vs C++20 ranges
 - ▶ Compatibility with existing STL code

ITERATION OPERATIONS WITH FLUX

- ▶ `flux :: first(seq)` returns a *cursor*, which represents a sequence position
 - ▶ For contiguous sequences, this is just an integer index
- ▶ `flux :: is_last(seq, cur)` returns whether the cursor is past-the-end
- ▶ `flux :: read_at(seq, cur)` accesses the element at the cursor position
- ▶ `flux :: inc(seq, cur)` increments the cursor

FLUX CONCEPTS

- ▶ A type which provides these four operations models the **sequence** concept
- ▶ Sequences are single-pass by default
- ▶ Multipass sequences are those which allow simultaneous iteration by independent cursors
- ▶ A sequence whose cursor is a *regular type* is assumed to be multipass
 - ▶ Regular = default constructible, copyable, equality comparable
 - ▶ Sequences can opt out of being multipass if necessary

FLUX CONCEPTS

- ▶ *Bidirectional sequences* are multipass sequences which additionally allow their cursors to be decremented:
 - ▶ `flux::dec(seq, cur)`
- ▶ *Random-access sequences* provide additional operations:
 - ▶ `flux::inc(seq, cur, offset)`
 - ▶ `flux::distance(seq, from_cur, to_cur)`
 - ▶ Cursors are `std::totally_ordered`
- ▶ *Contiguous sequences* are random-access, backed by an in-memory array

FLUX SEQUENCE ALGORITHMS

► As of October 2023

all
any
compare
contains
count
count_eq
count_if
contains
ends_with
equal
fill
find
find_if
find_if_not

find_max
find_min
find_minmax
fold
fold_first
for_each
for_each_while
inplace_reverse
max
min
minmax
none
output_to
product

search
sort
starts_with
sum
swap_elements
to
write_to
zip_find_if
zip_fold
zip_for_each
zip_for_each_while

FLUX PIPELINES

- ▶ Like C++20 Ranges, Flux provides a variety of *sequence adaptors* performing lazy evaluation
- ▶ Typically these are composed into pipelines
- ▶ A pipeline begins with some sequence (the *source*), passes through zero or more *adaptors*, and then finishes with a call to an *algorithm*

FLUX PIPELINE EXAMPLE

```
std::vector<int> vec = get_vector();

int max_even_sq = flux::ref(vec)           // iterate over a reference to vec
    .filter(flux::pred::even) // keep only even numbers
    .map([](int i) { return i * i; }) // square each int
    .max()                      // find the maximum value
    .value_or(0);               // handle empty sequences
```

FLUX PIPELINE EXAMPLE

```
std::vector<int> vec = get_vector();

int max_even_sq = flux::ref(vec)           // iterate over a reference to vec
    ._(flux::filter, flux::pred::even) // keep only even numbers
    ._(flux::map, [](int i) { return i * i; }) // square each int
    ._(flux::max)                  // find the maximum value
    .value_or(0);                 // handle empty sequences
```

FLUX SEQUENCE ADAPTORS

► As of October 2023

adjacent
adjacent_filter
adjacent_map
cache_last
cartesian_product
cartesian_product_with
chain
chunk
chunk_by
cursors
cycle
deduce
drop
drop_while

filter
flatten
map
mask
pairwise
pairwise_map
prescan
read_only
reverse
scan
scan_first
set_difference
set_intersection
set_symmetric_difference

set_union
slide
stride
split
take
take_while
unchecked
zip

IMPLEMENTING THE SEQUENCE PROTOCOL

```
template <typename T, std::size_t N>
struct my_array {
    T data[N];
};

template <typename T, std::size_t N>
struct flux::sequence_traits<my_array<T, N>> {
    using self_t = my_array<T, N>;

    static size_t first(const self_t&) { return 0; }
    static bool is_last(const self_t&, size_t cur) { return cur >= N; }
    static void inc(const self_t&, size_t& cur) { ++cur; }
    static decltype(auto) read_at(auto& self, size_t cur) {
        flux::bounds_check(cur < N);
        return self.data[cur];
    }
};
static_assert(flux::multipass_sequence<my_array<int, 10>>);
```

IMPLEMENTING THE SEQUENCE PROTOCOL

```
template <typename T, std::size_t N>
struct my_array {
    T data[N];

    struct flux_sequence_traits {
        static size_t first(const my_array&) { return 0; }
        static bool is_last(const my_array&, size_t cur) { return cur >= N; }
        static void inc(const my_array&, size_t& cur) { ++cur; }
        static decltype(auto) read_at(auto& self, size_t cur)
        {
            flux::bounds_check(cur < N);
            return self.data[cur];
        }
    };
};

static_assert(flux::multipass_sequence<my_array<int, 10>>);
```

IMPLEMENTING THE SEQUENCE PROTOCOL

- ▶ Considerably easier to implement vs STL ranges
 - ▶ Even more so for higher sequence/range categories
- ▶ No iterator/const_iterator confusion: the same cursor type is used for const and non-const access

RANGES COMPATIBILITY

- ▶ Every C++20 contiguous_range is automatically a flux::contiguous_sequence
 - ▶ Contiguous => “cheaply indexable”
 - ▶ Covers vector, array, string, string_view, span...
- ▶ Other ranges can be wrapped in a sequence implementation using flux::from_range(rng), with reduced safety guarantees
 - ▶ No worse than using ranges directly

RANGES COMPATIBILITY

- ▶ Conversely, every `flux :: sequence` provides STL-compatible iterators
 - ▶ So we can use them with C++20 algorithms and range adaptors
 - ▶ ...and range-for loops

```
template <flux :: sequence Seq>
struct sequence_iterator {
    Seq* seq;
    flux :: cursor_t<Seq> cur;

    sequence_iterator& operator++() {
        flux :: inc(*seq, cur);
        return *this;
    }

    // etc...
};
```



PERFORMANCE

- ▶ Compilers are **incredibly good** at eliding bounds checks
- ▶ Many (most?) new languages use bounds checking universally
- ▶ => Optimisers can recognise those patterns in very many cases
- ▶ For example: <https://flux.godbolt.org/z/v35dsd7sx>

Private < >

flux.godbolt.org

COMPILER EXPLORER Add... More Templates Watch C++ Weekly to learn new C++ features Sponsors intel CONAN Solid Sands Share Policies 41 Other

C++ source #1 x A + v C C++ Assembly Right: x86-64 gcc 13.1 -std=c++2b ... Assembly

#ifndef USE_FLUX

```
int count_evens(std::vector<int> const& vec)
{
    int count = 0;
    for (int i : vec) {
        if (i % 2 == 0) {
            ++count;
        }
    }
    return count;
}
```

30 } #endif

31

32 #endif

33

34

35

36 .L12:

37 jbe .L12

38 .L18:

39 movd eax, xmm1

40 je .L18

41 .L3:

42 mov ecx, DWORD PTR [rdx]

43 and ecx, 1

44 cmp ecx, 1

45 lea rcx, [rdx+4]

46 adc eax, 0

47 cmp rdi, rcx

48 .L4:

49 pxor xmm1, xmm1

50 movdqa xmm3, XMMWORD PTR .LC0[rip]

51 mov rax, rcx

52 shr rdx, 2

53 movdqa xmm2, xmm1

54 sal rdx, 4

55 add rdx, rcx

56 .L4:

57 movdqu xmm0, XMMWORD PTR [rax]

58 add rax, 16

59 pand xmm0, xmm3

60 pcmpeqd xmm0, xmm2

61 psubd xmm1, xmm0

62 cmp rax, rdx

63 jne .L4

64 movdqa xmm0, xmm1

65 mov rdx, rsi

66 psrldq xmm0, 8

67 and rdx, -4

68 paddd xmm1, xmm0

69 movdqa xmm0, xmm1

70 psrldq xmm0, 4

71 paddd xmm1, xmm0

72 movd eax, xmm1

73 test sil, 3

74 je .L18

75 .L3:

76 mov edi, DWORD PTR [rcx+rdx*4]

77 lea r8, [0+rdx*4]

78 and edi, 1

79 cmp edi, 1

80 lea rdi, [rdx+1]

81 adc eax, 0

82 cmp rsi, rdi

x86-64 gcc 13.1 (Editor #1) x

x86-64 gcc 13.1 -std=c++2b -O3 -DNDEBUG -DUSE_FLUX

x86-64 gcc 13.1 (Editor #1) x

x86-64 gcc 13.1 -std=c++2b -O3 -DNDEBUG

Private < >

flux.godbolt.org

COMPILER EXPLORER Add... More Templates Watch C++ Weekly to learn new C++ features Sponsors intel CONAN Solid Sands Share Policies 42 Other

C++ source #1

A C++

```
#ifndef USE_FLUX
int count_evens(std::vector<int> const& vec)
{
    int count = 0;
    for (int i : vec) {
        if (i % 2 == 0) {
            ++count;
        }
    }
    return count;
}
```

Diff Viewer x86-64 gcc 13.1 vs x86-64 gcc 13.1

A Left: x86-64 gcc 13.1 -std=c++2b ... Assembly Right: x86-64 gcc 13.1 -std=c++2b ... Assembly

| | Left: x86-64 gcc 13.1 -std=c++2b ... | Right: x86-64 gcc 13.1 -std=c++2b ... |
|---------|--------------------------------------|---------------------------------------|
| 13 | jbe .L12 | |
| 14 | mov rcx, rsi | |
| 15 | pxor xmm1, xmm1 | |
| 16 | movdqa xmm3, XMMWORD PTR .LC0[rip] | |
| 17 | shr rcx, 2 | |
| 18 | movdqa xmm2, xmm1 | |
| 19 | sal rdx, 4 | |
| 20 | add rdx, rdx | |
| 21 .L4: | movdqu xmm0, XMMWORD PTR [rax] | |
| 22 | | |
| 23 | | |
| 24 | | |
| 25 | | |
| 26 | | |
| 27 | | |
| 28 | | |
| 29 | | |
| 30 | | |
| 31 | | |
| 32 | | |
| 33 | | |
| 34 | | |
| 35 | | |
| 36 | | |
| 37 | | |
| 38 .L3: | | |
| 39 | | |
| 40 | | |
| 41 | | |
| 42 | | |
| 43 | | |
| 44 | | |
| 45 | | |
| 46 | adc eax, 0 | |
| 47 | cmp rdi, rcx | |

x86-64 gcc 13.1 (Editor #1)

x86-64 gcc 13.1 -std=c++2b -O3 -DNDEBUG -DUSE_FLUX

x86-64 gcc 13.1 (Editor #1)

x86-64 gcc 13.1 -std=c++2b -O3 -DNDEBUG

Private < >

flux.godbolt.org

COMPILER EXPLORER Add... More Templates Watch C++ Weekly to learn new C++ features Sponsors intel CONAN SolidSands Share Policies 43 Other

C++ source #1

```
#ifndef USE_FLUX
int count_evens(std::vector<int> const& vec)
{
    int count = 0;
    for (int i : vec) {
        if (i % 2 == 0) {
            ++count;
        }
    }
    return count;
}
#endif
int count_evens_flux(std::vector<int> const& vec)
{
    int count = 0;
    FLUX_FOR(int i, vec) {
        if (i % 2 == 0) {
            ++count;
        }
    }
    return count;
}
```

Diff Viewer x86-64 gcc 13.1 vs x86-64 gcc 13.1

A Left: x86-64 gcc 13.1 -std=c++2b ... Assembly Right: x86-64 gcc 13.1 -std=c++2b ... Assembly

| | Left: x86-64 gcc 13.1 -std=c++2b ... | Right: x86-64 gcc 13.1 -std=c++2b ... |
|---------|--------------------------------------|---------------------------------------|
| 13 | jbe .L12 | |
| 14 | mov rcx, rsi | |
| 15 | pxor xmm1, xmm1 | |
| 16 | movdqa xmm3, XMMWORD PTR .LC0[rip] | |
| 17 | shr rcx, 2 | |
| 18 | movdqa xmm2, xmm1 | |
| 19 | sal rcx, 4 | |
| 20 | add rcx, rdx | |
| 21 .L4: | | |
| 22 | movdqu xmm0, XMMWORD PTR [rax] | |
| 23 | add rax, 16 | |
| 24 | pand xmm0, xmm3 | |
| 25 | pcmpeqd xmm0, xmm2 | |
| 26 | psubd xmm1, xmm0 | |
| 27 | cmp rax, rcx | |
| 28 | jne .L4 | |
| 29 | movdqa xmm0, xmm1 | |
| 30 | mov rcx, rsi | |
| 31 | psrldq xmm0, 8 | |
| 32 | and rcx, -4 | |
| 33 | and esi, 3 | |
| 34 | paddl xmm1, xmm0 | |
| 35 | lea rdx, [rdx+rcx*4] | |
| 36 | movdqa xmm0, xmm1 | |
| 37 | psrldq xmm0, 4 | |
| 38 | paddl xmm1, xmm0 | |
| 39 | movd eax, xmm1 | |
| 40 | je .L18 | |
| 41 .L3: | | |
| 42 | mov ecx, DWORD PTR [rdx] | |
| 43 | and ecx, 1 | |
| 44 | cmp ecx, 1 | |
| 45 | lea rcx, [rdx+4] | |
| 46 | adc eax, 0 | |
| 47 | cmp rdi, rcx | |

x86-64 gcc 13.1 (Editor #1)

x86-64 gcc 13.1 -std=c++2b -O3 -DNDEBUG -DUSE_FLUX

x86-64 gcc 13.1 (Editor #1)

x86-64 gcc 13.1 -std=c++2b -O3 -DNDEBUG

PERFORMANCE

- ▶ This implementation is using the FLUX_FOR() macro to force iteration using Flux primitives (first, is_last, read_at, inc)
- ▶ Idiomatic use of Flux would be to use a named algorithm instead, i.e. `flux::count_if()`

Private < >

flux.godbolt.org

COMPILER EXPLORER Add... More Templates

Watch C++ Weekly to learn new C++ features

Sponsors intel CONAN SolidSands

Share Policies 45 Other

C++ source #1

A C++

```
#ifndef USE_FLUX
int count_evens(std::vector<int> const& vec)
{
    int count = 0;
    for (int i : vec) {
        if (i % 2 == 0) {
            ++count;
        }
    }
    return count;
}
#endif
int count_evens_flux(std::vector<int> const& vec)
{
    return flux::count_if(vec, flux::pred::even);
}
```

Diff Viewer x86-64 clang 16.0.0 vs x86-64 clang 16.0.0

Left: x86-64 clang 16.0.0 -std=c++17 | Right: x86-64 clang 16.0.0 -std=c++17

Assembly Assembly

```
.LCPI0_0:
    .long 1          # 0x1
    .long 1          # 0x1
    .long 1          # 0x1
    .long 1          # 0x1
    .LCPI0_0:        # count_evens_flux(std::vector<int>, std::allocator<int> const&): # @
    mov r8, qword ptr [rdi]
    mov rcx, qword ptr [rdi + 8]
    cmp r8, rcx
    je .LBB0_1
    mov rsi, rcx
    sub rsi, r8
    add rsi, -4
    xor eax, eax
    cmp rsi, 28
    jae .LBB0_4
    mov rdx, r8
    jmp .LBB0_7
.LBB0_1:
    xor eax, eax
    ret
.LBB0_4:
    shr rsi, 2
    inc rsi
    mov rdi, rsi
    and rdi, -8
    lea rdx, [r8 + 4*rdi]
    pxor xmm0, xmm0
    xor eax, eax
    movdqa xmm2, xmmword ptr [rip + .LCPI0_0] # xmm2 = [1,1,1,1
    pxor xmm1, xmm1
.LBB0_5:          # =>This Inner Loop Header: |
    movdqu xmm3, xmmword ptr [r8 + 4*rax]
    movdqu xmm4, xmmword ptr [r8 + 4*rax + 16]
    pandn xmm3, xmm2
    paddd xmm0, xmm3
    pandn xmm4, xmm2
    paddd xmm1, xmm4
.LBB0_5:          # =>This Inner Loop Header: |
    movdqu xmm3, xmmword ptr [r8 + 4*rax]
    movdqu xmm4, xmmword ptr [r8 + 4*rax + 16]
    pandn xmm3, xmm2
    paddd xmm0, xmm3
    pandn xmm4, xmm2
    paddd xmm1, xmm4
```

x86-64 clang 16.0.0 (Editor #1)

x86-64 clang 16.0.0 -std=c++2b -O3 -DNDEBUG -DUSE_FLUX

x86-64 clang 16.0.0 (Editor #1)

C 6-64 clang 16.0.0 -std=c++2b -O3 -DNDEBUG

INTERNAL ITERATION

- ▶ The Flux sequence protocol has an extra, optional customisation point:
`for_each_while(predicate)`
- ▶ This instructs the sequence to perform *internal iteration* until the predicate returns false
- ▶ Many sequence adaptors can perform internal iteration more efficiently than by a check/read/advance loop
- ▶ Many algorithms can be written in terms for `for_each_while` (e.g. `fold()`, `find_if()`, `min()`)
- ▶ => Flux pipelines can in many cases generate more efficient code than equivalent C++20 Ranges pipelines



Add... ▾ More ▾ Template

Do you have any suggestions, requests or bug reports? ×

Feel free to [contact us](#) at anytime

Sponsors **intel**  **sonarlint**

Share ▾ Policies ▾ Other ▾

The screenshot shows a debugger interface with two panes. The left pane displays the C++ source code for a function named `sum_sq_even_ranges`. The right pane shows the generated assembly code from `x86-64 clang 16.0.0`.

C++ Source Code:

```
1 #include <cstddef>
2 #include <stdint.h>
3 #include <vector>
4
5 #ifdef USE_FLUX
6
7 #include "https://raw.githubusercontent.com/tcbrindle/flux/main/single_i
8
9 int sum_sq_even_flux(std::vector<int> const& vec)
10 {
11     return flux::ref(vec)
12         .filter(flux::pred::even)
13
14
15 int sum_sq_even_ranges(std::vector<int> const& vec)
16 {
17     auto view = vec
18         | std::views::filter([](int i) { return i % 2 == 0; })
19         | std::views::transform([](int i) { return i * i; });
20
21     return std::accumulate(view.begin(), view.end(), 0);
22 }
23
24
25
26
27
28
29
30
31
32
33
```

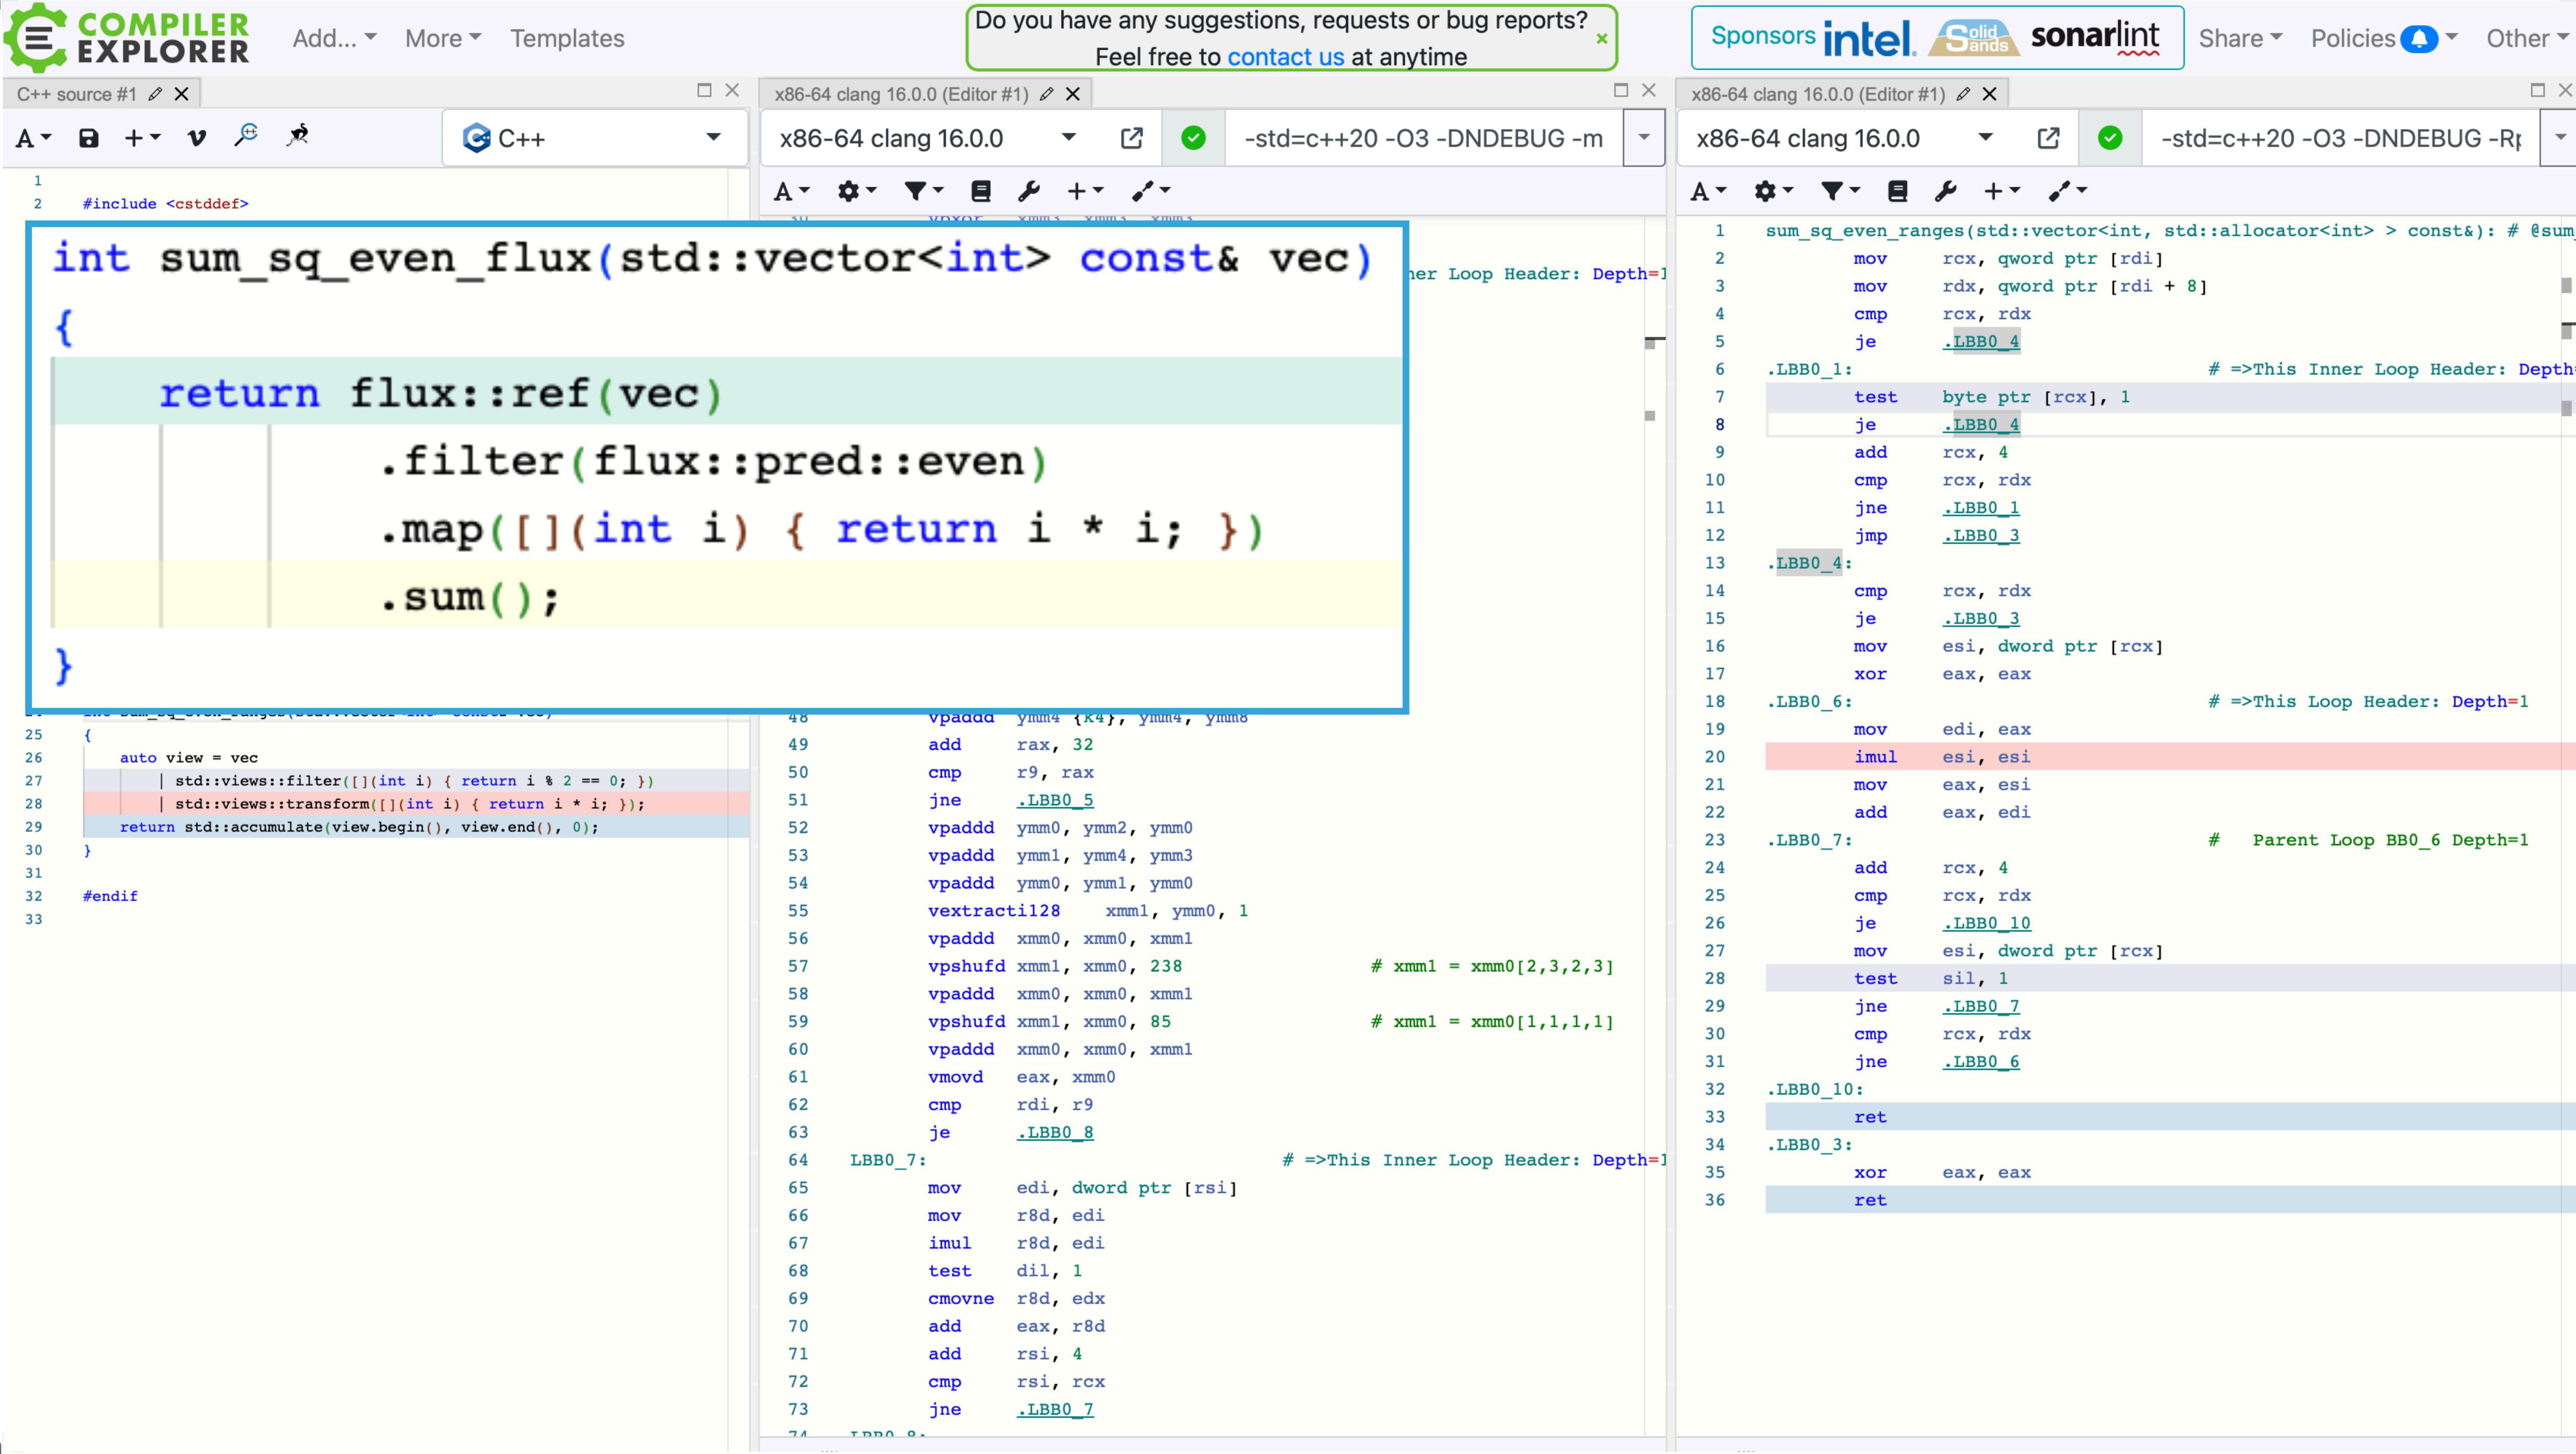
Assembly Output:

```
30 vpxor    xmm3, xmm3, xmm3
31 vpxor    xmm4, xmm4, xmm4
32 LBB0_5:                                # =>This Inner Loop Header: Depth=1
33 vmovdqu ymm5, ymmword ptr [r8 + 4*rax]
34 vmovdqu ymm6, ymmword ptr [r8 + 4*rax + 32]
35 vmovdqu ymm7, ymmword ptr [r8 + 4*rax + 64]
36 vmovdqu ymm8, ymmword ptr [r8 + 4*rax + 96]
37 vptestnmd      k1, ymm5, ymm1
38 vptestnmd      k2, ymm6, ymm1

57 vpshufd xmm1, xmm0, 238                 # xmm1 = xmm0[2,3,2,3]
58 vpadddd xmm0, xmm0, xmm1
59 vpshufd xmm1, xmm0, 85                  # xmm1 = xmm0[1,1,1,1]
60 vpadddd xmm0, xmm0, xmm1
61 vmovd eax, xmm0
62 cmp    rdi, r9
63 je     .LBB0_8
64 LBB0_7:                                # =>This Inner Loop Header: Depth=1
65 mov    edi, dword ptr [rsi]
66 mov    r8d, edi
67 imul   r8d, edi
68 test   dil, 1
69 cmovne r8d, edx
70 add    eax, r8d
71 add    rsi, 4
72 cmp    rsi, rcx
73 jne   .LBB0_7
74 tppn_o.
```

```
x86-64 clang 16.0.0 (Editor #1) ✎ X
x86-64 clang 16.0.0 ▾ ⌂ ✓ -std=c++20 -O3 -DNDEBUG -Rr
A ▾ ⌂ ▾ ⌂ ⌂ + ▾ ⌂

1 sum_sq_even_ranges(std::vector<int, std::allocator<int> > const&): # @sum
2     mov    rcx, qword ptr [rdi]
3     mov    rdx, qword ptr [rdi + 8]
4     cmp    rcx, rdx
5     je     .LBB0_4
6 .LBB0_1:                                # =>This Inner Loop Header: Depth=1
7     test   byte ptr [rcx], 1
8     je     .LBB0_4
9     add    rcx, 4
10    cmp   rcx, rdx
11    jne   .LBB0_1
12    jmp   .LBB0_3
13 .LBB0_4:
14    cmp   rcx, rdx
15    je     .LBB0_3
16    mov   esi, dword ptr [rcx]
17    xor   eax, eax
18 .LBB0_6:                                # =>This Loop Header: Depth=1
19    mov   edi, eax
20    imul  esi, esi
21    mov   eax, esi
22    add   eax, edi
23 .LBB0_7:                                # Parent Loop BB0_6 Depth=1
24    add   rcx, 4
25    cmp   rcx, rdx
26    je     .LBB0_10
27    mov   esi, dword ptr [rcx]
28    test   sil, 1
29    jne   .LBB0_7
30    cmp   rcx, rdx
31    jne   .LBB0_6
32 .LBB0_10:
33     ret
34 .LBB0_3:
35     xor   eax, eax
36     ret
```





Add... ▾ More ▾ Template

Do you have any suggestions, requests or bug reports? ×

Feel free to [contact us](#) at anytime

Sponsors **intel**  **sonarlint**

Share ▾ Policies 🔔 ▾ Other ▾

```
1 #include <cstddef>
2 #include <stdint.h>
3 #include <vector>
4
5
6 #ifdef USE_FLUX
7
8 #include "https://raw.githubusercontent.com/tcbrindle/flux/main/single.h"
9
10 int sum_sq_even_flux(std::vector<int> const& vec)
11 {
12     return flux::ref(vec)
13         .filter(flux::pred::even)
14         .map([](int i) { return i * i; })
15         .sum();
16 }
17
18 #else
19
20 #include <algorithm>
21 #include <numeric>
22 #include <ranges>
23
24 int sum_sq_even_ranges(std::vector<int> const& vec)
25 {
26     auto view = vec
27         | std::views::filter([](int i) { return i % 2 == 0; })
28         | std::views::transform([](int i) { return i * i; });
29     return std::accumulate(view.begin(), view.end(), 0);
30 }
31
32 #endif
```

```
30         vpxor    xmm3, xmm3, xmm3
31         vpxor    xmm4, xmm4, xmm4
32 LBB0_5:                                     # =>This Inner Loop Header: Depth=1
33         vmovdqu ymm5, ymmword ptr [r8 + 4*rax]
34         vmovdqu ymm6, ymmword ptr [r8 + 4*rax + 32]
35         vmovdqu ymm7, ymmword ptr [r8 + 4*rax + 64]
36         vmovdqu ymm8, ymmword ptr [r8 + 4*rax + 96]
37         vptestnmd      k1, ymm5, ymm1
38         vptestnmd      k2, ymm6, ymm1
39         vptestnmd      k3, ymm7, ymm1
40         vptestnmd      k4, ymm8, ymm1
41         vpmulld ymm5, ymm5, ymm5
42         vpmulld ymm6, ymm6, ymm6
43         vpmulld ymm7, ymm7, ymm7
44         vpmulld ymm8, ymm8, ymm8
45         vpadddd ymm0 {k1}, ymm0, ymm5
46         vpadddd ymm2 {k2}, ymm2, ymm6
47         vpadddd ymm3 {k3}, ymm3, ymm7
48         vpadddd ymm4 {k4}, ymm4, ymm8
49         add     rax, 32
50         cmp     r9, rax
51         jne     .LBB0_5
52         vpadddd ymm0, ymm2, ymm0
53         vpadddd ymm1, ymm4, ymm3
54         vpadddd ymm0, ymm1, ymm0
55         vextracti128  xmm1, ymm0, 1
56         vpadddd xmm0, xmm0, xmm1
57         vpshufd xmm1, xmm0, 238          # xmm1 = xmm0[2,3,2,3]
58         vpadddd xmm0, xmm0, xmm1
59         vpshufd xmm1, xmm0, 85           # xmm1 = xmm0[1,1,1,1]
60         vpadddd xmm0, xmm0, xmm1
61         vmove    eax, xmm0
62         cmp     rdi, r9
63         je      .LBB0_8
64 LBB0_7:                                     # =>This Inner Loop Header: Depth=1
65         mov     edi, dword ptr [rsi]
66         mov     r8d, edi
67         imul   r8d, edi
68         test    dil, 1
69         cmovne r8d, edx
70         add     eax, r8d
71         add     rsi, 4
72         cmp     rsi, rcx
73         jne     .LBB0_7
74 LBB0_8:
```

```
1 sum_sq_even_ranges(std::vector<int, std::allocator<int> > const&): # @sum
2     mov    rcx, qword ptr [rdi]
3     mov    rdx, qword ptr [rdi + 8]
4     cmp    rcx, rdx
5     je     .LBB0_4
6 .LBB0_1:                                # =>This Inner Loop Header: Depth=1
7     test   byte ptr [rcx], 1
8     je     .LBB0_4
9     add    rcx, 4
10    cmp    rcx, rdx
11    jne    .LBB0_1
12    jmp    .LBB0_3
13 .LBB0_4:
14    cmp    rcx, rdx
15    je     .LBB0_3
16    mov    esi, dword ptr [rcx]
17    xor    eax, eax
18 .LBB0_6:                                # =>This Loop Header: Depth=1
19    mov    edi, eax
20    imul  esi, esi
21    mov    eax, esi
22    add    eax, edi
23 .LBB0_7:                                # Parent Loop BB0_6 Depth=1
24    add    rcx, 4
25    cmp    rcx, rdx
26    je     .LBB0_10
27    mov    esi, dword ptr [rcx]
28    test   sil, 1
29    jne    .LBB0_7
30    cmp    rcx, rdx
31    jne    .LBB0_6
32 .LBB0_10:
33     ret
34 .LBB0_3:
35     xor    eax, eax
36     ret
```

AVOIDING BOUNDS CHECKING

- ▶ Implementations of `for_each_while()` typically don't need internal bounds-checking
- ▶ In most other cases, the compiler can optimise away bounds checks
- ▶ But what about corner cases?

AVOIDING BOUNDS CHECKING

- ▶ If bounds checking does lead to unacceptable overhead, Flux provides the `read_at_unchecked()` function which performs element access without additional checks
- ▶ This can be used ***selectively***, only where truly needed
 - ▶ Always measure!
- ▶ Think of it as the equivalent of unsafe blocks in other languages

CONCLUSION

- ▶ Flux is a modern C++20 library for *collection-orientated programming*
- ▶ It can be used everywhere C++20 ranges can be used, and offers
 - ▶ Much improved safety by default
 - ▶ An easy-to-use API for both producers and consumers
 - ▶ Compatibility with existing STL code
 - ▶ In many cases, better performance than equivalent ranges code

THANK YOU!

- ▶ Flux: github.com/tcbrindle/flux
- ▶ Docs: tristanbrindle.com/flux
- ▶ Twitter: @tristanbrindle

EXPLICIT REFERENCING IN FLUX

- ▶ C++20 range adaptors (“views”) take an implicit reference when passed a non-view, lvalue range
- ▶ This can lead to surprising behaviour and “dangling views” in some cases

```
auto filter_evens(std::vector<int> const& vec)
{
    return std::views::filter(vec, [](int i) { return i % 2 == 0; });
}

auto evens = filter_evens({1, 2, 3, 4, 5}); // uh oh
```

EXPLICIT REFERENCING IN FLUX

- ▶ In Flux, potentially “long-lived” references must be explicitly passed to sequence adaptors
- ▶ Makes dangling much more obvious in code

```
auto filter_evens(std::vector<int> const& vec)
{
    return flux::filter(vec, [](int i) { return i % 2 == 0; }); // Compile error!
```

EXPLICIT REFERENCING IN FLUX

- ▶ In Flux, potentially “long-lived” references must be explicitly passed to sequence adaptors
- ▶ Makes dangling much more obvious in code

```
auto filter_evens(std::vector<int> const& vec)
{
    // Now compiles, but more obviously problematic
    return flux::filter(flux::ref(vec), [](int i) { return i % 2 == 0; });
}
```