

+ 23

Object Introspection:

A Revolutionary Memory Profiler for C++ Objects

JONATHAN HASLAM
& ADITYA SARWADE



20
23



Scene Setting

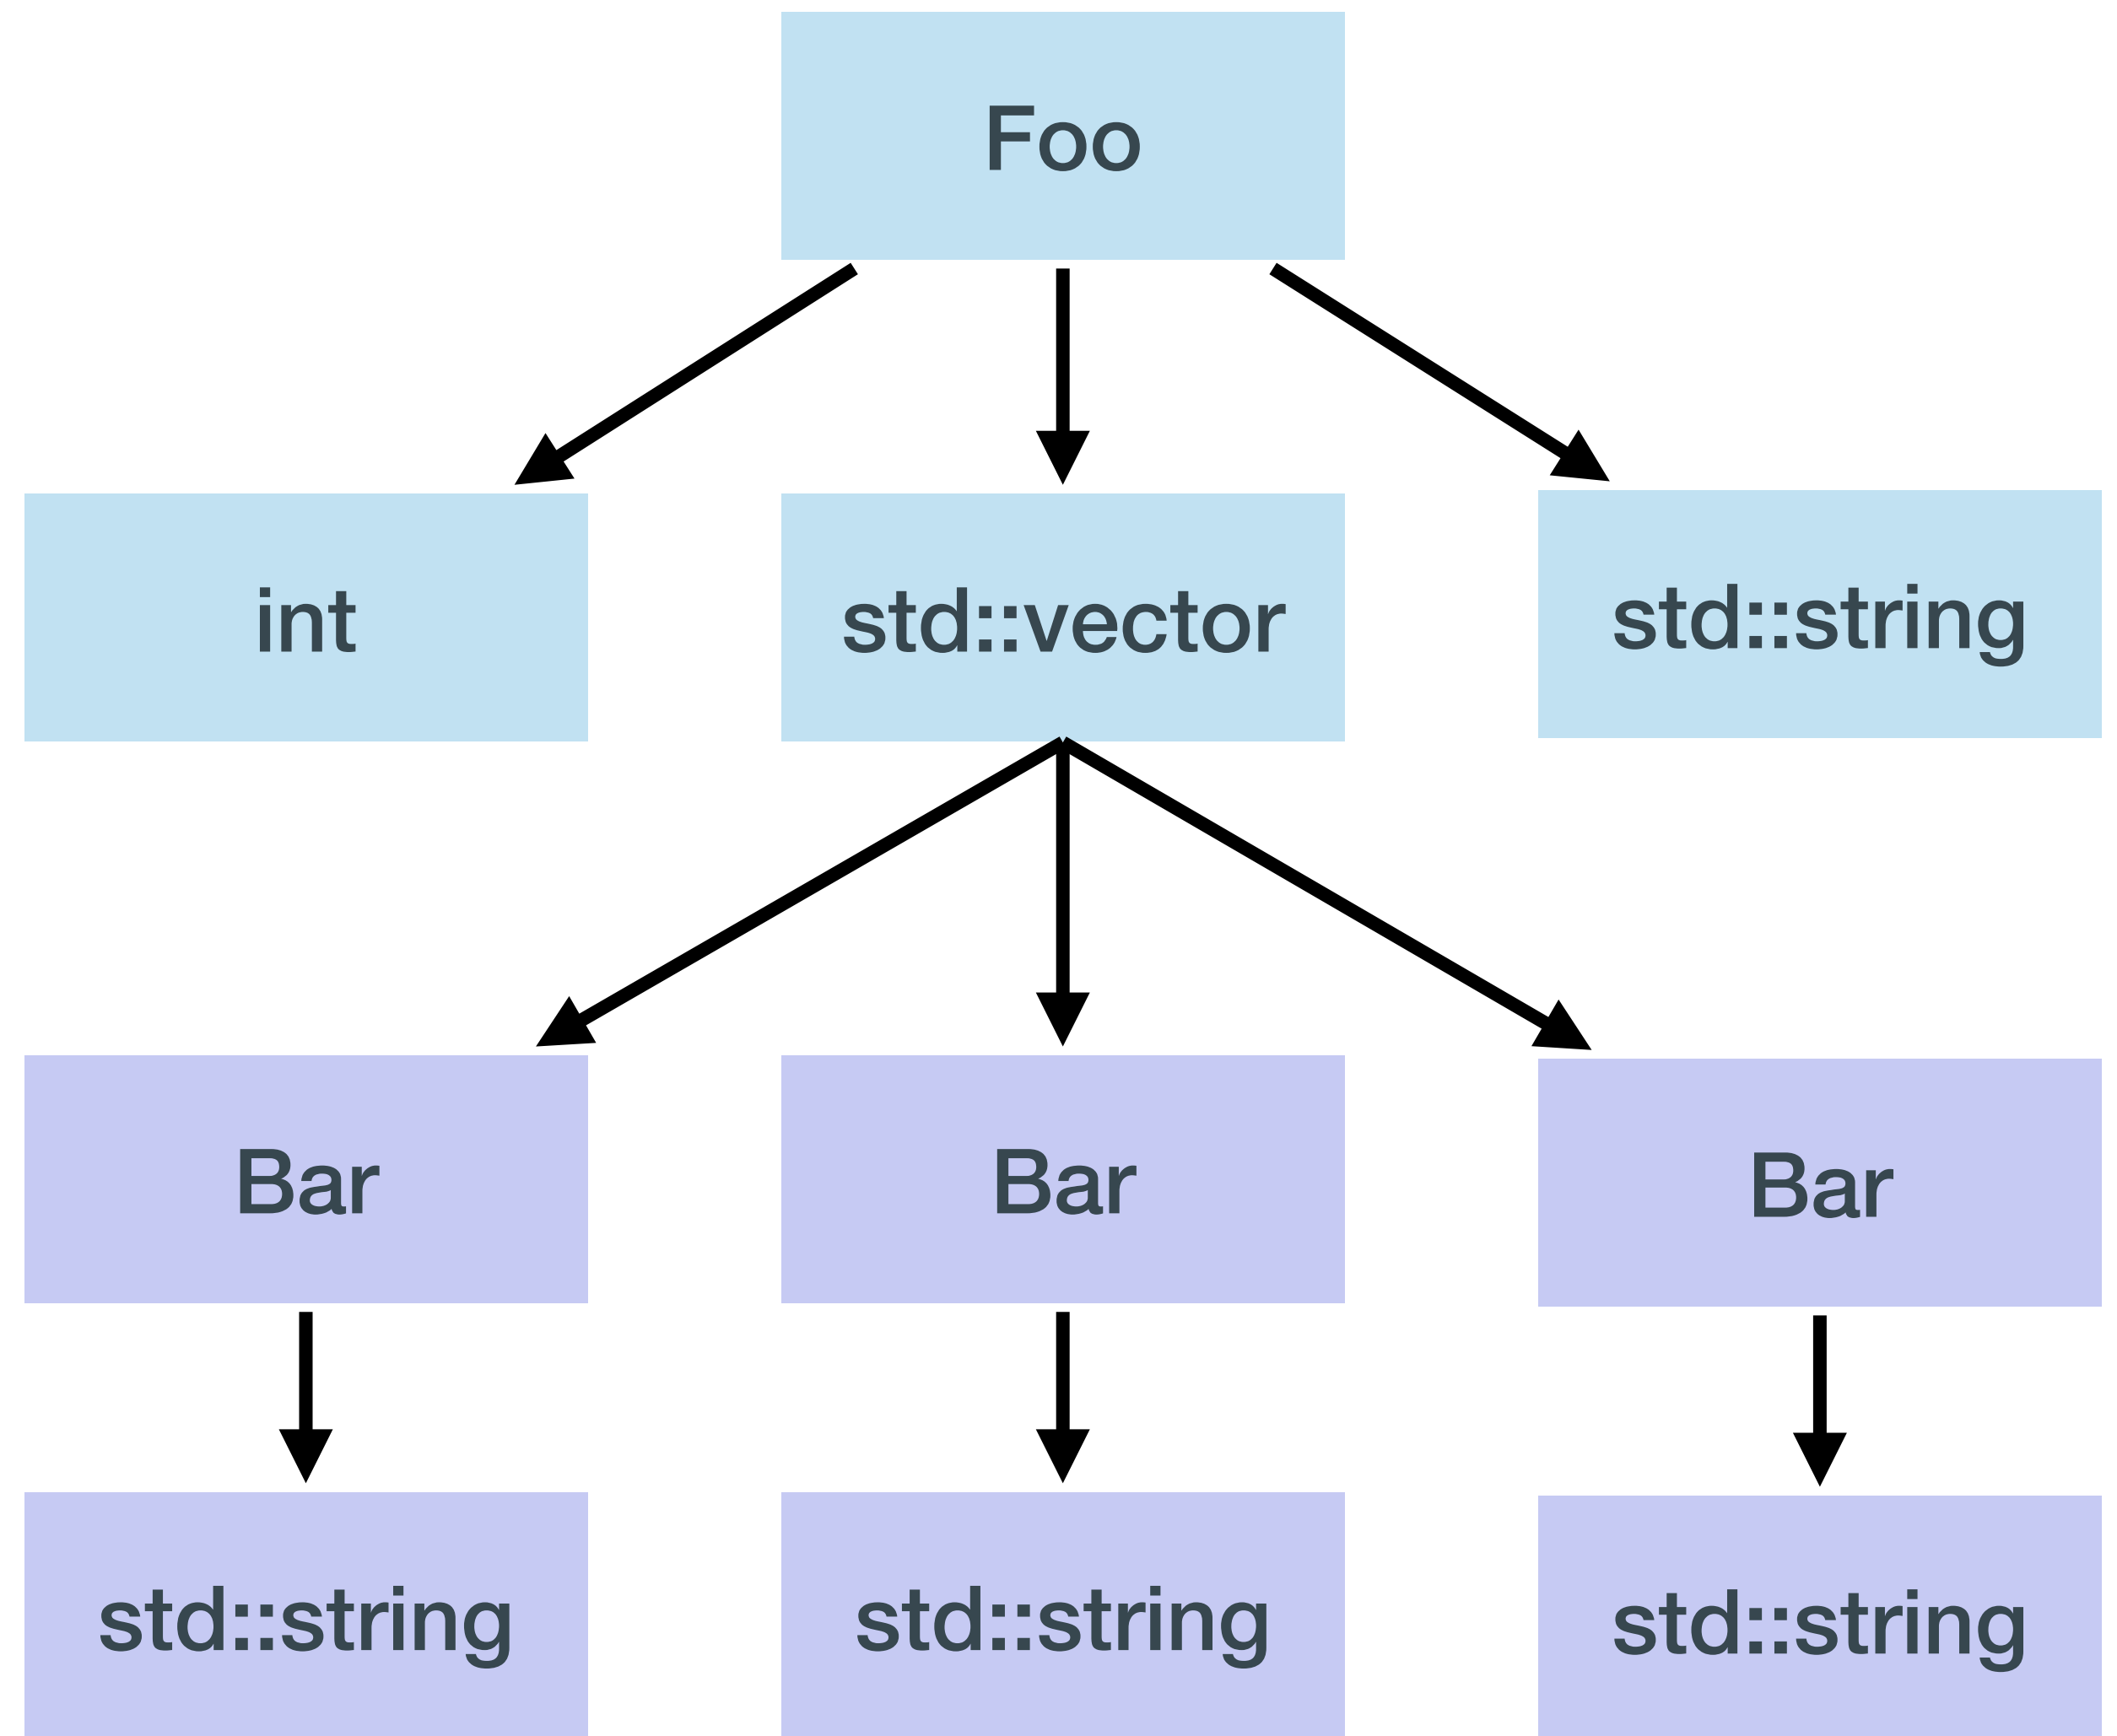
“... **functions** then provide the basic vocabulary of computation, just as **types** (built-in and user-defined) provide the basic vocabulary of data.”

Bjarne Stroustrup, A Tour of C++, Second Edition.

```
struct TimeStamps {  
    i64    insertionTime;  
    i64    processingTime;  
};
```

Scene Setting

```
struct Bar {  
    std::string str;  
};  
  
struct Foo {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```

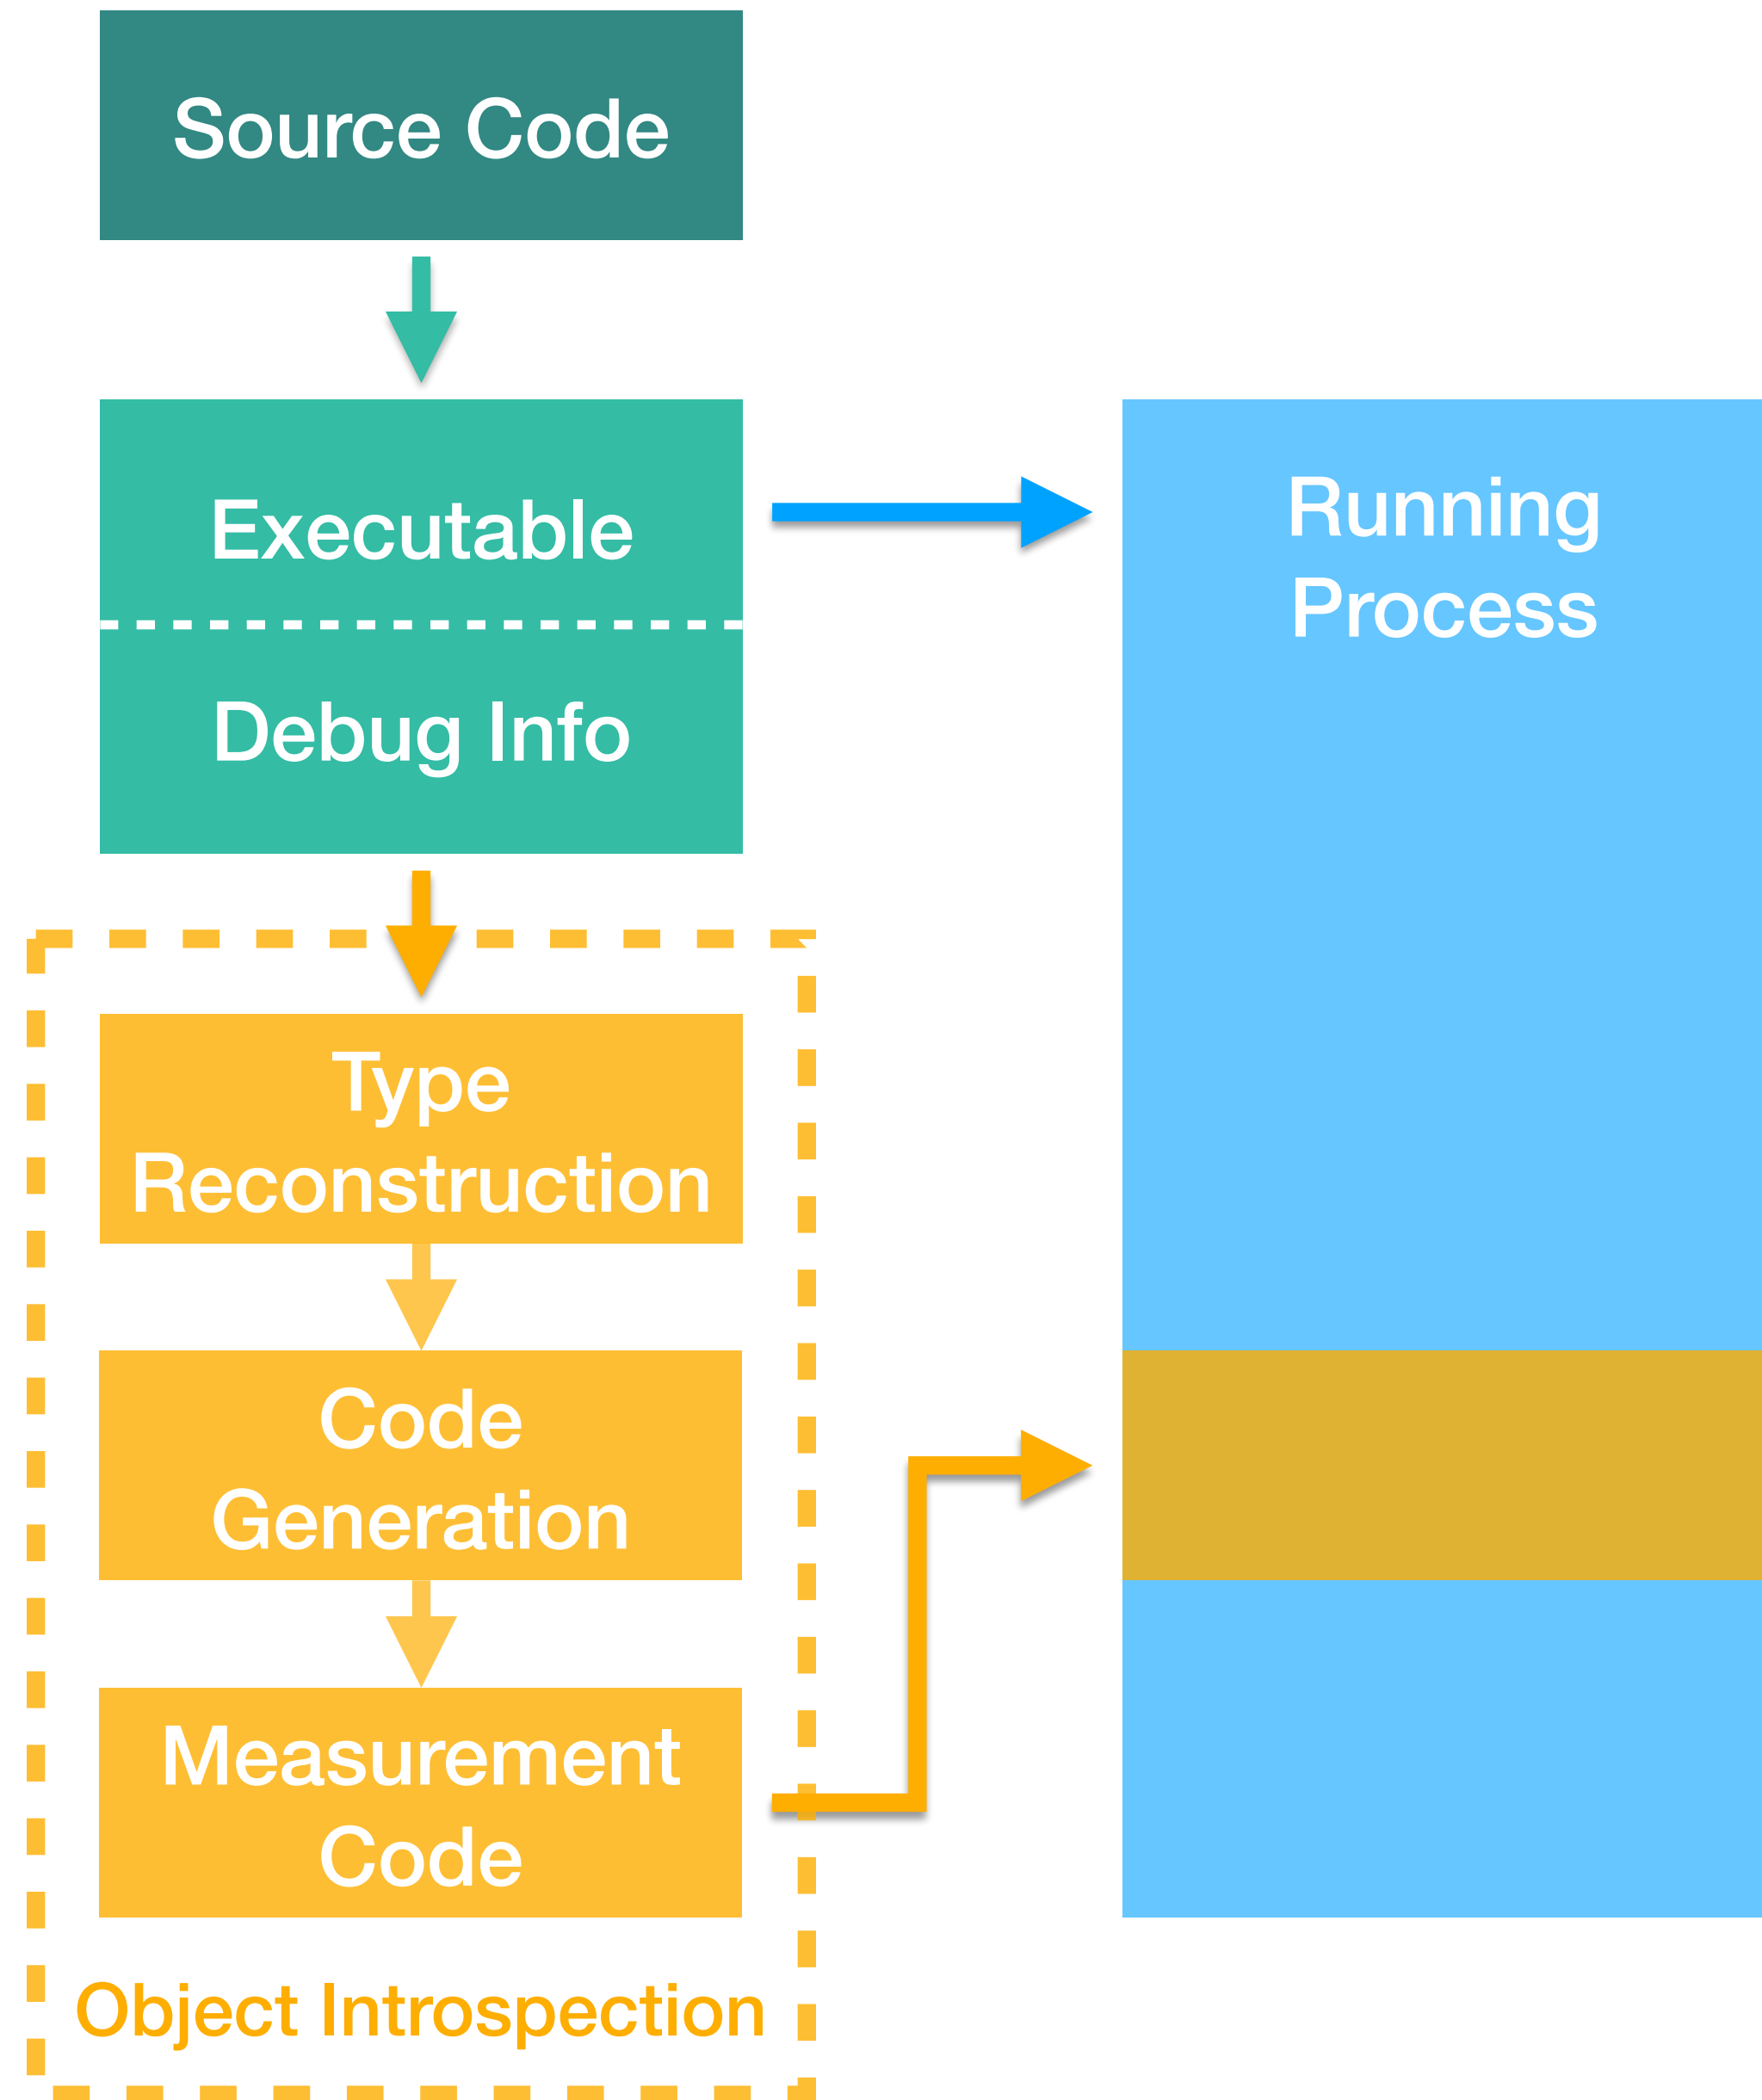


Object Introspection (OI) Goals

- Byte level memory footprint analysis for objects
 - Complete object type hierarchies
 - Dynamic allocations and containers
- Optimised, production applications in production environments
- Dynamic Profiler
 - Can be used with no recompilation or code modification
- APIs

Outline

- Debug Info Analysis
 - Type/Layout Reconstruction
 - Code Generation
- Object Introspection as a Library
- Object Introspection as a Profiler
- Object Introspection Applied



Type Reconstruction

- (Re)construct an entire object hierarchy from a given root type
 - Understand the layout in memory of the entire hierarchy
 - Understand how to interpret data at memory offsets
 - Understand containers
- Compiler generated debug information required
 - DWARF v4 (DebugFission supported)
 - drgn (<https://github.com/osandov/drgn>) and lldb

Type Graph

Source

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int                n;  
    std::vector<Bar>   bar_vec;  
    std::string        foo_str;  
};
```


Type Graph

.o

w/ DebugInfo

Source

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```

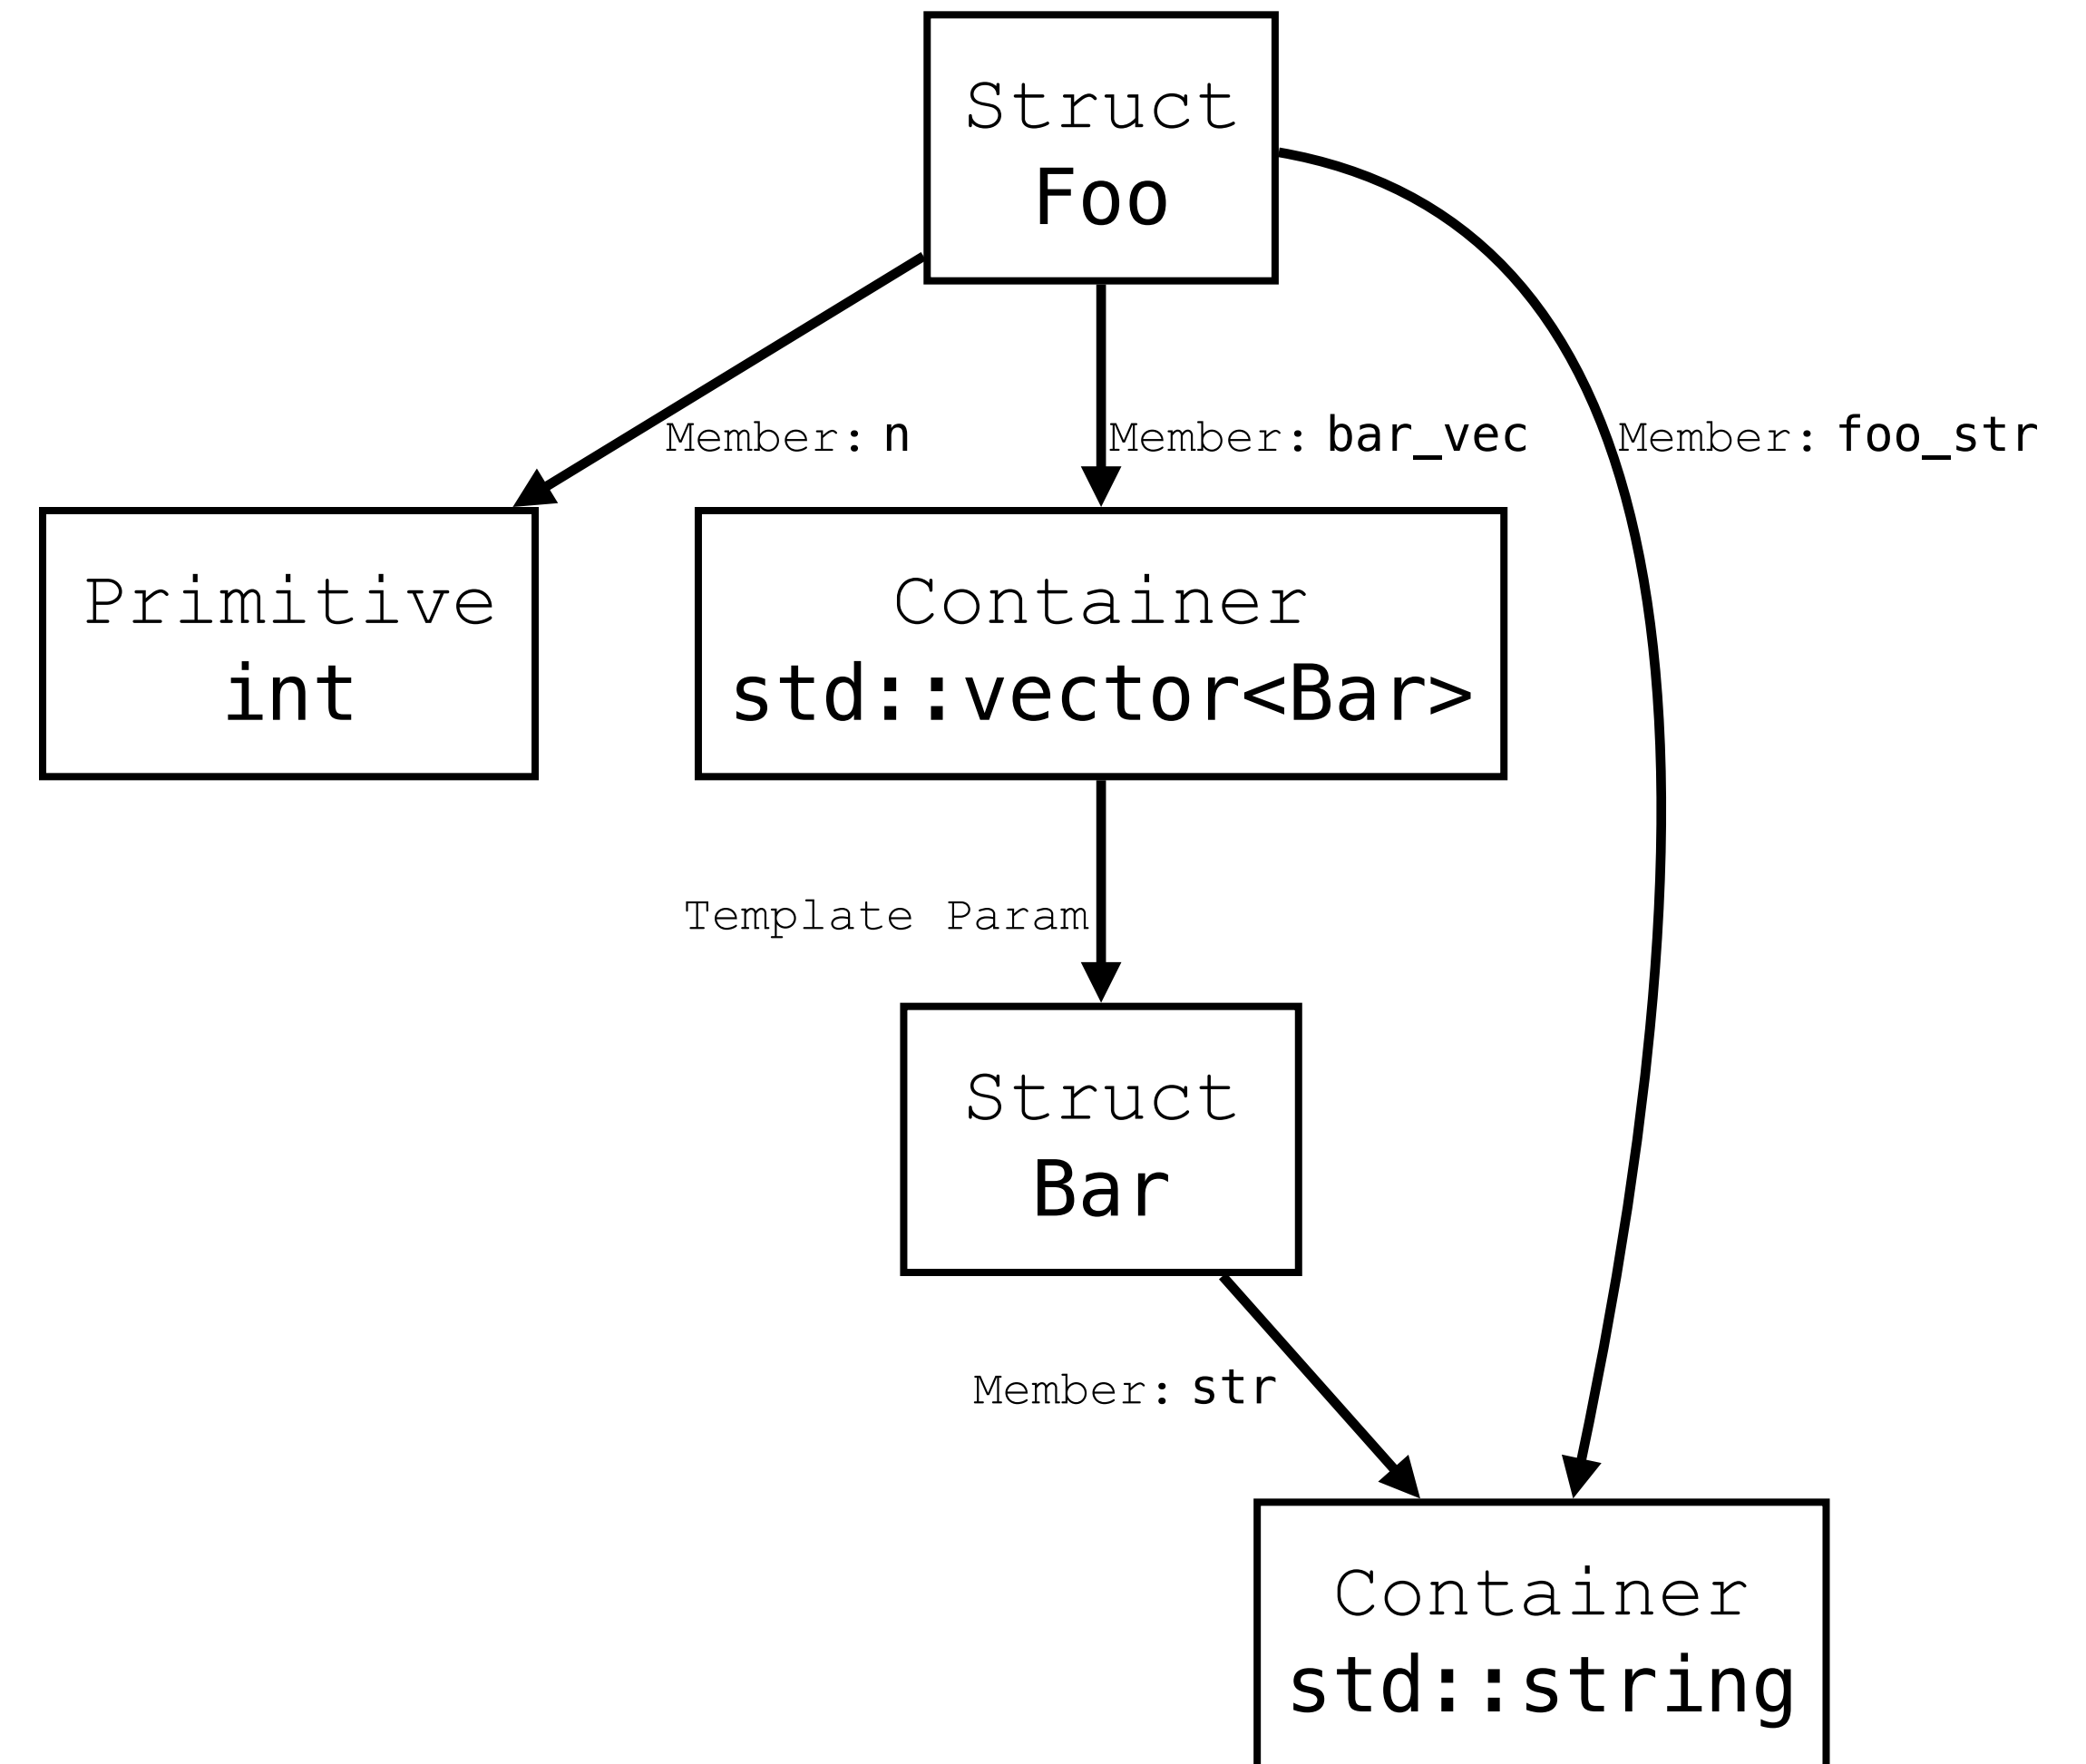
Type Graph

.o

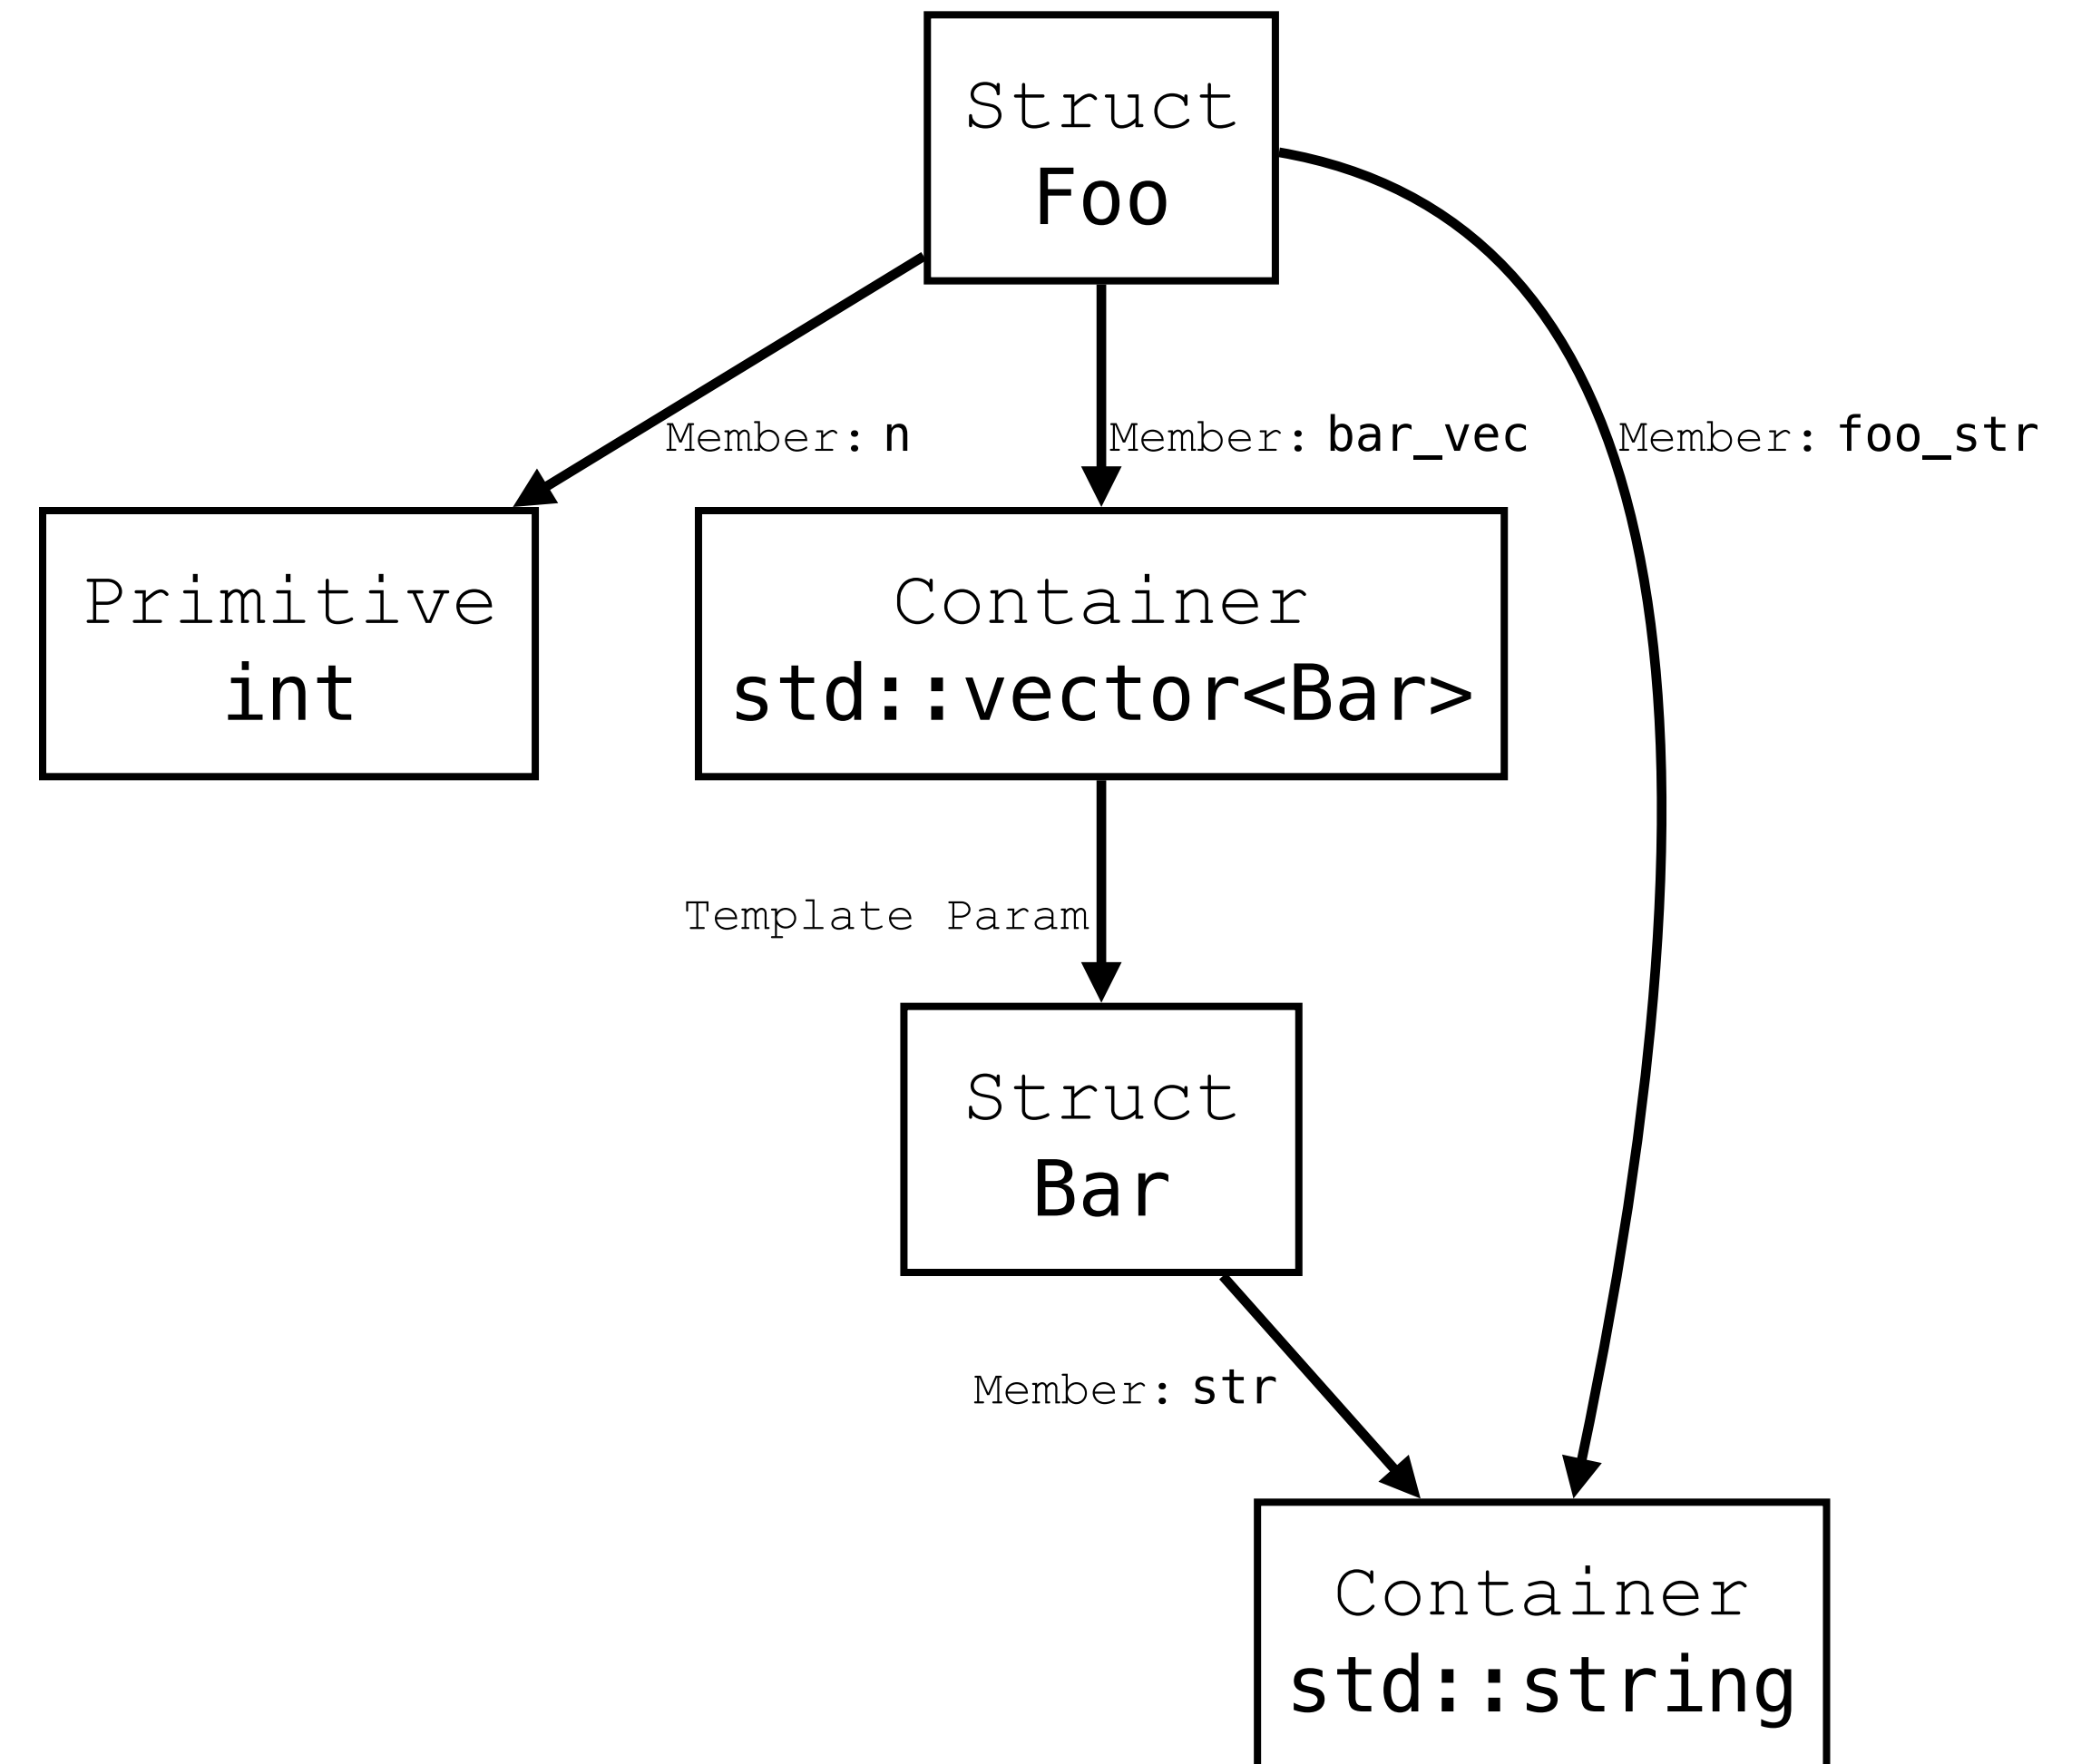
w/ DebugInfo

Type Graph

.o
w/ DebugInfo



Type Graph

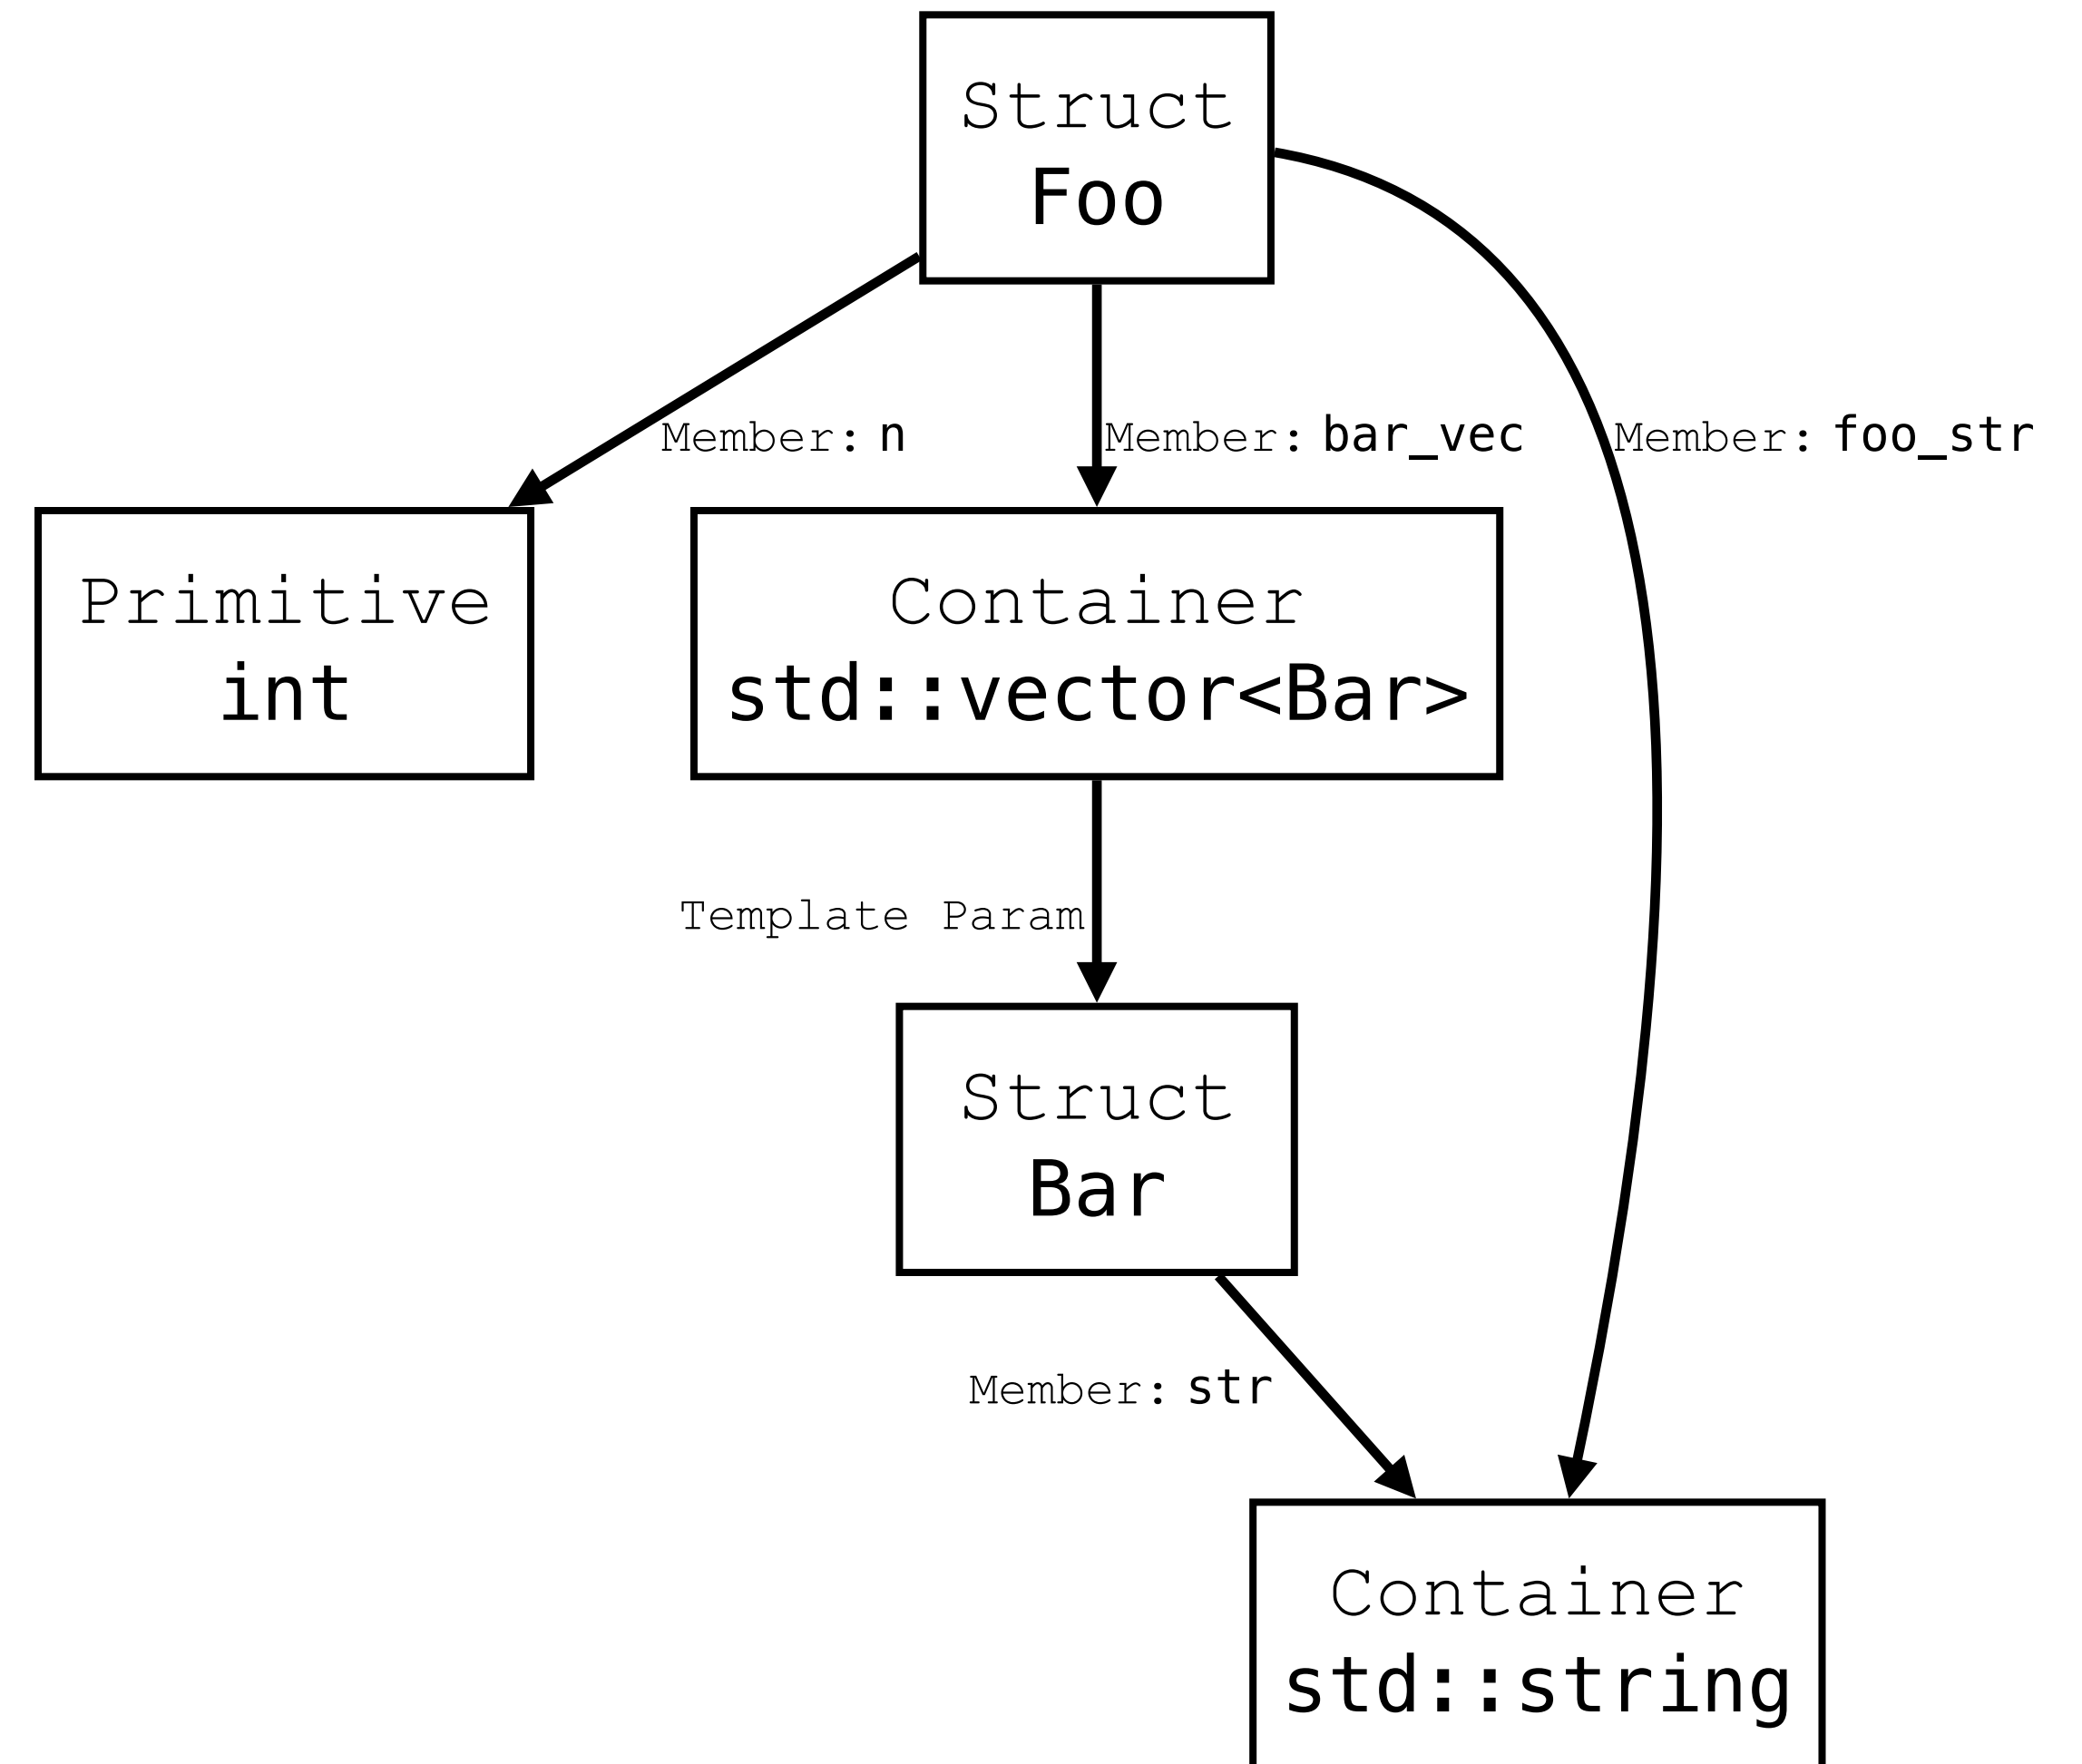


Type Graph

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```

Reconstructed

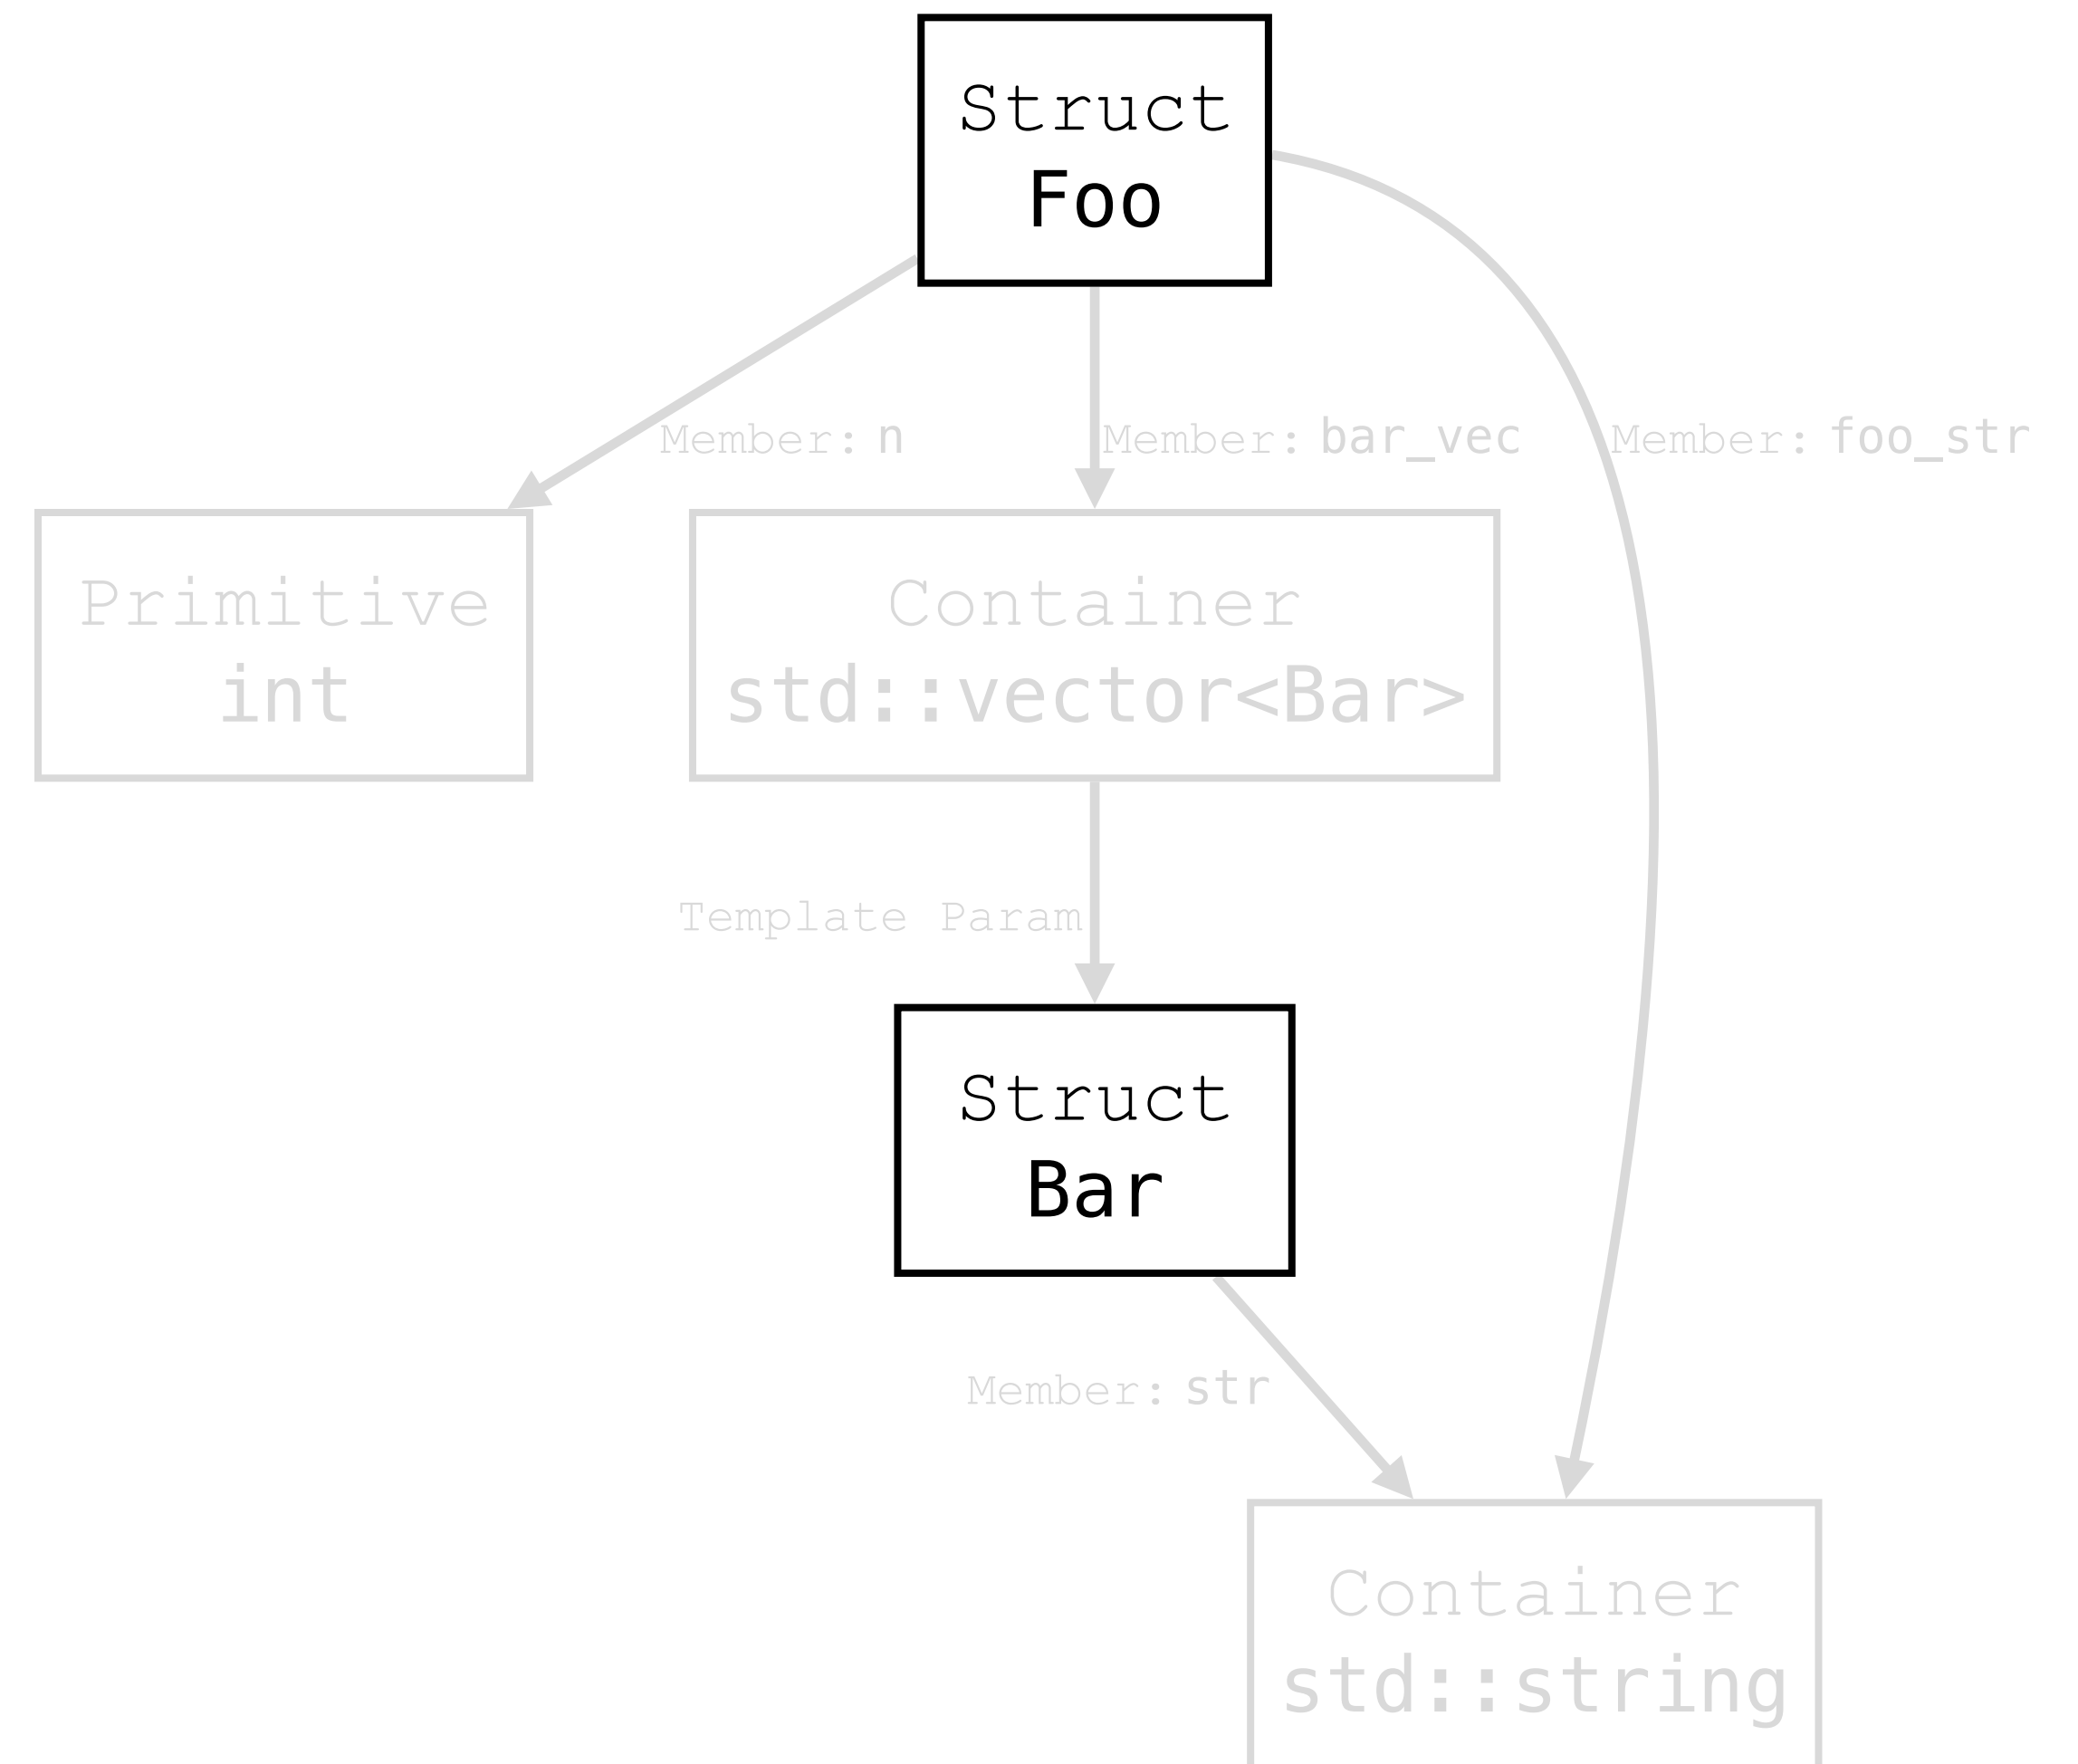


Type Graph

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```

Reconstructed

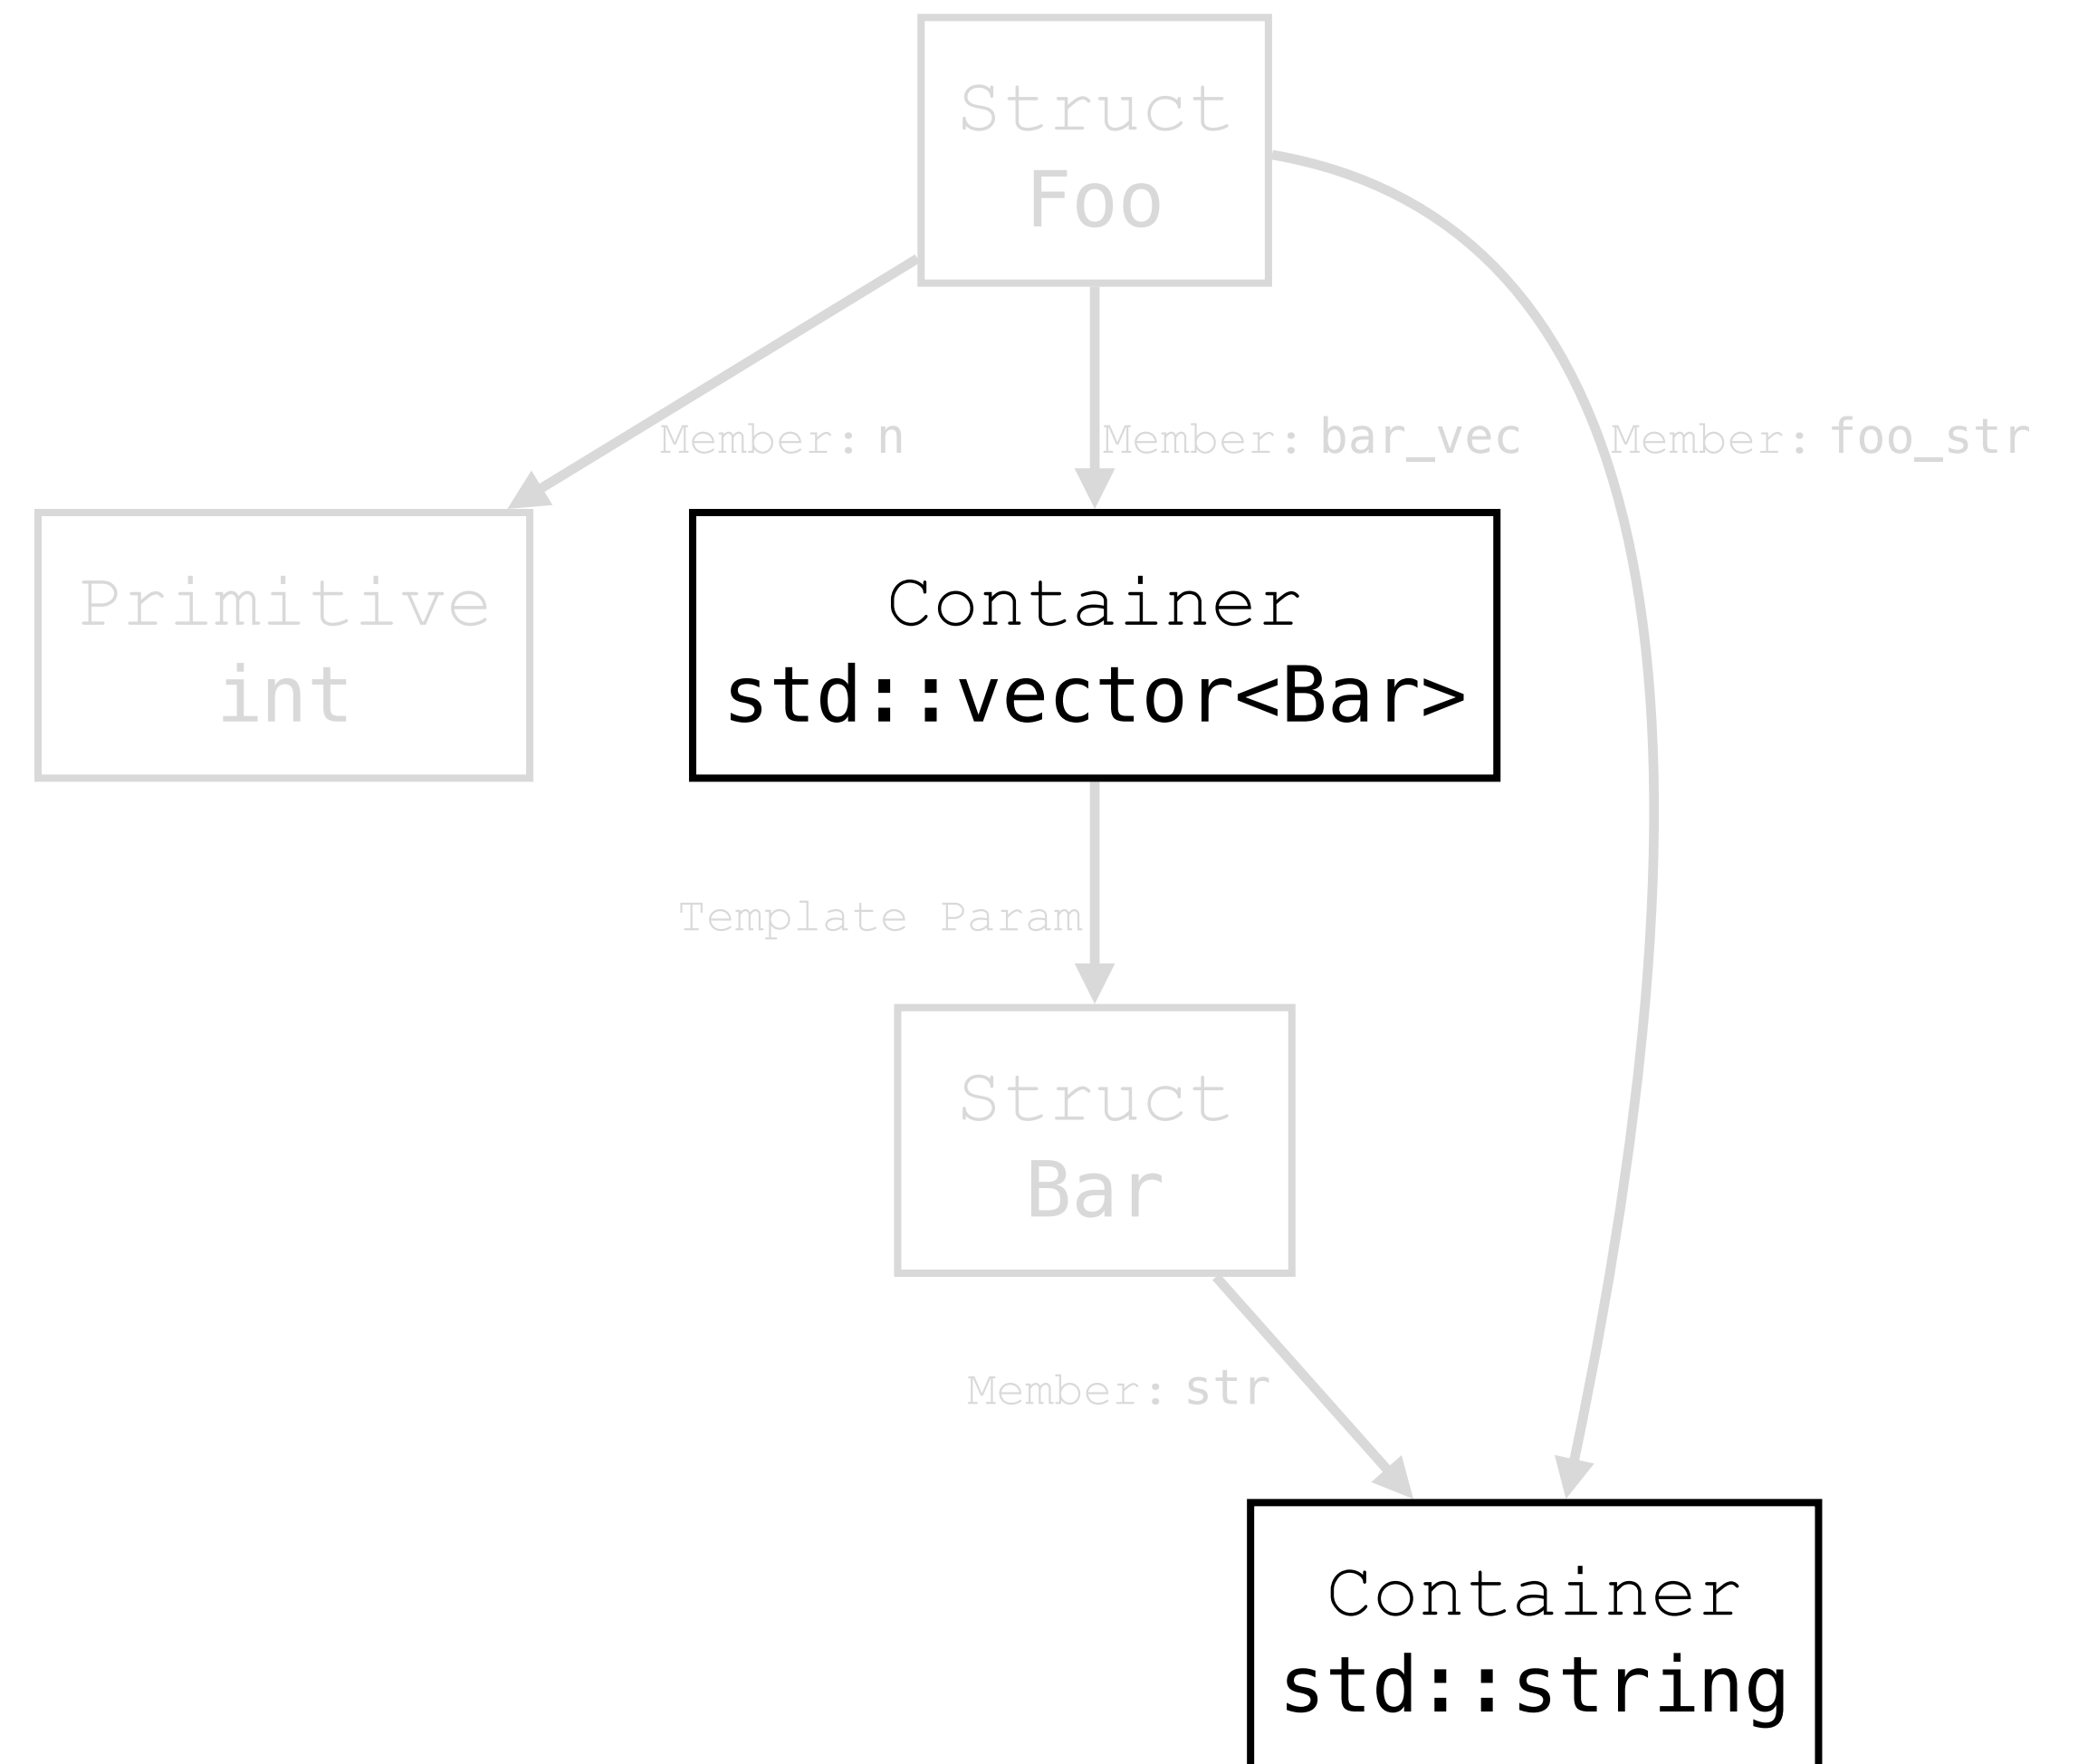


Type Graph

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```

Reconstructed



Type Graph Transformation: Inline Inheritance

Source

```
struct Bar {  
    std::string str;  
};
```

```
struct MyParent {  
    int a;  
    int b;  
};
```

```
struct Foo : public MyParent {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```

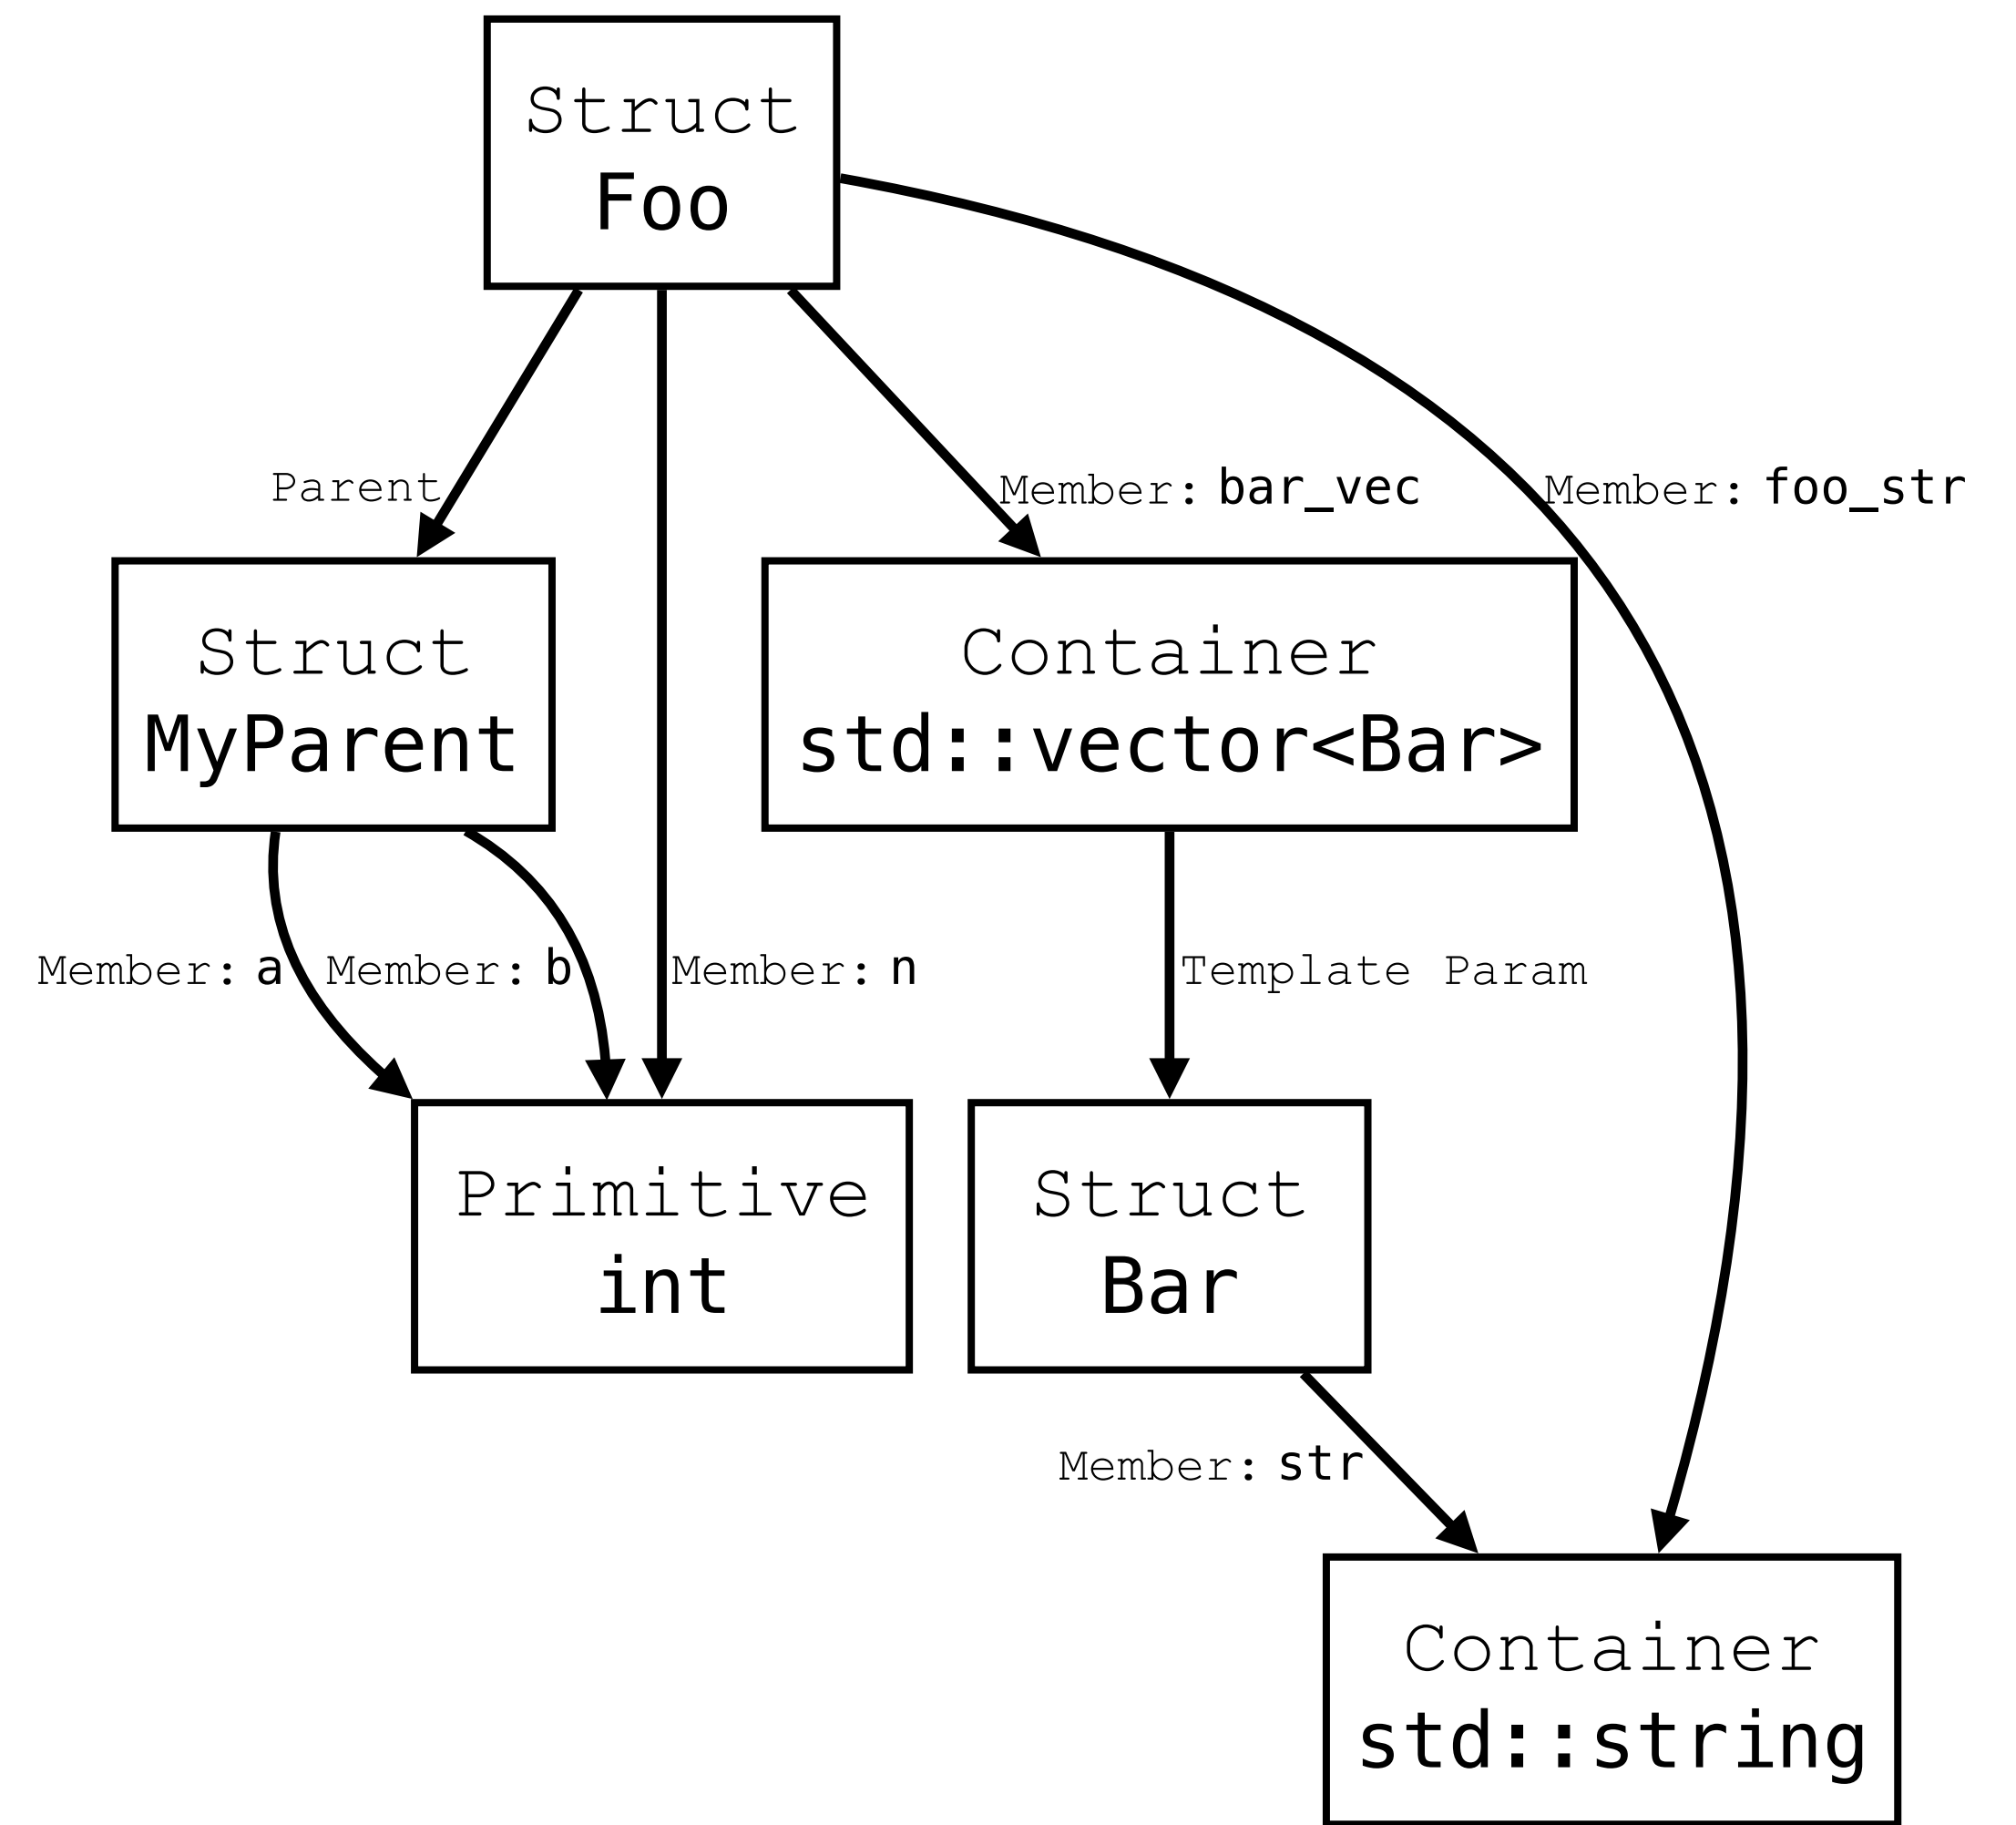

Type Graph Transformation: Inline Inheritance

Source

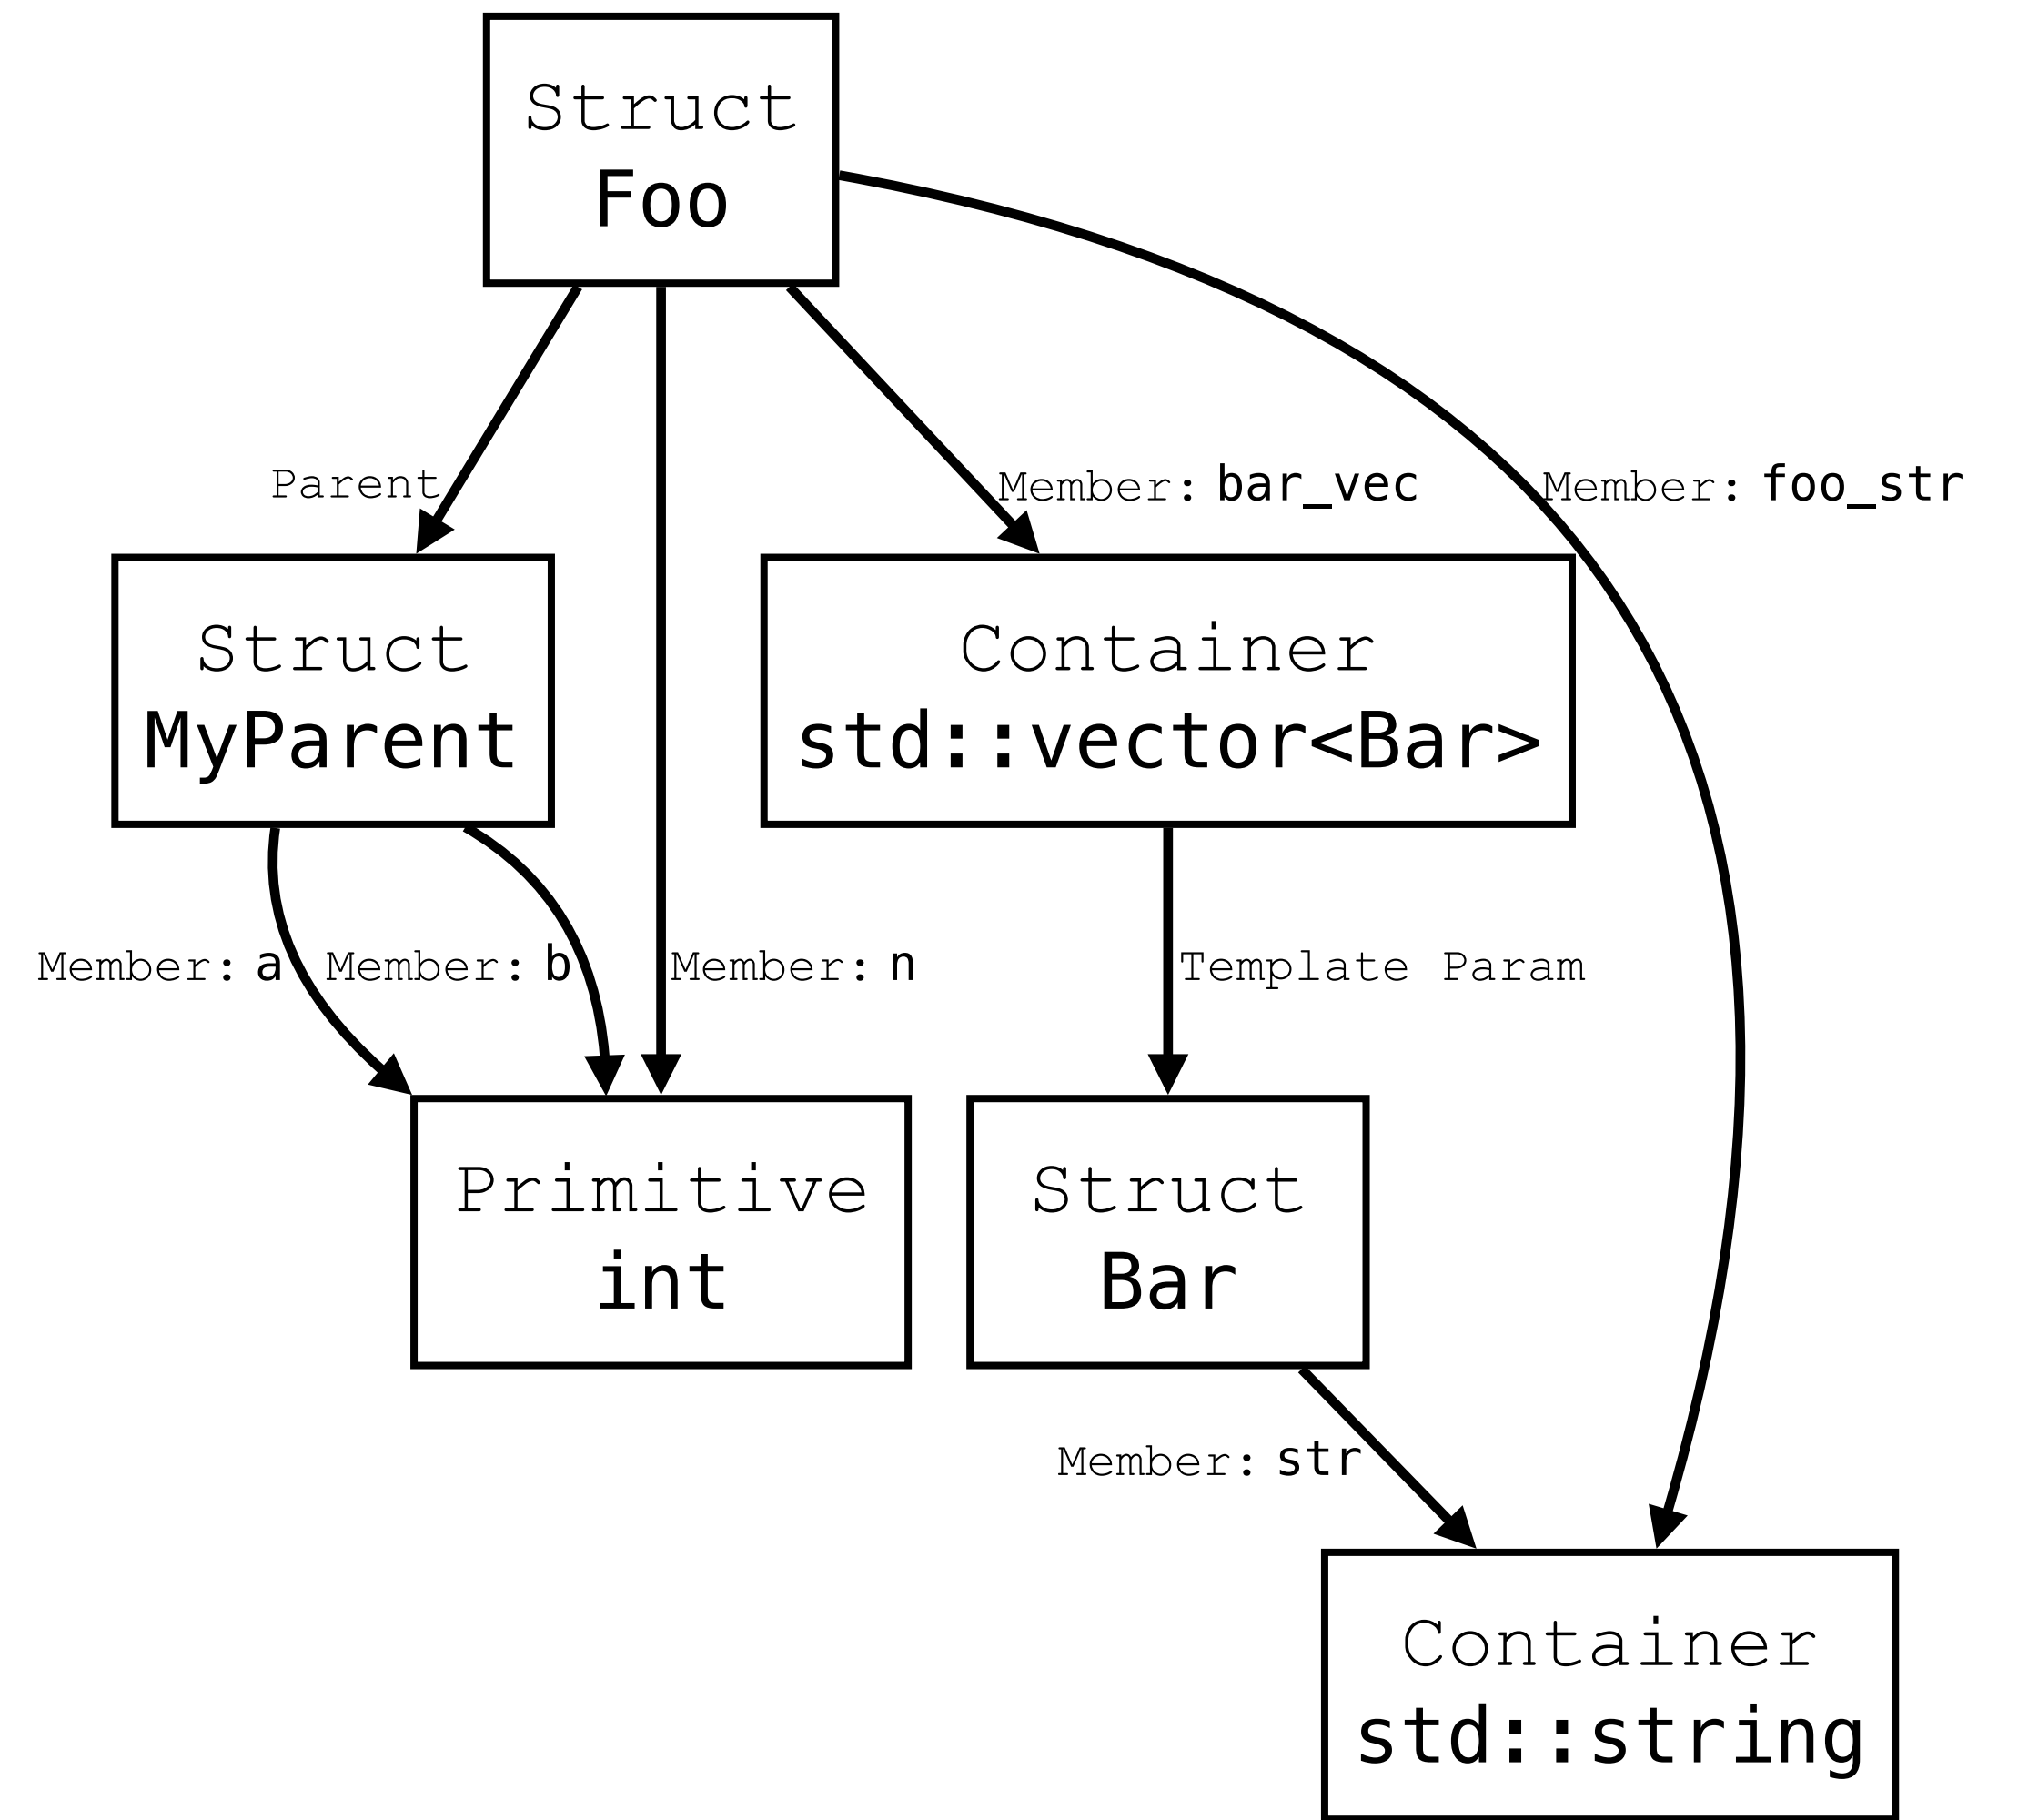
```
struct Bar {  
    std::string str;  
};
```

```
struct MyParent {  
    int a;  
    int b;  
};
```

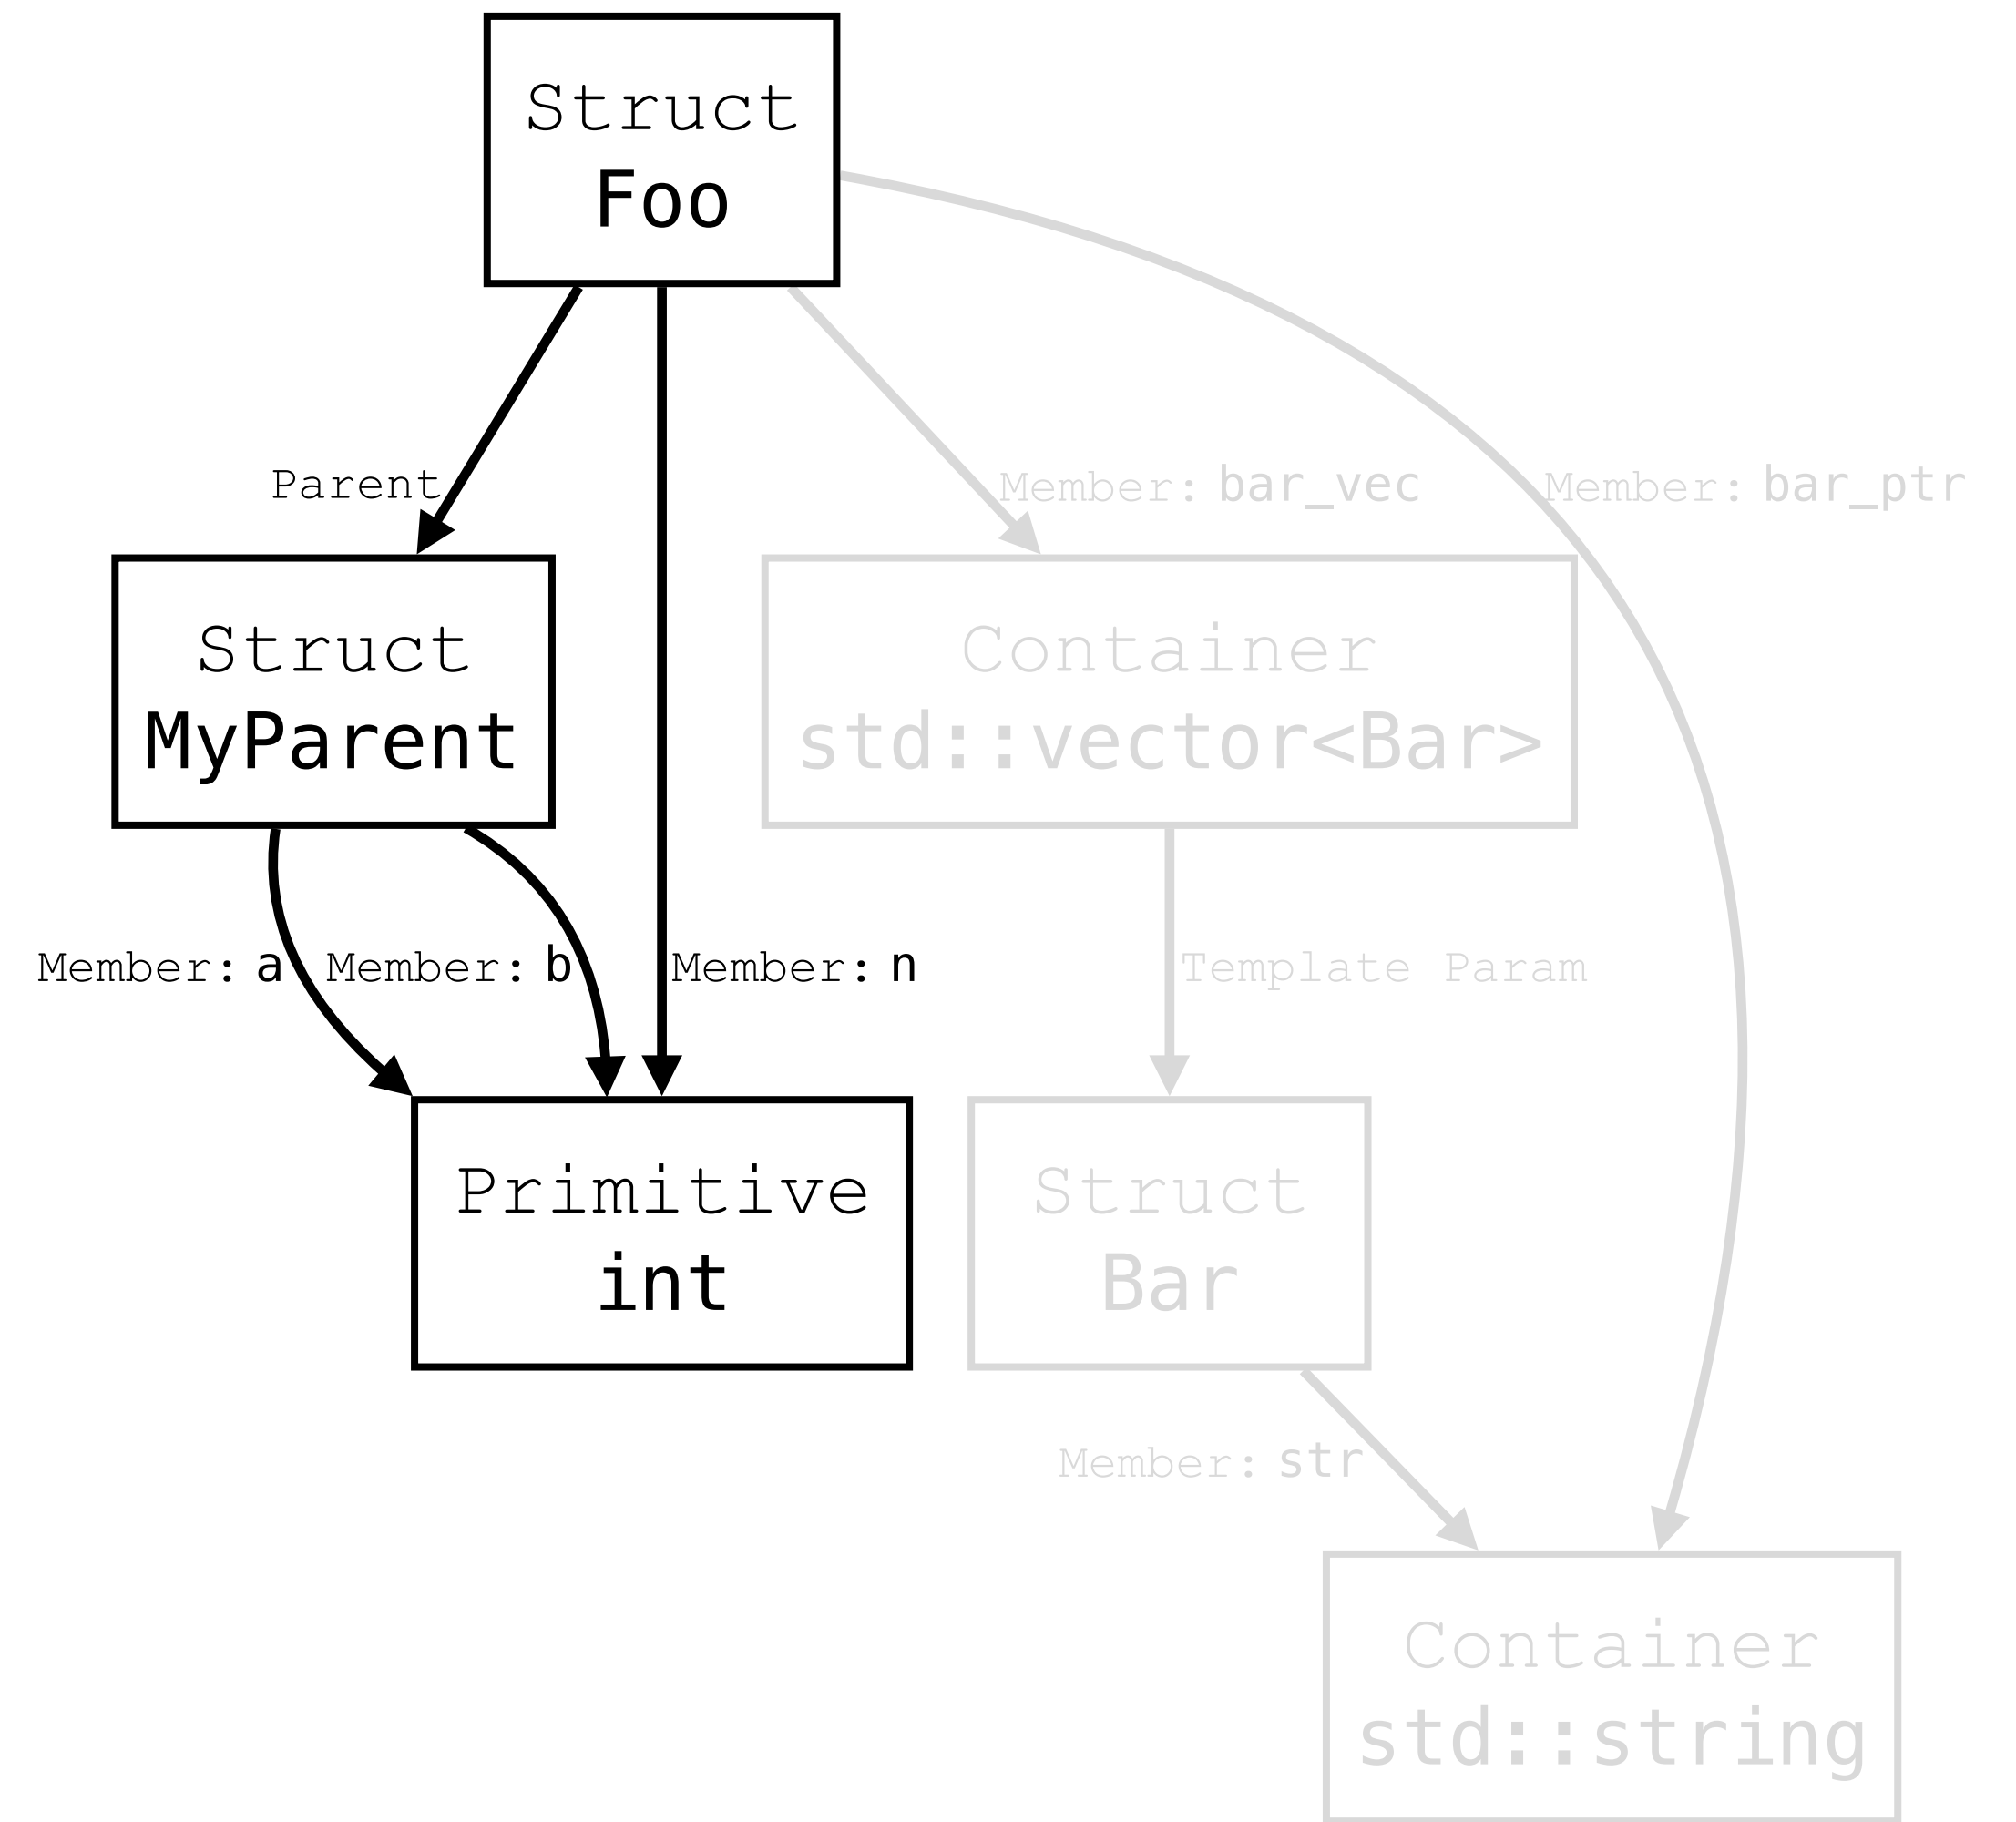
```
struct Foo : public MyParent {  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```



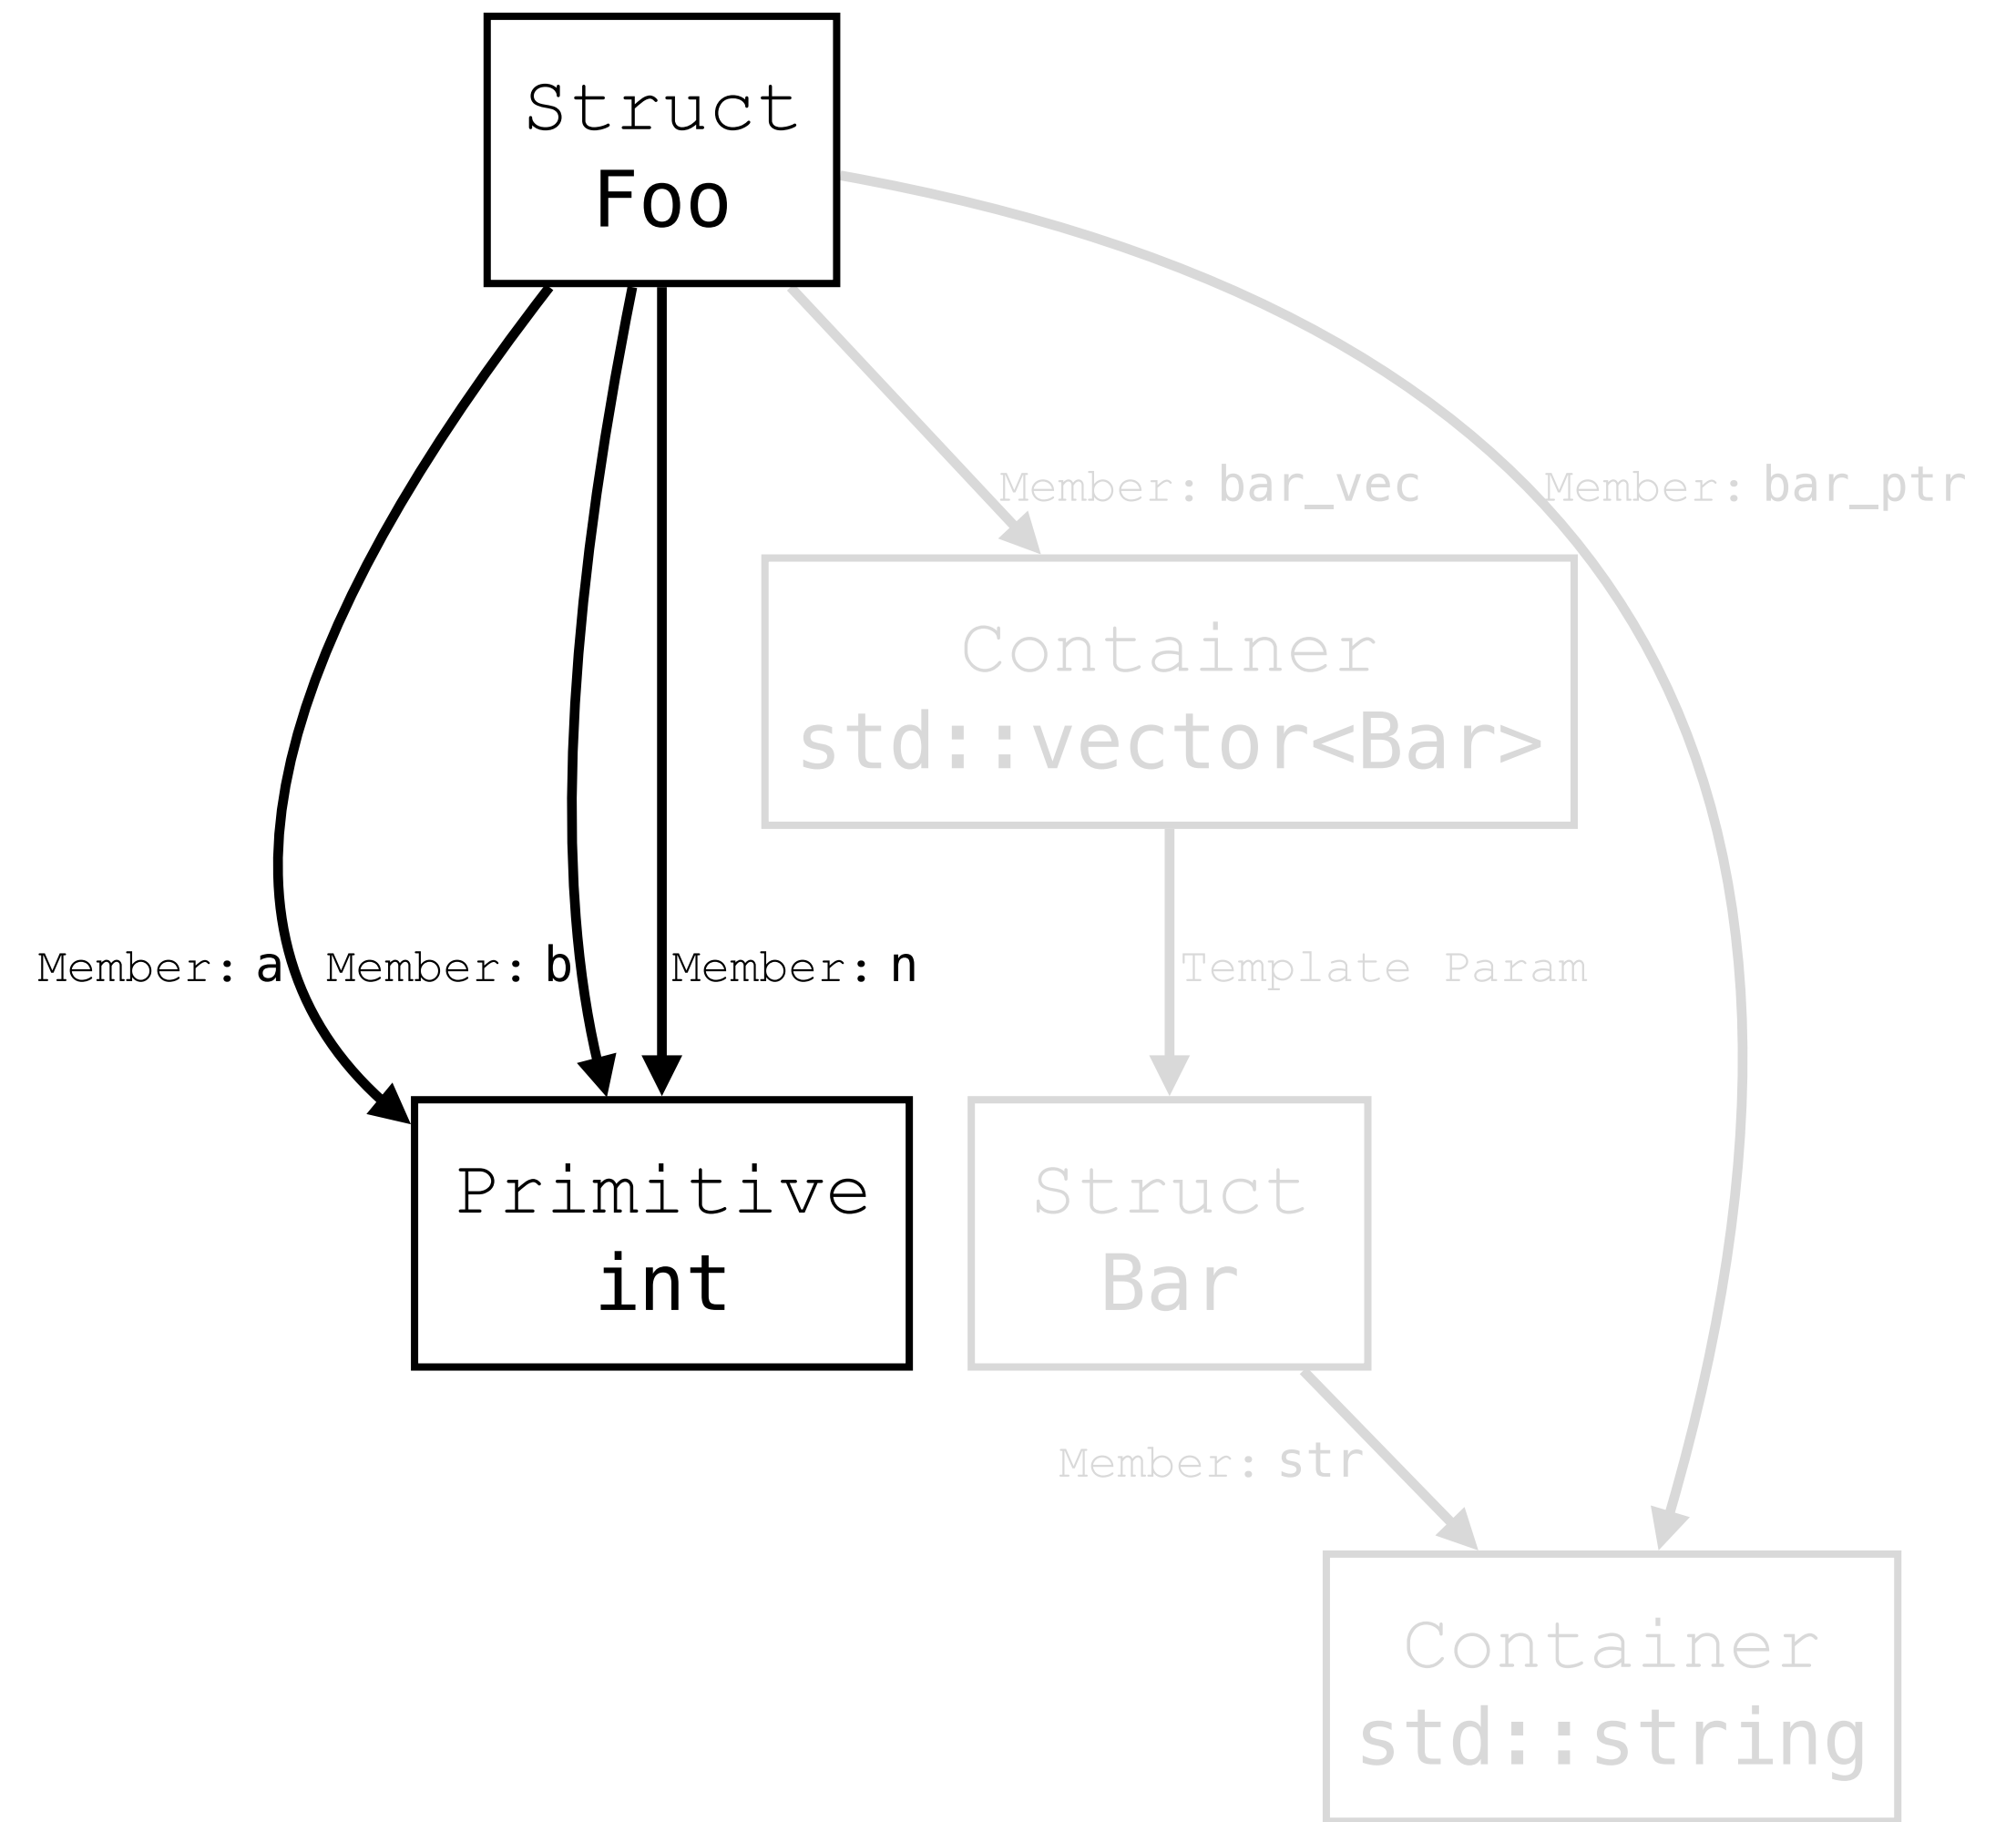
Type Graph Transformation: Inline Inheritance



Type Graph Transformation: Inline Inheritance



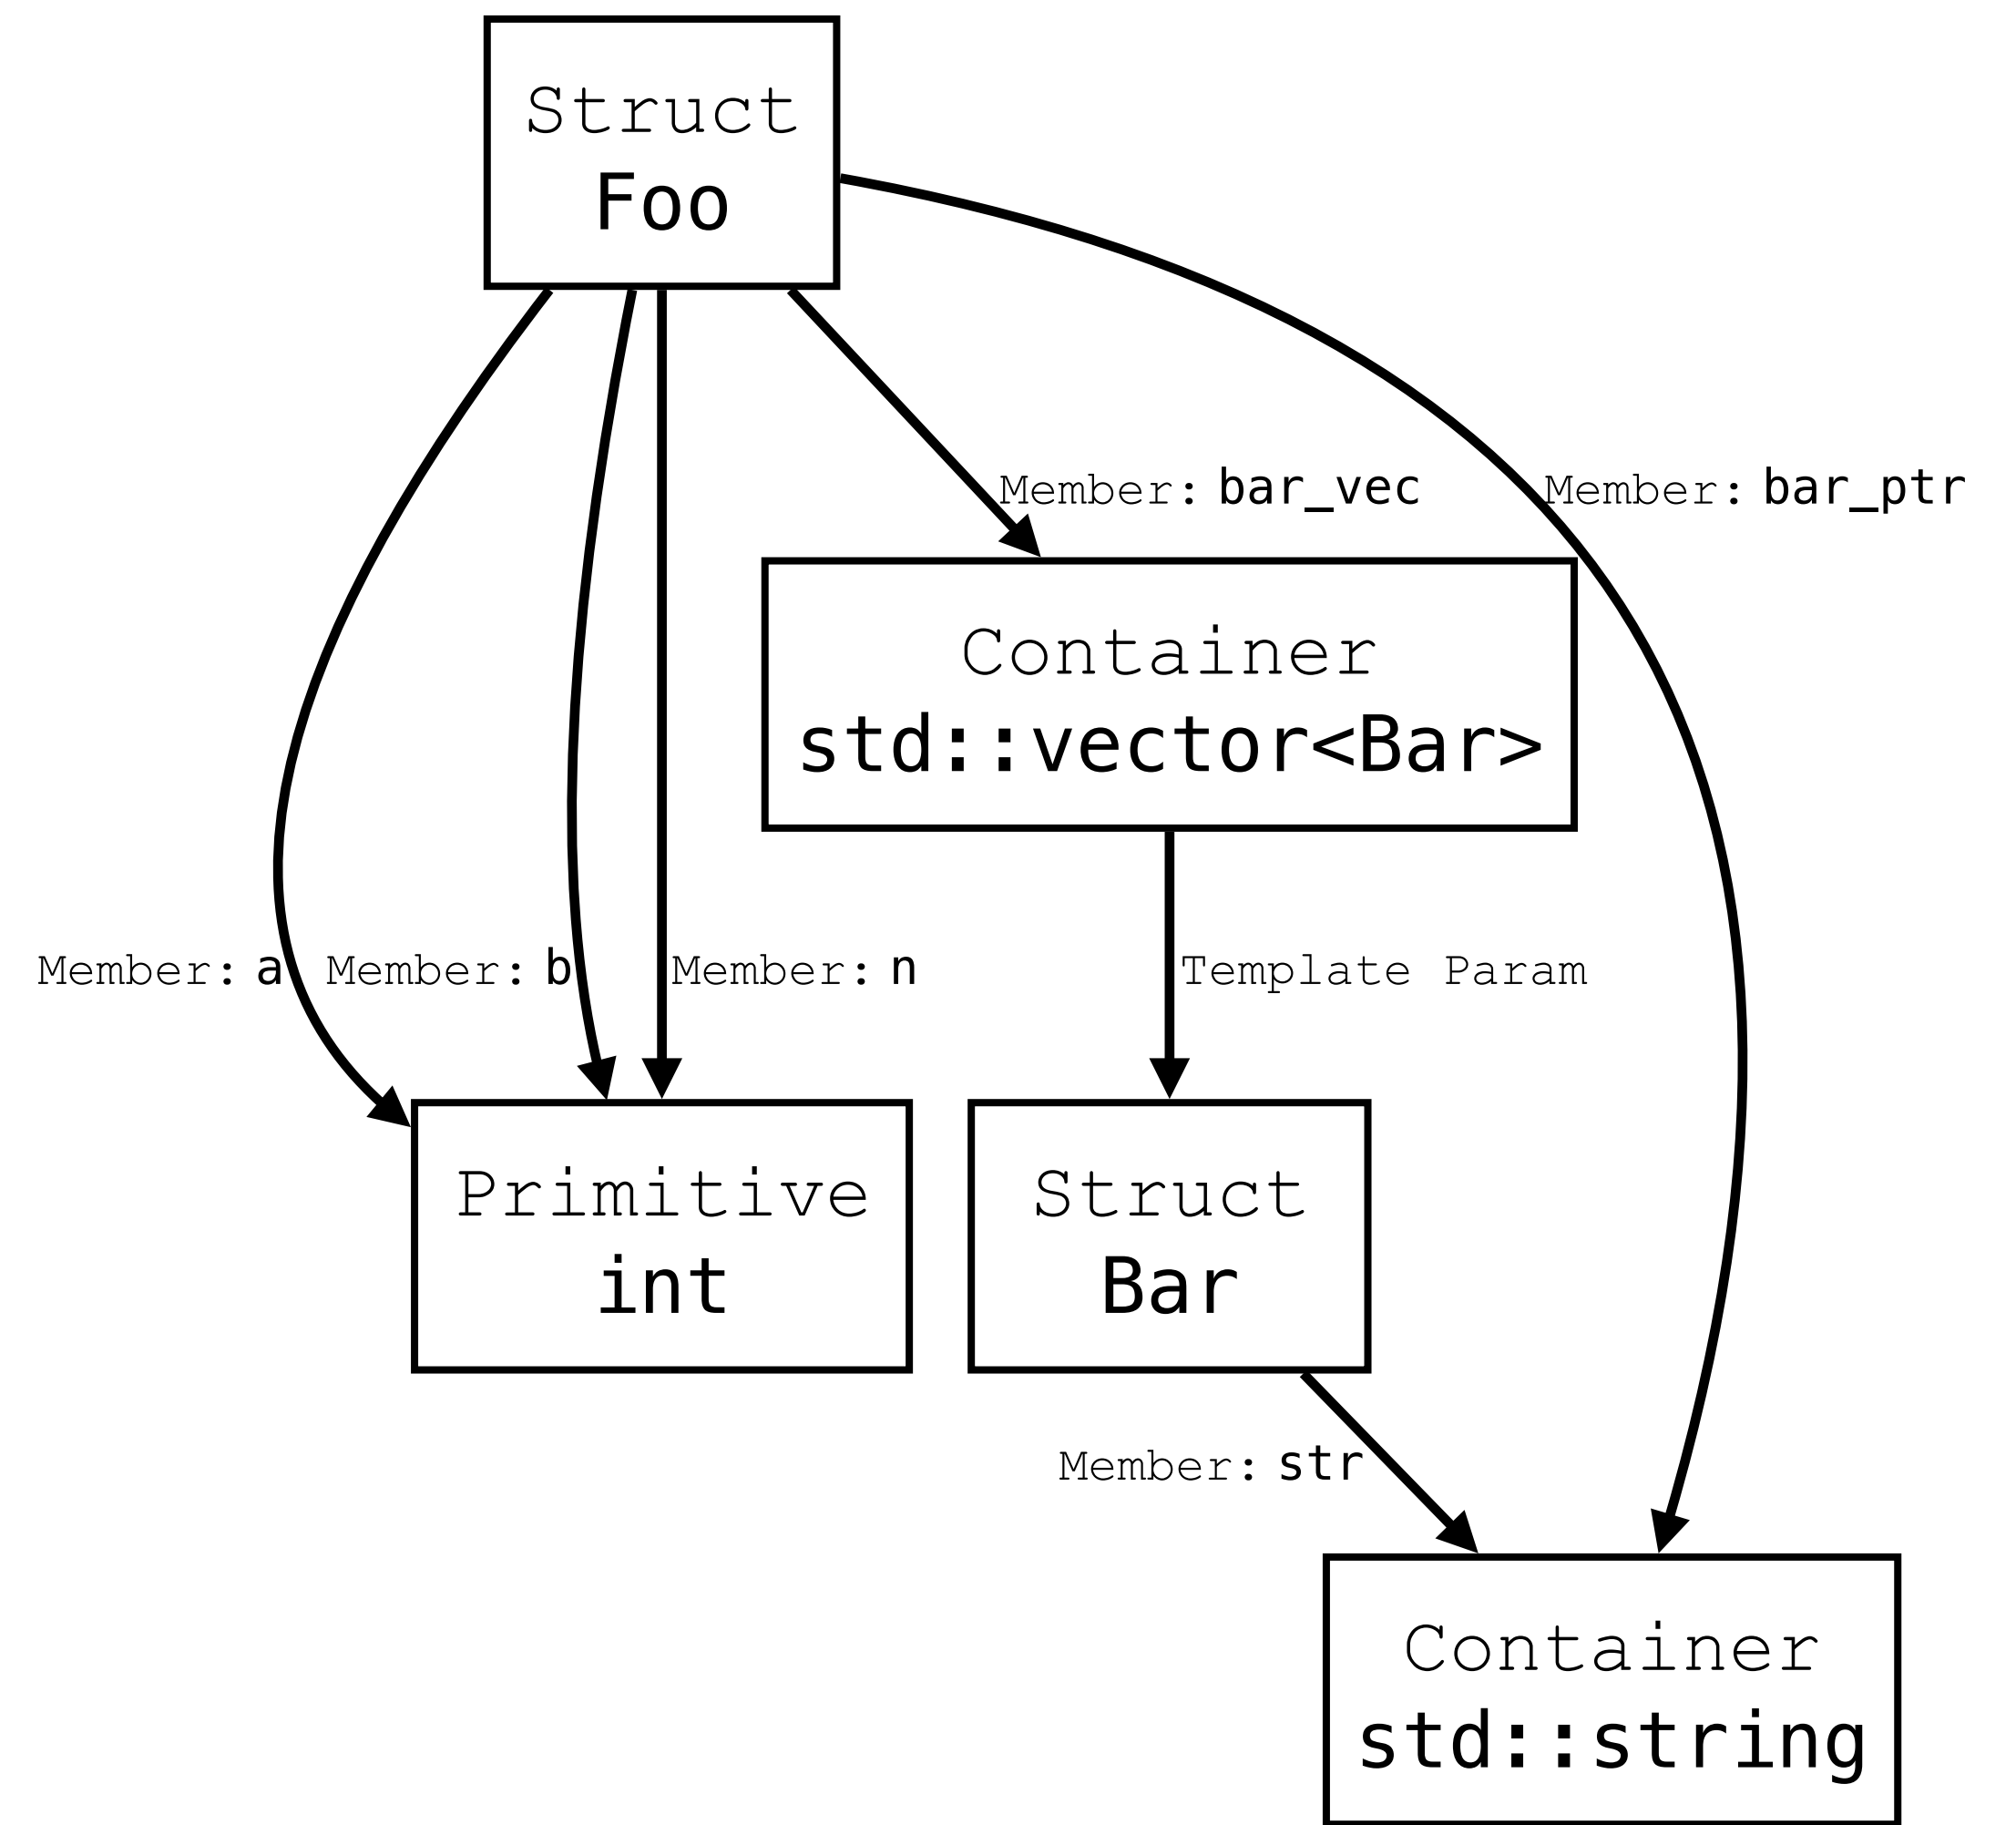
Type Graph Transformation: Inline Inheritance



Type Graph Transformation: Inline Inheritance

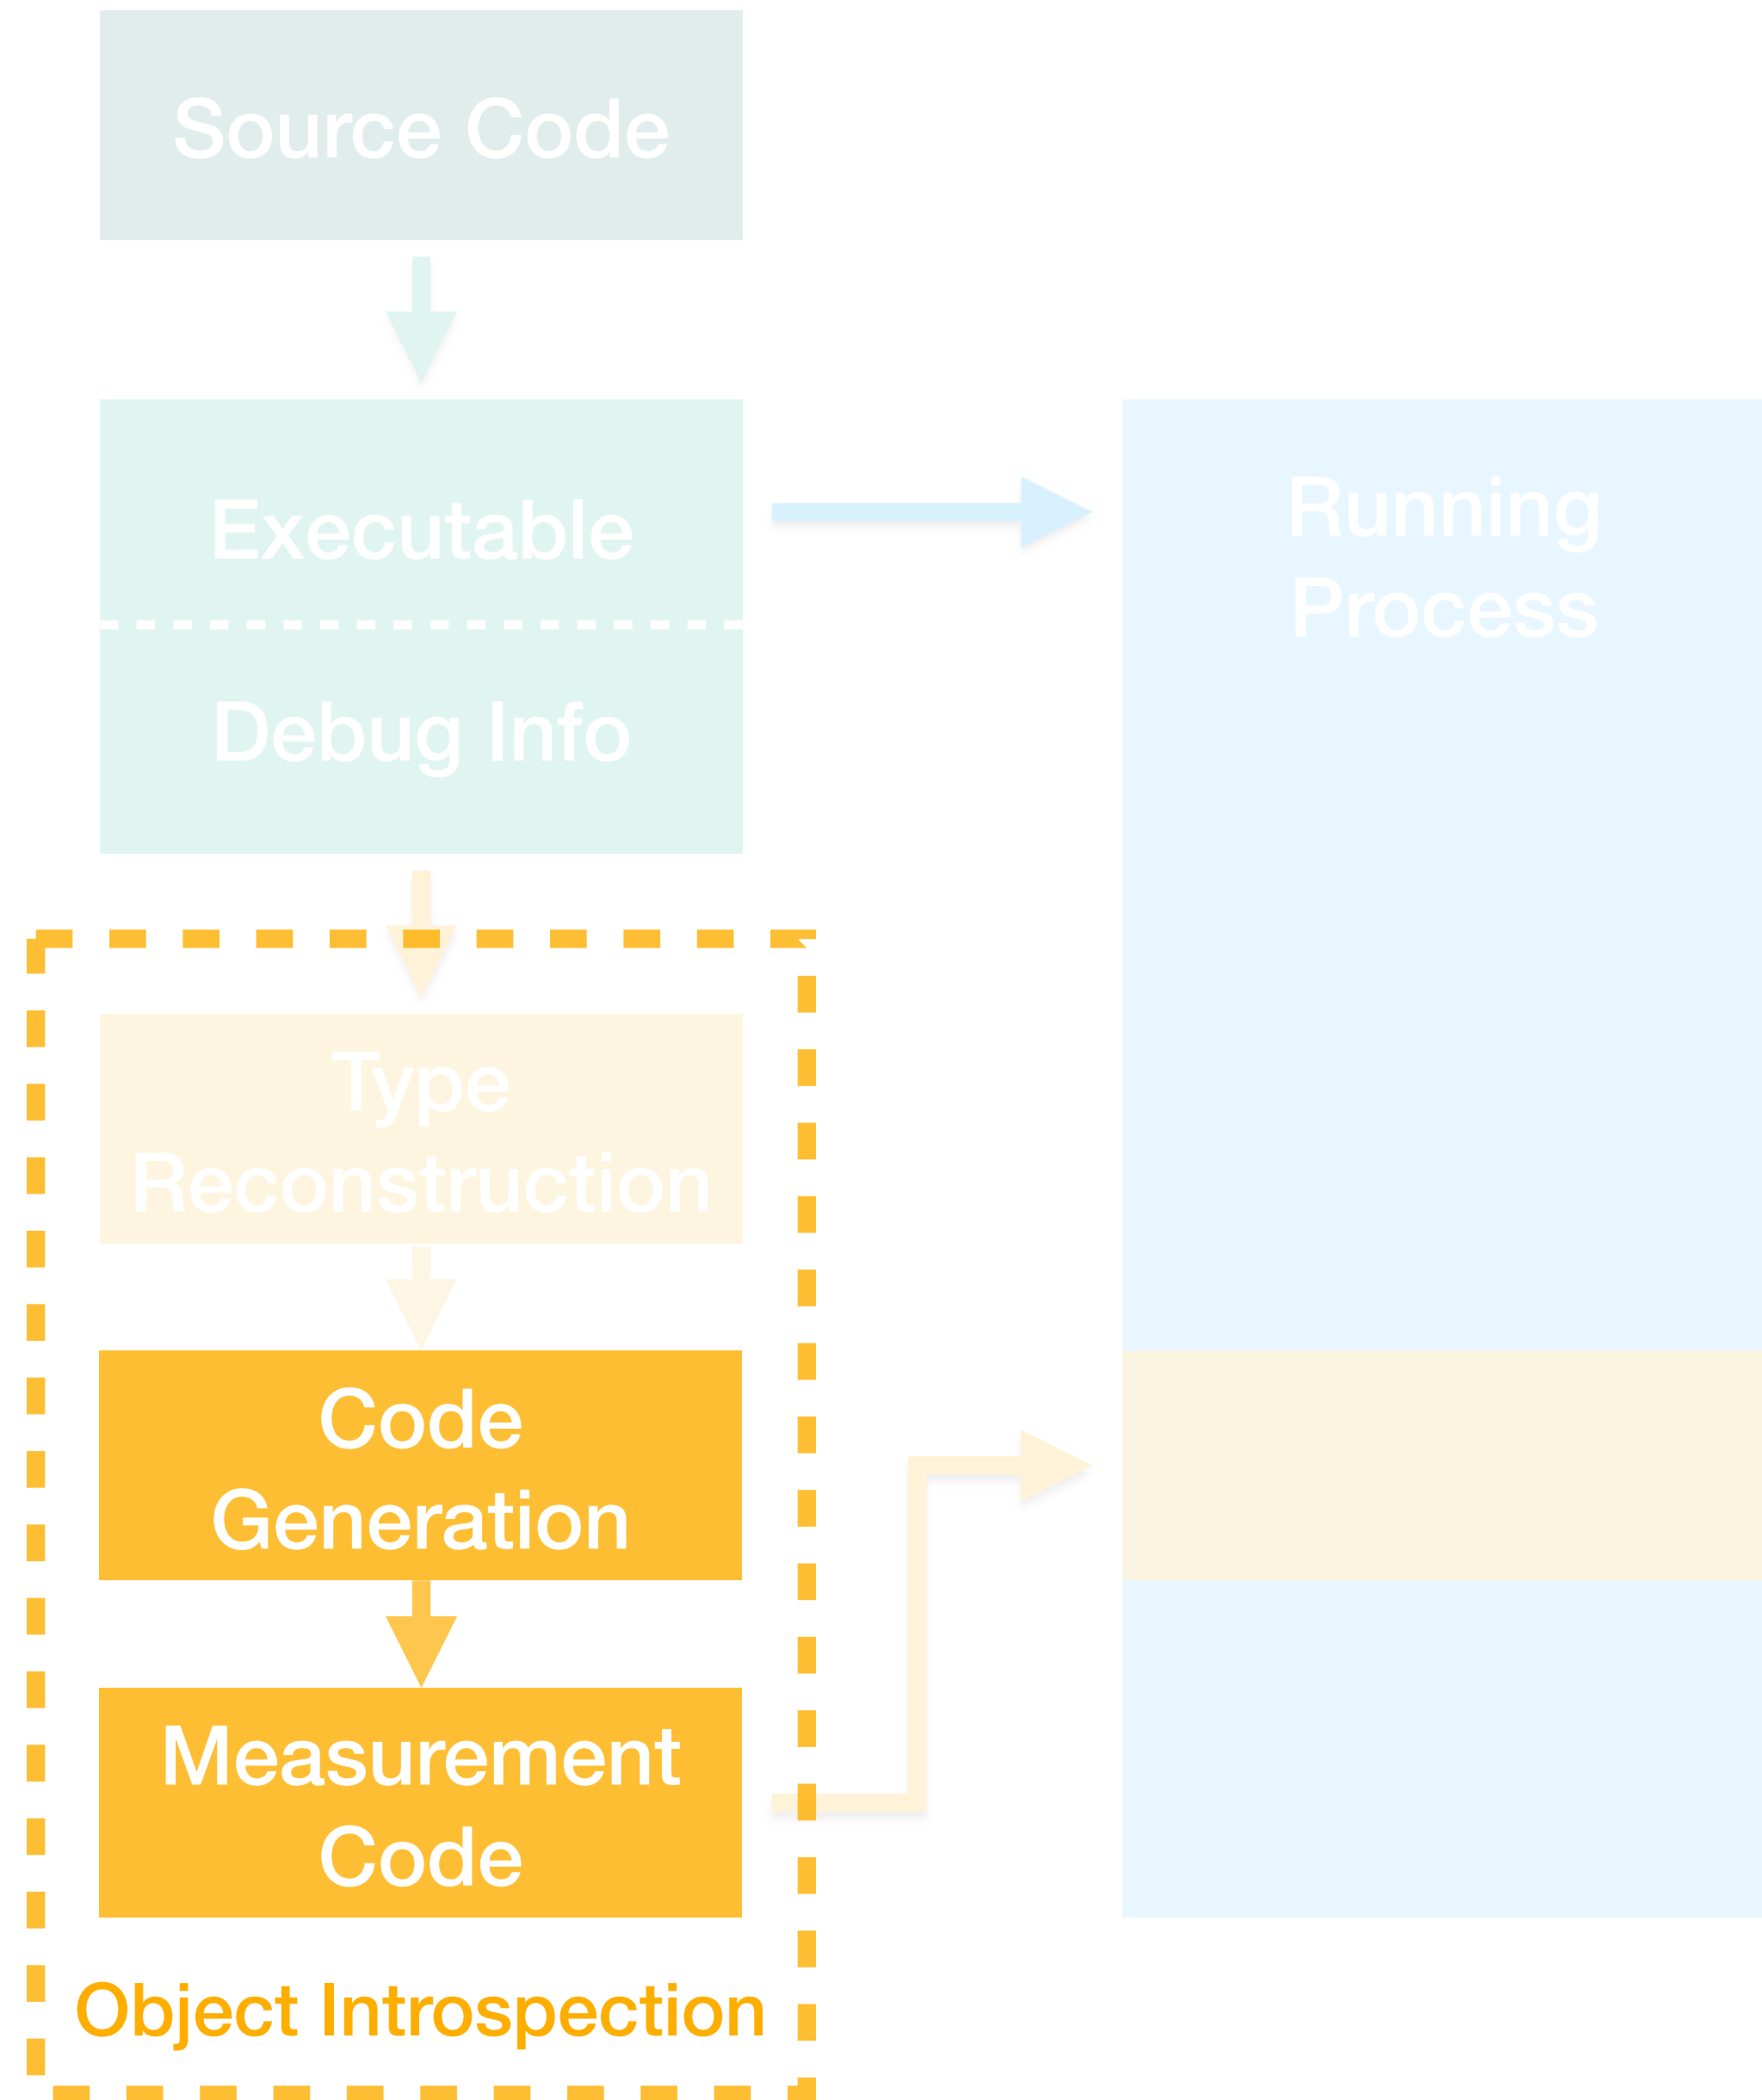
Reconstructed

```
struct Bar {  
    std::string str;  
};  
  
/* struct MyParent {}; */  
  
// Inline MyParent into Foo  
struct Foo {  
    int a; // from MyParent::a  
    int b; // from MyParent::b  
    int n;  
    std::vector<Bar> bar_vec;  
    std::string foo_str;  
};
```



Outline

- Debug Info Analysis
 - Type/Layout Reconstruction
- Code Generation
- Object Introspection as a Library
- Object Introspection as a Profiler
- Object Introspection Applied



CodeGen: class & struct

```
/* Generated from Debug Info */
static types::st::Unit<DB> getSizeType(
    const Foo& t,
    typename TypeHandler<DB, Foo>::type returnArg) {
    return returnArg
        .delegate( [&t](auto ret) { return getSizeType<DB>(t.a, ret); })
        .delegate( [&t](auto ret) { return getSizeType<DB>(t.b, ret); })
        .delegate( [&t](auto ret) { return getSizeType<DB>(t.n, ret); })
        .delegate( [&t](auto ret) { return getSizeType<DB>(t.bar_vec, ret); })
        .consume ( [&t](auto ret) { return getSizeType<DB>(t.foo_str, ret); });
}
```

```
struct Bar {
    std::string str;
};

struct Foo {
    int          a;
    int          b;
    int          n;
    std::vector<Bar> bar_vec;
    std::string  foo_str;
};
```

CodeGen: containers

```
template <class T0, class T1>  
class std::vector {...};
```

```
template <class T0, class T1>  
static types::st::Unit<DB> getSizeType(  
    const std::vector<T0, T1>& container,  
    typename TypeHandler<DB, std::vector<T0, T1>>::type returnArg) {  
    auto tail = returnArg  
        .write(container.capacity())  
        .write(container.size());  
    for (const auto &it : container) {  
        tail = tail.delegate([&it](auto ret) {  
            return getSizeType<DB>(it, ret);  
        });  
    }  
    return tail;  
}
```


CodeGen: class & struct

```
/* Generated from Debug Info */
static types::st::Unit<DB> getSizeType(
    const Bar& t,
    typename TypeHandler<DB, Bar>::type returnArg) {
    return returnArg
        .consume ([&t](auto ret) { return getSizeType<DB>(t.str, ret); });
}
```

```
struct Bar {
    std::string str;
};

struct Foo {
    int          a;
    int          b;
    int          n;
    std::vector<Bar> bar_vec;
    std::string  foo_str;
};
```

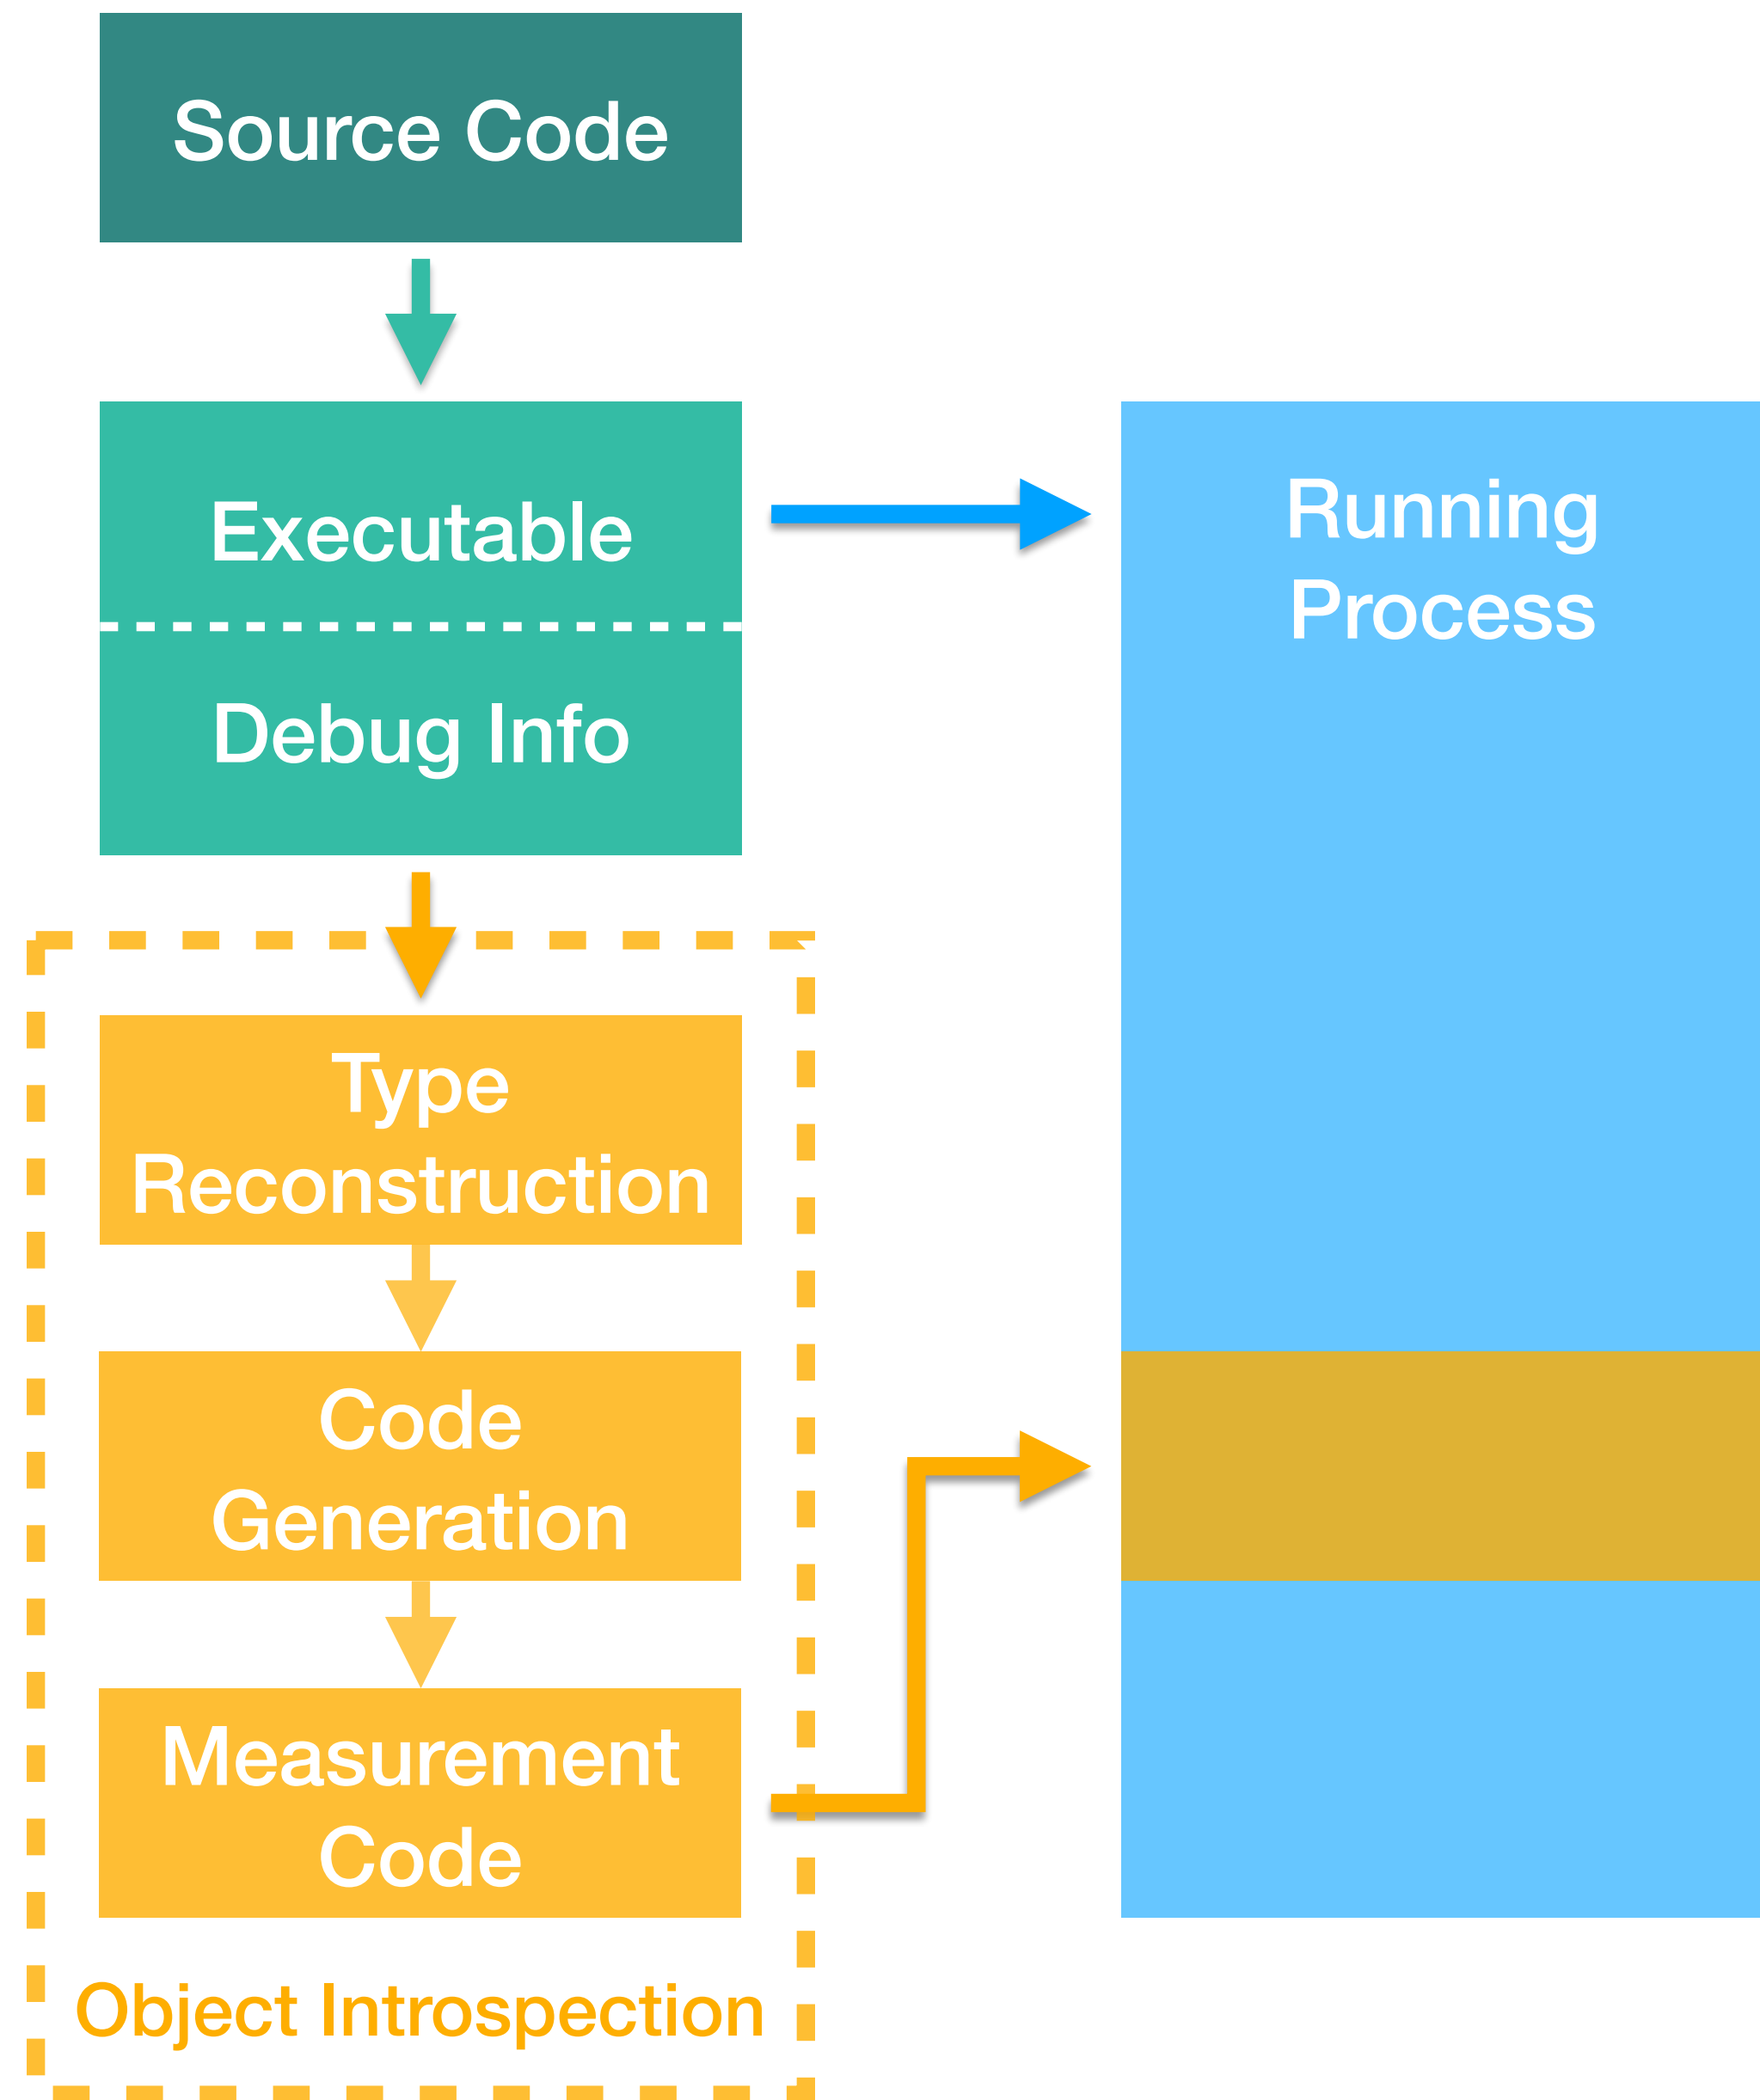
CodeGen: containers

```
template <class T0, class T1, class T2>
static types::st::Unit<DB> getSizeType(
    const std::basic_string<T0, T1, T2>& container,
    typename TypeHandler<DB, std::basic_string<T0, T1, T2>>::type returnArg) {
    bool sso = container.data() >= &container
        && container.data() < (&container + sizeof(container));
    return returnArg
        .write(container.capacity())
        .write(sso)
        .write(container.size());
}
```

```
template <class T0, class T1, class T2>
class std::basic_string {...};
```

Outline

- Debug Info Analysis
 - Type/Layout Reconstruction
 - Code Generation
- Object Introspection as a Library
- Object Introspection as a Profiler
- Object Introspection Applied



Object Introspection as a Library

```
Bar example(const Foo &f) {  
    auto result = oi::setupAndIntrospect(f, opts);  
    if (result) {  
        oi::exporters::Json json(std::cout);  
        json.setPretty(true);  
        json.print(*result);  
    }  
}
```

JSON Output

```
{ "name": "foo",  
  "typePath": ["foo"],  
  "typeNames": ["Foo"],  
  "staticSize": 80,  
  "exclusiveSize": 4,  
  "members": [  
    { "name": "a",  
      "typePath": ["foo", "a"],  
      "typeNames": ["int32_t"],  
      "staticSize": 4,  
      "exclusiveSize": 4,  
    }, {  
      "name": "b",  
      "typePath": ["foo", "b"],  
      [...]
  ]
}
```

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int          a;  
    int          b;  
    int          n;  
    std::vector<Bar> bar_vec;  
    std::string   foo_str;  
};
```

JSON Output

```
[...],
{ "name": "bar_vec",
  "typePath": ["foo", "bar_vec"],
  "typeNames": ["std::vector<Bar, std::allocator<Bar>>"],
  "length": 155,
  "capacity": 256,
  "staticSize": 24,
  "exclusiveSize": 3256,
  "members": [
    { "name": "[]",
      "typePath": ["foo", "bar_vec", "[]"],
      "typeNames": ["Bar"],
      "staticSize": 32,
      "exclusiveSize": 0,
      "members": [
        { "name": "str",
```

```
struct Bar {
    std::string str;
};
```

```
struct Foo {
    int          a;
    int          b;
    int          n;
    std::vector<Bar> bar_vec;
    std::string   foo_str;
};
```

JSON Output

```
{ "name": "[]",  
  "typePath": ["foo", "bar_vec", "[]"],  
  "typeNames": ["Bar"],  
  "staticSize": 32,  
  "exclusiveSize": 0,  
  "members": [  
    { "name": "str",  
      "typePath": ["foo", "bar_vec", "[]", "str"],  
      "typeNames": ["string", "std::basic_string<int8_t,  
std::char_traits<int8_t>, std::allocator<int8_t>>"],  
      "length": 12,  
      "capacity": 15,  
      "staticSize": 32,  
      "exclusiveSize": 32  
    }  
  ]  
}
```

```
struct Bar {  
    std::string str;  
};
```

```
struct Foo {  
    int          a;  
    int          b;  
    int          n;  
    std::vector<Bar> bar_vec;  
    std::string   foo_str;  
};
```

Result Iterator

```
IntrospectionResult result =  
    oi::setupAndIntrospect(f, opts);
```

```
struct Bar {  
    std::string    str;  
};  
  
struct Foo {  
    int            a;  
    int            b;  
    int            n;  
    std::vector<Bar> bar_vec;  
    std::string    foo_str;  
};
```

begin()

Element 0

Element 1

Element 2

Element 3

Element 4

...

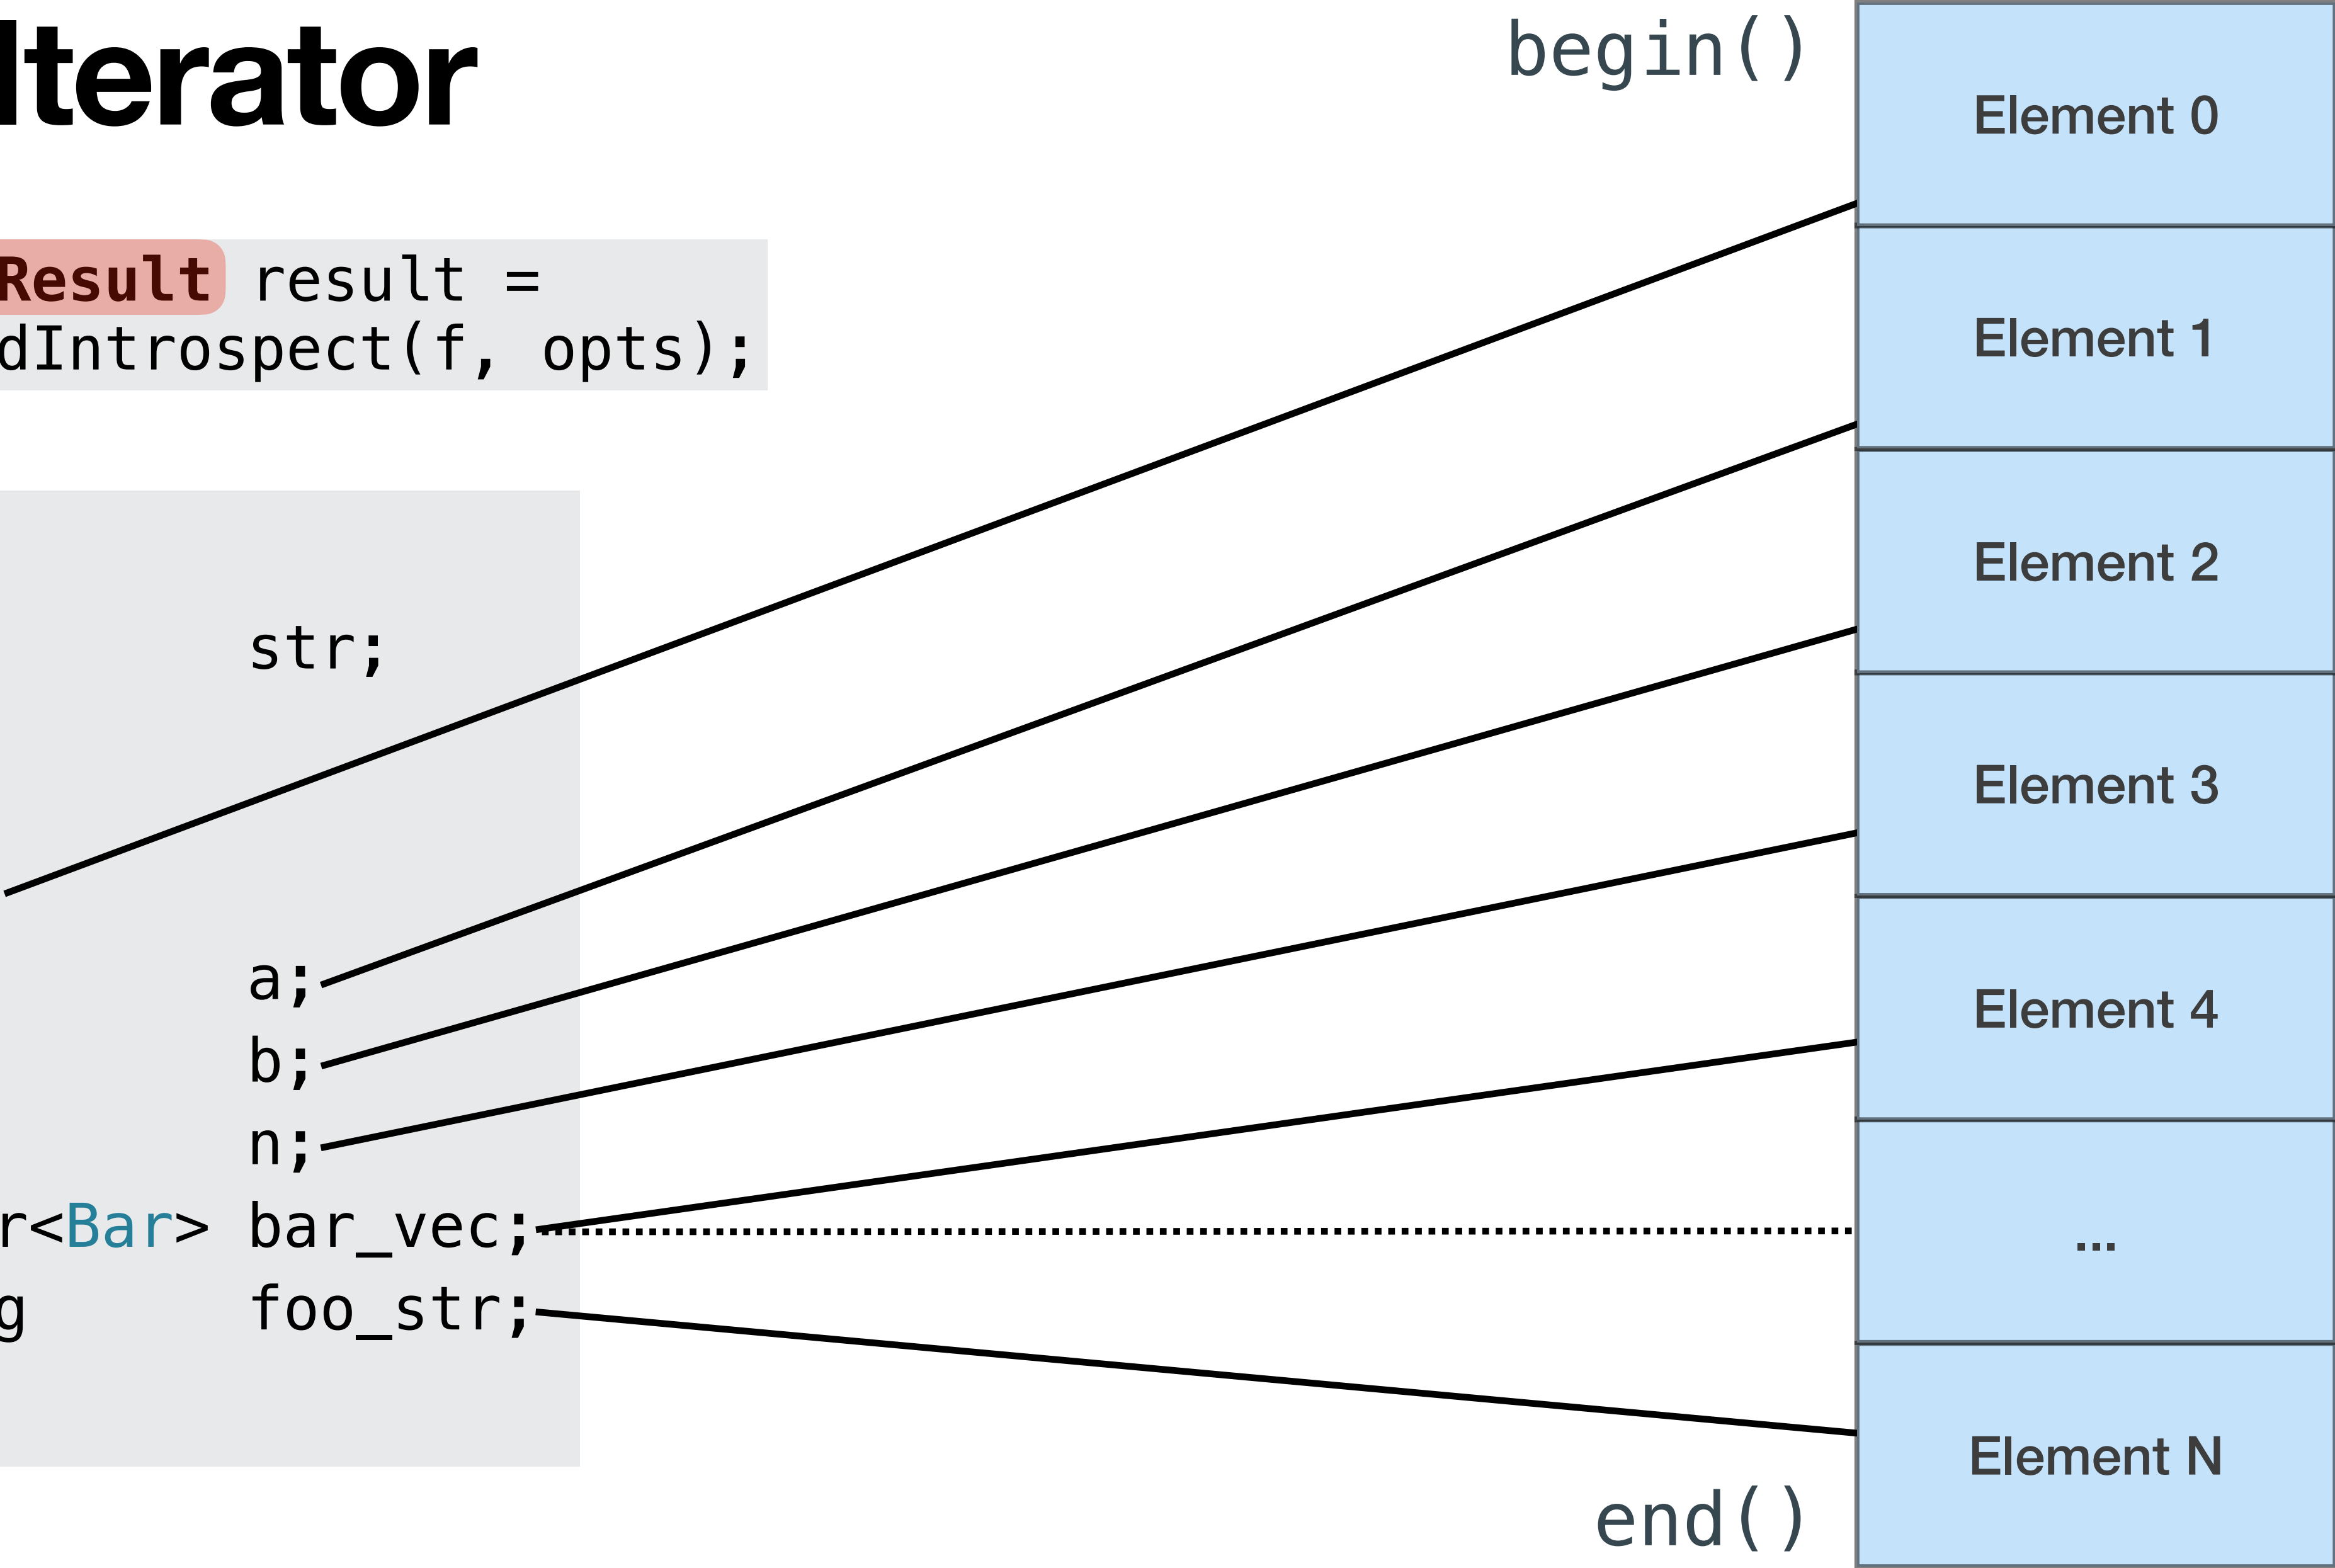
Element N

end()

Result Iterator

```
IntrospectionResult result =  
oi::setupAndIntrospect(f, opts);
```

```
struct Bar {  
    std::string    str;  
};  
  
struct Foo {  
    int            a;  
    int            b;  
    int            n;  
    std::vector<Bar> bar_vec;  
    std::string    foo_str;  
};
```



OIL: String Hunting

```
std::cout << "Field Name  Length Capacity\n";
std::cout << "=====  =====  =====\n";
for (const auto &el: results) {
    if (std::range::find(el.type_names, "std::string") != el.type_names.end()) {
        std::cout << std::format("{:<10}  {:>6} {:>8}\n",
            el.name, el.container_stats->length, el.container_stats->capacity);
    }
}
```

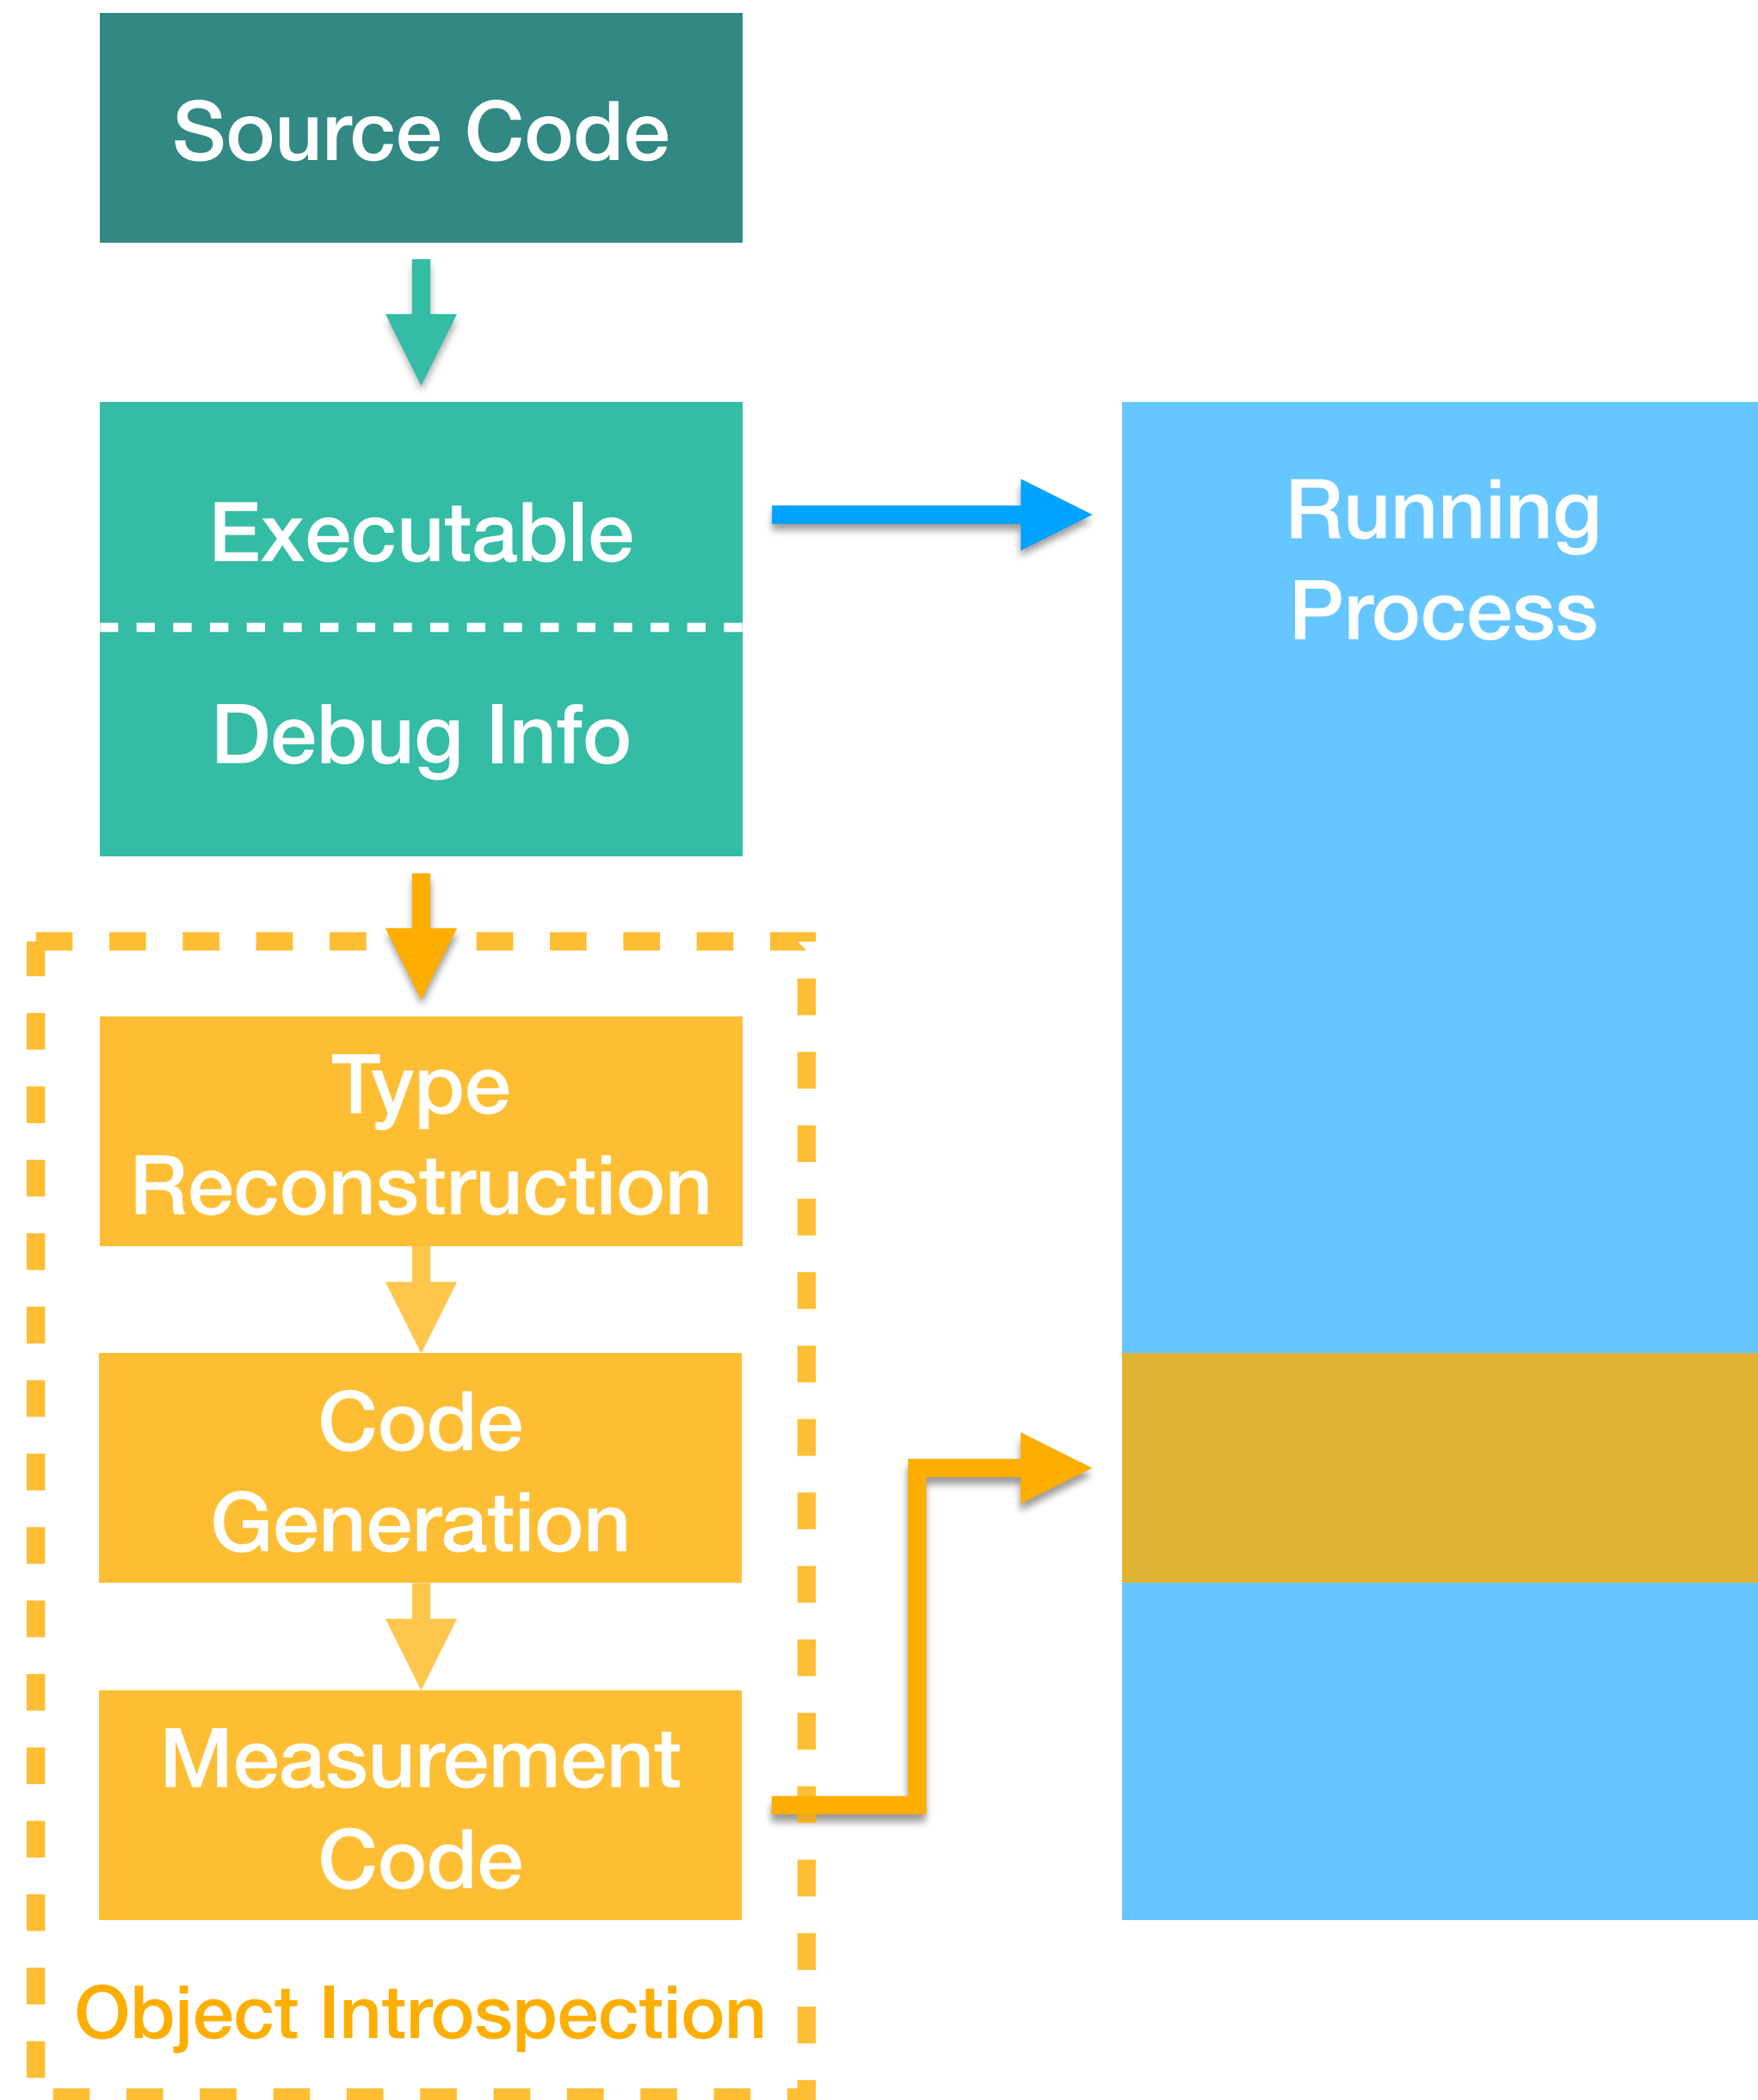
```
struct Bar {
    std::string str;
};

struct Foo {
    int          a;
    int          b;
    int          n;
    std::vector<Bar> bar_vec;
    std::string   foo_str;
};
```

Field Name	Length	Capacity
=====	=====	=====
str	12	15
str	85	85
str	20	20
str	57	57
str	46	46
str	69	69
foo_str	0	15

Outline

- Debug Info Analysis
 - Type/Layout Reconstruction
 - Code Generation
- Object Introspection as a Library
- **Object Introspection as a Profiler**
- Object Introspection Applied



Object Introspection as a Profiler

- Classic 'tracer-tracee' style application
 - Type reconstruction / Code Generation
 - Code injection
 - Process/thread management
 - Result processing
- Non-interactive, configuration driven

Running the Profiler

```
Bar example(Foo &f);
```

Function arguments at entry and return

```
entry:_Z7exampleR3Foo:arg0
```

```
return:_Z7exampleR3Foo:arg0
```

```
return:_Z7exampleR3Foo:retval
```

```
$ oid --json --pid `pidof target` -S "entry:_Z7exampleR3Foo:arg0"
```

Probe examples

Multiple arguments

```
return:_Z8doitR3FooS0_:arg0,arg1 // void doit(Foo&, Foo&)
```

‘this’ pointer

```
entry:_ZN3FooD1Ev:this           // Foo::~~Foo()
```

Global data

```
global:myGlobalState
```

Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data

oid Process



Target Process



Bar example(Foo&)

```
example+0    pushq  
example+2    pushq  
example+4    ...
```

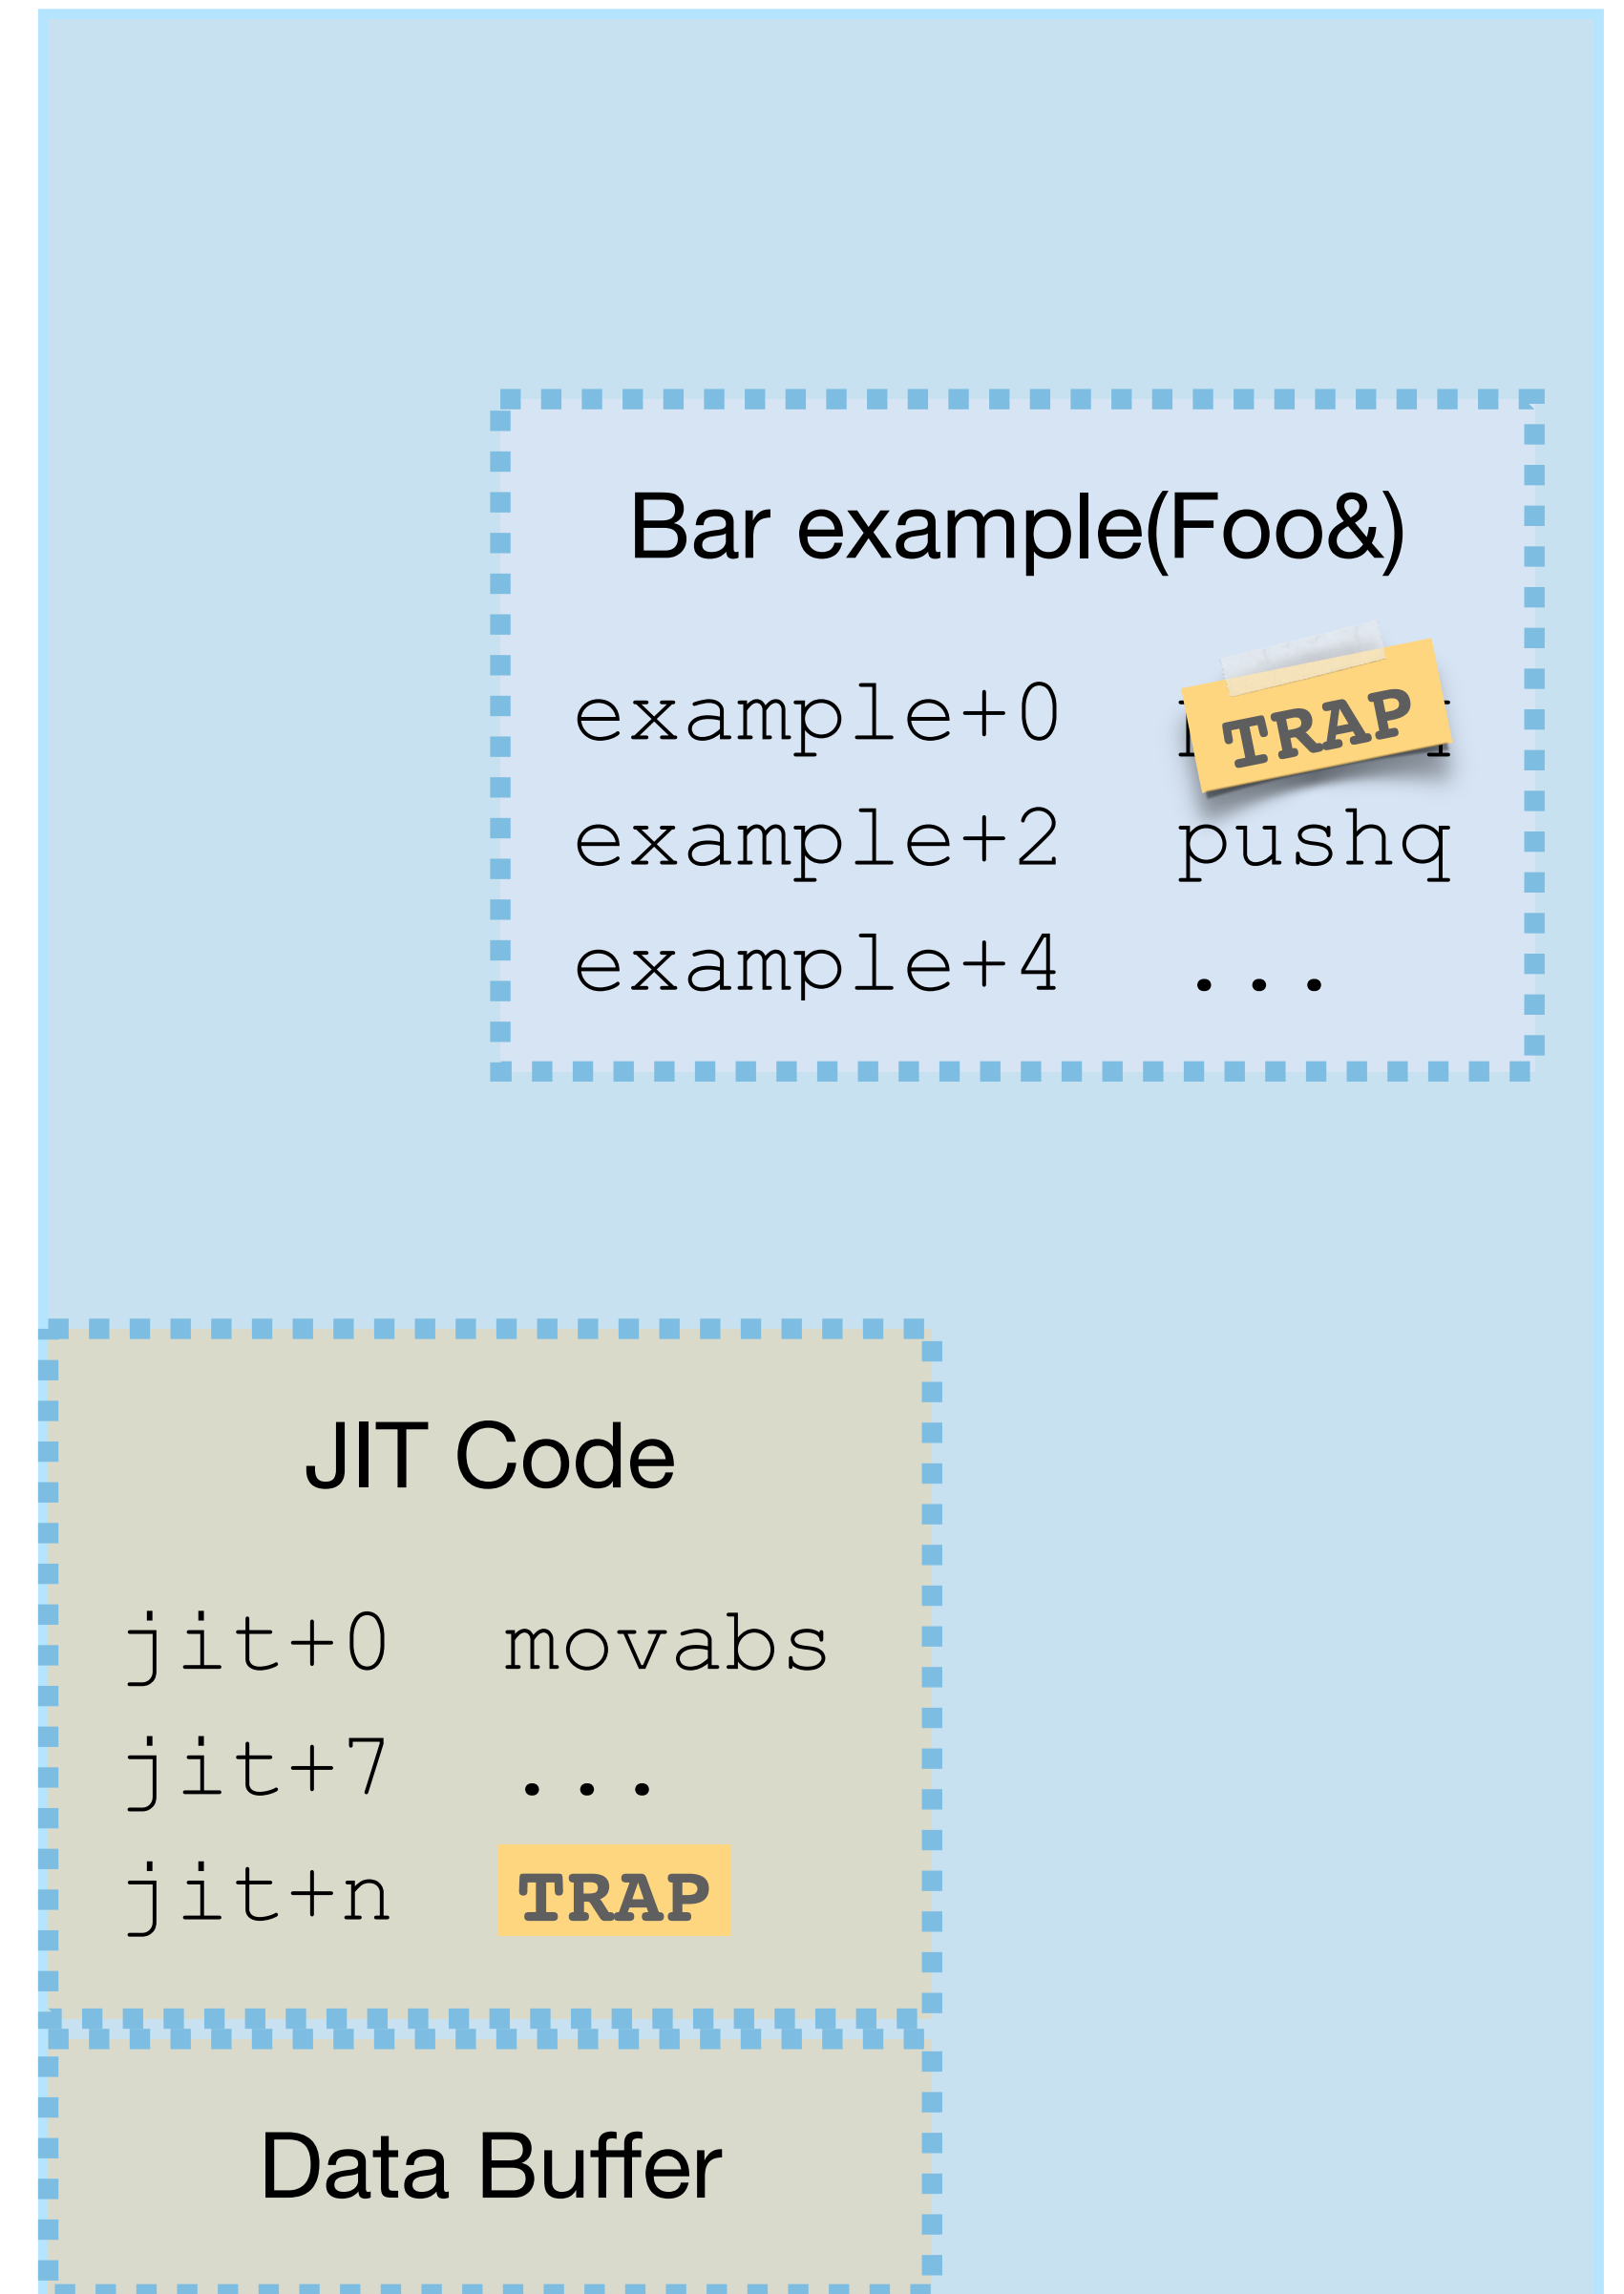
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data

oid Process



Target Process



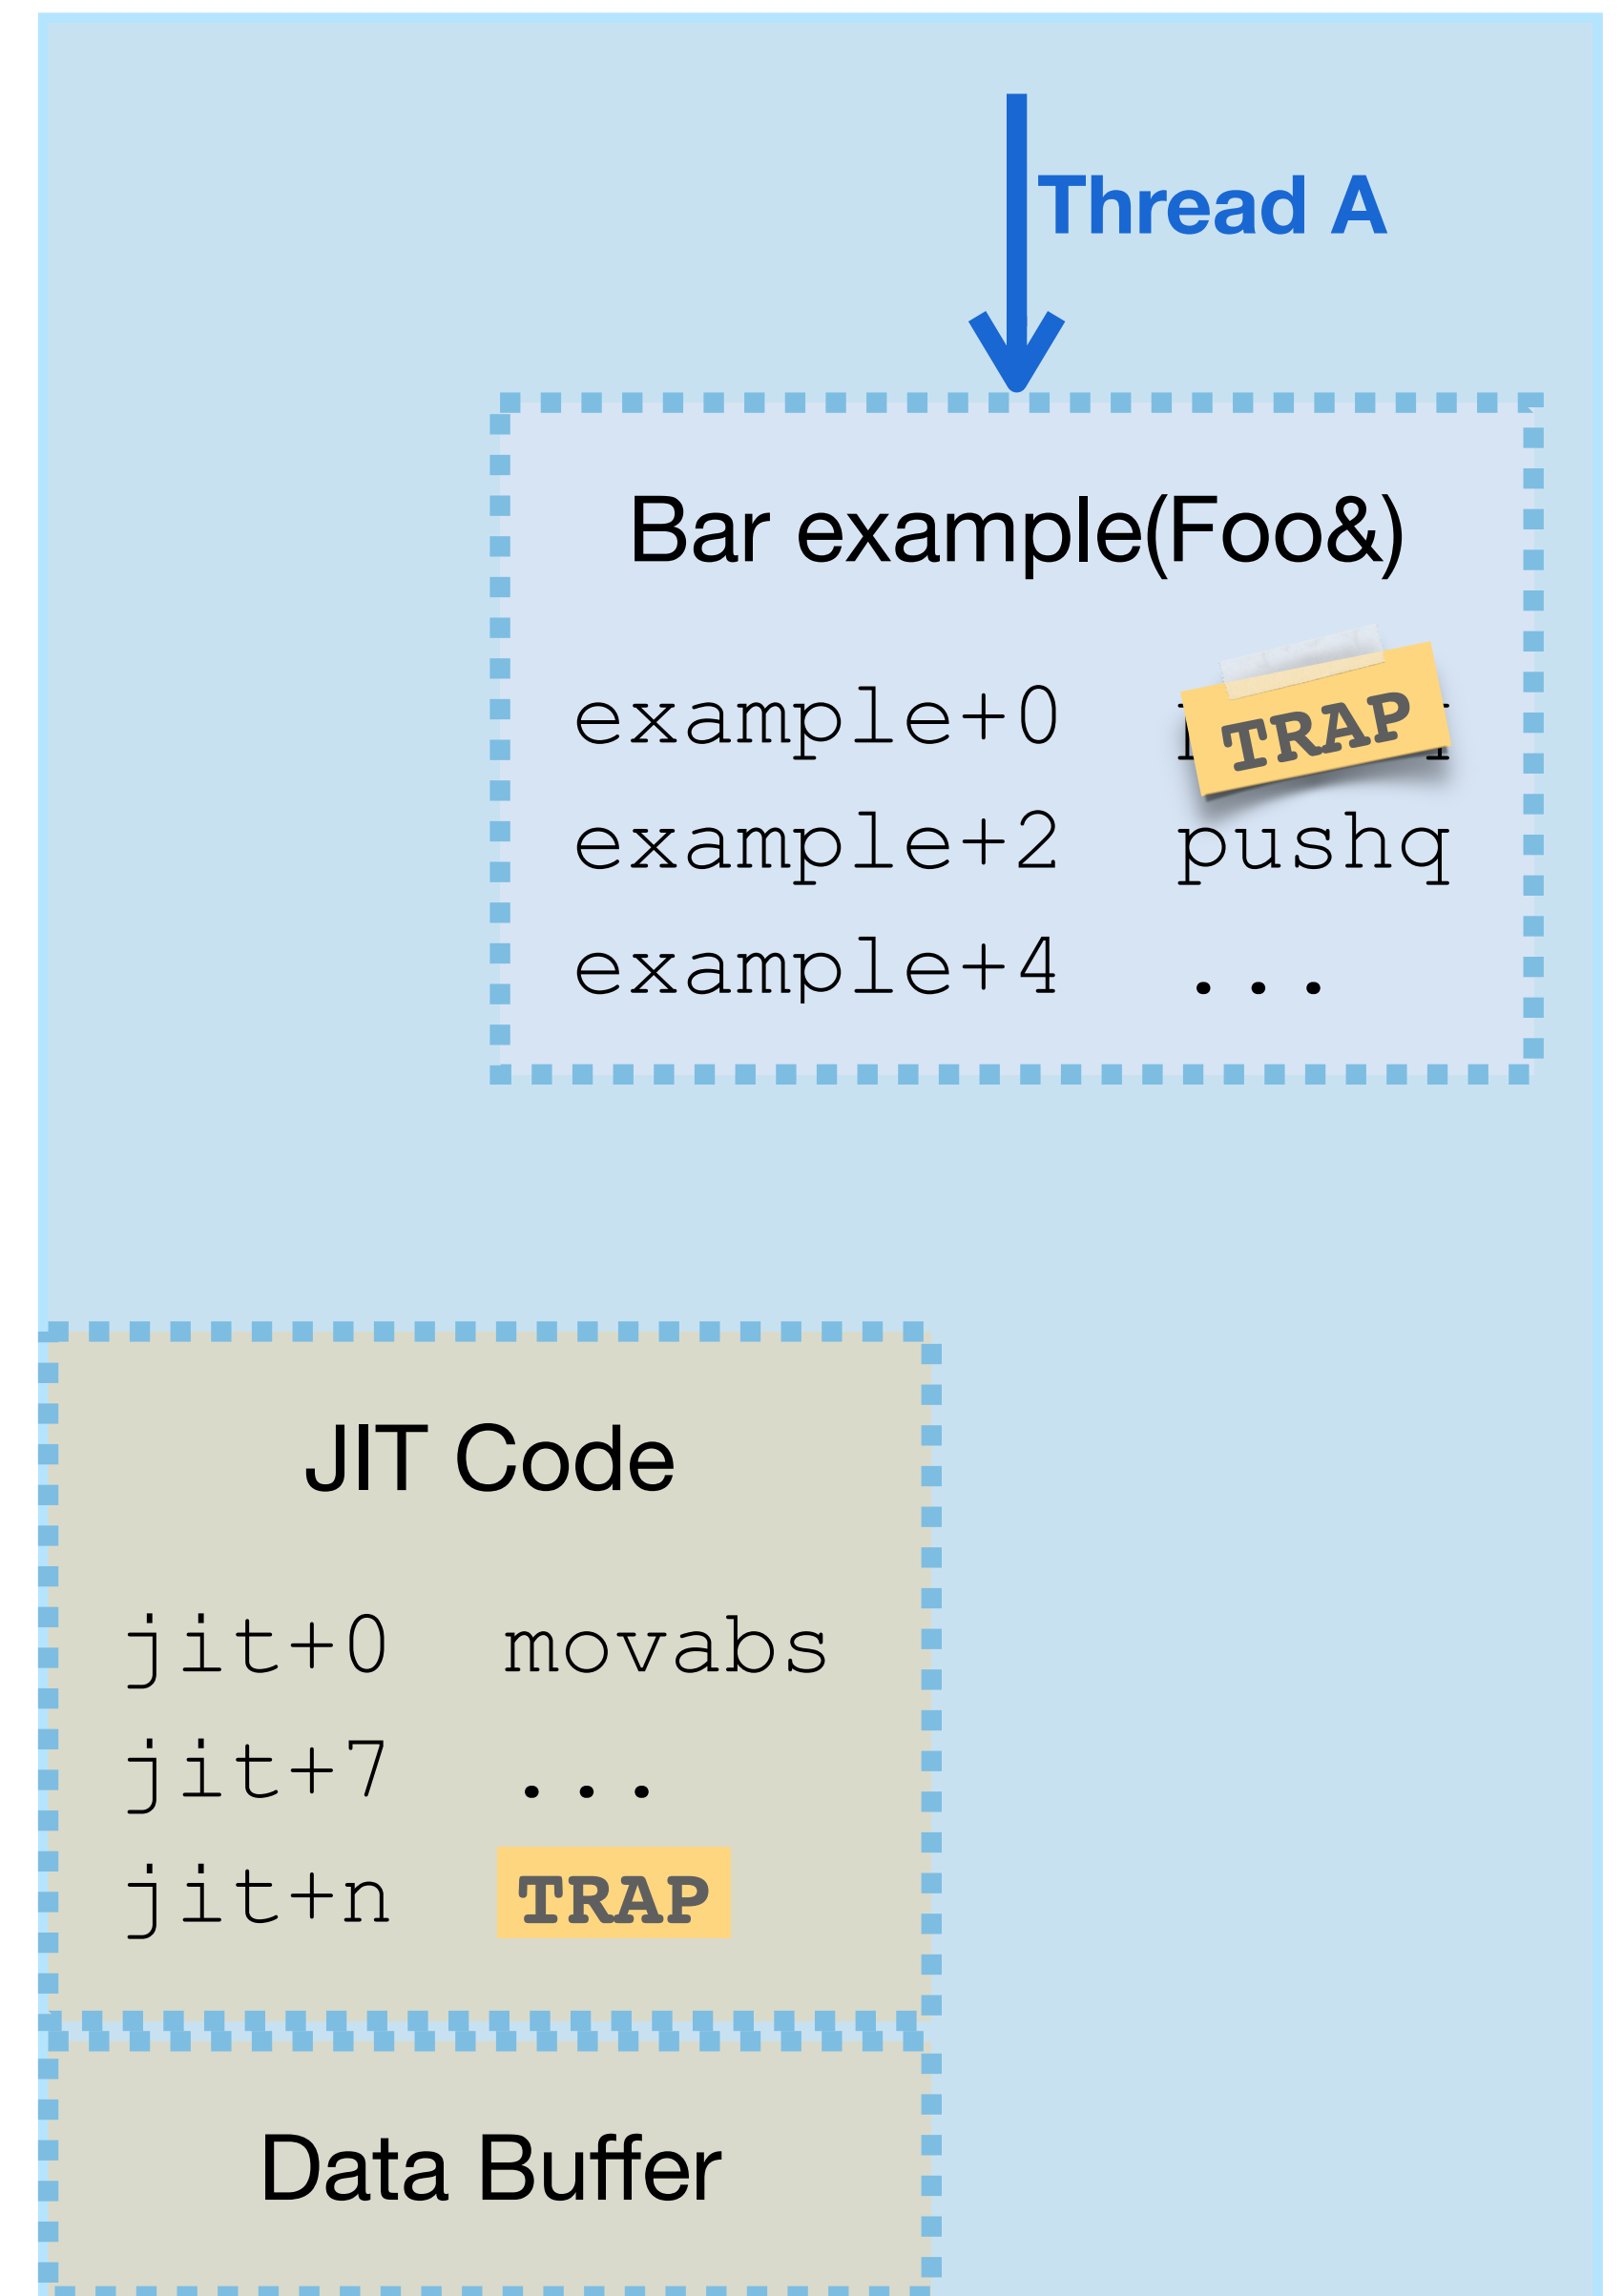
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data

oid Process

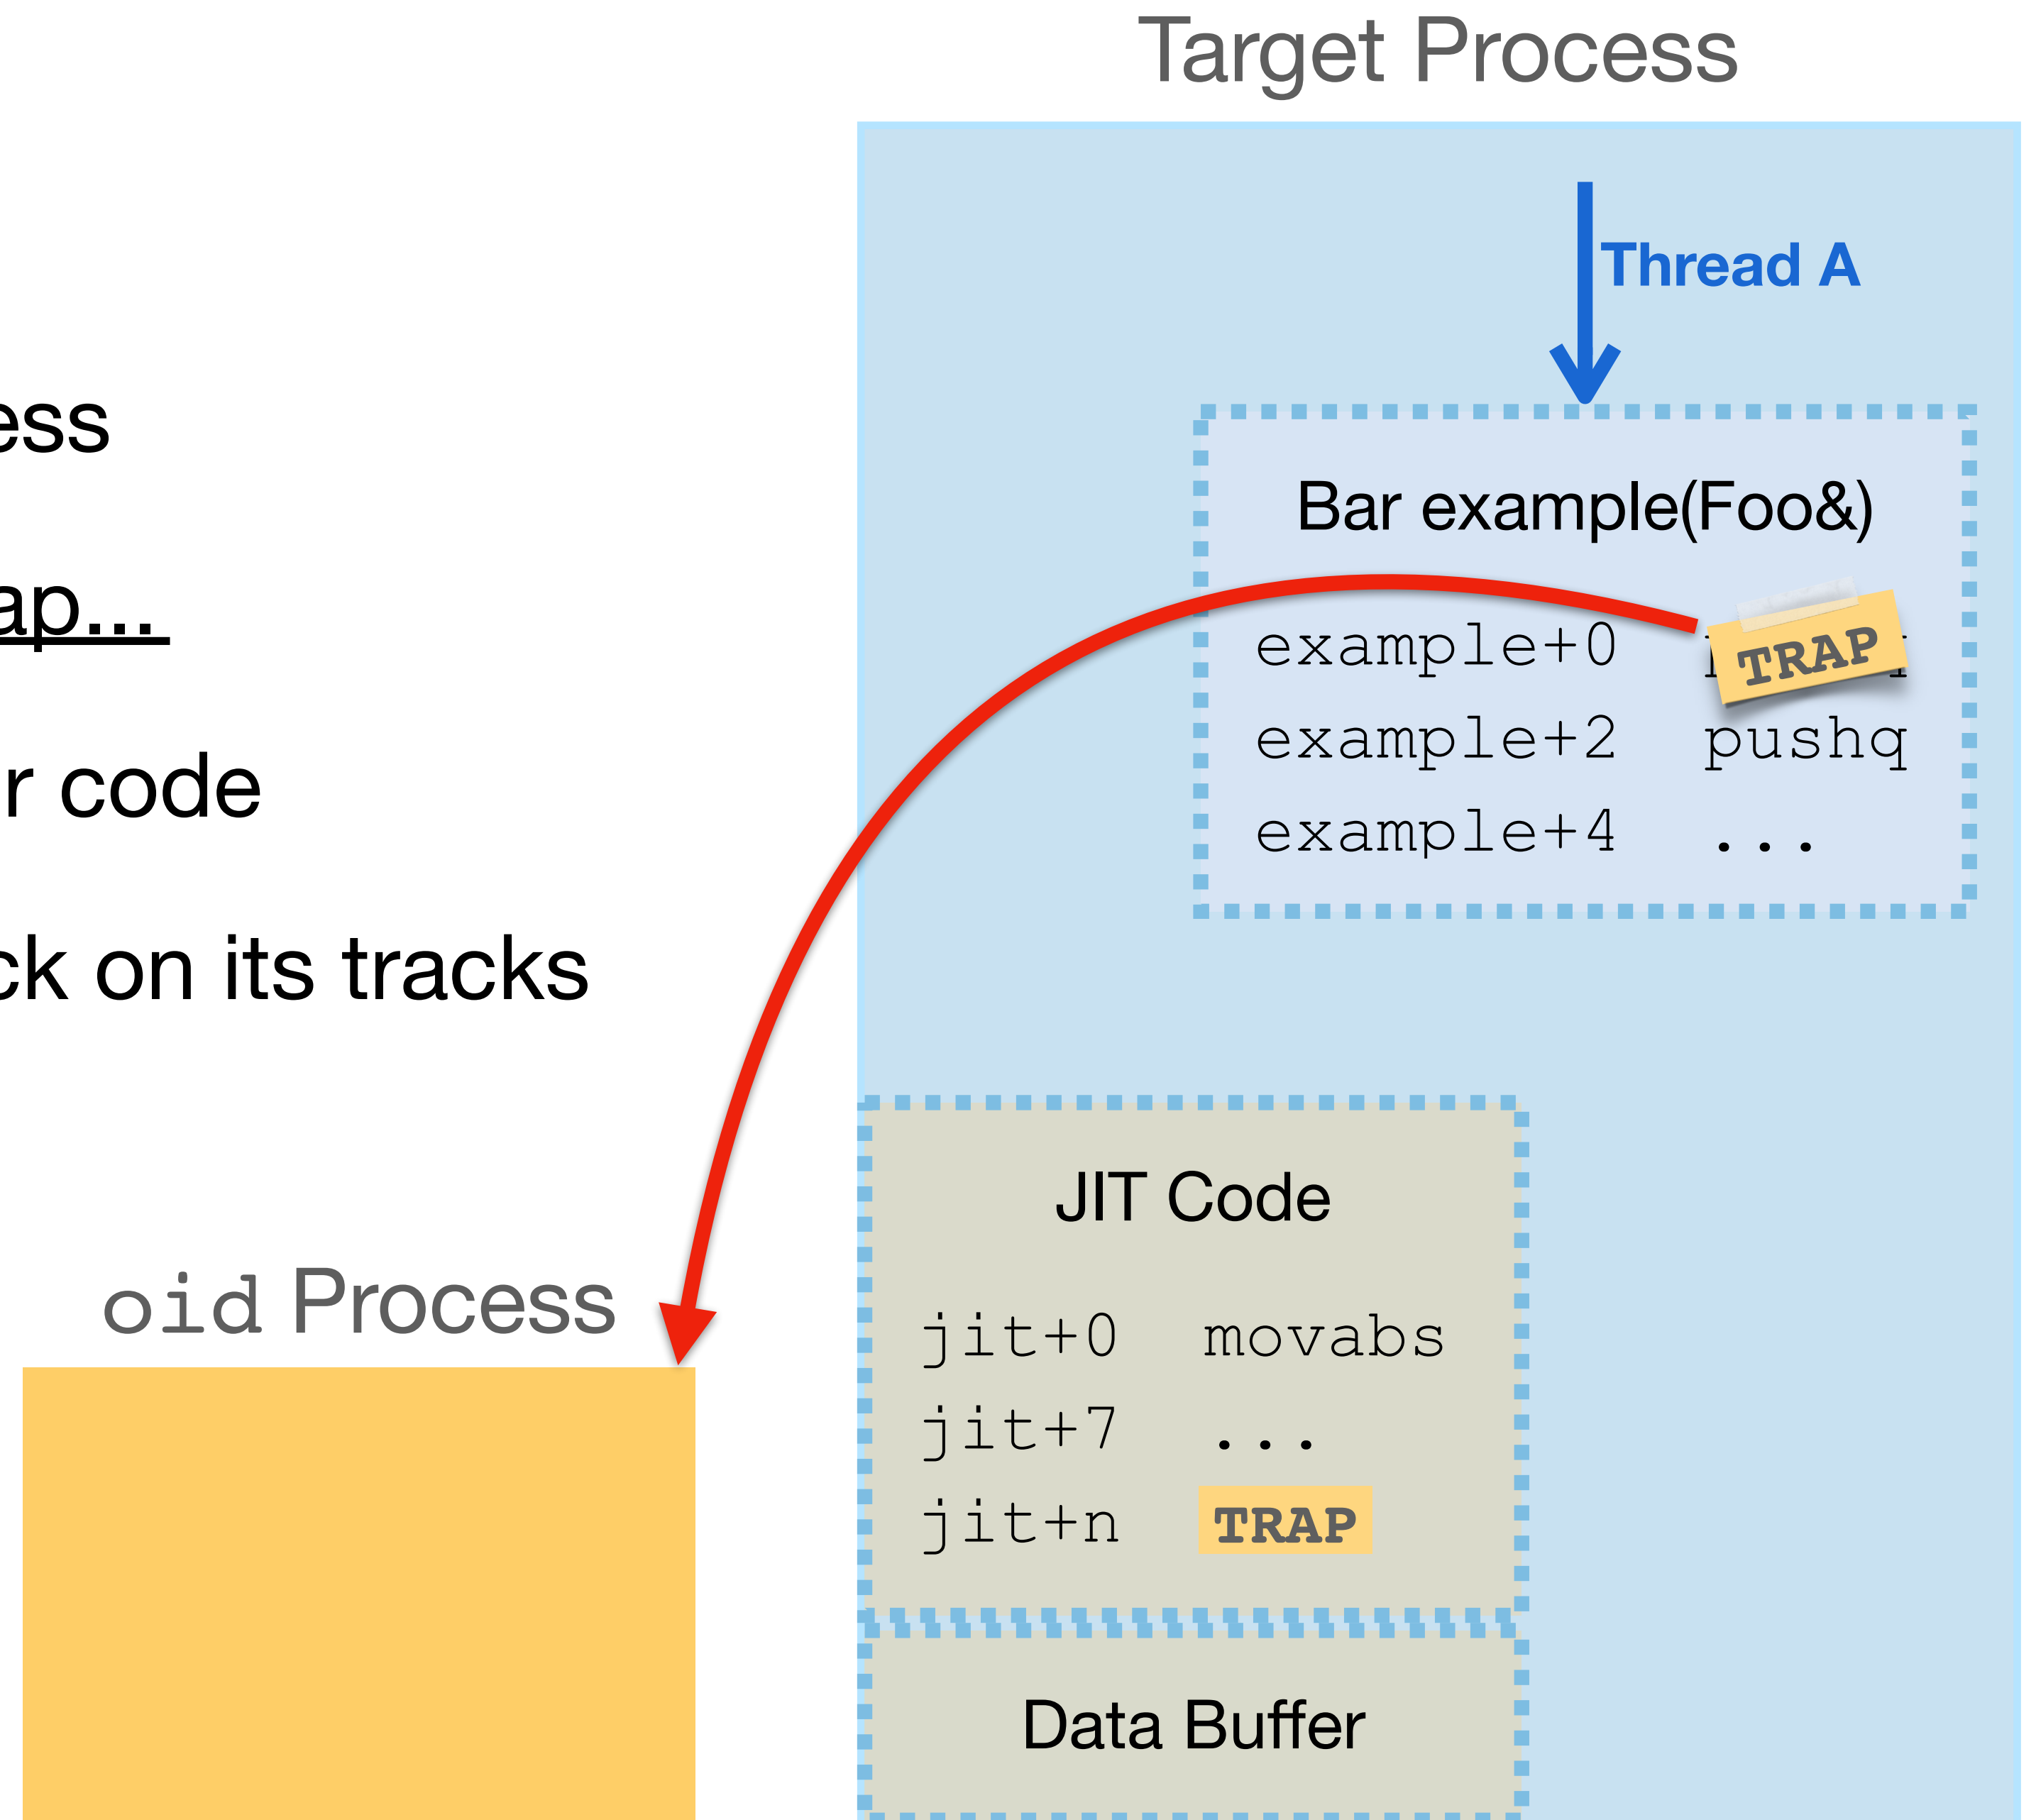


Target Process



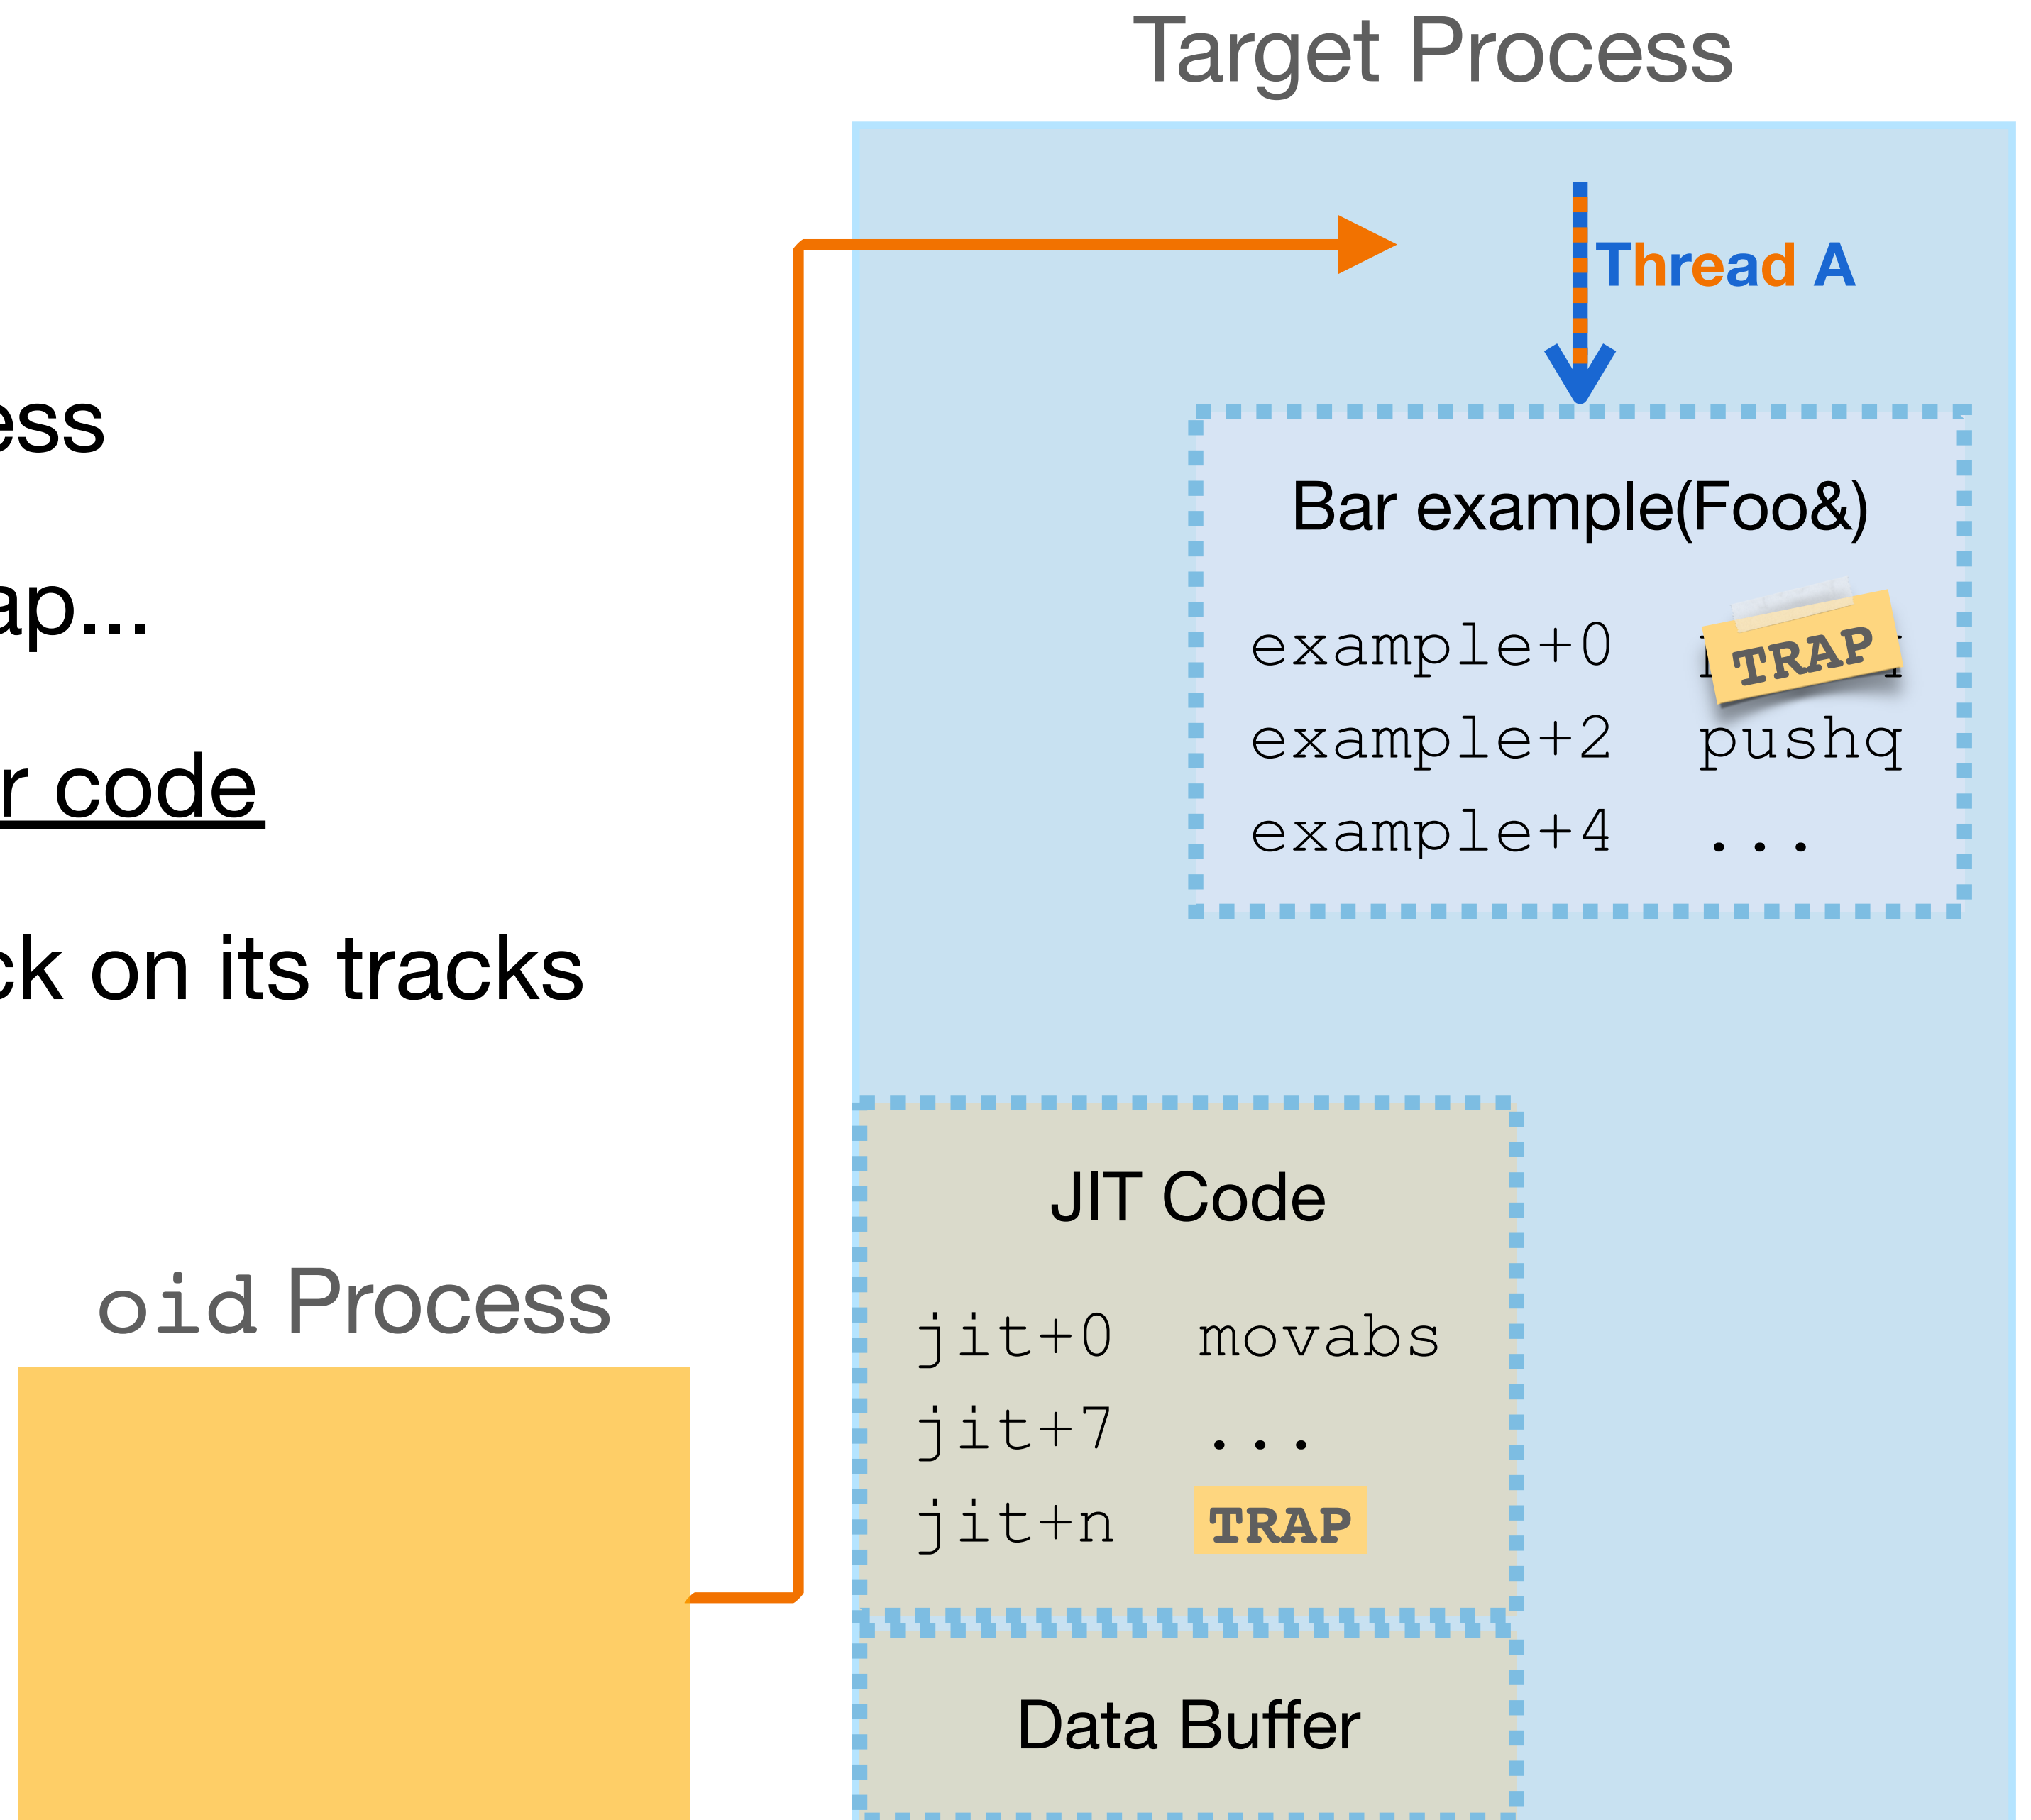
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data



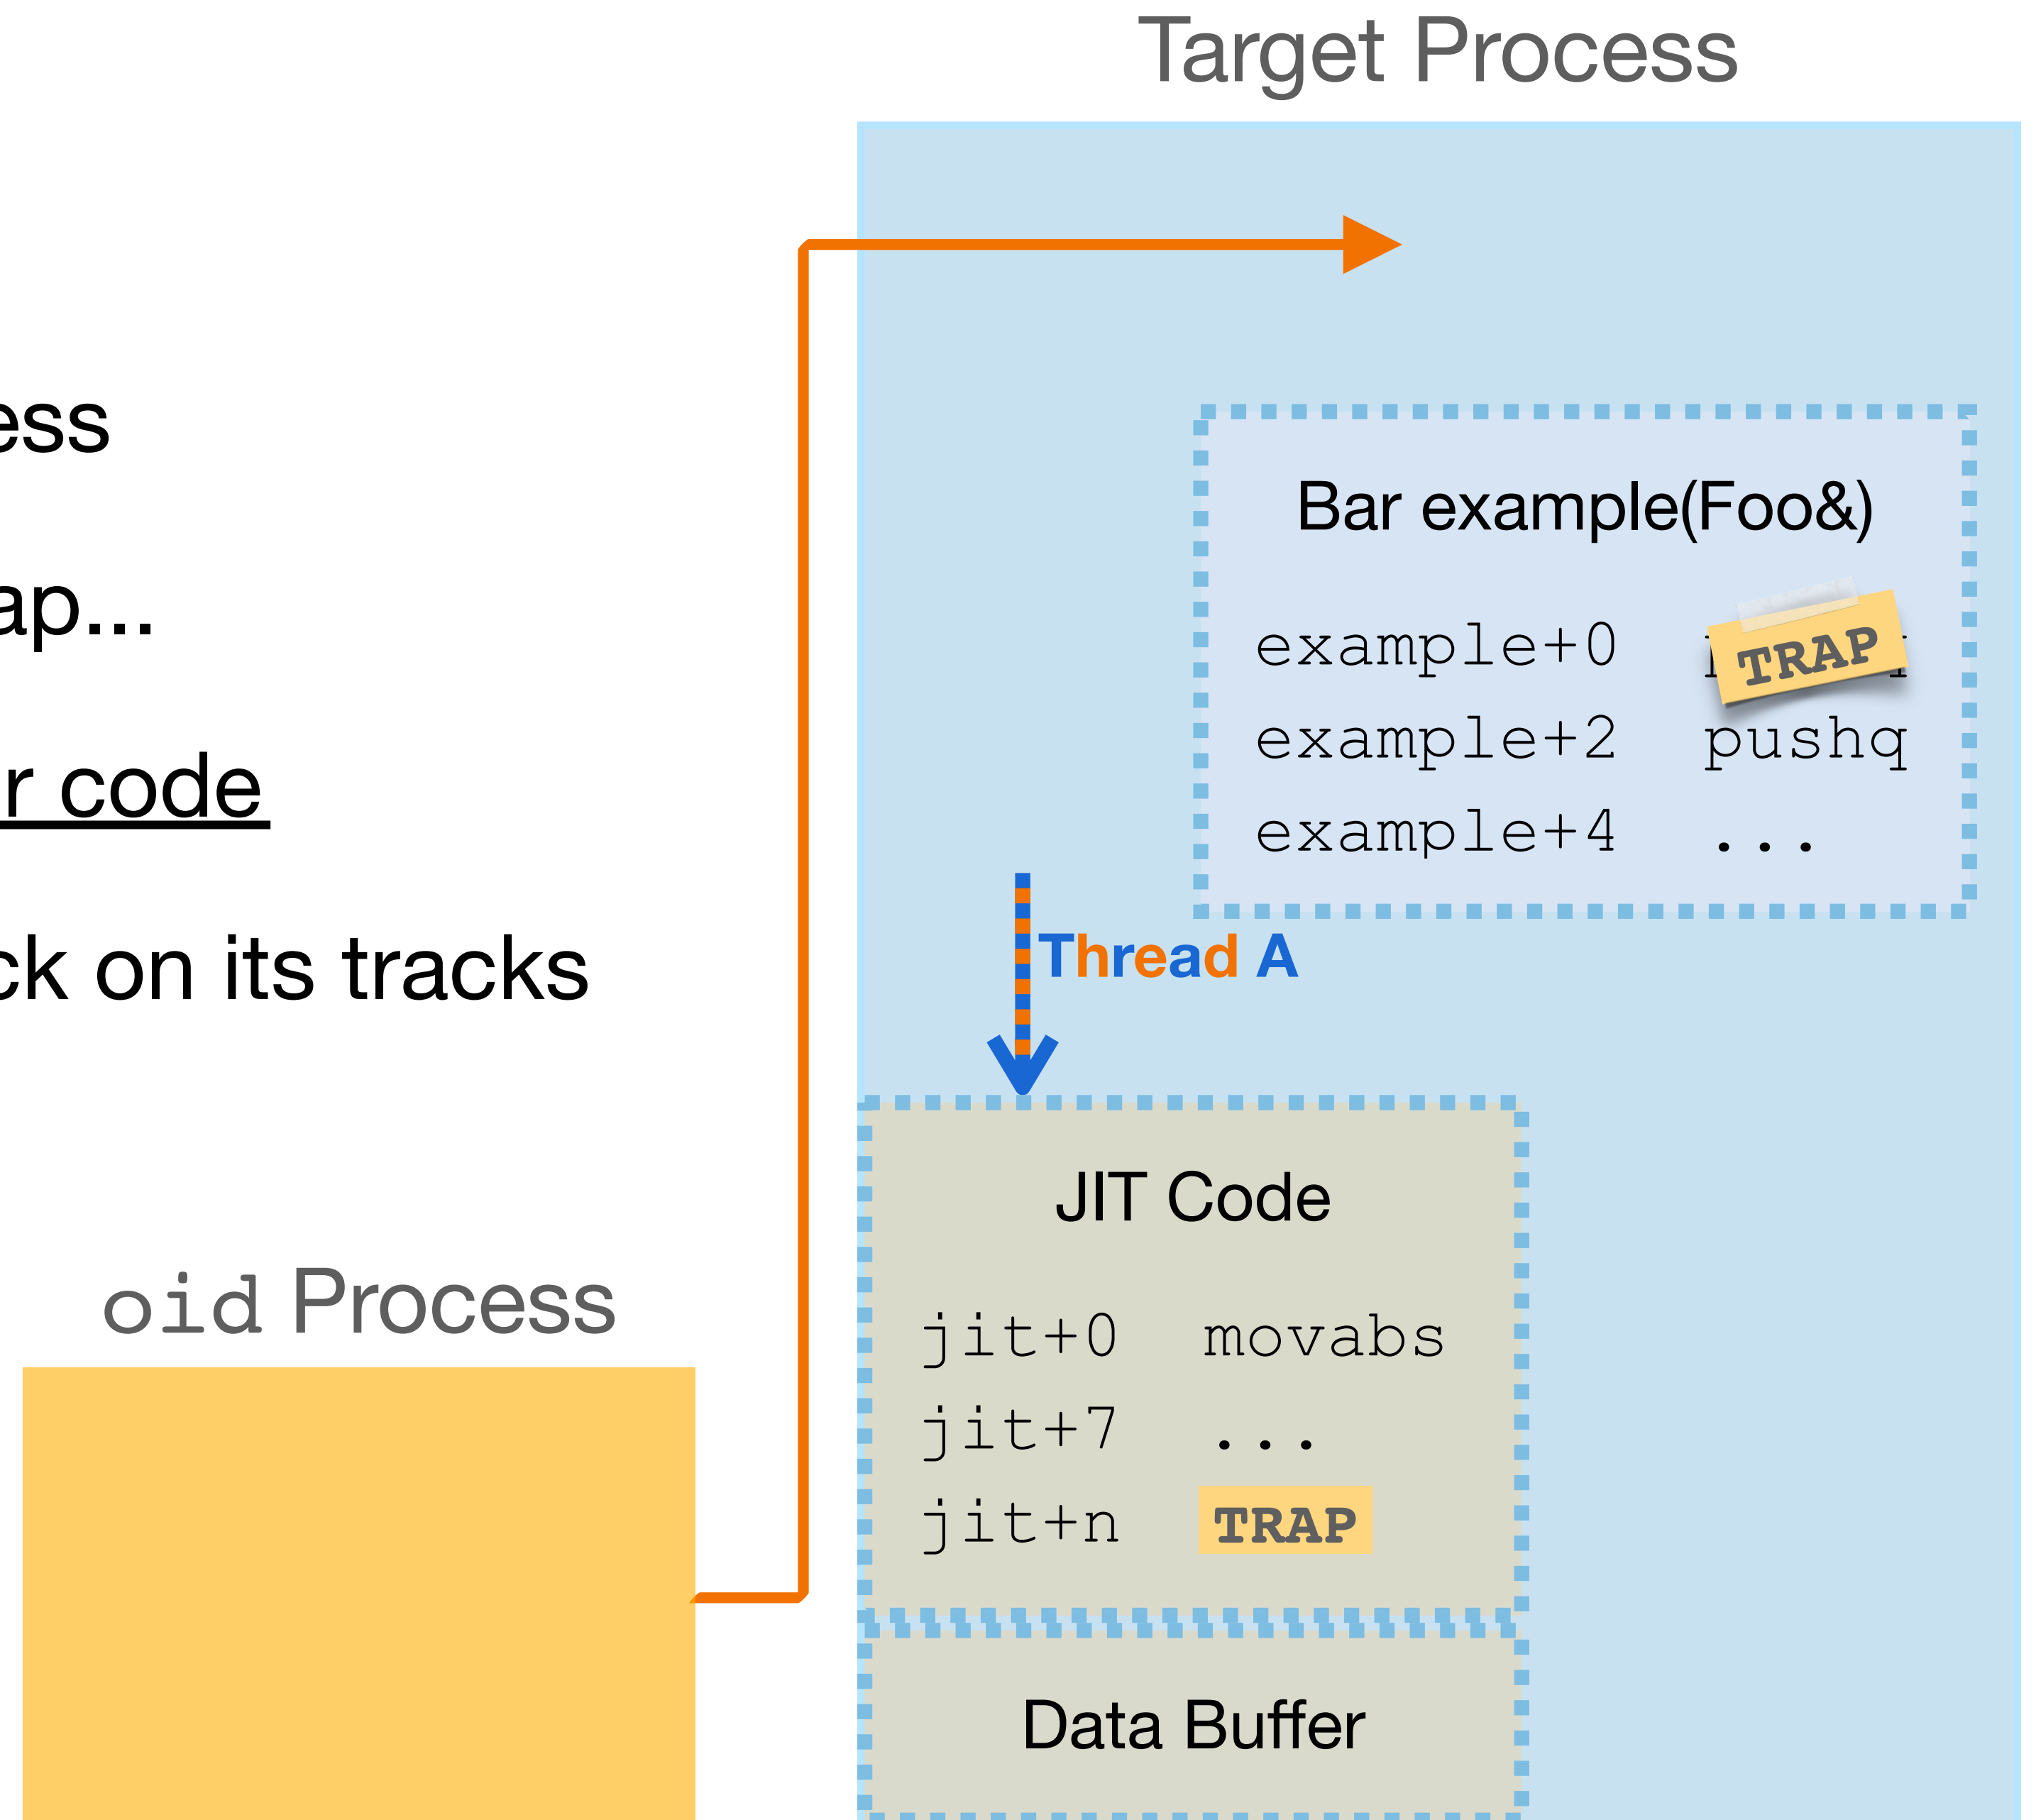
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data



Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data



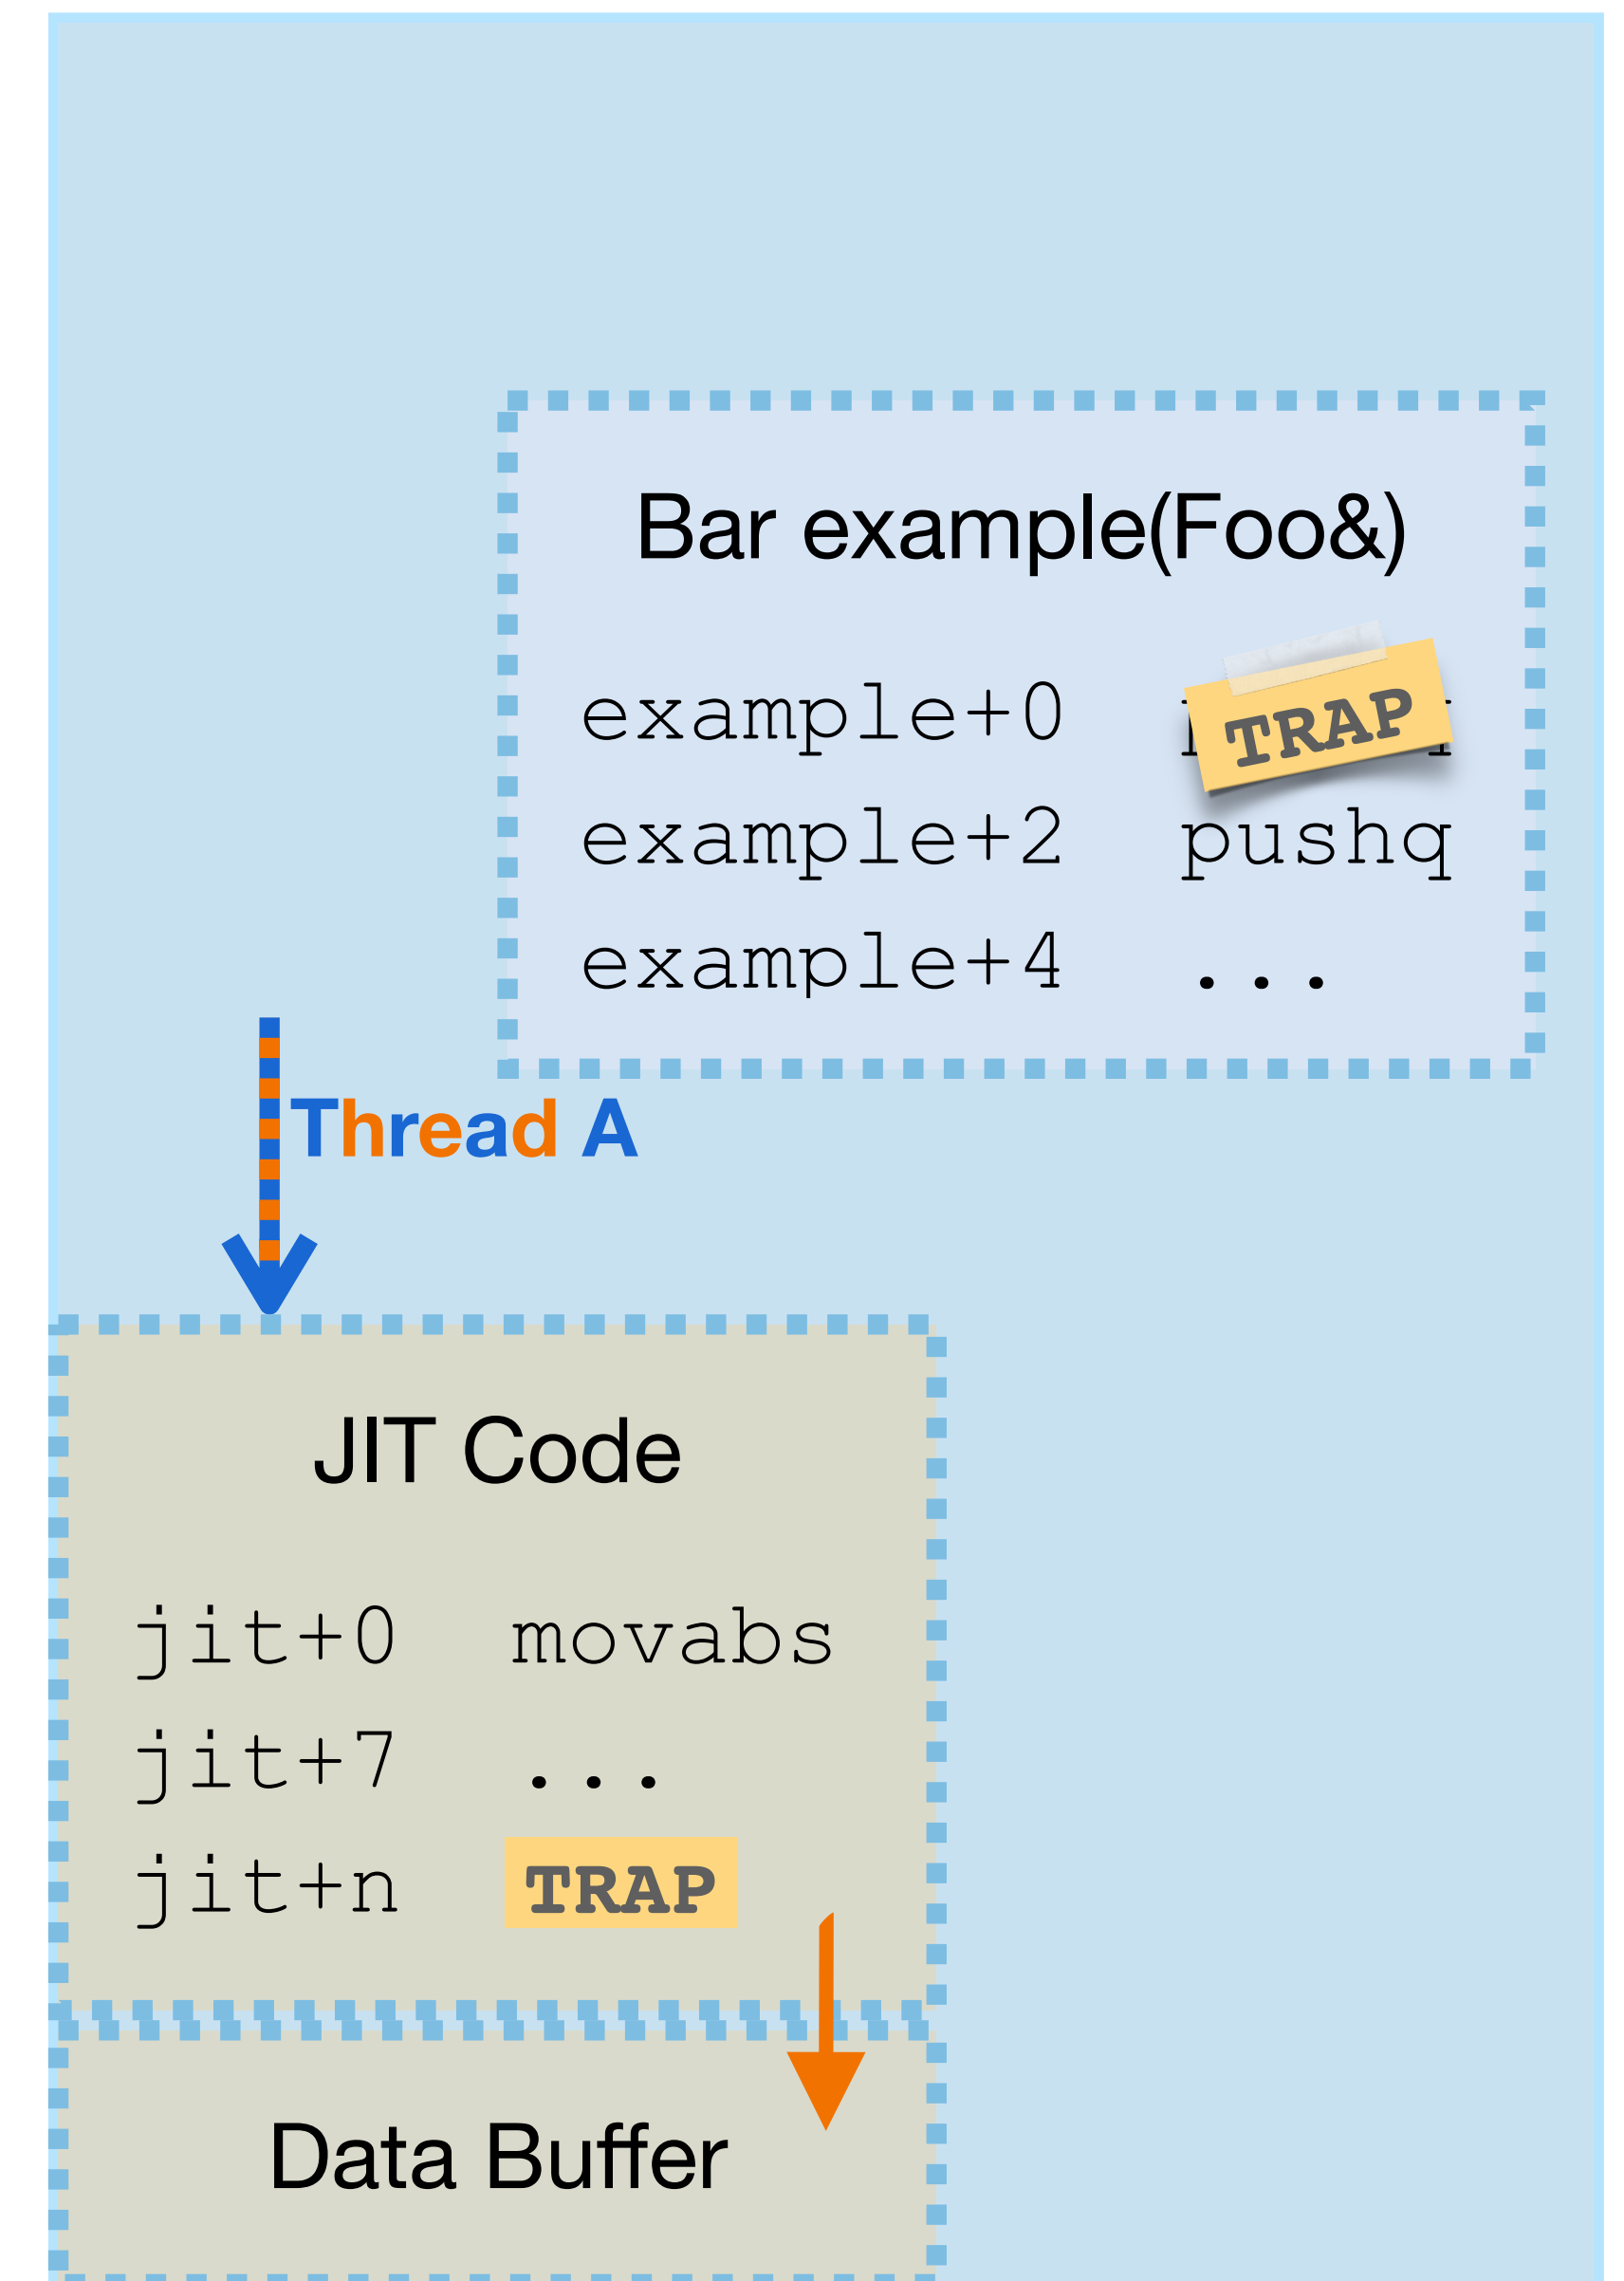
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data

oid Process

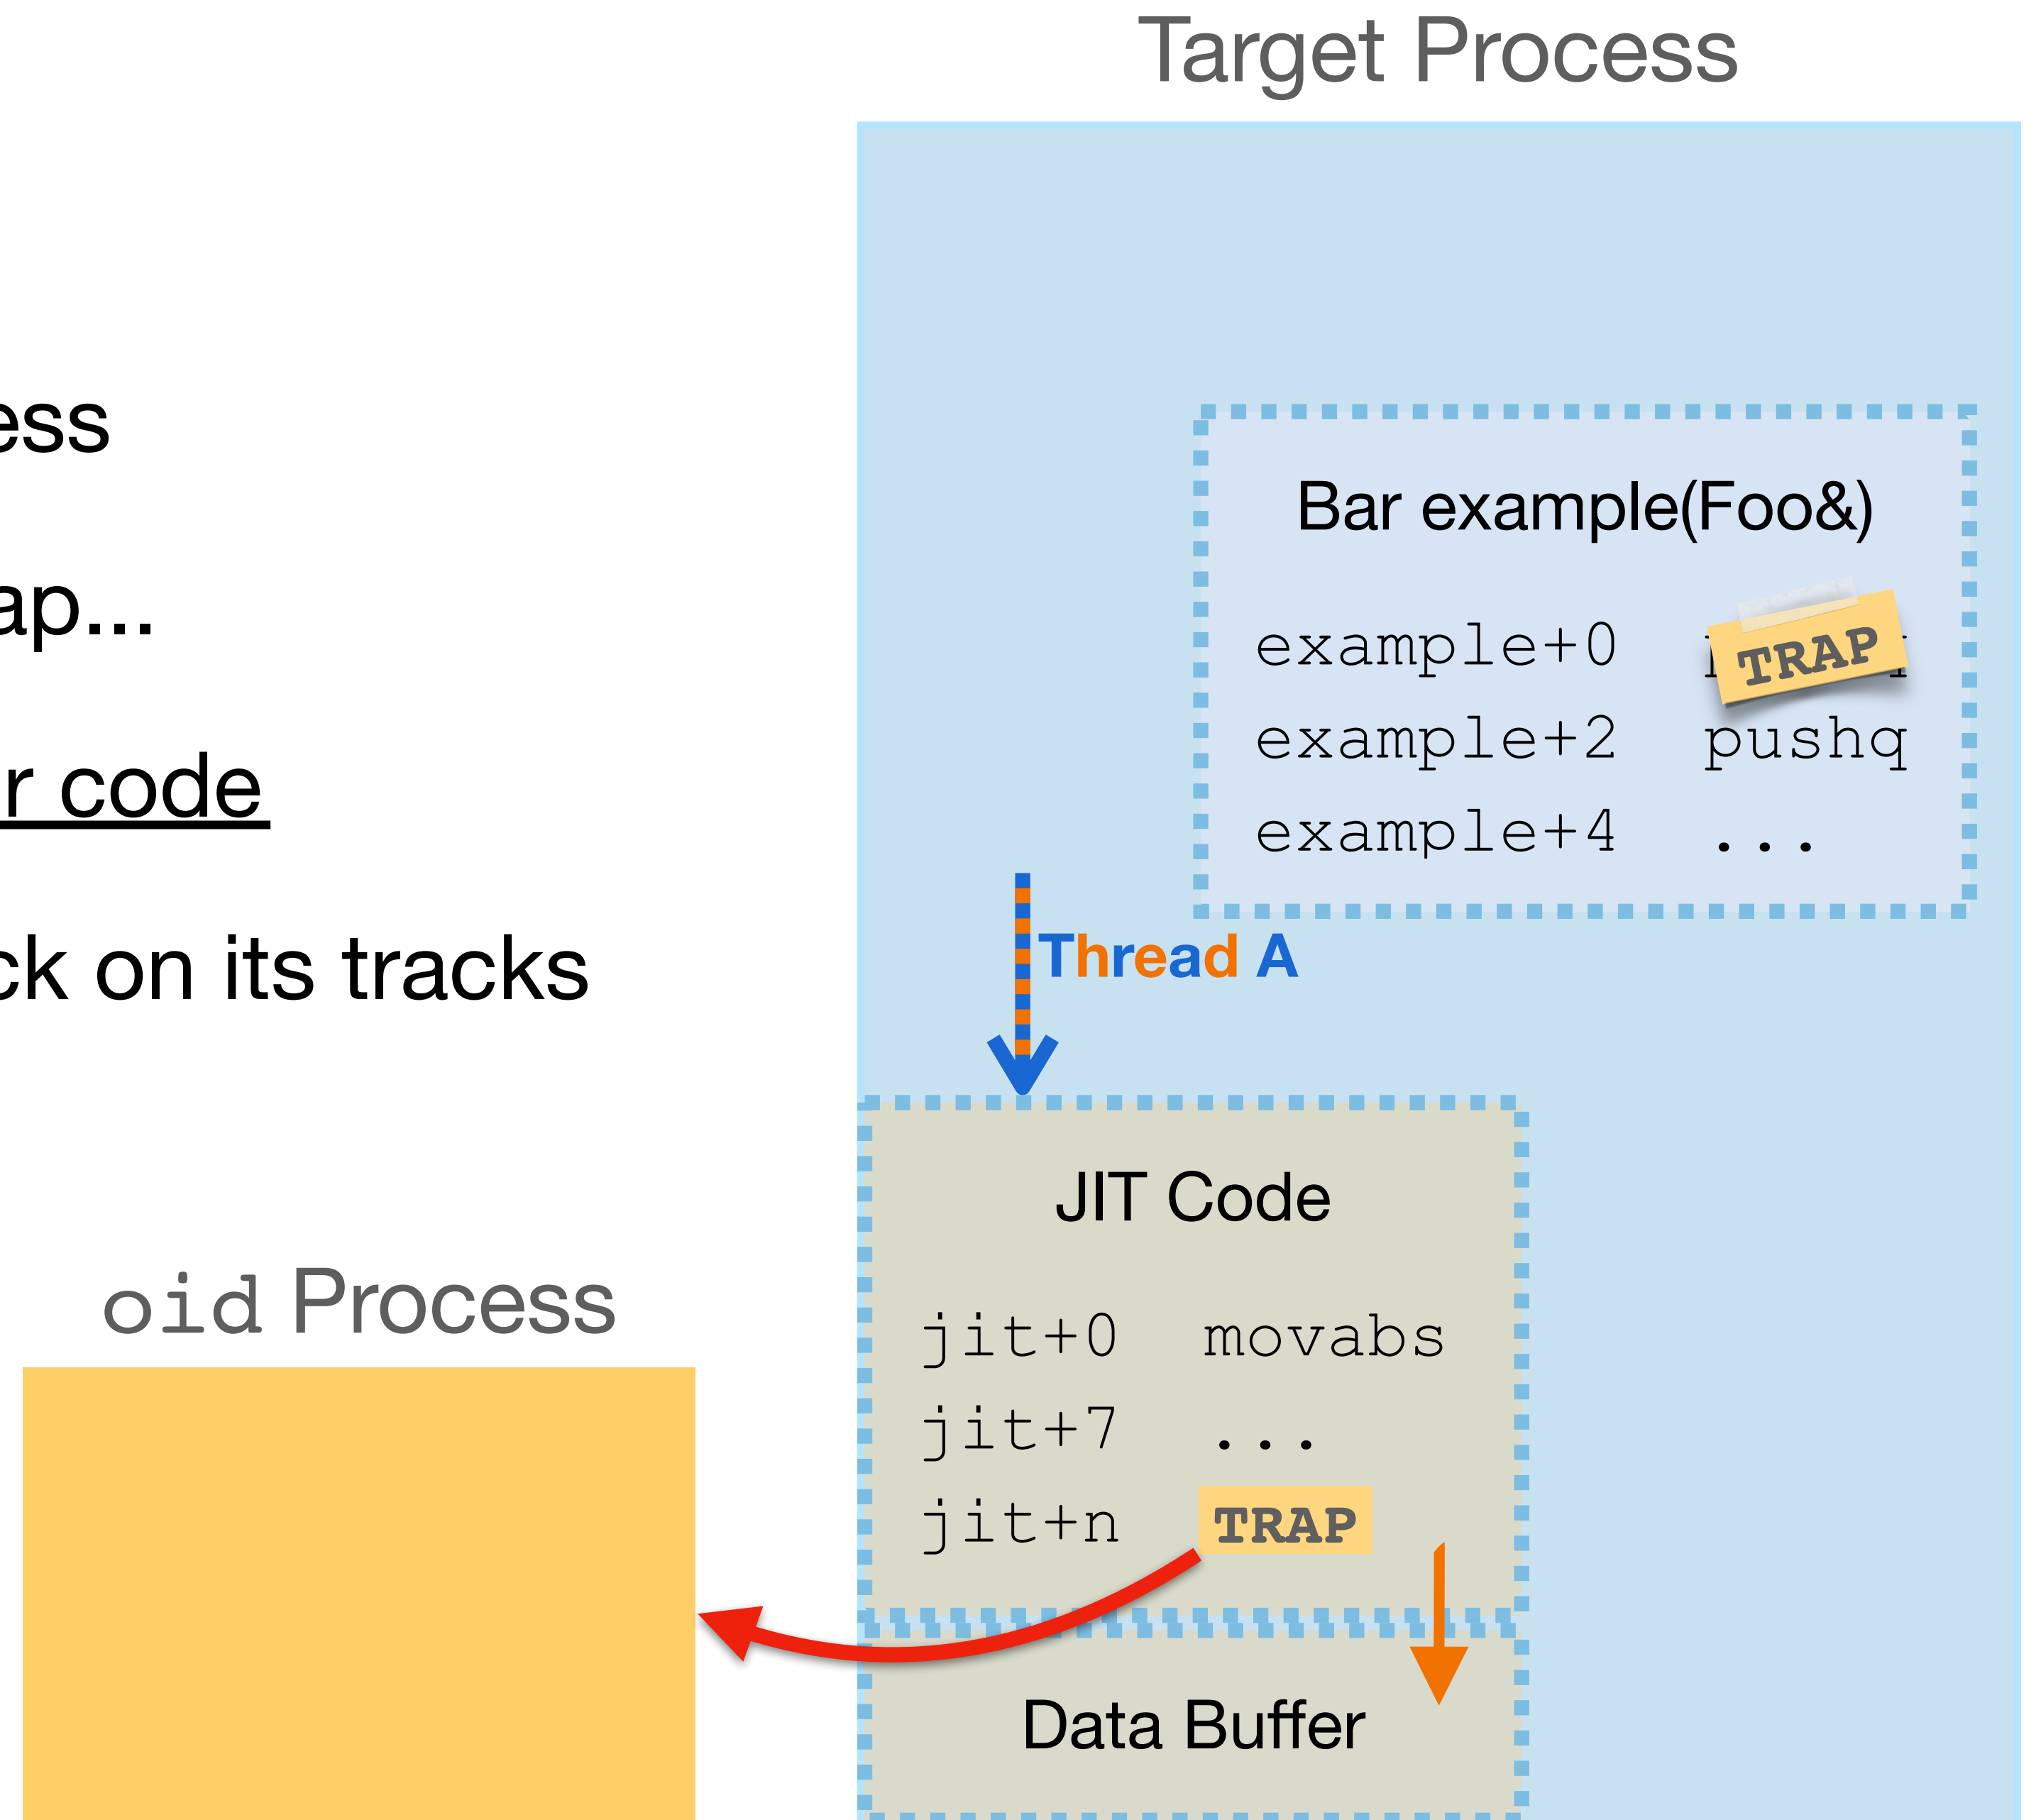


Target Process



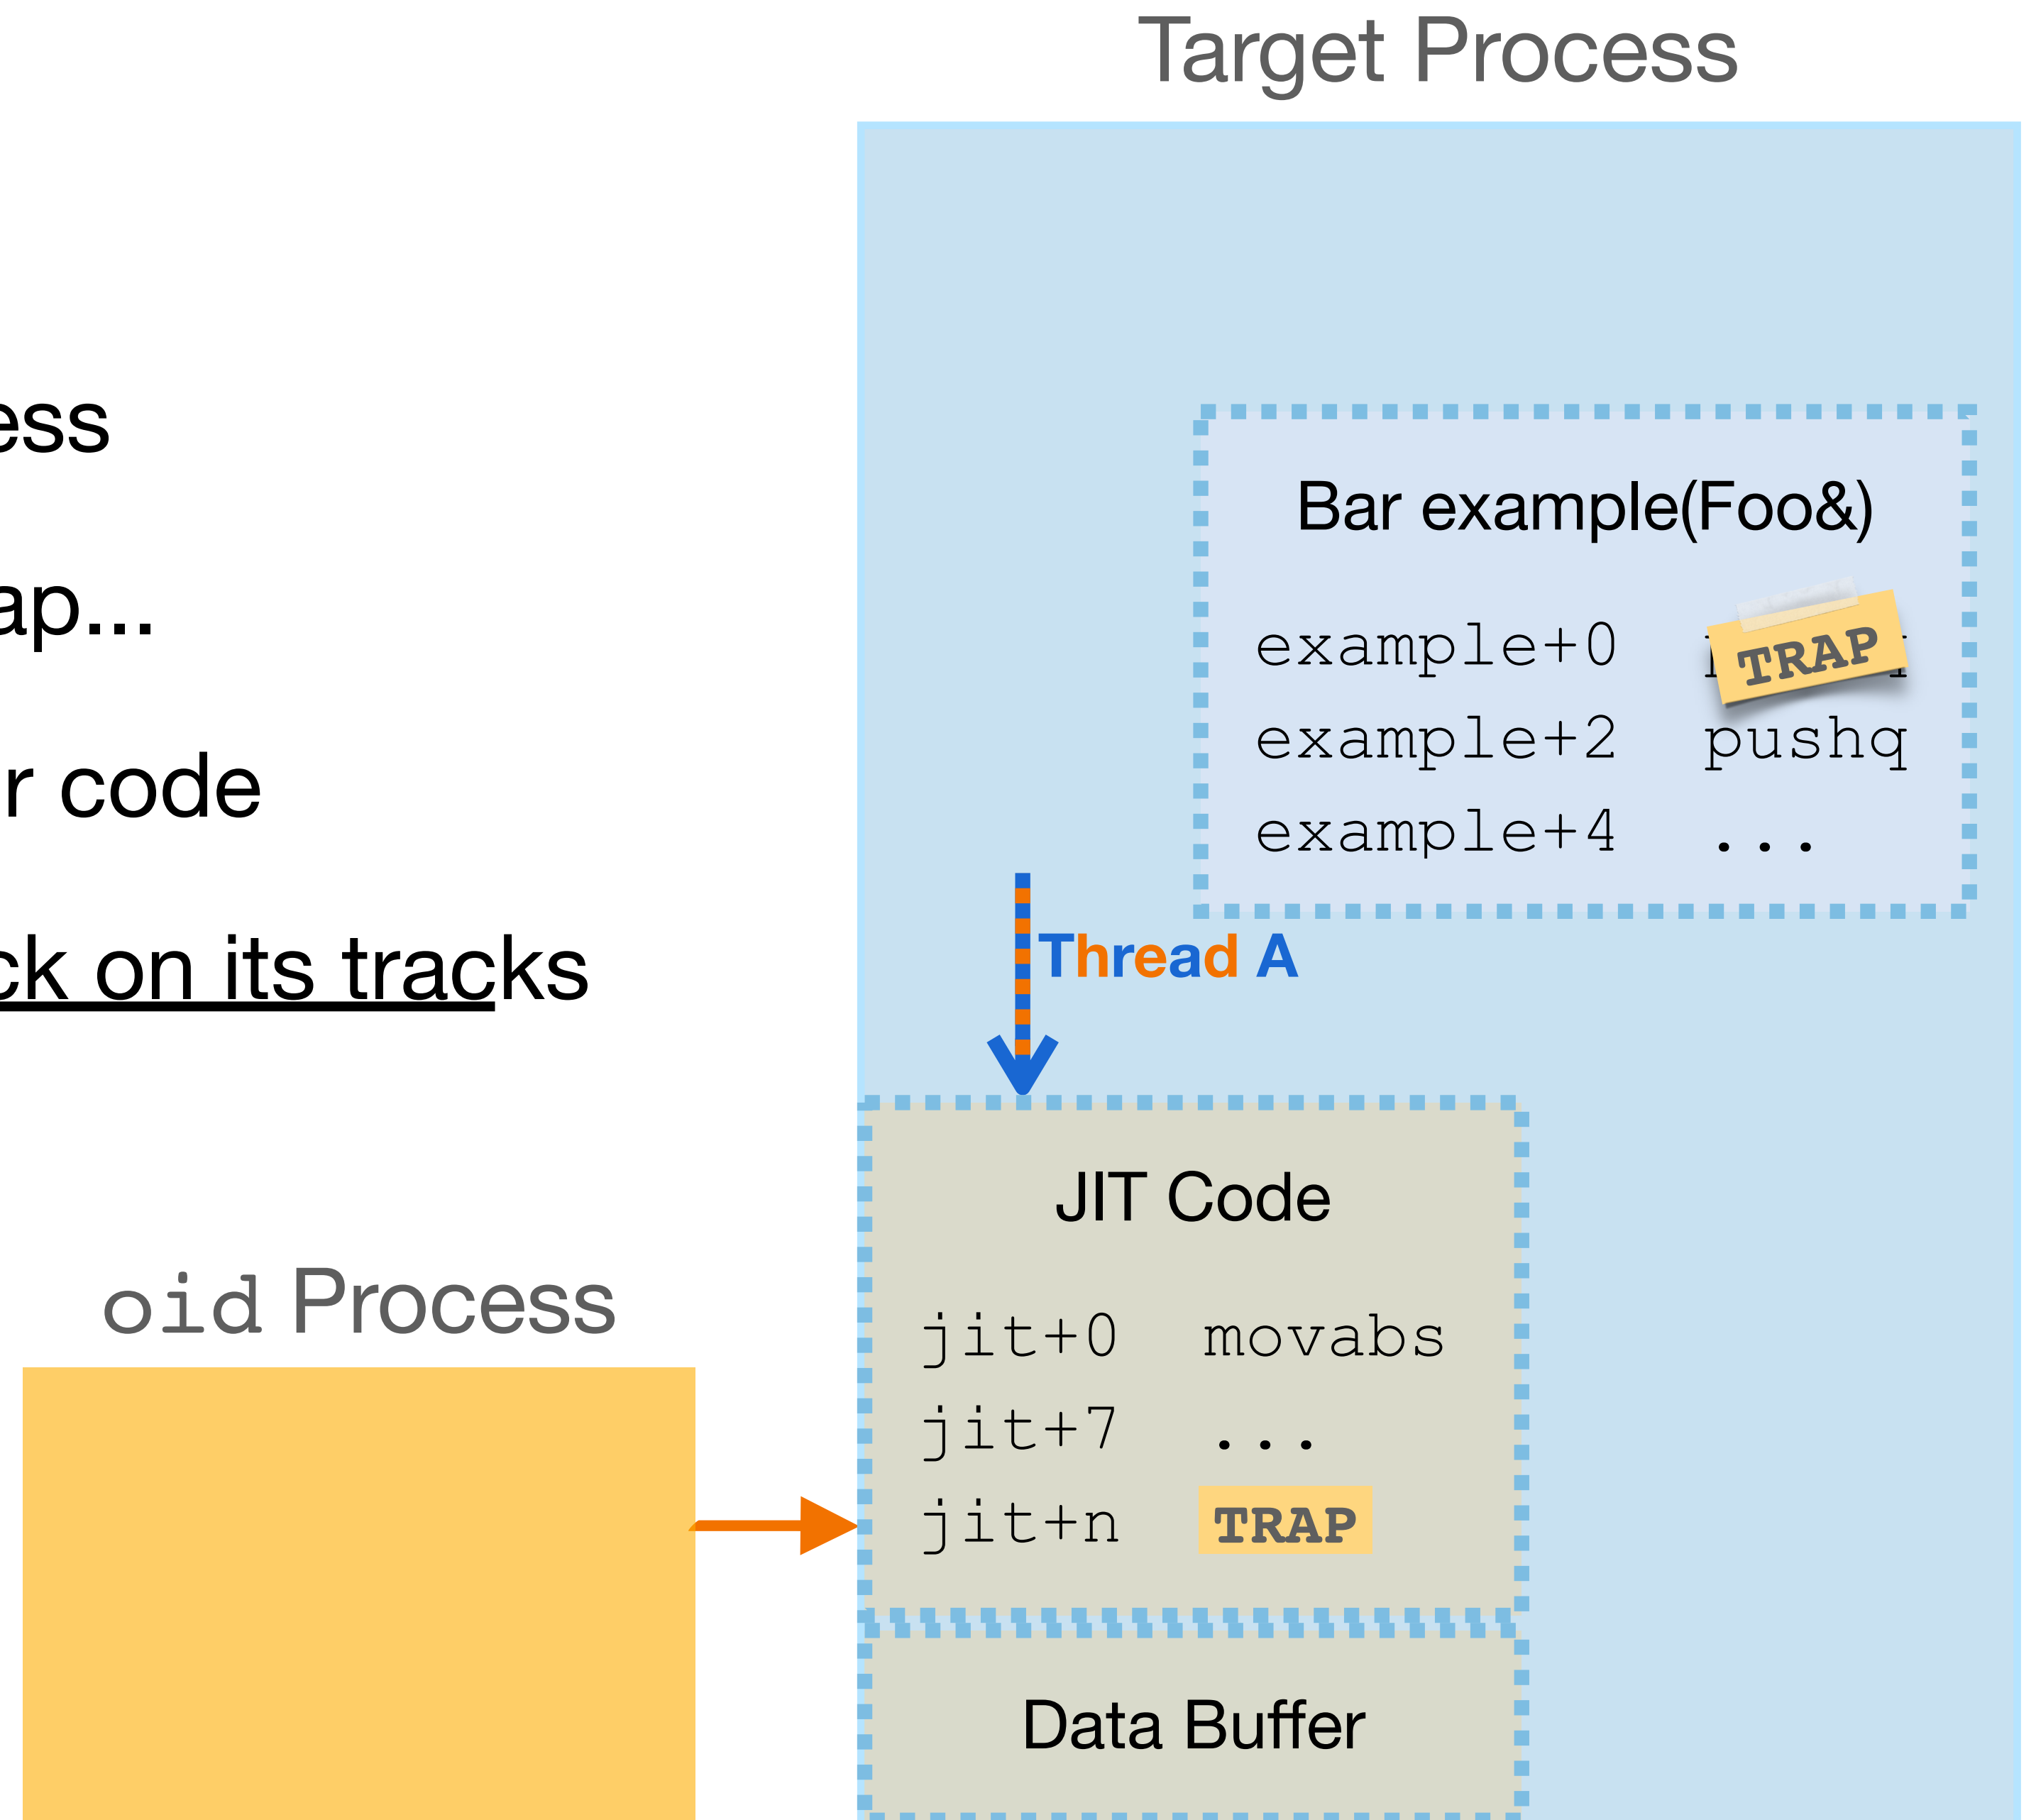
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data



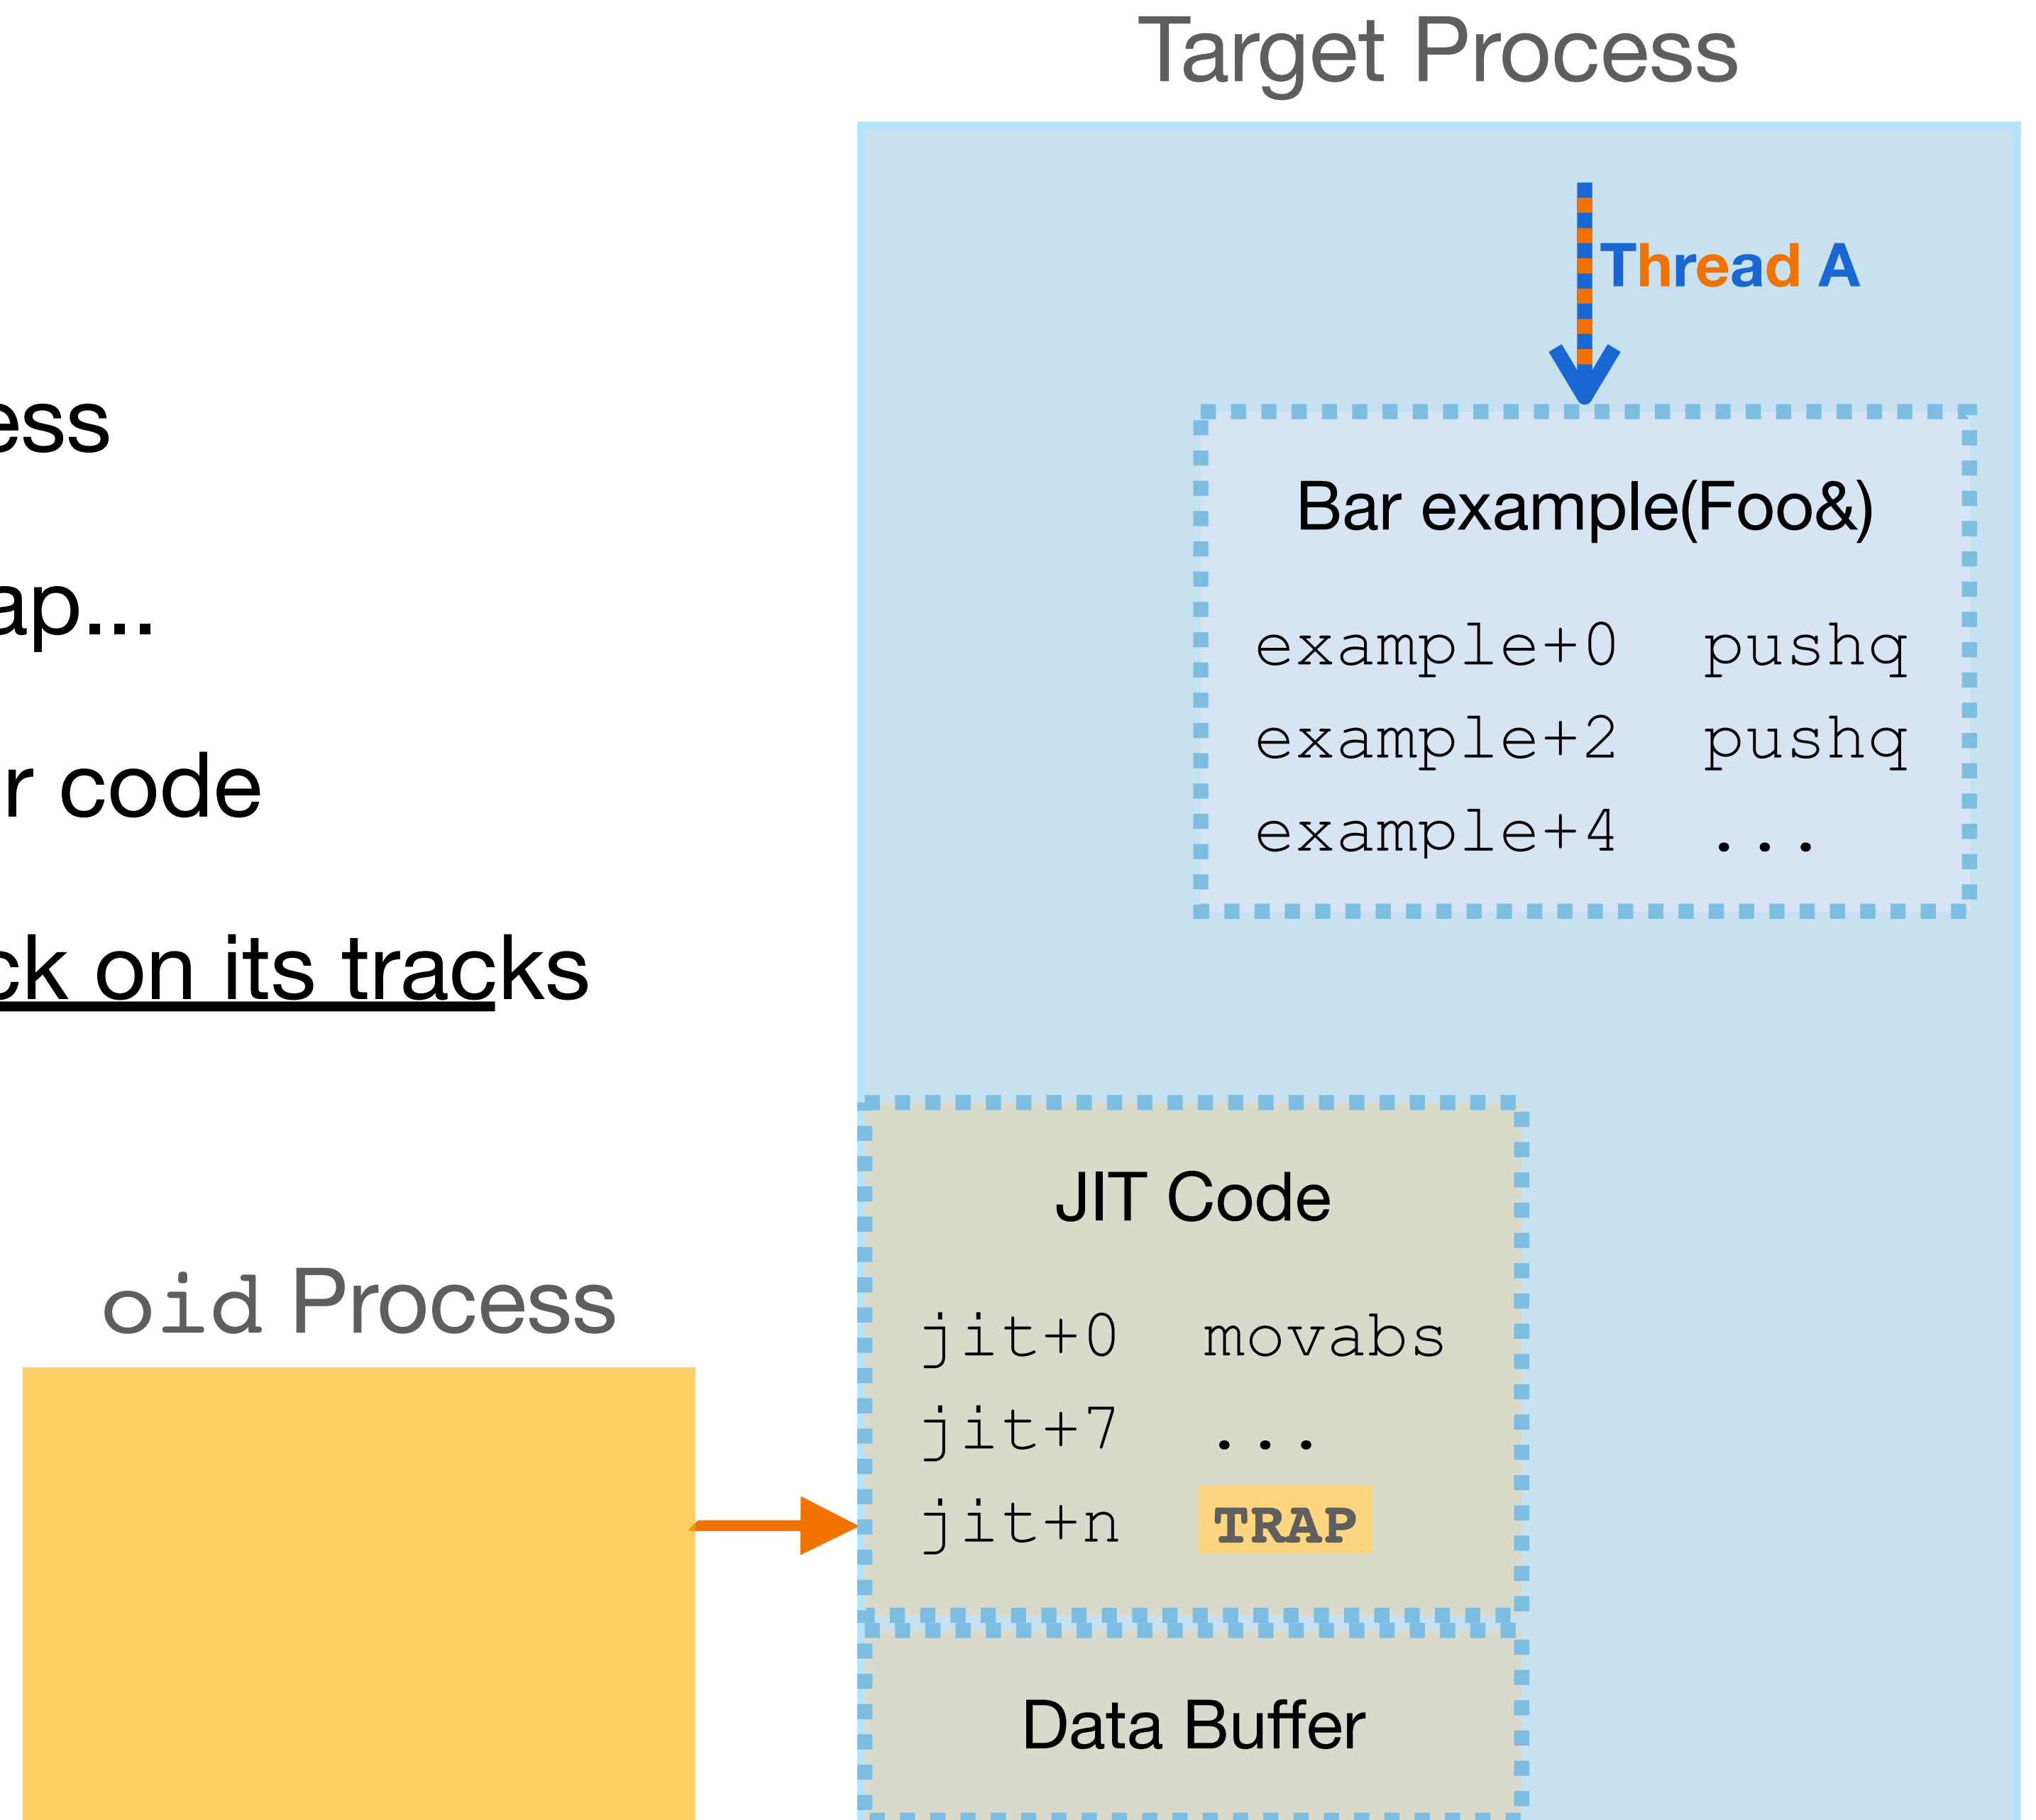
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data



Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data



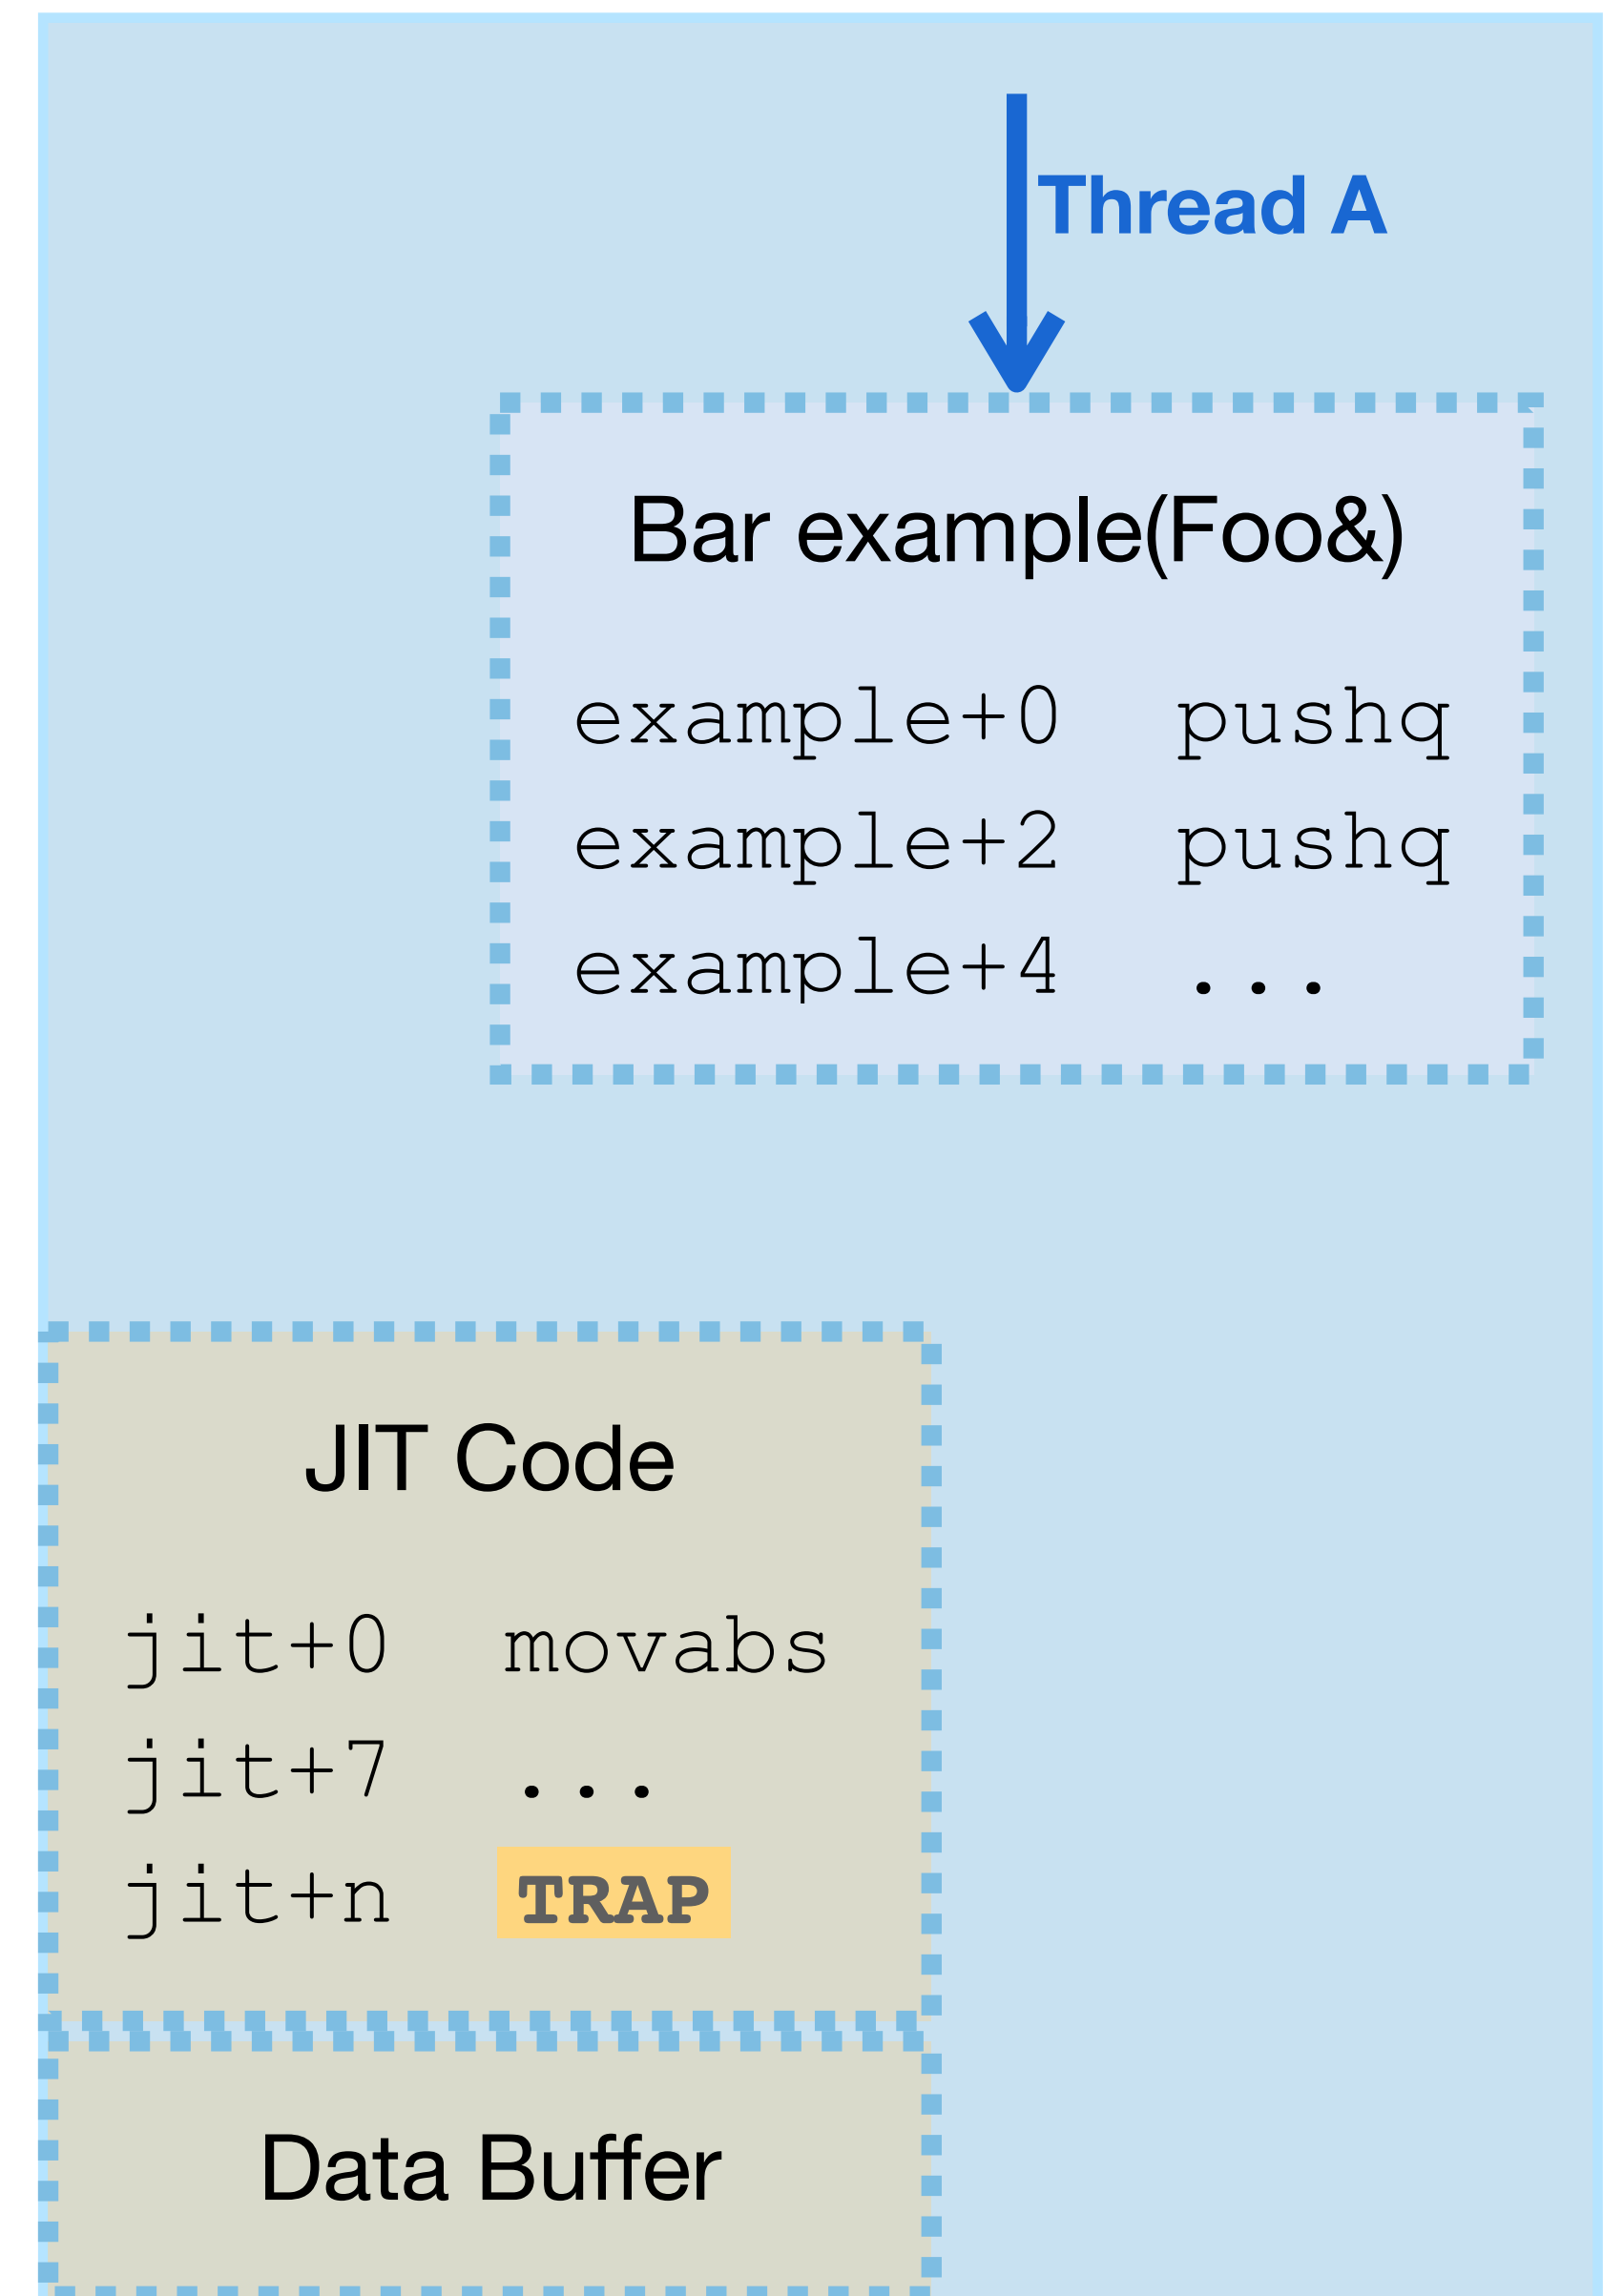
Ol as a Profiler

1. Inject Ol into the target process
2. Wait for a thread to hit our trap...
3. Redirect the thread to run our code
4. Put the redirected thread back on its tracks
5. Process the collected data

oid Process

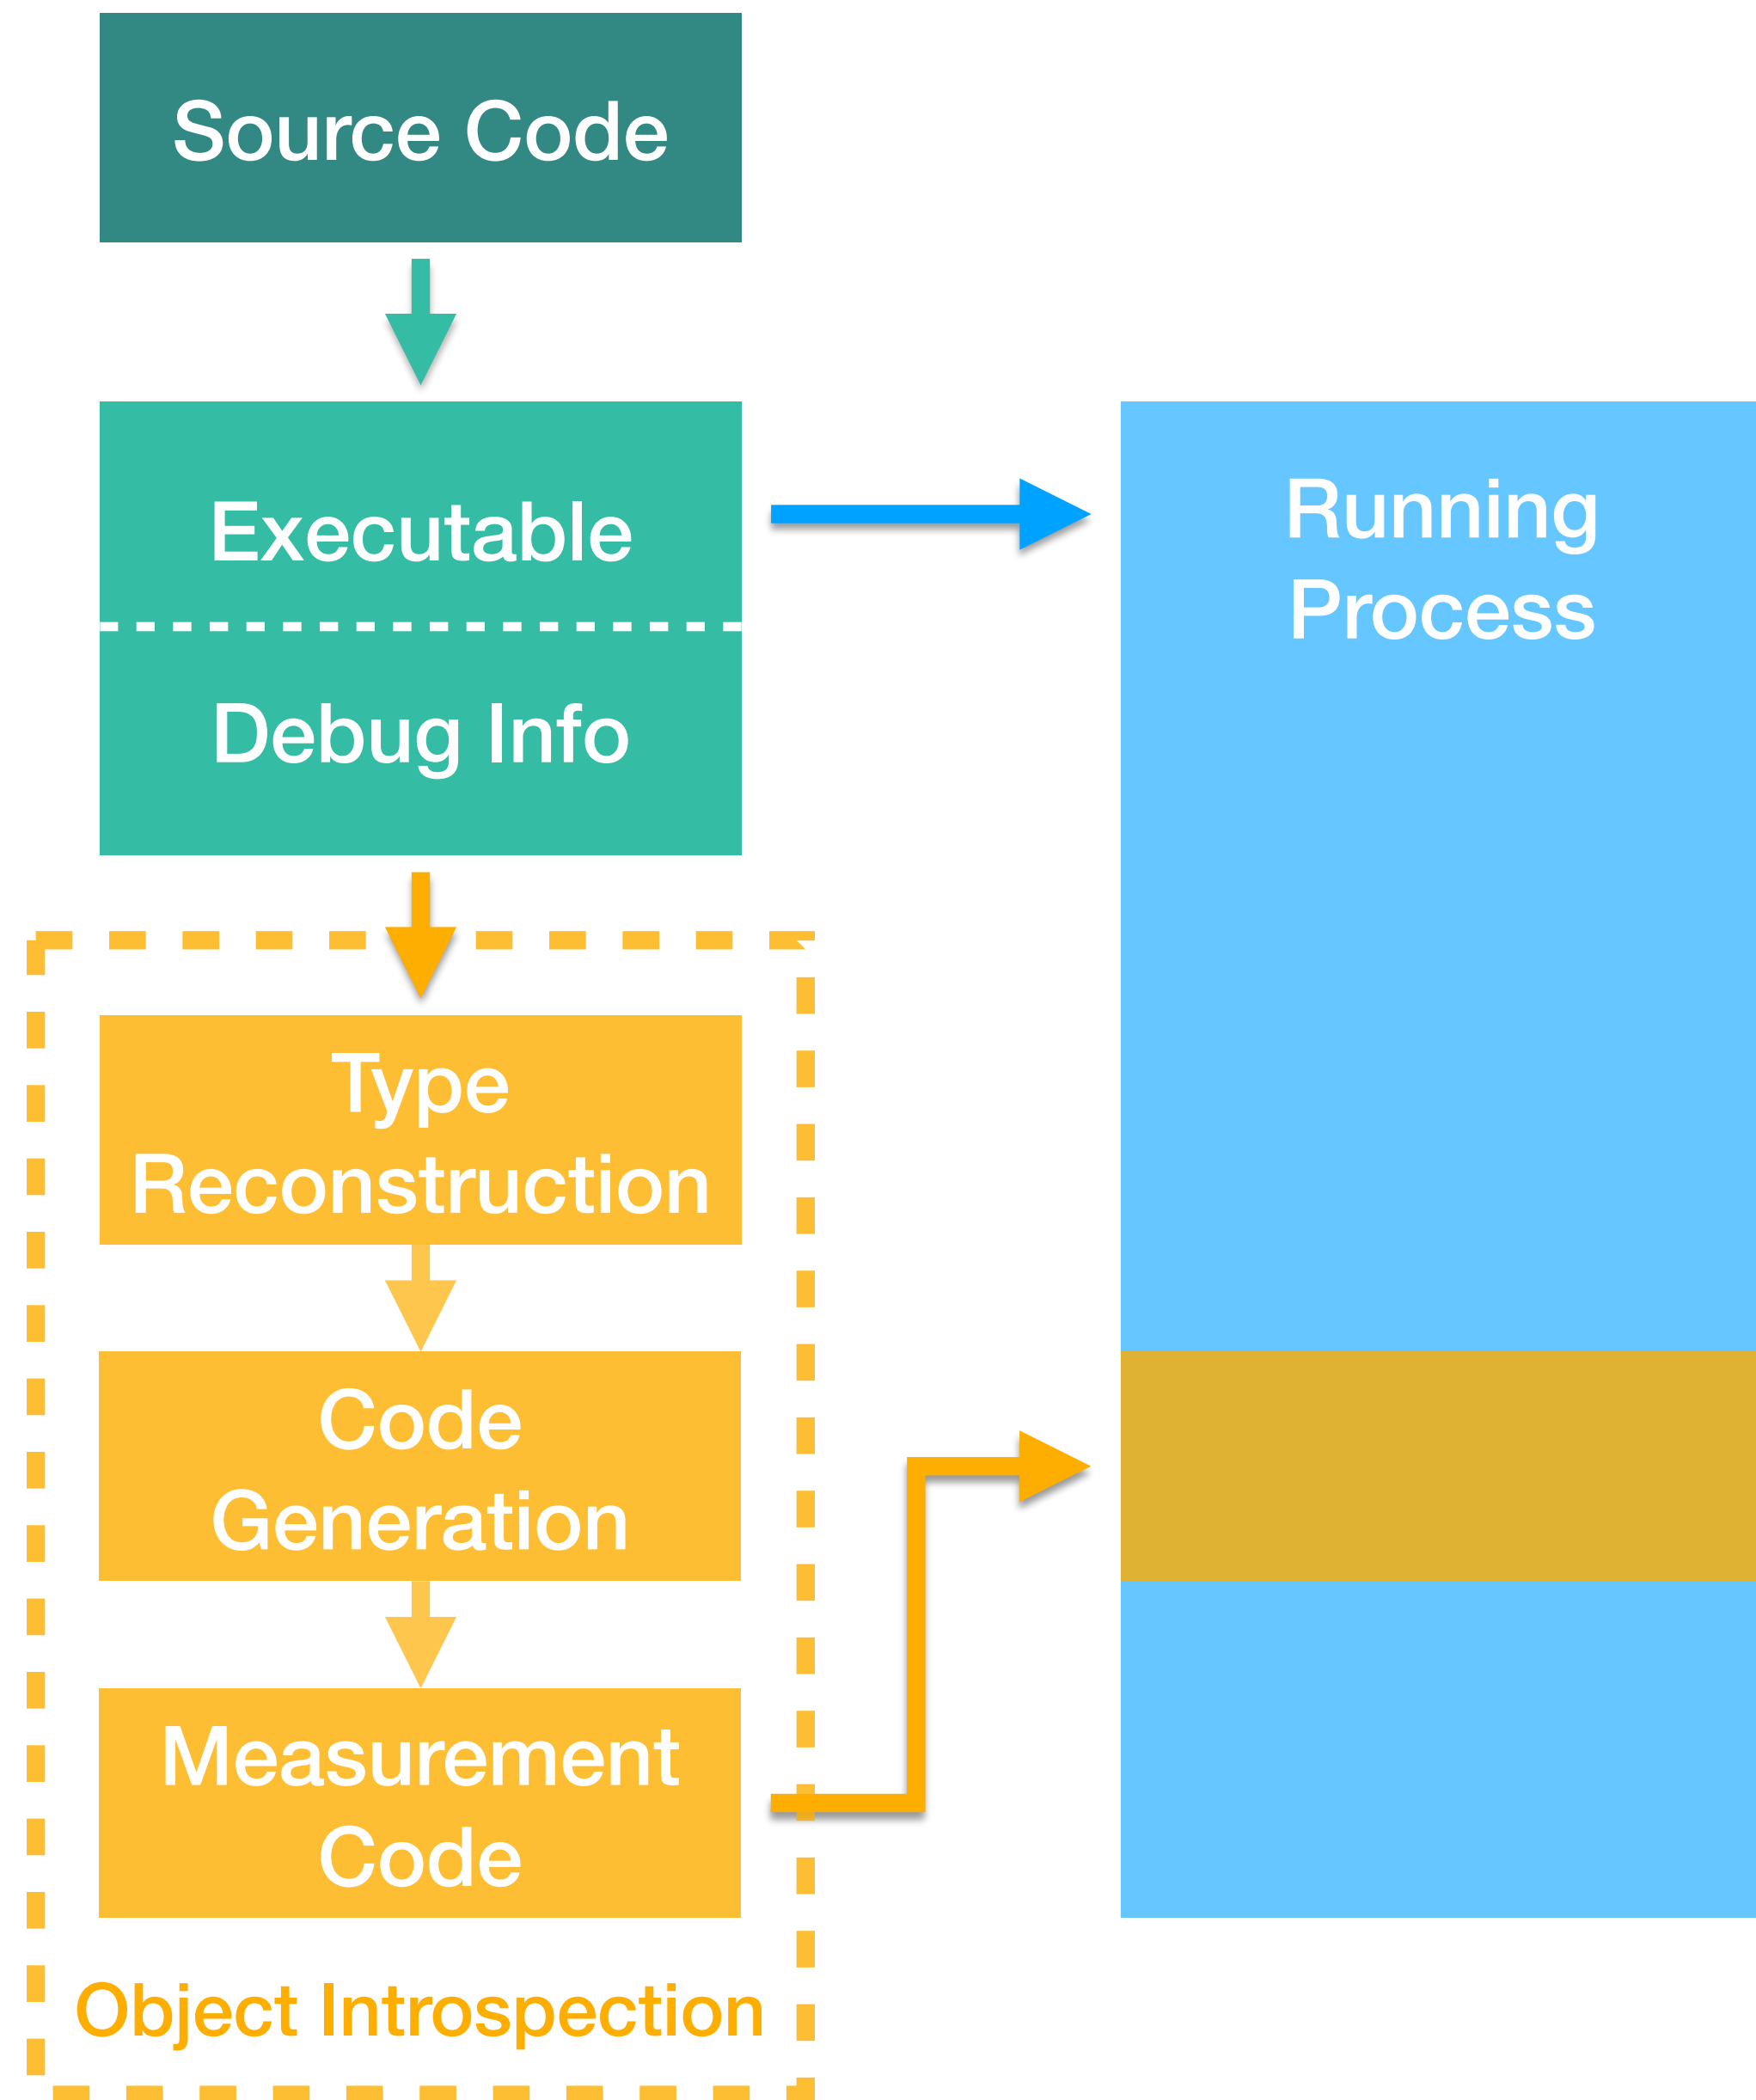


Target Process



Outline

- Debug Info Analysis
 - Type/Layout Reconstruction
 - Code Generation
- Object Introspection as a Library
- Object Introspection as a Profiler
- **Object Introspection Applied**



Applied Example 1

- Unused container memory:


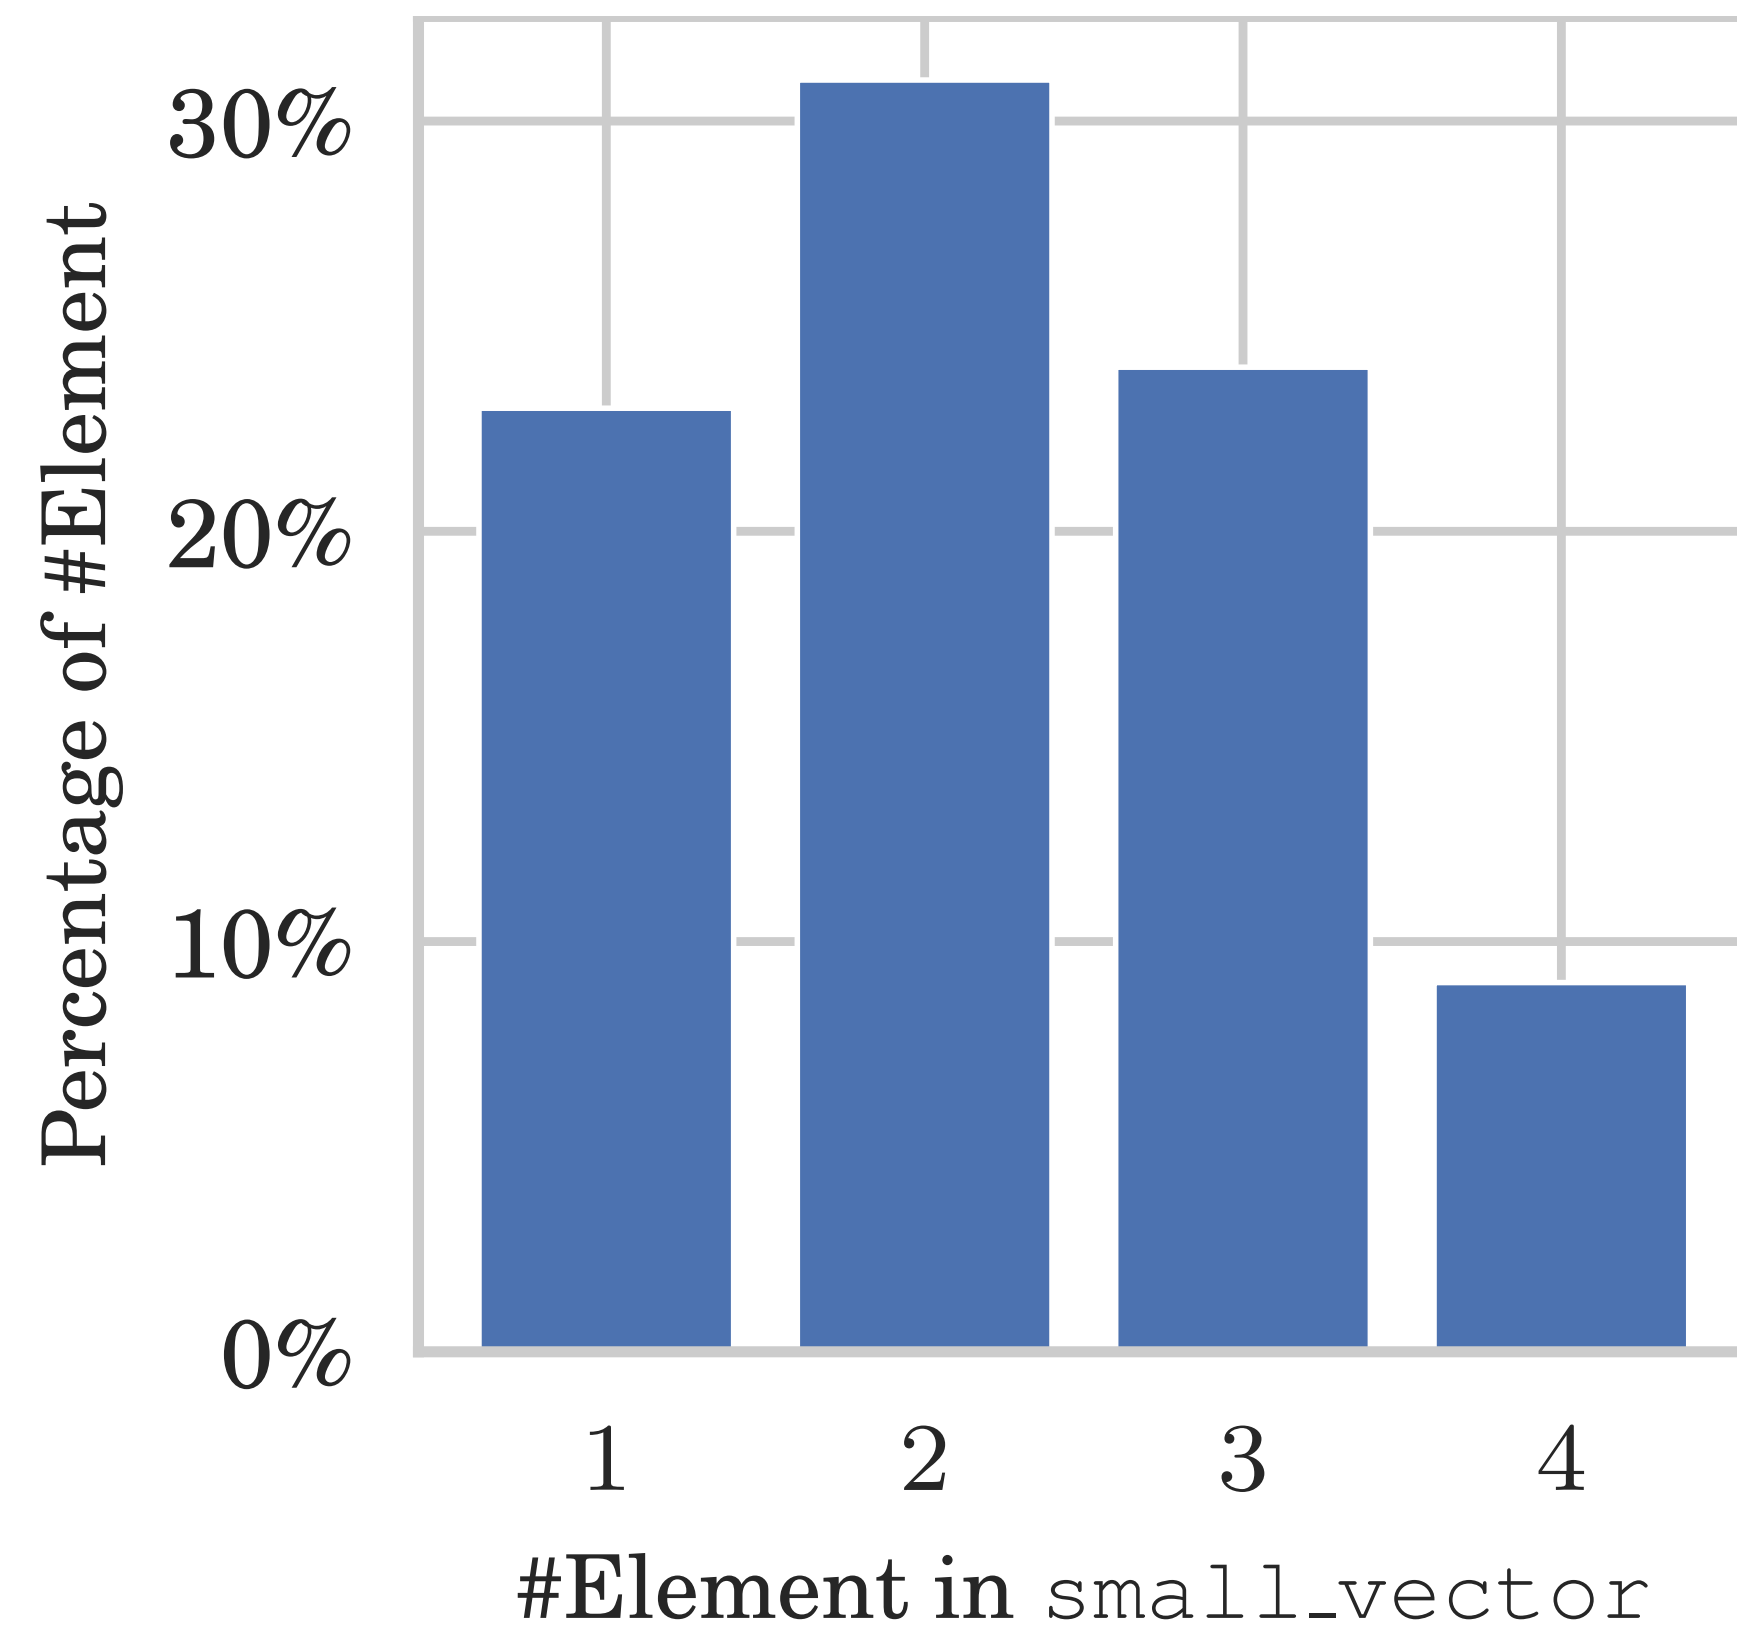
$$\textit{Unused Sz} = (\textit{C.capacity() - C.size()}) * \textit{sizeof(element)}$$

- 2800 nested objects of type ‘FooBar’ were analysed. Top of the list:

Name	TypeName	Number	ElemStatSz	Length	Capacity	DynSz	TotalSz	UnusedSz (Cumulative)
Feature	folly::small_vector<X, 10>	414132	8	2.359	10	0	88	25315061

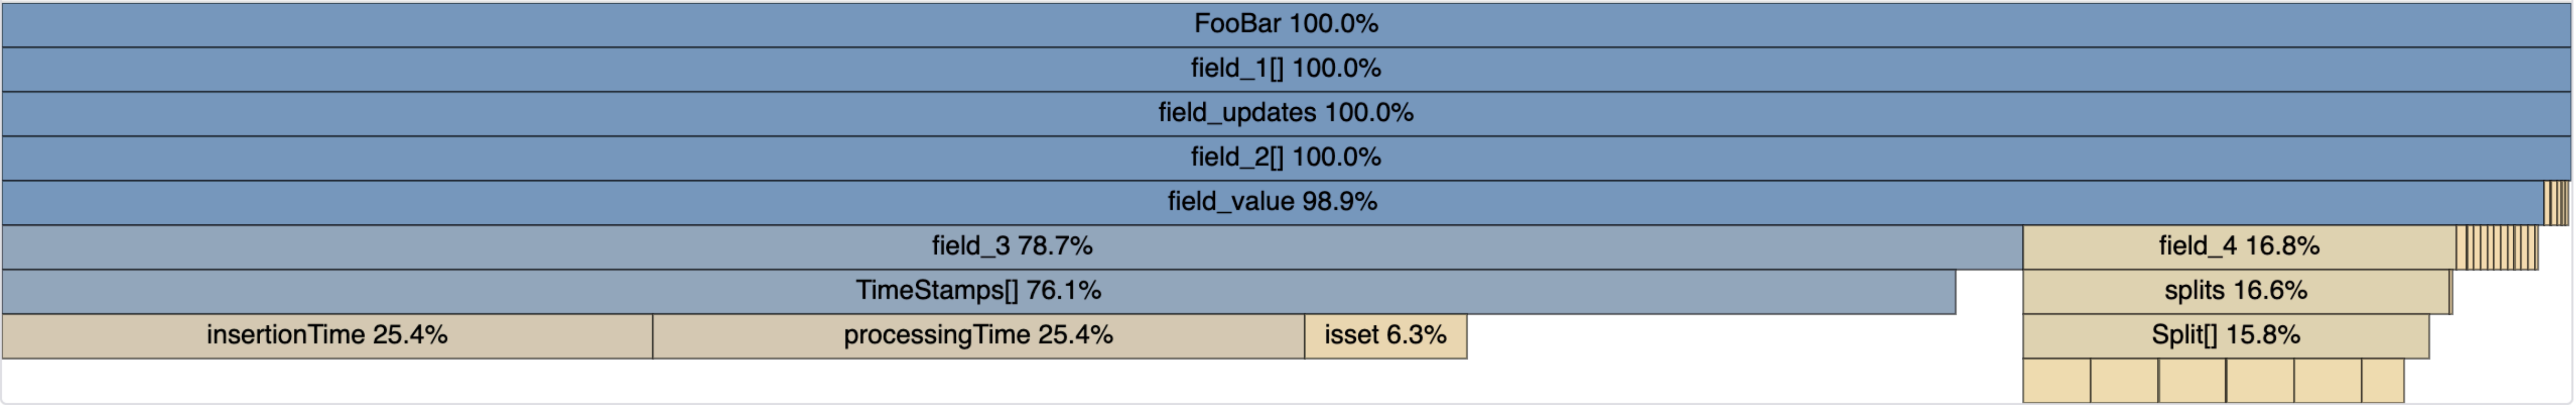
- UnusedSz (cumulative) = (10 - 2.359) * 8 * 414132

Applied Example 1: Results



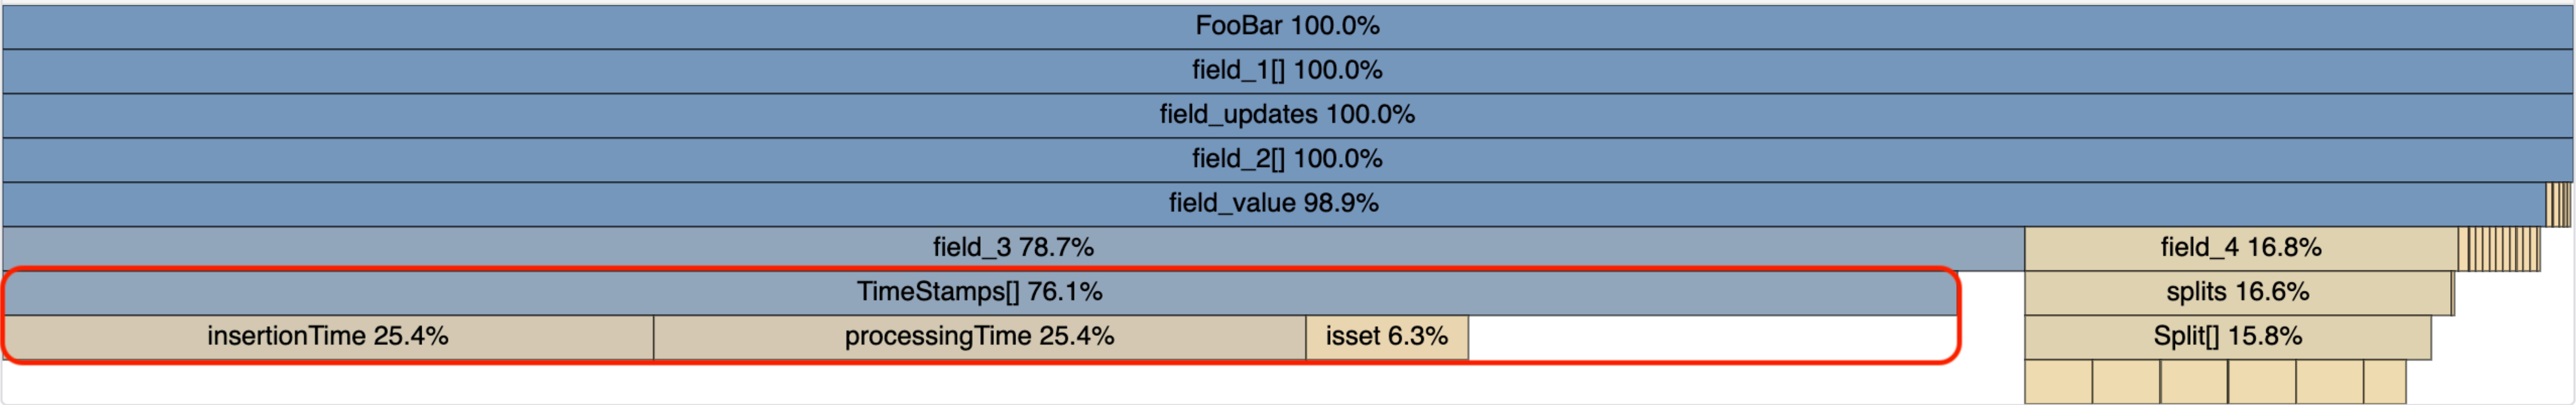
Applied Example 2

- 1000 objects of type ‘FooBar’ were sampled and introspected



Applied Example 2

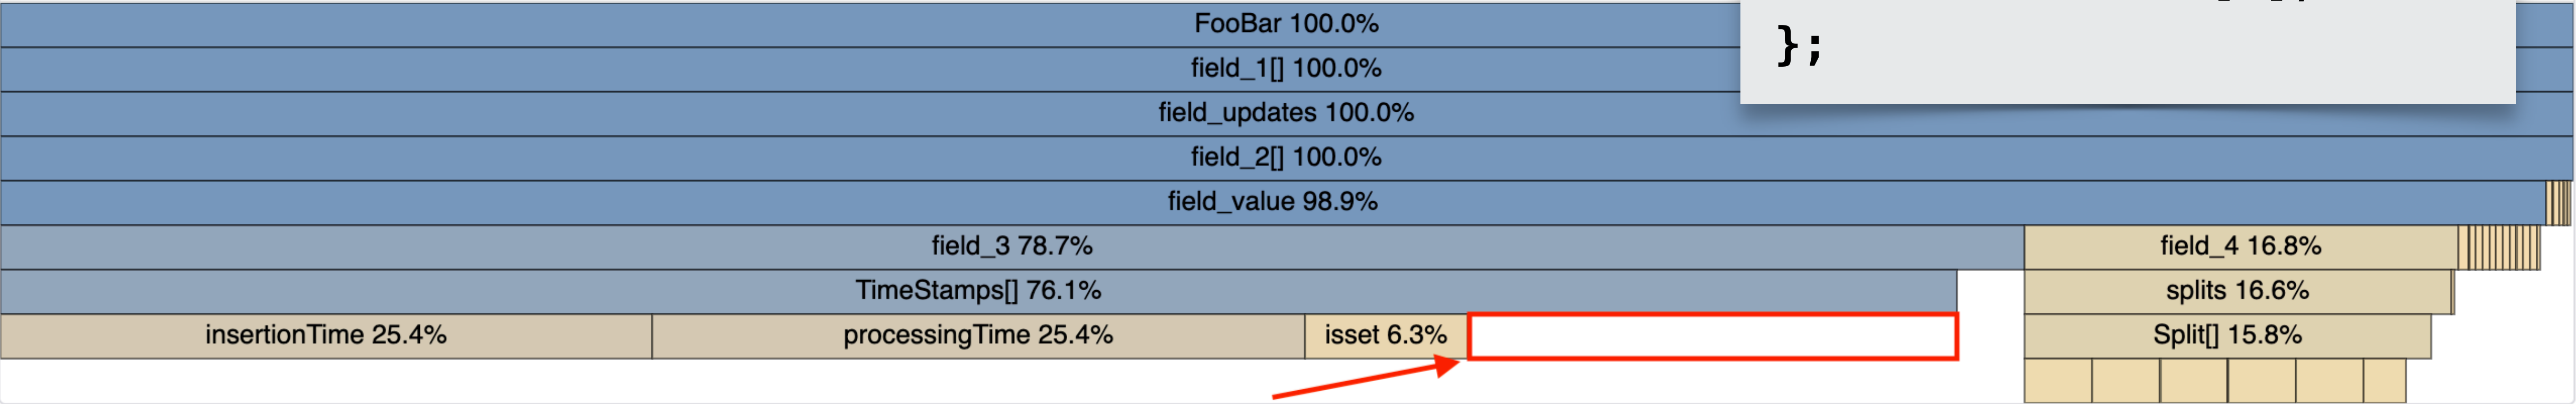
- 1000 objects of type 'FooBar' were sampled and introspected



Applied Example 2

- 1000 objects of type 'FooBar' were sampled and introspected

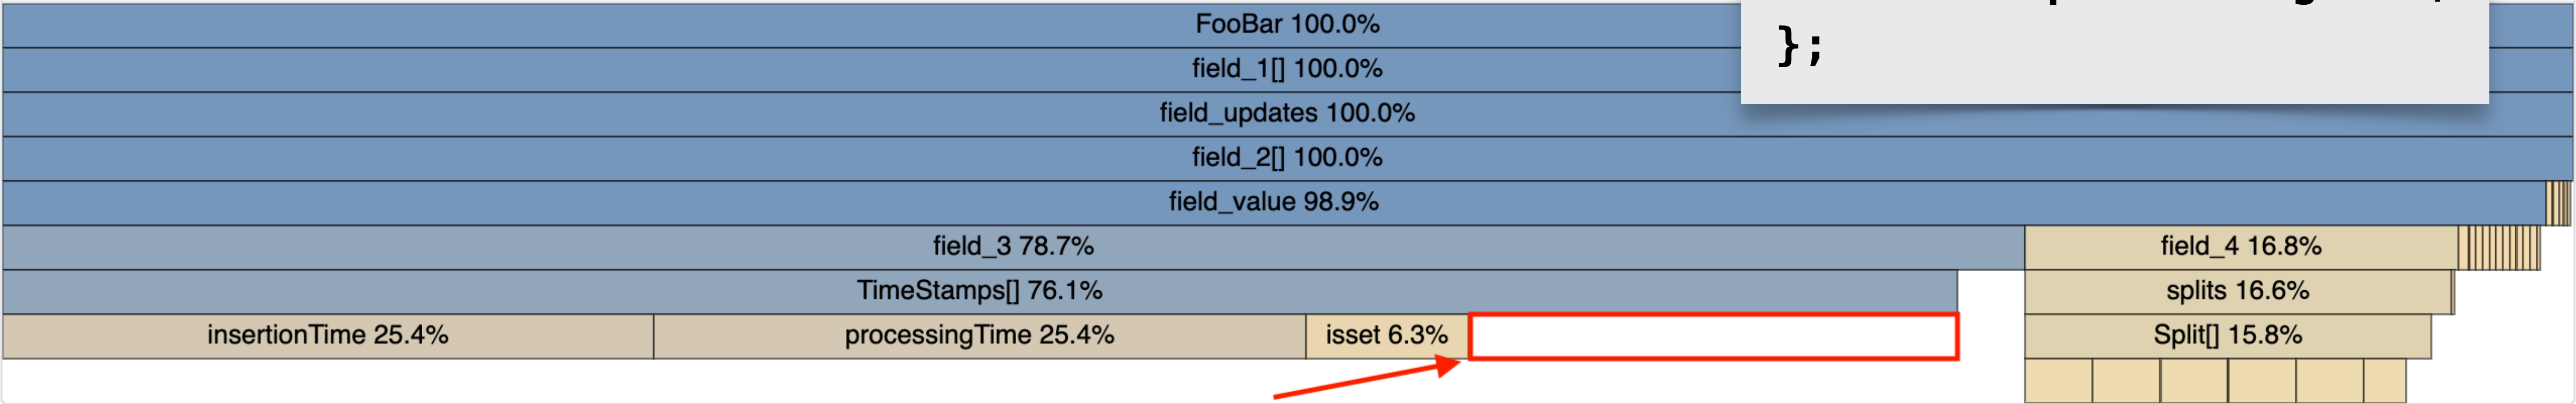
```
struct TimeStamps {  
    int64_t  insertionTime;  
    int64_t  processingTime;  
    bool     isset[2];  
};
```



Applied Example 2

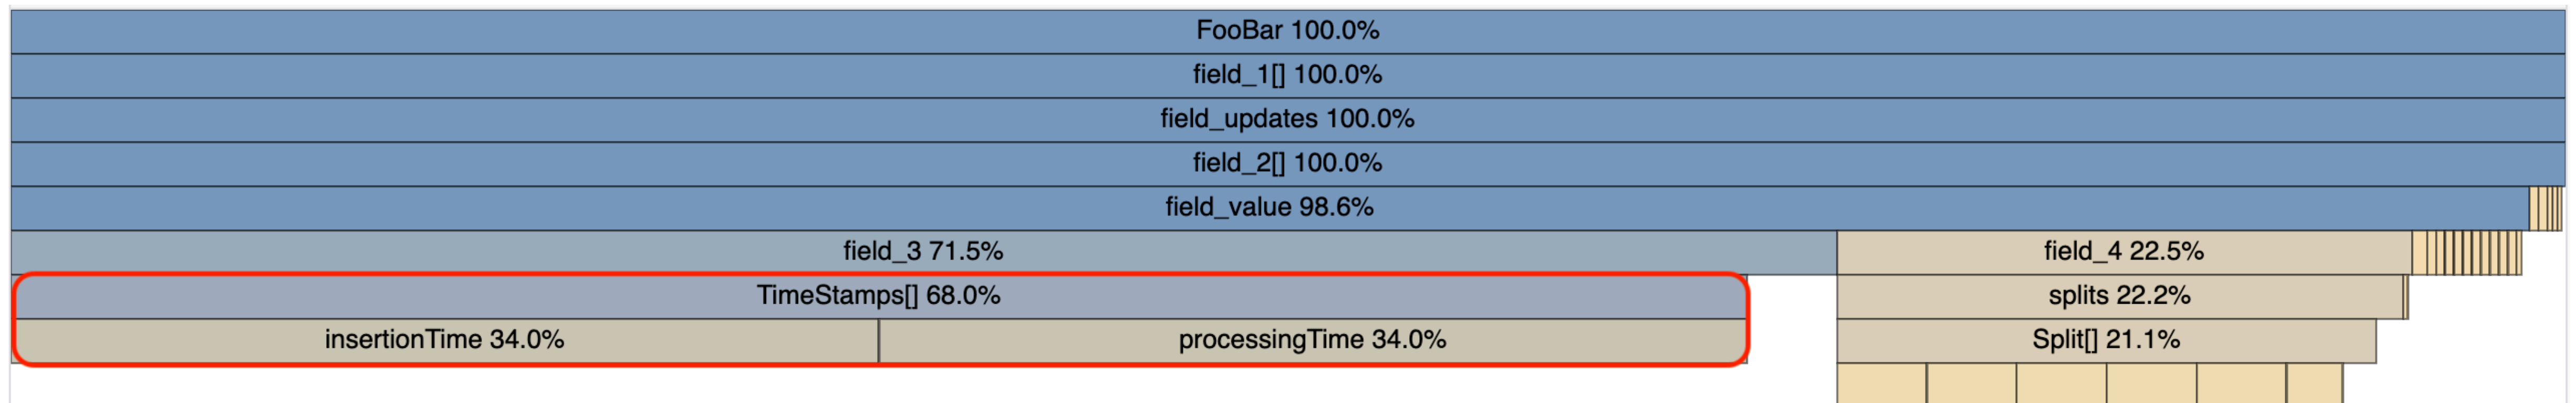
- 1000 objects of type 'FooBar' were sampled and introspected

```
// Thrift source
struct TimeStamps {
    1: i64 insertionTime;
    2: i64 processingTime;
};
```



Applied Example 2

- A thrift annotation, gets rid of isset field to avoid padding
- Profile collected after the change was applied
- Saved ~10% of p95 memory for that service



Gotchas / limitations

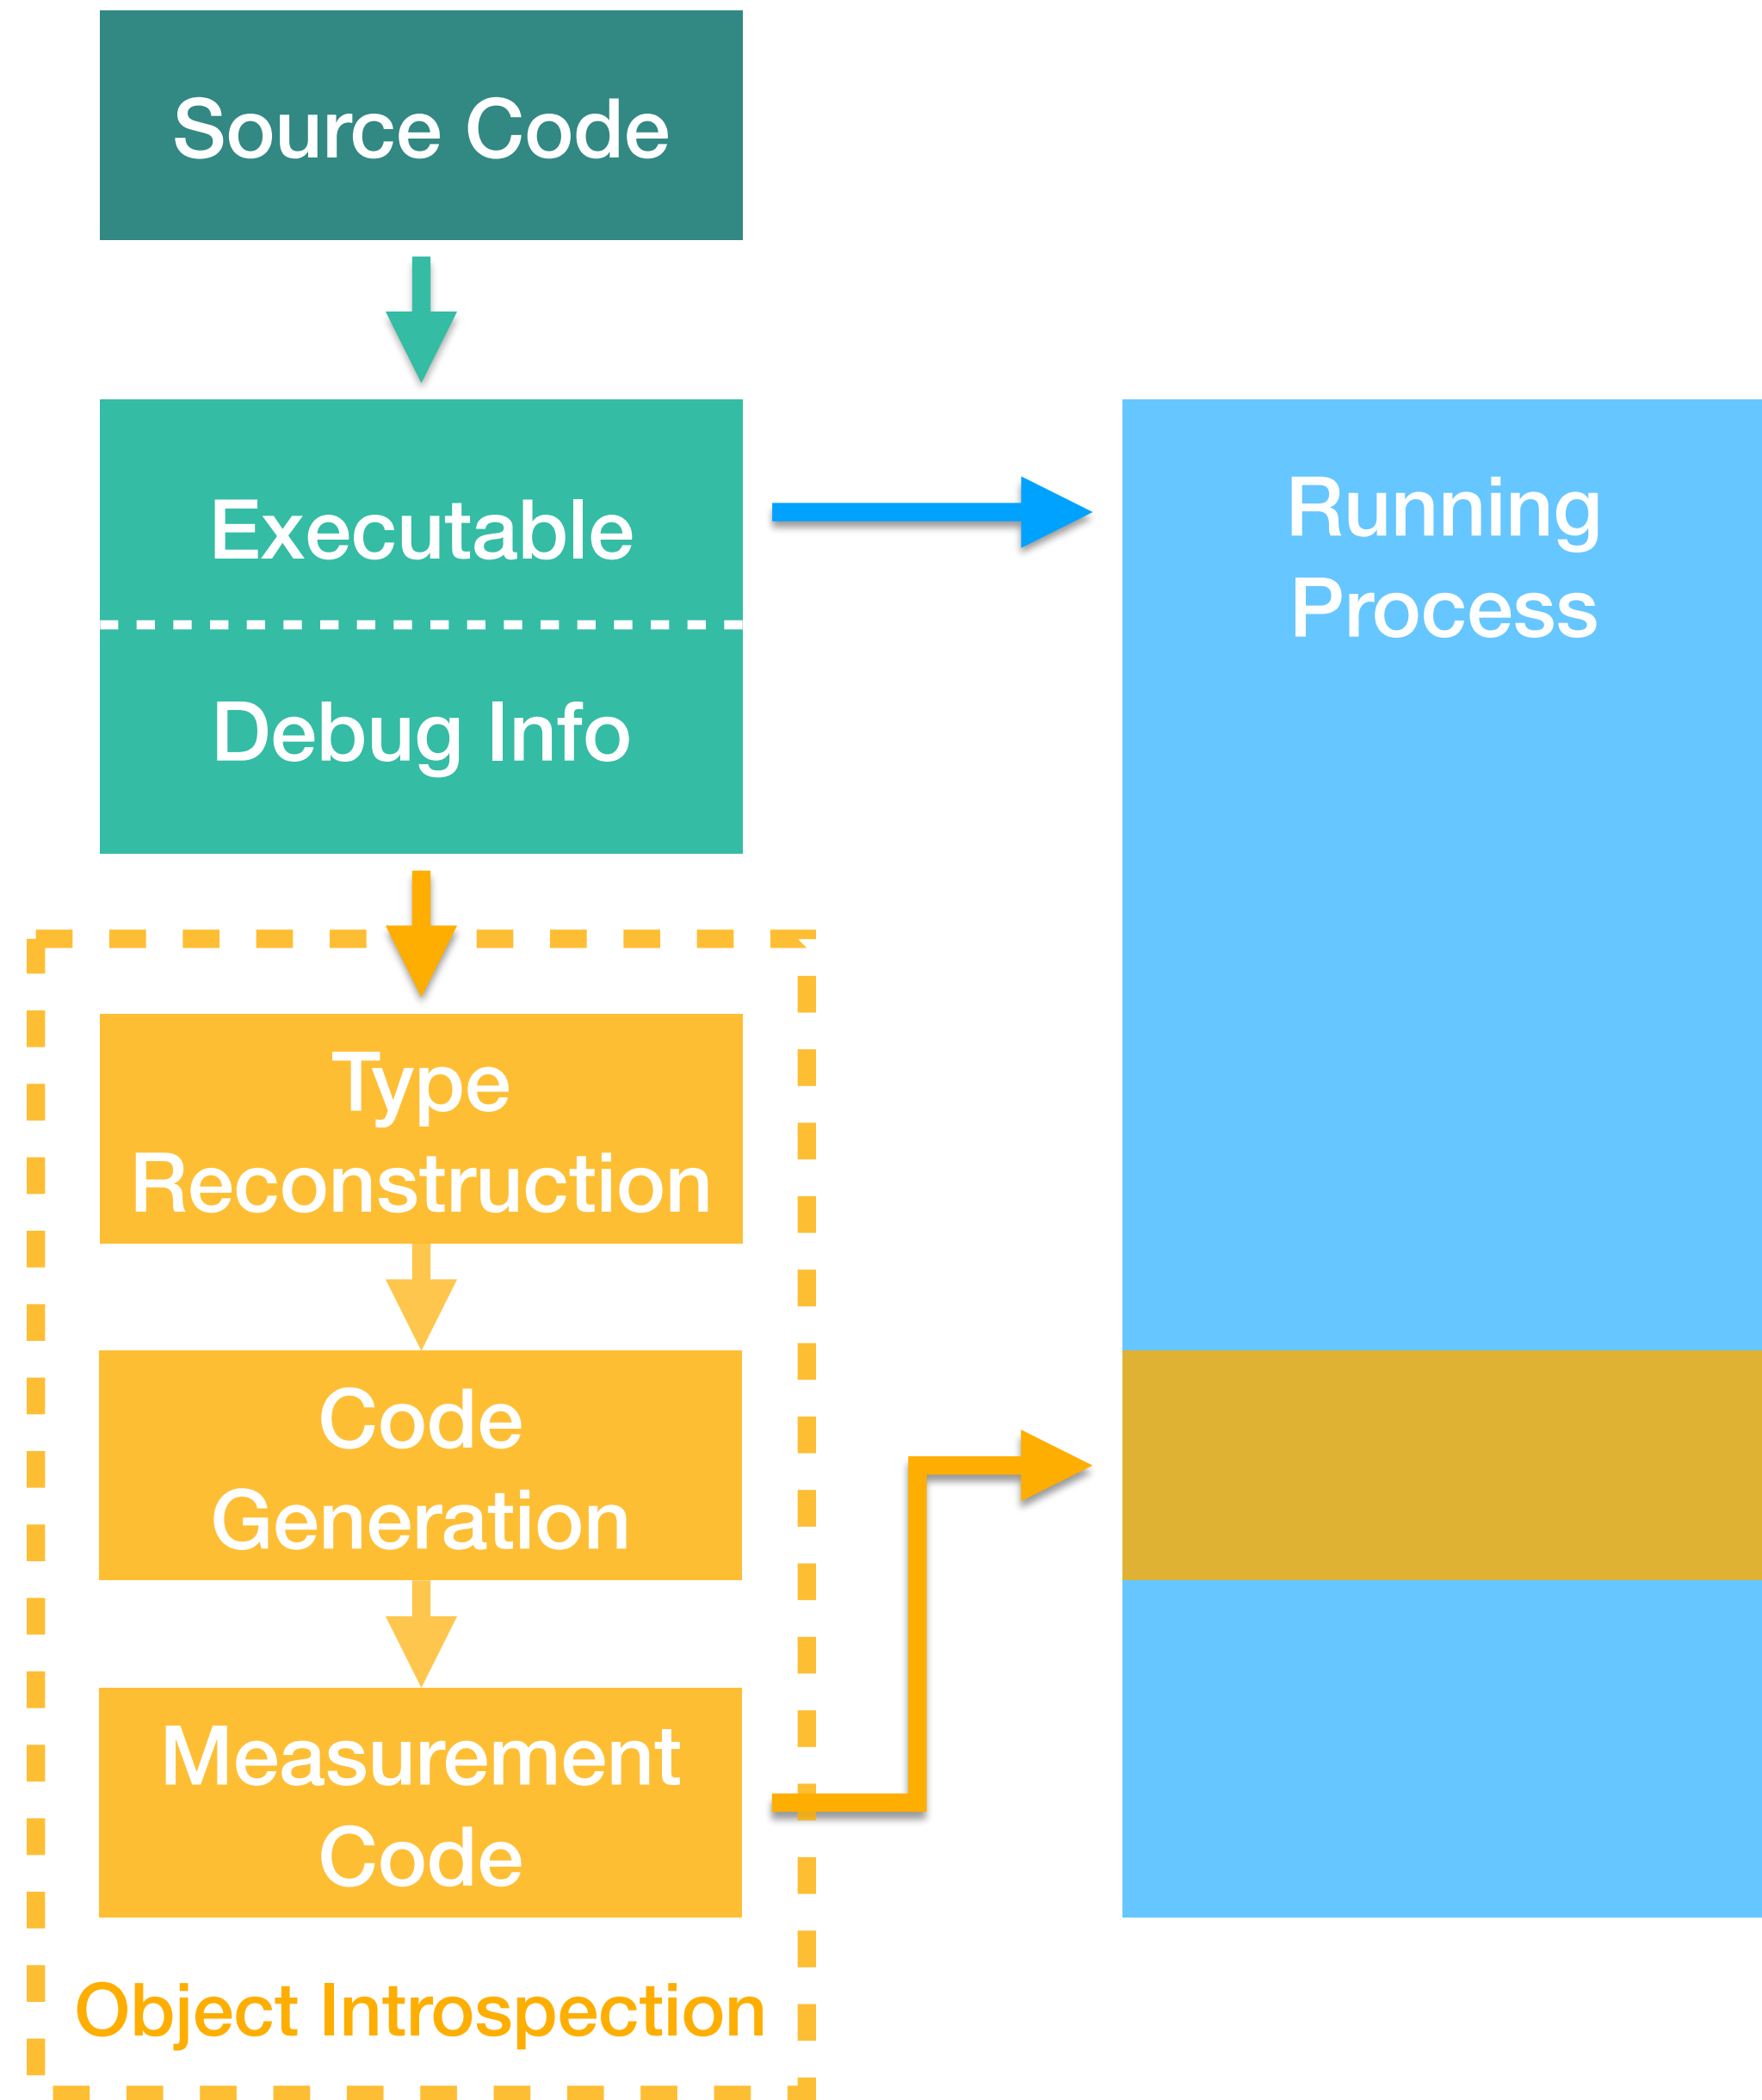
- Currently Linux, DWARF + x86_64 only.
- Dynamic linking
- Debug binaries
- (oid) Inlining!
- Data races

Gotchas / limitations

- User defined containers
- Template specialisation of containers
- Virtual Inheritance
- C style unions
- Pointers & `void*`

Outline

- Debug Info Analysis
 - Type/Layout Reconstruction
 - Code Generation
- Object Introspection as a Library
- Object Introspection as a Profiler
- Object Introspection Applied





**object
introspection.org**

+ 23

Object Introspection:

A Revolutionary Memory Profiler for C++ Objects

JONATHAN HASLAM
& ADITYA SARWADE



20
23

