

Delivering safe C++

Bjarne Stroustrup

Columbia University

www.stroustrup.com



Overview

- The challenges of safety
 - What is “safety”?
- C++ Evolution
 - with a focus on safety
- C++ Core Guidelines
 - How to write good contemporary C++
- Safety Profiles
 - How to guarantee safety









A cause for concern (not panic)

- The overarching software community across the private sector, academia, and the U.S. Government have begun initiatives to drive the culture of software development towards utilizing memory safe languages.
- ...
- NSA advises organizations to consider making a strategic shift from programming languages that provide little or no inherent memory protection, such as C/C++, to a memory safe language when possible. Some examples of memory safe languages are C#, Go, Java, Ruby™, and Swift®.
 - NSA: <https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2023/p2739r0.pdf>



To contrast (not a cause for complacency)

- **February 2023 Headline: C++ still unstoppable**
- Last month, C++ won the TIOBE programming language of the year award for 2022. C++ is continuing its success in 2023 so far. Its current year-over-year increase is 5.93%. This is far ahead of all other programming languages, of which the most popular ones only gain about 1%.


Feb 2023	Feb 2022	Change	Programming Language	Ratings	Change
1	1		 Python	15.49%	+0.16%
2	2		 C	15.39%	+1.31%
3	4	▲	 C++	13.94%	+5.93%
4	3	▼	 Java	13.21%	+1.07%
5	5		 C#	6.38%	+1.01%
6	6		 Visual Basic	4.14%	-1.09%

- But what does Tiobe measure?
- But this implies that what we do matters to billions of people – for good and bad

We must address the “safety” issue

- There is a real, serious problem for many uses and users
 - Bugs and security violations from bad code in various language (incl. C++) and other causes
 - Diversion of resources from C++ to other languages
 - Discouraging people from learning C++
- Massive improvements are possible in many areas
- C++ has a massive image problem (“C/C++”) 
 - And it is getting worse
- Governments and large corporations can coerce developers 

We == WG21 + community

There is no C/C++ language.
Write contemporary C++
- Ignoring the safety issues now would hurt large sections of the C++ community and undermine much of the other work we are doing to improve C++.
 - So would focusing exclusively on safety
- Offering guaranteed safety will be in the best tradition of C++ 

An opportunity

Complete type-and-resource safety

- Is an ideal (aim) of C++
 - From very early on (1979)
 - But “being careful” doesn’t scale
- Requires judicious programming techniques
 - Supported by libraries
 - Enforced by language rules and static analysis
 - The basic model for achieving that can be found in [A brief introduction to C++'s model for type- and resource-safety](#) (2015) and [Type-and-resource safety in modern C++](#) (2021).
- No limitations of what can be expressed
 - Compared to traditional C and C++ programming techniques
- No added run-time overhead
 - Except necessary range checking

Type-and-resource safety

- Every object is accessed according to the type with which it was defined (type safety)
 - Every object is properly constructed and destroyed (resource safety)
 - Every pointer either points to a valid object or is the **nullptr** (memory safety)
 - Every reference through a pointer is not through the **nullptr** (often a run-time check)
 - Every access through a subscripted pointer is in-range (often a run-time check)
 - That
 - Implies range checking and elimination of dangling pointers (“memory safety”)
 - Is just what C++ requires
 - Is what most programmers have tried to ensure since the dawn of time
 - The rules are more deduced than invented
- Enforcement rules are mutually dependent.
Don't judge individual rules in isolation

Constraints on a solution

- C++ must serve wide variety of users/areas
 - One size doesn't fit all
 - C++ is (also) a systems programming language – we can't “outsource” dangerous operations to some other language
 - We can't just break billions of lines of existing code
 - Even if we wanted to – major users would insist on compatibility (probably compatibility by default)
 - We can change the use of C++
 - We can't just “upgrade” millions of developers
 - And teaching material, courses, videos, books, articles
 - If you want a shiny new language, please go ahead
 - But it won't be C++ or the job of WG21
- Stability is a feature
- But we ***must*** improve

Challenges

- Describe a type-safe C++ use
 - No violations of the static type system
 - No resource leaks
- Convince developers to use that safe (or just safer) styles of use
 - Except where it is not appropriate
 - Direct use of system and hardware resources
 - Need for ultimate efficiency
 - Implementation of code that cannot be proven safe (e.g., some linked structures)
- Get this to work at scale
 - Not just “academic” examples
- Note: there is lots of great C++ “out there”
 - For any definition of “great” including “reliable over decades”

An opportunity
to finally accomplish
one of C++’s original aims

Safety is not just type safety

- **Logic errors:** perfectly legal constructs that don't reflect the programmer's intent, such as using `<` where a `<=` or a `>` was intended.
- **Resource leaks:** failing to delete resources (e.g., memory, file handles, and locks) potentially leading to the program grinding to a halt because of lack of available resources.
- **Concurrency errors:** failing to correctly take current activities into account leading to (typically) obscure problems (such as data races and deadlocks).
- **Memory corruption:** for example, through the result of a range error or by accessing and memory through a pointer to an object that no longer exists thereby changing a different object.
- **Type errors:** for example, using the result of an inappropriate cast or accessing a union through a member different from the one through which it was written.
- **Overflows and unanticipated conversions:** For example, an unanticipated wraparound of an unsigned integer loop variable or a narrowing conversion.
- **Timing errors:** for example, delivering a result in 1.2ms to a device supposedly responding to an external event in 1ms.
- **Allocation unpredictability:** for example, ban on free store allocation "after the engine starts."
- **Termination errors:** a library that terminates in case of "unanticipated conditions" being part of a program that is not allowed to unconditionally terminate.

Security is not just memory safety

- Physical break-ins
 - Spies (insider attacks)
 - Spear phishing
 - Door rattling
 - Denial of service attacks
 - SQL injection
 - Corrupted input/Data
-
- Rule of thumb: always first attack the weakest link



Languages are not safe – uses can be

- All “safe” general-purpose languages have “escape clauses”
 - To access system and hardware resources
 - E.g., the operating system
 - To improve efficiency of key abstractions
 - E.g., linked data structures
 - “Trusted” code segments to work with unsafe code segments
 - Libraries
 - Code written under less stringent rules (e.g., old code)
 - Code written in other languages
- Often the “escape clause” is C++
 - So, C++ needs to be able to do unsafe things
 - Do so efficiently
 - While being safe where it matters
 - guaranteed
 - Preferably by default.

This is **not** an excuse
to ignore safety issues

We **must** improve
where needed

C++ evolution – with a focus on safety

- It all started decades ago
 - 1979
- C++ was designed to be an evolving language
 - Relying on feedback is just good engineering
- Static type safety was an ideal from day #1
 - It still is
 - But an elusive ideal under real-world constraints
 - In general, “perfection” is an elusive concept



Being the world's best
at just one or two things
isn't enough

The earliest aims – Day #1 (1979)

- **Efficient use of hardware** – based on C
 - Direct access to hardware and system resources
 - No elaborate run-time support (e.g., no GC)
- **Manage complexity** – based on Simula
 - Classes
 - Strong static type checking

Overarching aims

I needed a tool for building a system
No language could do both



Only a few years earlier it was “well known”
that you couldn’t write good systems software
except in assembler

Argument type checking – 1980

- C in 1979 aka “Classic C”

```
double sqrt();  
double x = sqrt(2);           /* crash */
```

- C with classes in 1980 (C and C++ today)

```
double sqrt(double);          // argument type required  
double x = sqrt(2);
```

Strong static type checking was and is the ideal

Argument type checking enforced – 1983

- Somewhat controversial
 - “I have to look at a declaration to figure out what a call means?”
- Definitely incompatible
 - A major point against C++ for people who didn’t like something or other about C++
 - But upgrade was easy: “I convert 10,000 lines of C to C++ per day”

- Essential for
 - type checking
 - overloading
 - adding user-defined types
 - consistent linking

A few things are so important
that you must introduce incompatibilities

But you always pay dearly for that:
developers and teachers can be very conservative

- Type-safe linking
 - Ensuring much better consistency across separate compilations
 - Not perfect, but now we can use C++20 modules

Key idea: “Represent concepts in code”

- Direct representation of ideas in code
 - Focus on classes
- Make code more declarative
- Make more information available to compilers
- Early examples
 - Vector
 - String
 - File handle
 - Concurrent task
 - Message queue
 - Coroutine
 - Hash table
 - Graphical shape
 - Complex number
 - Infinite integer

Some are still not standard



immutability

- Constants
 - `const int x = 7;`
 - `const string s = "Immutable";`
- Interfaces
 - `const strcpy(char*, const char*);` *// today: don't use strcpy() ; unsafe (potential range error)*
- Historical factoid
 - My original design was for **readonly** and **writeonly**

RAII (1979 and later)

- From my 1979 lab book:
 - A “new function” creates the run-time environment for member functions
 - A “delete function” reverses that
- Later (1983)
 - “new function” -> constructor
 - “delete function” -> destructor
 - Because constructors and destructors were/are not just for the new and delete operators.
- A slightly later formulation (1980s)
 - A constructor establishes a class invariant (if any) for an object
 - A destructor releases all resources owned by the object
- And (1988)
 - “Resource Acquisition Is Initialization”
 - Apologies for that name

Memory isn't the only critical resource

- File handles
- Locks
- Sockets
- Shaders
- ...

Resources and Errors

- A resource is something that must be acquired and released after use
 - E.g., files, memory, locks, database transactions, communication channels, GUI connections, threads
 - Explicit release is error-prone
 - Resource exhaustion can render a system inoperative
 - I wouldn't call a system relying on explicit release safe

```
void f(const char* p)           // unsafe, naïve use
{
    FILE* f = fopen(p,"r");     // acquire
    // use f
    fclose(f);                  // release
}
```


pointer

Something
representing
a
resource

RAII (Resource Acquisition Is Initialization)

```
// use an object to represent a resource
class File_handle {           // belongs in some support library
    FILE* p;
public:
    File_handle(const char* pp, const char* r)
        { p = fopen(pp,r); if (p==0) throw File_error(pp,r); }
    File_handle(const string& s, const char* r)
        { p = fopen(s.c_str(),r); if (p==0) throw File_error(s,r); }
    ~File_handle() { fclose(p); } // destructor
    // copy operations
    // access functions
};

void f(string s)
{
    File_handle fh {s, "r"};    // now: ifstream fh{s}
    // use fh
}
```




OOP (1981 – 1984)

- Encapsulation
- Well-defined interfaces
 - Classes
 - Abstract base classes


Problem: Too many pointers in interfaces

- Overloading
 - Built-in and user-defined types
 - Operator overloading (initially for resource management)
 - References: simpler and safer argument passing



Later (1988 – 2011)

- Templates
 - Compile-time selection of implementations
 - C++20 finally got concepts – precisely defined interfaces
- Exceptions
 - Guaranteed error-handling – or termination
 - Proper interaction with resource management (RAII)
- Containers
 - No need to fiddle with arrays (and pointers)
 - Enable range checking
- Algorithms
 - C++98: **sort(begin(v),end(v))**
 - C++20: **sort(v)**
- “smart pointers” (resource management pointers)
 - **unique_ptr, shared_ptr**



Range-for and span (2011, 2017)

- Enable efficient range checking



```
void f(int* p, int n)    // what if p==nullptr? What if n is not the number of elements pointed to?  
{  
    for (int i =0; i<n; ++n)    // what if someone "messes with" the control variable?  
        cout << p[i] << ' ';  
}
```

```
void g(span<int> s)  
{  
    for (const auto& x : s)    // no loop variable to get wrong  
        cout << x << ' ';  
}
```

D&E (1994) – select language-technical rules


- No implicit violations of the static type system
- Provide as good support for user-defined types as for built-in types
- Say what you mean
 - Emphasizes declarative styles and abstraction
- Syntax matters (often in perverse ways)
 - In general, verbosity is to be avoided
- Leave no room for a lower-level language (except assembler)
- Preprocessor usage should be eliminated
- Missing
 - Make simple tasks simple
 - Make error handling regular

Not articulating those
caused problems that still linger



Benefits come from using C++ well

- Use
 - RAI
 - **const**
 - Containers
 - Resource management pointers
 - Algorithms
 - Range-**for**
 - **span**
 - ...
- Avoid
 - Owning raw pointers
 - Subscripting raw pointers
 - Unchecked pointer dereferencing
 - Uninitialized variables
 - ...




Published 2002

No pointers or arrays
until chapter 17.
Then, with warnings

If you want “safe C++”

- Don't write “C/C++”
 - There is no such language
 - But there is such use
 - Most documented safety and security violations are in such code
- Evolve your style towards what's provably safe
 - Encapsulate “messy low-level” code



“Being careful” doesn’t scale

- We must
 - Formulate rules for safe use
 - For a variety of safe uses
 - Provide ways of verifying that the rules are adhered to
- Articulate guidelines
 - A start: C++ Core Guidelines
- Enforce guidelines where needed
 - Profiles
 - Enforcing a variety of guidelines
 - Serving a variety of needs

Do you use the safer features?

Consistently?

If not, why not?

State of affairs

- The parts of what I am describing have been tried
 - Many “at scale” (e.g., range-checked **string**, **vector**, and **span**)
 - But nowhere has all been integrated into a single system and systematically enforced
- Much is influenced by work on the C++ core guidelines
 - But this is not just about guidelines
 - We need enforced rules
- The aim is guaranteed type-and-resource safe C++
 - And more, e.g., safe arithmetic
 - Paths to gradual adoption
 - Major improvements can be achieved today, e.g., consistent range checking)
- This is not just about safety
 - better use of the type system improves productivity
 - and often also improves performance



C++ Core Guidelines



General strategy

- Rely on static analysis to eliminate potential errors
 - Static analysis is impossible for arbitrary code
 - Global static analysis is very expensive
- Rely on rules to simplify the language used
 - to the point where local static analysis is possible
- Provide libraries to make relying on the rules practical
 - Pleasant to use
 - Efficient to run



High-level rules – “Philosophy”

- Provide a conceptual framework
 - Primarily for humans
- Many can't be checked completely or consistently
 - *P.1: Express ideas directly in code*
 - *P.2: Write in ISO Standard C++*
 - *P.3: Express intent*
 - *P.4: Ideally, a program should be statically type safe*
 - *P.5: Prefer compile-time checking to run-time checking*
 - *P.6: What cannot be checked at compile time should be checkable at run time*
 - *P.7: Catch run-time errors early*
 - *P.8: Don't leak any resource*
 - *P.9: Don't waste time or space*
 - *P.10: Prefer immutable data to mutable data*
 - *P.11: Encapsulate messy constructs, rather than spreading through the code*
 - *P.12: Use supporting tools as appropriate*
 - *P.13: Use support libraries as appropriate*



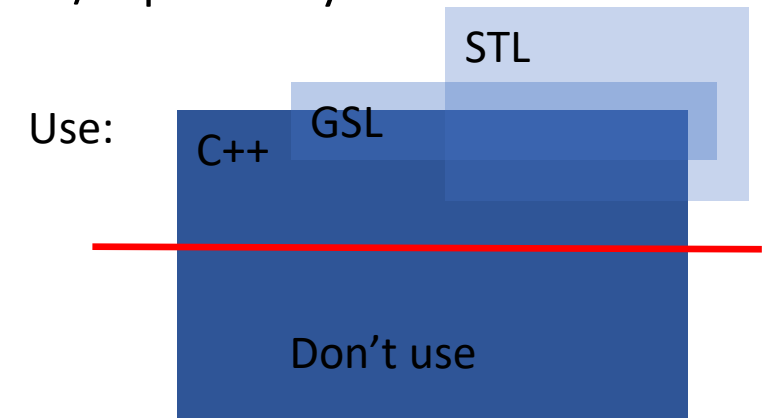
Lower-level rules

- Provide enforcement
 - Some complete
 - Some heuristics
 - Many rely on static analysis
 - Some beyond our current tools
 - Often easy to check “mechanically”
- Primarily for tools (static analysis)
 - To allow specific feedback to programmer
- Help to unify style
- Not minimal or orthogonal
 - *F.16: Use **T^*** or ***ownerT^**** to designate a single object*
 - *R.2: In interfaces, use raw pointers to denote individual objects (only)*
 - *ES.20: Always initialize an object*



Subset of superset

- Simple sub-setting doesn't work
 - We need the low-level/tricky/close-to-the-hardware/error-prone/expert-only features
 - For implementing higher-level facilities efficiently
 - Many low-level features can be used well
 - We need the standard library
- Extend language with a few abstractions
 - **Use** the STL
 - **Add** a small library (the GSL)
 - **No** new language features
 - Messy/dangerous/low-level features can be used to implement the GSL
 - **Then** subset
- What we want is “**C++ on steroids**”
 - Simple, safe, flexible, and fast
 - Not a neutered subset



No change of meaning:
The resulting code is ISO C++

We can

- Eliminate
 - Uninitialized variables
 - Range errors
 - Nullptr dereferencing
 - Resource leaks
 - Dangling references
 - Unions (use variants)
 - Casts
 - Underflow and overflow
 - Data races



Not discussed in this talk

- “Dangerous features” can be used in explicitly unverified code
- “Just test everywhere at run time” is **not** an acceptable answer
 - Hygiene rules + Static analysis + Run-time checks

Uninitialized variables

- Initialize every variable
 - Simple and effective
 - Implicit initialization of types with default constructors is OK

```
void f()
{
    int x;           // not OK
    int y = 7;       // OK
    std::string s;   // OK
    char* p;         // not OK
}
```

- Mark exceptional cases **[[uninitialized]]**
 - E.g., I/O buffers

Range checking

- Don't subscript raw pointers
 - `void f(int* p, int x) { p[x] = 7; } // not OK`
 - Except in the implementation of abstractions
- Use abstraction with sufficient data to range check
 - E.g., **vector** and **span**
 - `void f(span<int> s, int x) { s[x] = 7; } // OK for a checking span`
- Run-time range check every subscript operation
 - E.g., **vector** and **span** must be range checked
- Range-**for**
 - Prefer over C-style for loops
- Algorithms
 - Prefer range algorithms

Many more details
in the references

Null pointer dereferencing

- Don't dereference an unchecked pointer
 - `void f0(int* p) { *p = 7; }` *// not OK*
 - `void f1(int* p) { if (p) *p = 7; }` *// OK*
 - `void f2(not_null<int*> p) { *p = 7; }` *// OK (not_null constructor checks)*
 - `void f3(span<int> s) { s[2] = 2; }` *// OK (for checked span)*
- Except in the implementation of abstractions

No resource leaks

- We know how
 - Root every object in a scope
 - `vector<T>`
 - `string`
 - `ifstream`
 - `unique_ptr<T>`
 - `shared_ptr<T>`
 - RAII
 - “No naked **new**”
 - “No naked **delete**”



A resource leak is potentially damaging

Pointer Misuse

- Many (most?) uses of owning pointers in local scope are not resource safe

```
void f(int n, int x)    // not OK
{
    Gadget* p = new Gadget{n};    // look I'm a java programmer! ☺
    // ...
    if (x<100) throw std::runtime_error{"Weird!"};    // leak
    if (x<200) return;    // leak
    // ...
    delete p;    // I want my garbage collector! ☹
}
```

- But
 - garbage collection would not release non-memory resources
 - why use a “naked” pointer?

Resource Handles and Pointers

- Use a cheap, safe “smart” pointer
 - A `std::unique_ptr` releases its object at when it goes out of scope

```
void f(int n, int x)    // OK
{
    auto p = make_unique<Gadget>(n);           // manage that pointer!
    // ...
    if (x<100) throw std::runtime_error{"Weird!"}; // no leak
    if (x<200) return;                          // no leak
    // ...
}
```

- This is simple and cheap
 - No more expensive than a “plain old pointer” used correctly

Resource Handles and Pointers

- But why use a pointer at all?
 - If you can, just use a scoped variable


```
void f(int n, int x)    // OK, and better
{
    Gadget g {n};
    // ...
    if (x<100) throw std::runtime_error{"Weird!";}    // no leak
    if (x<200) return;                                // no leak
    // ...
}
```

A “naked new” is a code smell.
Encapsulate!

- No explicit resource management
 - Code not littered with (easy to forget) **try-catch**
 - No spurious allocations/deallocations
 - No pointers

Dangling pointers

- We ***must*** eliminate dangling pointers
 - Or type safety is compromised
 - Or memory safety is compromised
 - Or resource safety is compromised
 - By “pointer” I mean anything that directly refers to an object
- Eliminated by a combination of rules
 - Distinguish owners from non-owners
 - Assume raw pointers to be non-owners
 - Catch all attempts for a pointer to “escape” into a scope enclosing its owner’s scope
 - **return**, **throw**, out-parameters, long-lived containers, ...
- Statically verified



Dangling pointers

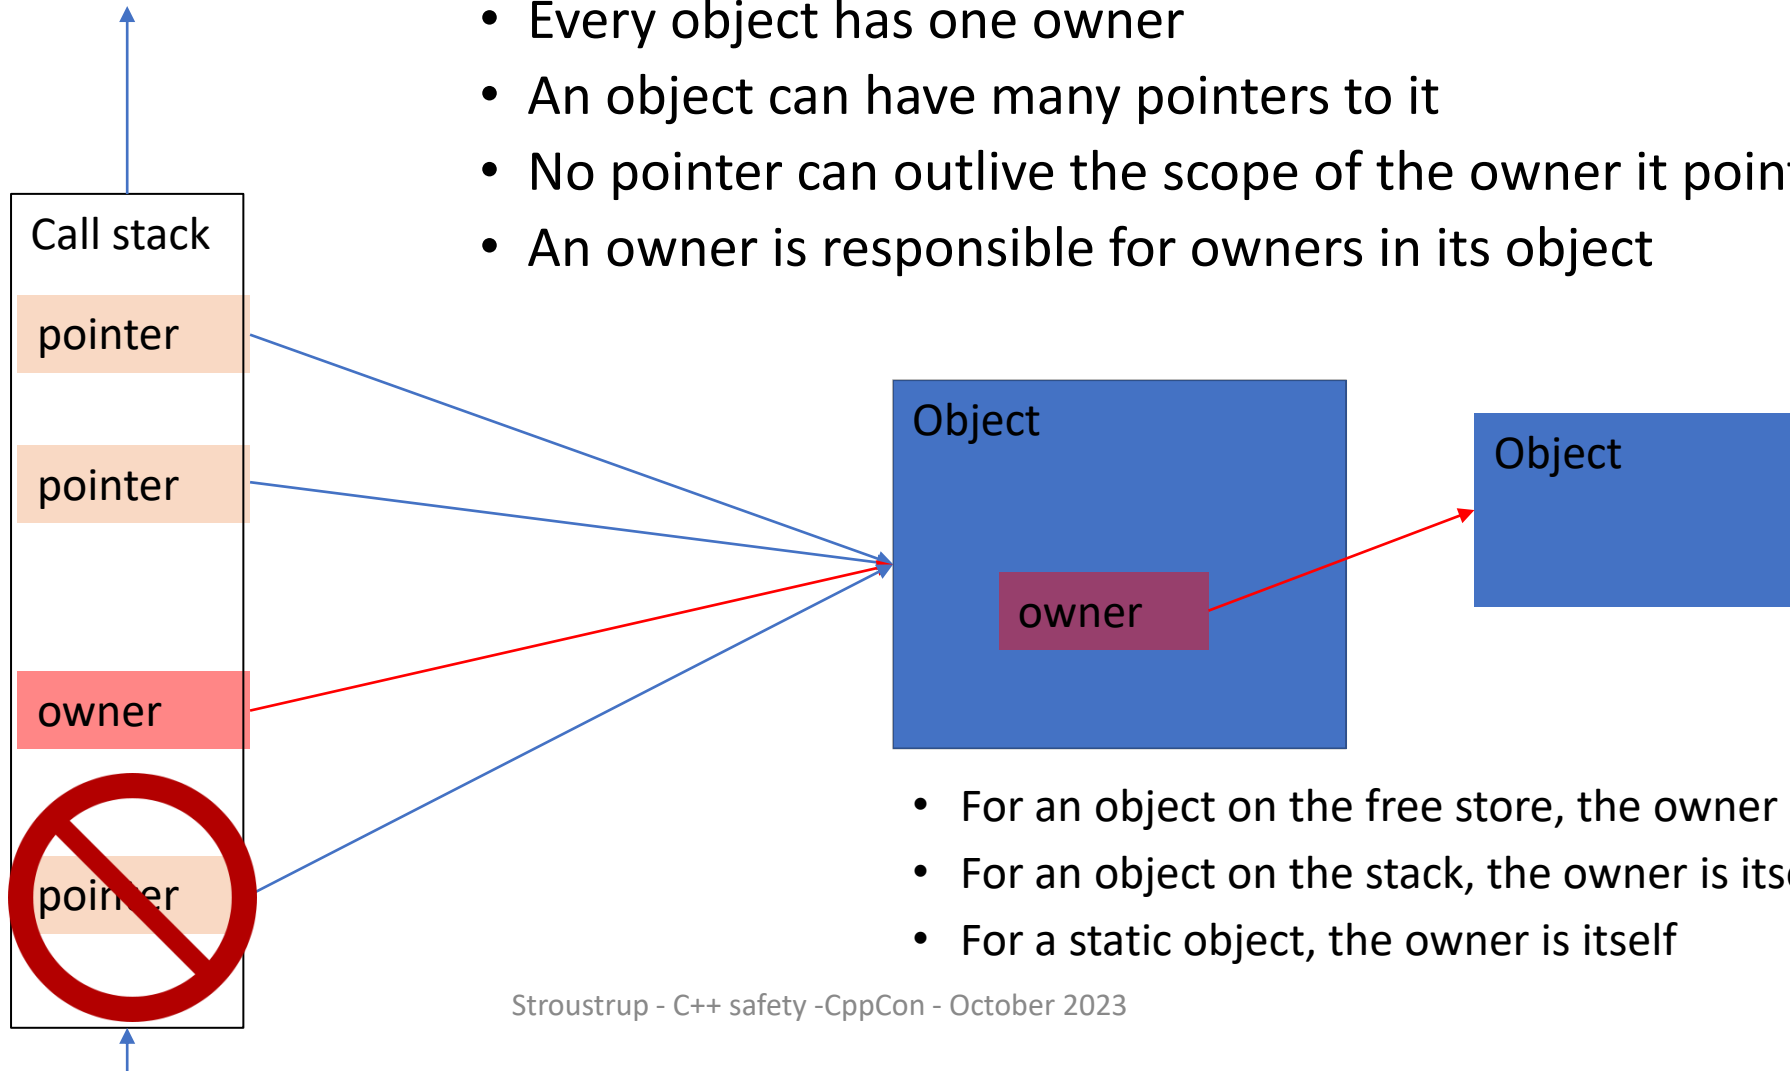
- Ensure that no pointer outlives the object it points to (statically)

```
void f(X* p)
{
    // ...
    delete p;           // not OK: delete non-owner
}

void g()
{
    X* q = new X;       // not OK: assign object to non-owner
    f(q);
    // ... do a lot of work here ...
    q->use();           // Make sure we never get here
}
```

Owners and pointers

- Every object has one owner
- An object can have many pointers to it
- No pointer can outlive the scope of the owner it points to
- An owner is responsible for owners in its object



- For an object on the free store, the owner is a pointer
- For an object on the stack, the owner is itself
- For a static object, the owner is itself

How to avoid/catch dangling pointers

- Any pointer can be passed as an argument, but not stored into a surrounding scope

```
int* glob = nullptr;
```

```
void g(int* p)
{
    glob = p;           // not OK: trying to store away the pointer
}
```

```
void f(int* p)
{
    int x = 4;
    g(&x);               // OK
    g(p);                // OK
}
```

- Enforcement is difficult if returned pointers are calculated in complex code
 - Don't write such code
 - Get a good static flow analyzer (perfect analysis is impossible in general)

How to avoid/catch dangling pointers

- Remember containers

```
void g(vector<int*>& v)
{
    // ...
}

vector<int*> f(int* p)
{
    int x = 4;
    vector<int*> res = {p, &x, new int{7}};    // OK (except for the ownership violation)
    g(res);                                    // OK
    return res;                                // not OK
}
```

Invalidation

- This must be prevented:

```
void f(vector<int>& vi)
{
    vi.push_back(9); // may relocate vi's elements
}

void g()
{
    vector<int> vi { 1,2 };
    auto p = vi.begin(); // point to first element of vi
    f(vi); // UB, may appear to work correctly
    *p = 7;
}
```

- Can be done statically
 - has been done

Invalidation

- “invalidated” == “when a pointer points to an element of a container that may have reallocated its elements”
- “container” == “any object that directly or indirectly could contain a pointer”
 - Classes with pointer members
 - E.g., **vector** and **thread**
 - Smart pointers (they are classes, e.g., **unique_ptr** and **shared_ptr**)
 - Threads (they are classes)
 - Lambdas (they are classes, and remember capture-by-reference)
 - Pointers to pointers
 - References to pointers
 - Arrays of pointers
- “pointer” == “anything that directly refers to an object”

Invalidation assumptions

- No **const** member function invalidates (easily validated)
- Any non-**const** member function invalidates
- Add **[[not_invalidating]]** to non-**const** member functions that do not invalidate
 - E.g., `vector::operator[]()`
 - Easily validated

Not yet Core Guideline

How do we represent ownership?

- High-level: Use an ownership abstraction
- Low-level: Primarily for communicating with C-style interfaces
 - mark owning pointers **owner**
 - Currently, **owner<T*>** is just an alias for **T***
 - An **owner** must be **deleted** or passed to another **owner**
 - A non-**owner** may not be **deleted**

Stay high level
except to implement
necessary abstractions
and interfaces

- Note
 - I talk about pointers
 - What I say applies to anything that refers to an object
 - References
 - Containers of pointers
 - Smart pointers
 - Pointers to pointers
 - Lambda captures
 - ..

Efficient communication
with C is essential

You can copy the pointer of an **owner** but not the ownership

- **owner** is low level, prefer **unique_ptr** or other ownership abstractions

```
void f1(owner<int*> p)      // f1() used_to_take_a_raw_pointer
{
    // ...
    delete p;  // required or f() must transfer ownership before returning
}
```

```
void f2(int* p)
{
    // ...
    delete p;  // error: p is not an owner
}
```

Enforced by static analyzer or compiler

```
void f3(int* p1, owner<int*> p2)
{
    f1(p1);  // error: p1 is not an owner
    f2(p1);  // call OK, as ever, but f2()'s definition is not
    f1(p2);  // transfers ownership; p2 is now invalid in f3() because f1() will delete it
    f2(p2);  // call OK: a non-owning copy of p2 is passed
}
```

Where do we go from here?

- We can write good C++
 - Too many developers don't
 - Guidelines are not enough, we need guarantees
 - We can (and do) use static analyzers but we need standards
- Alternatives: We can
 1. Change C++
 2. Start using another language
 3. Enforce a variety of guidelines: Profiles



Alternative 1: Fix C++

- How?
 - There are many incompatible ideas
 - => years of delay and chaos
 - A single “cleaned up” language cannot support the wide variety of safety notions
 - A “cleaned up C++” would have to interoperate with “classical C++” code “forever”
 - Serious design constraint
- There are billions of lines of C++
 - Much critical
 - Much high quality
 - Gradual adoption is essential
 - Partial adoption is essential (“safety critical code only”)
- Compatibility
 - C++ was meant to evolve
 - But we can’t break massive amounts of existing code

Alternative 2: Use another language

- The popular alternative with supporters of other languages
 - “safety” is used as an argument
 - Often mixing C and C++ in arguments
 - Often ignoring C++’s strengths
 - Often ignoring the weaknesses of alternatives
 - Often the safety mentioned is just memory safety
 - Often the need for unsafe constructs is unmentioned
 - Often the need to inter-operate with other languages is unmentioned
 - Often the cost of conversion is underestimated
- This is natural
 - Human nature
- Which other language?
 - Must interoperate with C and C++ code “forever”
 - Serious design constraint

Alternative 2: Use another language

- There are many “newer languages”
 - All different
- Each new tool chain require resources and expert users
 - Code from each language needs to communicate with code in all(?) other languages in a system
 - Each language has its own update schedule
- Every new language claims to be simpler, cleaner., safer, and more productive than C++
 - Languages grow significantly in size and complexity over time
- The old languages and systems don’t go away
- Often the claimed superiority of a new language over C++ is in a limited domain
 - Often compared to C/C++, rather than to C++
- Enthusiasts deliver better than average developers
 - Once the number of users of a language increases, their quality and enthusiasm converge towards the industry average

Consider converting a 10-million-line system

- Needing high reliability and high performance
 - i.e., the kind of system that's critical in some way
- A good developer completes N lines of tested production-quality code a day
 - What is N? 5?, 10?, 100?
 - Say – optimistically – 2,000 lines/year
 - Say – optimistically – that a reimplementation (without feature creep) could be 5 million lines
 - Then it would take 500 developers 5 years to complete the new system
 - The old system would have to be maintained for those 5 years say by 50 developers
- What is the loaded salary of a good developer?
 - Employee's compensation plus employer's other cost (e.g., buildings, computers, heating)
 - Say, \$500,000 in the US
- So, the cost would be $550 * 5 * \$500,000 == \sim \$1,400,000,000$
 - Vs. $\sim \$125,000,000$ for normal maintenance and development
 - Roughly \$1B added cost

For a 1-million line system divide by 10

For a 100-million line system multiply by 10

Consider converting a 10-million-line system

- Assumptions
 - US developers
 - developers with relevant experience in the domain and the new language can be found
 - maybe outsourcing could cut cost, but has its own problems and costs
 - The new system would be half the size of the old one (unlikely, but)
 - Better understanding from the start
 - Better language
 - Better tooling
 - The new language/languages can cope with the messy parts of the system
 - No feature creep
 - The new system would work and be delivered on time
 - Large projects often have time and cost overruns

Consider converting a 10-million-line system

- There are people for whom \$1B or \$100M are not scary numbers
- There are people who consider 10M line systems medium sized
- I consider these numbers an argument for an incremental and evolutionary approach.
 - Obviously, that could be in C++ or in a combination of C++ and other languages
 - Either way, C++ will play a major role for decades to come
 - We can and must improve C++ – as ever
- ***A complex system that works is invariably found to have evolved from a simple system that worked.*** – Gall's law

Alternative 3: C++ Profiles

- How to ***guarantee*** safety?
 - For a variety of definitions of safety
 - For a variety of users
 - Mostly statically
- Guidelines are not enough
 - “being careful” is not verifiable
 - The source code doesn’t directly express what’s to be guaranteed
 - Tool chains are complex and vary among systems
- How can we gradually improve safety?



Profiles summary

- The meaning of all constructs is defined by the ISO C++ standard
- The most fundamental guarantee offered is complete type-and-resource safety
- Ownership (that is, the obligation to delete/destroy) constitutes a DAG
- A pointer is the **nullptr** pointer, or is valid (i.e., points to an objects)
- A pointer (outside the implementation of abstractions) points to a single object
- There is a way to enforce **nullptr** checking and range checking
- Subscripting is done on abstractions such as **span** and **vector**, not on pointers
- Gradual conversion from older code to modern code offering guarantees is supported
- The set of guarantees is open
- A set of fundamental guarantees are standard
- There are rules for composing code fragments supporting different guarantees
- The set of guarantees assumed by and provided by a unit of code is stated in the code

Many notions of safety

- **Logic errors:** perfectly legal constructs that don't reflect the programmer's intent, such as using `<` where a `<=` or a `>` was intended.
- **Resource leaks:** failing to delete resources (e.g., memory, file handles, and locks) potentially leading to the program grinding to a halt because of lack of available resources.
- **Concurrency errors:** failing to correctly take current activities into account leading to (typically) obscure problems (such as data races and deadlocks).
- **Memory corruption:** for example, through the result of a range error or by accessing and memory through a pointer to an object that no longer exists thereby changing a different object.
- **Type errors:** for example, using the result of an inappropriate cast or accessing a union through a member different from the one through which it was written.
- **Overflows and unanticipated conversions:** For example, an unanticipated wraparound of an unsigned integer loop variable or a narrowing conversion.
- **Timing errors:** for example, delivering a result in 1.2ms to a device supposedly responding to an external event in 1ms.
- **Allocation unpredictability:** for example, ban on free store allocation "after the engine starts."
- **Termination errors:** a library that terminates in case of "unanticipated conditions" being part of a program that is not allowed to unconditionally terminate.

Type-and-resource safety

- Every object is accessed according to the type with which it was defined (type safety)
 - Every object is properly constructed and destroyed (resource safety)
 - Every pointer either points to a valid object or is the **nullptr** (memory safety)
 - Every reference through a pointer is not through the **nullptr** (often a run-time check)
 - Every access through a subscripted pointer is in-range (often a run-time check)
 - That
 - Implies range checking and elimination of dangling pointers (“memory safety”)
 - Is just what C++ requires
 - Is what most programmers have tried to ensure since the dawn of time
 - The rules are more deduced than invented
- Enforcement rules are mutually dependent.
Don't judge individual rules in isolation

What about C/C++?

- Arbitrary C or C++ code is too complex for static analysis
 - Halting problem
 - Dynamic linking
 - Cost of global analysis
 - Direct access to hardware
- Arbitrary C or C++ forces us to deal with too low an abstraction level
 - Hides complexity in messy old-style code
 - Pushes complexity into the applications code => bugs
 - Regression into the 1980s
 - “C/C++” will never be safe – for any definition of “safe”
- We care about performance as well as type-and-resource safety
 - Eventually much higher productivity

Strategy: Safety Profiles

- Our approach is “a cocktail of techniques” not a single neat miracle cure
- Static analysis (maybe in compilers)
 - to verify that no unsafe code is executed
- Coding rules
 - to simplify the code to make industrial-scale static analysis feasible
- Libraries
 - to make such simplified code reasonably easy to write
 - to guarantee run-time checks where needed




Safety Profiles

- Profiles: A coherent set of guarantees
 - Not just a lot of unrelated tests
 - Specified as a set of guarantees, not as a set of specific tests
- A profile: a coherent sets of rules yielding a guarantee
 - Current: bounds, type, memory
 - E.g., type-and-resource-safe, safe-embedded, safe-automotive, safe-medical, performance-games, performance-HPC, EU-government-regulation
 - Must be visible in code
 - To indicate intent
 - To trigger analysis
 - To trigger run-time checks (where needed)



Is this strategy “too novel”?

- “People are afraid of new things.
You should have taken an existing product and put a clock in it.”
– Homer Simpson
- Each individual technique has been tried many times before
 - Succeeded for specific tasks
 - E.g., smart pointers, libraries, static analyzers
 - Failed as general solutions
 - Static analysis – doesn’t scale to complete safety
 - Guidelines/rules – aren’t followed without enforcement
 - Foundation libraries – doesn’t give full access to the machine and system
 - Language subsetting – the most dangerous language features are essential (e.g., subscripting of pointers)
- A combined and coherent approach is necessary
 - Similar to Ada’s safety profiles: https://docs.adacore.com/gnathie_ug-docs/html/gnathie_ug/gnathie_ug/the_predefined_profiles.html#the-predefined-profiles



Profiles: beyond guidelines

- A type-and-resource safety profile can insert range checks
 - E.g., bind to a range-checked **string**, **vector**, **span**, etc.
- A performance profile can eliminate safety checks (after proving them redundant)
 - Allowing us to make safer defaults
- In-code annotations
 - To state the programmer's intention so that they can be verified/enforced
 - To detect mixing of profiles
 - We cannot upgrade billions of lines of code at once
 - We want to use several different profiles in larger programs
 - Mixing profiles is a generalization of the safe/unsafe oversimplification
 - “safe code” calling C code (e.g., the OS), assembler, etc.
 - “unsafe code” passing information (e.g., pointers) to “safe code”

Mixing profiles

- We can use subset/superset properties of profiles
 - The profile of an imported module cannot be stated (it might change, but the build-system knows)
- We can't express the relationships among arbitrary profiles
 - That would be a whole complex language
- Disjoint profiles combine cleanly
 - E.g., type-and-resource and arithmetic
- We mustn't require users to litter their code with many opt-outs of a stated profile
 - Needing/wanting so would be a sign of a bad program design (like the use of many casts)
 - We mustn't require users to mark every call to a different profile
- A single “foreign call” annotation is too coarse
 - We need to know which profile or aspect of a profile that's violated

Mixing profiles

Work in progress

- Still under development
 - Suggestions welcome
- Module-based controls
 - `export My_module [[provide(memory_safety)]];` *// enforce memory_safety for My_module*
 - `import std [[enable(memory_safety)]];` *// enforce memory_safety for uses of std*
 - `Import Mod [[suppress(type_safety)]];` *// don't check type_safety for uses of Mod*
- In-code controls
 - `[[suppress(type_safety)]] X` *// suppress type_safety check for declaration or scope X*
 - `[[enforce(type_safety)]] X` *// enforce type_safety check for declaration or scope X*
- Suggested initial standard profiles
 - **type_safety**: no type-or-resource violations.
 - **range**: no pointer arithmetic; no **nullptr** dereference, **span** and **vector** range **throw** or terminate on violations.
 - **arithmetic**: no overflow, no narrowing conversions, no implicit signed/unsigned conversions.

Summary

Work in progress


- `[[enforce(P)]]`
- `[[provide(P)]]`
- `[[enable(P)]]`
- `[[suppress(P)]]`
- `[[profile(P) = enforce(P1,P2)]]`
- Standard profiles
 - Type-safety
 - Arithmetic
 - Range
- Control of response to run-time violations
- **owner**
- **not_null**
- **not_end()**
- `[[not-invalidating]]`
- `[[uninitialized]]`
- **dynarray** – a vector where the size *s* fixed at construction
- Constraints on UB

Design Alternatives for
Type-and-Resource Safe C++
P2687

Safety Profiles:
Type-and-resource Safe
programming in ISO Standard C++
P2816


“Are we there yet?”

- We have come a long way
 - From “classic C”
 - From “C with Classes”
 - From C++11
- We can write type-and-resource safe C++
 - Use contemporary C++
 - Avoid C-style and 1980s style C++
- C++ Core Guidelines
 - Directions, rules, and some enforcement
 - But not standardized
 - Uneven enforcement across implementations
- We need to standardize “Profiles”
 - And get implementations deployed
 - Finally reach type-and-resource safe C++!



How can you help

- We need help to refine “profiles”
 - “memory safety” is not sufficient
- Which profiles do we need?
 - Which needs to be standardized?
 - Which first?
- How do we best formalize a profile?
 - By guarantees, not lists of rules.
 - We need work on key examples
- What can we get ready soon?
 - Faster than WG21 can move
 - With current compilers and tool chains
 - Subsets of “ideal profiles”
 - Industry?
- What library components do we need to simplify use?
- See <https://github.com/BjarneStroustrup/profiles>



References

- M. Wong, H. Hinnant, R. Orr, B. Stroustrup, D. Vandevoorde: [DG Opinion on Safety for ISO C++](#). P2759R1. 2023-01-22.
- B. Stroustrup: [A call to action: Think seriously about “safety”; then do something sensible about it](#). P2739R0. 2022-12-6.
- B. Stroustrup and G. Dos Reis: [Safety Profiles: Type-and-resource Safe programming in ISO Standard C++](#). P2816
- B. Stroustrup and G. Dos Reis: [Design Alternatives for Type-and-Resource Safe C++](#). P2687R0. 2022-20-15
- B. Stroustrup: [Type-and-resource safety in modern C++](#). P2410r0. 2021-07-12.
- B. Stroustrup, H. Sutter, and G. Dos Reis: [A brief introduction to C++'s model for type- and resource-safety](#). Isocpp.org. October 2015. Revised December 2015.
- B. Stroustrup: [Writing Good C++14](#) CppCon 2015.
- [The C++ Core Guidelines](#)
- [The Core Guidelines Support Library \(GSL\)](#)
- H. Sutter: [Lifetime safety: Preventing common dangling](#). P1179R1. 2019-11-22.
- [A Microsoft guide to using the Core Guidelines static analyzer in Visual Studio](#).
- T. Ramananandro, G. Dos Reis, and X. Leroy: [A mechanized semantics for C++ object construction and destruction, with applications to resource management](#). ACM/SIGPLAN Notices 2012/01/18.
- B. Stroustrup: [Thriving in a crowded and changing world: C++ 2006-2020](#). ACM/SIGPLAN History of Programming Languages conference, HOPL-IV. London. June 2020.
- B. Stroustrup: [C++ -- an Invisible Foundation of Everything](#). ACCU Overload No 161. Feb'21.

We didn't start yesterday