

+ 23

Shared Libraries in Windows, in Linux, and yes - in C++

OFEK SHILON



20
23

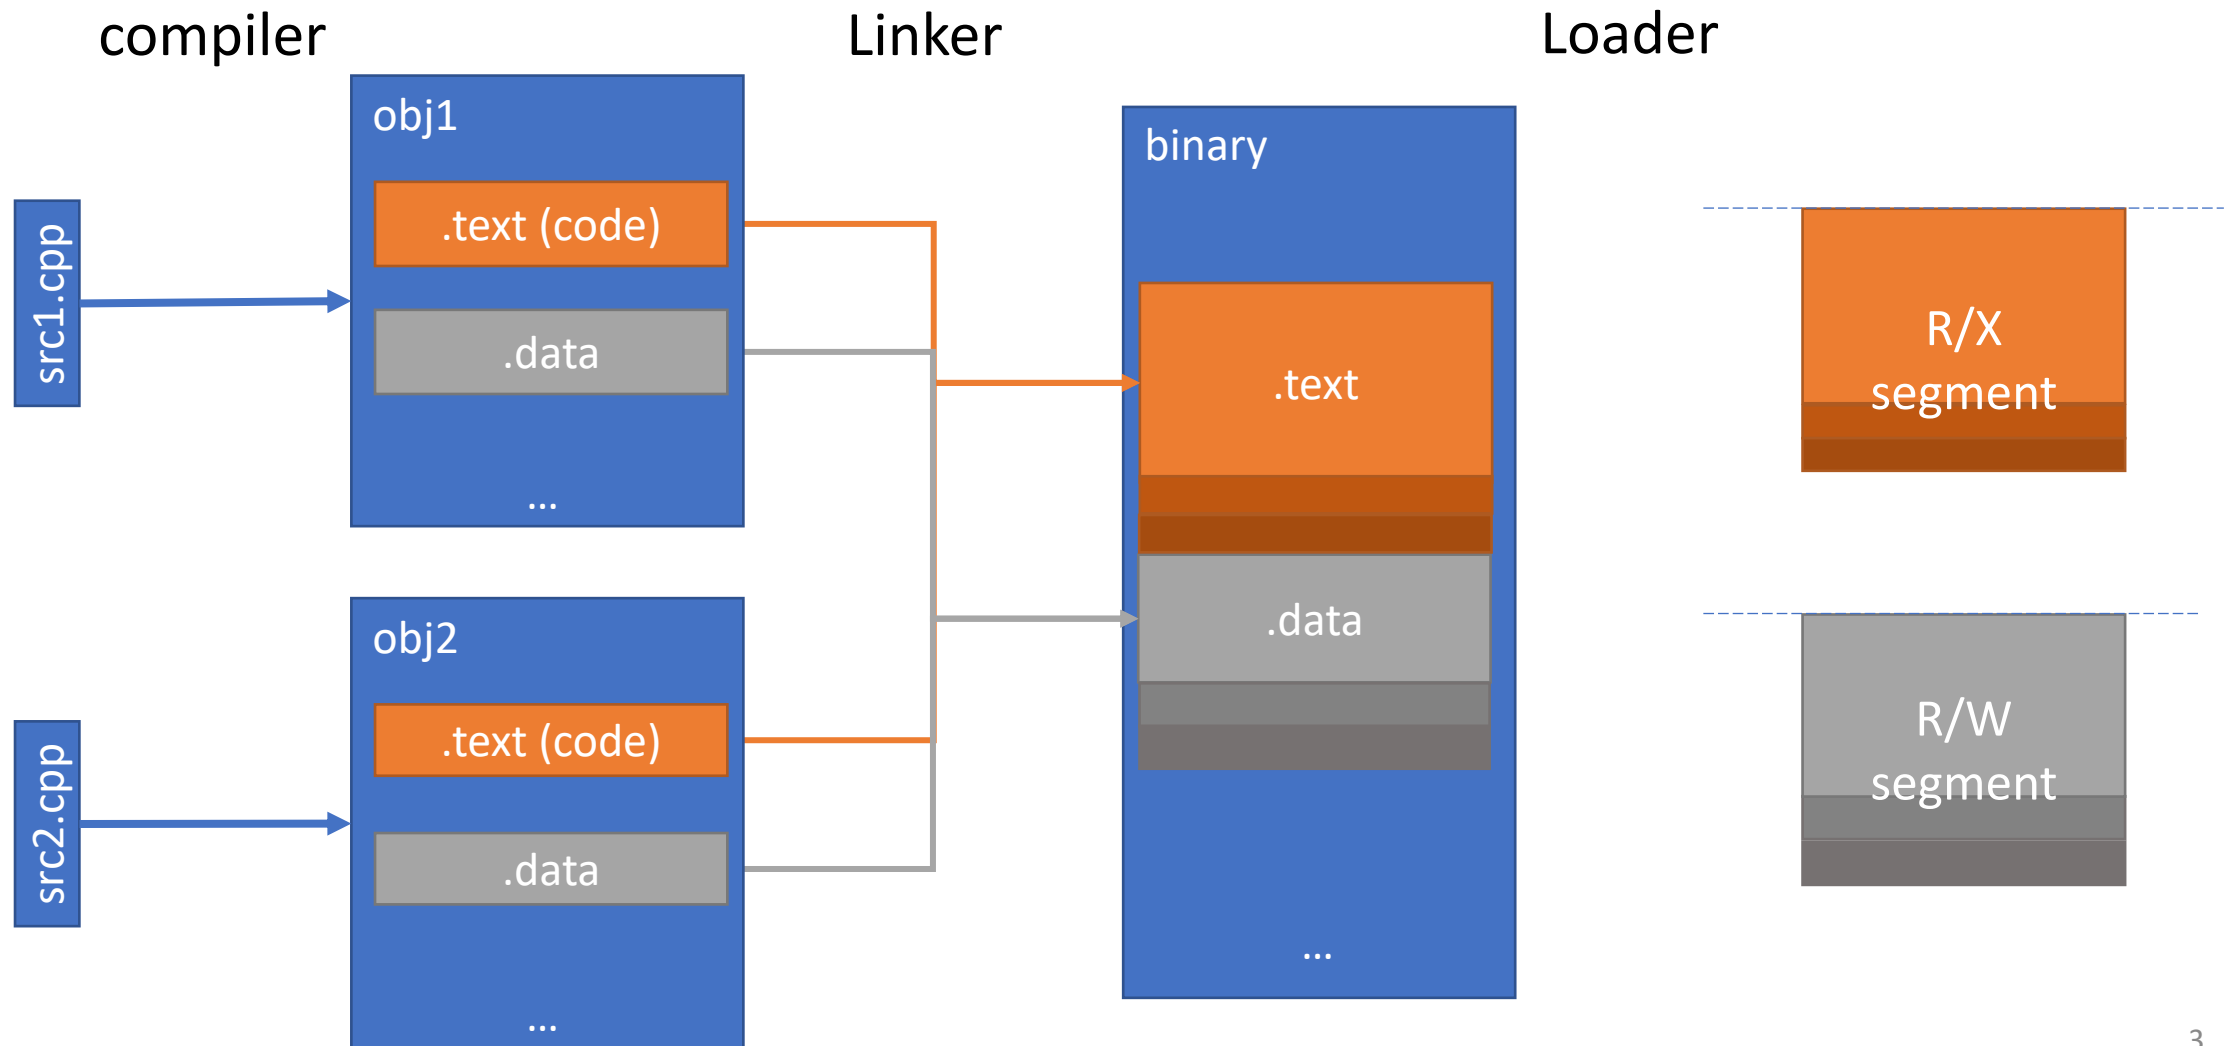


October 01 - 06

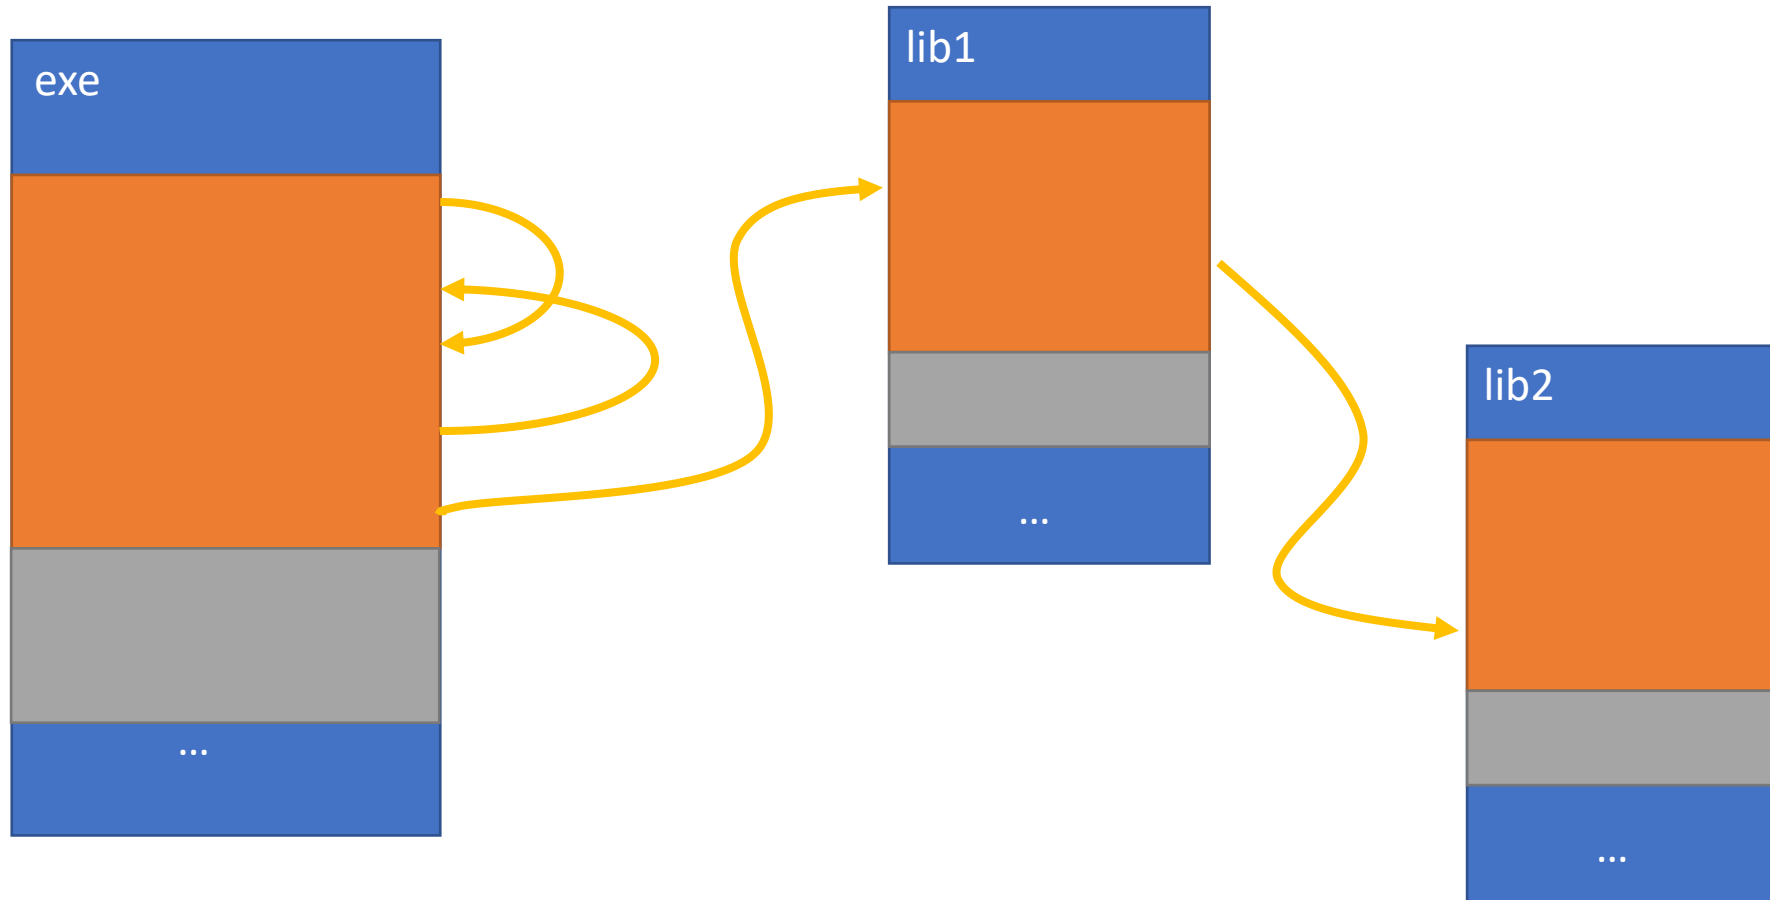
Terminology

- Shared Library:
 - Shared Object, Dynamic Object, Dynamic Shared Object (DSO), Dynamic Load Library (DLL), Dynamic Shared Library
- Binary:
 - Executable / Shared Library , Component, Module
- Symbol:
 - Function / Global variable
- “Linux”:
 - Unix-like systems
 - Mostly MacOS too (1 important distinction to come)

Intro to Linking in 3 slides, #1



Intro to Linking in 3 slides, #2



Intro to Linking in 3 slides, #3

code

```
foo();
```

```
...
```

Intro to Linking in 3 slides, #3

0x1000:

```
code
    call 0x00000000
    ...
```

```
.reloc
    "Find `foo` and write
    its address at 0x1001"
```

Intro to Linking in 3 slides, #3

Cross-binary calls are typically indirect – carry a virtual-call overhead

code

```
call [0x2000]  
...  
call [0x2000]  
...  
call [0x2000]  
...
```

0x2000:

Small lie 1

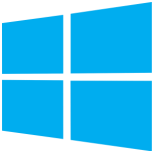
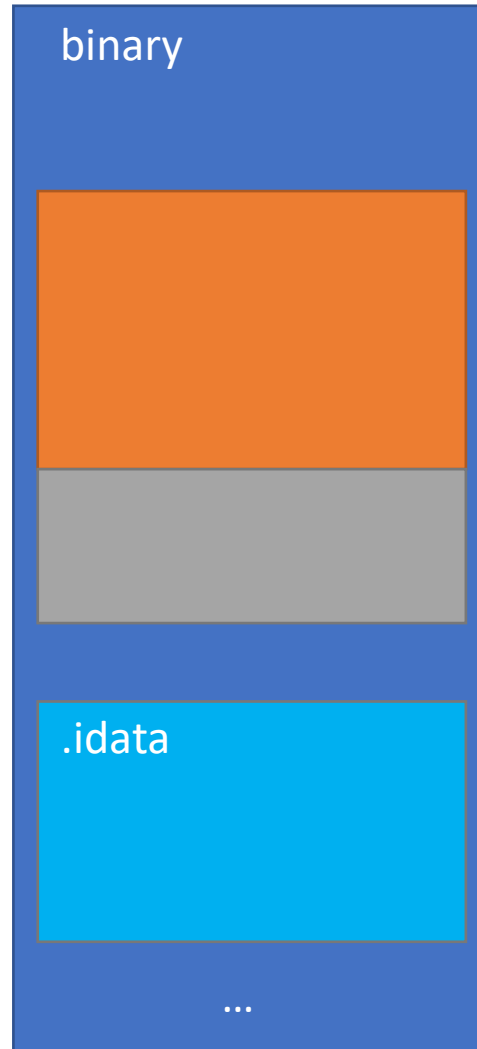
find `foo` and write
its address at 0x2000

Address-list section:
“IAT” in Windows,
“GOT” in Linux

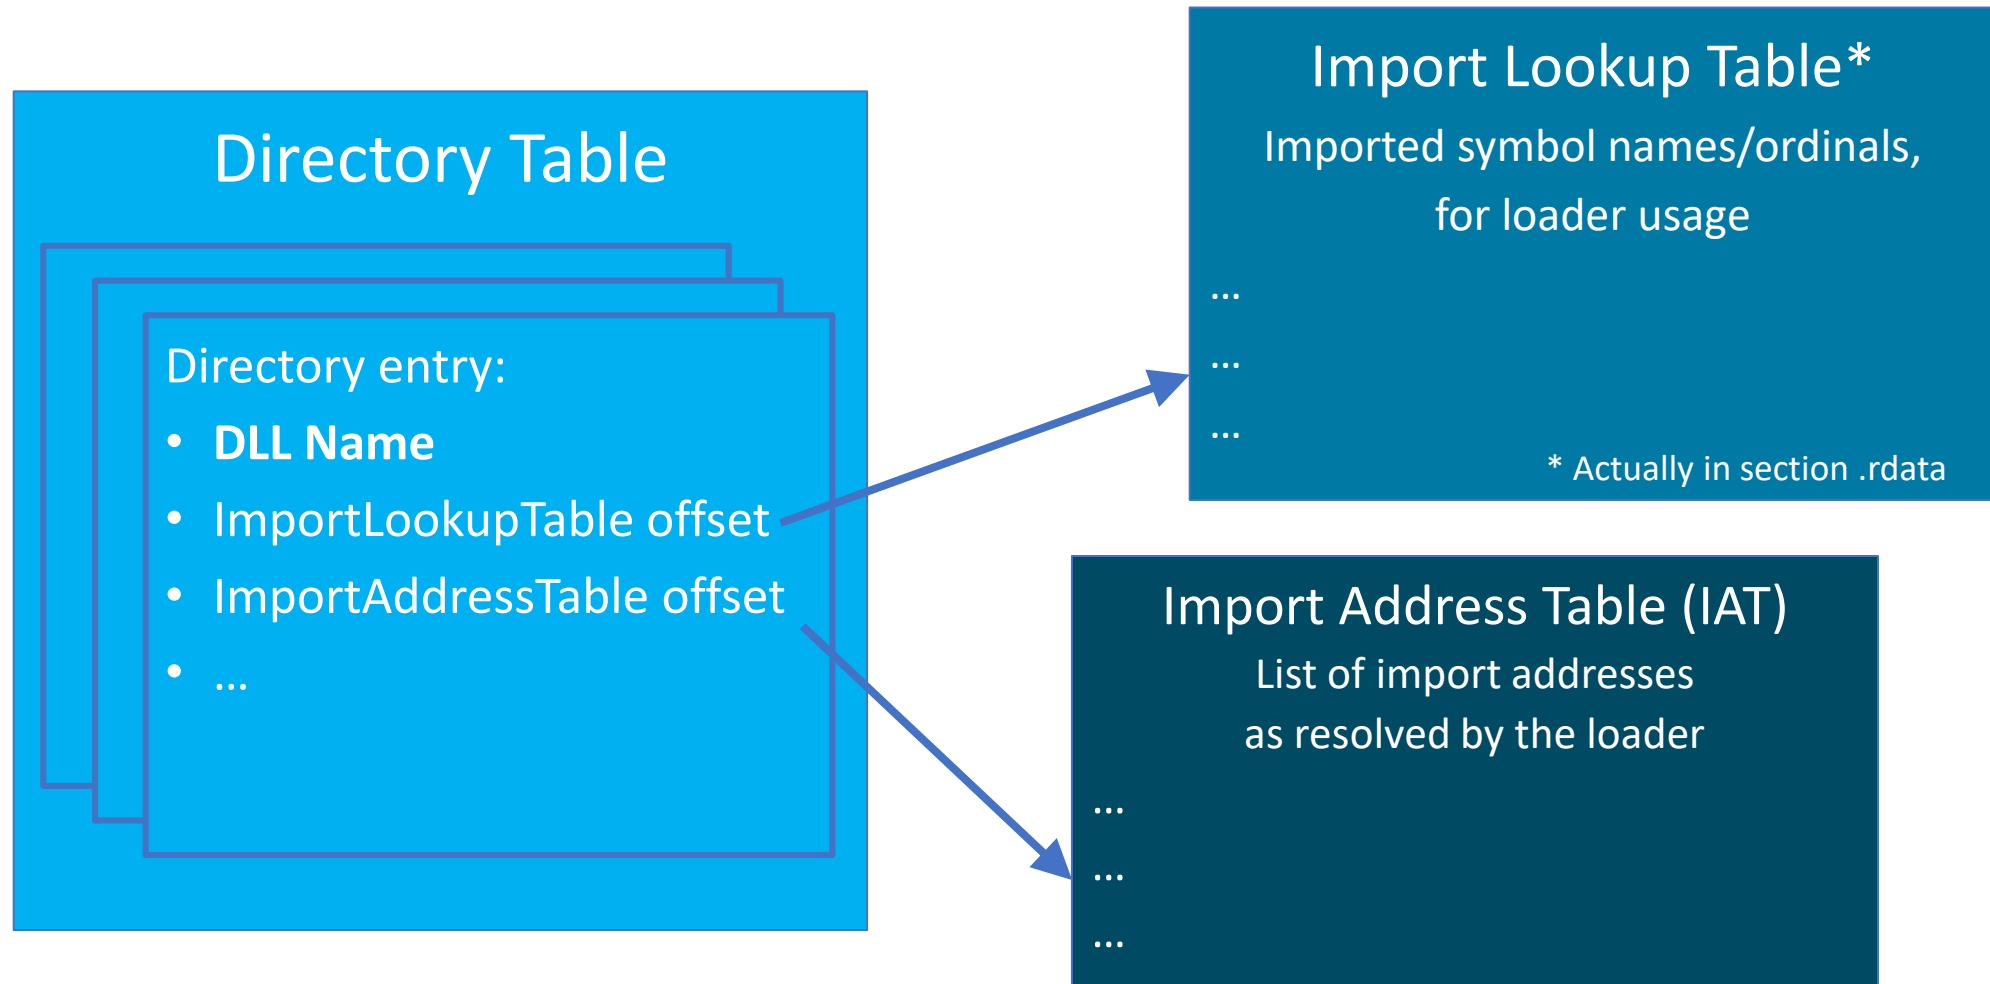
Windows



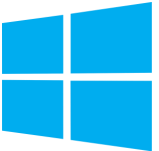
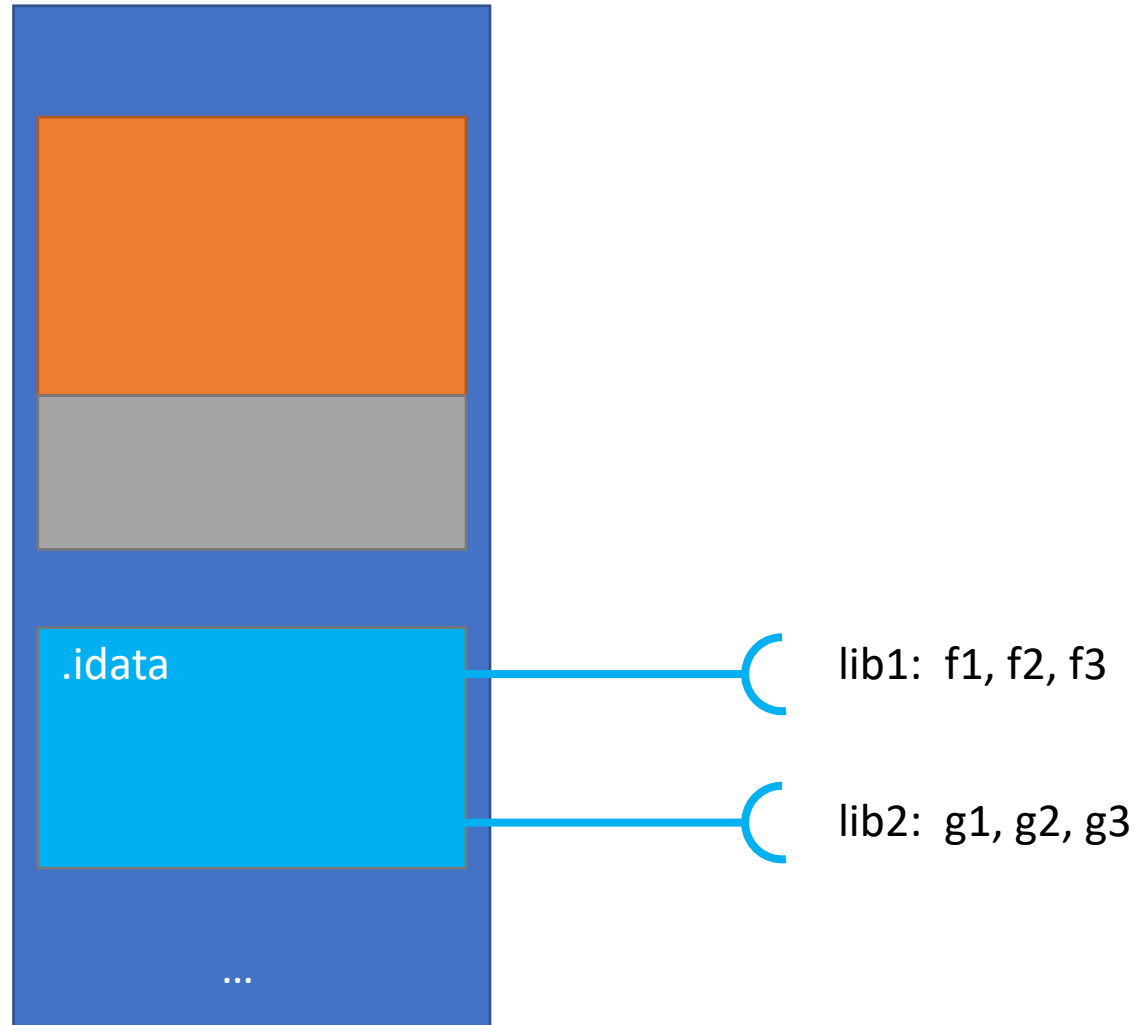
Import data section



.idata contents



Windows Schematic Interface



Linux



Linux import sections

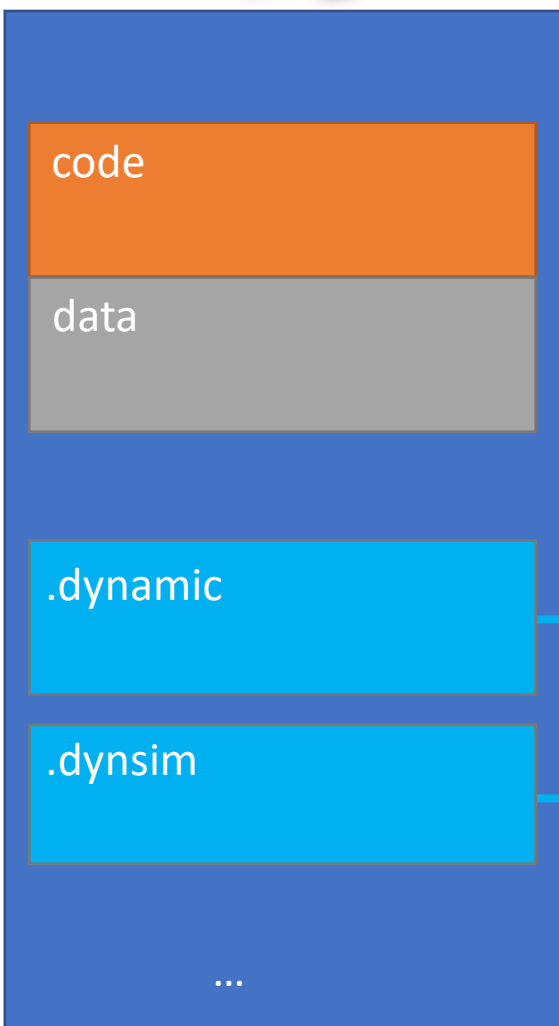
- .dynamic /.dynsym : separate buckets of lib names and symbol names

```
Dynamic section at offset 0x21a58 contains 28 entries:
  Tag                Type                Name/Value
  0x0000000000000001 (NEEDED)           Shared library: [libselinux.so.1]
  0x0000000000000001 (NEEDED)           Shared library: [libc.so.6]
```

```
Symbol table '.dynsym' contains 139 entries:
```

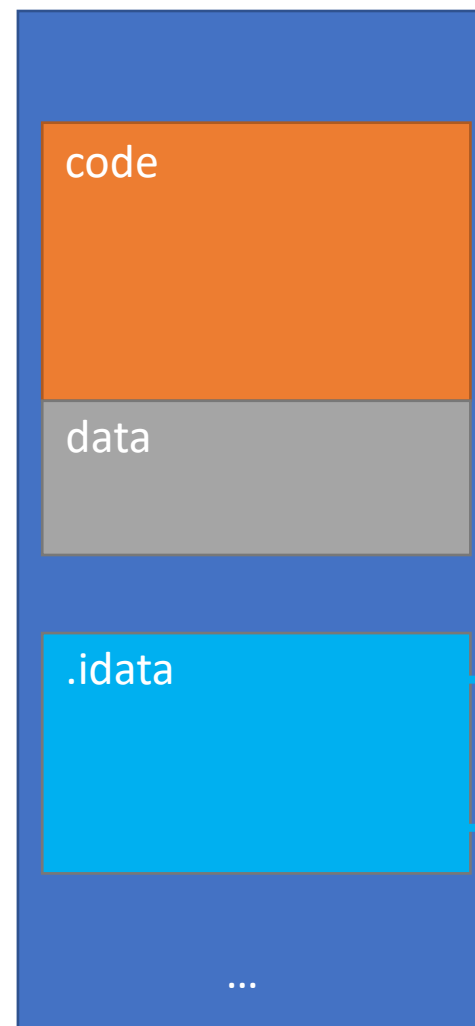
Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	__ctype_toupper_loc@GLIBC_2.3 (2)
2:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	getenv@GLIBC_2.2.5 (3)
3:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	sigprocmask@GLIBC_2.2.5 (3)
4:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	__snprintf_chk@GLIBC_2.3.4 (4)





lib1, lib2

f1, f2, f3,
g1,g2,g3



lib1: f1, f2, f3

lib2: g1, g2, g3

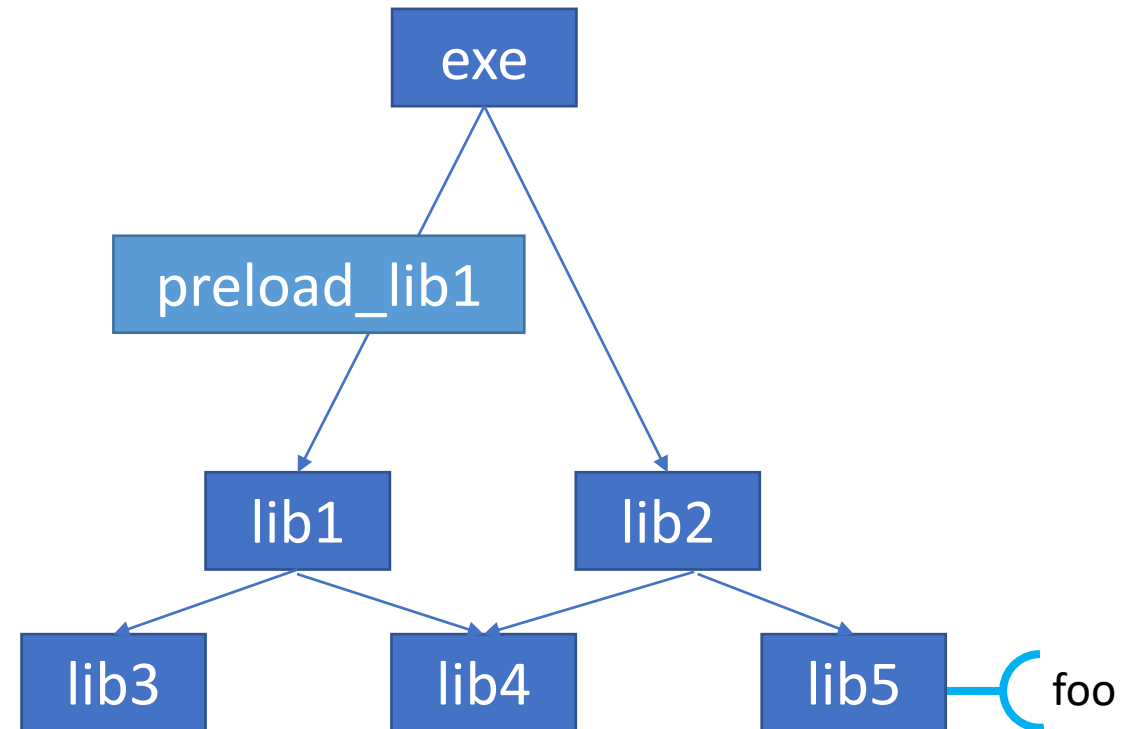
Interposition

- *Overriding a symbol in one binary from another.*
- Fundamental ABI design pillar!
- Alleged motivation:
 - Mimic the behavior of static libs: first definition shadows the later ones.
 - Not true: in static libs duplicate definitions can cause ODR violation errors.
- Speculated motivation:
 - In the elden days the canonical example of shared-lib was **libc**.
Users may reasonably wish to override implementations.



Library Search Order

- Breadth-first
- **exe before current lib** (By default)
- To search current lib first:
 - `-Bsymbolic*`,
 - `--dynamic-list*`
- **LD_PRELOAD:**
 - After exe, before any lib



C++ Implication #1: Can a shared-library symbol be overridden from an executable?

- Windows:
 - No.
- Linux:
 - Yes.
- Mac:
 - Yes, but requires non-default linker switches (eg `-flat_namespace`)



C++: new

[replacement.functions]: A C++ program may provide the definition for any of the following dynamic memory allocation function signatures declared in header `<new>` :

- `operator new(std::size_t)`
- `operator new(std::size_t, std::align_val_t)`
- ...

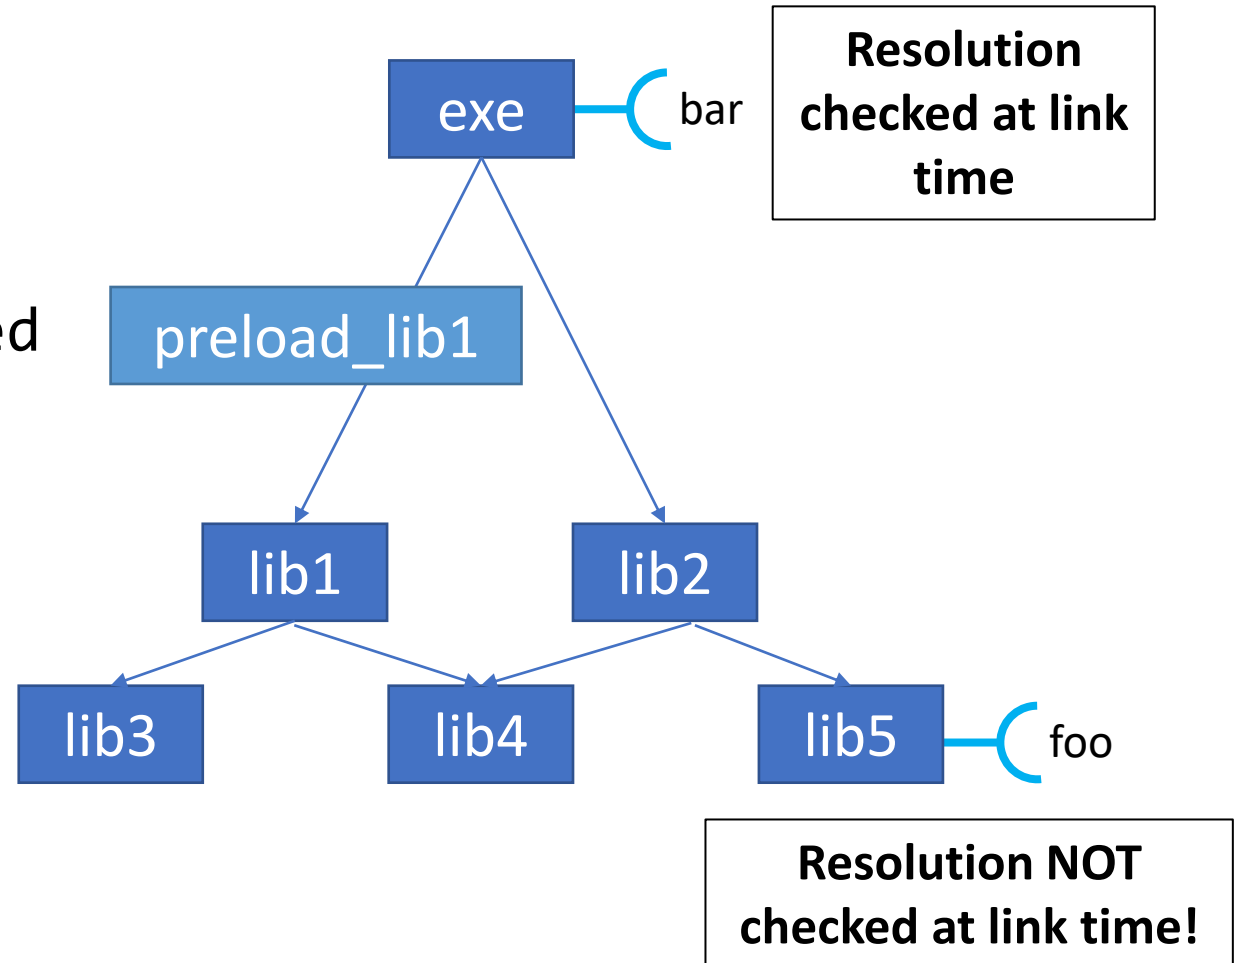
The program's definitions are used instead of the default versions supplied by the implementation ...

- *Windows doesn't (can't) do that.*
- Suggestion: drop this clause, add instead library hooks like `"set_new_override()"`, similar to the existing `set_terminate()` and others.
 - Credit: Thiago Macieira, <https://lists.isocpp.org/std-proposals/2023/07/7240.php>



Symbol Resolution Time

- Default:
 - `--allow-shlib-undefined`
- Can be controlled with –
 - `--no-allow-shlib-undefined` (on the exe)
 - Operates recursively on ld, not on gold/lld.
 - `-z defs` (on the shlib)
 - `--no-undefined` (on the shlib)



C++ Implication #2: How to form a process-wide singleton?

- Singleton: a single object usable by all the code in the process, from all binaries.
- Windows:
 - Usual singleton design patterns would create a *per-binary* singleton
 - Export the singleton variable from a shared-lib, link all your binaries against it.
- Linux:
 - Just put it in the executable



C++ Implication #3: Can you have circular library dependencies?

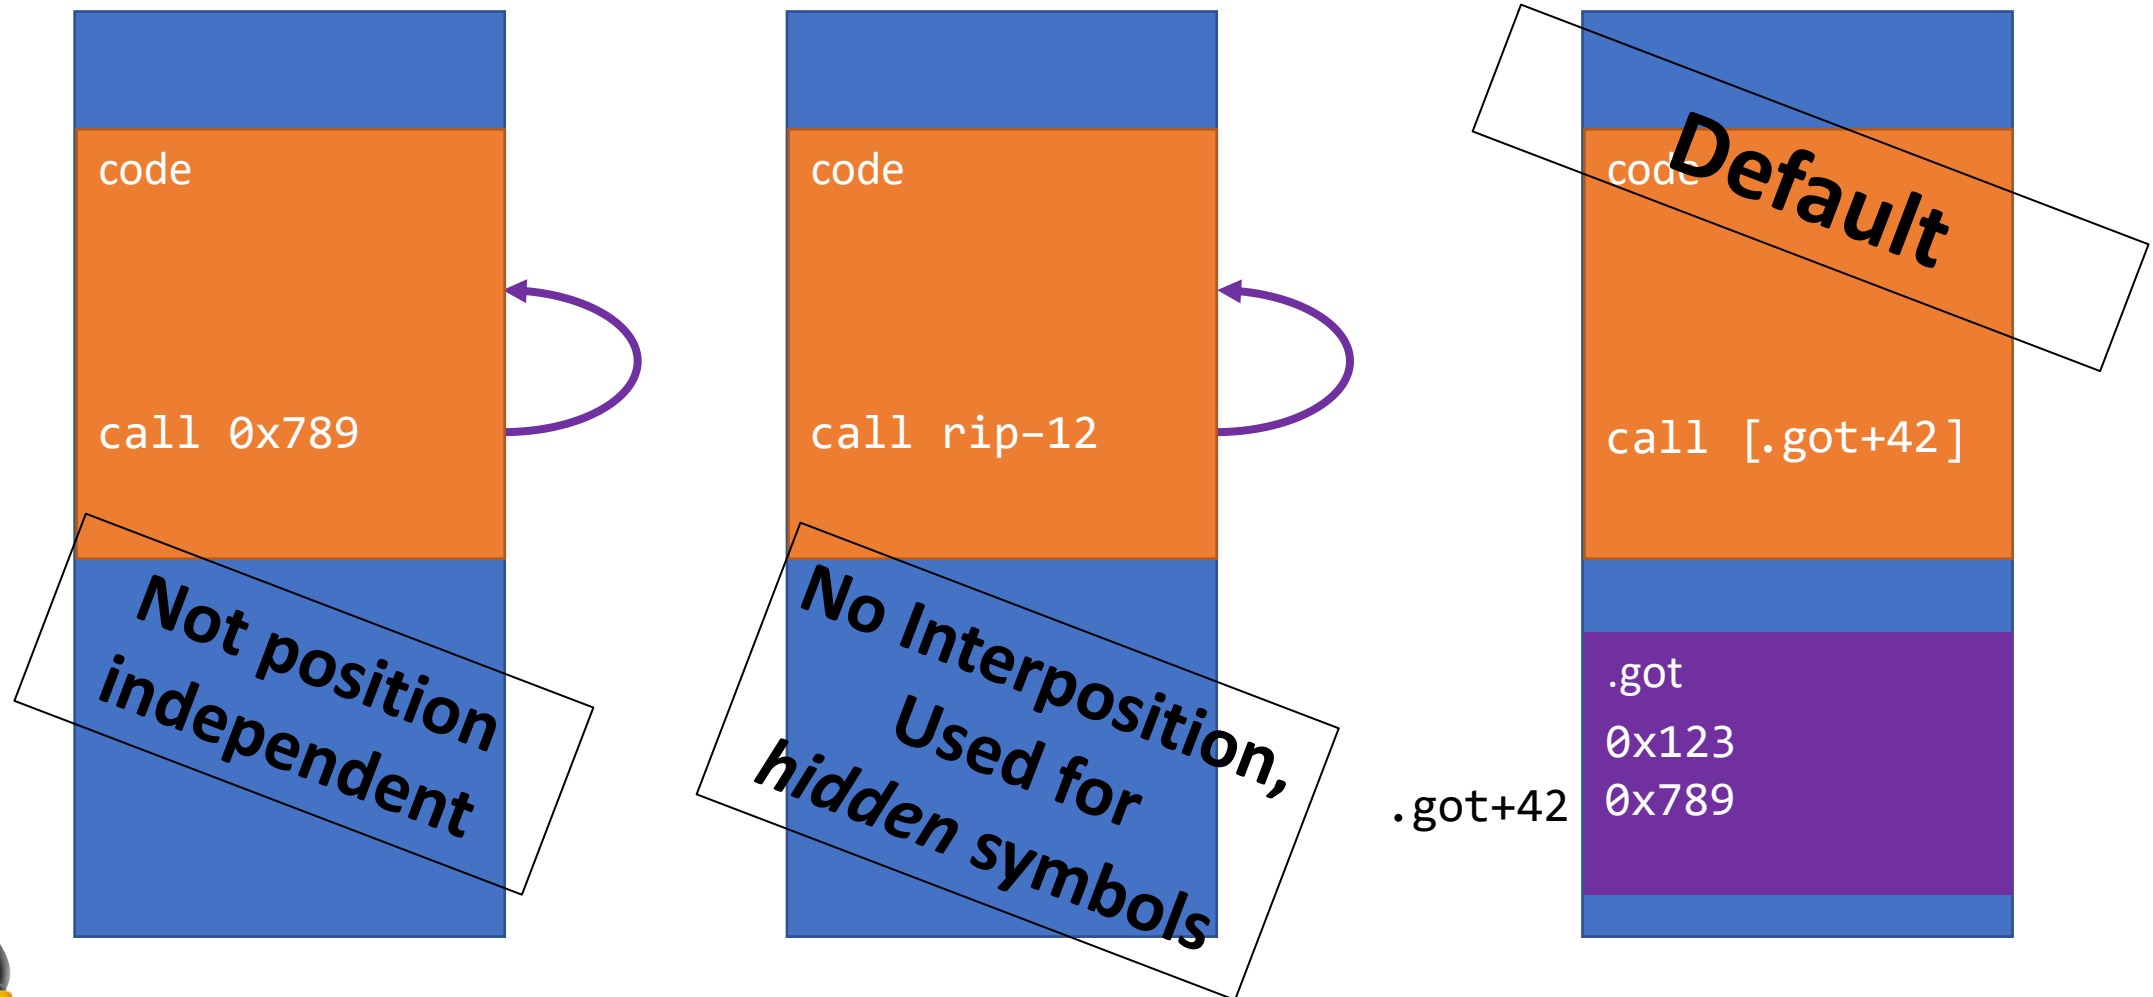
- Linux:
 - Yes
- Windows
 - No.
 - Well, you'd have to hack hard.
- The Linux design provides some flexibility, but ..
 - “This [allowed-shlib-undefined] is an unfortunate default for -shared. Changing it may be disruptive today. Mach-O and PE/COFF have many problems but this may be a place where they got right.”



Fangrui Song, LLD maintainer

<https://maskray.me/blog/2021-06-13-dependency-related-linker-options>

Position Independent Code: GOT



Position Independent Code: GOT

- *For all (default visibility) calls from a shared-lib!*

- This code -

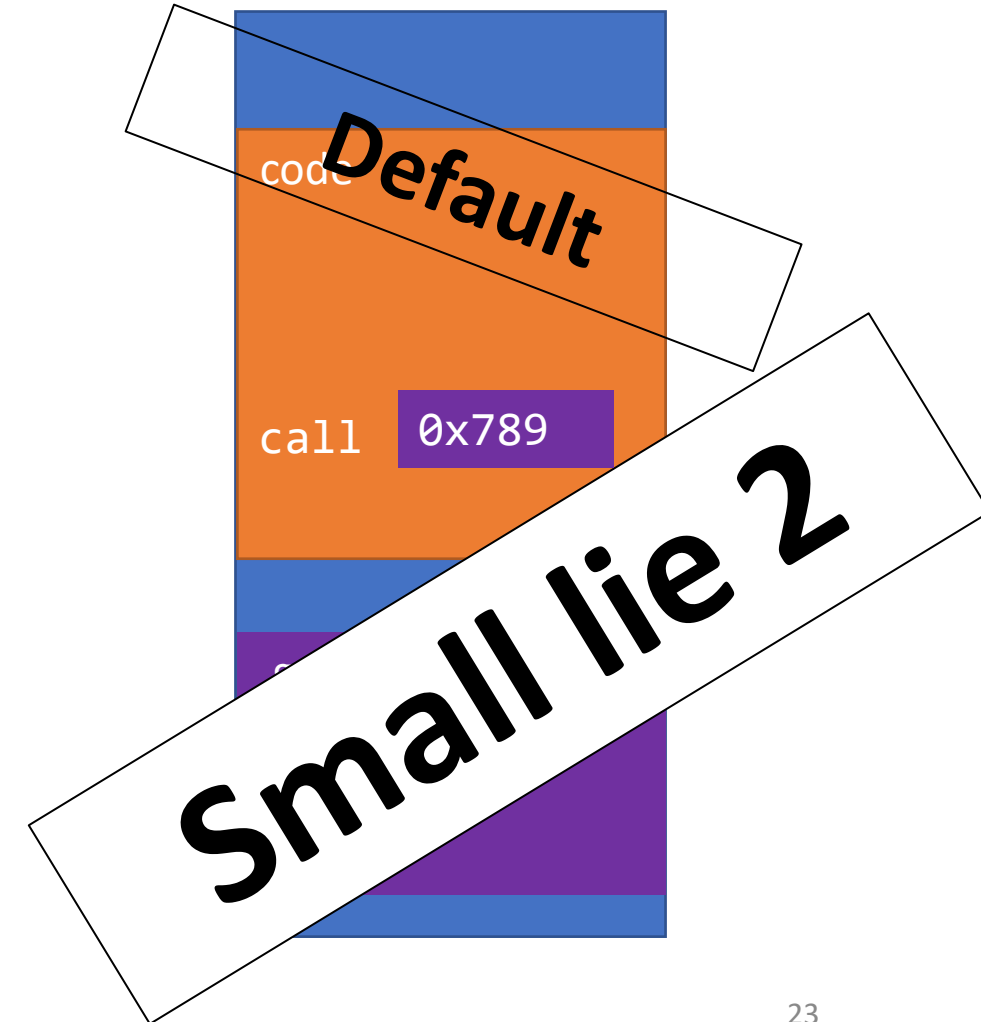
```
void f() { ... }  
void g() { f(); }
```

Will (probably) be inlined if put in an exe,
*but not if put in a shared library.**

- All inter-procedural optimizations are inhibited *



- * clang does it anyway...



Position Independent Code - switches

- **ALL** shared lib code must be position independent
 - Yet -fPIC is optional , not even implicit for -shared.
 - If you try to link a shared lib from obj files not built with fPIC *and* using global vars:

```
error: relocation R_X86_64_PC32 against symbol `global' can not be used when making a shared object; recompile with -fPIC
```
- -fpic vs -fpie:
 - EXE symbols are relocatable but NOT interposable. Interprocedural optimizations apply.



Lazy Binding

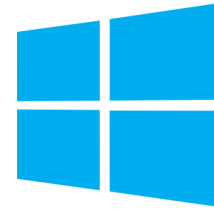
(a.k.a Delayed Loading)

Lazy Bind by Default?



Yes. *

Controllable with linker switches
-no-plt, -z now,
function attribute `noplt`
or env var LD_BIND_NOW



No.

Controllable with linker switch
/DELAYLOAD:<your_dll.dll>



.got+42

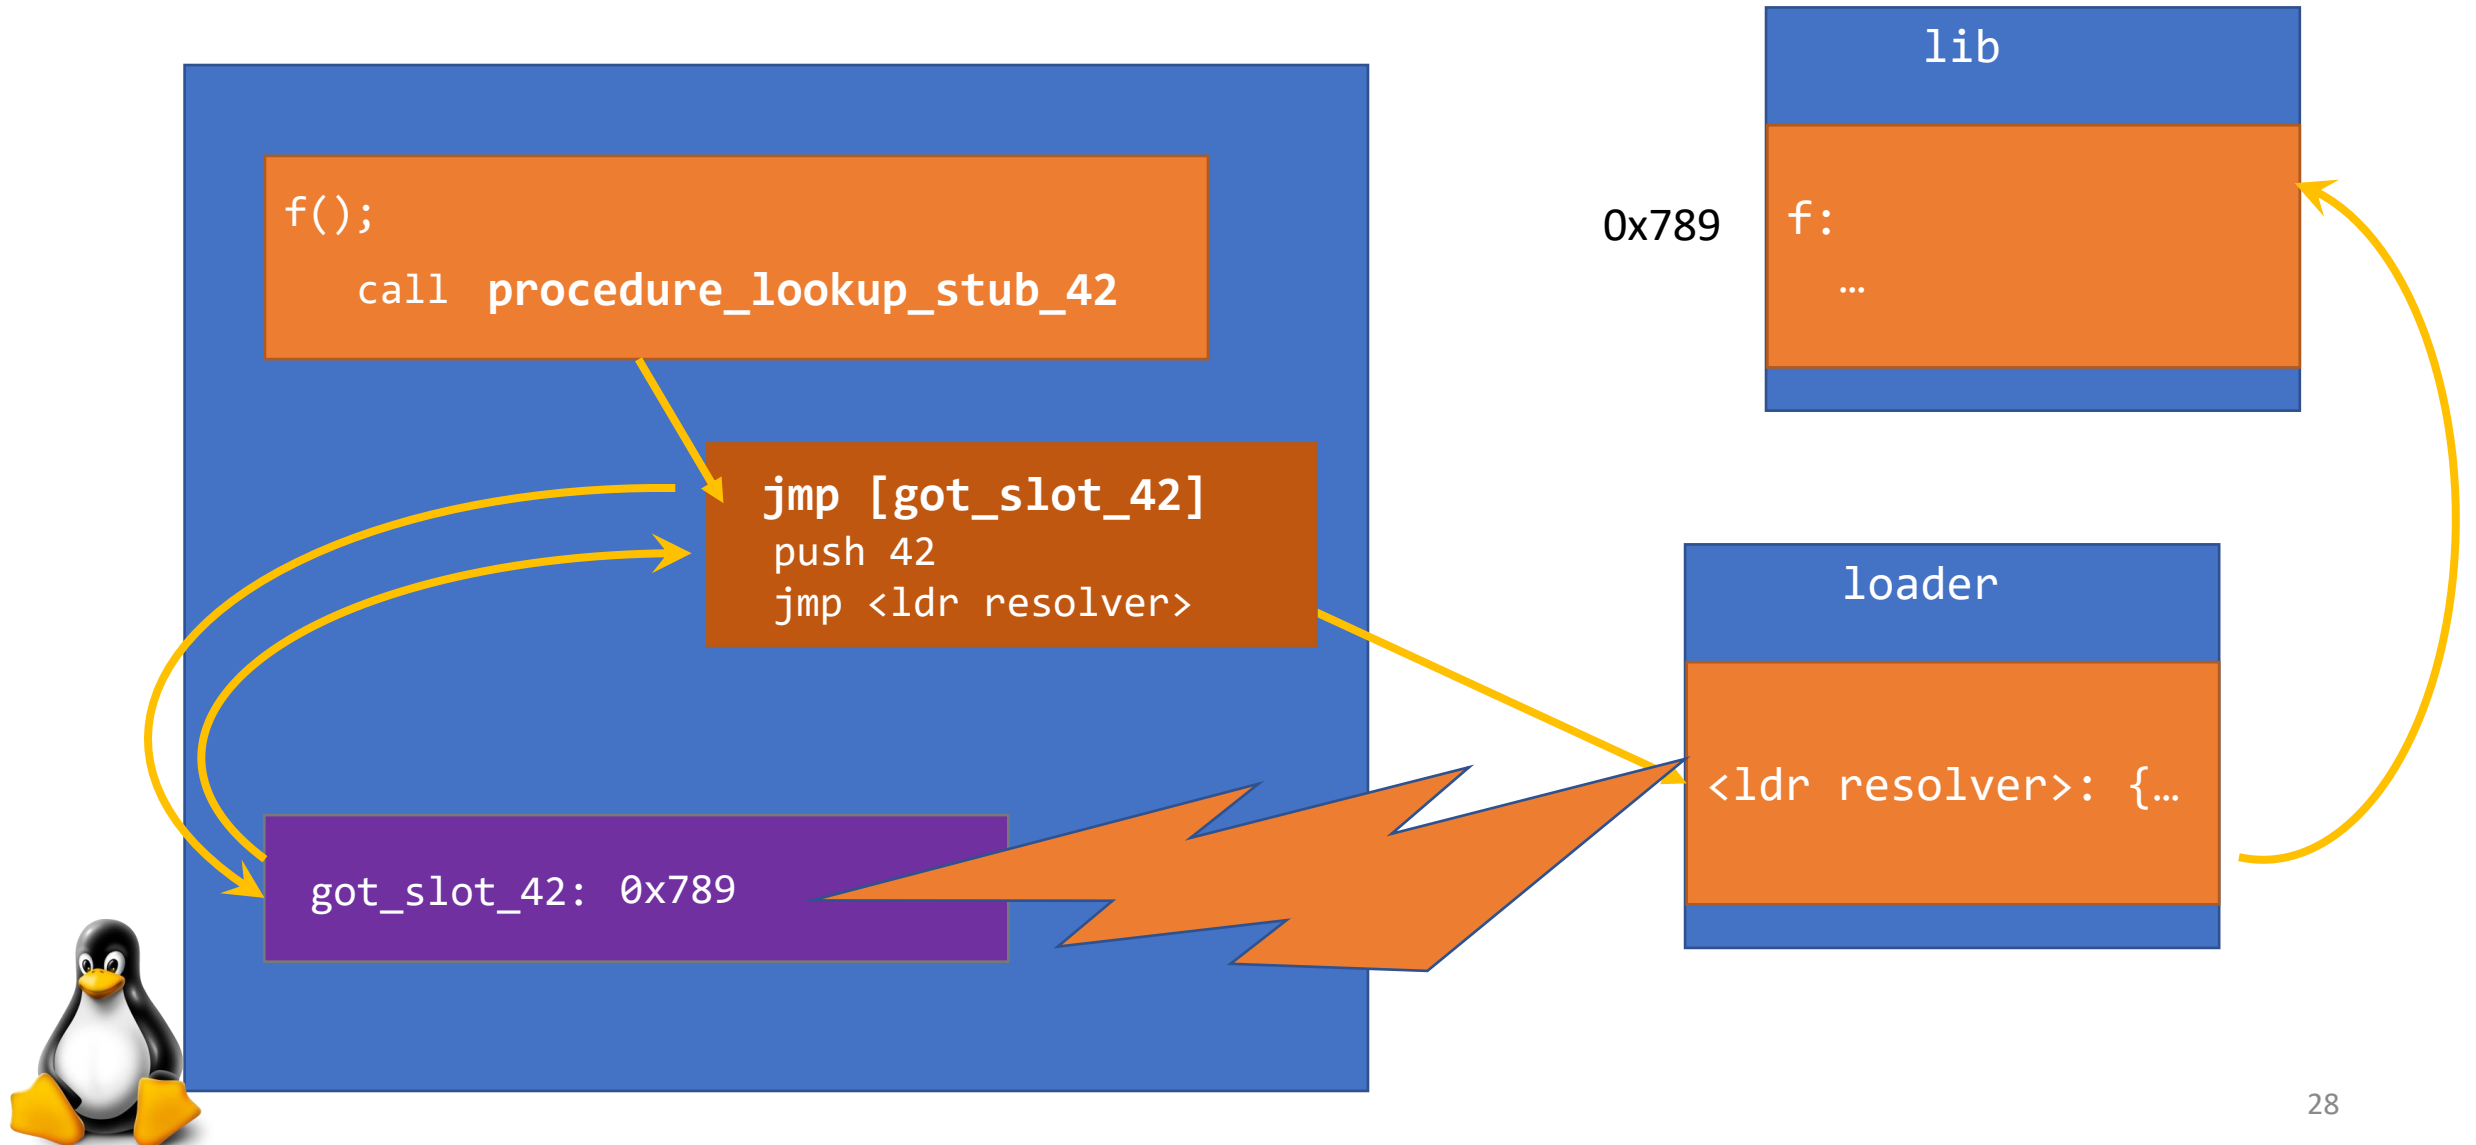


Small lie 1

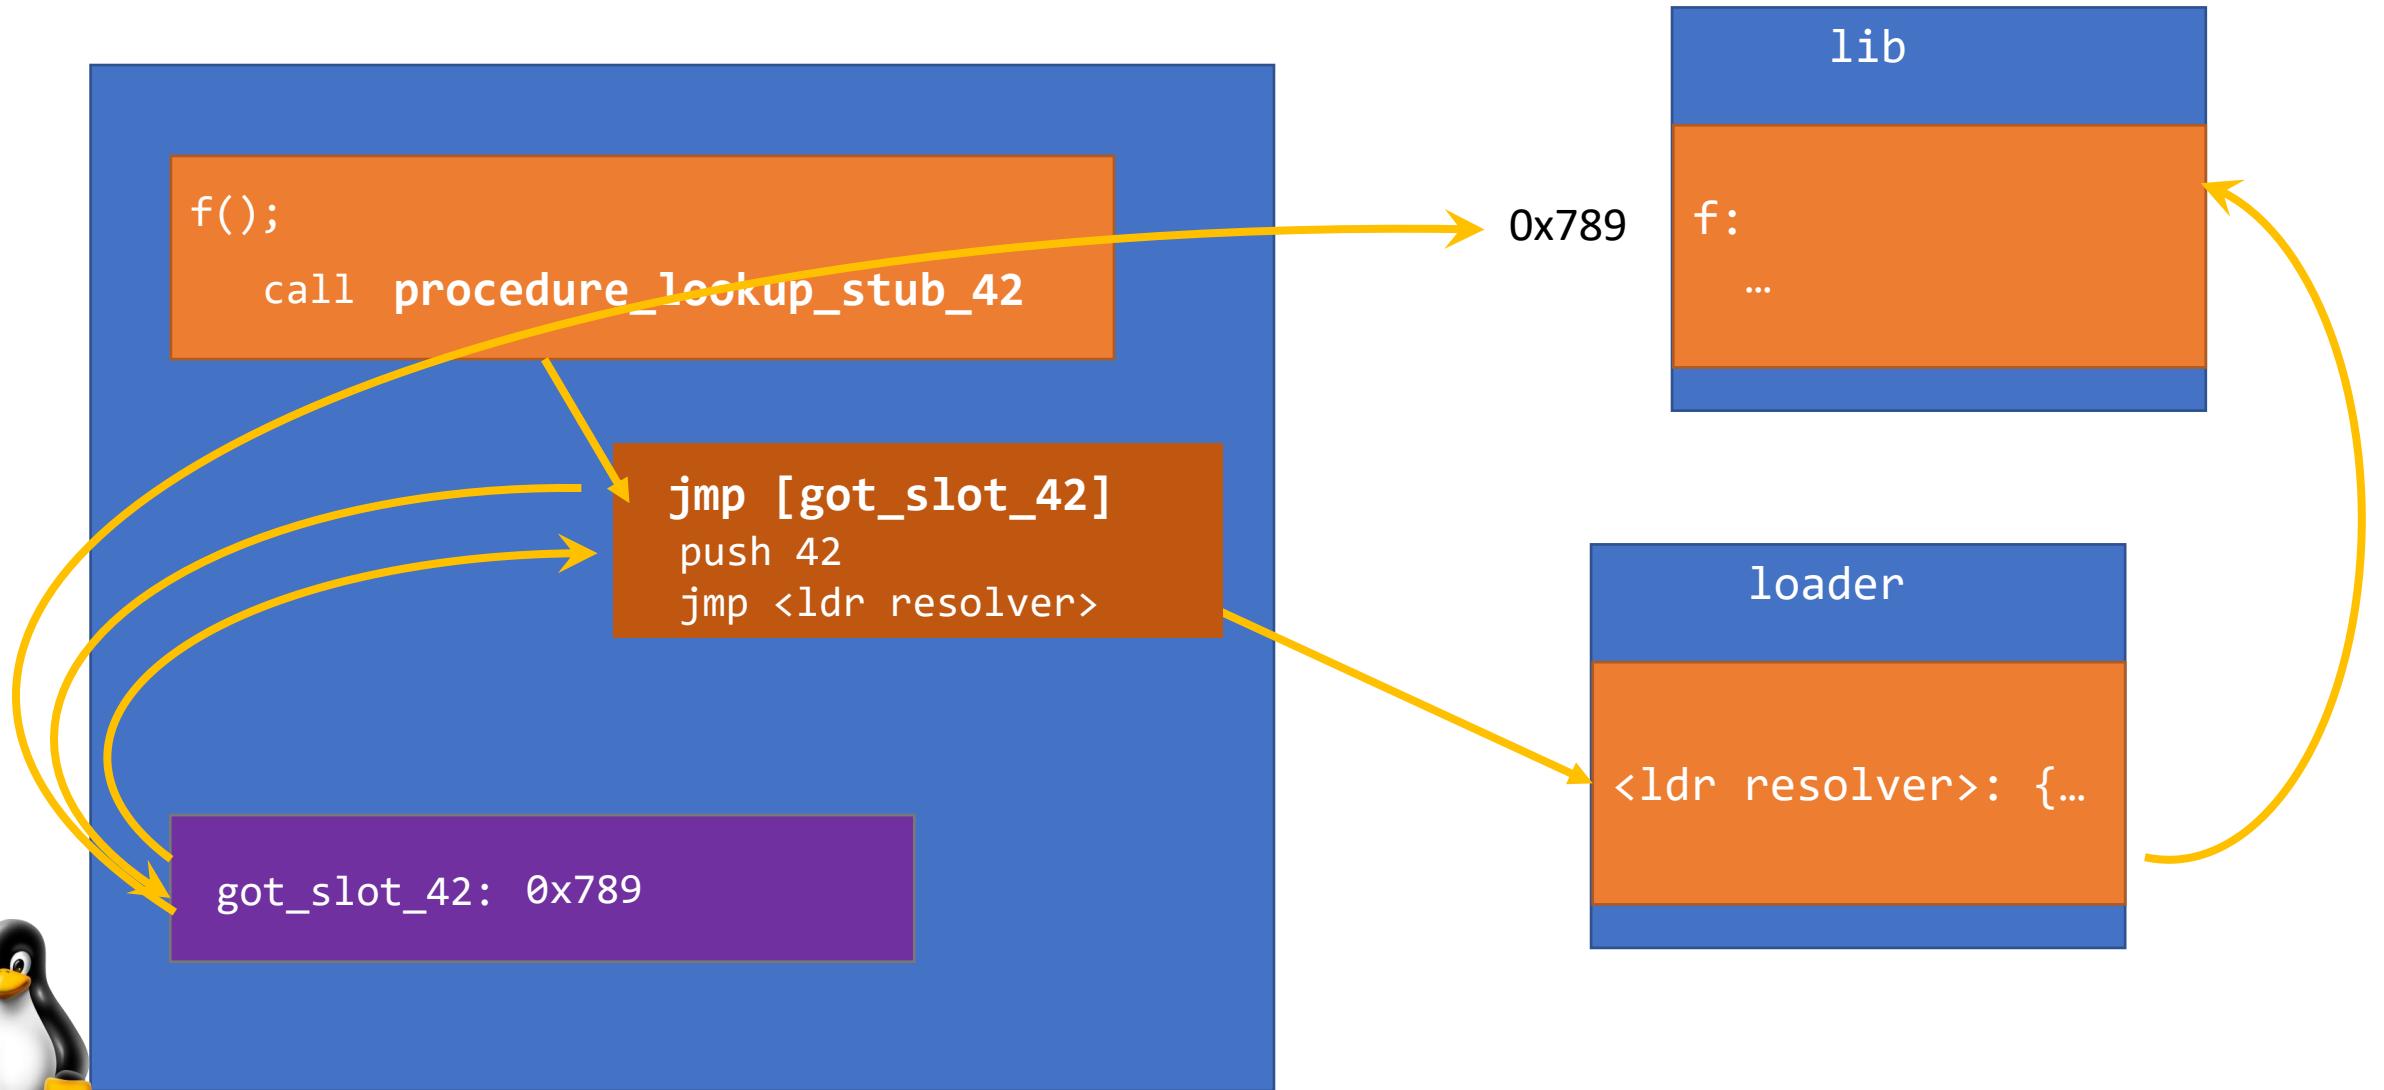
0x789



Lazy Binding Design



Lazy Binding Design



Lazy Binding Design

```
f();
```

```
call procedure
```

```
...
```

```
jmp [got_slot_41]
```

```
push 42
```

```
jmp <ldr resolver>
```

```
jmp [got_slot_42]
```

```
push 42
```

```
jmp <ldr resolver>
```

```
jmp [got_slot_43]
```

```
push 42
```

```
jmp <
```

```
got_slot_42: 0x78 ...
```

Procedure Link Table

Almost Full Truth



PLT: some extra details

- The 1st PLT entry is special and used in all PLT calls, add a module descriptor.
- **.got** section used for eager binding (usually global vars),
.got.plt used for lazy binding (functions).
- On modern x86/x64 compilers you'd see *another* level of indirection, through the section **.plt.sec**
 - Added because support for intel security features CET/MPX* didn't fit in 16-byte **.plt** entry scheme.
(* in a nutshell: branches marked with ``bnd`` prefix must land on an `endbr64` instruction)
 - Was another section really necessary? Probably a design error...
<https://reviews.llvm.org/D59780#1468080>



C++ Implication #4: Comparing Func Ptrs

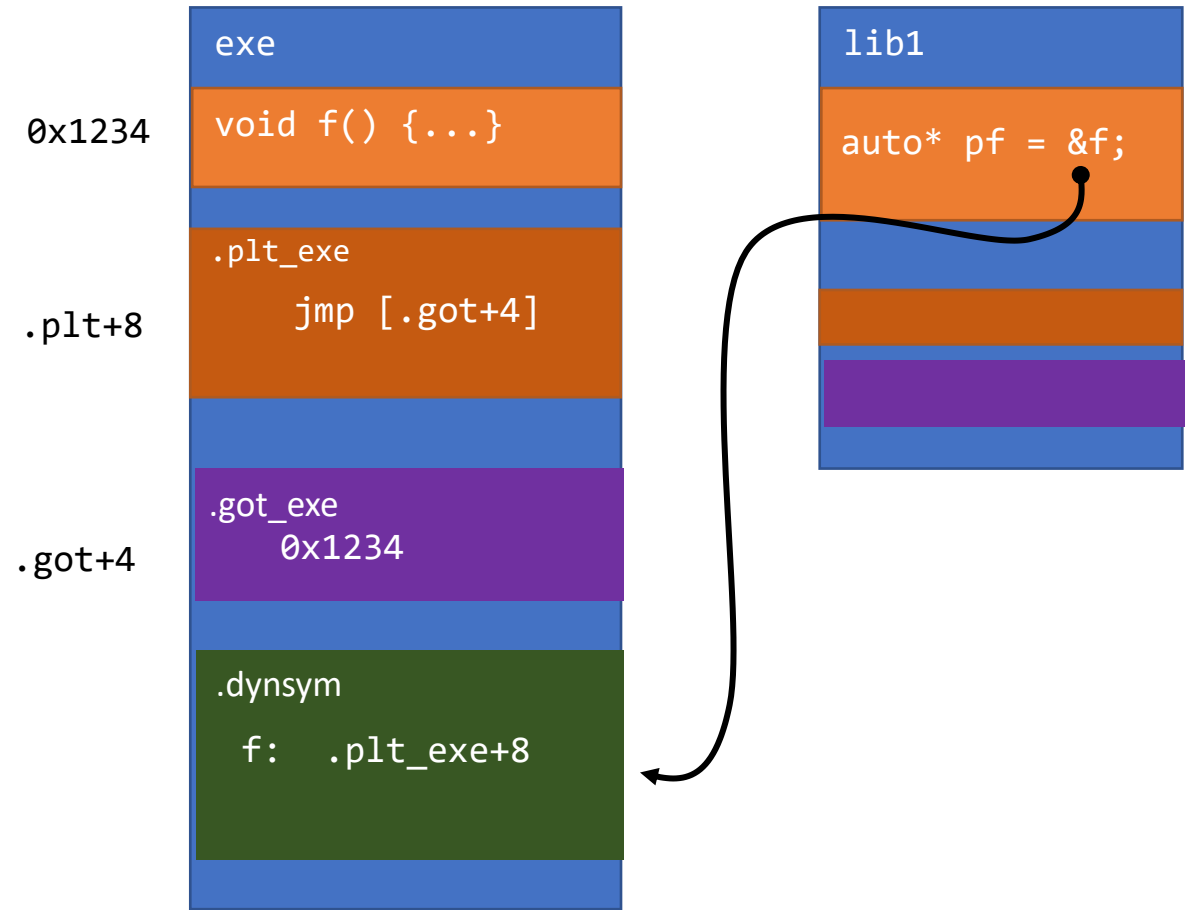
- C++ standard, [expr.eq]§3.2: “... if the pointers are both null, **both point to the same function**, or both represent the same address (6.8.2), they compare equal.”
- Actual calls are made to a PLT entry.
- Different among libraries!



C++ Implication #4: Comparing Func Ptrs

From the SystemV ABI doc:

To allow comparisons of function addresses to work as expected, if an executable file references a function defined in a shared object, the link editor will place the address of the procedure linkage table entry for that function in its associated symbol table entry. This will result in symbol table entries with section index of SHN_UNDEF but a type of STT_FUNC and a non-zero st_value. A reference to the address of a function from within a shared library will be satisfied by such a definition in the executable



C++ Implication #4: Comparing Func Ptrs

- Not comprehensive:
 - Different for `-fpie`
 - Different for comparison between pointers both taken at shared libs
- Anyway, Linux tries hard. And improves even today:
 - https://gcc.gnu.org/bugzilla/show_bug.cgi?id=100593#c13
 - (fix from 2023)



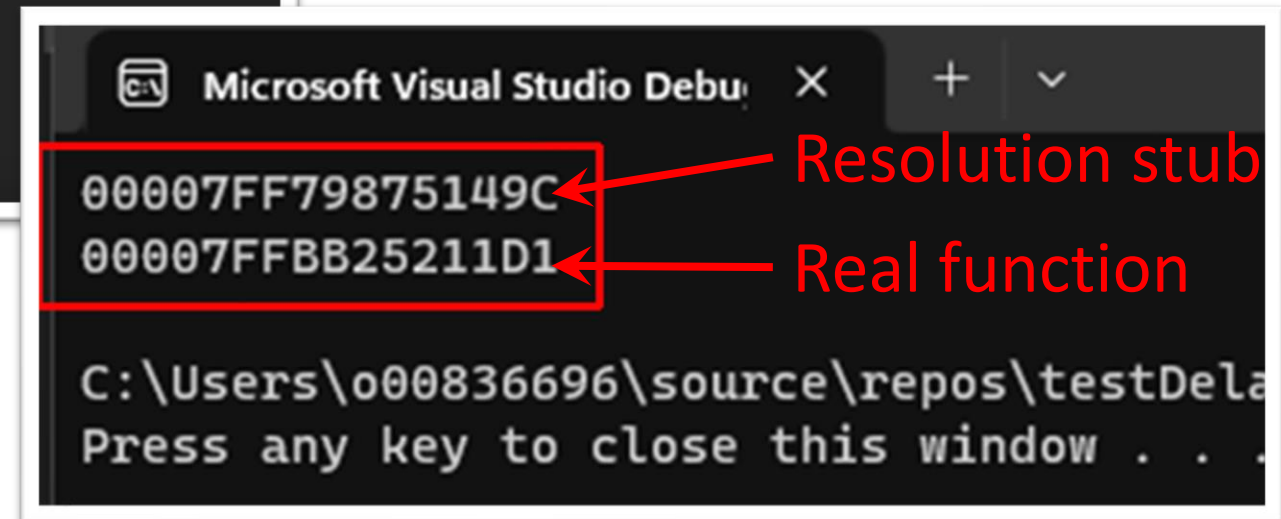
Windows doesn't even try

```
#include <iostream>
using std::cout;

__declspec(dllimport) void delayedLoadFunc();

int main()
{
    cout << &delayedLoadFunc << "\n";
    delayedLoadFunc();
    cout << &delayedLoadFunc << "\n";
}
```

<https://developercommunity.visualstudio.com/t/DELAYLOADs-implementation-breaks-the-C/10458529>



Intermediate Summary

- GOT is needed for Interposition
- PLT is needed for Lazy-Binding
- Both are questionable goals..!



Interposition, in retrospect

- Uneasy coexistence with C++ ODR,
- Rarely used,
- Takes a toll on *ALL* calls from shared libs, even those that are never interposed.
- Doubtful this would have been the design choice taken today.
 - <https://maskray.me/blog/2021-05-16-elf-interposition-and-bsymbolic>



Interposition, in retrospect

- “Outside of the GNU ELF world, many dynamic linking implementations have shifted to a direct binding and non-interposition by default.”
 - https://gcc.gnu.org/bugzilla/show_bug.cgi?id=100593#c13
- CPython got x1.3 faster by building with `-fno-semantic-interposition`
 - <https://github.com/python/cpython/issues/83161>



Daniel Colascione

April 3, 2021 · 🌐



Summary: Python is 1.3x faster when compiled in a way that re-examines shitty technical decisions from the 1990s.



Lazy Bind by Default?

- Lazy binding mandates having the GOT writable throughout the program execution
 - Glaring attack surface!
- As of Today:
 - `ld` still defaults to `-z lazy`
 - Many distros (Debian, Ubuntu, Fedora) configure compilers to build with `-z relro` : “relocate+read only”
 - Other languages (zig, rust) do too.
- Could consider `-no-plt` too
 - Not always possible: https://gcc.gnu.org/bugzilla/show_bug.cgi?id=100593#c13

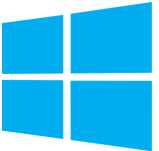
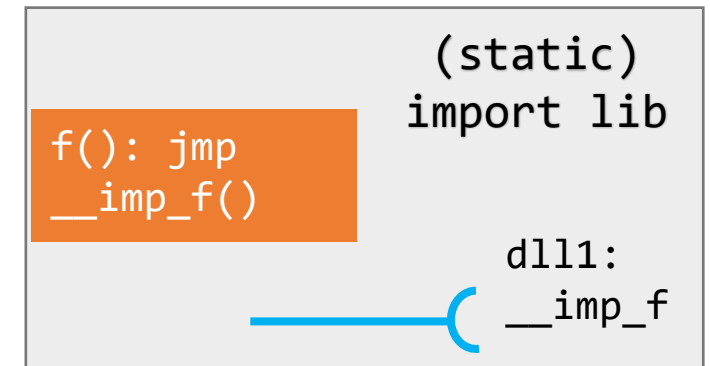
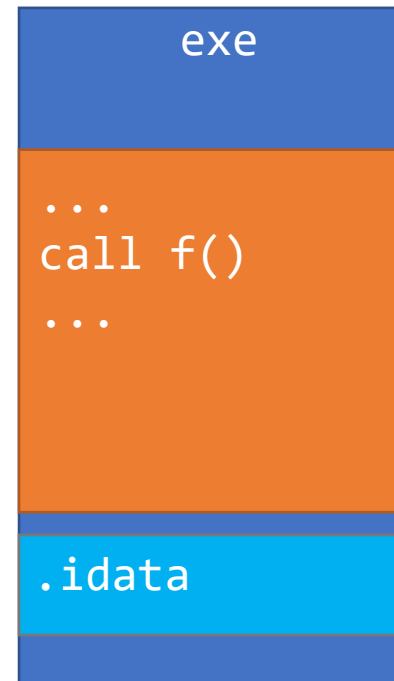


Yes. *

Symbol Visibility

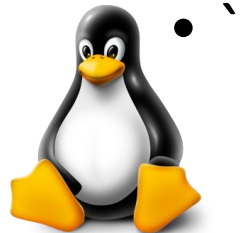
Symbol Visibility - Windows

- `__declspec(dllexport)` – add symbol to `.edata`
- `__declspec(dllimport)` – does *not* add to `.idata`!
 - That's achieved by statically linking against an *import library*
 - Minor optimization that happens in Release anyway
- Most symbols are neither.

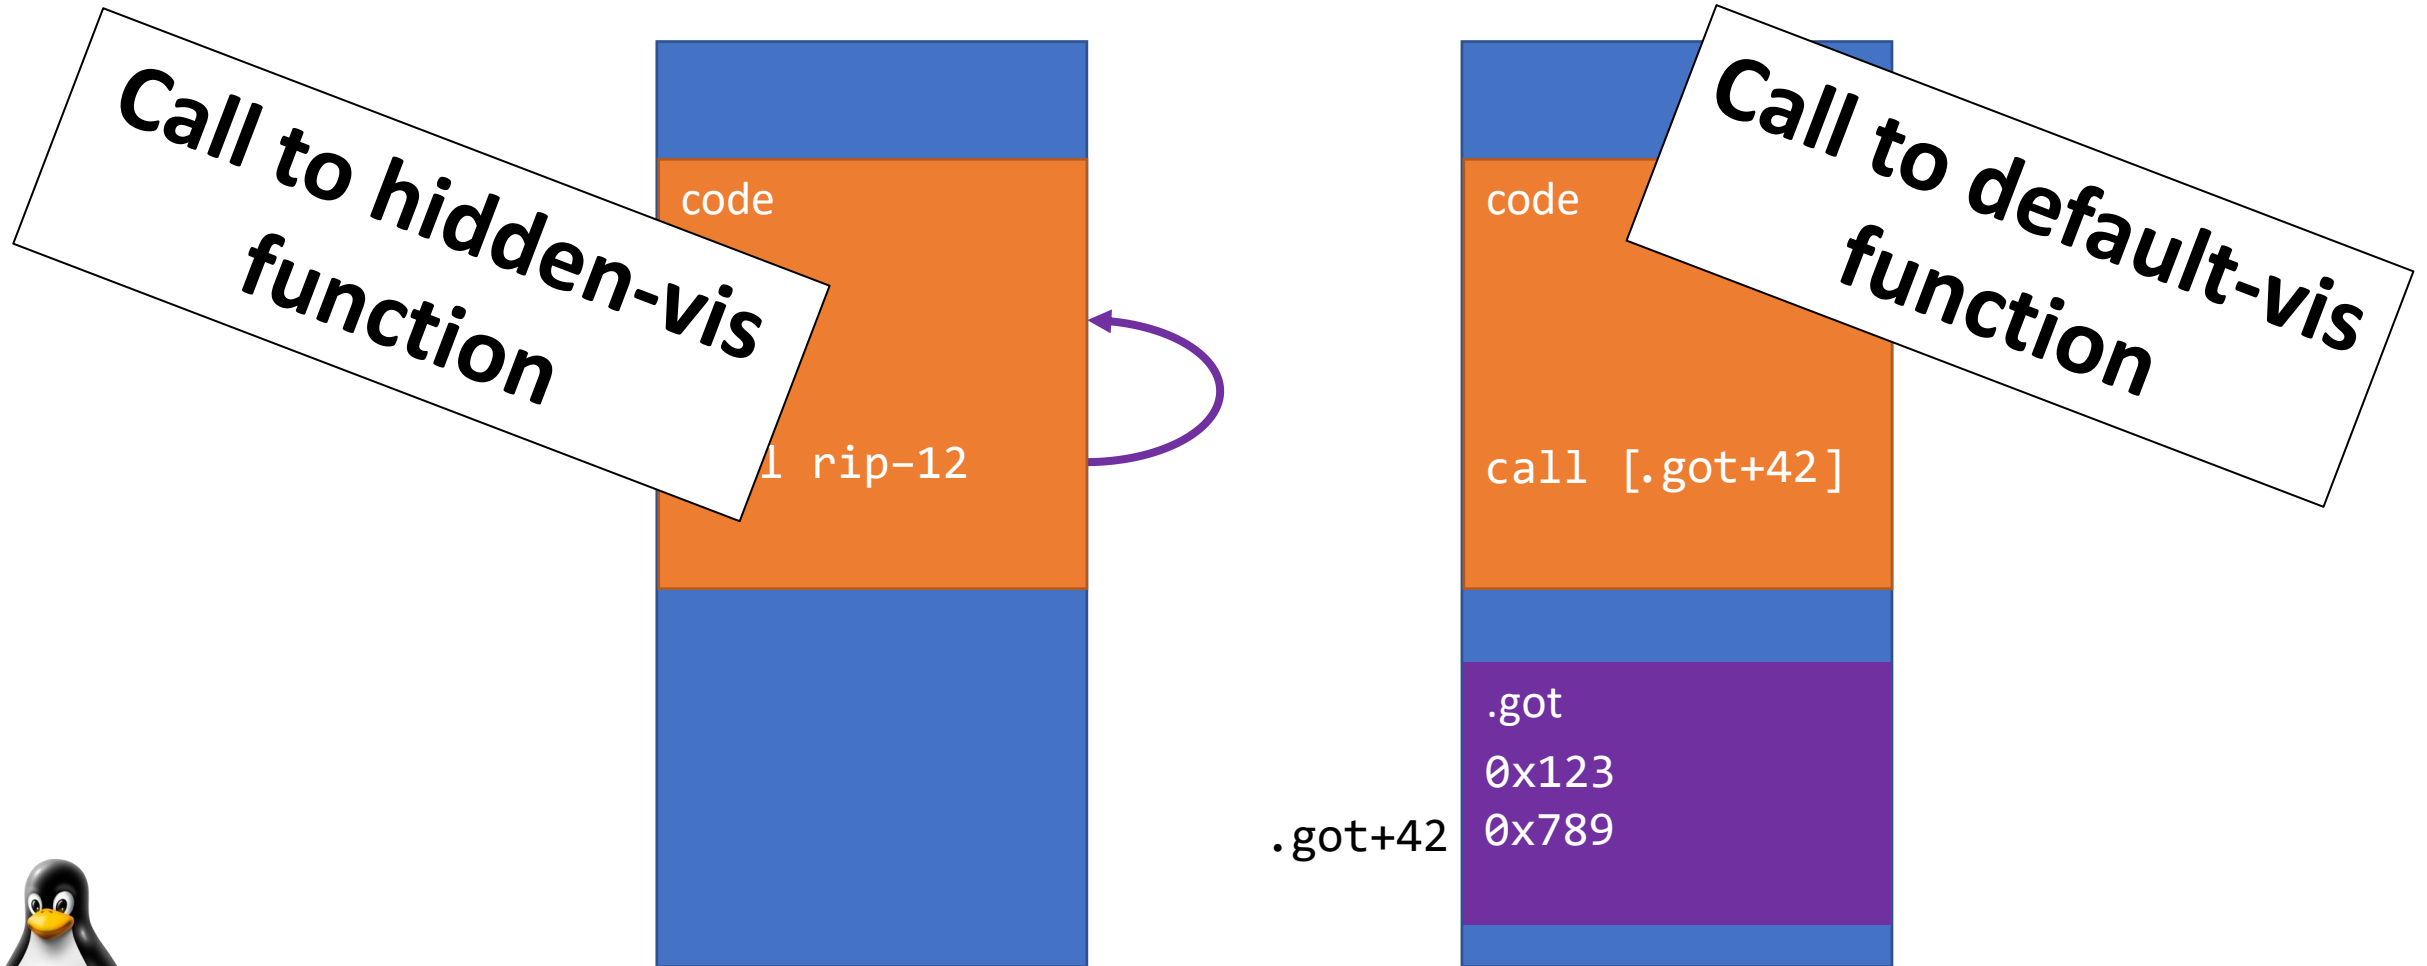


Symbol Visibility - Linux

- An ELF binary has a symbol table: section `.dynsym`
- Everything in the symbol table is accessible by other binaries
 - Potentially “exported”,
- Symbols can be marked UNDEFINED
 - Necessarily “*imported*” by the loader from other binaries.
- Not the end of the story – symbols have *visibility* :
 - Default / Protected / Hidden / ~~internal~~



Symbol Visibility – Who Cares?



Symbol Visibility - Linux

	Default	Protected	Hidden
Available to other binaries? (“exported”)	V	V	X
Subject to interposition? (potentially “imported”)	V	X	X
Has a GOT entry?	V	V	X
Appears in .dynsym?	V	V	X



Symbol Visibility - Linux

- “Using this feature [-fvisibility=hidden] can very substantially improve linking and load times of shared object libraries, produce more optimized code, provide near-perfect API export and prevent symbol clashes. It is strongly recommended that you use this in any shared objects you distribute.”

\$ man gcc

- Intervention hooks:
 - -fvisibility=hidden,
 - -fvisibility-inlines-hidden
 - -fvisibility-ms-compat
 - `__attribute__((visibility ("hidden")))`



Symbol Visibility vs Symbol Binding

- Binding (global/weak/local):
 - local binding →
 - static / anon-namespace symbols
 - weak binding →
 - Can be preempted even from a different file in the same binary. “Global” without ODR (eg, template instantiation).
 - Per file,
 - Used by the linker.
- Visibility (default/protected/hidden) :
 - Per binary,
 - Used by the Loader.



Recommendations for Shared Libraries on Linux

- Opt out of interposition!
 - Build with `-fvisibility=hidden`, `-Bsymbolic`
 - Mark only exported functions as `__attribute__((visibility("protected")))`
- Easier fallbacks:
 - `-fvisibility-inlines-hidden`
 - `-fvisibility-ms-compat`
 - `-fno-semantic-interposition`
- **Basically, link more like windows.**



C++ and Shared Libs – Preemptive Comment

Shared libraries are out of scope for C++ ?

“... Similarly, x86-64 instruction encoding is outside the scope of the C++ Standard. Does that mean a program that uses x86-64 instructions is not covered by the Standard? Of course not. It just means that that’s a low-level concern (relative to the semantics of C++), and it's up to the platform to make a conforming C++ implementation on top of that stuff. “

Arthur O’Dwyer, on the iso mailing list

<https://lists.isocpp.org/std-proposals/2023/07/7239.php>



Takeaways

Shared Libs - Takeaways

- Linux and Windows designs are fundamentally different.
- The Linux design is burdened by archaic design goals.
- Practical recommendations: `-fvisibility=hidden`, `-Bsymbolic`
- “Shared Libs are outside the scope of the standard” – is not a valid bailout!
 - Some fixes to the C++ standard are due.
- Linkers are messy, creative and wonderful.
And a *highly* active dev front.



Ofek Shilon

Senior Developer
@Toga Networks
(a Huawei Company)

ofekshilon@gmail.com



OfekShilon

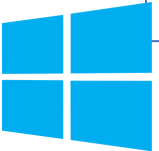
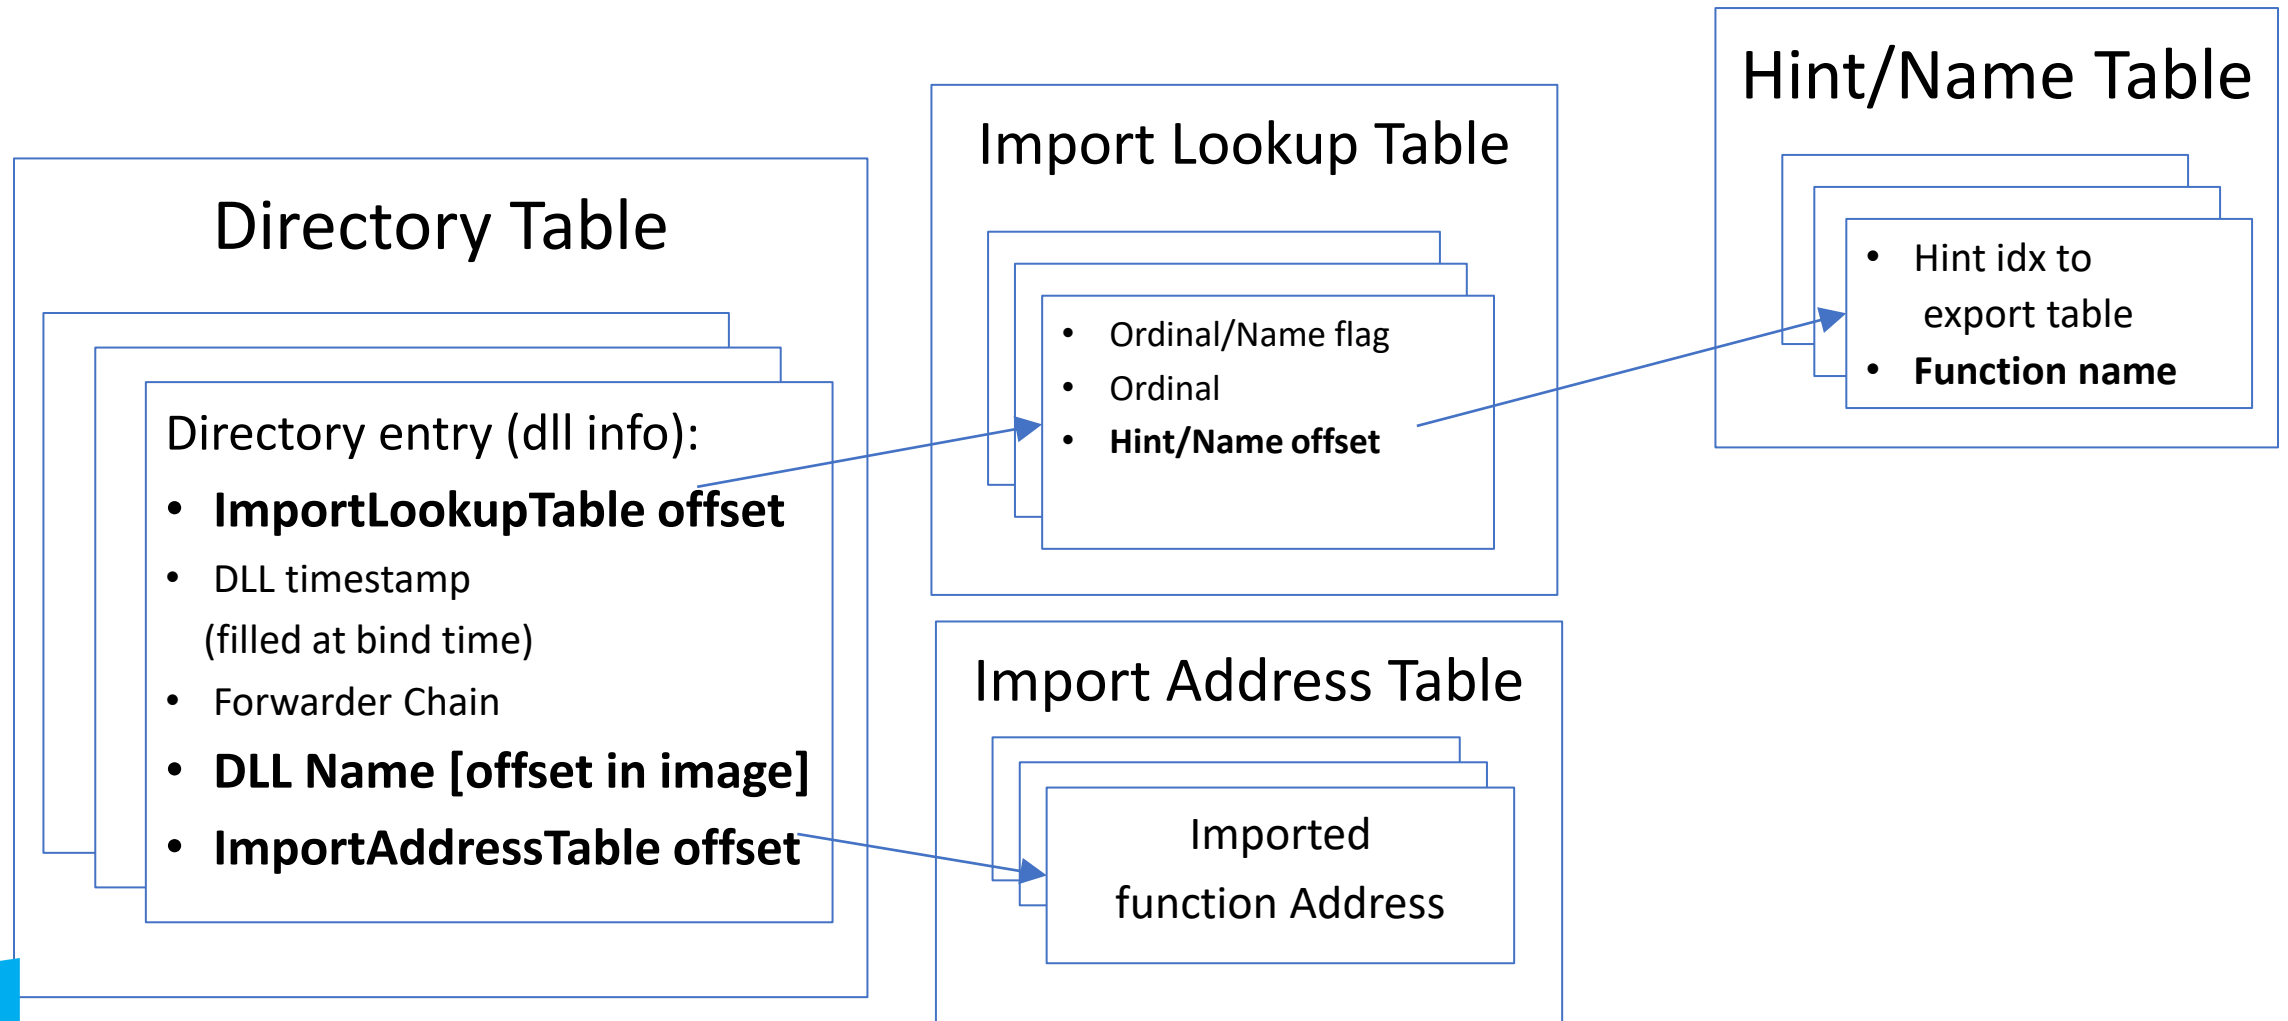
More Rabbit Holes for the Curious

- Relocation details
- Weak linkage
- Versioning (both of libraries and of API)
- Granularity
 - COMDAT, -ffunction-sections, -fdata-sections
- Optimizations
 - Identical Code Folding: /OPT:ICF, -icf=all, -icf=safe
 - Dead Code Elimination: /OPT:REF, -fvtable-gc, --gc-section
- Linker scripts

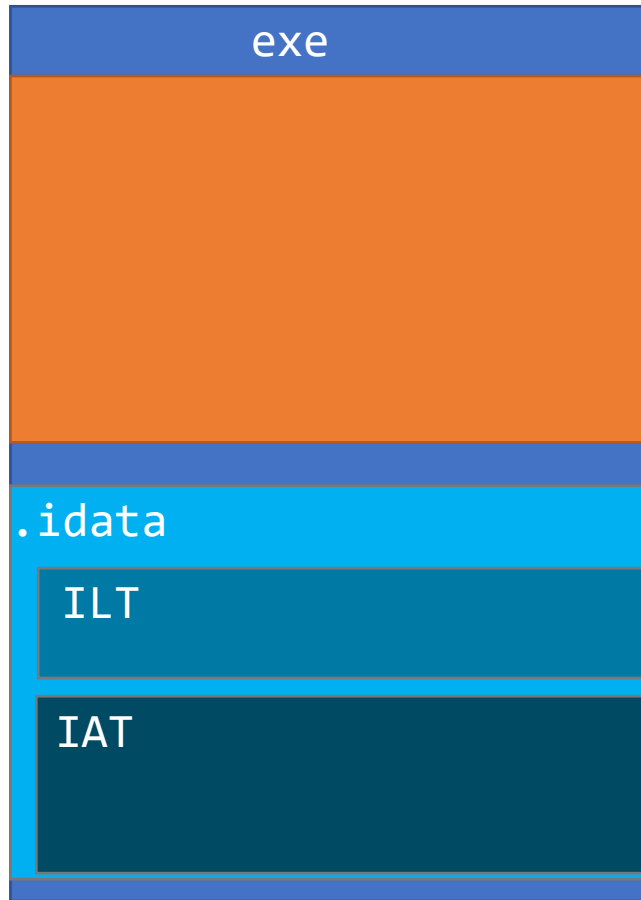
Resources

- Ulrich Drepper: “How to write shared libraries”
 - <http://library.bagrintsev.me/CPP/dsohowto.pdf>
- Eli Benderski:
 - <https://eli.thegreenplace.net/2011/08/25/load-time-relocation-of-shared-libraries/>
 - <https://eli.thegreenplace.net/2011/11/03/position-independent-code-pic-in-shared-libraries/>
- Fangrui Song, LLVM/LLD contributor:
 - <https://maskray.me/blog/>
- SystemV ABI spec and mailing list
 - <https://groups.google.com/g/x86-64-abi>
- Solaris documentation
 - https://docs.oracle.com/cd/E23824_01/html/819-0690/toc.html
- Ian Lance Taylor, author of GOLD:
 - <https://lwn.net/Articles/276782/>
- John Levine, “Linkers and Loaders” book:
 - <https://www.amazon.com/Linkers-Kaufmann-Software-Engineering-Programming/dp/1558604960>
- Michael Kerrisk online workshop:
 - <https://www.man7.org/training/shlib/index.html>

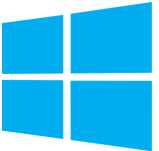
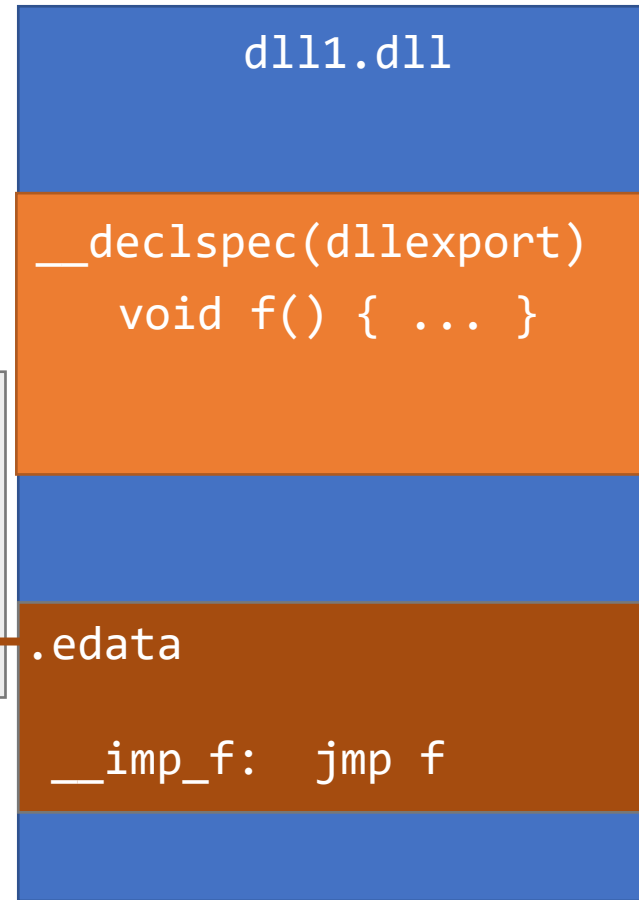
.idata section layout



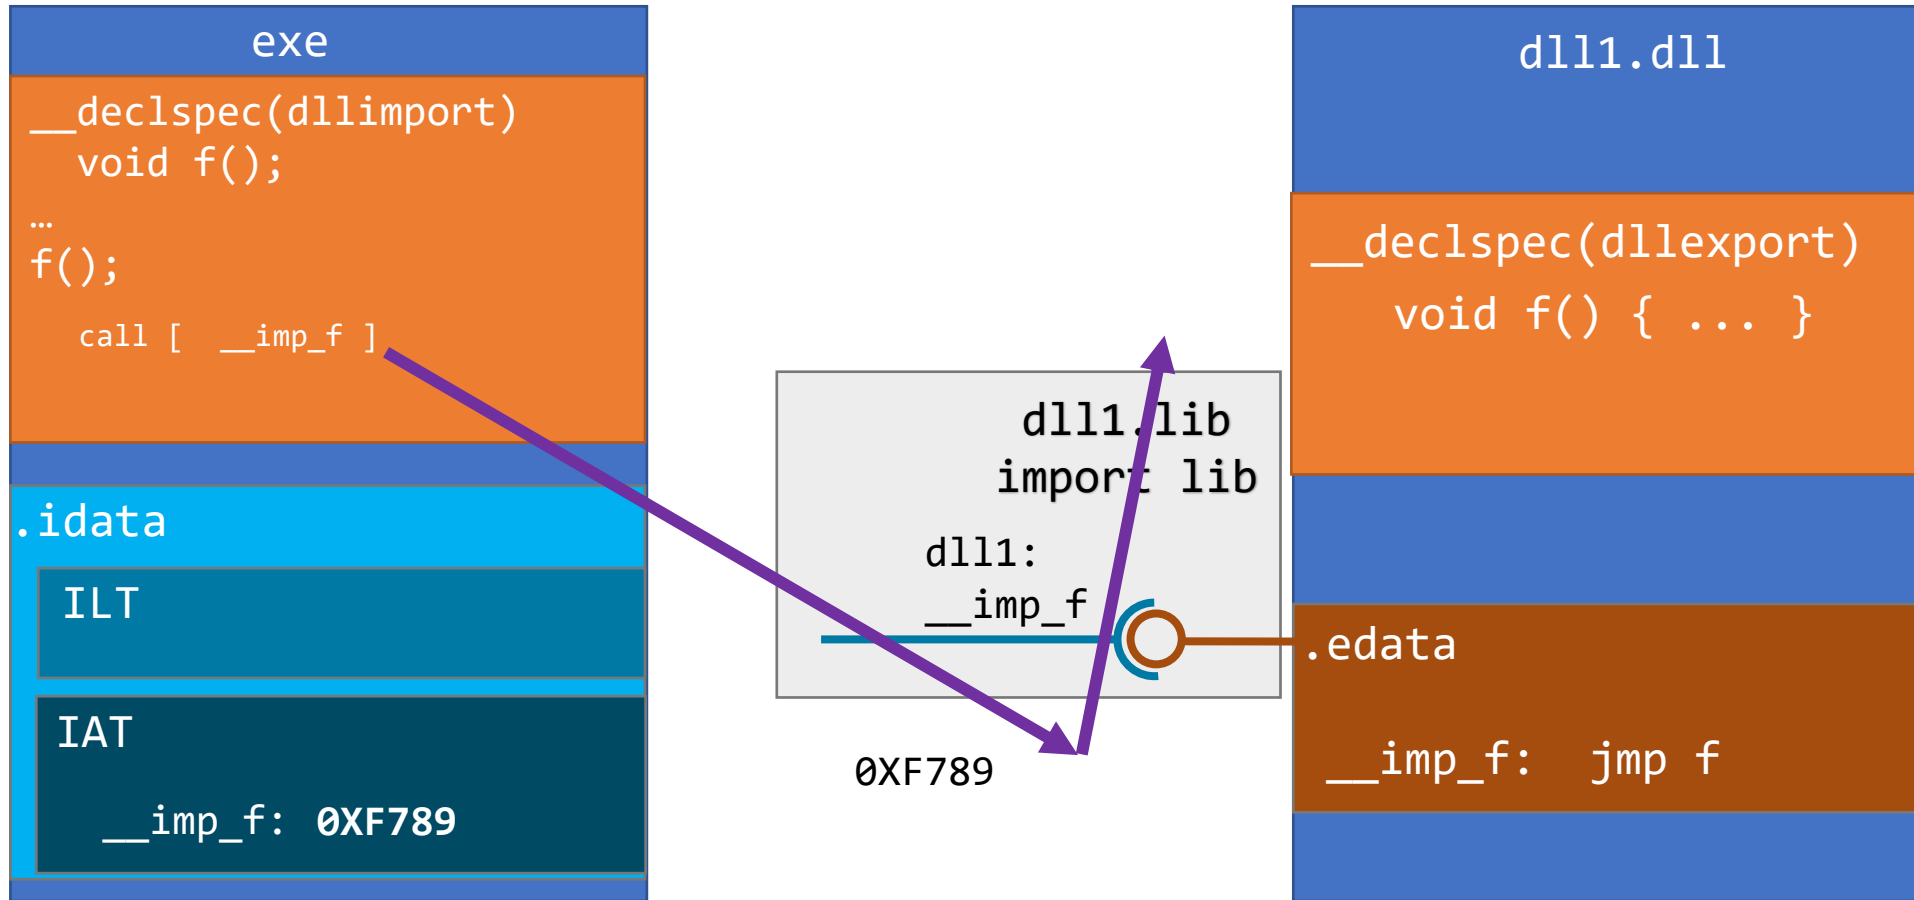
Import Library



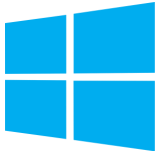
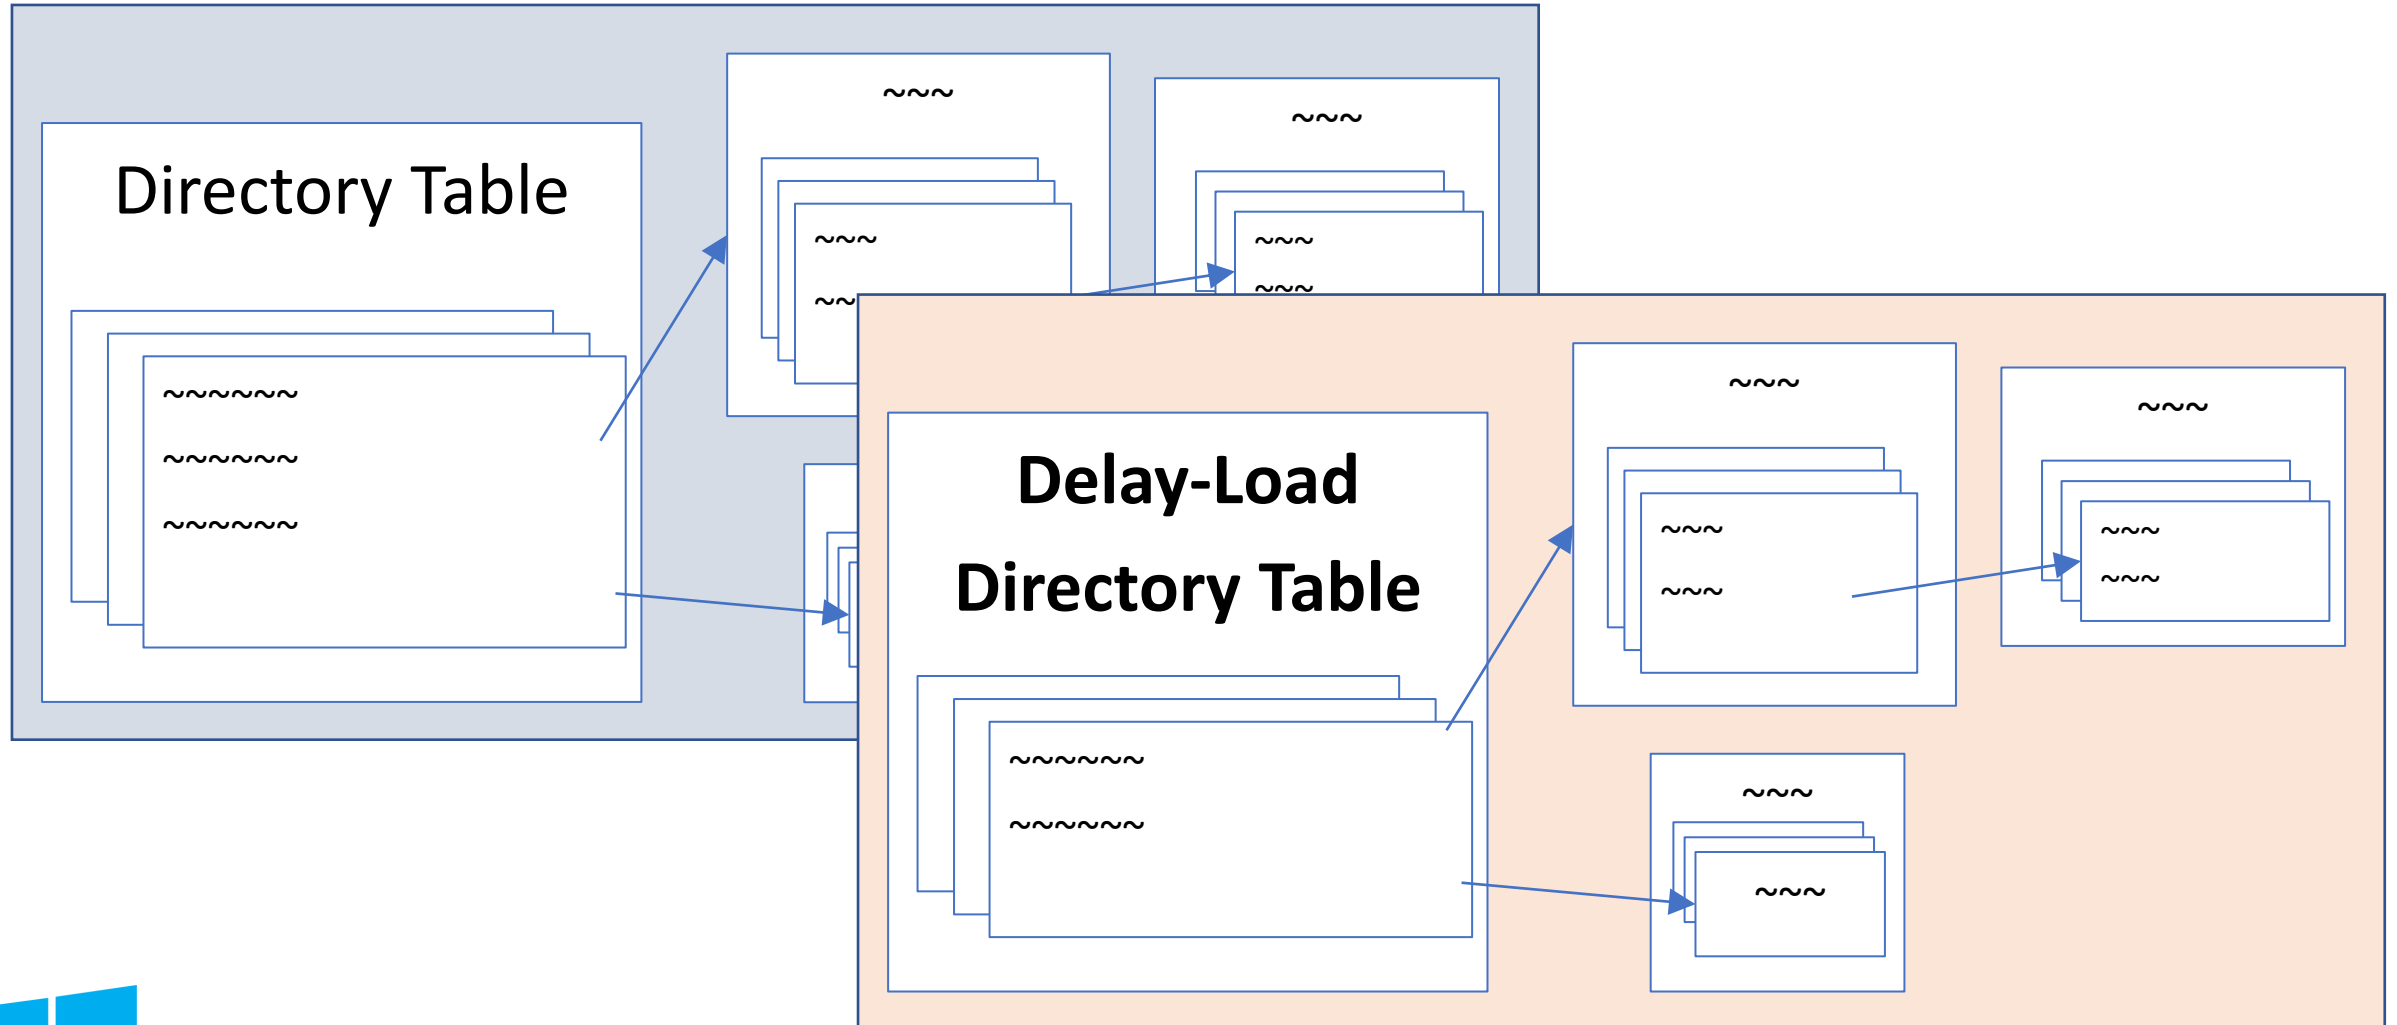
dll1.lib
import lib
dll1:
__imp_f



Import Library

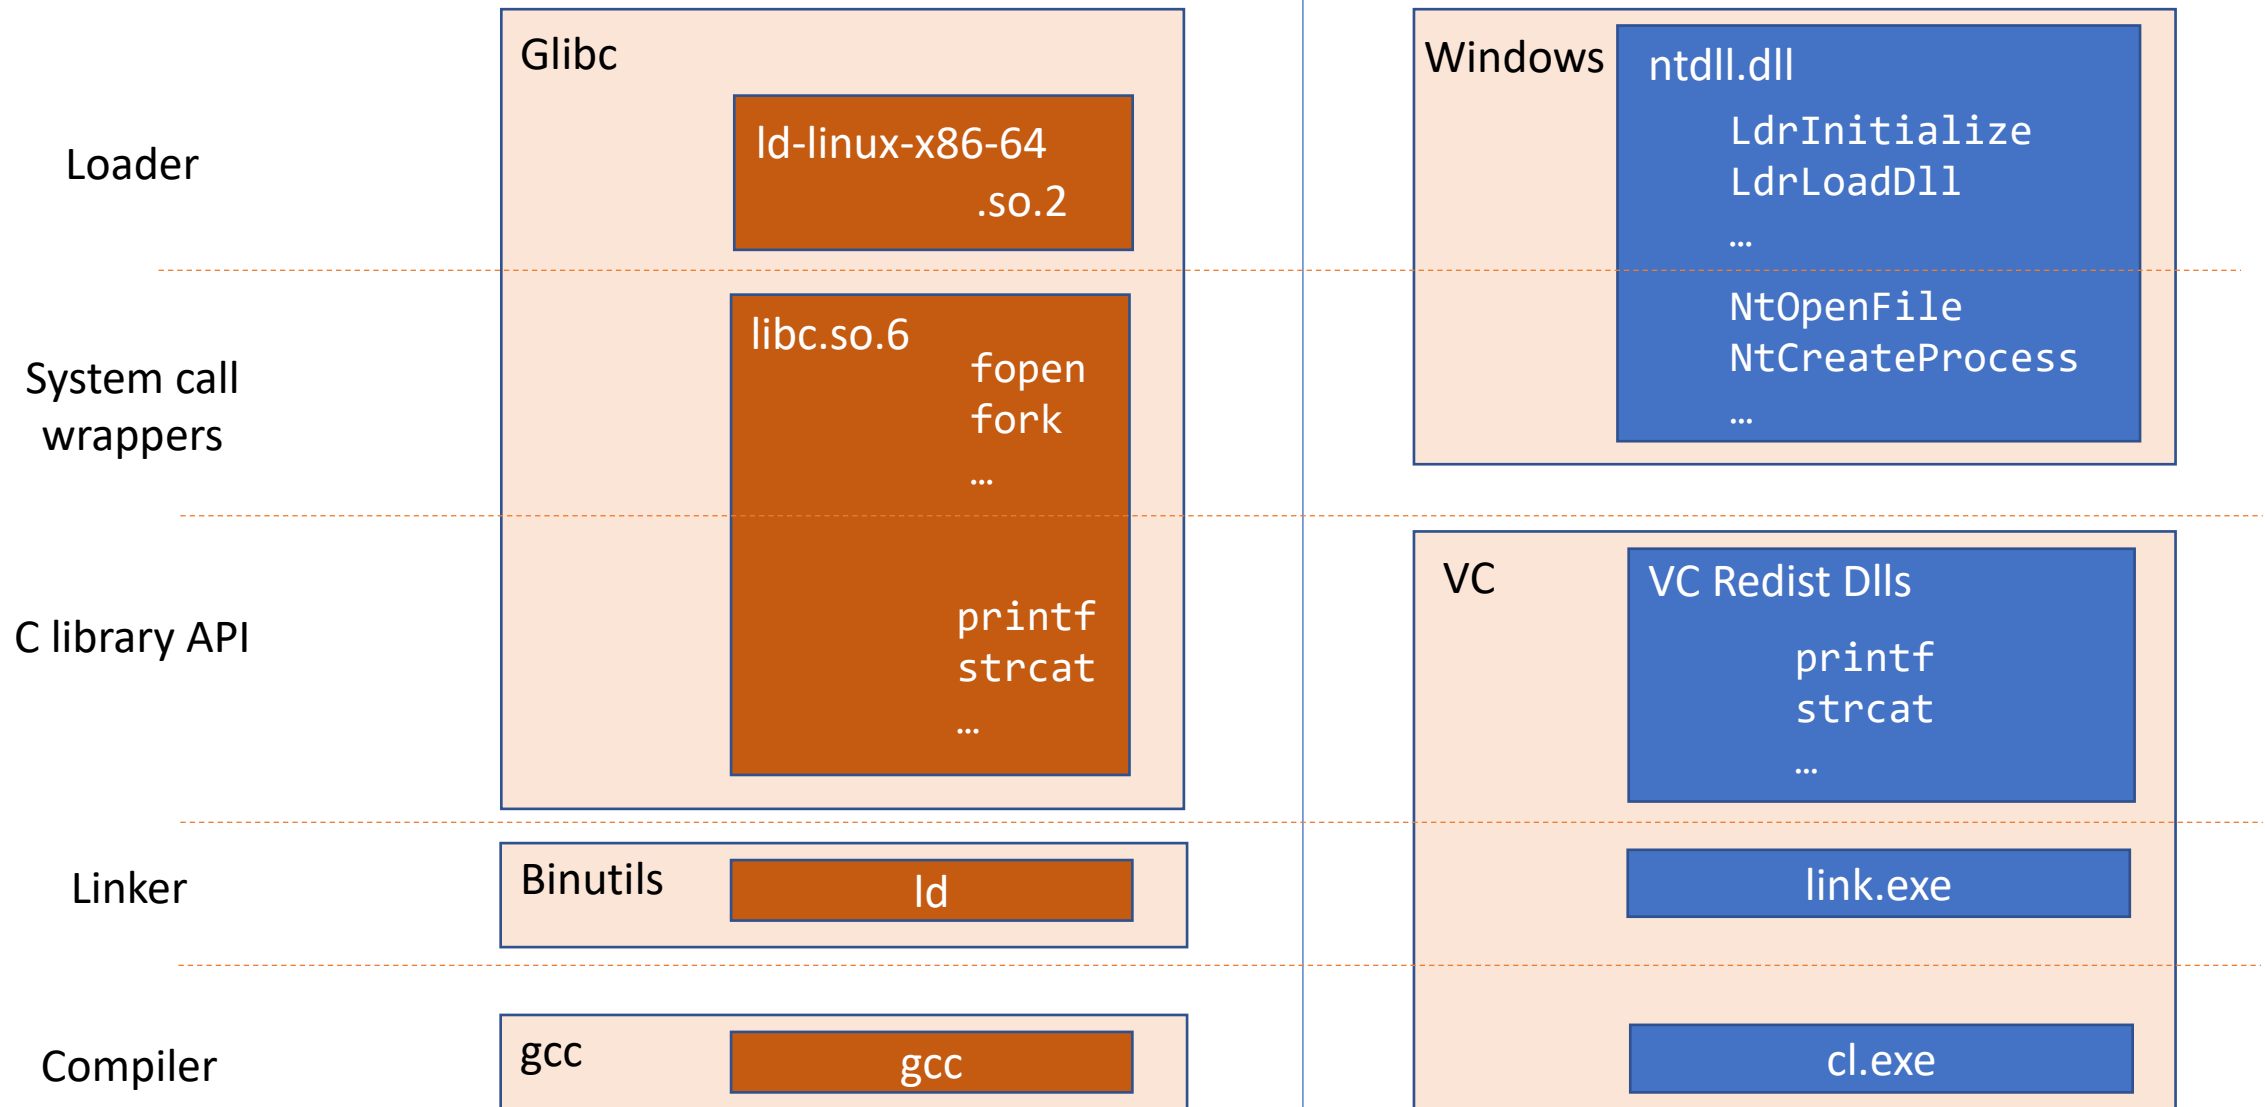


Windows .idata section - appendix



Loader

- A.k.a Dynamic Linker, a.k.a Interpreter, a.k.a Image Loader (esp in windows)
- Runs in user mode – operates on regular process address space
- In Linux: not invoked for executables linked with -static
- Not even with -static-pie!! <https://gcc.gnu.org/legacy-ml/gcc/2015-06/msg00008.html>
 - "... motivation for doing the relocations in the start file, rather than with an external program interpreter, is both to reduce runtime cost on very small systems, and to make deployment easier."
- So no LD_PRELOAD there



Selecting a different loader

```
$ gcc -v whatever.cpp
...
COLLECT_GCC_OPTIONS='-v' '-mtune=generic' '-march=x86-64'
 /usr/lib/gcc/x86_64-linux-gnu/9/collect2 -plugin /usr/lib/gcc/x86_64-linux-gnu/9/liblto_plugin.so -
plugin-opt=/usr/lib/gcc/x86_64-linux-gnu/9/lto-wrapper -plugin-opt=-fresolution=/tmp/ccfheLJQ.res -
plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lgcc_s -plugin-opt=-pass-through=-lc -
plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lgcc_s --build-id --eh-frame-hdr -m
elf_x86_64 --hash-style=gnu --as-needed -dynamic-linker /lib64/ld-linux-x86-64.so.2 -pie -z now -z
relro /usr/lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu/Scrt1.o /usr/lib/gcc/x86_64-linux-
gnu/9/../../../../x86_64-linux-gnu/crti.o /usr/lib/gcc/x86_64-linux-gnu/9/crtbeginS.o -
L/usr/lib/gcc/x86_64-linux-gnu/9 -L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu -
L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib -L/lib/x86_64-linux-gnu -L/lib/../../lib -
L/usr/lib/x86_64-linux-gnu -L/usr/lib/../../lib -L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../
/tmp/ccijfiIQ.o -lgcc --push-state --as-needed -lgcc_s --pop-state -lc -lgcc --push-state --as-
needed -lgcc_s --pop-state /usr/lib/gcc/x86_64-linux-gnu/9/crtendS.o /usr/lib/gcc/x86_64-linux-
gnu/9/../../../../x86_64-linux-gnu/crtn.o
...
```



Selecting a different loader

```
$ readelf --program-headers /usr/bin/ls
```

...

Program Headers:

Type	Offset FileSiz	VirtAddr MemSiz	PhysAddr Flags	Align
PHDR	0x0000000000000040	0x0000000000000040	0x0000000000000040	
	0x00000000000002d8	0x00000000000002d8	R	0x8
INTERP	0x0000000000000318	0x0000000000000318	0x0000000000000318	
	0x000000000000001c	0x000000000000001c	R	0x1

[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]

...



Observing the Loader in action - Linux

```
$ LD_DEBUG=help cat
```

Valid options for the LD_DEBUG environment variable are:

libs	display library search paths
reloc	display relocation processing
files	display progress for input file
symbols	display symbol table processing
bindings	display information about symbol binding
versions	display version dependencies
scopes	display scope information
all	all previous options combined
statistics	display relocation statistics
unused	determined unused DSOs
help	display this help message and exit



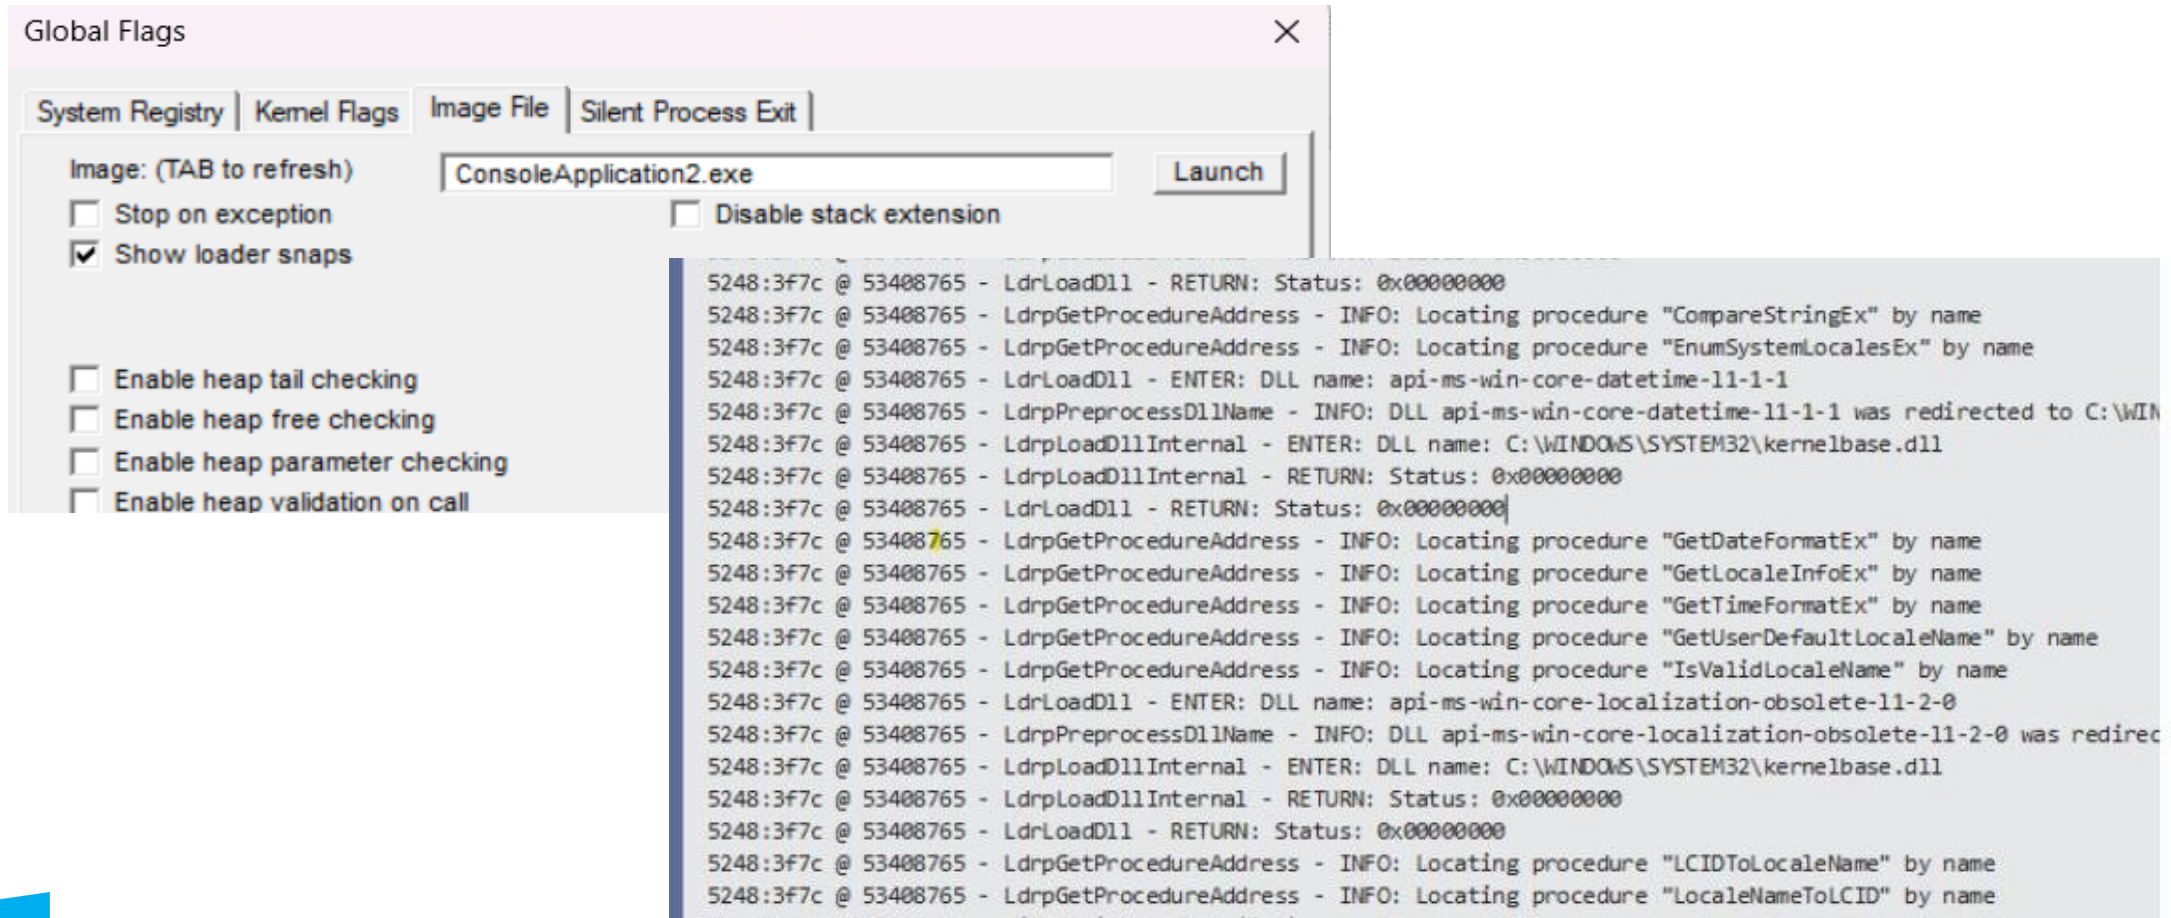
Observing the Loader in action - Linux

```
$ LD_AUDIT=1
```

Check!!



Observing the Loader in action - Windows



The image shows a Windows application window titled "Global Flags" with a close button (X) in the top right corner. The window has four tabs: "System Registry", "Kernel Flags", "Image File", and "Silent Process Exit". The "Image File" tab is selected. Below the tabs, there is a text box labeled "Image: (TAB to refresh)" containing the text "ConsoleApplication2.exe", and a "Launch" button to its right. Below this, there are two checkboxes: "Stop on exception" (unchecked) and "Disable stack extension" (unchecked). Further down, there is a checked checkbox labeled "Show loader snaps". At the bottom of the dialog, there are five unchecked checkboxes: "Enable heap tail checking", "Enable heap free checking", "Enable heap parameter checking", and "Enable heap validation on call".

Below the dialog box, a log window displays the following text:

```
5248:3f7c @ 53408765 - LdrLoadDll - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "CompareStringEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "EnumSystemLocalesEx" by name
5248:3f7c @ 53408765 - LdrLoadDll - ENTER: DLL name: api-ms-win-core-datetime-l1-1-1
5248:3f7c @ 53408765 - LdrpPreprocessDllName - INFO: DLL api-ms-win-core-datetime-l1-1-1 was redirected to C:\WIN
5248:3f7c @ 53408765 - LdrpLoadDllInternal - ENTER: DLL name: C:\WINDOWS\SYSTEM32\kernelbase.dll
5248:3f7c @ 53408765 - LdrpLoadDllInternal - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrLoadDll - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetDateFormatEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetLocaleInfoEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetTimeFormatEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetUserDefaultLocaleName" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "IsValidLocaleName" by name
5248:3f7c @ 53408765 - LdrLoadDll - ENTER: DLL name: api-ms-win-core-localization-obsolete-l1-2-0
5248:3f7c @ 53408765 - LdrpPreprocessDllName - INFO: DLL api-ms-win-core-localization-obsolete-l1-2-0 was redirec
5248:3f7c @ 53408765 - LdrpLoadDllInternal - ENTER: DLL name: C:\WINDOWS\SYSTEM32\kernelbase.dll
5248:3f7c @ 53408765 - LdrpLoadDllInternal - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrLoadDll - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "LCIDToLocaleName" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "LocaleNameToLCID" by name
```



Shared lib vs Static Lib

- Both are mechanisms for code sharing
- Static:
 1. Smaller single binary footprint (only used sections are imported)
 2. Better code locality
 3. Efficient calls
- Dynamic:
 1. Shared binary copies across processes (Linus disagrees)
 2. Save overall disk space (Linus disagrees)
 3. Separate lib versioning
 - Linus Torvalds disagrees with 1-3:
https://lore.kernel.org/lkml/CAHk-=whs8QZf3YnifdLv57+FhBi5_WeNTG1B-suOES=RcUSmQg@mail.gmail.com/
 4. Separate (smaller) linkage

Vague Linkage

- C++ constructs not clearly tied to a single translation unit:
 - Out-of-line copies of Inline Functions
 - String constants (shared across obj files)
 - VTables
 - type_info objects
 - Template Instantiations

Indirect Call Overhead

- Recent (2020) work on de-virtualization in clang:
 - “Modeling the Invariance of Virtual Pointers in LLVM” Piotr Padlewski, Krzysztof Pszeniczny, Richard Smith , <https://arxiv.org/pdf/2003.04228.pdf>
- “Our benchmarks show an average of 0.8% performance improvement on real-world C++ programs, with more than 30% speedup in some cases. ”

